

Librdkafka对kafka topic的封装



扫帚的影子
潜水三千只取一瓢

2 人赞同了该文章

- 上一节我们讲了librdkafka对topic-partition的封装, 任何一个partition都必须属于一下topic;
- 我们这节课来分析一下librdkafka对topic的封装

rd_kafka_itopic_s

- 所在文件: src/rdkafka_topic.h
- 这里还有一个类型 rd_kafka_topic_t, 定义: typedef struct rd_kafka_topic_s rd_kafka_topic_t; 这是个空定义没有实现, 其实就是 rd_kafka_itopic_s, 这个类型主要是面向librdkafka的使用者, sdk里称作 app topic, 它有自己的引用计数. 在librdkafka内部使用 rd_kafka_itopic, 它也有自己的引用计数, 有点罗嗦啊~
- 定义:

```
struct rd_kafka_itopic_s {
    // 定义成tailq的链表
    TAILQ_ENTRY(rd_kafka_itopic_s) rkt_link;

    // 引用计数
    rd_refcnt_t rkt_refcnt;

    rwlock_t rkt_lock;

    // 所属topic的名字
    rd_kafkap_str_t *rkt_topic;

    // 表示一个未assigned的partition
    shptr_rd_kafka_toppar_t *rkt_ua; /* unassigned partition */

    // 拥有的partition列表
    shptr_rd_kafka_toppar_t **rkt_p;
    int32_t rkt_partition_cnt;

    // 期望操作的partition, 但是没有从broker获取到该信息的partition列表
    rd_list_t rkt_desp; /* Desired partitions
                        * that are not yet seen
                        * in the cluster. */

    // 最近一次更新metadata的时间
    rd_ts_t rkt_ts_metadata; /* Timestamp of last metadata
                        * update for this topic. */

    mtx_t rkt_app_lock; /* Protects rkt_app_* */

    // 在application层一个rd_kafka_itopic_s对外表现为一个 rd_kafka_topic_t类型的对象,
    rd_kafka_topic_t *rkt_app_rkt; /* A shared topic pointer
                        * to be used for callbacks
                        * to the application. */

    int rkt_app_refcnt; /* Number of active rkt's new()ed
                        * by application. */

    // topic的三种状态:未知, 已存在, 不存在
    enum {
        RD_KAFKA_TOPIC_S_UNKNOWN, /* No cluster information yet */
        RD_KAFKA_TOPIC_S_EXISTS, /* Topic exists in cluster */
        RD_KAFKA_TOPIC_S_NOTEXISTS, /* Topic is not known in cluster */
    } rkt_state;

    int rkt_flags;
#define RD_KAFKA_TOPIC_F_LEADER_UNAVAIL 0x1 /* Leader lost/unavailable
                        * for at least one partition. */

    // 所属的rd_kafka_t
    rd_kafka_t *rkt_rk;

    shptr_rd_kafka_itopic_t *rkt_shptr_app; /* Application's topic_new() */

    rd_kafka_topic_conf_t rkt_conf;
};
```

- 创建一个 rd_kafka_itopic_s 对象 rd_kafka_topic_new0, 这是一个内部调用函数

```
shptr_rd_kafka_itopic_t *rd_kafka_topic_new0 (rd_kafka_t *rk,
                                              const char *topic,
                                              rd_kafka_topic_conf_t *conf,
                                              int *existingp,
                                              int do_lock) {
    rd_kafka_itopic_t *rkt;
    shptr_rd_kafka_itopic_t *s_rkt;
    const struct rd_kafka_metadata_cache_entry *rknce;

    // topic名字check, 长度不能超过512
    if (!topic || strlen(topic) > 512) {
        if (conf)
            rd_kafka_topic_conf_destroy(conf);
        rd_kafka_set_last_error(RD_KAFKA_RESP_ERR_INVALID_ARG,
                                EINVAL);
        return NULL;
    }

    if (do_lock)
        rd_kafka_wlock(rk);
    // 所有创建的rd_kafka_itopic对象都会加入到对应的topic的rk->rk_topics中, 先从中查找,
    if ((s_rkt = rd_kafka_topic_find(rk, topic, 0/*no lock*/))) {
        if (do_lock)
            rd_kafka_wunlock(rk);
        if (conf)
            rd_kafka_topic_conf_destroy(conf);
        if (existingp)
            *existingp = 1;
        return s_rkt;
    }

    if (existingp)
        *existingp = 0;

    // 分配对应的内存, 设置各属性
    rkt = rd_calloc(1, sizeof(*rkt));

    rkt->rkt_topic = rd_kafkap_str_new(topic, -1);
    rkt->rkt_rk = rk;

    if (!conf) {
        if (rk->rk_conf.topic_conf)
            conf = rd_kafka_topic_conf_dup(rk->rk_conf.topic_conf);
        else
            conf = rd_kafka_topic_conf_new();
    }
    rkt->rkt_conf = *conf;
    rd_free(conf); /* explicitly not rd_kafka_topic_destroy()
                    * since we dont want to rd_free internal members,
                    * just the placeholder. The internal members
                    * were copied on the line above. */

    /* Default partitioner: consistent_random */
    if (!rk->rk_conf.partitioner)
        rkt->rk_conf.partitioner = rd_kafka_msg_partitioner_consistent_random;

    if (rkt->rk_conf.compression_codec == RD_KAFKA_COMPRESSION_INHERIT)
        rkt->rk_conf.compression_codec = rk->rk_conf.compression_codec;

    rd_list_init(&rkt->rkt_desp, 16, NULL);
    rd_refcnt_init(&rkt->rkt_refcnt, 0);

    s_rkt = rd_kafka_topic_peek(rkt);

    rwlock_init(&rkt->rkt_lock);
    mtx_init(&rkt->rkt_app_lock, mtx_plain);

    /* Create unassigned partition */
    rkt->rkt_ua = rd_kafka_toppar_new(rkt, RD_KAFKA_PARTITION_UA);

    // 加入到对应的rk_kafka_t中的topic列表
    TAILQ_INSERT_TAIL(&rk->rk_topics, rkt, rkt_link);
    rk->rk_topic_cnt++;

    /* Populate from metadata cache. */
    // 加入或更新到metadata cache
    if ((rknce = rd_kafka_metadata_cache_find(rk, topic, 1/*valid*/))) {
        if (existingp)
            *existingp = 1;
```

```

        rd_kafka_topic_metadata_update(rkt, &rknce->rknce_mtopic,
                                         rknce->rknce_insert);
    }

    if (do_lock)
        rd_kafka_wrunlock(rk);

    return s_rkt;
}

```

- 创建 rd_kafka_topic_t 对象, 对外的接口 rd_kafka_topic_new

```

rd_kafka_topic_t *rd_kafka_topic_new(rd_kafka_t *rk, const char *topic,
                                       rd_kafka_topic_conf_t *conf) {
    shptr_rd_kafka_itopic_t *s_rkt;
    rd_kafka_itopic_t *rkt;
    rd_kafka_topic_t *app_rkt;
    int existing;

    // 创建一个 shptr_rd_kafka_itopic_t 对象
    s_rkt = rd_kafka_topic_new0(rk, topic, conf, &existing, 1/*lock*/);
    if (!s_rkt)
        return NULL;

    // 指针转换, 从 shptr_rd_kafka_itopic_t 到 rd_kafka_itopic_t, 引用计数不变
    rkt = rd_kafka_topic_s2i(s_rkt);

    /* Save a shared pointer to be used in callbacks. */
    // 引用计数加1, 指针转换成一个 rd_kafka_topic_t
    // app相对应的引用计数也加1
    app_rkt = rd_kafka_topic_keep_app(rkt);

    /* Query for the topic leader (async) */
    if (!existing)
        // 发metadata request, 获取leader等相关信息
        rd_kafka_topic_leader_query(rk, rkt);

    /* Drop our reference since there is already now a rkt_app_rkt */
    rd_kafka_topic_destroy0(s_rkt);

    return app_rkt;
}

```

- 获取当前 rd_kafka_t 对象持有的所有topic的名字, 保存在一个 rd_list 中

```

void rd_kafka_local_topics_to_list(rd_kafka_t *rk, rd_list_t *topics) {
    rd_kafka_itopic_t *rkt;

    rd_kafka_rdlck(rk);
    rd_list_grow(topics, rk->rk_topic_cnt);
    TAILQ_FOREACH(rkt, &rk->rk_topics, rkt_link)
        rd_list_add(topics, rd_strdup(rkt->rkt_topic->str));
    rd_kafka_rdlunlck(rk);
}

```

- 判断partition是否是有效的, 就是判断其leader是否有效

```

int rd_kafka_topic_partition_available(const rd_kafka_topic_t *app_rkt,
                                       int32_t partition) {
    int avail;
    shptr_rd_kafka_toppar_t *s_rkt;
    rd_kafka_toppar_t *rkt;
    rd_kafka_broker_t *rkb;

    s_rkt = rd_kafka_toppar_get(rd_kafka_topic_s2i(app_rkt),
                                 partition, 0/*no ua-on-miss*/);
    if (unlikely(!s_rkt))
        return 0;

    rkt = rd_kafka_toppar_s2i(s_rkt);
    rkb = rd_kafka_toppar_leader(rkt, 1/*proper broker*/);
    avail = rkb ? 1 : 0;
    if (rkb)
        rd_kafka_broker_destroy(rkb);
    rd_kafka_toppar_destroy(s_rkt);
    return avail;
}

```

- 扫描所有topic的pathions:
- 筛出 kafka message过期的, 回调application层
- 找出需要刷新metadata的, 发送metadata request

```

int rd_kafka_topic_scan_all(rd_kafka_t *rk, rd_ts_t now) {
    rd_kafka_itopic_t *rkt;
    rd_kafka_toppar_t *rkt;
    shptr_rd_kafka_toppar_t *s_rkt;
    int totcnt = 0;
    rd_list_t query_topics;

    rd_list_init(&query_topics, 0, rd_free);

    rd_kafka_rdlck(rk);
    TAILQ_FOREACH(rkt, &rk->rk_topics, rkt_link) {
        int p;

        int cnt = 0, tpcent = 0;
        rd_kafka_msgq_t timedout;
        int query_this = 0;

        rd_kafka_msgq_init(&timedout);

        rd_kafka_topic_wlck(rkt);

        /* Check if metadata information has timed out. */
        // metadata cache中没有缓存, 需要query metadata
        if (rkt->rkt_state != RD_KAFKA_TOPIC_S_UNKNOWN &&
            !rd_kafka_metadata_cache_topic_get(
                rk, rkt->rkt_topic->str, 1/*only valid*/)) {
            rd_kafka_topic_set_state(rkt, RD_KAFKA_TOPIC_S_UNKNOWN);

            query_this = 1;
        }

        /* Just need a read-lock from here on. */
        rd_kafka_topic_wrunlck(rkt);
        rd_kafka_topic_rdlck(rkt);

        if (rkt->rkt_partition_cnt == 0) {
            query_this = 1;
        }

        for (p = RD_KAFKA_PARTITION_UA ;
             p < rkt->rkt_partition_cnt ; p++) {
            int did_tmout = 0;

            if (!s_rkt = rd_kafka_toppar_get(rkt, p, 0))
                continue;

            rkt = rd_kafka_toppar_s2i(s_rkt);
            rd_kafka_toppar_lock(rkt);

            /* Check that partition has a leader that is up,
             * else add topic to query list. */
            // partition leader无效时, 要request metadata
            if (p != RD_KAFKA_PARTITION_UA &&
                (!rkt->rkt_leader ||
                 rkt->rkt_leader->rkb_source ==
                 RD_KAFKA_INTERNAL ||
                 rd_kafka_broker_get_state(rkt->rkt_leader) <
                 RD_KAFKA_BROKER_STATE_UP)) {
                query_this = 1;
            }

            /* Scan toppar's message queues for timeouts */
            if (rd_kafka_msgq_age_scan(&rkt->rkt_xmit_msgq,
                                       &timedout, now) > 0)
                did_tmout = 1;

            if (rd_kafka_msgq_age_scan(&rkt->rkt_com_msgq,
                                       &timedout, now) > 0)
                did_tmout = 1;

            tpcent += did_tmout;

            rd_kafka_toppar_unlock(rkt);
            rd_kafka_toppar_destroy(s_rkt);
        }

        rd_kafka_topic_rdlunlck(rkt);
    }
}

```

```

        if ((cnt = rd_atomic32_get(&timedout.rkmq_msg_cnt)) > 0) {
            totcnt += cnt;

            // kafka message过期, 则需要回到application层
            rd_kafka_dr_msgq(rkt, &timedout,
                            RD_KAFKA_RESP_ERR_TIMED_OUT);
        }

        /* Need to re-query this topic's leader. */
        if (query_this &&
            !rd_list_find(&query_topics, rkt->rkt_topic->str,
                          (void *)strcmp))
            rd_list_add(&query_topics,
                        rd_strdup(rkt->rkt_topic->str));
    }

    rd_kafka_rdonlock(rk);

    if (!rd_list_empty(&query_topics))
        // 发送 metadata request
        rd_kafka_metadata_refresh_topics(rk, NULL, &query_topics,
                                          1/*force even if cached
                                           * info exists*/,
                                          "refresh unavailable topics");

    rd_list_destroy(&query_topics);

    return totcnt;
}

```

- 更新topic的partition个数, partition个数可能增加, 也可能减

少 rd_kafka_topic_partition_cnt_update, 简单讲:

- 新增的partition, 创建:

- 老的partition, 删除:

```

static int rd_kafka_topic_partition_cnt_update(rd_kafka_itopic_t *rkt,
                                              int32_t partition_cnt) {
    rd_kafka_t *rk = rkt->rkt_rk;
    shptr_rd_kafka_toppar_t **rktps;
    shptr_rd_kafka_toppar_t *rktp_uu;
    shptr_rd_kafka_toppar_t *s_rktp;
    rd_kafka_toppar_t *rktp;
    rd_kafka_msgq_t tmpq = RD_KAFKA_MSGQ_INITIALIZER(tmpq);
    int32_t i;

    更新前后partition数量相同的id, 不作任何处理
    if (likely(rkt->rkt_partition_cnt == partition_cnt))
        return 0; /* No change in partition count */

    /* Create and assign new partition list */
    // 创建新的partition list, 分配内存
    if (partition_cnt > 0)
        rktps = rd_calloc(partition_cnt, sizeof(*rktps));
    else
        rktps = NULL;

    // 如果新个数大于老个数
    for (i = 0; i < partition_cnt; i++) {
        // 多出来的都是新扩容的partition
        if (i >= rkt->rkt_partition_cnt) {
            /* New partition. Check if its in the list of
             * desired partitions first. */
            // 检查是否在desired partition 列表中
            s_rktp = rd_kafka_toppar_desired_get(rkt, i);

            rktp = s_rktp ? rd_kafka_toppar_s2i(s_rktp) : NULL;
            if (rktp) {
                // 在desired partition 列表中, 则移除它
                rd_kafka_toppar_lock(rktp);
                rktp->rktp_flags &= ~RD_KAFKA_TOPPAR_F_UNKNOWN;

                /* Remove from desp list since the
                 * partition is now known. */
                rd_kafka_toppar_desired_unlink(rktp);
                rd_kafka_toppar_unlock(rktp);
            } else
                s_rktp = rd_kafka_toppar_new(rkt, i);
            // 赋值rktps[i]
            rktps[i] = s_rktp;
        } else {
            // 如果是已存在的partition, 放到rktps[i], 并且作引用计数的增减
            /* Existing partition, grab our own reference. */
            rktps[i] = rd_kafka_toppar_keep(
                rd_kafka_toppar_s2i(rkt->rkt_p[i]));
            /* Loose previous ref */
            rd_kafka_toppar_destroy(rkt->rkt_p[i]);
        }
    }

    rktp_uu = rd_kafka_toppar_get(rkt, RD_KAFKA_PARTITION_UA, 0);

    /* Propagate notexist errors for desired partitions */
    // 扫描desired partition 列表中, 还余下的都是无主的, 集群中不存在的partition, 返回
    RD_LIST_FOREACH(s_rktp, &rkt->rkt_desp, i) {
        rd_kafka_toppar_enq_error(rd_kafka_toppar_s2i(s_rktp),
                                RD_KAFKA_RESP_ERR_UNKNOWN_PARTITION);
    }

    /* Remove excessive partitions */
    // 处理更新后的partition个数小于更新前的情况, 需要删除一部分partition
    for (i = partition_cnt; i < rkt->rkt_partition_cnt; i++) {
        s_rktp = rkt->rkt_p[i];
        rktp = rd_kafka_toppar_s2i(s_rktp);
        rd_kafka_toppar_lock(rktp);

        if (rktp->rktp_flags & RD_KAFKA_TOPPAR_F_DESIRED) {
            rd_kafka_dbg(rkt->rkt_rk, TOPIC, "DESIRED",
                          "Topic %s [%i/%i] is desired "
                          "but no longer known: "
                          "moving back on desired list",
                          rkt->rkt_topic->str, rktp->rktp_partition);
            // 是DESIRED状态的话, 再放回desired列表
            rktp->rktp_flags |= RD_KAFKA_TOPPAR_F_UNKNOWN;
            rd_kafka_toppar_desired_unlink(rktp);

            if (!rd_kafka_terminating(rkt->rkt_rk))
                rd_kafka_toppar_enq_error(
                    rktp,
                    RD_KAFKA_RESP_ERR_UNKNOWN_PARTITION);
            // 解除和broker的联系, 实际上是关联到内部的UA broker
            rd_kafka_toppar_broker_delegate(rktp, NULL, 0);
        } else {
            /* Tell handling broker to let go of the toppar */
            rktp->rktp_flags |= RD_KAFKA_TOPPAR_F_REMOVE;
            rd_kafka_toppar_broker_leave_for_remove(rktp);
        }

        rd_kafka_toppar_unlock(rktp);

        rd_kafka_toppar_destroy(s_rktp);
    }

    if (rkt->rkt_p)
        rd_free(rkt->rkt_p);

    rkt->rkt_p = rktps;
    rkt->rkt_partition_cnt = partition_cnt;

    return 1;
}

```

- 将在UA partition上特发送的kafka message重新分配到有效的partition

上 rd_kafka_topic_assign_uus:

```

static void rd_kafka_topic_assign_uus(rd_kafka_itopic_t *rkt,
                                       rd_kafka_resp_err_t err) {
    rd_kafka_t *rk = rkt->rkt_rk;
    shptr_rd_kafka_toppar_t *s_rktp_uu;
    rd_kafka_toppar_t *rktp_uu;
    rd_kafka_msg_t *rkma, *tmp;
    rd_kafka_msgq_t uus = RD_KAFKA_MSGQ_INITIALIZER(uus);
    rd_kafka_msgn_t failed = RD_KAFKA_MSGQ_INITIALIZER(failed);
    int cnt;

    if (rkt->rkt_rk->rkt_type != RD_KAFKA_PRODUCER)
        return;

    // 没有UA partition, 就直接返回了
    s_rktp_uu = rd_kafka_toppar_get(rkt, RD_KAFKA_PARTITION_UA, 0);
    if (unlikely(!s_rktp_uu)) {
        return;
    }

```

```
    }

    rktp_uu = rd_kafka_toppar_s2i(s_rktp_uu);

    // 将ua partition上的msg移动到临时队列上
    rd_kafka_toppar_lock(rktp_uu);
    rd_kafka_msgq_move(&uas, &rktp_uu->rktp_msgq);
    cnt = rd_atomic32_get(&uas.rkmq_msg_cnt);
    rd_kafka_toppar_unlock(rktp_uu);

    TAILQ_FOREACH_SAFE(rkm, &uas.rkmq_msgs, rkm_link, tmp) {
        /* Fast-path for failing messages with forced partition */
        // 无效的msg放到failed 队列
        if (rkm->rkm_partition != RD_KAFKA_PARTITION_UA &&
            rkm->rkm_partition >= rkt->rkt_partition_cnt &&
            rkt->rkt_state != RD_KAFKA_TOPIC_S_UNKNOWN) {
            rd_kafka_msgq_enqueue(&failed, rkm);
            continue;
        }

        // 重新路由kafka message到相应的partition, 失败则放入failed 队列
        if (unlikely(rd_kafka_msg_partitioner(rkt, rkm, 0) != 0)) {
            /* Desired partition not available */
            rd_kafka_msgq_enqueue(&failed, rkm);
        }
    }

    // 失败的msg, 都回调给application
    if (rd_atomic32_get(&failed.rkmq_msg_cnt) > 0) {
        /* Fail the messages */
        rd_kafka_dr_msgq(rkt, &failed,
            rkt->rkt_state == RD_KAFKA_TOPIC_S_NOTEXISTS ?
                err :
                RD_KAFKA_RESP_ERR_UNKNOWN_PARTITION);
    }

    rd_kafka_toppar_destroy(s_rktp_uu); /* from get() */
}
```

* 关于metadata相关的操作, 我们介绍metadata时再来分析

Librdkafka源码分析-Content Table

发布于 2019-01-18 09:01

[Kafka](#) [分布式存储](#) [大数据](#)

还没有评论

写下你的评论...

文章被以下专栏收录

 消息系统kafka的那些事儿
分享分布式消息系统kafka的点点滴滴

推荐阅读



搞定Kafka的存储架构, 看这篇就够了

华仔聊技术

发表于kafka...

Kafka设计解析 (一) Kafka背景及架构介绍
原创文章, 转载请注明出处。 (已授权InfoQ中文站发布) 本文转自技术世界, 原文链接 Kafka设计解析 (一) - Kafka背景及架构介绍! 摘要Kafka是由LinkedIn开发并开...

郭俊 Jason



图解Kafka

Enjoy...

发表于Enjoy...



Kafka零拷贝

洞类er



× 登录即可查看 超5亿 专业优质内容
超 5 千万创作者的优质提问、专业回答、深度文章和精彩视频尽在知乎。
[立即登录/注册](#)

[赞同 2](#) [添加评论](#) [分享](#) [喜欢](#) [收藏](#) [申请转载](#) ...

