

# **Working Draft American National Standard**

# **Project T10/BSR INCITS 506**

Revision 10  
21 January 2016

---

## **Information technology - SCSI Block Commands – 4 (SBC-4)**

This is an internal working document of T10, a Technical Committee of Accredited Standards Committee INCITS (International Committee for Information Technology Standards). As such this is not a completed standard and has not been approved. The contents may be modified by the T10 Technical Committee. The contents are actively being modified by T10. This document is made available for review and comment only.

Permission is granted to members of INCITS, its technical committees, and their associated task groups to reproduce this document for the purposes of INCITS standardization activities without further permission, provided this notice is included. All other rights are reserved. Any duplication of this document for commercial or for-profit use is strictly prohibited.

T10 Technical Editor:

William Martin  
Samsung Semiconductor, Inc  
7213 Marblethorpe Drive  
Roseville, CA 95747  
USA

Telephone: (916) 765-6875  
Email: [bill.martin@ssi.samsung.com](mailto:bill.martin@ssi.samsung.com)

---

Reference number  
ISO/IEC 14776-324:201x  
BSR INCITS 506:201x

## Points of contact

### International Committee for Information Technology Standards (INCITS) T10 Technical Committee

#### T10 Chair

Ralph Weber  
Western Digital Technologies  
18484 Preston Road, Suite 102, PMB 178  
Dallas, TX 75252  
USA

Telephone: (214) 912-1373  
Email: Ralph.Weber@WDC.com

#### T10 Vice-Chair

William Martin  
Samsung Semiconductor, Inc  
7213 Marblethorpe Drive  
Roseville, CA 95747  
USA

Telephone: (916) 765-6875  
Email: bill.martin@ssi.samsung.com

T10 Web Site: <http://www.t10.org>

#### T10 E-mail reflector:

Server: majordomo@t10.org  
To subscribe send e-mail with 'subscribe' in message body  
To unsubscribe send e-mail with 'unsubscribe' in message body

#### INCITS Secretariat

1101 K Street, NW  
Suite 610  
Washington, DC 20005  
USA

Telephone: (202) 737-8888  
Web site: <http://www.incits.org>  
Email: [incits@itic.org](mailto:incits@itic.org)

#### Information Technology Industry Council

Web site: <http://www.itic.org>

#### Document Distribution

INCITS Online Store  
managed by Techstreet  
1327 Jones Drive  
Ann Arbor, MI 48105  
USA

Web site: <http://www.techstreet.com/incitsgate.tmpl>  
Telephone: (734) 302-7801 or (800) 699-9277

Global Engineering Documents, an IHS Company  
15 Inverness Way East  
Englewood, CO 80112-5704  
USA

Web site: <http://global.ihs.com>  
Telephone: (303) 397-7956 or (303) 792-2181 or (800) 854-7179

American National Standard  
for Information Technology

## **SCSI Block Commands – 4 (SBC-4)**

Secretariat  
**Information Technology Industry Council**

Approved mm.dd.yy

American National Standards Institute, Inc.

### **Abstract**

This standard specifies the functional requirements for the SCSI Block Commands - 4 (SBC-4) command set. SBC-4 permits SCSI block logical units such as rigid disks to attach to computers and provides the definition for their use.

This standard maintains a high degree of compatibility with the SCSI Block Commands (SBC-3) command set, INCITS 514-2014, and while providing additional functions, is not intended to require changes to presently installed devices or existing software.

## American National Standard

Approval of an American National Standard requires verification by ANSI that the requirements for due process, consensus, and other criteria for approval have been met by the standards developer. Consensus is established when, in the judgment of the ANSI Board of Standards Review, substantial agreement has been reached by directly and materially affected interests. Substantial agreement means much more than a simple majority, but not necessarily unanimity. Consensus requires that all views and objections be considered, and that effort be made towards their resolution.

The use of American National Standards is completely voluntary; their existence does not in any respect preclude anyone, whether he has approved the standards or not, from manufacturing, marketing, purchasing, or using products, processes, or procedures not conforming to the standards.

The American National Standards Institute does not develop standards and will in no circumstances give interpretation on any American National Standard. Moreover, no person shall have the right or authority to issue an interpretation of an American National Standard in the name of the American National Standards Institute. Requests for interpretations should be addressed to the secretariat or sponsor whose name appears on the title page of this standard.

**CAUTION NOTICE:** This American National Standard may be revised or withdrawn at any time. The procedures of the American National Standards Institute require that action be taken periodically to reaffirm, revise, or withdraw this standard. Purchasers of American National Standards may receive current information on all standards by calling or writing the American National Standards Institute.

**CAUTION:** The developers of this standard have requested that holders of patents that may be required for the implementation of the standard, disclose such patents to the publisher. However, neither the developers nor the publisher have undertaken a patent search in order to identify which, if any, patents may apply to this standard. As of the date of publication of this standard, following calls for the identification of patents that may be required for the implementation of the standard, no such claims have been made. No further patent search is conducted by the developer or the publisher in respect to any standard it processes. No representation is made or implied that licenses are not required to avoid infringement in the use of this standard.

Published by

**American National Standards Institute  
11 W. 42nd Street, New York, New York 10036**

Copyright © 2014 by Information Technology Industry Council (ITI).  
All rights reserved.

No part of this publication may be reproduced in any form, in an electronic retrieval system or otherwise, without prior written permission of ITI, 1101 K Street, NW Suite 610, Washington, DC 20005.

Printed in the United States of America

## Revision History

This revision history is not part of American National Standard INCITS 1799:200x.

### R.1 Revision 0 (27 February 2014)

Revision 0 of SBC-4 is substantially equal to revision 36 of SBC-3. The only differences arise from changes made in SBC-3 discovered during the ISO process and modification of the format of the definitions to use the new ISO style of "Note to Entry" for definitions that are longer than a single sentence.

### R.2 Revision 1 (14 March 2014)

Incorporated the following proposals:

- a) Fixed "LENGTH OF SENSE DATA" to "LENGTH OF THE SENSE DATA FIELD" in 5.22.2;
- b) 14-003r1 - Add reporting of unrecoverable errors that aren't reassigned
- c) 14-019r2 - Discussion of WCE intent
- d) 14-061r0 - 'Not defined' is not defined sufficiently;
- e) 14-018r3 - Update Block Device Characteristics VPD page for ZBC; and
- f) 13-094r4 - Supported Block lengths and protection types VPD page.

### R.3 Revision 2 (20 May 2014)

Incorporated the following proposals:

- a) Corrected editorial errors on incorporation 14-003r1;
- b) 14-112r0 - Clarifying token timeout;
- c) 14-103r1 - Obsolete EER, DCR bits in Read-Write Recovery mode page;
- d) 14-075r1 - Designed Utilization VPD data; and
- e) 14-043r4 - Atomic writes and reads.

### R.4 Revision 3 (27 August 2014)

Incorporated the following proposals:

- a) Corrected errors on incorporation of 13-094r5;
- b) 14-148r3 - Redefinition of Peripheral device type;
- c) 14-143r0 - Enhancing SBC-4 to allow ZBC to share its VPD page definitions;
- d) 14-163r1 - Make the WRITE SAME no Data-Out Buffer bit broadly useful; and
- e) 14-178r1 - Disable PI checking on write if PI is all Fs.

### R.5 Revision 4(4 December 2014)

Incorporated the following proposals:

- a) Corrected errors on incorporation of 14-178r1;
- b) 14-210r1 Minimal Workload Utilization logging
- c) 14-064r5 Granular Boundaries for ATOMIC commands
- d) 14-128r4 Unmap methods and text clarifications
- e) 14-262r1 Thin Provisioning and De-Dupe interactions
- f) 14-238r0 Remove Unused MMC Restrictions
- g) 14-188r3 Resolve error cases with supporting Application Tag mode page
- h) 14-107r8 Command Deadlines
- i) 14-281r1 (Zone Control command
- j) 14-233r3 (READ CAPACITY (16)
- k) 14-194r2 (Update Block Device Characteristics VPD page for Self Managed Devices

**R.6 Revision 5 (22 January 2015)**

Incorporated the following proposals:

- a) Corrected errors on incorporation of 14-107r8
- b) 15-024r1 Clarify SANITIZE command error case when AUSE=1;
- c) 14-264r3 Logical Block Markup descriptors;
- d) 15-019r0 Adding NO\_PI\_CHK to Supported Block Lengths And Protection Types VPD;
- e) 14-249r2 Obsolete TMC and ETC bits;
- f) bring all of clause 2 into conformance with the currently approved T10 Style Guide T10/14-006r7

**R.7 Revision 6 (28 April 2015)**

Incorporated the following proposals:

- a) 15-062r2 Revise Logical Block Model to support block device extension technologies
- b) 15-060r0 Change obsolete fields in FORMAT UNIT command to reserved
- c) 14-065r3 Per IO Advice (hints)
- d) Various editorial corrections for previously incorporated proposals including fixing cross references

**R.8 Revision 7(20 May 2015)**

Incorporated the following proposals:

- a) 15-007r5 Additional provisioning reporting
- b) 15-147r0 Remove ignored keyword
- c) 15-075r2 Obsolete READ LONG and WRITE LONG except for write uncorrectable
- d) 15-141r0 Storage Intelligence
- e) 15-072r3 Result of SANITIZE for ZBC devices
- f) modifications for officer changes in T10
- g) Various editorial corrections for previously incorporated proposals including fixing cross references

**R.9 Revision 7a(22 May 2015)**

Incorporated the following:

- a) Changed Background Operation Control mode page from 0Ah/05h to 0Ah/06h.

**R.10 Revision 7b (26 May 2015)**

Incorporated the following:

- a) Changed Background Operation log page from 15h/01h to 15h/02h.

**R.11 Revision 7c (26 May 2015)**

Incorporated the following:

- a) Changed SPC-5 to SPC-4 except:
  - A) CBCS references;
  - B) ACCESS CONTROL IN command;
  - C) ACCESS CONTROL OUT command.

**R.12 Revision 8 (1 September 2015)**

Incorporate the following:

- a) 14-265r3 Allow device discovery through sanitize;
- b) 15-162r1 Clarify FUA bit wording;
- c) 15-186r0 Stream Corrections;
- d) VPD page code for the Block Limits Extension VPD page overlapped the Supported Block Lengths And Protection Types VPD page. The VPD page code for the Block Limits Extension VPD page was changed to B7h

**R.13 Revision 9 (9 November 2015)**

Incorporate the following:

- a) 15-193r1 Expunge I\_T\_L\_Q nexus transaction;
- b) 15-203r1 More Stream Correction;
- c) 15-218r2 Definition of Stream Identifier zero;
- d) 15-219r1 Reset of stream identifiers; and
- e) Modified occurrences of “field in ...” and “bit in ...”.

**R.14 [Revision 9 \(9 November 2015\)](#)**

[Incorporate the following:](#)

- a) 15-264r1 Allow REPORT ZONES command during sanitize operation;
- b) [15-117r6 Add fast format feature to FORMAT UNIT command;](#)
- c) 15-081r4 IO Access Hint Data Class;
- d) 15-037r6 Detection handling and reporting of unaligned writes;
- e) 16-044r0 Correction for commands allowed during sanitize;
- f) 15-288r2 Sanitize fix for ZBC;
- g) 15-293r0 Revise NDOB bit wording;
- h) 15-163r4 Organization of 'things supported' tables; and
- i) 16-005r1 UUID descriptor usage.

## Contents

	Page
Points of contact .....	ii
Abstract .....	iii
Revision History .....	v
Contents .....	viii
Tables .....	xviii
Figures .....	xxiv
Foreword .....	xxv
Introduction .....	xxv
SCSI standards family .....	xxvi
 1 Scope .....	 1
 2 Normative references .....	 1
 3 Definitions, symbols, abbreviations, keywords, and conventions Introduction .....	 2
3.1 Definitions .....	2
3.2 Symbols .....	11
3.3 Abbreviations .....	11
3.4 Keywords .....	12
3.5 Editorial conventions .....	13
3.6 Numeric and character conventions .....	14
3.6.1 Numeric conventions .....	14
3.6.2 Units of measure .....	14
3.7 State machine conventions .....	15
 4 Direct access block device type model .....	 17
4.1 Direct access block device type model introduction .....	17
4.2 Direct access block device type model .....	18
4.2.1 Direct access block device type model overview .....	18
4.2.2 Logical block access command types .....	18
4.2.3 Logical block access operation types .....	18
4.3 Media examples .....	19
4.3.1 Media examples overview .....	19
4.3.2 Rotating media .....	19
4.3.3 Memory media .....	19
4.4 Removable media .....	20
4.5 Logical blocks .....	20
4.6 Physical blocks .....	21
4.6.1 Overview .....	21
4.6.2 Physical block misaligned write reporting .....	24
4.7 Logical block provisioning .....	25
4.7.1 Logical block provisioning overview .....	25
4.7.2 Fully provisioned logical unit .....	26
4.7.3 Logical block provisioning management .....	26
4.7.3.1 Logical block provisioning management overview .....	26
4.7.3.2 Resource provisioned logical unit .....	27
4.7.3.3 Thin provisioned logical unit .....	27
4.7.3.4 Unmapping LBAs .....	28
4.7.3.4.1 Unmapping overview .....	28
4.7.3.4.2 Processing unmap requests .....	28
4.7.3.4.3 Unmap operations .....	28
4.7.3.4.4 WRITE SAME command and unmap operations .....	29
4.7.3.5 Autonomous LBA transitions .....	30
4.7.3.6 Logical unit resource exhaustion considerations .....	30



4.7.3.6.1 Thin provisioned logical unit resource exhaustion considerations .....	30
4.7.3.6.2 Resource provisioned logical unit resource exhaustion considerations .....	30
4.7.3.7 Logical block provisioning thresholds .....	31
4.7.3.7.1 Logical block provisioning thresholds overview .....	31
4.7.3.7.2 Threshold sets .....	31
4.7.3.7.3 Threshold percentages .....	32
4.7.3.7.4 Logical block provisioning armed decreasing thresholds .....	33
4.7.3.7.5 Logical block provisioning armed increasing thresholds .....	33
4.7.3.7.6 Logical block provisioning threshold notification .....	34
4.7.4 LBP (logical block provisioning) state machine .....	35
4.7.4.1 LBP state machine overview .....	35
4.7.4.2 LBP state machine for logical units supporting anchored LBAs .....	35
4.7.4.3 LBP state machine for logical units not supporting anchored LBAs .....	36
4.7.4.4 Performing read operations with respect to logical block provisioning .....	36
4.7.4.5 LBP1:Mapped state .....	38
4.7.4.5.1 LBP1:Mapped state description .....	38
4.7.4.5.2 Transition LBP1:Mapped to LBP2:Deallocated .....	39
4.7.4.5.3 Transition LBP1:Mapped to LBP3:Anchored .....	39
4.7.4.6 LBP2:Deallocated state .....	39
4.7.4.6.1 LBP2:Deallocated state description .....	39
4.7.4.6.2 Transition LBP2:Deallocated to LBP1:Mapped .....	39
4.7.4.6.3 Transition LBP2:Deallocated to LBP3:Anchored .....	39
4.7.4.7 LBP3:Anchored state .....	40
4.7.4.7.1 LBP3:Anchored state description .....	40
4.7.4.7.2 Transition LBP3:Anchored to LBP1:Mapped .....	40
4.7.4.7.3 Transition LBP3:Anchored to LBP2:Deallocated .....	40
4.8 Data de-duplication .....	40
4.9 Ready state .....	40
4.10 Initialization .....	41
4.11 Sanitize operations .....	41
4.11.1 Sanitize operations overview .....	41
4.11.2 Commands allowed during sanitize .....	42
4.11.3 Performing a sanitize operation .....	43
4.11.4 Completing a sanitize operation .....	43
4.12 Write protection .....	44
4.13 Medium defects .....	45
4.13.1 Medium defects overview .....	45
4.13.2 Generation of defect lists .....	47
4.14 Write and unmap failures .....	48
4.15 Caches .....	48
4.15.1 Caches overview .....	48
4.15.2 Cache segments .....	48
4.15.3 Read caching .....	49
4.15.4 Write caching .....	49
4.15.5 Command interactions with caches .....	49
4.15.6 Write operation and write medium operation interactions with caches .....	49
4.15.7 Read operation and read medium operation interactions with caches .....	50
4.15.8 Verify medium operation interactions with caches .....	50
4.15.9 Unmap operation interactions with caches .....	50
4.15.10 Power loss effects on caches .....	50
4.16 Implicit head of queue command processing .....	51
4.17 Reservations .....	51
4.18 Error reporting .....	54
4.18.1 Error reporting overview .....	54
4.18.2 Processing pseudo unrecovered errors .....	56
4.18.3 Block commands sense data descriptor .....	57
4.18.4 User data segment referral sense data descriptor .....	57

4.18.5 Direct access block device sense data descriptor .....	60
4.19 Model for XOR commands .....	61
4.19.1 Model for XOR commands overview .....	61
4.19.2 SCSI storage array device supervised XOR operations .....	62
4.19.2.1 SCSI storage array device supervised XOR operations overview .....	62
4.19.2.2 Update write operation .....	62
4.19.2.3 Regenerate operation.....	62
4.19.2.4 Rebuild operation .....	63
4.19.3 Array subsystem considerations .....	63
4.19.3.1 Array subsystem considerations overview .....	63
4.19.3.2 Access to an inconsistent stripe .....	63
4.20 Rebuild assist mode .....	63
4.20.1 Rebuild assist mode overview .....	63
4.20.2 Enabling rebuild assist mode .....	64
4.20.3 Using the rebuild assist mode .....	64
4.20.3.1 Using rebuild assist mode overview .....	64
4.20.3.2 Unpredicted unrecovered read error .....	64
4.20.3.3 Predicted unrecovered read error .....	65
4.20.3.4 Unpredicted unrecovered write error.....	65
4.20.3.5 Predicted unrecovered write error .....	65
4.20.4 Disabling the rebuild assist mode .....	66
4.20.5 Testing rebuild assist mode .....	66
4.21 START STOP UNIT and power conditions.....	66
4.21.1 START STOP UNIT and power conditions overview.....	66
4.21.2 Processing of concurrent START STOP UNIT commands.....	66
4.21.3 Managing logical block access commands during a change to the active power condition .....	67
4.21.4 Stopped power condition .....	67
4.21.5 START STOP UNIT and power condition state machine .....	67
4.21.5.1 START STOP UNIT and power condition state machine overview.....	67
4.21.5.2 SSU_PC0:Powered_On state .....	70
4.21.5.2.1 SSU_PC0:Powered_On state description .....	70
4.21.5.2.2 Transition SSU_PC0:Powered_On to SSU_PC4:Active_Wait .....	70
4.21.5.2.3 Transition SSU_PC0:Powered_On to SSU_PC8:Stopped .....	70
4.21.5.3 SSU_PC1:Active state .....	70
4.21.5.3.1 SSU_PC1:Active state description .....	70
4.21.5.3.2 Transition SSU_PC1:Active to SSU_PC5:Wait_Idle .....	70
4.21.5.3.3 Transition SSU_PC1:Active to SSU_PC6:Wait_Standby .....	70
4.21.5.3.4 Transition SSU_PC1:Active to SSU_PC10:Wait_Stopped .....	71
4.21.5.4 SSU_PC2:Idle state .....	71
4.21.5.4.1 SSU_PC2:Idle state description .....	71
4.21.5.4.2 Transition SSU_PC2:Idle to SSU_PC4:Active_Wait .....	71
4.21.5.4.3 Transition SSU_PC2:Idle to SSU_PC5:Wait_Idle .....	71
4.21.5.4.4 Transition SSU_PC2:Idle to SSU_PC6:Wait_Standby .....	72
4.21.5.4.5 Transition SSU_PC2:Idle to SSU_PC7:Idle_Wait .....	72
4.21.5.4.6 Transition SSU_PC2:Idle to SSU_PC10:Wait_Stopped .....	72
4.21.5.5 SSU_PC3:Standby state .....	72
4.21.5.5.1 SSU_PC3:Standby state description .....	72
4.21.5.5.2 Transition SSU_PC3:Standby to SSU_PC4:Active_Wait .....	73
4.21.5.5.3 Transition SSU_PC3:Standby to SSU_PC6:Wait_Standby.....	73
4.21.5.5.4 Transition SSU_PC3:Standby to SSU_PC7:Idle_Wait .....	73
4.21.5.5.5 Transition SSU_PC3:Standby to SSU_PC9:Standby_Wait.....	74
4.21.5.5.6 Transition SSU_PC3:Standby to SSU_PC10:Wait_Stopped .....	74
4.21.5.6 SSU_PC4:Active_Wait state .....	74
4.21.5.6.1 SSU_PC4:Active_Wait state description .....	74
4.21.5.6.2 Transition SSU_PC4:Active_Wait to SSU_PC1:Active .....	75
4.21.5.7 SSU_PC5:Wait_Idle state .....	75
4.21.5.7.1 SSU_PC5:Wait_Idle state description .....	75

4.21.5.7.2 Transition SSU_PC5:Wait_Idle to SSU_PC2:Idle .....	75
4.21.5.8 SSU_PC6:Wait_Standby state .....	75
4.21.5.8.1 SSU_PC6:Wait_Standby state description .....	75
4.21.5.8.2 Transition SSU_PC6:Wait_Standby to SSU_PC3:Standby .....	75
4.21.5.9 SSU_PC7:Idle_Wait state .....	75
4.21.5.9.1 SSU_PC7:Idle_Wait state description .....	75
4.21.5.9.2 Transition SSU_PC7:Idle_Wait to SSU_PC2:Idle .....	76
4.21.5.10 SSU_PC8:Stopped state .....	76
4.21.5.10.1 SSU_PC8:Stopped state description .....	76
4.21.5.10.2 Transition SSU_PC8:Stopped to SSU_PC4:Active_Wait .....	77
4.21.5.10.3 Transition SSU_PC8:Stopped to SSU_PC7:Idle_Wait .....	77
4.21.5.10.4 Transition SSU_PC8:Stopped to SSU_PC9:Standby_Wait .....	77
4.21.5.11 SSU_PC9:Standby_Wait state .....	77
4.21.5.11.1 SSU_PC9:Standby_Wait state description .....	77
4.21.5.11.2 Transition SSU_PC9:Standby_Wait to SSU_PC3:Standby .....	78
4.21.5.12 SSU_PC10:Wait_Stopped state .....	78
4.21.5.12.1 SSU_PC10:Wait_Stopped state description .....	78
4.21.5.12.2 Transition SSU_PC10:Wait_Stopped to SSU_PC8:Stopped .....	78
4.22 Protection information model .....	78
4.22.1 Protection information overview .....	78
4.22.2 Protection types .....	79
4.22.2.1 Protection types overview .....	79
4.22.2.2 Type 0 protection .....	79
4.22.2.3 Type 1 protection .....	80
4.22.2.4 Type 2 protection .....	80
4.22.2.5 Type 3 protection .....	81
4.22.3 Protection information format .....	82
4.22.4 Logical block guard .....	86
4.22.4.1 Logical block guard overview .....	86
4.22.4.2 CRC generation .....	86
4.22.4.3 CRC checking .....	87
4.22.4.4 CRC test cases .....	87
4.22.5 Application of protection information .....	87
4.22.6 Protection information and commands .....	88
4.23 Grouping function .....	88
4.23.1 Grouping function overview .....	88
4.23.2 Grouping function extensions for IO advice hints .....	88
4.24 Background scan operations .....	89
4.24.1 Background scan overview .....	89
4.24.2 Background pre-scan operations .....	90
4.24.2.1 Enabling background pre-scan operations .....	90
4.24.2.2 Suspending and resuming background pre-scan operations .....	90
4.24.2.3 Halting background pre-scan operations .....	91
4.24.3 Background medium scan .....	91
4.24.3.1 Enabling background medium scan operations .....	91
4.24.3.2 Suspending and resuming background medium scan operations .....	91
4.24.3.3 Halting background medium scan operations .....	92
4.24.4 Interpreting the logged background scan results .....	92
4.25 Association between commands and CbCS permission bits .....	94
4.26 Deferred microcode activation .....	95
4.27 Model for uninterrupted sequences on LBA ranges .....	96
4.28 Referrals .....	96
4.28.1 Referrals overview .....	96
4.28.2 Discovering referrals .....	97
4.28.3 Referrals in sense data .....	98
4.29 ORWRITE commands .....	99
4.29.1 ORWRITE commands overview .....	99

4.29.2 ORWgeneration code .....	99
4.29.2.1 ORWgeneration code overview .....	99
4.29.2.2 ORWgeneration code processing .....	100
4.29.3 Change generation and clear operation .....	100
4.29.4 Set operation .....	101
4.30 Block device ROD token operations .....	102
4.30.1 Block device ROD token operations overview .....	102
4.30.2 POPULATE TOKEN command and WRITE USING TOKEN command completion .....	103
4.30.3 Block device specific ROD tokens .....	103
4.30.4 Block device zero ROD token .....	104
4.30.5 ROD token device type specific data .....	104
4.31 Atomic writes .....	105
4.31.1 Atomic writes overview .....	105
4.31.2 Atomic write operations that do not complete .....	106
4.31.3 Performing operations with respect to atomic operations .....	106
4.31.3.1 Performing operations before and after an atomic write operation .....	106
4.31.3.2 Performing operations during an atomic write operation .....	107
4.31.4 Processing ACA conditions during atomic write commands .....	107
4.32 IO Advice Hints .....	108
4.32.1 IO Advice Hints Overview .....	108
4.32.2 Specifying IO Advice Hints .....	108
4.33 Background Operation Control .....	108
4.34 Stream Control .....	110
4.35 Format operations .....	112
4.35.1 Format operations overview .....	112
4.35.2 Performing a format operation .....	112
4.35.3 Completing a format operation .....	113
4.35.3.1 Completing a format operation overview .....	113
4.35.3.2 Completing read commands after a successful format operation .....	113
4.35.3.2.1 Completing read commands overview .....	113
4.35.3.2.2 With ffmt field set to 00b .....	114
4.35.3.2.3 With ffmt field set to 01b .....	114
4.35.3.2.4 With ffmt field set to 10b .....	114
5 Commands for direct access block devices .....	115
5.1 Commands for direct access block devices overview .....	115
5.2 BACKGROUND CONTROL command .....	119
5.2.1 BACKGROUND CONTROL command overview .....	119
5.3 COMPARE AND WRITE command .....	120
5.4 FORMAT UNIT command .....	122
5.4.1 FORMAT UNIT command overview .....	122
5.4.2 FORMAT UNIT parameter list .....	126
5.4.2.1 FORMAT UNIT parameter list overview .....	126
5.4.2.2 Parameter list header .....	127
5.4.2.3 Initialization pattern descriptor .....	131
5.5 GET LBA STATUS command .....	133
5.5.1 GET LBA STATUS command overview .....	133
5.5.2 GET LBA STATUS parameter data .....	134
5.5.2.1 GET LBA STATUS parameter data overview .....	134
5.5.2.2 LBA status descriptor .....	135
5.5.2.3 LBA status descriptor relationships .....	135
5.6 GET STREAM STATUS command .....	136
5.6.1 GET STREAM STATUS command overview .....	136
5.6.2 GET STREAM STATUS parameter data .....	137
5.6.2.1 GET STREAM STATUS parameter data overview .....	137
5.6.2.2 Stream status descriptor .....	138
5.6.2.3 Stream status descriptor relationships .....	138

5.7 ORWRITE (16) command .....	139
5.8 ORWRITE (32) command .....	145
5.9 POPULATE TOKEN command .....	147
5.9.1 POPULATE TOKEN command overview .....	147
5.9.2 POPULATE TOKEN parameter list .....	148
5.9.3 Block device range descriptor .....	150
5.10 PRE-FETCH (10) command .....	151
5.11 PRE-FETCH (16) command .....	152
5.12 PREVENT ALLOW MEDIUM REMOVAL command .....	153
5.13 READ (10) command .....	154
5.14 READ (12) command .....	158
5.15 READ (16) command .....	160
5.16 READ (32) command .....	161
5.17 READ CAPACITY (10) command .....	162
5.17.1 READ CAPACITY (10) overview .....	162
5.17.2 READ CAPACITY (10) parameter data .....	163
5.18 READ CAPACITY (16) command .....	163
5.18.1 READ CAPACITY (16) command overview .....	163
5.18.2 READ CAPACITY (16) parameter data .....	165
5.19 READ DEFECT DATA (10) command .....	166
5.19.1 READ DEFECT DATA (10) command overview .....	166
5.19.2 READ DEFECT DATA (10) parameter data .....	168
5.20 READ DEFECT DATA (12) command .....	168
5.20.1 READ DEFECT DATA (12) command overview .....	168
5.20.2 READ DEFECT DATA (12) parameter data .....	170
5.21 REASSIGN BLOCKS command .....	171
5.21.1 REASSIGN BLOCKS command overview .....	171
5.21.2 REASSIGN BLOCKS parameter list .....	172
5.22 RECEIVE ROD TOKEN INFORMATION .....	174
5.22.1 RECEIVE ROD TOKEN INFORMATION overview .....	174
5.22.2 RECEIVE ROD TOKEN INFORMATION parameter data for POPULATE TOKEN command .....	175
5.22.3 RECEIVE ROD TOKEN INFORMATION parameter data for WRITE USING TOKEN command .....	178
5.23 REPORT REFERRALS command .....	179
5.23.1 REPORT REFERRALS command overview .....	179
5.23.2 REPORT REFERRALS parameter data .....	180
5.24 REPORT PROVISIONING INITIALIZATION PATTERN command .....	181
5.25 SANITIZE command .....	181
5.25.1 SANITIZE command overview .....	181
5.25.2 SANITIZE command service actions .....	182
5.25.2.1 SANITIZE command service actions overview .....	182
5.25.2.2 OVERWRITE service action .....	183
5.25.2.3 BLOCK ERASE service action .....	184
5.25.2.4 CRYPTOGRAPHIC ERASE service action .....	185
5.25.2.5 EXIT FAILURE MODE service action .....	185
5.26 START STOP UNIT command .....	186
5.27 STREAM CONTROL command .....	189
5.27.1 STREAM CONTROL command overview .....	189
5.27.2 STREAM CONTROL parameter data .....	191
5.28 SYNCHRONIZE CACHE (10) command .....	191
5.29 SYNCHRONIZE CACHE (16) command .....	193
5.30 UNMAP command .....	194
5.30.1 UNMAP command overview .....	194
5.30.2 UNMAP parameter list .....	195
5.31 VERIFY (10) command .....	196
5.32 VERIFY (12) command .....	209
5.33 VERIFY (16) command .....	210
5.34 VERIFY (32) command .....	211

5.35 WRITE (10) command .....	212
5.36 WRITE (12) command .....	216
5.37 WRITE (16) command .....	217
5.38 WRITE (32) command .....	218
5.39 WRITE AND VERIFY (10) command .....	219
5.40 WRITE AND VERIFY (12) command .....	220
5.41 WRITE AND VERIFY (16) command .....	221
5.42 WRITE AND VERIFY (32) command .....	222
5.43 WRITE ATOMIC (16) command .....	223
5.44 WRITE ATOMIC (32) command .....	224
5.45 WRITE LONG (10) command .....	225
5.46 WRITE LONG (16) command .....	226
5.47 WRITE SAME (10) command .....	226
5.48 WRITE SAME (16) command .....	229
5.49 WRITE SAME (32) command .....	230
5.50 WRITE STREAM (16) command .....	232
5.51 WRITE STREAM (32) command .....	232
5.52 WRITE USING TOKEN command .....	234
5.52.1 WRITE USING TOKEN command overview .....	234
5.52.2 WRITE USING TOKEN parameter list .....	235
5.53 XDWRITEREAD (10) command .....	237
5.54 XDWRITEREAD (32) command .....	239
5.55 XPWRITE (10) command .....	240
5.56 XPWRITE (32) command .....	241
6 Parameters for direct access block devices .....	243
6.1 Parameters for direct access block devices introduction .....	243
6.2 Address descriptors .....	243
6.2.1 Address descriptor overview .....	243
6.2.2 Short block format address descriptor .....	244
6.2.3 Extended bytes from index address descriptor .....	244
6.2.4 Extended physical sector format address descriptor .....	246
6.2.5 Long block format address descriptor .....	247
6.2.6 Bytes from index format address descriptor .....	247
6.2.7 Physical sector format address descriptor .....	248
6.3 Diagnostic parameters .....	249
6.3.1 Diagnostic parameters overview .....	249
6.3.2 Rebuild Assist Input diagnostic page .....	250
6.3.3 Rebuild Assist Output diagnostic page .....	251
6.3.4 Translate Address Input diagnostic page .....	252
6.3.5 Translate Address Output diagnostic page .....	254
6.4 Log parameters .....	255
6.4.1 Log parameters overview .....	255
6.4.1.1 Summary of log pages .....	255
6.4.1.2 Setting and resetting log parameters .....	257
6.4.2 Background Scan log page .....	257
6.4.2.1 Background Scan log page overview .....	257
6.4.2.2 Background Scan Status log parameter .....	259
6.4.2.3 Background Scan Results log parameter .....	261
6.4.3 Background Operation log page .....	264
6.4.3.1 Background Operation log page overview .....	264
6.4.3.2 Background Operation log parameter .....	265
6.4.4 Format Status log page .....	266
6.4.4.1 Format Status log page overview .....	266
6.4.4.2 Format Data Out log parameter .....	267
6.4.4.3 Grown Defects During Certification log parameter .....	268
6.4.4.4 Total Blocks Reassigned During Format log parameter .....	269

6.4.4.5 Total New Blocks Reassigned log parameter .....	270
6.4.4.6 Power On Minutes Since Format log parameter .....	271
6.4.5 Logical Block Provisioning log page .....	272
6.4.5.1 Logical Block Provisioning log page overview .....	272
6.4.5.2 Available LBA Mapping Resource Count log parameter .....	274
6.4.5.2.1 Available LBA Mapping Resource Count log parameter overview .....	274
6.4.5.2.2 RESOURCE COUNT field .....	275
6.4.5.3 Used LBA Mapping Resource Count log parameter .....	275
6.4.5.4 Available Provisioning Resource Percentage log parameter .....	276
6.4.5.4.1 Available Provisioning Resource Percentage log parameter overview .....	276
6.4.5.4.2 resource count field .....	276
6.4.5.5 De-duplicated LBA Resource Count log parameter .....	277
6.4.5.6 Compressed LBA Resource Count log parameter .....	278
6.4.5.7 Total Efficiency LBA Resource Count log parameter .....	279
6.4.6 LPS Misalignment log page .....	280
6.4.6.1 Overview .....	280
6.4.6.2 LPS Misalignment Count log parameter .....	281
6.4.6.3 LPS Misalignment log parameter .....	281
6.4.7 Non-volatile Cache log page .....	282
6.4.7.1 Non-volatile Cache log page overview .....	282
6.4.7.2 Remaining Nonvolatile Time log parameter .....	283
6.4.7.3 Maximum Nonvolatile Time log parameter .....	284
6.4.8 Pending Defects log page .....	285
6.4.8.1 Overview .....	285
6.4.8.2 Pending Defect Count log parameter .....	286
6.4.8.3 Pending Defect log parameter .....	287
6.4.9 Solid State Media log page .....	288
6.4.9.1 Solid State Media log page overview .....	288
6.4.9.2 Percentage Used Endurance Indicator log parameter .....	289
6.4.10 Utilization log page .....	290
6.4.10.1 Utilization log page overview .....	290
6.4.10.2 Workload Utilization log parameter .....	291
6.4.10.3 Utilization Usage Rate Based on Date and Time .....	292
6.5 Mode parameters .....	293
6.5.1 Mode pagesparameters overview .....	293
6.5.2 Mode parameter block descriptors .....	295
6.5.2.1 Mode parameter block descriptors overview .....	295
6.5.2.2 Short LBA mode parameter block descriptor .....	295
6.5.2.3 Long LBA mode parameter block descriptor .....	297
6.5.3 Application Tag mode page .....	298
6.5.3.1 Introduction .....	298
6.5.3.2 Application tag descriptor .....	300
6.5.4 Background Control mode page .....	301
6.5.5 Background Operation Control mode page .....	302
6.5.6 Caching mode page .....	304
6.5.7 IO Advice Hints Grouping mode page .....	308
6.5.8 Informational Exceptions Control mode page .....	309
6.5.9 Logical Block Provisioning mode page .....	314
6.5.9.1 Introduction .....	314
6.5.9.2 Threshold descriptor format .....	315
6.5.10 Read-Write Error Recovery mode page .....	316
6.5.11 Verify Error Recovery mode page .....	320
6.6 Vital product data (VPD) parameters .....	322
6.6.1 VPD pagesparameters overview .....	322
6.6.2 Block Device Characteristics VPD page .....	324
6.6.3 Block Device Characteristics Extension VPD page .....	327
6.6.4 Block Limits VPD page .....	329

6.6.5 Block Limits Extension VPD page.....	333
6.6.6 Logical Block Provisioning VPD page.....	335
6.6.7 Referrals VPD page.....	338
6.6.8 Third-Party Copy VPD page.....	339
6.6.8.1 Third-Party Copy VPD page overview.....	339
6.6.8.2 Block device third-party copy descriptor type codes.....	339
6.6.8.3 Block Device ROD Token Limits descriptor.....	340
6.6.9 Supported Block Lengths And Protection Types VPD page.....	341
6.7 Copy manager parameters.....	343
6.8 Logical Block Markup descriptors.....	344
6.8.1 Summary of LBM descriptors.....	344
6.8.2 LBM descriptor formats and types.....	344
6.8.3 Access patterns LBM descriptors.....	344
6.8.3.1 Access patterns LBM descriptor format.....	344
6.8.3.2 Accesses continue during low utilization (acdlu) bit.....	345
6.8.3.3 write sequentiality field and read sequentiality field.....	345
6.8.3.4 READ/WRITE FREQUENCY field.....	345
6.8.3.5 overall frequency field.....	345
6.8.3.6 subsequent i/o field.....	346
6.8.3.7 osi proximity field.....	346
6.8.3.8 IO CLASS field.....	347
6.8.3.9 The io class field (see table 275) specifies the type that is associated with this LBM descriptor....	347
6.8.3.10 Access patterns LBM descriptor usage considerations.....	347
Annex A (informative) Numeric order codes.....	348
A.1 Variable length CDBs.....	348
A.2 Service action CDBs.....	349
Annex B (informative) XOR command examples.....	350
B.1 XOR command examples overview.....	350
B.2 Update write operation.....	350
B.3 Regenerate operation.....	351
B.4 Rebuild operation.....	352
Annex C (informative) CRC example in C.....	354
Annex D (informative) Sense information for locked or encrypted logical units.....	356
Annex E (informative) Optimizing block access characteristics.....	357
E.1 Optimizing block access overview.....	357
E.2 Starting logical block offset.....	357
E.3 Optimal granularity sizes.....	357
E.4 Optimal stream granularity sizes.....	357
E.5 Optimizing transfers.....	358
E.5.1 Overview.....	358
E.5.2 Optimizing non-stream transfers.....	358
E.5.3 Optimizing stream transfers.....	359
E.6 Examples.....	359
Annex F (informative) Logical block provisioning reporting examples.....	361
F.1 Logical block provisioning reporting examples overview.....	361
F.2 Interpreting log parameter counts.....	361
F.3 Dedicated resource, threshold set tracked example.....	363
F.3.1 Dedicated resource, threshold set tracked example overview.....	363
F.3.2 Dedicated resource, threshold set tracked example configuration.....	363
F.3.3 Dedicated resource, threshold set tracked example sequence.....	364



F.3.4 Dedicated resource, threshold set tracked example initial conditions .....	365
F.3.5 Operations that occur .....	365
F.3.6 Dedicated resource, threshold set tracked example final log page values .....	366
F.4 Shared resource, logical block tracked example .....	366
F.4.1 Shared resource, logical block tracked example overview .....	366
F.4.2 Shared resource, logical block tracked example configuration .....	367
F.4.3 Shared resource, logical block tracked example time line .....	367
F.4.4 Shared resource, logical block tracked example initial conditions .....	368
F.4.5 Operations that occur .....	368
F.4.6 Shared resource, logical block tracked example final log page values .....	369
F.5 Shared available, dedicated used, logical block tracked example .....	370
F.5.1 Shared available, dedicated used, logical block tracked example overview .....	370
F.5.2 Shared available, dedicated used, logical block tracked example configuration .....	370
F.5.3 Shared available, dedicated used, logical block tracked example time line .....	370
F.5.4 Shared available, dedicated used, logical block tracked example initial conditions .....	371
F.5.5 Operations that occur .....	371
F.5.6 Shared available, dedicated used, example final log page values .....	372
Annex G (informative) Discovering referrals examples .....	373
G.1 Referrals example with no user data segment multiplier .....	373
G.2 Referrals example with non-zero user data segment multiplier .....	375
Annex H (informative) IO Advice Hints Usage .....	377
H.1 IO Advice Hints Overview .....	377
H.2 IO Advice Hints Grouping mode page .....	377
H.3 Issuing I/O commands .....	377
H.3.1 Group numbers and I/O commands .....	377
H.3.2 Possible constraints on IO advice hints .....	377
H.4 Logical Block Markup descriptor usage examples .....	378
H.4.1 Example usage in tiered storage device implementations .....	378
H.4.2 Example LBM descriptor values for software that sends read commands and write commands ..	378
Annex I (informative) Bibliography .....	381

## Tables

	Page
Table 1 — Direct access block device type mode topics and references .....	2
Table 2 — Numbering convention examples .....	14
Table 3 — Comparison of decimal prefixes and binary prefixes .....	15
Table 4 — Direct access block device type model topics .....	17
Table 5 — Logical block provisioning states supported by logical block provisioning type .....	26
Table 6 — WRITE SAME command and unmap operations .....	29
Table 7 — Threshold resource value, threshold type value, and threshold arming value for logical block provisioning thresholds .....	32
Table 8 — Threshold resource value, threshold type value, and threshold arming value for logical block provisioning percentages .....	32
Table 9 — Logical block data returned by a read operation from a mapped LBA .....	37
Table 10 — Logical block data returned by a read operation from an unmapped LBA .....	38
Table 11 — Defect lists (i.e., PLIST and GLIST) .....	46
Table 12 — Address descriptor formats .....	47
Table 13 — SBC-4 commands that are allowed in the presence of various reservations .....	52
Table 14 — Example error conditions .....	55
Table 15 — Sense data field usage for direct access block devices .....	56
Table 16 — Block commands sense data descriptor format .....	57
Table 17 — User data segment referral sense data descriptor format .....	58
Table 18 — User data segment referral descriptor format .....	59
Table 19 — Target port group descriptor .....	60
Table 20 — Direct access block device sense data descriptor format .....	61
Table 21 — Summary of states in the SSU_PC state machine .....	68
Table 22 — Logical block data format with a single protection information interval .....	82
Table 23 — An example of the logical block data for a logical block with more than one protection information interval .....	83
Table 24 — Content of the first LOGICAL BLOCK REFERENCE TAG field for the first logical block in the Data-In Buffer and/or Data-Out Buffer .....	84
Table 25 — Content of subsequent LOGICAL BLOCK REFERENCE TAG fields for a logical block in the Data-In Buffer and/or Data-Out Buffer .....	85
Table 26 — CRC polynomials .....	86
Table 27 — CRC test cases .....	87
Table 28 — Associations between commands and CbCS permissions .....	94
Table 29 — Commands that require uninterrupted sequences .....	96
Table 30 — Performing an ORWRITE set operation .....	101
Table 31 — ROD token type values .....	104
Table 32 — Block device zero ROD token format .....	104
Table 33 — Performing atomic operations with overlapping LBAs during current operations .....	107
Table 34 — Commands for direct access block devices .....	115
Table 35 — BACKGROUND CONTROL command .....	119
Table 36 — BO_CTL field .....	120
Table 37 — COMPARE AND WRITE command .....	121
Table 38 — FORMAT UNIT command .....	123
Table 39 — FORMAT UNIT command address descriptor support requirements .....	125
Table 40 — FFMT field description .....	126
Table 41 — FORMAT UNIT parameter list .....	126
Table 42 — Short parameter list header .....	127
Table 43 — Long parameter list header .....	127
Table 44 — FMTINFO field and PROTECTION FIELD USAGE field .....	128
Table 45 — Initialization pattern descriptor .....	131
Table 46 — INITIALIZATION PATTERN TYPE field .....	132
Table 47 — GET LBA STATUS command .....	133
Table 48 — GET LBA STATUS parameter data .....	134
Table 49 — LBA status descriptor format .....	135

Table 50 — PROVISIONING STATUS field .....	135
Table 51 — GET STREAM STATUS command .....	136
Table 52 — GET STREAM STATUS parameter data .....	137
Table 53 — Stream status descriptor format .....	138
Table 54 — ORWRITE (16) command .....	139
Table 55 — ORPROTECT field - checking protection information from the read operations .....	140
Table 56 — ORPROTECT field - checking protection information from the Data-Out Buffer .....	143
Table 57 — ORWRITE (32) command .....	145
Table 58 — BMOP field .....	146
Table 59 — POPULATE TOKEN command .....	147
Table 60 — POPULATE TOKEN parameter list .....	148
Table 61 — Block device range descriptor .....	150
Table 62 — PRE-FETCH (10) command .....	151
Table 63 — PRE-FETCH (16) command .....	152
Table 64 — PREVENT ALLOW MEDIUM REMOVAL command .....	153
Table 65 — PREVENT field .....	153
Table 66 — READ (10) command .....	154
Table 67 — RDPROTECT field .....	155
Table 68 — READ (12) command .....	159
Table 69 — READ (16) command .....	160
Table 70 — Duration limit value bits specifying command duration limit descriptor .....	160
Table 71 — READ (32) command .....	161
Table 72 — READ CAPACITY (10) command .....	162
Table 73 — READ CAPACITY (10) parameter data .....	163
Table 74 — READ CAPACITY (16) command .....	164
Table 75 — READ CAPACITY (16) parameter data .....	165
Table 76 — P_TYPE field and PROT_EN bit .....	165
Table 77 — LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT field .....	166
Table 78 — READ DEFECT DATA (10) command .....	167
Table 79 — REQ_PLIST bit and REQ_GLIST bit .....	167
Table 80 — READ DEFECT DATA (10) parameter data .....	168
Table 81 — READ DEFECT DATA (12) command .....	169
Table 82 — READ DEFECT DATA (12) parameter data .....	170
Table 83 — REASSIGN BLOCKS command .....	172
Table 84 — REASSIGN BLOCKS parameter list .....	172
Table 85 — REASSIGN BLOCKS short parameter list header .....	173
Table 86 — REASSIGN BLOCKS long parameter list header .....	173
Table 87 — Reassign LBA if the LONGLBA bit is set to zero .....	173
Table 88 — Reassign LBA if the LONGLBA bit is set to one .....	174
Table 89 — RECEIVE ROD TOKEN INFORMATION reference .....	175
Table 90 — RECEIVE ROD TOKEN INFORMATION parameter data for POPULATE TOKEN .....	176
Table 91 — RECEIVE ROD TOKEN INFORMATION parameter data for WRITE USING TOKEN .....	178
Table 92 — REPORT REFERRALS command .....	179
Table 93 — REPORT REFERRALS parameter data .....	180
Table 94 — REPORT PROVISIONING INITIALIZATION PATTERN command .....	181
Table 95 — SANITIZE command .....	181
Table 96 — SANITIZE service action codes .....	183
Table 97 — OVERWRITE service action parameter list .....	183
Table 98 — TEST field .....	184
Table 99 — START STOP UNIT command .....	186
Table 100 — POWER CONDITION and POWER CONDITION MODIFIER field .....	187
Table 101 — STREAM CONTROL command .....	190
Table 102 — STR_CTL field .....	190
Table 103 — STREAM CONTROL parameter data .....	191
Table 104 — SYNCHRONIZE CACHE (10) command .....	192
Table 105 — SYNCHRONIZE CACHE (16) command .....	193
Table 106 — UNMAP command .....	194

Table 107 — UNMAP parameter list .....	195
Table 108 — UNMAP block descriptor .....	196
Table 109 — Data-Out Buffer contents for the VERIFY (10) command .....	197
Table 110 — VERIFY (10) command .....	197
Table 111 — VRPROTECT field with the BYTCHK field set to 00b – checking protection information from the verify operations .....	199
Table 112 — VRPROTECT field with the BYTCHK field set to 01b or 11b – checking protection information from the verify operations .....	202
Table 113 — VRPROTECT field with the BYTCHK field set to 01b or 11b – checking protection information from the Data-Out Buffer .....	204
Table 114 — VRPROTECT field with the BYTCHK field set to 01b or 11b – compare operation requirements ..	206
Table 115 — VERIFY (12) command .....	209
Table 116 — VERIFY (16) command .....	210
Table 117 — VERIFY (32) command .....	211
Table 118 — WRITE (10) command .....	212
Table 119 — WRPROTECT field .....	213
Table 120 — WRITE (12) command .....	216
Table 121 — WRITE (16) command .....	217
Table 122 — WRITE (32) command .....	218
Table 123 — WRITE AND VERIFY (10) command .....	219
Table 124 — WRITE AND VERIFY (12) command .....	220
Table 125 — WRITE AND VERIFY (16) command .....	221
Table 126 — WRITE AND VERIFY (32) command .....	222
Table 127 — WRITE ATOMIC (16) command .....	223
Table 128 — WRITE ATOMIC (32) command .....	224
Table 129 — WRITE LONG (10) command .....	225
Table 130 — WR_UNCOR bit .....	225
Table 131 — WRITE LONG (16) command .....	226
Table 132 — WRITE SAME (10) command .....	228
Table 133 — UNMAP bit, ANCHOR bit, and ANC_SUP bit relationships .....	228
Table 134 — WRITE SAME (16) command .....	229
Table 135 — WRITE SAME (32) command .....	231
Table 136 — WRITE STREAM (16) command .....	232
Table 137 — WRITE STREAM(32) command .....	233
Table 138 — WRITE USING TOKEN command .....	234
Table 139 — WRITE USING TOKEN parameter list .....	235
Table 140 — XDWRITEREAD (10) command .....	238
Table 141 — XDWRITEREAD (32) command .....	239
Table 142 — XPWRITE (10) command .....	240
Table 143 — XPWRITE (32) command .....	241
Table 144 — Parameters for direct access block devices .....	243
Table 145 — Address descriptors .....	244
Table 146 — Short block format address descriptor (000b) .....	244
Table 147 — Extended bytes from index format address descriptor (001b) .....	245
Table 148 — Sorting order for extended bytes from index format address descriptors .....	245
Table 149 — Extended physical sector format address descriptor (010b) .....	246
Table 150 — Sorting order for extended physical sector format address descriptors .....	247
Table 151 — Long block format address descriptor (011b) .....	247
Table 152 — Bytes from index format address descriptor (100b) .....	247
Table 153 — Sorting order for bytes from index format address descriptors .....	248
Table 154 — Physical sector format address descriptor (101b) .....	248
Table 155 — Sorting order for physical sector format address descriptors .....	248
Table 156 — Diagnostic page codes for direct access block devices .....	249
Table 157 — Rebuild Assist Input diagnostic page .....	250
Table 158 — Rebuild Assist Output diagnostic page .....	251
Table 159 — Translate Address Input diagnostic page .....	252
Table 160 — Translate Address Output diagnostic page .....	254

Table 161 — Log page codes and subpage codes for direct access block devices .....	256
Table 162 — Keywords for resetting or changing log parameters .....	257
Table 163 — Background Scan log page parameter codes .....	258
Table 164 — Background Scan log page .....	258
Table 165 — Background Scan Status log parameter format .....	259
Table 166 — BACKGROUND SCAN STATUS field .....	260
Table 167 — Background Scan Results log parameter format .....	261
Table 168 — REASSIGN STATUS field .....	262
Table 169 — Background Operation log page .....	264
Table 170 — Background Operation log page parameter codes .....	264
Table 171 — Background Operation log parameter format .....	265
Table 172 — bo_status definitions .....	265
Table 173 — Format Status log page parameter codes .....	266
Table 174 — Format Status log page .....	266
Table 175 — Format Data Out log parameter format .....	267
Table 176 — Grown Defects During Certification log parameter format .....	268
Table 177 — Total Blocks Reassigned During Format log parameter format .....	269
Table 178 — Total New Blocks Reassigned log parameter format .....	270
Table 179 — Power On Minutes Since Format log parameter format .....	271
Table 180 — Logical Block Provisioning log parameters .....	272
Table 181 — Logical Block Provisioning log page .....	273
Table 182 — Available LBA Mapping Resource Count log parameter format .....	274
Table 183 — SCOPE field .....	274
Table 184 — Used LBA Mapping Resource Count log parameter format .....	275
Table 185 — Available Provisioning Resource Percentage log parameter format .....	276
Table 186 — resource count field .....	276
Table 187 — De-duplicated LBA Resource Count log parameter format .....	277
Table 188 — Compressed LBA Resource Count log parameter format .....	278
Table 189 — Total Efficiency LBA Resource Count log parameter format .....	279
Table 190 — LPS Misalignment log page parameter codes .....	280
Table 191 — LPS Misalignment log page .....	280
Table 192 — LPS Misalignment Count log parameter format .....	281
Table 193 — LPS Misalignment log parameter format .....	282
Table 194 — Nonvolatile Cache log parameters .....	282
Table 195 — Nonvolatile Cache log page .....	283
Table 196 — Remaining Nonvolatile Time log parameter format .....	283
Table 197 — REMAINING NONVOLATILE TIME field .....	284
Table 198 — Maximum Nonvolatile Time log parameter format .....	284
Table 199 — MAXIMUM NONVOLATILE TIME field .....	285
Table 200 — Pending Defects log page parameter codes .....	285
Table 201 — Pending Defects log page .....	286
Table 202 — Pending Defect Count log parameter format .....	286
Table 203 — Pending Defect log parameter format .....	287
Table 204 — Solid State Media log page .....	288
Table 205 — Solid State Media log parameters .....	289
Table 206 — Percentage Used Endurance Indicator log parameter format .....	289
Table 207 — Utilization log page parameter codes .....	290
Table 208 — Utilization log page .....	290
Table 209 — Workload Utilization log parameter format .....	291
Table 210 — workload utilization field .....	291
Table 211 — Utilization Rate Based on Date and Time log parameter format .....	292
Table 212 — date and time based utilization rate field .....	292
Table 213 — Mode page codes and subpage codes for direct access block devices .....	293
Table 214 — DEVICE-SPECIFIC PARAMETER field for direct access block devices .....	294
Table 215 — Short LBA mode parameter block descriptor .....	295
Table 216 — Long LBA mode parameter block descriptor .....	297
Table 217 — Application Tag mode page .....	299

Table 218 — Application tag descriptor format .....	300
Table 219 — Background Control mode page .....	301
Table 220 — Background Operation Control mode page .....	302
Table 221 — bo_mode field .....	303
Table 222 — Caching mode page .....	304
Table 223 — DEMAND READ RETENTION PRIORITY field .....	305
Table 224 — WRITE RETENTION PRIORITY field .....	306
Table 225 — SYNC_PROG field .....	307
Table 226 — IO Advice Hints Grouping mode page .....	308
Table 227 — IO advice hints group descriptor .....	309
Table 228 — IO ADVICE HINTS MODE field .....	309
Table 229 — Informational Exceptions Control mode page .....	310
Table 230 — Definitions for the combinations of values in EWASC, DEXCPT, and TEST .....	311
Table 231 — Method of reporting informational exceptions (MRIE) field .....	312
Table 232 — Use of the INTERVAL TIMER field and the REPORT COUNT field based on the MRIE field .....	313
Table 233 — Logical Block Provisioning mode page .....	314
Table 234 — Threshold descriptor format .....	315
Table 235 — THRESHOLD TYPE field .....	315
Table 236 — THRESHOLD ARMING field .....	315
Table 237 — Read-Write Error Recovery mode page .....	316
Table 238 — Error recovery bit combinations .....	318
Table 239 — MWR field .....	319
Table 240 — Verify Error Recovery mode page .....	320
Table 241 — VPD page codes for direct access block devices .....	323
Table 242 — Block Device Characteristics VPD page .....	324
Table 243 — MEDIUM ROTATION RATE field .....	324
Table 244 — PRODUCT TYPE field .....	325
Table 245 — WABEREQ field .....	325
Table 246 — WACEREQ field .....	326
Table 247 — NOMINAL FORM FACTOR field .....	326
Table 248 — zoned field .....	326
Table 249 — Block Device Characteristics Extension VPD page .....	327
Table 250 — UTILIZATION TYPE field .....	328
Table 251 — utilization units field .....	328
Table 252 — utilization interval field .....	328
Table 253 — Block Limits VPD page .....	329
Table 254 — Transfer limits for commands .....	331
Table 255 — Block Limits Extension VPD page .....	333
Table 256 — Logical Block Provisioning VPD page .....	335
Table 257 — threshold percentage field .....	336
Table 258 — minimum percentage field .....	336
Table 259 — LBPRZ field .....	336
Table 260 — PROVISIONING TYPE field .....	337
Table 261 — Referrals VPD page .....	338
Table 262 — Block device third-party copy descriptor type codes .....	339
Table 263 — Block Device ROD Token Limits descriptor .....	340
Table 264 — Supported Block Lengths And Protection Types VPD page .....	341
Table 265 — Logical block length and protection types descriptor format .....	342
Table 266 — ROD token device type specific data .....	343
Table 267 — Parameters for direct access block devices .....	344
Table 268 — lbm descriptor type field .....	344
Table 269 — Access patterns LBM descriptor format .....	344
Table 270 — WRITE SEQUENTIALITY field and READ SEQUENTIALITY field .....	345
Table 271 — READ/WRITE FREQUENCY field .....	345
Table 272 — OVERALL FREQUENCY field .....	346
Table 273 — SUBSEQUENT I/O field .....	346
Table 274 — OSI PROXIMITY field .....	346

Table 275 — io class field .....	347
Table A.1 — Variable length command service action code assignments .....	348
Table A.2 — SERVICE ACTION IN (16) service actions .....	349
Table A.3 — SERVICE ACTION OUT (16) service actions .....	349
Table D.1 — Sense information for locked or encrypted logical units .....	356
Table F.1 — Dedicated resource, threshold set tracked example capacity information .....	363
Table F.2 — Dedicated resource, threshold set tracked example capacity information .....	364
Table F.3 — Dedicated resource, threshold set tracked example initial conditions .....	365
Table F.4 — Dedicated resource, threshold set tracked example final log page values .....	366
Table F.5 — Shared resource, logical block tracked example capacity information .....	367
Table F.6 — Shared resource, logical block tracked example initial conditions .....	368
Table F.7 — Shared resource, logical block tracked example final log page values .....	369
Table F.8 — Shared available, dedicated used example capacity information .....	370
Table F.9 — Shared resource, logical block tracked example initial conditions .....	371
Table F.10 — Shared available, dedicated used example final log page values .....	372
Table G.1 — Referrals application client information with no user data segment multiplier .....	374
Table G.2 — User data segment calculations with no user data segment multiplier .....	374
Table G.3 — Referrals application client information with non-zero user data segment multiplier .....	376
Table G.4 — User data segment calculations with non-zero user data segment multiplier .....	376
Table H.1 — Tiered product access patterns LBM descriptor examples .....	378
Table H.2 — Sending device access patterns LBM descriptor examples .....	379

## Figures

	Page
Figure 0 — SCSI document relationships .....	xxvi
Figure 1 — Example state machine figure .....	15
Figure 2 — One or more physical blocks per logical block examples .....	22
Figure 3 — One or more logical blocks per physical block examples .....	23
Figure 4 — Two logical blocks per physical block alignment examples .....	23
Figure 5 — Four logical blocks per physical block alignment examples .....	24
Figure 6 — Examples of the relationship between mapped and unmapped LBAs and physical blocks .....	25
Figure 7 — Armed decreasing threshold operation .....	33
Figure 8 — Armed increasing threshold operation .....	34
Figure 9 — LBP state machine (anchored LBAs supported and deallocated LBAs supported) .....	36
Figure 10 — LBP state machine (anchored LBAs not supported) .....	36
Figure 11 — SSU_PC state machine .....	69
Figure 12 — Referrals .....	97
Figure 13 — Stream Block Relationships .....	110
Figure 14 — Multiple streams example .....	111
Figure B.1 — Update write operation (SCSI storage array device supervised) .....	351
Figure B.2 — Regenerate operation (SCSI storage array device supervised) .....	352
Figure B.3 — Rebuild operation (SCSI storage array device supervised) .....	353
Figure G.1 — Referrals example with no user data segment multiplier .....	373
Figure G.2 — Referrals example with non-zero user data segment multiplier .....	375



## Foreword

This foreword is not part of American National Standard BSR INCITS 514.

This purpose of this standard is to define the model and command set extensions to be used in conjunction with the SCSI Primary Command Set standard – 5 (SPC-5) to facilitate operation of SCSI direct access block devices (e.g., hard disk drives).

Requests for interpretation, suggestions for improvement and addenda, or defect reports are welcome. They should be sent to the INCITS Secretariat, International Committee for Information Technology Standards, Information Technology Institute Council, Suite 610 K Street, NW, Washington, DC 20005.

This standard was processed and approved for submittal to ANSI by the International Committee for Information Technology Standards (INCITS). Committee approval of the standard does not necessarily imply that all committee members voted for approval. At the time it approved this standard, INCITS had the following members:

Karen Higginbottom, Chair

David Michael, Vice-Chair

INCITS Technical Committee T10 on SCSI Storage Interfaces, which developed and reviewed this standard, had the following members:

Ralph Weber, Chair

William Martin, Vice-Chair

John Geldman, Secretary

*Organization Represented*

*Name of Representative*

.....

## Introduction

The standard is organized as follows:

Clause 1 (Scope) describes the relationship of this standard to the SCSI family of standards.

Clause 2 (Normative references) provides references to other standards and documents.

Clause (T10/BSR INCITS 536, Block Commands (ZBC) (planned as ISO/IEC 14776-345)) defines terms and conventions used throughout this standard.

Clause 4 (Direct access block device type model) provides an overview of the direct access block device type.

Clause 5 (Commands for direct access block devices) defines commands specific to direct access block devices.

Clause 6 (Parameters for direct access block devices) defines address descriptors, diagnostic pages, mode parameters and pages, log pages, VPD pages, and copy manager parameters specific to direct access block devices.

Informative Annex A (Numeric order codes) summarizes service action assignments for variable-length commands and commands using the SERVICE ACTION IN operation code and SERVICE ACTION OUT operation code.

Informative Annex B (XOR command examples) provides examples of XOR command usage.

Informative Annex C (CRC example in C) provides example C code for generating the CRC contained in the protection information LOGICAL BLOCK GUARD field.

Informative Annex D (Sense information for locked or encrypted logical units) describes the conditions relative to the sense key and the additional sense code returned by the device server with the CHECK CONDITION status for a SCSI target device that is locked or encrypted.

Informative Annex E (Optimizing block access characteristics) describes an example method that application clients may use to achieve optimal performance for logical block access.

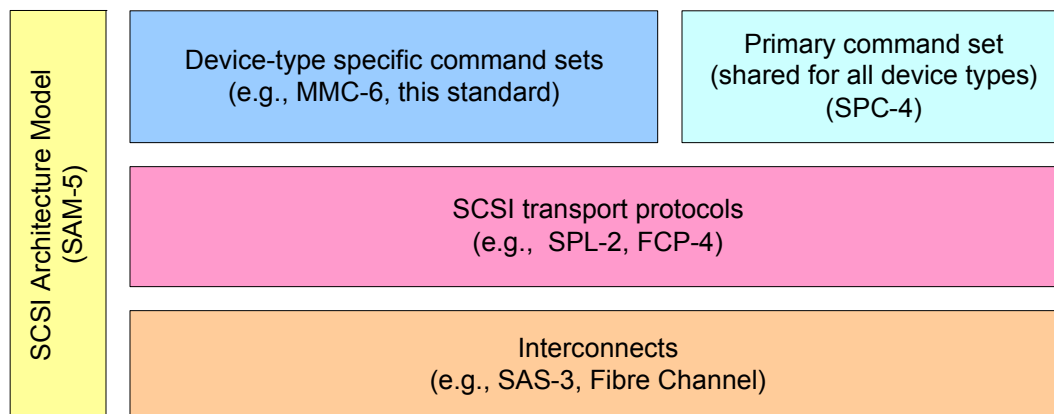
Informative Annex F (Logical block provisioning reporting examples) provides examples of implementations for different methods of reporting logical block provisioning.

Informative Annex G (Discovering referrals examples) provides examples for referrals with no user data segment multiple and referrals with a non-zero user data segment multiplier.

Informative Annex I (Bibliography) provides a list of informative references for this standard.

## SCSI standards family

Figure 0 shows the relationship of this standard to the other standards and related projects in the SCSI family of standards as of the publication of this standard.



**Figure 0 — SCSI document relationships**

Figure 0 is intended to show the general relationship of the documents to one another and is not intended to imply a relationship such as a hierarchy, protocol stack, or system architecture.

The set of SCSI standards specifies the interfaces, functions, and operations necessary to ensure interoperability between conforming SCSI implementations. This standard is a functional description. Conforming implementations may employ any design technique that does not violate interoperability. See SAM-5 for more information about the relationships between the SCSI standards.

This standard makes obsolete the following concepts from SBC-3:

- a) The eer and dcr bit (see 6.5.10);
- b) The TMC and ETC bit in various log parameters;
- c) READ LONG command and WRITE LONG command except for write uncorrectable

**American National Standard  
for Information Technology -****SCSI Block Commands – 4 (SBC-4)****1 Scope**

This standard defines the command set extensions to facilitate operation of SCSI direct access block devices. The clauses in this standard, implemented in conjunction with the applicable clauses of SPC-5, specify the standard command set for SCSI direct access block devices.

The objectives of this standard are to:

- a) permit an application client to communicate over a SCSI service delivery subsystem (see SAM-5) with a logical unit that declares itself to be a direct access block device in the PERIPHERAL DEVICE TYPE field of the standard INQUIRY data (see SPC-5); and
- b) define commands and parameters unique to the direct access block device type.

**2 Normative references**

The following documents, in whole or in part, are normatively referenced in this document and are indispensable for its application. For dated references, only the edition cited applies. For undated references, the latest edition of the referenced document (including any amendments) applies.

ISO/IEC 14776-262, *SAS Protocol Layer-2 (SPL-2)*

ISO/IEC 14776-342, *SCSI-3 Controller Commands - 2 (SCC-2)*

INCITS 468-2010, *Multi-Media Commands - 6 (MMC-6)*

INCITS 468-2010/AM 1 *MultiMedia Command Set - 6 - Amendment 1 (MMC-6/AM 1)*

T10/BSR INCITS 502, *SCSI Primary Commands - 5 (SPC-5)* (planned as ISO/IEC 14776-455)

T10/BSR INCITS 513, *SCSI Primary Commands - 4 (SPC-4)* (planned as ISO/IEC 14776-454)

T10/BSR INCITS 515, *SCSI Architecture Model - 5 (SAM-5)* (planned as ISO/IEC 14776-415)

T10/BSR INCITS 517, *SCSI / ATA Translation - 3 (SAT-3)* (planned as ISO/IEC 14776-923)

T10/BSR INCITS 518, *SCSI Enclosure Services - 3 (SES-3)* (planned as ISO/IEC 14776-373)

T10/BSR INCITS 536, *Block Commands (ZBC)* (planned as ISO/IEC 14776-345)

### 3 Definitions, symbols, abbreviations, keywords, and conventions Introduction

Table 1 shows the topics in clause and a reference to the subclause where each topic is described.

**Table 1 — Direct access block device type mode topics and references**

Topic	Reference
Definitions	3.1
Symbols	3.2
Abbreviations	3.3
Keywords	3.4
Editorial conventions	3.5
Numeric conventions	3.6
State machine conventions	3.7

## 3.1 Definitions

### 3.1.1 additional sense code

combination of the ADDITIONAL SENSE CODE field and the ADDITIONAL SENSE CODE QUALIFIER field in sense data

Note 1 to entry: See SPC-5.

### 3.1.2 advanced background operation

background operation that may impact device server response time to affected LBAs and may include garbage collection operations

### 3.1.3 AND

Boolean arithmetic function (see 3.1.15) on two binary input values that results in an output value of one if both of the input values are one or zero if either of the input values is zero

### 3.1.4 AND operation

performance of an AND (see 3.1.3) bitwise on two multiple-bit input values both having the same number of bits (e.g., on the current content of a logical block and the content contained in the Data-Out Buffer having the same number of bytes)

### 3.1.5 anchored

logical block provisioning state of an LBA (see 4.7.1) in which physical capacity has been reserved for the referenced logical block (see 4.7.4.7)

### 3.1.6 application client

object that is the source of SCSI commands

Note 1 to entry: See SAM-5.

### 3.1.7 atomic command

atomic write command

Note 1 to entry: See 4.31.

### 3.1.8 atomic operation

atomic write operation

**3.1.9 automatic read reassignment**

sequence after the device server detects a recovered read error during which the device server, without intervention from an application client, performs a reassign operation on the LBA for which the error was detected

**3.1.10 atomic write command**

command that performs one or more atomic write operations

Note 1 to entry: See 4.31.

**3.1.11 atomic write operation**

process by which a device server performs a write operation that is either completed in its entirety or has no effects on stored logical block data as described in 4.30

**3.1.12 automatic write reassignment**

sequence after the device server detects a recovered error or an unrecovered error during which the device server, without intervention from an application client, performs a reassign operation on the LBA for which the error was detected

**3.1.13 background function**

either a background scan operation (see 4.24) or a device specific background function (see 3.1.29)

**3.1.14 bitmap buffer**

temporary buffer within a device server (e.g., for one or more bytes of the result of an AND operation (see 3.1.4) or an OR operation (see 3.1.63))

**3.1.15 Boolean arithmetic function**

function that produces an output from one or more inputs according to the rules of AND (see 3.1.3), XOR (see 3.1.33), and OR (see 3.1.62)

**3.1.16 byte**

sequence of eight contiguous bits considered as a unit

**3.1.17 cache**

temporary data storage area that is capable of containing a subset of the logical block data stored by the logical unit and is either volatile or non-volatile

**3.1.18 check data**

information contained within a redundancy group (see 3.1.85) that may allow lost or destroyed XOR-protected data (see 3.1.119) to be recreated

**3.1.19 command**

request describing a unit of work to be performed by a device server

Note 1 to entry: See SAM-5.

**3.1.20 command descriptor block (CDB)**

structure used to communicate commands from an application client to a device server

Note 1 to entry: See SPC-5.

**3.1.21 compare operation**

process by which a device server compares two sets of data for equality

**3.1.22 cyclic redundancy check (CRC)**

error checking mechanism that checks data integrity by computing a polynomial algorithm based checksum (see 4.22.4)

**3.1.23 Data-In Buffer**

buffer specified by the application client to receive data from the device server during the processing of a command

Note 1 to entry: See SAM-5 and SPC-5.

**3.1.24 Data-Out Buffer**

buffer specified by the application client to supply data that is sent from the application client to the device server during the processing of a command

Note 1 to entry: See SAM-5 and SPC-5.

**3.1.25 deallocated**

logical block provisioning state of an LBA (see 4.7.1) in which physical capacity has not been reserved for the referenced logical block (see 4.7.4.6)

**3.1.26 defect list**

GLIST (see 4.13) or PLIST (see 4.13)

**3.1.27 device managed zoned block device**

device server that implements commands as defined in this standard, implements sequential write preferred zones as described in ZBC (see ZBC), and does not implement commands as defined in ZBC

**3.1.28 device server**

object within a logical unit (see 3.1.50) that processes SCSI commands according to the rules of command management

Note 1 to entry: See SAM-5.

**3.1.29 device specific background functions**

SCSI target device specific functions that a SCSI target device may perform that have no specific association with application client-initiated operations

Note 1 to entry: See SPC-5.

**3.1.30 device type**

device model

implemented by the logical unit and indicated to the application client by the PERIPHERAL DEVICE TYPE field of the standard INQUIRY data (see SPC-5)

**3.1.31 direct access block device**

device that is capable of containing data stored in logical blocks that each have a unique LBA (see 4.2)

**3.1.32 error correcting code (ECC)**

error checking mechanism that checks data integrity and enables some errors in the logical block data to be corrected

**3.1.33 exclusive-or (XOR)**

boolean arithmetic function on two binary input values that results in an output value of one if one and only one of the input values is one, or zero if both of the input values are either zero or one

**3.1.34 extent**

set of logical blocks occupying contiguous LBAs on a logical unit

**3.1.35 field**

group of one or more contiguous bits, a part of a larger structure (e.g., a CDB (see 3.1.20) or sense data (see SPC-5))

**3.1.36 format corrupt**

vendor specific condition in which the device server may not be able to perform logical block access commands-

Note 1 to entry: ~~(S~~See 4.10 [and 4.35.](#)~~)~~

**3.1.37 format operation**

process by which a device server initializes the medium in a logical unit

Note 1 to entry: See 4.10 [and 4.35.](#)

**3.1.38 fully provisioned logical unit**

logical unit that stores logical block data for every LBA and has assigned physical capacity for every LBA

Note 1 to entry: See 4.7.2.

**3.1.39 garbage collection operation**

process that prepares resources for future allocation to LBAs

**3.1.40 grown defect list (GLIST)**

list of physical blocks that the device server has detected as containing medium defects or that the application client has specified as containing medium defects

Note 1 to entry: See 4.13.

**3.1.41 hard reset**

condition resulting from the events defined by SAM-5 during which the SCSI device performs the hard reset operations described in SAM-5, this standard, and other applicable command standards (see table 34)

**3.1.42 I\_T nexus**

relationship between a SCSI initiator port and a SCSI target port

Note 1 to entry: See SAM-5.

**3.1.43 I\_T nexus loss**

condition resulting from the events defined by SAM-5 during which the SCSI device performs the I\_T nexus loss operations described in SAM-5, this standard, and other applicable command standards (see table 34)

**3.1.44 LBA mapping resource**

resource (e.g., a physical block or a data structure associated with tracking resource usage) used by a logical unit that supports logical block provisioning management

**3.1.45 logical block**

set of data bytes accessed and referenced as a unit (see 4.5)

**3.1.46 logical block access command**

command that requests access to one or more logical blocks that may require access to the medium

Note 1 to entry: See 4.2.2.

**3.1.47 logical block address (LBA)**

value used to reference a logical block (see 4.5)

**3.1.48 logical block data**

user data and protection information, if any

**3.1.49 logical block length**

number of bytes of user data in a logical block (see 4.5)

**3.1.50 logical unit**

externally addressable entity within a SCSI target device (see 3.1.88) that implements a SCSI device model

Note 1 to entry: See SAM-5.

**3.1.51 logical unit reset**

condition resulting from the events defined by SAM-5 in which the logical unit performs the logical unit reset operations described in SAM-5, this standard, and other applicable command standards (see table 34)

**3.1.52 mapped**

logical block provisioning state of an LBA (see 4.7.1) in which physical capacity has been assigned to the referenced logical block (see 4.7.4.5)

**3.1.53 media**

plural of medium

**3.1.54 medium**

material that is not cache on which data is stored (e.g., a magnetic disk)

**3.1.55 medium defect**

area of the medium that results in a recovered error or an unrecovered error when a read medium operation or a write medium operation is performed

Note 1 to entry: See 4.13.

**3.1.56 [misaligned write command](#)**

[write command with fields set as described in 4.6.2.](#)

**3.1.57 non-advanced background operation**

background operation that does not impact device server response time to affected LBAs and may include garbage collection operations

**3.1.58 non-atomic command**

command that is not an atomic command

**3.1.59 non-atomic operation**

operation that is not an atomic operation

**3.1.60 non-volatile cache**

cache that retains logical block data through any power cycle

**3.1.61 non-volatile medium**

medium that retains logical block data through any power cycle

**3.1.62 OR**

Boolean arithmetic function (see 3.1.15) on two binary input values that results in an output value of one if either of the input values are one or zero if both of the input values are zero

**3.1.63 OR operation**

performance of an OR (see 3.1.62) bitwise on two multiple-bit input values both having the same number of bits (e.g., on the current content of a logical block and the content contained in the Data-Out Buffer having the same number of bytes)

**3.1.64 point in time ROD token**

ROD token with a ROD type that is a point in time copy ROD

Note 1 to entry: See SPC-5.



**3.1.65 physical block**

set of data bytes accessed as a unit by the device server (see 4.6)

**3.1.66 physical block length**

number of bytes of logical block data in a physical block (see 4.6)

**3.1.67 physical element**

component that provides non-volatile storage for an associated group of logical blocks (see 4.20)

**3.1.68 power cycle**

sequence of power being removed followed by power being applied to a SCSI device

**3.1.69 power on**

condition resulting from the events defined by SAM-5 during which a SCSI device performs the power on operations described in SAM-5, this standard, and other applicable command standards (see table 34)

**3.1.70 primary defect list (PLIST)**

list of physical blocks containing medium defects that are considered permanent

Note 1 to entry: See 4.13.

**3.1.71 protection information**

group of fields at the end of each logical block or at specified intervals within each logical block that contain a logical block guard, an application tag, and a reference tag

Note 1 to entry: See 4.22.

**3.1.72 protection information interval**

length of user data that occurs within a logical block before each protection information

**3.1.73 provisioning initialization pattern**

non-zero pattern that is the length of one logical block

**3.1.74 pseudo read data**

indeterminate logical block data

**3.1.75 pseudo unrecovered error**

simulated error (e.g., created by a WRITE LONG command (see 5.45 and 5.46)) for which a device server reports that it is unable to read or write a logical block, regardless of whether the data on the medium is valid, recoverable, or unrecoverable

**3.1.76 pseudo unrecovered error with correction disabled**

pseudo unrecovered error for which a device server performs no error recovery

Note 1 to entry: See 4.18.2.

**3.1.77 read cache operation**

process by which a device server reads logical blocks for one or more LBAs from cache as described in 4.15

**3.1.78 read command**

command that requests read operations

Note 1 to entry: See 4.2.2.

**3.1.79 read medium operation**

process by which a device server reads logical blocks for one or more LBAs from the medium using the parameters specified in the Read-Write Error Recovery mode page (see 6.5.10)

**3.1.80 read operation**

process by which a device server performs operations (e.g., read cache operations and read medium operations) as described in this standard

Note 1 to entry: See 4.2.3.

**3.1.81 reassign**

perform a reassign operation

**3.1.82 reassign operation**

operation during which the device server changes the assignment of an LBA from a specified physical block to another physical block and adds the specified physical block to the GLIST

**3.1.83 recovered error**

error for which a device server is able to read or write a logical block within the recovery limits specified in the Read-Write Error Recovery mode page (see 6.5.10) or the Verify Error Recovery mode page (see 6.5.11)

**3.1.84 recovered read error**

recovered error that occurs during a read medium operation

**3.1.85 redundancy group**

grouping of XOR-protected data (see 3.1.119) and associated check data (see 3.1.18) into a single type of data redundancy (see SCC-2)

**3.1.86 resource provisioned logical unit**

logical unit that may or may not store logical block data for every LBA and that provides enough LBA mapping resources (see 3.1.44) to map every LBA

Note 1 to entry: See 4.7.3.2.

**3.1.87 sanitize operation**

process by which a device server alters information on a logical unit such that recovery of previous logical block data from the cache and the medium is not possible

Note 1 to entry: See 4.11.

**3.1.88 SCSI target device**

SCSI device containing logical units and SCSI target ports that receives device service requests and task management requests for processing and sends device service responses and task management responses to SCSI initiator devices

Note 1 to entry: See SAM-5

**3.1.89 sense data** (see SPC-5)

data describing command completed information that a device server delivers to an application client in the same I\_T nexus transaction as the status or as parameter data in response to a REQUEST SENSE command

**3.1.90 sense key**

contents of the SENSE KEY field in the sense data

Note 1 to entry: See SAM-5

**3.1.91 status**

one byte of response information that contains a coded value defined in SAM-5, transferred from a device server to an application client upon completion of each command

Note 1 to entry: See SAM-5

**3.1.92 stopped power condition**

power condition in which a device server terminates TEST UNIT READY commands and logical block access commands

Note 1 to entry: See 4.21.4.

**3.1.93 stream block**

stream granularity size (see 6.6.5) area of physical media that is able to contain logical block data

**3.1.94 stream identifier**

identifier supplied by the device server and used by the application client to identify a stream

Note 1 to entry: See 4.34.

**3.1.95 synchronize cache operation**

process by which a device server synchronizes logical blocks within the volatile cache with the non-volatile cache or the medium

**3.1.96 thin provisioned logical unit**

logical unit that may or may not store logical block data for every LBA and that may or may not provide enough LBA mapping resources (see 3.1.44) to map every LBA

Note 1 to entry: See 4.7.3.3.

**3.1.97 threshold set**

set of two or more logical blocks used for tracking logical block provisioning thresholds

Note 1 to entry: See 4.7.3.7.

**3.1.98 threshold set size**

number of LBAs in a threshold set

Note 1 to entry: See 4.7.3.7.

**3.1.99 token**

representation of a collection of data

Note 1 to entry: See SPC-5.

**3.1.100 unit attention condition**

state that a logical unit (see 3.1.50) maintains while the logical unit has asynchronous status information to report to the SCSI initiator ports associated with one or more I\_T nexuses (see 3.1.42)

Note 1 to entry: See SAM-5.

**3.1.101 unmap command**

command that requests an unmap operation

Note 1 to entry: See 4.2.2.

**3.1.102 unmap operation**

process by which a device server either deallocates or anchors a single LBA

Note 1 to entry: See 4.2.3 and 4.7.3.4.

**3.1.103 unmapped**

logical block provisioning state of an LBA (see 4.7.1) in which the LBA is either anchored or deallocated

**3.1.104 unrecovered error**

error for which a device server is unable to read a logical block or write a logical block within the recovery limits specified in the Read-Write Error Recovery mode page (see 6.5.10) and/or the Verify Error Recovery mode page (see 6.5.11)

**3.1.105 unrecovered read error**

unrecovered error that occurs during a read medium operation

**3.1.106 unrecovered write error**

unrecovered error that occurs during a write medium operation

**3.1.107 user data**

data contained in logical blocks that is accessible by an application client and is neither protection information nor other information that may not be accessible to the application client

**3.1.108 user data segment**

contiguous sequence of logical blocks

Note 1 to entry: See 4.28.

**3.1.109 verify command**

command that requests verify operations

Note 1 to entry: See 4.2.2.

**3.1.110 verify medium operation**

process by which a device server reads logical blocks for one or more LBAs from the medium using the parameters specified in the Verify Error Recovery mode page (see 6.5.11)

**3.1.111 verify operation**

process by which the device server performs operations (e.g., verify medium operations) as described in this standard

Note 1 to entry: See 4.2.3.

**3.1.112 volatile cache**

cache that does not retain logical block data between power cycles

**3.1.113 volatile medium**

medium that does not retain logical block data between power cycles (e.g., a silicon memory device that loses data written to it if device power is lost)

**3.1.114 write cache operation**

process by which a device server writes logical blocks for one or more LBAs to the cache (see 4.7.1)

**3.1.115 write command**

command that requests write operations

Note 1 to entry: See 4.2.2.

**3.1.116 write medium operation**

process by which a device server writes logical blocks for one or more LBAs to the medium using the parameters specified in the Read-Write Error Recovery mode page (see 6.5.10)

**3.1.117 write operation**

process by which a device server performs operations (e.g., write cache operations and write medium operations) as described in this standard

Note 1 to entry: See 4.2.3.

**3.1.118 XOR operation**

processing of an XOR bitwise on two identical-sized multiple-bit input values (e.g., the current value of a logical block and the new value for that logical block)

**3.1.119 XOR-protected data**

logical blocks (i.e., including logical block data) that are part of a redundancy group

**3.1.120 zoned block device**

host aware zoned block device (see ZBC) or host managed zoned block device (see ZBC)

**3.2 Symbols**

Symbols used in this standard include:

<b>Symbol</b>	<b>Meaning</b>
+	plus
–	minus
×	multiplied by
÷	divided by
=	equals
≠	not equal
<	less than
>	greater than

**3.3 Abbreviations**

Abbreviations used in this standard include:

<b>Abbreviation</b>	<b>Meaning</b>
CbCS	Capability based Command Security
CDB	command descriptor block (see 3.1.20)
CRC	cyclic redundancy check (see 3.1.22)
ECC	error correcting code (see 3.1.32)
GLIST	grown defect list (see 3.1.40)
LBA	logical block address (see 3.1.47)
LBM	Logical Block Markup (see 6.8)
LBP	logical block provisioning (see 4.7.4)
LSB	least significant bit
LUN	logical unit number
M	implementation is mandatory
MMC-6	SCSI Multimedia Commands - 6 (see 3.1)
MSB	most significant bit
O	implementation is optional
PLIST	primary defect list (see 3.1.70)
n/a	not applicable
ROD	representation of data (see SPC-5)
SAM-5	SCSI Architecture Model - 5 (see 3.2)
SCSI	Small Computer System Interface family of standards
SCC-2	SCSI-3 Controller Commands - 2 (see 3.1)
SES-3	SCSI Enclosure Services - 3 (see 3.2)

SPC-4	SCSI Primary Commands - 4 (see 3.2)
SPC-5	SCSI Primary Commands - 5 (see 3.2)
SPL-2	SAS Protocol Layer-2 (see 3.1)
VPD	Vital product data (see 6.6)
XOR	exclusive-or (see 3.1.33)
ZBC	Zoned Block Commands

## 3.4 Keywords

### 3.4.1 invalid

keyword used to describe an illegal or unsupported bit, byte, word, field or code value

Note 1 to entry: Receipt of an invalid bit, byte, word, field or code value shall be reported as an error.

### 3.4.2 mandatory

keyword indicating an item that is required to be implemented as defined in this standard

### 3.4.3 may

keyword that indicates flexibility of choice with no implied preference

Note 1 to entry: "May" is equivalent to "may or may not".

### 3.4.4 may not

keywords that indicate flexibility of choice with no implied preference

Note 1 to entry: "May not" is equivalent to "may or may not".

### 3.4.5 obsolete

keyword indicating that an item was defined in prior SCSI standards but has been removed from this standard

### 3.4.6 optional

keyword that describes features that are not required to be implemented by this standard; however, if any optional feature defined in this standard is implemented, then it shall be implemented as defined in this standard

### 3.4.7 prohibited

keyword used to describe a feature, function, or coded value that is defined in a non-SCSI standard (i.e., a standard that is not a member of the SCSI family of standards) to which this standard makes a normative reference where the use of said feature, function, or coded value is not allowed for implementations of this standard

### 3.4.8 reserved

keyword referring to bits, bytes, words, fields and code values that are set aside for future standardization

Note 1 to entry: A reserved bit, byte, word or field shall be set to zero, or in accordance with a future extension to this standard.

Note 2 to entry: Recipients are not required to check reserved bits, bytes, words or fields for zero values.

Note 3 to entry: Receipt of reserved code values in defined fields shall be reported as an error.

### 3.4.9 restricted

keyword referring to bits, bytes, words, and fields that are set aside for other identified standardization purposes

Note 1 to entry: A restricted bit, byte, word, or field shall be treated as a reserved bit, byte, word or field in the context where the restricted designation appears.

**3.4.10 shall**

keyword indicating a mandatory requirement

Note 1 to entry: Designers are required to implement all such mandatory requirements to ensure interoperability with other products that conform to this standard.

**3.4.11 should**

keyword indicating flexibility of choice with a strongly preferred alternative

Note 1 to entry: "Should" is equivalent to the phrase "it is strongly recommended".

**3.4.12 vendor specific**

something (e.g., a bit, field, code value) that is not defined by this standard

Note 1 to entry: Specification of the referenced item is determined by the SCSI device vendor and may be used differently in various implementations.

**3.5 Editorial conventions**

Certain words and terms used in this standard have a specific meaning beyond the normal English meaning. These words and terms are defined either in this clause or in the text where they first appear.

Normal case is used for words having the normal English meaning.

Upper case is used when referring to names of commands, status codes, sense keys, and additional sense codes (e.g., REQUEST SENSE command).

If there is more than one CDB length for a particular command (e.g., ORWRITE (16) and ORWRITE (32)), and the name of the command is used in a sentence without any CDB length descriptor (e.g., ORWRITE), then the condition described in the sentence applies to all CDB lengths for that command.

Names of fields and state variables are in small uppercase (e.g. NAME). When a field or state variable name contains acronyms, uppercase letters may also be used for readability (e.g., the LOGERR bit). Normal case is used when the contents of a field or state variable are being discussed. Fields or state variables containing only one bit are usually referred to as the NAME bit instead of the NAME field.

Lists sequenced by lowercase or uppercase letters show no ordering relationship between the listed items.

EXAMPLE 1 - The following list shows no relationship between the listed items:

- a) red (i.e., one of the following colors):
  - A) crimson; or
  - B) amber;
- b) blue; or
- c) green.

Lists sequenced by numbers show an ordering relationship between the listed items.

EXAMPLE 2 -The following list shows an ordered relationship between the named items:

- 1) top;
- 2) middle; and
- 3) bottom.

Lists are associated with an introductory paragraph or phrase, and are numbered relative to that paragraph or phrase (i.e., all lists begin with an a) or 1) entry).

In the event of conflicting information the precedence for requirements defined in this standard is:

- 1) text;
- 2) tables; then
- 3) figures.

Tables show data format and values. Not all tables or figures are fully described in the text.

Notes and examples do not constitute any requirements for implementers, and notes are numbered consecutively throughout this standard.

## 3.6 Numeric and character conventions

### 3.6.1 Numeric conventions

A binary number is represented in this standard by any sequence of digits comprised of only the Western-Arabic numerals 0 and 1 immediately followed by a lower-case b (e.g., 0101b). Underscores are included between characters in binary number representations to increase readability or delineate field boundaries (e.g., 0\_0101\_1010b).

A hexadecimal number is represented in this standard by any sequence of digits comprised of only the Western-Arabic numerals 0 to 9 and/or the upper-case English letters A to F immediately followed by a lower-case h (e.g., FA23h). Underscores are included in hexadecimal number representations to increase readability or delineate field boundaries (e.g., B\_FD8C\_FA23h).

A decimal number is represented in this standard by any sequence of digits comprised of only the Western-Arabic numerals 0 to 9 not immediately followed by a lower-case b or lower-case h (e.g., 25).

A range of numeric values is represented in this standard in the form “a to z”, where a is the first value included in the range, all values between a and z are included in the range, and z is the last value included in the range (e.g., the representation “0h to 3h” includes the values 0h, 1h, 2h, and 3h).

This standard uses the following conventions for representing decimal numbers:

- a) the decimal separator (i.e., separating the integer and fractional portions of the number) is a period;
- b) the thousands separator (i.e., separating groups of three digits in a portion of the number) is a space;  
and
- c) the thousands separator is used in both the integer portion and the fraction portion of a number.

Table 2 shows some examples of decimal numbers represented using various conventions.

**Table 2 — Numbering convention examples**

French	English	This standard
0,6	0.6	0.6
3,141 592 65	3.14159265	3.141 592 65
1 000	1,000	1 000
1 323 462,95	1,323,462.95	1 323 462.95

A decimal number represented in this standard with an overline over one or more digits following the decimal point is a number where the overlined digits are infinitely repeating (e.g., 666. $\overline{6}$  means 666.666 666... or  $666 \frac{2}{3}$ , and 12.142 857 $\overline{}$  means 12.142 857 142 857... or  $12 \frac{1}{7}$ ).

### 3.6.2 Units of measure

This standard represents values using both decimal units of measure and binary units of measure. Values are represented by the following formats:

- a) for values based on decimal units of measure:
  - 1) numerical value (e.g., 100);
  - 2) space;
  - 3) prefix symbol and unit:
    - 1) decimal prefix symbol (e.g., M) (see table 3); and
    - 2) unit abbreviation (e.g., B);

and
- b) for values based on binary units of measure:



- 1) numerical value (e.g., 1 024);
- 2) space;
- 3) prefix symbol and unit:
  - 1) binary prefix symbol (e.g., Gi) (see table 3); and
  - 2) unit abbreviation (e.g., b).

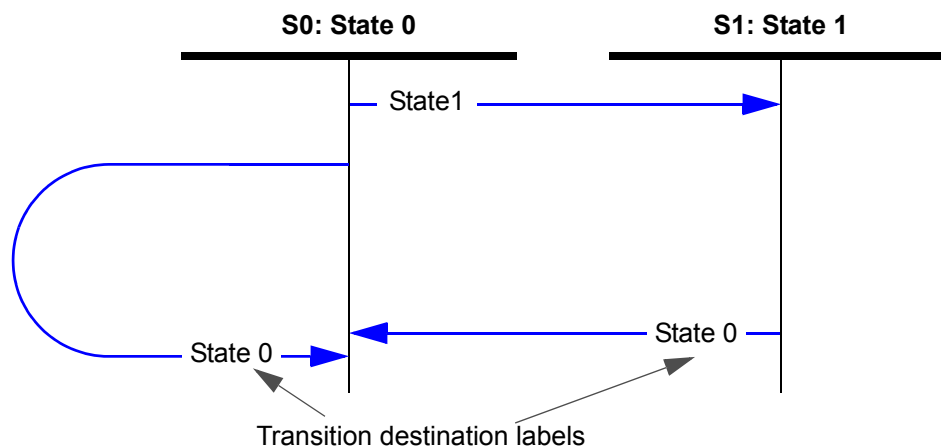
Table 3 compares the prefix, symbols, and power of the binary and decimal units.

**Table 3 — Comparison of decimal prefixes and binary prefixes**

Decimal			Binary		
Prefix name	Prefix symbol	Power (base-10)	Prefix name	Prefix symbol	Power (base-2)
kilo	k	$10^3$	kibi	Ki	$2^{10}$
mega	M	$10^6$	mebi	Mi	$2^{20}$
giga	G	$10^9$	gibi	Gi	$2^{30}$
tera	T	$10^{12}$	tebi	Ti	$2^{40}$
peta	P	$10^{15}$	pebi	Pi	$2^{50}$
exa	E	$10^{18}$	exbi	Ei	$2^{60}$
zetta	Z	$10^{21}$	zebi	Zi	$2^{70}$
yotta	Y	$10^{24}$	yobi	Yi	$2^{80}$

### 3.7 State machine conventions

Figure 1 shows how state machines are described in this standard.



**Figure 1 — Example state machine figure**

The state machine figure is followed by subclauses describing the states and state transitions.

Each state and state transition is described in the list with particular attention to the conditions that cause the transition to occur and special conditions related to the transition.

A system specified in this manner has the following properties:

- a) time elapses only within discrete states; and
- b) state transitions are logically instantaneous.

## 4 Direct access block device type model

### 4.1 Direct access block device type model introduction

Table 4 shows the topics in clause 4 and a reference to the subclause where each topic is described.

**Table 4 — Direct access block device type model topics**

Topic	Reference
Direct access block device type model overview	4.2
Media examples	4.3
Removable media	4.4
Logical blocks	4.5
Physical blocks	4.6
Logical block provisioning	4.7
Data de-duplication	4.8
Ready state	4.9
Initialization	4.10
Sanitize operations	4.11
Write protection	4.12
Medium defects	4.13
Write and unmap failures	4.14
Caches	4.15
Implicit HEAD OF QUEUE command processing	4.16
Reservations	4.17
Error reporting	4.18
Model for XOR commands	4.19
Rebuild assist mode	4.20
START STOP UNIT and power conditions	4.21
Protection information model	4.22
Grouping function	4.23
Background scan operations	4.24
Association between commands and CbCS permission bits	4.25

**Table 4 — Direct access block device type model topics**

<b>Topic</b>	<b>Reference</b>
Deferred microcode activation	4.26
Model for uninterrupted sequences on LBA ranges	4.27
Referrals	4.28
ORWRITE commands	4.29
Block device ROD token operations	4.30
Atomic Writes	4.31
IO Advice Hints	4.32
Background Operation Control	4.33
Stream Control	4.34
Format Operations	4.35

## **4.2 Direct access block device type model**

### **4.2.1 Direct access block device type model overview**

SCSI devices that conform to this standard are referred to as direct access block devices (e.g., hard disk drives, removable rigid disks, and solid state drives).

This standard is intended to be used in conjunction with SAM-5, SPC-5, SCC-2, and SES-3.

Direct access block devices store data in logical blocks for later retrieval.

Logical blocks are stored by a process that causes localized changes or transitions within a medium. The changes made to the medium to store the logical blocks may be volatile (i.e., not retained through power cycles) or non-volatile (i.e., retained through power cycles).

### **4.2.2 Logical block access command types**

The following are logical block access command types:

- a) read commands;
- b) unmap commands;
- c) verify commands;
- d) write commands; and
- e) other commands (e.g., a FORMAT UNIT command or a SANITIZE command).

See table 34 for a list of commands for direct access block devices, including the logical block access command type. Some commands may be more than one type of logical block access command (e.g., a COMPARE AND WRITE command is both a read command and a write command).

### **4.2.3 Logical block access operation types**

Each named command type (see 4.2.2) is processed by performing one or more of the following:

- a) read operations;
- b) unmap operations;
- c) verify operations; and
- d) write operations.

A device server that supports optional features (e.g., caches (see 4.15)) may be required to support additional requirements (e.g., cache coherency, LBA mapping resource allocations) that are related to those features for specific operations (e.g., read operations, write operations).

In a device server that does not support any optional features:

- a) any read operation causes only read medium operations to be performed;
- b) any verify operation causes only verify medium operations to be performed; and
- c) any write operation causes only write medium operations to be performed.

The requirements for any optional feature (e.g., caches) may include additional requirements for the operations described in this subclause.

If an optional feature (e.g., caches) defines requirements for read operations, then the device server shall support those requirements for both verify operations and read operations.

## 4.3 Media examples

### 4.3.1 Media examples overview

Examples of types of media used by the direct access block device are:

- a) a rotating medium (see 4.3.2); and
- b) a memory medium (see 4.3.3).

Other types of media are possible.

### 4.3.2 Rotating media

A rotating medium is one or more spinning disks, each coated with a magnetic material that allows flux changes to be induced and recorded. An actuator positions a read-write head radially across the spinning disk, allowing the device to randomly read or write the information at any radial position. Data is stored by using the write portion of the head to record flux changes and the recorded data is read by using the read portion of the head.

The circular path followed by the read-write head at a particular radius is called a track. A track is divided into sectors each containing blocks of stored data. If there is more than one disk spinning on a single axis and the actuator has a read-write head to access each of the disk surfaces, then the collection of tracks at a particular radius is called a cylinder.

A logical block is stored in one or more sectors, or a sector may store more than one logical block. Sectors may also contain information for accessing, synchronizing, and protecting the integrity of the logical blocks.

A rotating medium direct access block device is ready if:

- a) the disks are rotating at the correct speed; and
- b) the read-write circuitry is powered and ready to access the data.

A START STOP UNIT command (see 5.26) may be required to bring the logical unit to the ready state.

The rotating medium in a direct access block device is non-volatile.

### 4.3.3 Memory media

A memory medium is solid state, random access memory (RAM) (e.g., static RAM (SRAM), dynamic RAM (DRAM), magnetoresistive RAM (MRAM), ferroelectric RAM (FeRAM), or flash memory).

A memory medium direct access block device may be ready after power on and may not require a START STOP UNIT command (see 5.26) to bring the logical unit to a ready state.

These logical units may be nonmechanical, and logical blocks may be accessed with similar access times regardless of their location on the medium. Memory medium direct access block devices may store less data than disks or tapes and may be volatile.

A memory medium may be volatile (e.g., SRAM or DRAM) or non-volatile (e.g., SRAM or DRAM with battery backup, MRAM, FeRAM, or flash memory).

#### 4.4 Removable media

The medium may be removable or non-removable. A removable medium may be contained within a cartridge or jacket to prevent damage to the recording surfaces.

A removable medium has an attribute of being mounted or demounted on a suitable transport mechanism in a direct access block device. A removable medium is mounted when the direct access block device is capable of accessing its medium (e.g., performing read medium operations and write medium operations). A removable medium is demounted at any other time (e.g., during loading, unloading, or storage).

An application client may check whether a removable medium is mounted by sending a TEST UNIT READY command (see SPC-5). A direct access block device containing a removable medium may not be accessible for read operations, unmap operations, and write operations until it receives a START STOP UNIT command with the START bit set to one (see 5.26).

If a direct access block device implements cache, either volatile or non-volatile, then the device server ensures that all logical blocks on the medium contain the most recent logical block data prior to permitting demounting of the removable medium.

If the medium in a direct access block device is removable, and the medium is removed, then the device server shall establish a unit attention condition with the additional sense code set to the appropriate value (e.g., MEDIUM NOT PRESENT). The PREVENT ALLOW MEDIUM REMOVAL command (see 5.12) allows an application client to restrict the demounting of the removable medium.

If an application client sends a START STOP UNIT command to request that the removable medium to be ejected and the direct access block device is prevented from demounting the medium by a previous PREVENT ALLOW MEDIUM REMOVAL command, then the START STOP UNIT command is terminated by the device server.

#### 4.5 Logical blocks

Logical blocks are stored on the medium. Logical blocks:

- a) contain logical block data that contains:
  - A) user data; and
  - B) protection information, if any;and
- b) may contain additional information (e.g., an ECC which may be used for medium defect management (see 4.13)), which may not be accessible to the application client.

The number of bytes of user data contained in each logical block is the logical block length. The logical block length is greater than or equal to one byte and should be an even number of bytes (e.g., 512 bytes, 520 bytes, 4 096 bytes, or 4 104 bytes). The logical block length does not include the length of protection information, if any, and additional information, if any, that are contained in the logical block. The logical block length is the same for all logical blocks in the logical unit. The LOGICAL BLOCK LENGTH field (see 5.17.2) and the READ CAPACITY (16) parameter data (see 5.18.2) indicates the logical block length. The FORMAT UNIT command (see 5.4) and the mode parameter block descriptor (see 6.5.2) are used together by an application client to change the logical block length in direct access block devices that support changeable logical block lengths.

Each logical block is referenced by a unique LBA, which is either four bytes in length or eight bytes in length. The LBAs on a logical unit shall begin with zero and shall be contiguous up to the last LBA on the logical unit. The last LBA is  $[n-1]$ , where  $[n]$  is the number of logical blocks accessible by an application client.

For this standard, the RETURNED LOGICAL BLOCK ADDRESS field in (see 5.17.2) and the READ CAPACITY (16) parameter data (see 5.18.2) indicates the value of  $[n-1]$ .

Other command standards (e.g., ZBC) may define the contents of the RETURNED LOGICAL BLOCK ADDRESS field to be less than  $[n-1]$ .

Each LBA has a logical block provisioning state (see 4.7) of mapped, deallocated, or anchored.

Some commands support only four-byte LOGICAL BLOCK ADDRESS fields (e.g., READ (10), and WRITE (10)).

If the capacity of the logical unit exceeds that accessible with four-byte LBAs, then the device server returns the RETURNED LOGICAL BLOCK ADDRESS field set to FFFF\_FFFFh in the READ CAPACITY (10) parameter data, indicating that an application client should:

- a) enable descriptor format sense data (see SPC-5) in the Control mode page (see SPC-5) and in any REQUEST SENSE commands (see SPC-5) it sends; and
- b) use commands with eight-byte LOGICAL BLOCK ADDRESS fields (e.g., READ (16), and WRITE (16)).

NOTE 1 - If a command requests access to an LBA greater than FFFF\_FFFFh, fixed format sense data is used, and an error occurs for that LBA, then there is no field in the sense data large enough to report that LBA as having an error (see 4.18).

If a command is received that references or attempts to access a logical block that exceeds the capacity of the medium, then the device server shall terminate the command (e.g., with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to LOGICAL BLOCK ADDRESS OUT OF RANGE). The device server:

- a) should terminate the command before processing; and
- b) may terminate the command after the device server has transferred some, all, or none of the data.

The location of a logical block on the medium is not required to have a relationship to the location of any other logical block. However, in a direct access block device with rotating media (see 4.3.2), the time to access a logical block at LBA  $[x+1]$  after accessing LBA  $[x]$  is often less than the time to access some other logical block.

## 4.6 Physical blocks

### 4.6.1 [Overview](#)

A physical block is a set of data bytes on the medium accessed by the device server as a unit. A physical block may contain:

- a) a portion of a logical block (i.e., there are multiple physical blocks in the logical block)(e.g., a physical block length of 512 bytes with a logical block length of 2 048 bytes);
- b) a single complete logical block; or
- c) more than one logical block (i.e., there are multiple logical blocks in the physical block)(e.g., a physical block length of 4 096 bytes with a logical block length of 512 bytes).

Each physical block may include additional information (e.g., an ECC which may be used for medium defect management (see 4.13)), which may not be accessible to the application client.

If the device server supports the creation of pseudo unrecovered errors (see 4.18.2), then the device server shall have the capability of marking individual logical blocks as containing pseudo unrecovered errors.

Logical blocks may or may not be aligned to physical block boundaries. A mechanism for establishing the alignment is not defined by this standard.

Figure 2 shows examples of where there are one or more physical blocks per logical block, and LBA 0 is aligned to a physical block boundary. The LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT field and the LOWEST ALIGNED LOGICAL BLOCK ADDRESS field (see 5.18.2) indicate the alignment.

The LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT field is set to 0h (i.e., indicating one or more physical blocks per logical block).

The LOWEST ALIGNED LOGICAL BLOCK ADDRESS field is set to 0000h (i.e., indicating that LBA 0 is located at the beginning of a physical block).

---

4 physical blocks per logical block:

LBA 0				LBA 1				...
PB	PB	PB	PB	PB	PB	PB	PB	...

---

3 physical blocks per logical block:

LBA 0			LBA 1			LBA 2			...
PB	PB	PB	PB	PB	PB	PB	PB	PB	...

---

2 physical blocks per logical block:

LBA 0		LBA 1		LBA 2		LBA 3		LBA 4		...
PB	PB	PB	PB	PB	PB	PB	PB	PB	PB	...

---

1 physical block per logical block:

LBA 0	LBA 1	LBA 2	LBA 3	LBA 4	LBA 5	LBA 6	LBA 7	LBA 8	LBA 9	LBA 10	...
PB	PB	PB	PB	PB	PB	PB	PB	PB	PB	PB	...

**Key:**

LBA n = logical block with LBA n

PB = physical block

**Figure 2 — One or more physical blocks per logical block examples**



Figure 3 shows examples of where there are one or more logical blocks per physical block, and LBA 0 is aligned to a physical block boundary. The LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT field and the LOWEST ALIGNED LOGICAL BLOCK ADDRESS field (see 5.18.2) indicate the alignment.

The LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT field is set to a non-zero value (i.e., indicating more than one logical block per physical block).

The LOWEST ALIGNED LOGICAL BLOCK ADDRESS field is set to 0000h (i.e., indicating that LBA 0 is located at the beginning of a physical block).

LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT field set to 1h (i.e.,  $2^1$  logical blocks per physical block):

LBA 0	LBA 1	LBA 2	LBA 3	LBA 4	LBA 5	LBA 6	LBA 7	LBA 8	LBA 9	LBA 10	LBA 11	...
PB		PB		PB		PB		PB		PB		...

LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT field set to 2h (i.e.,  $2^2$  logical blocks per physical block):

LBA 0	LBA 1	LBA 2	LBA 3	LBA 4	LBA 5	LBA 6	LBA 7	LBA 8	LBA 9	LBA 10	LBA 11	...
PB				PB				PB				...

LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT field set to 3h (i.e.,  $2^3$  logical blocks per physical block):

LBA 0	LBA 1	LBA 2	LBA 3	LBA 4	LBA 5	LBA 6	LBA 7	...
PB								...

**Key:**

LBA n = logical block with LBA n

PB = physical block

**Figure 3 — One or more logical blocks per physical block examples**

Figure 4 shows examples of where there are two logical blocks per physical block, and different LBAs are aligned to physical block boundaries. The LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT field and the LOWEST ALIGNED LOGICAL BLOCK ADDRESS field (see 5.18.2) indicate the alignment.

LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT field set to 1h (i.e.,  $2^1$  logical blocks per physical block):

LOWEST ALIGNED LOGICAL BLOCK ADDRESS field set to 0000h:

LBA 0	LBA 1	LBA 2	LBA 3	LBA 4	LBA 5	LBA 6	LBA 7	LBA 8	LBA 9	...
PB		PB		PB		PB		PB		...

LOWEST ALIGNED LOGICAL BLOCK ADDRESS field set to 0001h:

	LBA 0	LBA 1	LBA 2	LBA 3	LBA 4	LBA 5	LBA 6	LBA 7	LBA 8	LBA 9	LBA 10	...
PB		PB		PB		PB		PB		PB		...

**Key:**

LBA n = logical block with LBA n

PB = physical block

**Figure 4 — Two logical blocks per physical block alignment examples**

Figure 5 shows examples of where there are four logical blocks per physical block, and different LBAs are aligned to physical block boundaries. The LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT field and the LOWEST ALIGNED LOGICAL BLOCK ADDRESS field (see 5.18.2) indicate the alignment.

LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT field set to 2h (i.e.,  $2^2$  logical blocks per physical block):

LOWEST ALIGNED LOGICAL BLOCK ADDRESS field set to 0000h:

LBA 0	LBA 1	LBA 2	LBA 3	LBA 4	LBA 5	LBA 6	LBA 7	...
PB				PB				...

LOWEST ALIGNED LOGICAL BLOCK ADDRESS field set to 0001h:

	LBA 0	LBA 1	LBA 2	LBA 3	LBA 4	LBA 5	LBA 6	LBA 7	LBA 8	...
PB			PB			PB				...

LOWEST ALIGNED LOGICAL BLOCK ADDRESS field set to 0002h:

	LBA 0	LBA 1	LBA 2	LBA 3	LBA 4	LBA 5	LBA 6	LBA 7	LBA 8	LBA 9	...
PB			PB			PB				...	

LOWEST ALIGNED LOGICAL BLOCK ADDRESS field set to 0003h:

LBA 0	LBA 1	LBA 2	LBA 3	LBA 4	LBA 5	LBA 6	LBA 7	LBA 8	LBA 9	LBA 10	...
PB			PB			PB				...	

**Key:**

LBA n = logical block with LBA n

PB = physical block

**Figure 5 — Four logical blocks per physical block alignment examples**

If there is more than one logical block per physical block, then not all of the logical blocks are aligned to the physical block boundaries. When using logical block access commands (see 4.2.2), application clients should:

- specify an LBA that is aligned to a physical block boundary; and
- access an integral number of physical blocks, provided that the access does not go beyond the last LBA on the medium.

See annex E for an example method in which application clients may use alignment information to determine optimal performance for logical block access.

#### 4.6.2 [Physical block misaligned write reporting](#)

[A misaligned write command is a write command that specifies a:](#)

- [LOGICAL BLOCK ADDRESS field that does not correspond to the first LBA of a physical block; or](#)
- [TRANSFER LENGTH field that is not a multiple of the number of logical blocks per physical block \(see table 77\) and the number of logical blocks per physical block is greater than one.](#)

[LPS Misalignment log parameters in the LPS Misalignment log page \(see 6.4.6\) report misaligned write command information. The oldest reported misaligned write command is identified by a parameter code value of 0001h, with each successive time ordered misaligned write command identified by successively numbered parameter codes.](#)

The maximum number of reportable LPS Misalignment log parameters supported by the device server is indicated in the MAX\_LPSM field of the LPS Misalignment Count log parameter. If the parameter code value for the newest LPS Misalignment log parameter (i.e. parameter code with the greatest value) is equal to the value of the MAX\_LPSM field in the LPS Misalignment Count log parameter (i.e. the log is full), then the device server shall not append additional LPS Misalignment log parameters to the LPS Misalignment log page.

The MWR field (see 6.5.10) specifies how misaligned write commands are processed. If the device server processes a misaligned write command, and the MWR field is set to:

- a) 00b (i.e. DISABLED), then the device server shall process that write command without adding a log parameter in the LPS Misalignment log page (see 6.4.6);
  - b) 01b (i.e. ENABLED), then the device server shall:
    - 1) process that write command;
    - 2) add a log parameter in the LPS Misalignment log page, if the log is not full; and
    - 3) if that write command completes without error, complete the command with GOOD status with the sense key set to COMPLETED and the additional sense code set to MISALIGNED WRITE COMMAND;
- or
- c) 10b (i.e. TERMINATE), then the device server shall:
    - 1) add a log parameter in the LPS Misalignment log page, if the page is not full; and
    - 2) terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to MISALIGNED WRITE COMMAND.

## 4.7 Logical block provisioning

### 4.7.1 Logical block provisioning overview

Each LBA in a logical unit is either mapped or unmapped. For LBAs that are mapped, there is a known relationship between the LBA and one or more physical blocks that contain logical block data. For LBAs that are unmapped, the relationship between the LBA and a physical block is not defined. Figure 6 shows two examples of the relationship between mapped and unmapped LBAs and physical blocks in a logical unit. One example shows one LBA per physical block and one example shows two LBAs per physical block. The LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT field is defined in the READ CAPACITY (16) parameter data (see 5.18.2).

LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT field set to 0h (i.e.,  $2^0$  logical blocks per physical block):

LBA 0	LBA 1	LBA 2	LBA 3	LBA 4	LBA 5	LBA 6	LBA 7
PB	PB	Unmapped	PB	Unmapped	Unmapped	PB	PB

LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT field set to 1h (i.e.,  $2^1$  logical blocks per physical block):

LBA 0	LBA 1	LBA 2	LBA 3	LBA 4	LBA 5	LBA 6	LBA 7
PB		Unmapped		PB		Unmapped	

**Key:**

LBA n = logical block with LBA n

PB = physical block

Unmapped = the relationship between the LBA(s) and a physical block is not defined

**Figure 6 — Examples of the relationship between mapped and unmapped LBAs and physical blocks**

Each unmapped LBA is either anchored or deallocated. Anchored and deallocated are states in the LBP state machine (see 4.7.4) that have the following properties:

- a) a write command that specifies an anchored LBA does not require allocation of additional LBA mapping resources for that LBA; and
- b) a write command that specifies a deallocated LBA may require allocation of LBA mapping resources.

Depending on the logical block provisioning types (see table 5), the quantity of LBA mapping resources available to a logical unit may be greater than, equal to, or less than the quantity required to store logical block data for every LBA.

Table 5 list the logical block provisioning states supported by each type of logical block provisioning.

**Table 5 — Logical block provisioning states supported by logical block provisioning type**

Type	Logical block provisioning states			Reference
	Mapped	Unmapped		
		Anchored	Deallocated	
Fully provisioned	Mandatory	Prohibited	Prohibited	4.7.2
Resource provisioned	Mandatory	Optional	Optional	4.7.3.2
Thin provisioned	Mandatory	Optional	Mandatory	4.7.3.3

#### 4.7.2 Fully provisioned logical unit

The device server shall map every LBA in a fully provisioned logical unit. A fully provisioned logical unit shall provide enough LBA mapping resources to contain all logical blocks for the logical unit's capacity as reported in the READ CAPACITY (10) parameter data (see 5.17.2) and the READ CAPACITY (16) parameter data (see 5.18.2). The device server shall not cause any LBA on a fully provisioned logical unit to become unmapped (i.e., anchored or deallocated).

A fully provisioned logical unit does not support logical block provisioning management (see 4.7.3). A fully provisioned logical unit may support the GET LBA STATUS command (see 5.5).

The device server in a fully provisioned logical unit shall set the LBPME bit to zero in the READ CAPACITY (16) parameter data (see 5.18.2).

#### 4.7.3 Logical block provisioning management

##### 4.7.3.1 Logical block provisioning management overview

A logical unit that supports logical block provisioning management (i.e., implements unmapped LBAs, unmap operations, and related actions) shall be either:

- a) resource provisioned (see 4.7.3.2); or
- b) thin provisioned (see 4.7.3.3).

A logical unit that supports logical block provisioning management may change from resource provisioned to thin provisioned as resources become unavailable. If the logical unit transitions from resource provisioned to thin provisioned, then the logical unit shall change the PROVISIONING TYPE field to 010b (i.e., the logical unit is thin provisioned) in the Logical Block Provisioning VPD page (see 6.6.6) and establish a unit attention condition with the additional sense code set to INQUIRY DATA HAS CHANGED as described in SPC-5.

A logical unit that supports logical block provisioning management shall implement the LBP state machine (see 4.7.4) for each LBA.

The device server in a logical unit that supports logical block provisioning management:

- a) shall support the Logical Block Provisioning VPD page (see 6.6.6);

- b) may supply a provisioning group designation descriptor as defined in the Logical Block Provisioning VPD page;
- c) may support logical block provisioning thresholds (see 4.7.3.7.1);
- d) may support the GET LBA STATUS command (see 5.5);
- e) should support the Block Limits VPD page (see 6.6.4); and
- f) shall support at least one of the following unmap mechanisms:
  - A) the UNMAP command (see 5.30);
  - B) the UNMAP bit in the WRITE SAME (10) command (see 5.47);
  - C) the UNMAP bit in the WRITE SAME (16) command (see 5.48); or
  - D) the UNMAP bit in the WRITE SAME (32) command (see 5.49).

If the device server supports:

- a) the UNMAP bit in the WRITE SAME (10) command or in the WRITE SAME (16) command; and
- b) the WRITE SAME (32) command (see 5.49),

then the device server shall support the UNMAP bit in the WRITE SAME (32) command.

If a device server supports the UNMAP command and the Block Limits VPD page, then the device server shall:

- a) set the MAXIMUM UNMAP LBA COUNT field (see 6.6.4) to a value greater than or equal to one; and
- b) set the MAXIMUM UNMAP DESCRIPTOR COUNT field (see 6.6.4) to a value greater than or equal to one.

#### 4.7.3.2 Resource provisioned logical unit

A resource provisioned logical unit shall support logical block provisioning management (see 4.7.3.1).

The device server shall map, anchor, or deallocate each LBA in a resource provisioned logical unit. A resource provisioned logical unit shall provide LBA mapping resources sufficient to map all LBAs for the logical unit's capacity as indicated in the RETURNED LOGICAL BLOCK ADDRESS field of the READ CAPACITY (10) parameter data (see 5.17.2) and the READ CAPACITY (16) parameter data (see 5.18.2). A resource provisioned logical unit may provide resources in excess of this requirement.

The device server in a resource provisioned logical unit:

- a) shall set the LBPME bit to one in the READ CAPACITY (16) parameter data (see 5.18.2);
- b) shall set the PROVISIONING TYPE field to 001b (i.e., resource provisioned) in the Logical Block Provisioning VPD page (see 6.6.6); and
- c) may set the ANC\_SUP bit to one in the Logical Block Provisioning VPD page.

The initial condition of every LBA in a resource provisioned logical unit is anchored (see 4.7.4.1) or dellocated.

#### 4.7.3.3 Thin provisioned logical unit

A thin provisioned logical unit shall support logical block provisioning management (see 4.7.3.1).

The device server in a thin provisioned logical unit may indicate a larger capacity in the RETURNED LOGICAL BLOCK ADDRESS field (see 5.17.2) and the READ CAPACITY (16) parameter data (see 5.18.2) than the number of LBA mapping resources available for mapping LBAs in the logical unit.

The device server shall map, anchor, or deallocate each LBA in a thin provisioned logical unit (see table 5). A thin provisioned logical unit is not required to provide LBA mapping resources sufficient to map all LBAs for the logical unit's capacity as indicated in the RETURNED LOGICAL BLOCK ADDRESS field of the READ CAPACITY (10) parameter data (see 5.17.2) and the READ CAPACITY (16) parameter data (see 5.18.2).

If the logical unit does not support anchored LBAs (i.e., the ANC\_SUP bit is set to zero in the Logical Block Provisioning VPD page (see 6.6.6)), then:

- a) every unmapped LBA in the logical unit shall be deallocated; and
- b) the device server shall terminate every command that specifies anchoring an LBA (e.g., a WRITE SAME command with the ANCHOR bit set to one (see 5.47)).

The device server in a thin provisioned logical unit shall set:

- a) the LBPME bit to one in the READ CAPACITY (16) parameter data (see 5.18.2); and
- b) the PROVISIONING TYPE field to 010b (i.e., thin provisioned) in the Logical Block Provisioning VPD page (see 6.6.6).

The initial condition of every LBA in a thin provisioned logical unit is deallocated (see 4.7.4.1).

#### **4.7.3.4 Unmapping LBAs**

##### **4.7.3.4.1 Unmapping overview**

A logical unit that supports logical block provisioning management shall support unmapping of LBAs.

The logical block provisioning state of an LBA may change to unmapped (i.e., deallocated (see 4.7.4.6) or anchored (see 4.7.4.7)) as a result of:

- a) an unmap request (see 4.7.3.4.2);
- b) an autonomous LBA transition (see 4.7.3.5) (e.g., following a FORMAT UNIT command (see 5.4) or a write command that sets the logical block data to zero); or
- c) other commands that result in LBAs being initialized to unmapped (e.g., a SANITIZE command (see 5.25)).

##### **4.7.3.4.2 Processing unmap requests**

Application clients use unmap commands (see 4.2.2) to request that LBAs be unmapped. For each LBA that is requested to be unmapped, the device server shall:

- a) perform an unmap operation (see 4.7.3.4.3) on the LBA; or
- b) make no change to the logical block provisioning state of the LBA.

The application client determines the logical block provisioning state of LBAs using the GET LBA STATUS command (see 5.5).

Application clients should not rely on an UNMAP command (see 5.30) to cause specific data (e.g., zeros) to be returned by subsequent read operations on the specified LBAs. To produce consistent results for subsequent read operations, a write command (e.g., the WRITE SAME command) should be used to write user data.

EXAMPLE - To ensure that subsequent read operations return all zeros in a logical block, use the WRITE SAME (16) command with the NDOB bit set to one. If the UNMAP bit is set to one, then the device server may unmap the logical blocks specified by the WRITE SAME (16) command as described in 4.7.3.4.4.

##### **4.7.3.4.3 Unmap operations**

An unmap operation:

- a) results in a single LBA becoming either deallocated or anchored;
- b) may change the relationship between one LBA and one or more physical blocks; and
- c) may change the logical block data that is returned in response to a subsequent read command specifying that LBA.

The data in all other mapped LBAs on the medium shall be preserved. Performing an unmap operation (e.g., to change from anchored to deallocated, or remain in the same logical block provisioning state) on an unmapped LBA shall not be considered an error.

An unmap operation may or may not release LBA mapping resources.

An application client may use an unmap command (see 4.2.2) to request that the device server perform an unmap operation on each specified LBA. A single unmap command may result in zero or more unmap operations.

#### 4.7.3.4.4 WRITE SAME command and unmap operations

A WRITE SAME command (see 5.47, 5.48, and 5.49) may be used to request unmap operations that deallocate or anchor the specified LBAs. If unmap operations are requested in a WRITE SAME command, then for each specified LBA:

- a) if the Data-Out Buffer of the WRITE SAME command is the same as the logical block data returned by a read operation from that LBA while in the unmapped state (see 4.7.4.4), then:
  - 1) the device server performs the actions described in table 6; and
  - 2) if an unmap operation is not performed in step 1), then the device server shall perform the specified write operation to that LBA;
- or
- b) if the Data-Out Buffer of the WRITE SAME command is not the same as the logical block data returned by a read operation from that LBA while in the unmapped state (see 4.7.4.4), then the device server shall perform the specified write operation to that LBA.

**Table 6 — WRITE SAME command and unmap operations**

Logical block provisioning management type	Unmap operations that request to deallocate the specified LBA	Unmap operations that request to anchor the specified LBA
Thin provisioned logical unit	a) should perform an unmap operation to deallocate the LBA (see 4.7.4.6.1); or b) may perform an unmap operation to anchor the LBA (see 4.7.4.7.1)	should perform an unmap operation to anchor the LBA
Resource provisioned logical unit	a) should perform an unmap operation to anchor the LBA, or b) may perform an unmap operation to deallocate the LBA	should perform an unmap operation to anchor the LBA

A WRITE SAME command shall not cause an LBA to become unmapped if unmapping that LBA creates a case in which a subsequent read of that unmapped LBA is able to return logical block data that differs from the Data-Out Buffer for that WRITE SAME command (see 4.7.4.4).

If the device server does not support allowing a WRITE SAME command to request unmap operations, then the device server shall:

- a) perform the write operations specified by the WRITE SAME command; and
- b) not perform any unmap operations.

The device server shall perform the write operations specified by a WRITE SAME command and shall not perform any unmap operations if the device server sets the LBPRZ field to xx1b in the Logical Block Provisioning VPD page (see 6.6.6), and:

- a) any bit in the user data transferred from the Data-Out Buffer is not zero; or
- b) the protection information, if any, transferred from the Data-Out Buffer is not set to FFFF\_FFFF\_FFFF\_FFFFh.

The device server shall perform the write operations specified by a WRITE SAME command and shall not perform any unmap operations if the device server sets the LBPRZ field to 010b in the Logical Block Provisioning VPD page, and:

- a) the user data transferred from the Data-Out Buffer is not set to the provisioning initialization pattern; or
- b) the protection information, if any, transferred from the Data-Out Buffer is not set to FFFF\_FFFF\_FFFF\_FFFFh.

#### 4.7.3.5 Autonomous LBA transitions

A device server may perform the following actions at any time:

- a) transition any deallocated LBA to mapped;
- b) transition any anchored LBA to mapped; or
- c) transition any deallocated LBA to anchored.

If the LBPRZ field (see 6.6.6) is set to:

- a) xx1b, and a mapped LBA references a logical block that contains:
  - A) user data with all bits set to zero; and
  - B) protection information, if any, set to FFFF\_FFFF\_FFFF\_FFFFh;
- or
- a) 010b, and a mapped LBA references a logical block that contains:
  - A) user data set to the provisioning initialization pattern; and
  - B) protection information, if any, set to FFFF\_FFFF\_FFFF\_FFFFh,

then the device server may transition that mapped LBA to anchored or deallocated at any time.

The logical block provisioning state machine (see 4.7.4) specifies additional requirements for the transitions specified in this subclause.

#### 4.7.3.6 Logical unit resource exhaustion considerations

##### 4.7.3.6.1 Thin provisioned logical unit resource exhaustion considerations

If:

- a) a write operation is requested by an application client, and a temporary lack of LBA mapping resources prevents the logical unit from performing the write operation; or
- b) an unmap operation is requested by an application client to transition an LBA to the anchored state and a temporary lack of LBA mapping resources prevents the logical unit from anchoring the LBA,

then the device server shall terminate the command requesting the operation with CHECK CONDITION status with the sense key set to NOT READY and the additional sense code set to LOGICAL UNIT NOT READY, SPACE ALLOCATION IN PROGRESS. In this case, the application client should resend the command.

If:

- a) a write operation is requested by an application client, and a persistent lack of LBA mapping resources prevents the logical unit from performing the write operation; or
- b) an unmap operation is requested by an application client to transition an LBA to the anchored state and a persistent lack of LBA mapping resources prevents the logical unit from anchoring the LBA,

then the device server shall terminate the command requesting the unmap operation with CHECK CONDITION status with the sense key set to DATA PROTECT and the additional sense code set to SPACE ALLOCATION FAILED WRITE PROTECT. This condition shall not cause the device server to set the WP bit in the DEVICE-SPECIFIC PARAMETER field of the mode parameter header to one (see 6.5.1). In this case, recovery actions by the application client are outside the scope of this standard.

A logical block provisioning threshold may be available to monitor the availability of LBA mapping resources (see 4.7.3.7). A logical block provisioning log parameter that reports available LBA mapping resources may be available in the Logical Block Provisioning log page (see 6.4.5).

##### 4.7.3.6.2 Resource provisioned logical unit resource exhaustion considerations

If:

- a) a write operation is requested by an application client, and a temporary lack of LBA mapping resources prevents the logical unit from performing the write operation; or



- b) an unmap operation is requested by an application client to transition an LBA to the anchored state and a temporary lack of LBA mapping resources prevents the logical unit from anchoring the LBA,

then the device server shall terminate the command requesting the operation with CHECK CONDITION status with the sense key set to NOT READY and the additional sense code set to LOGICAL UNIT NOT READY, SPACE ALLOCATION IN PROGRESS. In this case, the application client should resend the command.

In a resource provisioned logical unit a write operation requested by an application client shall not result in a persistent lack of LBA mapping resources that prevents the logical unit from performing the write operation. Therefore, a resource provisioned logical unit shall not terminate any command with CHECK CONDITION status with the sense key set to DATA PROTECT and the additional sense code set to SPACE ALLOCATION FAILED WRITE PROTECT.

A logical block provisioning threshold may be available to monitor the availability of LBA mapping resources (see 4.7.3.7). A parameter that reports available LBA mapping resources may be available in the page (see 6.4.5).

#### 4.7.3.7 Logical block provisioning thresholds

##### 4.7.3.7.1 Logical block provisioning thresholds overview

Logical block provisioning thresholds provide a mechanism for the device server to establish a unit attention condition to notify application clients when thresholds related to logical block provisioning are crossed. Logical block provisioning thresholds may operate on an armed increasing basis or an armed decreasing basis.

If a device server supports logical block provisioning thresholds, then the device server:

- a) shall support the Logical Block Provisioning mode page (see 6.5.9); and
- b) may support the Logical Block Provisioning log page (see 6.4.5).

Logical block provisioning thresholds may be based on threshold sets (see 4.7.3.7.2) or percentages (see 4.7.3.7.3).

##### 4.7.3.7.2 Threshold sets

The end points of the range over which a logical block provisioning threshold operates are defined as follows:

$$\text{threshold minimum} = ((\text{threshold count} \times \text{threshold set size}) - (\text{threshold set size} \times 0.5))$$

$$\text{threshold maximum} = ((\text{threshold count} \times \text{threshold set size}) + (\text{threshold set size} \times 0.5))$$

where:

threshold minimum	is the lowest number of LBAs in the range for this threshold;
threshold maximum	is the highest number of LBAs in the range for this threshold;
threshold count	is the center of the threshold range for this threshold (i.e., the threshold count value as specified in the threshold descriptor in the Logical Block Provisioning mode page); and
threshold set size	is the number of LBAs in each threshold set (i.e., $2^{(\text{threshold exponent})}$ LBAs where the threshold exponent is indicated in the Logical Block Provisioning VPD page (see 6.6.6)).

Table 7 defines the meaning of the combinations of values for the THRESHOLD RESOURCE field, the THRESHOLD TYPE field, and the THRESHOLD ARMING field that are used for logical block provisioning thresholds. See the Logical Block Provisioning mode page (see 6.5.9) for the definition of these fields.

**Table 7 — Threshold resource value, threshold type value, and threshold arming value for logical block provisioning thresholds**

Threshold resource value	Threshold type value	Threshold arming value	Description
01h	000b	000b	The device server applies the threshold to the availability of LBA mapping resources and performs notifications as the availability of those resources decreases. <sup>a</sup>
02h	000b	001b	The device server applies the threshold to the usage of LBA mapping resources and performs notifications as the usage of those resources increases.
All other combinations			Reserved
<sup>a</sup> The point when availability of LBA mapping resources reaches zero corresponds to the persistent lack of LBA mapping resources described in 4.7.3.6.1.			

**4.7.3.7.3 Threshold percentages**

The end points of the range over which a logical block provisioning threshold percentages operates are defined as follows:

$$\text{threshold minimum} = (\text{threshold count} - (\text{threshold percentage size} \times 0.5))$$

$$\text{threshold maximum} = (\text{threshold count} + (\text{threshold percentage size} \times 0.5))$$

where:

threshold minimum	is the lowest percentage of device resources available for allocation to logical blocks in the range for this threshold;
threshold maximum	is the highest percentage of device resources available for allocation to logical blocks in the range for this threshold;
threshold count	is the center of the threshold range for this threshold (i.e., the threshold count value as specified in the threshold descriptor in the Logical Block Provisioning mode page); and
threshold set size	is the percentage of allocation resources in each threshold percentage (i.e., the percentage indicated by the THRESHOLD PERCENTAGE field (see 6.6.6)).

Table 4 defines the meaning of the combinations of values for the THRESHOLD RESOURCE field, the THRESHOLD TYPE field, and the THRESHOLD ARMING field that are used for logical block provisioning threshold percentages. See the Logical Block Provisioning mode page (see 6.5.9) for the definition of these fields.


**Table 8 — Threshold resource value, threshold type value, and threshold arming value for logical block provisioning percentages**

Threshold resource value	Threshold type value	Threshold arming value	Description
01h	001b	000b	The device server applies the threshold to the availability of LBA mapping resources and performs notifications as the availability of those resources decreases.
All other combinations			Reserved

#### 4.7.3.7.4 Logical block provisioning armed decreasing thresholds

Figure 7 shows the operation of a logical block provisioning armed decreasing threshold. Figure 7 represents the entire range of possible values over which the threshold is being applied (e.g., for an available resource, the lowest value represents zero available resources and the highest value represents the maximum possible number of available resources).

If enabled, reporting of armed decreasing threshold events (i.e., the THRESHOLD ARMING field is set to 000b in the threshold descriptor in the Logical Block Provisioning mode page (see 6.5.9)) operates as shown in figure 7.



Notes:


- a) if the value to which the threshold is being applied drops below the threshold maximum for the threshold range, then the notification trigger shall be enabled;
- b) if the value to which the threshold is being applied increases above the threshold maximum for the threshold range, then the notification trigger shall be disabled;
- c) if the notification trigger is enabled, then the device server may disable the notification trigger and perform logical block provisioning threshold notification (see 4.7.3.7.6); and
- d) if the notification trigger is enabled and the value to which the threshold is being applied drops below the threshold minimum for the threshold range, then the device server shall disable the notification trigger and perform logical block provisioning threshold notification as defined in 4.7.3.7.6.

**Figure 7 — Armed decreasing threshold operation**

#### 4.7.3.7.5 Logical block provisioning armed increasing thresholds

Figure 8 shows the operation of a logical block provisioning armed increasing threshold. Figure 8 represents the entire range of possible values over which the threshold is being applied (e.g., for tracking usage of a resource, the lowest value represents zero resources being used and the highest value represents the maximum possible number of resources being used).

If enabled, reporting of armed increasing threshold events (i.e., the THRESHOLD ARMING field is set to 001b in the threshold descriptor in the Logical Block Provisioning mode page (see 6.5.9)) operates as shown in figure 8.



Notes:

- a) if the value to which the threshold is being applied increases above the threshold minimum for the threshold range, then the notification trigger shall be enabled;
- b) if the value to which the threshold is being applied decreases below the threshold minimum for the threshold range, then the notification trigger shall be disabled;
- c) if the notification trigger is enabled, then the device server may disable the notification trigger and perform logical block provisioning threshold notification (see 4.7.3.7.6); and
- d) if the notification trigger is enabled and the value to which the threshold is being applied increases above the threshold maximum for the threshold range, then the device server shall disable the notification trigger and perform logical block provisioning threshold notification as defined in 4.7.3.7.6.

**Figure 8 — Armed increasing threshold operation**

#### 4.7.3.7.6 Logical block provisioning threshold notification

If the LBPERE bit is set to one in the Read-Write Error Recovery mode page (see 6.5.10), then logical block provisioning threshold notification is enabled and the device server shall perform notification for thresholds with the THRESHOLD TYPE field set to 000b in the threshold descriptor in the Logical Block Provisioning mode page (see 6.5.9) as follows:

- a) if the SITUA bit is set to one in the Logical Block Provisioning mode page, then:
  - A) if the device server has not established a unit attention condition as a result of this threshold being crossed since the last logical unit reset (see SAM-5) and a command through which the device server is able to report a unit attention condition arrives on any I\_T nexus, then the device server shall establish a unit attention condition with the additional sense code set to THIN PROVISIONING SOFT THRESHOLD REACHED for only the SCSI initiator port associated with the I\_T nexus on which that command was received before processing that command; or
  - B) if the device server has established a unit attention condition as a result of this threshold being crossed since the last logical unit reset and a command through which the device server is able to report a unit attention condition arrives on any I\_T nexus, then the device server should establish a unit attention condition with the additional sense code set to THIN PROVISIONING SOFT THRESHOLD REACHED for only the SCSI initiator port associated with the I\_T nexus on which that command was received before processing that command unless establishment of the unit attention condition causes a vendor specific frequency of unit attention conditions for this threshold to be exceeded;
- or
- b) if the SITUA bit is set to zero, then:
  - A) if the device server has not established a unit attention condition for the SCSI initiator port associated with all I\_T nexuses as a result of this threshold being crossed since the last logical unit reset (see SAM-5), then the device server shall establish a unit attention condition with the additional sense code set to THIN PROVISIONING SOFT THRESHOLD REACHED for the SCSI initiator port associated with every I\_T nexus; or
  - B) if the device server has established a unit attention condition for the SCSI initiator ports associated with all I\_T nexuses as a result of this threshold being crossed since the last logical unit reset, then the device server should establish a unit attention condition with the additional sense code set to THIN PROVISIONING SOFT THRESHOLD REACHED for the SCSI initiator port associated with every I\_T nexus, unless establishment of the unit attention condition causes a vendor specific frequency of unit attention conditions for this threshold to be exceeded.

If a unit attention condition is established as described in this subclause, then the device server shall report the following value in the INFORMATION field in the sense data (see SPC-5):

- a) the byte offset in the Logical Block Provisioning mode page of the first byte of the threshold descriptor to which this threshold notification applies.

If a unit attention condition with the additional sense code set to THIN PROVISIONING SOFT THRESHOLD REACHED is received by the application client, then the application client should reissue the command and take further recovery actions (e.g., administrator notification or other administrator actions). These recovery actions are outside the scope of this standard.

If the LBPERE bit is set to zero, then logical block provisioning threshold notification is disabled and the device server shall not establish any unit attention condition with the additional sense code set to THIN PROVISIONING SOFT THRESHOLD REACHED.

An additional sense code set to THIN PROVISIONING SOFT THRESHOLD REACHED is applicable to both thin provisioned logical units (see 4.7.3.3) and resource provisioned logical units (see 4.7.3.2).

#### **4.7.4 LBP (logical block provisioning) state machine**

##### **4.7.4.1 LBP state machine overview**

The LBP (logical block provisioning) state machine describes the mapping and unmapping of a single LBA by the device server for a thin provisioned logical unit (see 4.7.3.3) or a resource provisioned logical unit (see 4.7.3.2). This state machine does not apply to fully provisioned logical units (see 4.7.2).

There is one instance of this state machine for each LBA. If a command requests mapping or unmapping of more than one LBA, then there may be an independent transition in each instance of the state machine (e.g., each LBA may individually transition from mapped to deallocated or from anchored to deallocated).

##### **4.7.4.2 LBP state machine for logical units supporting anchored LBAs**


If the logical unit supports anchored LBAs (i.e., the ANC\_SUP bit is set to one) and deallocated LBAs (i.e., is a thin provisioned logical unit, or a resource provisioned logical unit), then this state machine consists of the following states:

- a) LBP1:Mapped state (see 4.7.4.5);
- b) LBP2:Deallocated state (see 4.7.4.6) (initial state for thin provisioned logical units); and
- c) LBP3:Anchored state (see 4.7.4.7) (initial state for resource provisioned logical units supporting anchored LBAs).

For thin provisioned logical units and resource provisioned logical units not supporting anchored LBAs the initial state of the state machine associated with each LBA is the LBP2:Deallocated state.

For resource provisioned logical units supporting anchored LBAs the initial state of the state machine associated with each LBA is the LBP3:Anchored state.

Figure 9 describes the LBP state machine for a logical unit that supports anchored LBAs and deallocated LBAs.



**Figure 9 — LBP state machine (anchored LBAs supported and deallocated LBAs supported)**


#### 4.7.4.3 LBP state machine for logical units not supporting anchored LBAs

If the logical unit does not support anchored LBAs (i.e., is a thin provisioned logical unit and the ANC\_SUP bit is set to zero or a resource provisioned logical unit and the ANC\_SUP bit is set to zero), then this state machine consists of the following states:

- a) LBP1:Mapped state (see 4.7.4.5); and
- b) LBP2:Deallocated state (see 4.7.4.6) (initial state).

The initial state of the state machine associated with each LBA is the LBP2:Deallocated state.

Figure 10 describes the LBP state machine for a logical unit that does not support anchored LBAs.



**Figure 10 — LBP state machine (anchored LBAs not supported)**

#### 4.7.4.4 Performing read operations with respect to logical block provisioning

Table 9 defines the logical block data that a read operation shall return for a mapped LBA.

**Table 9 — Logical block data returned by a read operation from a mapped LBA**

<b>Condition</b>	<b>Logical block data returned</b>
The LBA became mapped as the result of a format operation or sanitize operation and no write command has specified that LBA since the LBA became mapped.	The logical block data that was written to that LBA by the format operation or the sanitize operation.
The LBA became mapped as the result of a write command and no additional write command has specified that LBA since the LBA was mapped.	The logical block data that was written to that LBA by that write command.
The LBA became mapped as the result of an autonomous transition, and no write command has specified that LBA since the LBA was mapped.	The logical block data that would be returned if that autonomous transition had not occurred and the LBA had remained unmapped (see table 10).
A write command has specified that LBA since that LBA was mapped.	The logical block data that was most recently written to that LBA.

Table 10 defines the logical block data that a read operation shall return for an unmapped LBA.

**Table 10 — Logical block data returned by a read operation from an unmapped LBA**

LBPRZ field <sup>a</sup>	Method used to unmap the LBA	Logical block data returned
000b	see <sup>b</sup>	The value specified to be written if the write operation had been performed.
	see <sup>c</sup>	Logical block data containing: a) user data set to a vendor-specific value that is not obtained from any other LBA; and b) protection information, if any, set to FFFF_FFFF_FFFF_FFFFh.
xx1b	Any	Logical block data containing: a) user data set to zero; and b) protection information, if any, set to FFFF_FFFF_FFFF_FFFFh.
010b	Any	Logical block data containing: a) user data set to the provisioning initialization pattern; and b) protection information, if any, set to FFFF_FFFF_FFFF_FFFFh.
<sup>a</sup> The LBPRZ field is in the Logical Block Provisioning VPD page (see 6.6.6). <sup>b</sup> A command for which the device server is allowed to perform either: a) a write with specified logical block data; or b) an unmap operation if the logical block data returned by read operations from unmapped LBAs matches the logical block data specified for the command that resulted in the unmap operation. These commands are: a) a FORMAT UNIT command specifying an initialization pattern; b) a SANITIZE command specifying a sanitize overwrite operation; and c) a WRITE SAME command with the UNMAP bit set to one. <sup>c</sup> These methods include but are not limited to: a) a FORMAT UNIT command not specifying an initialization pattern; b) a REASSIGN BLOCKS command; c) a SANITIZE command specifying a sanitize block erase operation or a sanitize cryptographic erase operation; and d) an UNMAP command.		

After a read operation returns a value for an LBA, subsequent read operations from that LBA shall return the same value until a subsequent command alters the logical block data in that LBA (e.g., a write command or an unmap command (see table 34)).

#### 4.7.4.5 LBP1:Mapped state

##### 4.7.4.5.1 LBP1:Mapped state description

Upon entry into this state, the relationship between the LBA and the physical block(s) that contains the logical block for that LBA shall be established.

If this state was entered from the LBP2:Deallocated state (see 4.7.4.6), then the device server shall allocate LBA mapping resources, if any, required to map the LBA.

If this state was entered from the LBP3:Anchored state (see 4.7.4.7), then:

- a) the device server shall not allocate LBA mapping resources; and
- b) the resource exhaustion conditions described in 4.7.3.6.1 shall not occur.



**4.7.4.5.2 Transition LBP1:Mapped to LBP2:Deallocated**

This transition shall occur after:

- a) an unmap operation that results in the deallocation of the LBA that was mapped.

This transition may occur at any time if the LBPRZ field in the Logical Block Provisioning VPD page (see 6.6.6) is set to:

- a) xx1b, and the mapped LBA references a logical block that contains:
  - A) user data with all bits set to zero; and
  - B) protection information, if any, set to FFFF\_FFFF\_FFFF\_FFFFh;
- or
- b) 010b, and a mapped LBA references a logical block that contains;
  - A) user data set to the provisioning initialization pattern; and
  - B) protection information, if any, set to FFFF\_FFFF\_FFFF\_FFFF.

**4.7.4.5.3 Transition LBP1:Mapped to LBP3:Anchored**

This transition shall occur after:

- a) an unmap operation that results in the anchoring of the LBA that was mapped.

This transition may occur at any time if the LBPRZ field (see 6.6.6) is set to:

- a) xx1b, and the mapped LBA references a logical block that contains:
  - A) user data with all bits set to zero; and
  - B) protection information, if any, set to FFFF\_FFFF\_FFFF\_FFFFh
- or
- b) 010b, and a mapped LBA references a logical block that contains;
  - A) user data set to the provisioning initialization pattern; and
  - B) protection information, if any, set to FFFF\_FFFF\_FFFF\_FFFF.

**4.7.4.6 LBP2:Deallocated state****4.7.4.6.1 LBP2:Deallocated state description**

While in this state:

- a) there shall be no relationship between the LBA and any physical block(s); and
- b) the device server should deallocate LBA mapping resources after they are no longer in use.

For an LBA in this state, an unmap operation that specifies deallocation of the LBA shall not cause a transition from this state.

The device server shall process a read command specifying an LBA in this state (i.e., a deallocated LBA) as described in table 10.

**4.7.4.6.2 Transition LBP2:Deallocated to LBP1:Mapped**

This transition:

- a) shall occur after a write operation to the LBA that was deallocated; or
- b) may occur at any time for reasons outside the scope of this standard.

**4.7.4.6.3 Transition LBP2:Deallocated to LBP3:Anchored**

This transition:

- a) shall occur after an unmap operation that results in anchoring of the LBA that was deallocated; or
- b) may occur at any time for reasons outside the scope of this standard.

#### **4.7.4.7 LBP3:Anchored state**

##### **4.7.4.7.1 LBP3:Anchored state description**

Upon entry into this state:

- a) LBA mapping resources shall be associated with the LBA; and
- b) there may or may not be a relationship between the LBA and physical block(s).

If this state was entered from the LBP2:Deallocated state, then the device server shall allocate LBA mapping resources, if any, required to anchor the LBA.

If this state was entered from the LBP1:Mapped state, then:

- a) the device server shall not allocate LBA mapping resources;
- b) the device server relies on LBA mapping resources already allocated to the LBA; and
- c) the resource exhaustion conditions described in 4.7.3.6.1 shall not occur.

For an LBA in this state, an unmap operation that specifies anchoring of the LBA shall not cause a transition from of this state.

The device server shall process a read command specifying an LBA in this state (i.e., an anchored LBA) as described in table 10.

##### **4.7.4.7.2 Transition LBP3:Anchored to LBP1:Mapped**

This transition:

- a) shall occur after a write operation to the LBA that was anchored; or
- b) may occur at any time for reasons outside the scope of this standard.

##### **4.7.4.7.3 Transition LBP3:Anchored to LBP2:Deallocated**

This transition shall occur after:

- a) an unmap operation that results in the deallocation of the LBA that was anchored.

## **4.8 Data de-duplication**

Data de-duplication is the ability of a device server to recognize redundant or duplicate data and reduce the number of duplicate or redundant copies of the data while maintaining the application client supplied LBAs of the duplicate or redundant copies of the data. De-duplication shall not affect protection information, if any. Logical units that support data de-duplication may report a count of LBA resources that have been made available as a result of being de-duplicated (see 6.4.5.5).

Data de-duplication may impact the number of LBA Mapping resources indicated in the Logical Block Provisioning log page (see 6.4.5) (e.g., the value indicated in the Available LBA Mapping Resource Count log parameter (see 6.4.5.2) and the Used LBA Mapping Resource Count log parameter (see 6.4.5.3) may not change as a result of successful write operations that contain data that is de-duplicated or successful unmap operations on LBAs that contain data that has been de-duplicated).

## **4.9 Ready state**

A direct access block device is ready when the device server is capable of processing logical block access commands that require access to the medium (see 4.2.2).

A direct access block device using a removable medium (see 4.4) is not ready until a volume is mounted and other conditions are met (see 4.3). While a direct access block device is not ready the device server shall

terminate logical block access commands with CHECK CONDITION status with the sense key set to NOT READY and the appropriate additional sense code for the condition.

Some direct access block devices may be switched from being ready to being not ready by using the START STOP UNIT command (see 5.26).

To make a direct access block device ready, an application client may be required to issue a START STOP UNIT command:

- a) with a START bit set to one and the POWER CONDITION field set to 0h (i.e., START\_VALID); or
- b) with the POWER CONDITION field set to 1h (i.e., ACTIVE).

## 4.10 Initialization

A direct access block device may require initialization of its medium prior to processing logical block access commands. This initialization is requested by an application client using a FORMAT UNIT command (see 5.4). Parameters related to the format (e.g., logical block length) may be set with a MODE SELECT command (see SPC-5 and 6.5.2) prior to the format operation. Some direct access block devices are initialized by means outside the scope of this standard. The time when the initialization occurs is vendor specific.

Direct access block devices using a non-volatile medium may save the parameters related to the format and only require initialization once. However, some mode parameters may require initialization after each logical unit reset. A catastrophic failure of the direct access block device may require that an application client send a FORMAT UNIT command to recover from the failure.

Direct access block devices that use a volatile medium may require initialization after each logical unit reset prior to the processing of logical block access commands (see 4.2.2). Mode parameters may also require initialization after logical unit resets.

NOTE 2 - It is possible that mode parameter block descriptors read with a MODE SENSE command before a FORMAT UNIT completes contain information not reflecting the true state of the medium.

A direct access block device may become format corrupt after processing a MODE SELECT command that changes parameters (e.g., the logical block length) related to the medium format. During this time, the device server may terminate logical block access commands with CHECK CONDITION status with the sense key set to NOT READY and the appropriate additional sense code for the condition.

Any time the READ CAPACITY (10) parameter data (see 5.17.2) or the READ CAPACITY (16) parameter data (see 5.18.2) changes (e.g., when a FORMAT UNIT command or a MODE SELECT command causes a change to the logical block length or protection information, or when a vendor specific mechanism causes a change), then the device server shall establish a unit attention condition for the SCSI initiator port (see SAM-5) associated with each I\_T nexus, except the I\_T nexus on which the command causing the change was received with the additional sense code set to CAPACITY DATA HAS CHANGED.

NOTE 3 - Logical units compliant with SBC were not required to establish a unit attention condition with the additional sense code set to CAPACITY DATA HAS CHANGED.

## 4.11 Sanitize operations

### 4.11.1 Sanitize operations overview

A sanitize operation causes the device server to:

- a) affect the information on the logical unit's medium such that recovery of logical block data from the medium is not possible;
- b) affect the information in the cache by a method that is outside the scope of this standard such that previously existing data in cache is unable to be accessed; and

- c) prevent future access by an application client to cache or medium where the device server is unable to alter the information.

One of the following sanitize operations may be requested on the logical unit's medium:

- a) a sanitize overwrite operation that causes the device server to alter information by writing a data pattern to the medium one or more times;
- b) a sanitize block erase operation that causes the device server to alter information by setting the physical blocks to a vendor specific value; or
- c) a sanitize cryptographic erase operation that causes the device server to change encryption keys to prevent correct decryption of previously stored information, which may cause protection information, if any, to be indeterminate.

For zoned block devices the ZNR bit (see 5.25) controls whether a reset write pointer operation (see ZBC) is performed on each write pointer zone after successful completion of a sanitize operation.

An application client may request that a sanitize operation be performed in the restricted completion mode or the unrestricted completion mode (see 4.11.4) using the AUSE bit (see 5.25).

In the unrestricted completion mode, a SANITIZE command with the EXIT FAILURE MODE service action exits a failed sanitize operation.

In the restricted completion mode, the only method to exit a failed sanitize operation is for a SANITIZE command to request another sanitize operation and for that operation to complete without error. If a sanitize operation in the restricted completion mode completes with an error, and a subsequent SANITIZE command requests the unrestricted completion mode (i.e., the AUSE bit set to one), then the device server shall terminate that SANITIZE command as described in 5.25.1.

All sanitize operations shall be performed on:

- a) the medium that is being used to store logical block data;
- b) the medium that is not being used to store logical block data (e.g., areas previously used to store logical block data, areas available for allocation, and physical blocks that have become inaccessible); and
- c) all cache.

An application client requests that the device server perform a sanitize operation using the SANITIZE command. While the medium is write protected (see 4.12) the device server shall terminate a SANITIZE command with CHECK CONDITION status with the sense key set to DATA PROTECT and the appropriate additional sense code for the condition.

#### 4.11.2 Commands allowed during sanitize

Those of the following that the device server supports while not performing a sanitize operation are allowed during sanitize:

- a) INQUIRY commands;
- b) LOG SENSE commands that specify the Temperature log page (see SPC-5);
- c) MODE SENSE commands that specify (see SPC-5):
  - A) ~~the mode page header;~~
  - B) the Informational Exceptions Control mode page;
  - C) the Caching mode page;
  - D) the Control mode page;
  - E) the Protocol Specific Port mode page; or
  - F) the Protocol Specific Logical Unit mode page
- d) READ CAPACITY (16) commands (see 5.16);
- e) REPORT LUNS commands (see SPC-5);
- f) REPORT SUPPORTED OPERATION CODES commands (see SPC-5);
- g) REPORT SUPPORTED TASK MANAGEMENT FUNCTIONS commands (see SPC-5);
- h) REPORT ZONES commands (see ZBC) with:
  - A) the ZONE START LBA field set to zero;
  - B) the REPORTING OPTIONS field set to 3Fh;

- C) [the PARTIAL bit set to one; and](#)
- D) [the ALLOCATION LENGTH field set to a value less than or equal to 64;](#)

and

- i) REQUEST SENSE commands.

#### 4.11.3 Performing a sanitize operation

Before performing a sanitize operation, the device server shall:

- a) terminate all commands in all task sets except commands allowed during sanitize (see 4.11.2) with CHECK CONDITION status with the sense key set to NOT READY, the additional sense code set to LOGICAL UNIT NOT READY, SANITIZE IN PROGRESS, and the PROGRESS INDICATION field in the sense data set to indicate that the sanitize operation is beginning;
- b) stop all enabled power condition timers (see SPC-5);
- c) stop all timers for enabled background scan operations (see 4.24);
- d) stop all timers or counters enabled for device specific background functions (see SPC-5);
- e) discard partially downloaded microcode, if any; and
- f) close open streams (see 4.34), if any.

While performing a sanitize operation, the device server shall:

- a) process commands allowed during sanitize (see 4.11.2) and terminate all other commands with CHECK CONDITION status with the sense key set to NOT READY, the additional sense code set to LOGICAL UNIT NOT READY, SANITIZE IN PROGRESS, and the PROGRESS INDICATION field in the sense data set to indicate the progress of the sanitize operation;
- b) provide pollable sense data (see SPC-5) with the sense key set to NOT READY, the additional sense code set to LOGICAL UNIT NOT READY, SANITIZE IN PROGRESS, and the PROGRESS INDICATION field set to indicate the progress of the sanitize operation;
- c) suspend the sanitize operation while processing the following conditions (see SAM-5):
  - A) a power on;
  - B) a hard reset;
  - C) a logical unit reset; or
  - D) a power loss expected;
- d) not suspend the sanitize operation while processing an I\_T nexus loss;
- e) resume performing the sanitize operation after processing:
  - A) a logical unit reset; or
  - B) a power loss expected condition in which no power loss occurs within constraints defined by the applicable SCSI transport protocol standard (e.g., power loss timeout in SPL-3);
- f) process task management functions without affecting the processing of the sanitize operation (e.g., an ABORT TASK task management function aborts the SANITIZE command and has no effect on performing the sanitize operation);
- g) not alter mode data, INQUIRY data, or READ CAPACITY (16) parameter data (e.g., the number of logical blocks, logical block length, or protection information settings for the logical unit); and
- h) identify inaccessible physical blocks and in a vendor specific manner prevent future access to these blocks following a successful sanitize operation.

#### 4.11.4 Completing a sanitize operation

If a sanitize operation completes without error, and logical block provisioning management (see 4.7.3) is supported, then:

- a) the initial condition for every LBA should be anchored (see 4.7.3.2) or deallocated (see 4.7.3.3); and
- b) read operations and write operations should complete without error.

If a sanitize operation completes without error and logical block provisioning management is not supported, then:

- a) read commands are processed as described in 5.25.2.2, 5.25.2.3, 5.25.2.4, and 5.25.2.5; and
- b) write operations should complete without error.

If a sanitize operation completes without error on a zoned block device, then:

- a) the initial condition for each LBA in write pointer zones (see ZBC) is determined by the ZNR bit (see 5.25); ~~and~~
- b) the zone condition (see ZBC) shall be set to offline for every write pointer zone that was not able to be sanitized (e.g., due to an inability to write to that zone);
- c) the zone condition shall be set to not write pointer for every conventional zone; and
- d) unless otherwise specified, the zone condition for each write pointer zone shall be set to the zone condition associated with the initial state for that write pointer zone after power on (see ZBC).

If the sanitize operation completes with an error in restricted completion mode, then the device server shall:

- a) terminate the SANITIZE command being performed, if any (e.g., the IMMED bit was set to zero in the CDB, and the failure occurs before status is returned for the command), with CHECK CONDITION status with the sense key set to MEDIUM ERROR and the additional sense code set to SANITIZE COMMAND FAILED; and
- b) until completion of a successful sanitize operation has occurred, terminate all commands, except SANITIZE commands allowed in the restricted completion mode and commands allowed during sanitize (see 4.11.2) with CHECK CONDITION status with the sense key set to MEDIUM ERROR and the additional sense code set to SANITIZE COMMAND FAILED.

If a sanitize operation completes with an error in unrestricted completion mode, then the device server shall:

- a) terminate the SANITIZE command being performed, if any (e.g., the IMMED bit was set to zero in the CDB, and the failure occurs before status is returned for the command), with CHECK CONDITION status with the sense key set to MEDIUM ERROR and the additional sense code set to SANITIZE COMMAND FAILED; and
- b) until completion of a successful sanitize operation has occurred, terminate all commands, except SANITIZE commands and commands allowed during sanitize (see 4.11.2) with CHECK CONDITION status with the sense key set to MEDIUM ERROR and the additional sense code set to SANITIZE COMMAND FAILED.

A sanitize operation that completed with error and was cleared with a SANITIZE command with the service action of EXIT FAILURE MODE may have not performed a complete sanitize operation (e.g., this action may enable the recovery of logical block data from the cache and medium for those logical blocks that were not sanitized).

After the sanitize operation completes the device server shall:

- 1) initialize all enabled timers and counters; and
- 2) start all enabled timers and counters for power conditions and background functions.

## 4.12 Write protection

Write protection prevents the alteration of the medium by logical block access commands issued to the device server. Write protection is controlled by:

- a) the user of the medium through manual intervention (e.g., a mechanical lock on the SCSI target device);
- b) hardware controls (e.g., tabs on the medium's housing); or
- c) software write protection.

All sources of write protection are independent. If present, any write protection shall cause otherwise valid logical block access commands that request alteration of the medium to be terminated by the device server with CHECK CONDITION status with the sense key set to DATA PROTECT and the appropriate additional sense code for the condition. Only when all write protections are disabled shall the device server process logical block access commands that request alteration of the medium.

Hardware write protection results when a physical attribute of the SCSI target device or its medium is changed to specify that writing shall be prohibited. Changing the state of the hardware write protection requires physical intervention, either with the SCSI target device or its medium. If allowed by the SCSI target device,

then changing the hardware write protection while the medium is mounted results in vendor specific behavior that may include the writing of previously buffered data (e.g., data in cache).

Software write protection results when the device server is marked as write protected by the application client using the SWP bit in the Control mode page (see SPC-5). Changing the state of software write protection shall not prevent previously accepted logical block data (e.g., logical block data in cache) from being written to the medium.

The device server reports the status of write protection in the device server and on the medium with the DEVICE-SPECIFIC PARAMETER field (see 6.5.1).

## 4.13 Medium defects

### 4.13.1 Medium defects overview

Any medium has the potential for medium defects that cause data to be lost. Therefore, physical blocks and/or logical blocks may contain additional information that allows the detection of changes to the logical block data caused by medium defect or other phenomena. The additional information may also allow the logical block data to be reconstructed following the detection of such a change (e.g., ECC bytes).

A medium defect causes:

- a) a recovered error if the device server is able to read or write a logical block within the logical unit's recovery limits; or
- b) an unrecovered error if the device server is unable to read or write a logical block within the logical unit's recovery limits,

where the logical unit's recovery limits are:

- a) specified in the Read-Write Error Recovery mode page (see 6.5.10);
- b) specified in the Verify Error Recovery mode page (see 6.5.11); or
- c) vendor specific, if the device server does not implement the Read-Write Error Recovery mode page or the Verify Error Recovery mode page.

Direct access block devices may allow an application client to use the features of the WRITE LONG commands (see 5.45 and 5.46) to create a pseudo uncorrectable error. Processing and clearing pseudo uncorrectable errors is described in 4.18.2.

The device server maintains the defect lists defined in table 11.

**Table 11 — Defect lists (i.e., PLIST and GLIST)**

Defect list	Source	Content
PLIST (i.e., primary defect list)	Manufacturer	Address descriptors (see 6.2) for physical blocks that contain permanent medium defects and never contain logical block data
GLIST (i.e., grown defect list)	FORMAT UNIT commands (see 5.4)	Address descriptors for physical blocks detected by the device server to have medium defects during an optional certification process performed during a format operation
		Address descriptors for physical blocks specified in the FORMAT UNIT parameter list (see 5.4.2)
	REASSIGN BLOCKS commands (see 5.21)	Address descriptors for physical blocks referenced by the LBAs specified in the reassign LBA list (see 5.21.2)
	Read medium operations	Address descriptors for physical blocks that have been reassigned as the result of automatic read reassignment
	Write medium operations	Address descriptors for physical blocks that have been reassigned as the result of automatic write reassignment

The READ DEFECT DATA commands (see 5.19 and 5.20) allow an application client to request that the device server return the PLIST and/or the GLIST.

The FORMAT UNIT command allows an application client to request that the device server clear the GLIST.

During a format operation, the device server shall not assign LBAs to any physical block in:

- a) the PLIST, if the PLIST is specified to be used; or
- b) the GLIST, if the GLIST is specified to be used.

A device server performs automatic reassignment of defects as specified by the settings in the Read-Write Error Recovery mode page (see 6.5.10).

The device server does not perform automatic read reassignment for an LBA referencing a logical block on which an unrecovered error has occurred. If the application client is notified by the device server that an unrecovered error occurred (e.g., as indicated by a read command being terminated with CHECK CONDITION status with the sense key set to MEDIUM ERROR and the additional sense code set to UNRECOVERED READ ERROR) and:

- a) the application client is able to regenerate the logical block data for the LBA (e.g., in a redundancy group (see 4.19.1), the application client regenerates logical block data from the logical block data on the other logical units in the redundancy group) and the AWRE bit is set to one in the Read-Write Error Recovery mode page, then the application client may send a write command with that regenerated logical block data to trigger automatic write reassignment;
- b) the application client is able to regenerate the logical block data for the LBA and the AWRE bit is set to zero in the Read-Write Error Recovery mode page, then the application client may:
  - 1) send a REASSIGN BLOCKS command to perform a reassign operation on the LBA; and
  - 2) send a write command with that regenerated logical block data;
- or
- c) the application client is unable to regenerate the logical block data for the LBA, then the application client may send a REASSIGN BLOCKS command to request that the device server perform a reassign operation on the LBA.



#### 4.13.2 Generation of defect lists

This standard defines address descriptor formats for describing defects (see 6.2). Table 12 lists the defects that each address descriptor format is capable of describing.

**Table 12 — Address descriptor formats**

Format	Single physical block	Multiple sequential physical blocks		Reference
		Entire track	Range	
Short block format	yes	no	no	6.2.2
Extended bytes from index format	yes <sup>a</sup>	yes <sup>e</sup>	yes <sup>i</sup>	6.2.3
Extended physical sector format	yes <sup>b</sup>	yes <sup>f</sup>	yes <sup>i</sup>	6.2.4
Long block format	yes	no	no	6.2.5
Bytes from index format	yes <sup>c</sup>	yes <sup>g</sup>	no	6.2.6
Physical sector format	yes <sup>d</sup>	yes <sup>h</sup>	no	6.2.7
<sup>a</sup> Describes a single physical block with the MADS bit set to zero and the BYTES FROM INDEX field set to a value other than FFF_FFFFh. <sup>b</sup> Describes a single physical block with the MADS bit set to zero and the SECTOR NUMBER field set to a value other than FFF_FFFFh. <sup>c</sup> Describes a single physical block with the BYTES FROM INDEX field set to a value other than FFFF_FFFFh. <sup>d</sup> Describes a single physical block with the SECTOR NUMBER field set to a value other than FFFF_FFFFh. <sup>e</sup> Describes an entire track with the BYTES FROM INDEX field set to FFF_FFFFh. <sup>f</sup> Describes an entire track with the SECTOR NUMBER field set to FFF_FFFFh. <sup>g</sup> Describes an entire track with the BYTES FROM INDEX field set to FFFF_FFFFh. <sup>h</sup> Describes an entire track with the SECTOR NUMBER field set to FFFF_FFFFh. <sup>i</sup> Describes a range with a pair of address descriptors using the same address descriptor format in which: a) the first address descriptor describes the starting location and has the MADS bit set to one; b) the second address descriptor describes the ending location and has the MADS bit set to zero; and c) the ending location is after the starting location.				

For a direct access block device using rotating media (see 4.3.2), to represent two or more sequential physical blocks on the same track using a pair of address descriptors:

- the MADS bit shall be set to one in the first address descriptor;
- the MADS bit shall be set to zero in the second address descriptor;
- the CYLINDER NUMBER field in the first address descriptor shall be equal to the CYLINDER NUMBER field in the second address descriptor;
- the HEAD NUMBER field in the first address descriptor shall be equal to the HEAD NUMBER field in the second address descriptor;
- for a pair of extended bytes from index format address descriptors, the BYTES FROM INDEX field in the first address descriptor shall be less than the BYTES FROM INDEX field in the second address descriptor; and
- for a pair of extended physical sector format address descriptors, the SECTOR NUMBER field in the first address descriptor shall be less than the SECTOR NUMBER field in the second address descriptor.

For a direct access block device using rotating media, to represent two or more sequential tracks on the same head using a pair of address descriptors:

- the MADS bit shall be set to one in the first address descriptor;
- the MADS bit shall be set to zero in the second address descriptor;

- c) the CYLINDER NUMBER field in the first address descriptor shall be less than the CYLINDER NUMBER field in the second address descriptor;
- d) the HEAD NUMBER field in the first address descriptor shall be equal to the HEAD NUMBER field in the second address descriptor;
- e) for a pair of extended bytes from index format address descriptors, the BYTES FROM INDEX field in the first address descriptor and the second address descriptor shall be equal to FFF\_FFFFh; and
- f) for a pair of extended physical sector format address descriptors, the SECTOR NUMBER field in the first address descriptor and the second address descriptor shall be equal to FFF\_FFFFh.

## 4.14 Write and unmap failures

If one or more write commands have not completed when a power loss, a medium error, or hardware error occurs (e.g., a removable medium was incorrectly demounted), then any data in the logical blocks referenced by the LBAs specified by any of those commands is indeterminate. Before sending a read command or verify command specifying any LBAs that were specified by one of the write commands that did not complete, the application client should resend that write command. If an application client sends a read command or verify command specifying any LBAs that were specified by one of the write commands that did not complete before resending that write command, then the device server may return old data, new data, vendor-specific data, or any combination thereof for the logical blocks referenced by the specified LBAs.

If logical block provisioning (see 4.7) is supported and one or more unmap commands have not completed when a power loss, medium error, or hardware error occurs, then the logical block provisioning state of the LBAs requested to be unmapped by any of those commands may or may not have changed. ~~Before sending a read command or verify command specifying any LBAs that were specified by one of the unmap commands that did not complete, the application client should resend that unmap command.~~

## 4.15 Caches

### 4.15.1 Caches overview

Direct access block devices may implement caches. A cache is an area of temporary storage in the direct access block device (e.g., to enhance performance) separate from the medium that is not directly accessible by the application client.

A cache stores logical block data.

A cache may be volatile or non-volatile. A volatile cache does not retain data through power cycles. A non-volatile cache retains data through power cycles. There may be a limit on the amount of time a non-volatile cache is able to retain data without power (see 4.15.10).

### 4.15.2 Cache segments

The cache may be divided into cache segments. The device server may allow application client control over cache segments using the following in the Caching mode page (see 6.5.6):

- a) the IC bit;
- b) the NUMBER OF CACHE SEGMENTS field;
- c) the CACHE SEGMENT SIZE field; and
- d) the SIZE bit.

Cache segments may be grouped where each group is identified by a cache ID. An application client may request that a specific group of cache segments be used for an IO through the use of IO advice hints (see 4.32). The GROUP NUMBER field in each command specifies the cache ID of the cache segments to use during the processing of that command. The algorithms used to manage the data in each cache segment is outside the scope of this standard. How cache segments are associated with cache IDs is outside the scope of this standard.

#### 4.15.3 Read caching

While processing read commands and verify commands, the device server may use the cache to store logical blocks that the application client may request at some future time. The algorithm used to manage the cache is not part of this standard. However, parameters are provided (see 6.5.6) to advise the device server about future requests, or to restrict the use of cache for a particular request.

#### 4.15.4 Write caching

While processing write commands, the device server may perform a write cache operation to store logical block data that is to be written to the medium at a later time with a write medium operation. This is called writeback caching. A write command may complete prior to logical blocks being written to the medium. As a result of using writeback caching there is a period of time during which the logical block data may be lost if:

- a) power to the SCSI target device is lost and a volatile cache is being used; or
- b) a hardware failure occurs.

There is also the possibility of an error occurring during the subsequent write medium operation. If an error occurs during the write medium operation, then the error may be reported as a deferred error on a later command. The application client may request that writeback caching be disabled with the Caching mode page (see 6.5.6) to prevent detected write errors from being reported as deferred errors. Even with writeback caching disabled, undetected write errors may occur. Verify commands (e.g., VERIFY and WRITE AND VERIFY) may be used to detect those errors.

If processing a write command results in logical block data in cache that is different from the logical block data on the medium, then the device server shall retain that logical block data in cache until a write medium operation is performed using that logical block data. After the write medium operation is complete, the device server may retain that logical block data in cache.

#### 4.15.5 Command interactions with caches

The application client may affect behavior of the cache with:

- a) the PRE-FETCH commands (see 5.10 and 5.11);
- b) the SYNCHRONIZE CACHE commands (see 5.28 and 5.29); and
- c) the Caching mode page (see 6.5.6).

When the cache becomes full of logical block data, the device server may replace the logical block data in the cache with new logical block data. The disable page out (DPO) bit in the CDBs of read commands, verify commands, and write commands allows the application client to influence the replacement of logical block data in the cache. A read command, verify command, or a write command with a DPO bit set to one is a hint to the device server that the logical blocks specified by that command are not likely to be accessed in the near future and should not be put in the cache or retained by the cache.

Application clients may use the force unit access (FUA) bit in the CDBs of read commands (e.g., see 5.13) or write commands (e.g., see 5.35) to specify that the device server shall access:

- a) the medium;
- b) non-volatile cache; or
- c) the specified data pattern for that LBA (e.g., the data pattern for unmapped data (see 4.7.4.4)).

Setting the DPO bit to one (e.g., see 5.13) and the FUA bit to one in all read commands and all write commands has the same effect as bypassing the volatile cache.

#### 4.15.6 Write operation and write medium operation interactions with caches

For each LBA accessed by a write operation:

- 1) if a cache contains more recent logical block data for that LBA than the medium, then the device server shall:
  - A) perform a write cache operation to that LBA to update the logical block data in the cache;

- B) invalidate that LBA in the cache and perform a write medium operation to that LBA; or
- C) perform a write cache operation to that LBA to update the logical block data in the cache and perform a write medium operation to that LBA.

For each LBA accessed by a write medium operation that is not part of a write operation:

- 1) if a cache contains more recent logical block data for that LBA than the medium, then the device server shall:
  - A) perform a write cache operation to that LBA to update the logical block data in the cache; or
  - B) invalidate that LBA in the cache, before the device server performs the write medium operation to that LBA.

#### 4.15.7 Read operation and read medium operation interactions with caches

For each LBA accessed by a read operation:

- 1) if a cache contains more recent logical block data for that LBA than the medium, then the device server shall perform a read cache operation from that LBA; or
- 2) the device server shall perform a read medium operation from that LBA.

For each LBA accessed by a read medium operation that is not part of a read operation:

- 1) if a cache contains more recent logical block data for the LBA than the medium, then the device server shall perform a write medium operation to that LBA; and
- 2) the device server may invalidate that LBA in the cache, before the device server performs the read medium operation from that LBA.

#### 4.15.8 Verify medium operation interactions with caches

For each LBA accessed by a verify medium operation:

- 1) if a cache contains more recent logical block data for the LBA than the medium, then the device server shall perform a write medium operation to that LBA;
- 2) the device server may invalidate that LBA in the cache; and
- 3) before the device server performs the verify medium operation from that LBA.

#### 4.15.9 Unmap operation interactions with caches

During an unmap operation, the device server changes any logical block data in the cache for the LBA unmapped by the operation so that any logical block data transferred by the device server to the Data-In Buffer while processing a subsequent read command reflects the results of the unmap operation (see 4.7.4.6.1 and 4.7.4.7.1).

#### 4.15.10 Power loss effects on caches

The power, if any, needed to maintain a non-volatile cache may decrease to the point that the device server is unable to ensure the non-volatility of the cache for a vendor specific interval of time (e.g., the battery voltage becomes too low to sustain cache contents beyond a vendor specific time). If this occurs and the Extended INQUIRY Data VPD page (see SPC-5) indicates that the device server contains non-volatile cache (i.e., NV\_SUP bit set to one), then:

- a) if the reporting of informational exceptions control warnings is enabled (i.e., the EWASC bit is set to one in the Information Exceptions Control mode page (see 6.5.8)), then the device server shall report the degraded non-volatile cache as specified in the Information Exceptions Control mode page with an additional sense code set to WARNING - DEGRADED POWER TO NON-VOLATILE CACHE; or
- b) if the reporting of informational exceptions control warnings is disabled (i.e., the EWASC bit is set to zero in the Information Exceptions Control mode page), then the device server shall establish a unit attention condition (see SAM-5) for the SCSI initiator port associated with every I\_T nexus with the additional sense code set to WARNING - DEGRADED POWER TO NON-VOLATILE CACHE.

Non-volatile caches may become volatile (e.g., battery voltage becomes too low to sustain cache contents when power is lost). If non-volatile caches become volatile, then logical block data transferred for read commands or write commands in which the force unit access (FUA) bit in the CDB is set to one may bypass the cache.

If a non-volatile cache becomes volatile, then the device server shall set the REMAINING NON-VOLATILE TIME field to zero in the Non-volatile Cache log page (see 6.4.7).

If non-volatile cache becomes volatile and the Extended INQUIRY Data VPD page (see SPC-5) indicates that the device server contains non-volatile cache (i.e., the NV\_SUP bit is set to one), then:

- a) if the reporting of informational exceptions control warnings is enabled (i.e., the EWASC bit is set to one in the Information Exceptions Control mode page (see 6.5.8)), then the device server shall report the change in the cache as specified in the Information Exceptions Control mode page with the additional sense code set to WARNING - NON-VOLATILE CACHE NOW VOLATILE; or
- b) if the reporting of informational exceptions control warnings is disabled (i.e., the EWASC bit is set to zero in the Information Exceptions Control mode page), then the device server shall establish a unit attention condition (see SAM-5) for the SCSI initiator port associated with every I\_T nexus with the additional sense code set to WARNING - NON-VOLATILE CACHE NOW VOLATILE.

If:

- a) a power on or hard reset occurs;
- b) the Extended INQUIRY Data VPD page indicates that the device server contains a non-volatile cache (i.e., the NV\_SUP bit is set to one); and
- c) the non-volatile cache is currently volatile,

then the device server shall establish a unit attention condition for the SCSI initiator port associated with every I\_T nexus with the additional sense code set to WARNING - NON-VOLATILE CACHE NOW VOLATILE.

## 4.16 Implicit head of queue command processing

Each of the following commands defined by this standard may be processed by the task manager as if it has a HEAD OF QUEUE task attribute (see SAM-5), even if the command is received with a SIMPLE task attribute or an ORDERED task attribute:

- a) the READ CAPACITY (10) command (see 5.17); and
- b) the READ CAPACITY (16) command (see 5.18).

The following command defined by this standard shall be processed by the task manager as if it has a HEAD OF QUEUE task attribute, even if the command is received with a SIMPLE task attribute or an ORDERED task attribute:

- a) the SANITIZE command (see 5.25).

See SPC-5 for additional commands subject to implicit HEAD OF QUEUE command processing. See SAM-5 for additional rules on implicit head of queue processing.

## 4.17 Reservations

Reservation restrictions are placed on commands as a result of access qualifiers associated with the type of reservation. See SPC-5 for a description of reservations. The details of commands that are allowed under what types of reservations are described in table 13.

Commands from I\_T nexuses holding a reservation should complete normally. Table 13 specifies the behavior of commands from registered I\_T nexuses when a registrants only or all registrants type persistent reservation is present.

For each command, this standard or SPC-5 defines the conditions that result in the device server completing the command with RESERVATION CONFLICT status.

**Table 13 — SBC-4 commands that are allowed in the presence of various reservations (part 1 of 3)**

Command	Addressed logical unit has this type of persistent reservation held by another I_T nexus				
	From any I_T nexus		From registered I_T nexus (RR all types)	From I_T nexus not registered	
	Write Exclusive	Exclusive Access		Write Exclusive - RR	Exclusive Access - RR
BACKGROUND CONTROL	Conflict	Conflict	Allowed	Conflict	Conflict
COMPARE AND WRITE	Conflict	Conflict	Allowed	Conflict	Conflict
FORMAT UNIT	Conflict	Conflict	Allowed	Conflict	Conflict
GET LBA STATUS	Allowed	Conflict	Allowed	Allowed	Conflict
GET STREAM STATUS	Allowed	Conflict	Allowed	Allowed	Conflict
ORWRITE (16) / (32)	Conflict	Conflict	Allowed	Conflict	Conflict
POPULATE TOKEN	Allowed	Conflict	Allowed	Allowed	Conflict
PRE-FETCH (10) / (16)	Allowed	Conflict	Allowed	Allowed	Conflict
PREVENT ALLOW MEDIUM REMOVAL (Prevent=0)	Allowed	Allowed	Allowed	Allowed	Allowed
PREVENT ALLOW MEDIUM REMOVAL (Prevent ≠ 0)	Conflict	Conflict	Allowed	Conflict	Conflict
READ (10) / (12) / (16) / (32)	Allowed	Conflict	Allowed	Allowed	Conflict
<p>Key:</p> <p>RR = Registrants Only or All Registrants</p> <p>Allowed = Commands received from I_T nexuses not holding the reservation or from I_T nexuses not registered when a registrants only or all registrants type persistent reservation is present should complete normally.</p> <p>Conflict = Commands received from I_T nexuses not holding the reservation or from I_T nexuses not registered when a registrants only or all registrants type persistent reservation is present shall not be performed, and the device server shall complete the command with RESERVATION CONFLICT status.</p>					
<p><sup>a</sup> The device server in logical units claiming compliance with SBC-2 may complete the command with RESERVATION CONFLICT status. Device servers may report whether certain commands are allowed in the PERSISTENT RESERVE IN command REPORT CAPABILITIES service action parameter data ALLOW COMMANDS field (see SPC-5).</p>					

Table 13 — SBC-4 commands that are allowed in the presence of various reservations (part 2 of 3)

Command	Addressed logical unit has this type of persistent reservation held by another I_T nexus				
	From any I_T nexus		From registered I_T nexus (RR all types)	From I_T nexus not registered	
	Write Exclusive	Exclusive Access		Write Exclusive - RR	Exclusive Access - RR
READ CAPACITY (10) / (16)	Allowed	Allowed	Allowed	Allowed	Allowed
READ DEFECT DATA (10) / (12)	Allowed <sup>a</sup>	Conflict	Allowed	Allowed <sup>a</sup>	Conflict
REASSIGN BLOCKS	Conflict	Conflict	Allowed	Conflict	Conflict
REPORT REFERRALS	Allowed	Allowed	Allowed	Allowed	Allowed
REPORT PROVISIONING INITIALIZATION PATTERN	Allowed	Allowed	Allowed	Allowed	Allowed
SANITIZE	Conflict	Conflict	Allowed	Conflict	Conflict
START STOP UNIT with START bit set to one and POWER CONDITION field set to 0h	Allowed	Allowed	Allowed	Allowed	Allowed
START STOP UNIT with START bit set to zero or POWER CONDITION field set to a value other than 0h	Conflict	Conflict	Allowed	Conflict	Conflict
STREAM CONTROL	Conflict	Conflict	Allowed	Conflict	Conflict
SYNCHRONIZE CACHE (10) / (16)	Conflict	Conflict	Allowed	Conflict	Conflict
UNMAP	Conflict	Conflict	Allowed	Conflict	Conflict
VERIFY (10) / (12) / (16) / (32)	Allowed	Conflict	Allowed	Allowed	Conflict
WRITE (10) / (12) / (16) / (32)	Conflict	Conflict	Allowed	Conflict	Conflict
WRITE AND VERIFY (10) / (12) / (16) / (32)	Conflict	Conflict	Allowed	Conflict	Conflict
WRITE ATOMIC (16) / (32)	Conflict	Conflict	Allowed	Conflict	Conflict
WRITE LONG (10) / (16)	Conflict	Conflict	Allowed	Conflict	Conflict
WRITE SAME (10) / (16) / (32)	Conflict	Conflict	Allowed	Conflict	Conflict
WRITE STREAM (16) / (32)	Conflict	Conflict	Allowed	Conflict	Conflict
<p>Key:</p> <p>RR = Registrants Only or All Registrants</p> <p>Allowed = Commands received from I_T nexuses not holding the reservation or from I_T nexuses not registered when a registrants only or all registrants type persistent reservation is present should complete normally.</p> <p>Conflict = Commands received from I_T nexuses not holding the reservation or from I_T nexuses not registered when a registrants only or all registrants type persistent reservation is present shall not be performed, and the device server shall complete the command with RESERVATION CONFLICT status.</p>					
<p><sup>a</sup> The device server in logical units claiming compliance with SBC-2 may complete the command with RESERVATION CONFLICT status. Device servers may report whether certain commands are allowed in the PERSISTENT RESERVE IN command REPORT CAPABILITIES service action parameter data ALLOW COMMANDS field (see SPC-5).</p>					

**Table 13 — SBC-4 commands that are allowed in the presence of various reservations (part 3 of 3)**

Command	Addressed logical unit has this type of persistent reservation held by another I_T nexus				
	From any I_T nexus		From registered I_T nexus (RR all types)	From I_T nexus not registered	
	Write Exclusive	Exclusive Access		Write Exclusive - RR	Exclusive Access - RR
WRITE STREAM (16) / (32)	Conflict	Conflict	Allowed	Conflict	Conflict
WRITE USING TOKEN	Conflict	Conflict	Allowed	Conflict	Conflict
XDWRITEREAD (10) / (32)	Conflict	Conflict	Allowed	Conflict	Conflict
XPWRITE (10) / (32)	Conflict	Conflict	Allowed	Conflict	Conflict
<p>Key:</p> <p>RR = Registrants Only or All Registrants</p> <p>Allowed = Commands received from I_T nexuses not holding the reservation or from I_T nexuses not registered when a registrants only or all registrants type persistent reservation is present should complete normally.</p> <p>Conflict = Commands received from I_T nexuses not holding the reservation or from I_T nexuses not registered when a registrants only or all registrants type persistent reservation is present shall not be performed, and the device server shall complete the command with RESERVATION CONFLICT status.</p>					
<p><sup>a</sup> The device server in logical units claiming compliance with SBC-2 may complete the command with RESERVATION CONFLICT status. Device servers may report whether certain commands are allowed in the PERSISTENT RESERVE IN command REPORT CAPABILITIES service action parameter data ALLOW COMMANDS field (see SPC-5).</p>					

## 4.18 Error reporting

### 4.18.1 Error reporting overview

If any of the conditions listed in table 14 occur during the processing of a command, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to the specified value and the additional sense code set to the appropriate value for the condition. Some errors may occur after the completion status has already been reported. For such errors, SPC-5 defines a deferred error reporting mechanism. Table 14 lists some error conditions and the applicable sense keys. The list does not provide a complete list of all conditions that may cause CHECK CONDITION status.



**Table 14 — Example error conditions**

<b>Condition</b>	<b>Sense key</b>
Invalid LBA	ILLEGAL REQUEST
Unsupported option requested	ILLEGAL REQUEST
Logical unit reset, I_T nexus loss, or medium change since last command from this application client	UNIT ATTENTION
Logical block provisioning threshold notification	UNIT ATTENTION
Self diagnostic failed	HARDWARE ERROR
Unrecovered error	MEDIUM ERROR or HARDWARE ERROR
Recovered read error	RECOVERED ERROR
Pseudo unrecovered error	MEDIUM ERROR
Over-run or other error that may be resolved by repeating the command	ABORTED COMMAND
Attempt to write on write-protected medium	DATA PROTECT

Direct access block devices compliant with this standard shall support both the fixed and descriptor formats of sense data (see SPC-5).

Table 15 summarizes use of the sense data fields.

**Table 15 — Sense data field usage for direct access block devices**

Field	Usage	Reference
INFORMATION field <sup>a</sup>		
	REASSIGN BLOCKS command	5.21
	Any command that accesses the medium, based on the Read-Write Error Recovery mode page	4.20.3 and 6.5.10
	Any command that accesses the medium, based on the Verify Error Recovery mode page	6.5.11
	Any command that is terminated with a logical block provisioning threshold notification	4.7.3.7.6
	COMPARE AND WRITE command	5.3
COMMAND-SPECIFIC INFORMATION field	EXTENDED COPY command	SPC-5
	REASSIGN BLOCKS command	5.21
	If rebuild assist mode is enabled (see 4.20), then any command that accesses the medium, based on the Read-Write Error Recovery mode page <sup>b</sup>	4.20.3
<sup>a</sup> See SPC-5 for a description of how the VALID bit interacts with the INFORMATION field. <sup>b</sup> If fixed format sense data is used but the value to be placed in the COMMAND-SPECIFIC INFORMATION field is greater than FFFF_FFFFh (e.g., an 8-byte LBA), then the device server shall report no value in the INFORMATION field (see SPC-5) and shall report no value in the COMMAND-SPECIFIC INFORMATION field (see SPC-5).		

If a command attempts to access or reference an invalid LBA, then the device server shall report the first invalid LBA (e.g., lowest numbered LBA) in the INFORMATION field of the sense data (see SPC-5).

If a recovered read error is reported, then the device server shall report the last LBA (e.g., highest numbered LBA) on which a recovered read error occurred for the command in the INFORMATION field of the sense data.

If an unrecovered error is reported, then the device server shall report the LBA of the logical block on which an unrecovered error occurred in the INFORMATION field of the sense data.

#### 4.18.2 Processing pseudo unrecovered errors

If a pseudo unrecovered error with correction disabled is encountered on a logical block (e.g., by a command, a background scan(see 4.24.1), or a background self-test (see SPC-5)), then the device server shall:

- perform no error recovery on the affected logical blocks, including any read error recovery enabled by the Read-Write Error Recovery mode page (see 6.5.10) or the Verify Error Recovery mode page (see 6.5.11);
- perform no automatic read reassignment or automatic write reassignment for the affected logical blocks, regardless of the settings of the AWRE bit and the ARRE bit in the Read-Write Error Recovery mode page;
- not consider errors on the affected logical blocks to be informational exception conditions as defined in the Information Exceptions Control mode page (see 6.5.8);

- d) not log errors on the affected logical blocks in any log page that contain error counters (see SPC-5); and
- e) in any information returned for the error (e.g., in sense data or in the Background Scan Results log page (see 6.4.2)), set the sense key to MEDIUM ERROR and either:
  - A) should set the additional sense code to READ ERROR – LBA MARKED BAD BY APPLICATION CLIENT; or
  - B) may set the additional sense code to UNRECOVERABLE READ ERROR.

The logical unit shall clear a pseudo unrecovered error if it processes or performs one of the following for that LBA:

- a) a format operation;
- b) a reassign operation;
- c) a sanitize overwrite operation;
- d) a sanitize block erase operation; or
- e) a write command that is not a WRITE LONG command specifying a pseudo unrecovered error.

The logical unit may clear a pseudo unrecovered error if it processes or performs one of the following for that LBA:

- a) a sanitize cryptographic erase operation;
- b) an unmap operation; or
- c) a MODE SELECT command that uses the mode parameter block descriptor (see 6.5.2) to change the capacity to be lower than that LBA.

The logical unit shall not clear a pseudo unrecovered error if it processes one of the following for that LBA:

- a) a read command.

#### 4.18.3 Block commands sense data descriptor

Table 16 defines the block commands sense data descriptor used in descriptor format sense data (see SPC-5) for direct access block devices.

**Table 16 — Block commands sense data descriptor format**

Byte	Bit	7	6	5	4	3	2	1	0
0		DESCRIPTOR TYPE (05h)							
1		ADDITIONAL LENGTH (02h)							
2		Reserved							
3		Reserved		Obsolete	Reserved				

The DESCRIPTOR TYPE field and the ADDITIONAL LENGTH field are defined in SPC-5 and shall be set to the values shown in table 16 for the block commands sense data descriptor.

#### 4.18.4 User data segment referral sense data descriptor

Table 17 defines the user data segment referral sense data descriptor used in descriptor format sense data for direct access block devices. The user data segment referral sense data descriptor contains descriptors indicating the user data segment(s) on the logical unit and the SCSI target port groups through which those user data segments may be accessed (see 4.28).

Table 17 — User data segment referral sense data descriptor format

Byte	Bit	7	6	5	4	3	2	1	0
0	DESCRIPTOR TYPE (0Bh)								
1	ADDITIONAL LENGTH (y -1)								
2	Reserved								NOT_ALL_R
3	Reserved								
User data segment referral descriptor list									
4	User data segment referral descriptor [first]								
...									
4 + n									
⋮									
y - m	User data segment referral descriptor [last]								
...									
y									

The DESCRIPTOR TYPE field is defined in SPC-5 and shall be set to the value shown in table 17 for the user data segment referral sense data descriptor.

The ADDITIONAL LENGTH field indicates the number of bytes that follow in the logical block referrals sense data descriptor.

A not all referrals (NOT\_ALL\_R) bit set to zero indicates that the list of user data segment referral descriptors is a complete list of user data segments. A NOT\_ALL\_R bit set to one indicates that there are more user data segments than are able to be indicated by the user data segment referral sense data.

Each user data segment referral descriptor (see table 18) indicates information identifying:

- a user data segment that is accessible through the SCSI target port groups indicated by this descriptor; and
- one or more SCSI target port groups through which the user data segment indicated by this descriptor is able to be accessed.

User data segment referral descriptors shall be listed in ascending LBA order. If a user data segment referral descriptor describes the last user data segment (i.e., points to the largest LBA) and the preceding user data segment descriptors do not represent the complete list of user data segments, then the next user data segment referral descriptor, if any, shall describe the first user data segment (i.e., the user data segments may wrap).

Table 18 defines the user data segment referral descriptor.

**Table 18 — User data segment referral descriptor format**

Byte	Bit	7	6	5	4	3	2	1	0
0		Reserved							
...									
2									
3		NUMBER OF TARGET PORT GROUP DESCRIPTORS							
4	(MSB)	FIRST USER DATA SEGMENT LBA							
...									
11		(LSB)							
12	(MSB)	LAST USER DATA SEGMENT LBA							
...									
19		(LSB)							
Target port group descriptor list									
20		Target port group descriptor [first]							
...									
23									
⋮									
m-3		Target port group descriptor [last]							
...									
m									

The NUMBER OF TARGET PORT GROUP DESCRIPTORS field indicates the number of target port group descriptors that follow.

The FIRST USER DATA SEGMENT LBA field indicates the first LBA of the first user data segment (see 4.28) indicated by this descriptor.

The LAST USER DATA SEGMENT LBA field indicates the last LBA of the last user data segment (see 4.28) indicated by this descriptor.

The target port group descriptor (see table 19) specifies the target port group and the asymmetric access state of the target port group (see SPC-5). The device server shall return one target port group descriptor for each target port group in a target port asymmetric access state of active/optimized, active/non-optimized, or transitioning. The device server may return one target port group descriptor for each target port group in a target port asymmetric access state of unavailable.

**Table 19 — Target port group descriptor**

Byte	Bit	7	6	5	4	3	2	1	0
0		Reserved				ASYMMETRIC ACCESS STATE			
1		Reserved							
2		(MSB) _____							
3		TARGET PORT GROUP _____ (LSB)							

The ASYMMETRIC ACCESS STATE field (see SPC-5) contains the asymmetric access state of the user data segment(s) specified by this descriptor that may be accessed through this target port group.

The TARGET PORT GROUP field specifies a target port group (see SPC-5) that the application client uses when issuing commands associated with the user data segments specified by this descriptor.

#### **4.18.5 Direct access block device sense data descriptor**

Table 20 defines the direct access block device sense data descriptor, which may be used in descriptor format sense data (see SPC-5) instead of any of the following sense data descriptors:

- a) information (see SPC-5);
- b) command-specific information (see SPC-5);
- c) sense key specific (see SPC-5);
- d) field replaceable unit (see SPC-5); and
- e) block commands (see 4.18.3).

If the device server includes the direct access block device sense data descriptor in a sense data descriptor list, then it shall not include any of those sense data descriptors in the same sense data descriptor list.

**Table 20 — Direct access block device sense data descriptor format**

Byte	Bit	7	6	5	4	3	2	1	0
0	DESCRIPTOR TYPE (0Dh)								
1	ADDITIONAL LENGTH (1Eh)								
2	VALID	Reserved	Obsolete	Reserved					
3	Reserved								
4	SKSV	Sense key specific information							
...									
6									
7	FIELD REPLACEABLE UNIT CODE								
8	(MSB)	INFORMATION							
...									
15									
16	(MSB)	COMMAND-SPECIFIC INFORMATION							
...									
23	(LSB)								

The DESCRIPTOR TYPE field is described in SPC-5, and shall be set as shown in table 20 for the direct access block device sense data descriptor.

A VALID bit set to zero indicates that the INFORMATION field is not defined in this standard or any other command standard. A VALID bit set to one indicates the INFORMATION field contains valid information as defined in this standard or a command standard.

A sense-key specific valid (SKSV) bit set to one indicates the sense key specific information contains valid information as defined in SPC-5. An SKSV bit set to zero indicates that the sense key specific information is not as defined by SPC-5.

The sense key specific information is described in the sense key specific sense data descriptor (see SPC-5).

The FIELD REPLACEABLE UNIT CODE field is described in the field replaceable unit sense data descriptor (see SPC-5).

The INFORMATION field is described in the information sense data descriptor (see SPC-5).

The COMMAND-SPECIFIC INFORMATION field is described in the command-specific information sense data descriptor (see SPC-5). The COMMAND-SPECIFIC INFORMATION field should be ignored in sense data for a command or operation for which the COMMAND-SPECIFIC INFORMATION field is not defined and in sense data that is not related to a command or operation (e.g., pollable sense data).

## 4.19 Model for XOR commands

### 4.19.1 Model for XOR commands overview

In storage arrays, a SCSI storage array device (see SCC-2) organizes a group of direct access block devices into objects. The type of object supported by this model is the redundancy group, where some of the logical blocks on the direct access block devices are used for XOR-protected data and some of the logical blocks are used for check data. The check data is generated by performing a cumulative XOR operation of the

XOR-protected data. The XOR operation may be performed by the SCSI storage array device or by the direct access block devices.

A direct access block device containing XOR-protected data is called a data disk. A direct access block device containing check data is called a parity disk.

Performing the XOR operation in the direct access block devices may result in a reduced number of data transfers across a service delivery subsystem.

For example, when the XOR operation is performed within the SCSI storage array device, the following commands are used for a typical update write sequence:

- 1) a read command to the data disk;
- 2) a write command to the data disk;
- 3) a read command to the parity disk; and
- 4) a write command to the parity disk.

The SCSI storage array device also does two internal XOR operations in this sequence.

In contrast, during SCSI storage array device supervised XOR operations (see 4.19.2) only the following operations are used:

- a) a write command to the data disk;
- b) a read command to the data disk; and
- c) a write command to the parity disk.

#### **4.19.2 SCSI storage array device supervised XOR operations**

##### **4.19.2.1 SCSI storage array device supervised XOR operations overview**

A SCSI storage array device (see SCC-2) supervises three basic operations that require XOR functionality:

- a) update write operation (see 4.19.2.2);
- b) regenerate operation (see 4.19.2.3); and
- c) rebuild operation (see 4.19.2.4).

Command sequences for each of these operations use the device servers in the direct access block devices to perform the necessary XOR functions.

XDWRITEREAD commands (see 5.53 and 5.54) and XPWRITE commands (see 5.55 and 5.56) are required to implement SCSI storage array device supervised XOR operations. The SCSI storage array device also uses READ commands and WRITE commands for certain operations.

##### **4.19.2.2 Update write operation**

The update write operation writes new XOR-protected data to a data disk and updates the check data on the parity disk. The sequence performed by the SCSI storage array device (see SCC-2) is as follows:

- 1) send an XDWRITEREAD command to the data disk with the Data-Out Buffer containing the new XOR-protected data; and
- 2) send an XPWRITE command to the parity disk with the Data-Out Buffer for this command containing the logical block data from the Data-In Buffer of the XDWRITEREAD command.

##### **4.19.2.3 Regenerate operation**

The regenerate operation is used to recreate one or more logical blocks that have an error. This is accomplished by reading the associated logical block from each of the other direct access block devices within the redundancy group and performing an XOR operation on each of these logical blocks. The last XOR result is the data that should have been present on the unreadable direct access block device. The number of steps is dependent on the number of direct access block devices in the redundancy group, but the sequence performed by the SCSI storage array device (see SCC-2) is as follows:

- 1) send a READ command to the first direct access block device;



- 2) send an XDWRITEREAD command with the DISABLE WRITE bit set to one to the next direct access block device using the logical block data from the Data-In Buffer of the previous command as the logical block data in the Data-Out Buffer for this command; and
- 3) repeat step 2) until all direct access block devices in the redundancy group except the failed device have been accessed. The logical block data in the Data-In Buffer for the last command is the regenerated data.

#### **4.19.2.4 Rebuild operation**

The rebuild operation is similar to the regenerate operation, except that the last XOR result is written to the replacement device. This function is used when a failed device is replaced and the SCSI storage array device (see SCC-2) is writing the rebuilt data to the replacement device. The number of steps is dependent on the number of direct access block devices in the redundancy group, but the sequence performed by the SCSI storage array device (see SCC-2) is as follows:

- 1) perform the regenerate operation (see 4.19.2.3); and
- 2) send a WRITE command to the replacement device using the regenerated data.

### **4.19.3 Array subsystem considerations**

#### **4.19.3.1 Array subsystem considerations overview**

This subclause lists considerations that apply to any array subsystem and describes how use of the XOR commands may affect handling of those situations.

#### **4.19.3.2 Access to an inconsistent stripe**

A stripe is a set of corresponding extents from two or more direct access block devices.

When the SCSI storage array device (see SCC-2) issues an update write to a data disk, the data in the data disk has been updated when successful status is returned for the command. Until the parity disk has been updated the associated stripe in the redundancy group is not consistent (i.e., performing an XOR operation on the XOR-protected data does not produce the check data).

The SCSI storage array device shall keep track of this window of inconsistency and ensure that a regenerate or rebuild operation for any data extent within the stripe is not attempted until after the parity disk has been updated, making the stripe consistent again. For multi-initiator systems, tracking the updates may be more complex because each SCSI storage array device is required to ensure that a second SCSI storage array device is not writing to a stripe that the first SCSI storage array device is regenerating or rebuilding. The coordination between SCSI storage array devices is system specific and is beyond the scope of this standard.

If a device server terminates any of the XOR commands with CHECK CONDITION status and an unrecovered error is indicated, then an inconsistent stripe may result. It is the SCSI storage array device's responsibility to identify the failing device, identify the scope of the failure, and then to limit access to the inconsistent stripe. The recovery procedures that the SCSI storage array device implements are outside the scope of this standard.

## **4.20 Rebuild assist mode**

### **4.20.1 Rebuild assist mode overview**

The rebuild assist mode provides a method for a SCSI storage array device (see SCC-2) to read recovered logical blocks from a failed logical unit in a storage array instead of rebuilding the logical blocks from other logical units in the storage array. This mode allows the failed logical unit to report logical blocks that are unreadable without requiring the SCSI storage array device to read every LBA in the failed logical unit to determine the unrecovered logical blocks. The SCSI storage array device then copies the logical blocks recovered from the failed logical unit to a replacement logical unit and only rebuilds the failed logical blocks.

Enabling the rebuild assist mode:

- a) may cause the device server to initiate a self test to identify the scope of failures, if any;
- b) modifies READ command recovery behavior by the device server based on the setting of the RARC bit (see 4.20.3 and 5.13); and
- c) may cause sense data to be returned by the device server that indicates the location of multiple failing logical blocks on read commands and write commands.

The self-test operations performed by the device server while rebuild assist mode is enabled may result in detection of failed physical elements. A predicted unrecovered error is an unrecovered error that is the result of an attempt by the device server to access an LBA associated with a failed physical element. An unpredicted unrecovered error is an unrecovered error that is the result of a device server accessing an LBA that is not associated with a failed physical element.

#### **4.20.2 Enabling rebuild assist mode**

An application client should enable rebuild assist mode after the application client determines that a rebuild is required. The application client enables the rebuild assist mode by setting the ENABLED bit to one and setting the DISABLED PHYSICAL ELEMENT field to all zeros in the Rebuild Assist Output diagnostic page (see 6.3.3).

If a SEND DIAGNOSTIC command requests the enabling of the rebuild assist mode, then the device server:

- 1) shall enable the rebuild assist mode;
- 2) may perform a diagnostic test of the physical elements contained within the logical unit; and
- 3) should disable any physical elements that are not functional if a diagnostic test of the physical elements is performed.

The application client may verify that rebuild assist mode is enabled by verifying that the ENABLED bit is set to one in the Rebuild Assist Input diagnostic page (see 6.3.2).

#### **4.20.3 Using the rebuild assist mode**

##### **4.20.3.1 Using rebuild assist mode overview**

After rebuild assist mode is enabled, the application client should issue read commands to read the available logical block data from the failed logical unit. If the device server does not detect an unrecovered error while processing a read command, then the device server should complete the read command without error.

The rebuild assist mode allows the device server to report unrecovered read errors or unrecovered write errors that are either predicted (i.e., predicted unrecovered errors) or unpredicted (i.e., unpredicted unrecovered errors).

##### **4.20.3.2 Unpredicted unrecovered read error**

If a device server receives a read command with the RARC bit set to one, then rebuild assist mode shall not affect processing of the read command.

If rebuild assist mode is enabled and a device server receives a read command with the RARC bit set to zero and the device server detects an unpredicted unrecovered error that is not a pseudo unrecovered read error (see 4.18.2), then the device server:

- 1) shall perform limited read recovery that is vendor specific;
- 2) shall transfer the data for all recovered logical blocks, if any, from the logical block referenced by the starting LBA of the failed read command up to the first unrecovered logical block (i.e., the lowest numbered LBA) to the Data-in Buffer;
- 3) shall terminate the command with CHECK CONDITION status with the sense key set to MEDIUM ERROR, the additional sense code set to UNRECOVERED READ ERROR, and report the LBA referencing the unrecovered logical block in the INFORMATION field (see SPC-5); and
- 4) may use this failure in a vendor specific manner to predict other logical blocks that may be unrecovered.

If the application client receives sense data with the sense key set to MEDIUM ERROR, the additional sense code set to UNRECOVERED READ ERROR, and the INFORMATION field indicating a valid LBA (see SPC-5), then the application client should issue the next read command with the starting LBA set to the contents of the INFORMATION field plus one.

#### 4.20.3.3 Predicted unrecovered read error

If the device server receives a read command with the RARC bit set to one, then rebuild assist mode shall not affect the processing of the read command.

If rebuild assist mode is enabled and the device server receives a read command with the RARC bit set to zero, and the device server detects a predicted unrecovered error, then the device server:

- 1) shall perform limited read recovery that is vendor specific;
- 2) shall transfer the data for all recovered logical blocks, if any, from the logical block referenced by the starting LBA of the failed read command up to the first predicted unrecovered LBA (i.e., the lowest numbered LBA) to the Data-in Buffer; and
- 3) shall terminate the read command with CHECK CONDITION status with the sense key set to ABORTED COMMAND, the additional sense code set to MULTIPLE READ ERRORS, and:
  - A) report the following value in the INFORMATION field (see SPC-5):
    - a) the LBA referencing the first unrecovered logical block (i.e., the lowest numbered LBA);
 and
  - B) report the following value in the COMMAND-SPECIFIC INFORMATION field (see SPC-5):
    - a) the LBA referencing the last unrecovered logical block (i.e., the highest numbered LBA) in a sequence of contiguous unrecovered logical blocks that started with the LBA indicated in the INFORMATION field.

If the application client receives sense data with the sense key set to ABORTED COMMAND, the additional sense code set to MULTIPLE READ ERRORS, and the INFORMATION field indicating a valid LBA (see SPC-5), then the application client should issue the next read command with the starting LBA set to the contents of the COMMAND-SPECIFIC INFORMATION field plus one to continue recovering data from the logical unit. This process should be repeated until all of the LBAs have been scanned.

#### 4.20.3.4 Unpredicted unrecovered write error

If rebuild assist mode is enabled and the device server detects an unpredicted unrecovered error while processing a write command, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to MEDIUM ERROR, the additional sense code set to WRITE ERROR, and report the LBA referencing the first logical block (i.e., the lowest numbered LBA) in error in the INFORMATION field (see SPC-5).

#### 4.20.3.5 Predicted unrecovered write error

If rebuild assist mode is enabled and the device server detects a predicted unrecovered error while processing a write command, then the device server:

- 1) transfers the write data from the Data-Out Buffer;
- 2) writes the transferred data up to the logical block referenced by the failing LBA; and
- 3) shall terminate the write command with CHECK CONDITION status with the sense key set to ABORTED COMMAND, the additional sense code set to MULTIPLE WRITE ERRORS, and:
  - A) report the following value in the INFORMATION field (see SPC-5):
    - a) the LBA referencing the first logical block (i.e., the lowest numbered LBA) in error;
 and
  - B) report the following value in the COMMAND-SPECIFIC INFORMATION field (see SPC-5):
    - a) the LBA referencing the last logical block (i.e., the highest numbered LBA) in error in a sequence of contiguous logical blocks that started with the LBA indicated in the INFORMATION field.

If the application client receives sense data with the sense key set to ABORTED COMMAND, the additional sense code set to MULTIPLE WRITE ERRORS, and the INFORMATION field indicating a valid LBA (see SPC-5), then the application client should issue the next write command with the starting LBA set to the contents of the COMMAND-SPECIFIC INFORMATION field plus one to continue writing to the logical unit.

#### 4.20.4 Disabling the rebuild assist mode

Rebuild assist mode shall be disabled after a power on.

Rebuild assist mode shall not be affected by a hard reset, an I\_T nexus loss, or any task management functions (see SAM-5).

The application client disables rebuild assist mode by setting the ENABLED bit to zero in the Rebuild Assist Output diagnostic page (see 6.3.3).

#### 4.20.5 Testing rebuild assist mode

The Rebuild Assist Output diagnostic page (see 6.3.3) provides a method to test the application client's rebuild process.

An application client places a logical unit into a simulated failing condition by setting the ENABLED bit to one and setting one or more bits in the DISABLED PHYSICAL ELEMENT field to one in the Rebuild Assist Output diagnostic page.

Each bit in the DISABLED PHYSICAL ELEMENT field represents a physical element that is associated with a group of LBAs that are treated as having predicted unrecovered read errors. The correlation of bits in the DISABLED PHYSICAL ELEMENT field to LBAs in the logical unit is vendor specific.

An application client ends a test by disabling the rebuild assist mode (see 4.20.4).

### 4.21 START STOP UNIT and power conditions

#### 4.21.1 START STOP UNIT and power conditions overview

The START STOP UNIT command (see 5.26) allows an application client to control the power condition of a logical unit. This method includes specifying that the logical unit transition to a specific power condition.

In addition to the START STOP UNIT command, the power condition of a logical unit may be controlled by the Power Condition mode page (see SPC-5). If both the START STOP UNIT command and the Power Condition mode page methods are being used to control the power condition of the same logical unit, then the power condition specified by any START STOP UNIT command shall override the Power Condition mode page's power control.

#### 4.21.2 Processing of concurrent START STOP UNIT commands

If a START STOP UNIT command is being processed by the device server, and a subsequent START STOP UNIT command for which the CDB is validated requests that the logical unit change to a different power condition than was specified by the START STOP UNIT command being processed, then the device server shall terminate the subsequent START STOP UNIT command with CHECK CONDITION status with the sense key set to NOT READY and the additional sense code set to LOGICAL UNIT NOT READY, START STOP UNIT COMMAND IN PROGRESS.

The constraints on concurrent START STOP UNIT commands apply only to commands that have the IMMED bit set to zero. The effects of concurrent power condition changes requested by START STOP UNIT commands with the IMMED bit set to one are vendor specific.

#### 4.21.3 Managing logical block access commands during a change to the active power condition

Application clients may minimize the return of BUSY status or TASK SET FULL status during a change to the active power condition by:

- a) polling the power condition using the REQUEST SENSE command (see SPC-5); or
- b) sending a START STOP UNIT command with the IMMED bit set to zero and the START bit set to one and waiting for GOOD status to be returned.

#### 4.21.4 Stopped power condition

In addition to the active power condition, idle power conditions, and standby power conditions described in SPC-5, this standard describes the stopped power condition.

While in the stopped power condition:

- a) the device server shall terminate TEST UNIT READY commands and logical block access commands with CHECK CONDITION status with the sense key set to NOT READY and the additional sense code set to LOGICAL UNIT NOT READY, INITIALIZING COMMAND REQUIRED;
- b) the power consumed by the SCSI target device while in the stopped power condition should be less than the power consumed when the logical unit is in the active power condition or any of the idle power conditions (e.g., for direct access block devices that have a rotating medium, the medium is stopped in the stopped power condition);
- c) the peak power consumption during a change from the stopped power condition to the active power condition or an idle power condition is not limited by this standard; and
- d) the peak power consumption during a change from the stopped power condition to a standby power condition shall be no more than the typical peak power consumption in the active power condition.

No power condition defined in this standard shall affect the supply of any power required for proper operation of a service delivery subsystem.

#### 4.21.5 START STOP UNIT and power condition state machine

##### 4.21.5.1 START STOP UNIT and power condition state machine overview

The SSU\_PC (START STOP UNIT and power condition) state machine for logical units implementing the START STOP UNIT command describes the:

- a) logical unit power states and transitions resulting from specifications in the START STOP UNIT command;
- b) settings in the Power Condition mode page (see SPC-5); and
- c) the processing of commands.

NOTE 4 - The SSU\_PC state machine is an enhanced version of the PC (power condition) state machine described in SPC-5.

The SSU\_PC state machine consists of the states shown in table 21.

**Table 21 — Summary of states in the SSU\_PC state machine**

State	Reference	PC state machine state with additional definition (see SPC-5)
SSU_PC0:Powered_On <sup>a</sup>	4.21.5.2	PC0:Powered_On
SSU_PC1:Active	4.21.5.3	PC1:Active
SSU_PC2:Idle	4.21.5.4	PC2:Idle
SSU_PC3:Standby	4.21.5.5	PC3:Standby
SSU_PC4:Active_Wait	4.21.5.6	PC4:Active_Wait
SSU_PC5:Wait_Idle	4.21.5.7	PC5:Wait_Idle
SSU_PC6:Wait_Standby	4.21.5.8	PC6:Wait_Standby
SSU_PC7:Idle_Wait	4.21.5.9	n/a
SSU_PC8:Stopped	4.21.5.10	n/a
SSU_PC9:Standby_Wait	4.21.5.11	n/a
SSU_PC10:Wait_Stopped	4.21.5.12	n/a
<sup>a</sup> SSU_PC0:Powered_On is the initial state.		

While in the following SSU\_PC states, the logical unit may be increasing power usage to enter a higher power condition:

- a) SSU\_PC4:Active\_Wait;
- b) SSU\_PC7:Idle\_Wait; and
- c) SSU\_PC9:Standby\_Wait.

While in the following SSU\_PC states, the logical unit may be decreasing power usage to enter a lower power condition:

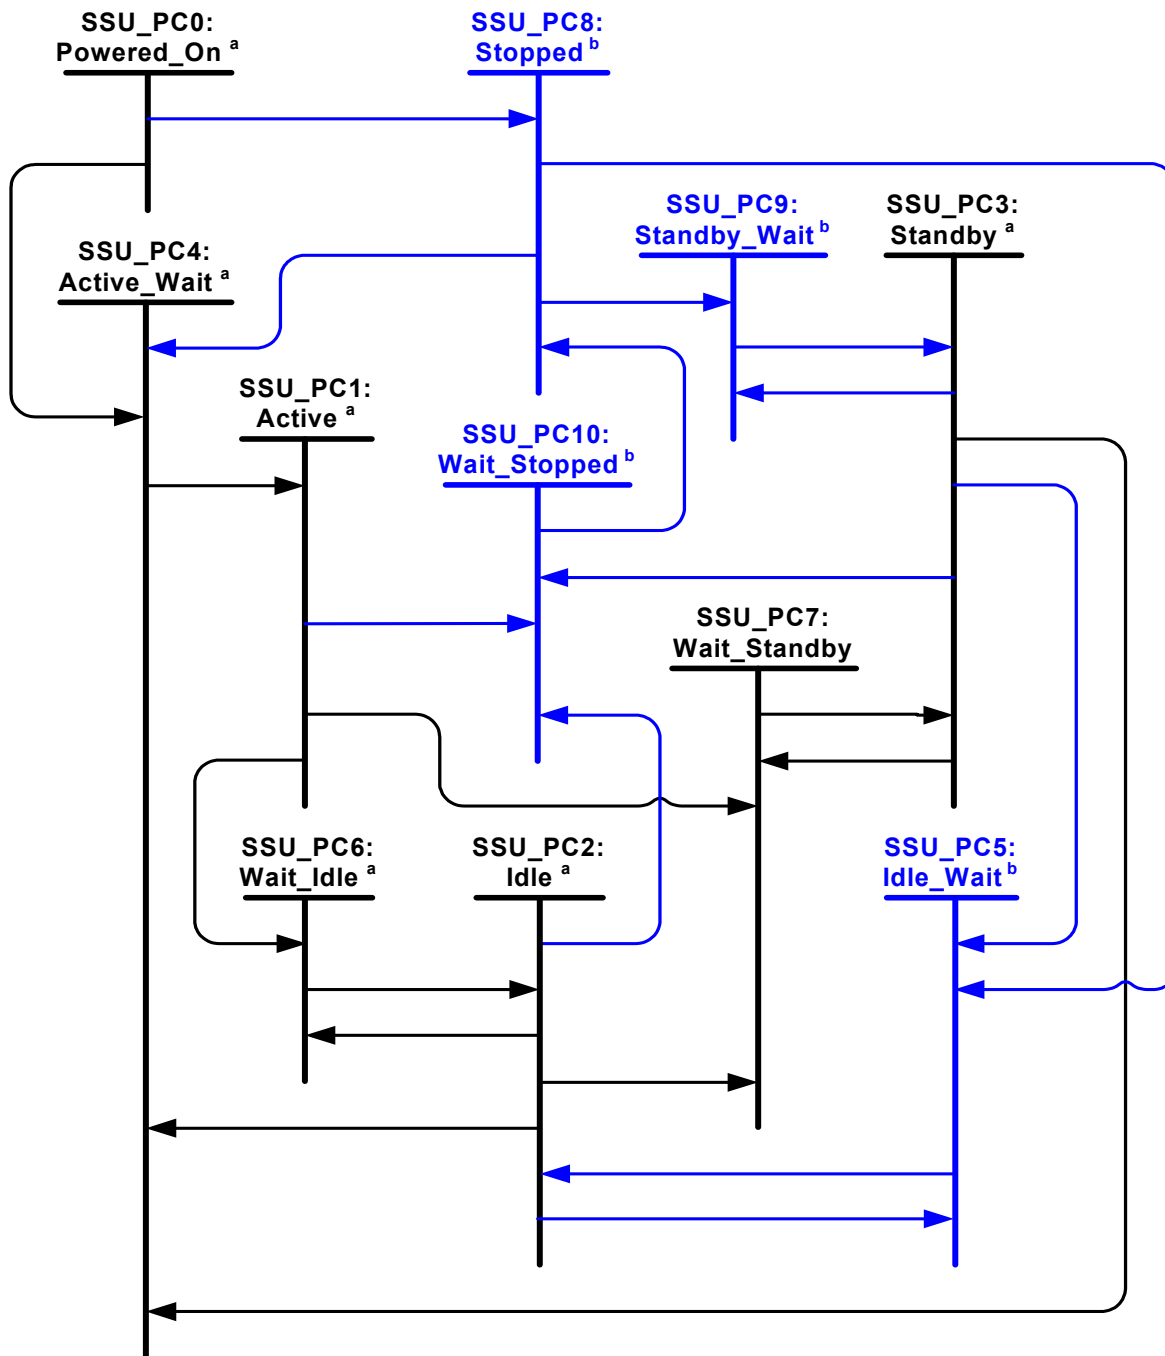
- a) SSU\_PC5:Wait\_Idle;
- b) SSU\_PC6:Wait\_Standby; and
- c) SSU\_PC10:Wait\_Stopped.

Any command causing a state machine transition (e.g., a START STOP UNIT command with the IMMED bit set to zero) shall not complete with GOOD status until this state machine reaches the state (i.e., power condition) required or specified by the command.

The SSU\_PC state machine shall start in the SSU\_PC0:Powered\_On state after power on. The SSU\_PC state machine shall be configured to transition to the SSU\_PC4:Active\_Wait state or the SSU\_PC8:Stopped state after power on by a mechanism outside the scope of this standard.

This state machine references timers controlled by the Power Condition mode page (see SPC-5) and refers to the START STOP UNIT command (see 5.26).

Figure 11 describes the SSU\_PC state machine.



Notes:

<sup>a</sup> This state or transition is also described in SPC-4, but may have additional characteristics unique to this standard (e.g., a transition to or from a state described in this standard).

<sup>b</sup> This state or transition is described in this standard.

Figure 11 — SSU\_PC state machine

**4.21.5.2 SSU\_PC0:Powered\_On state****4.21.5.2.1 SSU\_PC0:Powered\_On state description**

See the PC0:Powered\_On state in SPC-5 for details about this state.

**4.21.5.2.2 Transition SSU\_PC0:Powered\_On to SSU\_PC4:Active\_Wait**

This transition shall occur if:

- a) the logical unit is ready to begin power on initialization; and
- b) the logical unit has been configured to transition to the SSU\_PC4:Active\_Wait state.

The transition shall include a Transitioning From Powered On argument.

**4.21.5.2.3 Transition SSU\_PC0:Powered\_On to SSU\_PC8:Stopped**

This transition shall occur if:

- a) the logical unit has been configured to transition to the SSU\_PC8:Stopped state.

The transition shall include a Transitioning From Powered On argument.

**4.21.5.3 SSU\_PC1:Active state****4.21.5.3.1 SSU\_PC1:Active state description**

See the PC1:Active state in SPC-5 for details about this state.

**4.21.5.3.2 Transition SSU\_PC1:Active to SSU\_PC5:Wait\_Idle**

This transition shall occur if:

- a) the device server processes a START STOP UNIT command (see 5.26) with the POWER CONDITION field set to 2h (i.e., IDLE); or
- b) an idle condition timer (see SPC-5) is enabled, and that timer has expired.

The transition shall include a:

- a) Transitioning To Idle\_a argument, if:
    - A) the highest priority timer that expired is the idle\_a condition timer; or
    - B) the START STOP UNIT command being processed has the POWER CONDITION MODIFIER field set to 0h (i.e., idle\_a power condition);
  - b) Transitioning To Idle\_b argument, if:
    - A) the highest priority timer that expired is the idle\_b condition timer; or
    - B) the START STOP UNIT command being processed has the POWER CONDITION MODIFIER field set to 1h (i.e., idle\_b power condition);
- or
- c) Transitioning To Idle\_c argument, if:
    - A) the highest priority timer that expired is the idle\_c condition timer; or
    - B) the START STOP UNIT command being processed has the POWER CONDITION MODIFIER field set to 2h (i.e., idle\_c power condition).

**4.21.5.3.3 Transition SSU\_PC1:Active to SSU\_PC6:Wait\_Standby**

This transition shall occur if:

- a) the device server processes a START STOP UNIT command (see 5.26) with the POWER CONDITION field set to 3h (i.e., STANDBY); or
- b) a standby condition timer (see SPC-5) is enabled and that timer has expired.



The transition shall include a:

- a) Transitioning To Standby\_z argument, if:
  - A) the highest priority timer that expired is the standby\_z condition timer; or
  - B) the START STOP UNIT command being processed has the POWER CONDITION MODIFIER field set to 0h (i.e., standby\_z power condition);
- or
- b) Transitioning To Standby\_y argument, if:
  - A) the highest priority timer that expired is the standby\_y condition timer; or
  - B) the START STOP UNIT command being processed has the POWER CONDITION MODIFIER field set to 1h (i.e., standby\_y power condition).

#### **4.21.5.3.4 Transition SSU\_PC1:Active to SSU\_PC10:Wait\_Stopped**

This transition shall occur if:

- a) the device server processes a START STOP UNIT command (see 5.26) with the START bit set to zero and the POWER CONDITION field set to 0h (i.e., START\_VALID).

#### **4.21.5.4 SSU\_PC2:Idle state**

##### **4.21.5.4.1 SSU\_PC2:Idle state description**

See the PC2:Idle state in SPC-5 for details about this state.

##### **4.21.5.4.2 Transition SSU\_PC2:Idle to SSU\_PC4:Active\_Wait**

This transition shall occur if:

- a) the device server processes a START STOP UNIT command (see 5.26) with the START bit set to one and the POWER CONDITION field set to 0h (i.e., START\_VALID);
- b) the device server processes a START STOP UNIT command with the POWER CONDITION field set to 1h (i.e., ACTIVE); or
- c) the device server processes a command that requires the logical unit to be in the SSU\_PC1:Active state to continue processing that command.

The transition shall include a:

- a) Transitioning From Idle argument; and
- b) Transitioning From Idle\_c argument if the current power condition is the idle\_c power condition.

##### **4.21.5.4.3 Transition SSU\_PC2:Idle to SSU\_PC5:Wait\_Idle**

This transition shall occur if:

- a) the following occur:
  - A) an idle condition timer is enabled and that idle condition timer has expired; and
  - B) the priority of that idle condition timer is greater than the priority of the idle condition timer associated with the current idle power condition (see SPC-5);
- or
- b) the device server processes a START STOP UNIT command (see 5.26) with the POWER CONDITION field set to 2h (i.e., IDLE) and the POWER CONDITION MODIFIER field set to a value that specifies that the logical unit transition to a lower idle power condition.

The transition shall include a:

- a) Transitioning To Idle\_b argument, if:
  - A) the highest priority timer that expired is the idle\_b condition timer; or

- B) the START STOP UNIT command being processed has the POWER CONDITION MODIFIER field set to 1h (i.e., idle\_b power condition);

or

- b) Transitioning To Idle\_c argument, if:
  - A) the highest priority timer that expired is the idle\_c condition timer; or
  - B) the START STOP UNIT command being processed has the POWER CONDITION MODIFIER field set to 2h (i.e., idle\_c power condition).

#### 4.21.5.4.4 Transition SSU\_PC2:Idle to SSU\_PC6:Wait\_Standby

This transition shall occur if:

- a) the device server processes a START STOP UNIT command (see 5.26) with the POWER CONDITION field set to 3h (i.e., STANDBY); or
- b) a standby condition timer is enabled and that timer has expired.

The transition shall include a:

- a) Transitioning To Standby\_z argument, if:
  - A) the highest priority timer that expired is the standby\_z condition timer; or
  - B) the START STOP UNIT command being processed has the POWER CONDITION MODIFIER field set to 0h (i.e., standby\_z power condition);
- or
- b) Transitioning To Standby\_y argument, if:
  - A) the highest priority timer that expired is the standby\_y condition timer; or
  - B) the START STOP UNIT command being processed has the POWER CONDITION MODIFIER field set to 1h (i.e., standby\_y power condition).

#### 4.21.5.4.5 Transition SSU\_PC2:Idle to SSU\_PC7:Idle\_Wait

This transition shall occur if:

- a) the device server processes a START STOP UNIT command (see 5.26) with the POWER CONDITION field set to 2h (i.e., IDLE) and the POWER CONDITION MODIFIER field set to a value that specifies that the logical unit transition to a higher idle power condition.

The transition shall include Transitioning From Idle argument and a:

- a) Transitioning To Idle\_a argument, if the START STOP UNIT command being processed has the POWER CONDITION MODIFIER field set to 0h (i.e., idle\_a power condition); or
- b) Transitioning To Idle\_b argument, if the START STOP UNIT command being processed has the POWER CONDITION MODIFIER field set to 1h (i.e., idle\_b power condition).

#### 4.21.5.4.6 Transition SSU\_PC2:Idle to SSU\_PC10:Wait\_Stopped

This transition shall occur if:

- a) the device server processes a START STOP UNIT command (see 5.26) with the START bit set to zero and the POWER CONDITION field set to 0h (i.e., START\_VALID).

#### 4.21.5.5 SSU\_PC3:Standby state

##### 4.21.5.5.1 SSU\_PC3:Standby state description

See the PC3:Standby state in SPC-5 for details about this state.

**4.21.5.5.2 Transition SSU\_PC3:Standby to SSU\_PC4:Active\_Wait**

This transition shall occur if:

- a) the device server processes a START STOP UNIT command (see 5.26) with the START bit set to one and the POWER CONDITION field set to 0h (i.e., START\_VALID);
- b) the device server processes a START STOP UNIT command with the POWER CONDITION field set to 1h (i.e., ACTIVE); or
- c) the device server processes a command that requires the logical unit to be in the SSU\_PC1:Active state to continue processing that command.

The transition shall include a Transitioning From Standby argument.

**4.21.5.5.3 Transition SSU\_PC3:Standby to SSU\_PC6:Wait\_Standby**

This transition shall occur if:

- a) the following occur:
  - A) the standby\_z condition timer is enabled and that timer expires; and
  - B) the priority of that standby condition timer is greater than the priority of the standby condition timer associated with the current standby power condition (see SPC-5);
- or
- b) the device server processes a START STOP UNIT command with the POWER CONDITION field set to 3h (i.e., STANDBY) and the POWER CONDITION MODIFIER field set to a value that specifies that the logical unit transition to a lower standby power condition.

The transition shall include Transitioning To Standby\_z argument.

**4.21.5.5.4 Transition SSU\_PC3:Standby to SSU\_PC7:Idle\_Wait**

This transition shall occur if:

- a) the device server processes a START STOP UNIT command (see 5.26) with the POWER CONDITION field set to 2h (i.e., IDLE); or
- b) the device server processes a command and determines that the device server is capable of continuing the processing of that command, when the logical unit is in the SSU\_PC2:Idle state.

The transition shall include a Transitioning From Standby argument and a:

- a) Transitioning To Idle\_a argument, if:
  - A) the device server processes a command and determines that the device server is capable of continuing the processing of that command, when the logical unit is in the idle\_a power condition;
  - or
  - B) the device server processes a START STOP UNIT command with the POWER CONDITION MODIFIER field set to 0h (i.e., idle\_a power condition);
- b) Transitioning To Idle\_b argument, if:
  - A) the device server processes a command and determines that the device server is capable of continuing the processing of that command, when the logical unit is in the idle\_b power condition;
  - or
  - B) the device server processes a START STOP UNIT command with the POWER CONDITION MODIFIER field set to 1h (i.e., idle\_b power condition);
- or
- c) Transitioning To Idle\_c argument, if:
  - A) the device server processes a command and determines that the device server is capable of continuing the processing of that command, when the logical unit is in the idle\_c power condition;
  - or
  - B) the device server processes a START STOP UNIT command with the POWER CONDITION MODIFIER field set to 2h (i.e., idle\_c power condition).

**4.21.5.5.5 Transition SSU\_PC3:Standby to SSU\_PC9:Standby\_Wait**

This transition shall occur if:

- a) the device server processes a START STOP UNIT command (see 5.26) with the POWER CONDITION field set to 3h (i.e., STANDBY) and the POWER CONDITION MODIFIER field set to a value that specifies that the logical unit transition to a higher standby power condition.

The transition shall include a Transitioning To Standby\_y argument.

**4.21.5.5.6 Transition SSU\_PC3:Standby to SSU\_PC10:Wait\_Stopped**

This transition shall occur if:

- a) the device server processes a START STOP UNIT command (see 5.26) with the START bit set to zero and the POWER CONDITION field set to 0h (i.e., START\_VALID).

**4.21.5.6 SSU\_PC4:Active\_Wait state****4.21.5.6.1 SSU\_PC4:Active\_Wait state description**

While in this state:

- a) each idle condition timer that is enabled and not expired is running;
- b) each standby condition timer that is enabled and not expired is running;
- c) the device server shall provide power management pollable sense data (see SPC-5) with the sense key set to NO SENSE and the additional sense code set to LOGICAL UNIT TRANSITIONING TO ANOTHER POWER CONDITION; and
- d) the logical unit is performing the operations required for it to be in the SSU\_PC1:Active state (e.g., a disk drive spins up its medium).

If this state was entered with a Transitioning From Idle argument, then:

- a) the device server is capable of processing and completing the same commands, except START STOP UNIT commands with the IMMED bit set to zero (see 5.26), that the device server is able to process and complete while in the SSU\_PC2:Idle state;
- b) the peak power consumed in this state shall be no more than the typical peak power consumed in the SSU\_PC1:Active state; and
- c) if:
  - A) this state was entered with a Transitioning From Idle\_c argument; and
  - B) the CCF IDLE field in the Power Condition mode page (see SPC-5) is set to 10b (i.e., enabled),
 then the device server shall terminate any command, except a START STOP UNIT command, that requires the logical unit be in the SSU\_PC1:Active state to continue processing, with CHECK CONDITION status, with the sense key set to NOT READY and the additional sense code set to LOGICAL UNIT IS IN PROCESS OF BECOMING READY.

If this state was entered with a Transitioning From Standby argument, then:

- a) the device server is capable of processing and completing the same commands, except START STOP UNIT commands with the IMMED bit set to zero, that the device server is able to process and complete while in the SSU\_PC3:Standby state;
- b) the peak power consumption in this state is not limited by this standard; and
- c) if the CCF STANDBY field in the Power Condition mode page (see SPC-5) is set to 10b (i.e., enabled), then the device server shall terminate any command, except a START STOP UNIT command, that requires the logical unit be in the SSU\_PC1:Active state or SSU\_PC2:Idle state to continue processing, with CHECK CONDITION status, with the sense key set to NOT READY and the additional sense code set to LOGICAL UNIT IS IN PROCESS OF BECOMING READY.

If this state was entered with a Transitioning From Stopped argument, then:

- a) the device server is capable of processing and completing the same commands, except START STOP UNIT commands with the IMMED bit set to zero, that the device server is able to process and complete while in the SSU\_PC8:Stopped state;
- b) the peak power consumption in this state is not limited by this standard; and
- c) if the CCF STOPPED field in the Power Condition mode page (see SPC-5) is set to 10b (i.e., enabled), then the device server shall terminate any TEST UNIT READY command or medium access command, with CHECK CONDITION status, with the sense key set to NOT READY and the additional sense code set to LOGICAL UNIT IS IN PROCESS OF BECOMING READY.

If this state was entered with a Transitioning From Powered On argument, then:

- a) the device server is capable of processing and completing the same commands, except START STOP UNIT commands with the IMMED bit set to zero or TEST UNIT READY command, that the device server is able to process and complete while in the SSU\_PC8:Stopped state;
- b) the peak power consumption in this state is not limited by this standard; and
- c) the device server shall terminate any TEST UNIT READY command or medium access command with CHECK CONDITION status with the sense key set to NOT READY and the additional sense code set to LOGICAL UNIT IS IN PROCESS OF BECOMING READY.

If an idle condition timer or a standby condition timer is enabled and expires, then that timer is ignored in this state.

#### **4.21.5.6.2 Transition SSU\_PC4:Active\_Wait to SSU\_PC1:Active**

See the PC4:Active\_Wait to PC1:Active transition in SPC-5 for details about this transition.

#### **4.21.5.7 SSU\_PC5:Wait\_Idle state**

##### **4.21.5.7.1 SSU\_PC5:Wait\_Idle state description**

See the PC5:Wait\_Idle state in SPC-5 for details about this state.

##### **4.21.5.7.2 Transition SSU\_PC5:Wait\_Idle to SSU\_PC2:Idle**

See the PC5:Wait\_Idle to PC2:Idle transition in SPC-5 for details about this transition.

#### **4.21.5.8 SSU\_PC6:Wait\_Standby state**

##### **4.21.5.8.1 SSU\_PC6:Wait\_Standby state description**

See the PC6:Wait\_Standby state in SPC-5 for details about this state.

##### **4.21.5.8.2 Transition SSU\_PC6:Wait\_Standby to SSU\_PC3:Standby**

See the PC6:Wait\_Standby to PC3:Standby transition in SPC-5 for details about this transition.

#### **4.21.5.9 SSU\_PC7:Idle\_Wait state**

##### **4.21.5.9.1 SSU\_PC7:Idle\_Wait state description**

While in this state:

- a) each idle condition timer that is enabled and not expired is running;
- b) each standby condition timer that is enabled and not expired is running;
- c) the device server shall provide power management pollable sense data (see SPC-5) with the sense key set to NO SENSE and the additional sense code set to LOGICAL UNIT TRANSITIONING TO ANOTHER POWER CONDITION; and

- d) the logical unit is performing the operations required for it to be in the SSU\_PC2:Idle state (e.g., a disk drive spins up its medium).

If this state was entered with a Transitioning From Idle argument, then:

- a) the device server is capable of processing and completing the same commands, except START STOP UNIT commands with the IMMED bit set to zero (see 5.26), that the device server is able to process and complete while in the SSU\_PC2:Idle state; and
- b) the peak power consumed in this state shall be no more than the typical peak power consumed in the SSU\_PC1:Active state.

If this state was entered with a Transitioning From Standby argument, then:

- a) the device server is capable of processing and completing the same commands, except START STOP UNIT commands with the IMMED bit set to zero, that the device server is able to process and complete while in the SSU\_PC3:Standby state;
- b) the peak power consumption in this state is not limited by this standard; and
- c) the CCF STANDBY field in the Power Condition mode page (see SPC-5) is set to 10b (i.e., enabled), then the device server shall terminate any command, except a START STOP UNIT command, that requires the logical unit be in the SSU\_PC1:Active state or SSU\_PC2:Idle state to continue processing, with CHECK CONDITION status, with the sense key set to NOT READY and the additional sense code set to LOGICAL UNIT IS IN PROCESS OF BECOMING READY.

If this state was entered with a Transitioning From Stopped argument, then:

- a) the device server is capable of processing and completing the same commands, except START STOP UNIT commands with the IMMED bit set to zero, that the device server is able to process and complete while in the SSU\_PC8:Stopped state;
- b) the peak power consumption in this state is not limited by this standard; and
- c) if the CCF STOPPED field in the Power Condition mode page (see SPC-5) is set to 10b (i.e., enabled), then the device server shall terminate any TEST UNIT READY command or medium access command, with CHECK CONDITION status, with the sense key set to NOT READY and the additional sense code set to LOGICAL UNIT IS IN PROCESS OF BECOMING READY.

If an idle condition timer or a standby condition timer is enabled and expires, then that timer is ignored in this state.

#### **4.21.5.9.2 Transition SSU\_PC7:Idle\_Wait to SSU\_PC2:Idle**

This transition shall occur when the logical unit meets the requirements for being in the:

- a) idle\_a power condition, if this state was entered with a Transitioning To Idle\_a argument;
- b) idle\_b power condition, if this state was entered with a Transitioning To Idle\_b argument; or
- c) idle\_c power condition, if this state was entered with a Transitioning To Idle\_c argument.

#### **4.21.5.10 SSU\_PC8:Stopped state**

##### **4.21.5.10.1 SSU\_PC8:Stopped state description**

While in this state:

- a) the logical unit is in the stopped power condition (see 4.21.4);
- b) the idle condition timers and the standby condition timers are disabled;
- c) the device server shall provide power management pollable sense data (see SPC-5); and
- d) the device server shall terminate each medium access command or TEST UNIT READY command (see SPC-5) as described in 4.21.4.

**4.21.5.10.2 Transition SSU\_PC8:Stopped to SSU\_PC4:Active\_Wait**

This transition shall occur if:

- a) the device server processes a START STOP UNIT command (see 5.26) with the START bit set to one and the POWER CONDITION field set to 0h (i.e., START\_VALID); or
- b) the device server processes a START STOP UNIT command with the POWER CONDITION field set to 1h (i.e., ACTIVE).

The transition shall include a Transitioning From Stopped argument.

**4.21.5.10.3 Transition SSU\_PC8:Stopped to SSU\_PC7:Idle\_Wait**

This transition shall occur if:

- a) the device server processes a START STOP UNIT command (see 5.26) with the POWER CONDITION field set to 2h (i.e., IDLE).

The transition shall include a Transitioning From Stopped argument and a:

- a) Transitioning To Idle\_a argument, if the START STOP UNIT command being processed has the POWER CONDITION MODIFIER field set to 0h (i.e., idle\_a power condition);
- b) Transitioning To Idle\_b argument, if the START STOP UNIT command being processed has the POWER CONDITION MODIFIER field set to 1h (i.e., idle\_b power condition); or
- c) Transitioning To Idle\_c argument, if the START STOP UNIT command being processed has the POWER CONDITION MODIFIER field set to 2h (i.e., idle\_c power condition).

**4.21.5.10.4 Transition SSU\_PC8:Stopped to SSU\_PC9:Standby\_Wait**

This transition shall occur if:

- a) the device server processes a START STOP UNIT command (see 5.26) with the POWER CONDITION field set to STANDBY.

The transition shall include a Transitioning From Stopped argument and a:

- a) Transitioning To Standby\_z argument, if the START STOP UNIT command being processed has the POWER CONDITION MODIFIER field set to 0h (i.e., standby\_z power condition); or
- b) Transitioning To Standby\_y argument, if the START STOP UNIT command being processed has the POWER CONDITION MODIFIER field set to 1h (i.e., standby\_y power condition).

**4.21.5.11 SSU\_PC9:Standby\_Wait state****4.21.5.11.1 SSU\_PC9:Standby\_Wait state description**

While in this state:

- a) the device server shall provide power management pollable sense data (see SPC-5) with the sense key set to NO SENSE and the additional sense code set to LOGICAL UNIT TRANSITIONING TO ANOTHER POWER CONDITION;
- b) the peak power consumed in this state shall be no more than the typical peak power consumed in the SSU\_PC1:Active state; and
- c) the logical unit is performing the operations required for it to be in the SSU\_PC3:Standby state ((e.g., a direct access block device is activating circuitry).

If this state was entered with a Transitioning From Standby argument, then the device server is capable of processing and completing the same commands, except START STOP UNIT commands with the IMMED bit set to zero (see 5.26), that the device server is able to process and complete in the SSU\_PC3:Standby state.

If this state was entered with a Transitioning From Stopped argument, then:

- a) the device server is capable of processing and completing the same commands, except START STOP UNIT commands with the IMMED bit set to zero, that the device server is able to process and complete while in the SSU\_PC8:Stopped state; and
- b) if the CCF STOPPED field in the Power Condition mode page (see SPC-5) is set to 10b (i.e., enabled), then the device server shall terminate any TEST UNIT READY command or medium access command, with CHECK CONDITION status, with the sense key set to NOT READY and the additional sense code set to LOGICAL UNIT IS IN PROCESS OF BECOMING READY.

#### **4.21.5.11.2 Transition SSU\_PC9:Standby\_Wait to SSU\_PC3:Standby**

This transition shall occur when the logical unit meets the requirements for being in the:

- a) standby\_y power condition, if this state was entered with a Transitioning To Standby\_y argument; or
- b) standby\_z power condition, if this state was entered with a Transitioning To Standby\_z argument.

#### **4.21.5.12 SSU\_PC10:Wait\_Stopped state**

##### **4.21.5.12.1 SSU\_PC10:Wait\_Stopped state description**

While in this state:

- a) the device server shall provide power management pollable sense data (see SPC-5) with the sense key set to NO SENSE and the additional sense code set to LOGICAL UNIT TRANSITIONING TO ANOTHER POWER CONDITION;
- b) the device server is capable of processing and completing the same commands, except START STOP UNIT commands with the IMMED bit set to zero (see 5.26), that the device server is able to process and complete in the SSU\_PC8:Stopped state;
- c) the logical unit is performing the operations required for it to be in the SSU\_PC8:Stopped state (e.g., a disk drive spins down its medium); and
- d) the device server shall terminate any TEST UNIT READY command or medium access command with CHECK CONDITION status with the sense key set to NOT READY and the additional sense code set to LOGICAL UNIT NOT READY, INITIALIZING COMMAND REQUIRED.

##### **4.21.5.12.2 Transition SSU\_PC10:Wait\_Stopped to SSU\_PC8:Stopped**

This transition shall occur when:

- a) the logical unit meets the requirements for being in the SSU\_PC8:Stopped state.

## **4.22 Protection information model**

### **4.22.1 Protection information overview**

The protection information model provides for protection of user data while user data is being transferred between a sender and a receiver. Protection information is generated at the application layer and may be checked by any object associated with the I\_T\_L nexus (see SAM-5). Once received, protection information is retained (e.g., written to the medium, stored in non-volatile memory, or recalculated on read back) by the device server until overwritten. Power loss, hard reset, logical unit reset, and I\_T nexus loss shall have no effect on the retention of protection information.

Support for protection information shall be indicated in the PROTECT bit in the standard INQUIRY data (see SPC-5).

If the logical unit is formatted with protection information, and the EMDP bit is set to one in the Disconnect-Reconnect mode page (see SPC-5), then checking of the logical block reference tag within a



service delivery subsystem without accounting for modified data pointers and data alignments may cause false errors when logical blocks are transmitted out of order.

Protection information is also referred to as the data integrity field (DIF).

#### 4.22.2 Protection types

##### 4.22.2.1 Protection types overview

The content of protection information is dependent on the type of protection to which a logical unit has been formatted.

The type of protection supported by the logical unit shall be indicated in the SPT field in the Extended INQUIRY Data VPD page (see SPC-5). The current protection type shall be indicated in the P\_TYPE field ~~in the READ-GAPACITY (16) parameter data~~ (see 5.18.2).

An application client may format the logical unit to a specific type of protection using the FMTPINFO field and the PROTECTION FIELD USAGE field ~~in the FORMAT UNIT command~~ (see 5.4).

An application client may format the logical unit to place protection information at intervals other than on logical block boundaries using the PROTECTION INTERVAL EXPONENT field ~~in the FORMAT UNIT command~~.

A medium access command is processed in a different manner by a device server depending on the type of protection in effect. When used in relation to types of protection, the term “medium access command” is defined as any one of the following commands:

- a) COMPARE AND WRITE;
- b) ORWRITE (16);
- c) ORWRITE (32);
- d) READ (10);
- e) READ (12);
- f) READ (16);
- g) READ (32);
- h) VERIFY (10);
- i) VERIFY (12);
- j) VERIFY (16);
- k) VERIFY (32);
- l) WRITE (10);
- m) WRITE (12);
- n) WRITE (16);
- o) WRITE (32);
- p) WRITE AND VERIFY (10);
- q) WRITE AND VERIFY (12);
- r) WRITE AND VERIFY (16);
- s) WRITE AND VERIFY (32);
- t) WRITE ATOMIC (16);
- u) WRITE ATOMIC (32);
- v) WRITE SAME (10);
- w) WRITE SAME (16);
- x) WRITE SAME (32);
- y) WRITE STREAM (16);
- z) WRITE STREAM (32);
- aa) XDWRITEREAD (10); and
- ab) XDWRITEREAD (32).

##### 4.22.2.2 Type 0 protection

Type 0 protection defines no protection over that which is defined within the transport protocol.

A logical unit that has been formatted with protection information disabled (see 5.3) or a logical unit that does not support protection information (i.e., the PROTECT bit set to zero in the standard INQUIRY data (see SPC-5)) has type 0 protection.

If type 0 protection is enabled and the RDPROTECT field, the WRPROTECT field, the VRPROTECT field, or the ORPROTECT field is set to a non-zero value, then medium access commands are invalid and may be terminated by the device server with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

If type 0 protection is enabled, then the following medium access commands are invalid and shall be terminated by the device server with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID COMMAND OPERATION CODE:

- a) READ (32);
- b) VERIFY (32);
- c) WRITE (32);
- d) WRITE AND VERIFY (32);
- e) WRITE ATOMIC (32);
- f) WRITE SAME (32); and
- g) WRITE STREAM (32).

#### 4.22.2.3 Type 1 protection

Type 1 protection:

- a) defines the content of each LOGICAL BLOCK GUARD field;
- b) does not define the content of any LOGICAL BLOCK APPLICATION TAG field; and
- c) defines the content of each LOGICAL BLOCK REFERENCE TAG field.

If type 1 protection is enabled, then the following medium access commands are invalid and shall be terminated by the device server with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID COMMAND OPERATION CODE:

- a) READ (32);
- b) VERIFY (32);
- c) WRITE (32);
- d) WRITE AND VERIFY (32);
- e) WRITE ATOMIC (32);
- f) WRITE SAME (32); and
- g) WRITE STREAM (32).

For valid medium access commands in which the RDPROTECT field, the WRPROTECT field, the VRPROTECT field, or the ORPROTECT field is set to:

- a) zero, the Data-In Buffer and/or Data-Out Buffer associated with those commands shall consist of logical block data containing only user data; or
- b) a non-zero value, the Data-In Buffer and/or Data-Out Buffer shall consist of logical block data containing both user data and protection information.

#### 4.22.2.4 Type 2 protection

Type 2 protection:

- a) defines the content of each LOGICAL BLOCK GUARD field;
- b) does not define the content of any LOGICAL BLOCK APPLICATION TAG field; and
- c) defines, except for the first logical block addressed by the command, the content of each LOGICAL BLOCK REFERENCE TAG field.

If type 2 protection is enabled and the RDPROTECT field, the WRPROTECT field, the VRPROTECT field, or the ORPROTECT field is set to a non-zero value, then the following medium access commands are invalid and shall

be terminated by the device server with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID COMMAND OPERATION CODE:

- a) COMPARE AND WRITE;
- b) ORWRITE (16);
- c) ORWRITE (32);
- d) READ (10);
- e) READ (12);
- f) READ (16);
- g) VERIFY (10);
- h) VERIFY (12);
- i) VERIFY (16);
- j) WRITE (10);
- k) WRITE (12);
- l) WRITE (16);
- m) WRITE AND VERIFY (10);
- n) WRITE AND VERIFY (12);
- o) WRITE AND VERIFY (16);
- p) WRITE ATOMIC (16);
- q) WRITE SAME (10);
- r) WRITE SAME (16);
- s) WRITE STREAM (16);
- t) XDWRITEREAD (10); and
- u) XDWRITEREAD (32).

For valid medium access commands in which the RDPROTECT field, the WRPROTECT field, the VRPROTECT field, or the ORPROTECT field is set to:

- a) zero, the Data-In Buffer and/or Data-Out Buffer associated with those commands shall consist of logical block data containing only user data; or
- b) a non-zero value, the Data-In Buffer and/or Data-Out Buffer shall consist of logical block data containing both user data and protection information.

#### 4.22.2.5 Type 3 protection

Type 3 protection:

- a) defines the content of each LOGICAL BLOCK GUARD field;
- b) does not define the content of any LOGICAL BLOCK APPLICATION TAG field; and
- c) does not define the content of any LOGICAL BLOCK REFERENCE TAG field.

If type 3 protection is enabled, then the following medium access commands are invalid and shall be terminated by the device server with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID COMMAND OPERATION CODE:

- a) READ (32);
- b) VERIFY (32);
- c) WRITE (32);
- d) WRITE AND VERIFY (32);
- e) WRITE ATOMIC (32);
- f) WRITE SAME (32); and
- g) WRITE STREAM (32).

For valid medium access commands in which the RDPROTECT field, the WRPROTECT field, the VRPROTECT field, or the ORPROTECT field is set to:

- a) zero, the Data-In Buffer and/or Data-Out Buffer associated with those commands shall consist of logical block data containing only user data; or
- b) a non-zero value, the Data-In Buffer and/or Data-Out Buffer shall consist of logical block data containing both user data and protection information.

### 4.22.3 Protection information format

Table 22 defines the placement of protection information in a logical block with a single protection information interval (i.e., the PROTECTION INTERVAL EXPONENT field is set to zero in the parameter list header for a FORMAT UNIT command (see 5.4.2.2))

**Table 22 — Logical block data format with a single protection information interval**

Byte	Bit	7	6	5	4	3	2	1	0
0		USER DATA							
...									
n - 1									
n	(MSB)	LOGICAL BLOCK GUARD							
n + 1		(LSB)							
n + 2	(MSB)	LOGICAL BLOCK APPLICATION TAG							
n + 3		(LSB)							
n + 4	(MSB)	LOGICAL BLOCK REFERENCE TAG							
...									
n + 7		(LSB)							

Table 23 shows an example of the placement of protection information in a logical block with more than one protection information interval (i.e., the PROTECTION INTERVAL EXPONENT field is set to a non-zero value in the parameter list header for a FORMAT UNIT command (see 5.4.2.2)).

**Table 23 — An example of the logical block data for a logical block with more than one protection information interval**

Bit Byte	7	6	5	4	3	2	1	0
0	USER DATA [first]							
...								
n - 1								
n	(MSB)	LOGICAL BLOCK GUARD [first]						(LSB)
n + 1								
n + 2	(MSB)	LOGICAL BLOCK APPLICATION TAG [first]						(LSB)
n + 3								
n + 4	(MSB)	LOGICAL BLOCK REFERENCE TAG [first]						(LSB)
...								
n + 7								(LSB)
n + 8		USER DATA [second]						
...								
m - 1								
m	(MSB)	LOGICAL BLOCK GUARD [second]						(LSB)
m + 1								
m + 2	(MSB)	LOGICAL BLOCK APPLICATION TAG [second]						(LSB)
m + 3								
m + 4	(MSB)	LOGICAL BLOCK REFERENCE TAG [second]						(LSB)
...								
m + 7								(LSB)
		⋮						
...								
z - 1		USER DATA [last]						
z	(MSB)	LOGICAL BLOCK GUARD [last]						(LSB)
z + 1								
z + 2	(MSB)	LOGICAL BLOCK APPLICATION TAG [last]						(LSB)
z + 3								
z + 4	(MSB)	LOGICAL BLOCK REFERENCE TAG [last]						(LSB)
...								
z + 7								(LSB)

Each USER DATA field shall contain user data.

Each LOGICAL BLOCK GUARD field contains a CRC (see 4.22.4). Only the contents of the USER DATA field immediately preceding THE LOGICAL BLOCK GUARD field (i.e., the user data between the preceding logical block reference tag, if any, and the current logical block guard) shall be used to generate and check the CRC contained in the LOGICAL BLOCK GUARD field.

Each LOGICAL BLOCK APPLICATION TAG field is set by the application client. If the device server detects a:

- a) LOGICAL BLOCK APPLICATION TAG field set to FFFFh and type 1 protection (see 4.22.2.3) is enabled;
- b) LOGICAL BLOCK APPLICATION TAG field set to FFFFh and type 2 protection (see 4.22.2.4) is enabled; or
- c) LOGICAL BLOCK APPLICATION TAG field set to FFFFh, LOGICAL BLOCK REFERENCE TAG field set to FFFF\_FFFFh, and type 3 protection (see 4.22.2.5) is enabled,

then the device server disables checking of all protection information for the associated protection information interval when performing:

- a) a read operation; and
- b) a write operation if the NO\_PI\_CHK bit is set to one in the Extended INQUIRY Data VPD page (see SPC-5).

Otherwise, if the ATMPE bit in the Control mode page (see SPC-5) is:

- a) set to one, then the logical block application tags are defined by the Application Tag mode page (see 6.5.3); or
- b) set to zero, then the logical block application tags are not defined by this standard.

The LOGICAL BLOCK APPLICATION TAG field may be modified by a device server if the ATO bit is set to zero in the Control mode page (see SPC-5). If the ATO bit is set to one in the Control mode page, then the device server shall not modify the LOGICAL BLOCK APPLICATION TAG field.

The contents of a LOGICAL BLOCK APPLICATION TAG field shall not be used to generate or check the CRC contained in the LOGICAL BLOCK GUARD field.

The first LOGICAL BLOCK REFERENCE TAG field of the first logical block in the Data-In Buffer and/or Data-Out Buffer shall contain the value specified in table 24.

**Table 24 — Content of the first LOGICAL BLOCK REFERENCE TAG field for the first logical block in the Data-In Buffer and/or Data-Out Buffer**

Protection Type	Content of the first LOGICAL BLOCK REFERENCE TAG field for the first logical block in the Data-In Buffer and/or Data-Out Buffer
Type 1 <sup>a</sup> protection (see 4.22.2.3)	The least significant four bytes of the LBA contained in the LOGICAL BLOCK ADDRESS field of the CDB.
Type 2 protection (see 4.22.2.4)	The value in the EXPECTED INITIAL LOGICAL BLOCK REFERENCE TAG field of the CDB.
Type 3 protection (see 4.22.2.5)	Not defined in this standard. If the ATO bit is set to zero in the Control mode page (see SPC-5), then this field may be modified by the device server. If the ATO bit is set to one in the Control mode page, then the device server shall not modify this field.
<sup>a</sup> The length of the protection information interval is equal to the logical block length (see 5.4.2).	

Subsequent LOGICAL BLOCK REFERENCE TAG fields for a logical block in the Data-In Buffer and/or Data-Out Buffer shall be set as specified in table 25.

**Table 25 — Content of subsequent LOGICAL BLOCK REFERENCE TAG fields for a logical block in the Data-In Buffer and/or Data-Out Buffer**

Protection Type	The content of subsequent LOGICAL BLOCK REFERENCE TAG fields in the Data-In Buffer and/or Data-Out Buffer
Type 1 protection (see 4.22.2.3) and Type 2 protection (see 4.22.2.4)	The previous logical block reference tag plus one. If the contents of the previous LOGICAL BLOCK REFERENCE TAG field is FFFF_FFFFh, then the contents of the subsequent LOGICAL BLOCK REFERENCE TAG field is 0000_0000h.
Type 3 protection (see 4.22.2.5)	Not defined in this standard. If the ATO bit is set to zero in the Control mode page (see SPC-5), then this field may be modified by the device server. If the ATO bit is set to one in the Control mode page, then the device server shall not modify this field.

The contents of a LOGICAL BLOCK REFERENCE TAG field shall not be used to generate or check the CRC contained in the LOGICAL BLOCK GUARD field.

#### 4.22.4 Logical block guard

##### 4.22.4.1 Logical block guard overview

A LOGICAL BLOCK GUARD field shall contain a CRC that is generated from the contents of only the USER DATA field immediately preceding the LOGICAL BLOCK GUARD field.

Table 26 defines the CRC polynomials used to generate the logical block guard from the contents of the USER DATA field.

**Table 26 — CRC polynomials**

Function	Definition
$F(x)$	A polynomial representing the transmitted USER DATA field, which is covered by the CRC. For the purposes of the CRC, the coefficient of the highest order term shall be byte zero bit seven of the USER DATA field and the coefficient of the lowest order term shall be bit zero of the last byte of the USER DATA field.
$F'(x)$	A polynomial representing the received USER DATA field.
$G(x)$	The generator polynomial: $G(x) = x^{16} + x^{15} + x^{11} + x^9 + x^8 + x^7 + x^5 + x^4 + x^2 + x + 1$ (i.e., in finite field notation $G(x) = 1\_8BB7h$ )
$R(x)$	The remainder polynomial calculated during CRC generation by the transmitter, representing the transmitted LOGICAL BLOCK GUARD field.
$R'(x)$	A polynomial representing the received LOGICAL BLOCK GUARD field.
$RB(x)$	The remainder polynomial calculated during CRC checking by the receiver. $RB(x) = 0$ indicates no error was detected.
$RC(x)$	The remainder polynomial calculated during CRC checking by the receiver. $RC(x) = 0$ indicates no error was detected.
$QA(x)$	The quotient polynomial calculated during CRC generation by the transmitter. The value of $QA(x)$ is not used.
$QB(x)$	The quotient polynomial calculated during CRC checking by the receiver. The value of $QB(x)$ is not used.
$QC(x)$	The quotient polynomial calculated during CRC checking by the receiver. The value of $QC(x)$ is not used.
$M(x)$	A polynomial representing the transmitted USER DATA field followed by the transmitted LOGICAL BLOCK GUARD field.
$M'(x)$	A polynomial representing the received USER DATA field followed by the received LOGICAL BLOCK GUARD field.

##### 4.22.4.2 CRC generation

The equations that are used to generate the CRC from  $F(x)$  are as follows. All arithmetic is modulo 2.

The transmitter shall calculate the CRC by appending 16 zeros to  $F(x)$  and dividing by  $G(x)$  to obtain the remainder  $R(x)$ :

$$\frac{(x^{16} \times F(x))}{G(x)} = QA(x) + \frac{R(x)}{G(x)}$$

$R(x)$  is the CRC value, and is transmitted in the LOGICAL BLOCK GUARD field.



$M(x)$  is the polynomial representing the USER DATA field followed by the LOGICAL BLOCK GUARD field (i.e.,  $F(x)$  followed by  $R(x)$ ):

$$M(x) = (x^{16} \times F(x)) + R(x)$$

#### 4.22.4.3 CRC checking

$M'(x)$  (i.e., the polynomial representing the received USER DATA field followed by the received LOGICAL BLOCK GUARD field) may differ from  $M(x)$  (i.e., the polynomial representing the transmitted USER DATA field followed by the transmitted LOGICAL BLOCK GUARD field) if there are transmission errors.

The receiver may check  $M'(x)$  validity by appending 16 zeros to  $F'(x)$  and dividing by  $G(x)$  and comparing the calculated remainder  $RB(x)$  to the received CRC value  $R'(x)$ :

$$\frac{(x^{16} \times F'(x))}{G(x)} = QB(x) + \frac{RB(x)}{G(x)}$$

In the absence of errors in  $F'(x)$  and  $R'(x)$ , the remainder  $RB(x)$  is equal to  $R'(x)$ .

The receiver may check  $M'(x)$  validity by dividing  $M'(x)$  by  $G(x)$  and comparing the calculated remainder  $RC(x)$  to zero:

$$\frac{M'(x)}{G(x)} = QC(x) + \frac{RC(x)}{G(x)}$$

In the absence of errors in  $F'(x)$  and  $R'(x)$ , the remainder  $RC(x)$  is equal to zero.

Both methods of checking  $M'(x)$  validity are mathematically equivalent.

#### 4.22.4.4 CRC test cases

Several CRC test cases are shown in table 27.

**Table 27 — CRC test cases**

Pattern	CRC
32 bytes each set to 00h	0000h
32 bytes each set to FFh	A293h
32 bytes of an incrementing pattern from 00h to 1Fh	0224h
2 bytes each set to FFh followed by 30 bytes set to 00h	21B8h
32 bytes of a decrementing pattern from FFh to E0h	A0B7h

#### 4.22.5 Application of protection information

Before an application client transmits or receives logical block data with protection information, the application client:

- 1) determines if a logical unit supports protection information using the INQUIRY command (see the PROTECT bit in the standard INQUIRY data in SPC-5);
- 2) if protection information is supported, then determines if the logical unit is formatted to accept protection information using the READ CAPACITY (16) command (e.g., see the PROT\_EN bit and the P\_TYPE field in the returned parameter data (see 5.18.2)); and
- 3) if the logical unit supports protection information and is not formatted to accept protection information, then formats the logical unit (see 5.4) with protection information enabled.

If the logical unit supports protection information and is formatted to accept protection information, then the application client may use read commands that support protection information and should use verify commands and write commands that support protection information.

#### 4.22.6 Protection information and commands

The enabling of protection information enables fields in medium access commands that instruct the device server on the handling of protection information. The detailed definitions of each command's protection information fields are in the individual command descriptions.

The commands that are affected while protection information is enabled are listed in table 34.

Commands that cause a device server to return the length in bytes of each logical block (e.g., the MODE SENSE (10) command and the READ CAPACITY (16) command) shall cause the device server to return the combined length of the USER DATA field(s) contained in the logical block, not including the length of any protection information (i.e., the LOGICAL BLOCK GUARD field(s), the LOGICAL BLOCK APPLICATION TAG field(s), and the LOGICAL BLOCK REFERENCE TAG field(s)) (e.g., if the user data plus the protection information is equal to 520 bytes and there is one protection information interval, then 512 is returned).

### 4.23 Grouping function

#### 4.23.1 Grouping function overview

A grouping function is a function that collects information about attributes associated with commands (i.e., information about commands with the same group value are collected into the specified group). The definition of the attributes and the groups is outside the scope of this standard. Groups are identified with the GROUP NUMBER field in the CDBs of certain commands (e.g., the PRE-FETCH (10) command (see 5.10)).

Support for the grouping function is indicated in the GROUP\_SUP bit in the Extended INQUIRY Data VPD page (see SPC-5).

The collection of this information is outside the scope of this standard (e.g., the information may not be transmitted using any SCSI protocols).

EXAMPLE - In a SCSI domain in which two applications are using a subsystem where one application streams data and another accesses data randomly, if:

- a) the streaming application groups all of its commands with one group number (e.g., x); and
  - b) the random application groups all of its commands with another group number (e.g., y),
- then the applications use those group numbers (e.g., x and y) to collect separate performance metrics for each application.

A management application then reads the performance metrics and determines if the performance of a specific group is acceptable.

Device servers that support the grouping function shall support at least group 0 to group 31.

#### 4.23.2 Grouping function extensions for IO advice hints

The IO Advice Hints Grouping mode page (see 6.5.7) may be used to enable or disable the collection of information, and define the IO advice hints associated with each group. The IO Advice hints Grouping mode page contains 64 IO advice hints descriptors (i.e., one descriptor for each possible group number) (see table 227) that are in ascending order of group number. The IO advice hints group number associated with each IO advice hints descriptor is calculated as follows:

$$\text{IO advice hints group number} = (\text{offset} / 16) - 1$$

where:

offset      the offset in bytes of the IO advice hints descriptor from the beginning of the IO Advice Hints Grouping mode page

If the device server implements the IO Advice Hints Grouping mode page, then for all commands that contain a GROUP NUMBER field, the GROUP NUMBER field specifies:

- a) the IO advice hints group number that the device server shall use to associate IO advice hints with all operations associated with that command; and
- b) the cache ID (see 4.15.2) that the device server shall use for all operations associated with that command.

If the IO ADVICE HINTS MODE field is set to 00b (see table 228) in the IO advice hints group descriptor for the specified group number in the IO Advice Hints Grouping mode page, then the device server shall use the IO advice hints in the IO advice hints group descriptor specified by the IO advice hints group number to associate IO advice hints (see 4.32) with each operation associated with the command. If the IO ADVICE HINTS MODE field is set to 01b (see table 228) in the IO advice hints group descriptor for the specified group number in the IO Advice Hints Grouping mode page, then the device server may associate vendor specific IO advice hints with each operation associated with the command.

If the CS\_ENABLE bit is set to one in the IO advice hints descriptor in the IO Advice Hints Grouping mode page, then the device server should use the cache segments associated with the specified cache ID.

If the IC\_ENABLE bit is set to one in the IO advice hints descriptor in the IO Advice Hints Grouping mode page, then the device server shall perform collection of information as described in 4.23.1.

Device servers that support the IO Advice Hints Grouping mode page shall support group 0 to group 63.

## 4.24 Background scan operations

### 4.24.1 Background scan overview

A background scan operation is either a background pre-scan operation (see 4.24.2) or a background medium scan operation (see 4.24.3).

During a background scan operation, the device server performs read medium operations for the purpose of:

- a) identifying logical blocks that are difficult to read (i.e., recoverable) or unreadable (i.e., unrecoverable);
- b) logging problems encountered during the background scan operation; and
- c) when allowed, taking a vendor specific action to repair recoverable logical blocks or perform automatic read reallocation of recoverable logical blocks.

During a background scan operation, if a read medium operation encounters a recovered error (i.e., a logical block is readable but requires extra actions (e.g., retries or application of a correction algorithm) to be read), then the device server may resolve the problem using vendor specific means. The value of the ARRE bit ~~in the Read-Write-Error-Recovery mode page~~ (see 6.5.10) determines whether or not the device server performs automatic read reassignment.

During a background scan operation, if a read medium operation encounters an unrecovered error (i.e., a logical block is unreadable), then the device server may mark the logical block unrecoverable. The value of the AWRE bit ~~in the Read-Write-Error-Recovery mode page~~ (see 6.5.10) determines whether or not the device server performs automatic write reassignment. If the AWRE bit is set to one, then the device server performs automatic write reassignment at the start of the next write medium operation accessing that logical block.

During a background scan operation, the device server:

- a) may scan the logical blocks in any order (e.g., based on physical block layout);
- b) should not retain any data from logical blocks in cache memory after the logical blocks are read;
- c) shall ignore pseudo unrecovered errors with correction disabled (see 4.18.2); and
- d) shall process pseudo unrecovered errors with correction enabled.

## 4.24.2 Background pre-scan operations

### 4.24.2.1 Enabling background pre-scan operations

A background pre-scan operation is enabled after:

- 1) the EN\_PS bit ~~in the Background Control mode page~~ (see 6.5.4) is set to zero;
- 2) the EN\_PS bit is set to one; and
- 3) the SCSI device is power cycled if;
  - A) the S\_L\_FULL bit ~~in the Background Control mode page~~ (see 6.5.4) is:
    - a) set to zero; or
    - b) set to one and the Background Scan log parameters (see 6.4.2) are not all used;
  - and
  - B) the saved value of the EN\_PS bit is set to one.

After a background pre-scan operation is enabled, the device server shall:

- a) initialize the Background Pre-scan Time Limit timer to the time specified in the BACKGROUND PRE-SCAN TIME LIMIT field (see 6.5.4) and start the timer;
- b) initialize the Background Medium Scan Interval timer to the time specified in the BACKGROUND MEDIUM SCAN INTERVAL TIME field (see 6.5.4) and start the timer; and
- c) begin the background pre-scan operation (i.e., begin scanning the medium).

### 4.24.2.2 Suspending and resuming background pre-scan operations

A background pre-scan operation shall be suspended when any of the following occurs:

- a) a command or task management function is processed that requires the background pre-scan operation to be suspended;
- b) a SCSI event (e.g., a hard reset) (see SAM-5) is processed that requires the background pre-scan operation to be suspended;
- c) a power condition timer expires (see the Power Condition mode page in SPC-5), and the PM\_BG\_PRECEDENCE field in the Power Condition mode page is set to 10b; or
- d) the S\_L\_FULL bit (see 6.5.4) is set to one, and the Background Scan log parameters (see 6.4.2) are all used.

If a command is received that requires a background pre-scan operation to be suspended, then the following should occur within the time specified in the MAXIMUM TIME TO SUSPEND BACKGROUND SCAN field (see 6.5.4):

- a) the logical unit suspends the background medium scan operation; and
- b) the device server begins processing the command.

If a background pre-scan operation is suspended, then the device server shall not stop:

- a) the Background Pre-scan Time Limit timer;
- b) the Background Medium Scan Interval timer; and
- c) any process that results in an event that causes a background function to occur (e.g., not stop any timers or counters associated with background functions).

While a background pre-scan operation is suspended and not halted (see 4.24.3.2), the device server shall convert each write operation accessing a logical block that has not been scanned during the background pre-scan operation into a write medium operation followed by a verify medium operation in order to verify that the logical block data just written was read back without error. If a write medium operation accesses a logical block that has already been scanned during the background pre-scan operation, then the device server shall not perform the additional verify medium operation.

A background pre-scan operation shall be resumed from where the operation was suspended when:

- a) there are no commands in any task set to be processed;
- b) there are no task management functions to be processed;
- c) there are no SCSI events to be processed;

- d) no ACA condition exists;
- e) the PM\_BG\_PRECEDENCE field in the Power Condition mode page is set to 10b (see SPC-5), but no power condition timer defined in the Power Condition mode page has expired;
- f) the S\_L\_FULL bit is set to zero (see 6.5.4), or the Background Medium Scan log parameters (see 6.4.2) are not all used;
- g) the logical unit has been idle for the time specified in the MINIMUM IDLE TIME BEFORE BACKGROUND SCAN field (see 6.5.4); and
- h) the background pre-scan operation has not been halted (see 4.24.3.2).

#### 4.24.2.3 Halting background pre-scan operations

The device server shall halt a background pre-scan operation if any of the following occurs:

- a) the background pre-scan operation completes scanning all logical blocks on the medium;
- b) an application client sets the EN\_PS bit to zero in the Background Control mode page (see 6.5.4);
- c) the Background Pre-scan Time Limit timer expires;
- d) the device server detects a fatal error;
- e) the device server detects a vendor specific pattern of errors;
- f) the device server detects a medium formatted without a PLIST (see 4.13); or
- g) the device server detects temperature out of range.

After a background pre-scan operation has been halted, the device server shall not enable a background operation until the conditions in 4.24.2.1 are met.

### 4.24.3 Background medium scan

#### 4.24.3.1 Enabling background medium scan operations

Background medium scan operations are enabled if:

- a) a background pre-scan operation (see 4.24.2) is not in progress;
- b) the S\_L\_FULL bit (see 6.5.4) is:
  - A) set to zero; or
  - B) set to one and the Background Scan log parameters (see 6.4.2) are not all used;
 and
- c) the EN\_BMS bit (see 6.5.4) is set to one.

If background medium scan operations are enabled, then the device server shall begin a background medium scan operation (i.e., begin scanning the medium) when:

- a) the Background Medium Scan Interval timer has expired; and
- b) the logical unit has been idle for the time specified in the MINIMUM IDLE TIME BEFORE BACKGROUND SCAN field (see 6.5.4).

After power on, if background pre-scan operations are not enabled (see 4.24.2.1), then the device server shall set the Background Medium Scan Interval timer to zero (i.e., expired).

Whenever a background medium scan operation begins, the device server shall set the Background Medium Scan Interval timer to the time specified in the BACKGROUND MEDIUM SCAN INTERVAL TIME field (see 6.5.4) and start the timer.

#### 4.24.3.2 Suspending and resuming background medium scan operations

The logical unit shall suspend a background medium scan operation if any of the following occurs:

- a) a command or task management function is processed that requires the background medium scan operation to be suspended;
- b) a SCSI event (e.g., a hard reset) (see SAM-5) is processed that requires the background medium scan operation to be suspended;

- c) a power condition timer expires (see the Power Condition mode page in SPC-5), and the PM\_BG\_PRECEDENCE field in the Power Condition mode page is set to 10b;
- d) the S\_L\_FULL bit (see 6.5.4) is set to one, and the Background Scan log parameters (see 6.4.2) are all used; or
- e) an application client sets the EN\_BMS bit to zero (see 6.5.4).

If a command is received that requires a background medium scan operation to be suspended, then the following should occur within the time specified in the MAXIMUM TIME TO SUSPEND BACKGROUND SCAN field (see 6.5.4):

- a) the logical unit suspends the background medium scan operation; and
- b) and the device server begins processing the command.

If a background pre-scan operation is suspended, then the device server shall not stop:

- a) the Background Medium Scan Interval timer; and
- b) any process that results in an event that causes a background function to occur (e.g., not stop any timers or counters associated with background functions).

The logical unit shall resume a suspended background medium scan operation from where the operation was suspended when:

- a) there are no commands in any task set to be processed;
- b) there are no task management functions to be processed;
- c) there are no SCSI events to be processed;
- d) the PM\_BG\_PRECEDENCE field in the Power Condition mode page is set to 10b (see SPC-5), but no power condition timer defined in the Power Condition mode page has expired;
- e) the S\_L\_FULL bit (see 6.5.4) is set to zero, or the Background Medium Scan log parameters in the Background Scan Results log page are not all used;
- f) the EN\_BMS (see 6.5.4) is set to one; and
- g) the logical unit has been idle for the time specified in the MINIMUM IDLE TIME BEFORE BACKGROUND SCAN field (see 6.5.4).

#### 4.24.3.3 Halting background medium scan operations

The device server shall halt background medium scan operations if any of the following occurs:

- a) the background medium scan operation completes scanning all logical blocks on the medium;
- b) the device server detects a fatal error;
- c) the device server detects a vendor specific pattern of errors;
- d) the device server detects a medium formatted without a PLIST (see 4.13); or
- e) the device server detects temperature out of range.

After background medium scan operations have been halted, the device server shall not enable a background medium scan operation until the conditions in 4.24.3.1 are met.

#### 4.24.4 Interpreting the logged background scan results

An application client may:

- a) poll the Background Scan Results log page (see 6.4.2) to get information about background pre-scan and background medium scan activity; or
- b) use the EBACKERR bit and the MRIE field (see 6.5.8) to select a method of indicating that a medium error was detected.

If the EBACKERR bit is set to one and a medium error was detected, then the device server shall return the following additional sense codes using the method defined by the value in the MRIE field:

- a) WARNING - BACKGROUND PRE-SCAN DETECTED MEDIUM ERROR, if the failure occurs during a background pre-scan operation; or
- b) WARNING - BACKGROUND MEDIUM SCAN DETECTED MEDIUM ERROR, if the failure occurs during a background medium scan operation.

The Background Scan Status log parameter (see 6.4.2.2) in the Background Scan Results log page (see 6.4.2) indicates:

- a) whether or not a background scan operation is active or halted;
- b) the number of background scan operations that have been performed on the medium; and
- c) the progress of a background scan operation, if active.

This information may be used by an application client to monitor the background scan operations and should be used by an application client after notification via an informational exception (see 6.5.8).

The Background Scan Results log parameters (see 6.4.2.3), if any, in the Background Scan Results log page describe the LBA and the reassignment status of each logical block that generated recovered errors or unrecovered errors during the background scan's read medium operations.

After an application client analyzes the Background Scan Results log parameters and has completed actions, if any, to repair any of the indicated LBAs, the application client may delete all Background Scan Results parameters by issuing a LOG SELECT command (e.g., with the PCR bit set to one in the CDB or with the PC field set to 11b and the PARAMETER LIST LENGTH field set to zero in the CDB) (see SPC-5).

A background medium scan operation may continue to run during log page accesses. To ensure that the values in the Background Scan Results log page do not change during a sequence of accesses, the application client:

- 1) sets the EN\_BMS bit to zero in the Background Control mode page in order to suspend the background medium scan operation;
- 2) reads the Background Scan Results log page with a LOG SENSE command;
- 3) processes the Background Scan Results log page;
- 4) deletes the Background Scan Results log page entries with the LOG SELECT command (e.g., with the PCR bit set to one in the CDB); and
- 5) sets the EN\_BMS bit to one in the Background Control mode page in order to re-enable the background scan operation.

## 4.25 Association between commands and CbCS permission bits

Table 28 defines the CbCS permissions required for each command defined in this standard. The permissions shown in table 28 are defined in the PERMISSIONS BIT MASK field in the CbCS capability descriptor in a CbCS extension descriptor (see SPC-4). This standard does not define any permissions specific to block commands.

**Table 28 — Associations between commands and CbCS permissions** (part 1 of 2)

Command	Permissions bit mask bits <sup>a</sup>				
	DATA READ	DATA WRITE	PARM READ	PARM WRITE	PHY ACC
COMPARE AND WRITE		1			
FORMAT UNIT		1		1	
GET LBA STATUS			1		
ORWRITE (16) / (32)		1			
POPULATE TOKEN	1				
PRE-FETCH (10) / (16)	1				
PREVENT ALLOW MEDIUM REMOVAL					1
READ (10) / (12) / (16) / (32)	1				
READ CAPACITY (10) / (16)			1		
READ DEFECT DATA (10) / (12)			1		
READ LONG (10) / (16)	1				
REASSIGN BLOCKS					1
RECEIVE ROD TOKEN INFORMATION	See SPC-4				
REPORT REFERRALS			1		
START STOP UNIT					1
SYNCHRONIZE CACHE (10) / (16)		1			
UNMAP		1			
VERIFY (10) / (12) / (16) / (32)	1				
WRITE (10) / (12) / (16) / (32)		1			
WRITE AND VERIFY (10) / (12) / (16) / (32)		1			
WRITE ATOMIC (16) / (32)		1			
<sup>a</sup> A device server shall only process a command shown in this table as specified by the CDB field of an extended CDB (see SPC-4) that contains a CbCS capability descriptor when all of the bits marked with a "1" in the row for that command are set to one in the PERMISSIONS BIT MASK field in that descriptor. The permissions bits represented by the empty cells in a row are ignored. If a device server receives a command specified by the CDB field of an extended CDB that does not contain the CbCS capability descriptor with all of the bits set to one as defined in this table, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.					



**Table 28 — Associations between commands and CbCS permissions (part 2 of 2)**

Command	Permissions bit mask bits <sup>a</sup>				
	DATA READ	DATA WRITE	PARM READ	PARM WRITE	PHY ACC
WRITE LONG (10) / (16)		1			
WRITE SAME (10) / (16) / (32)		1			
WRITE STREAM (16) / (32)		1			
WRITE USING TOKEN		1			
XDWRITEREAD (10) / (32)	1	1			
XPWRITE (10) / (32)		1			
<sup>a</sup> A device server shall only process a command shown in this table as specified by the CDB field of an extended CDB (see SPC-4) that contains a CbCS capability descriptor when all of the bits marked with a "1" in the row for that command are set to one in the PERMISSIONS BIT MASK field in that descriptor. The permissions bits represented by the empty cells in a row are ignored. If a device server receives a command specified by the CDB field of an extended CDB that does not contain the CbCS capability descriptor with all of the bits set to one as defined in this table, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.					

#### 4.26 Deferred microcode activation

After receiving a FORMAT UNIT command (see 5.4) or a START STOP UNIT command (see 5.26), a device server shall, prior to processing the command, activate any deferred microcode that has been downloaded as a result of a WRITE BUFFER command with the MODE field set to 0Eh (see SPC-5).

## 4.27 Model for uninterrupted sequences on LBA ranges

Direct access block devices may perform commands that require an uninterrupted sequence of actions to be performed on a specified range of LBAs. The uninterrupted sequence requirements are described in table 29. The uninterrupted sequences do not impact the processing of commands that access logical blocks other than those specified in the command requiring an uninterrupted sequence. The task attribute (see SAM-5) controls interactions between multiple commands. Commands with uninterrupted sequences on LBA ranges are shown in table 29.

**Table 29 — Commands that require uninterrupted sequences**

Command	Consistency enforcement	Reference
ORWRITE (16) ORWRITE (32)	The device server shall not perform any operations requested by any other command in the task set on logical blocks in the range specified by the command that requires an uninterrupted sequence of actions while performing the specified uninterrupted sequence of actions.	4.29
XDWRITEREAD (10) XDWRITEREAD (32)		5.53
XPWRITE (10) XPWRITE (32)		5.55
COMPARE AND WRITE	<p>The device server shall not perform:</p> <ul style="list-style-type: none"> <li>a) any operations requested by any COMPARE AND WRITE command in the task set on logical blocks in the range specified by the command that requires an uninterrupted sequence of actions while performing the specified uninterrupted sequence of actions;</li> <li>b) any write operations to or unmap operations on logical blocks in the range specified by the command that requires an uninterrupted sequence of actions while performing the read operations specified in the uninterrupted sequence of actions; and</li> <li>c) any read operations or verify operations from logical blocks in the range specified by the command that requires an uninterrupted sequence of actions while performing the write operations specified in the uninterrupted sequence of actions.</li> </ul>	5.3

## 4.28 Referrals

### 4.28.1 Referrals overview

Referrals allow a logical unit to inform an application client that one or more user data segments (i.e., ranges of logical blocks) are accessible through target port group(s).

Support for referrals is indicated by the device server setting the R\_SUP bit to one in the Extended INQUIRY Data VPD page (see SPC-5).

An application client may determine information on referrals by:

- a) issuing commands; or
- b) monitoring sense data returned as part of a completed command or a terminated command.

Figure 12 shows an example of how a logical unit informs an application client that one or more user data segments are accessible through target port groups.

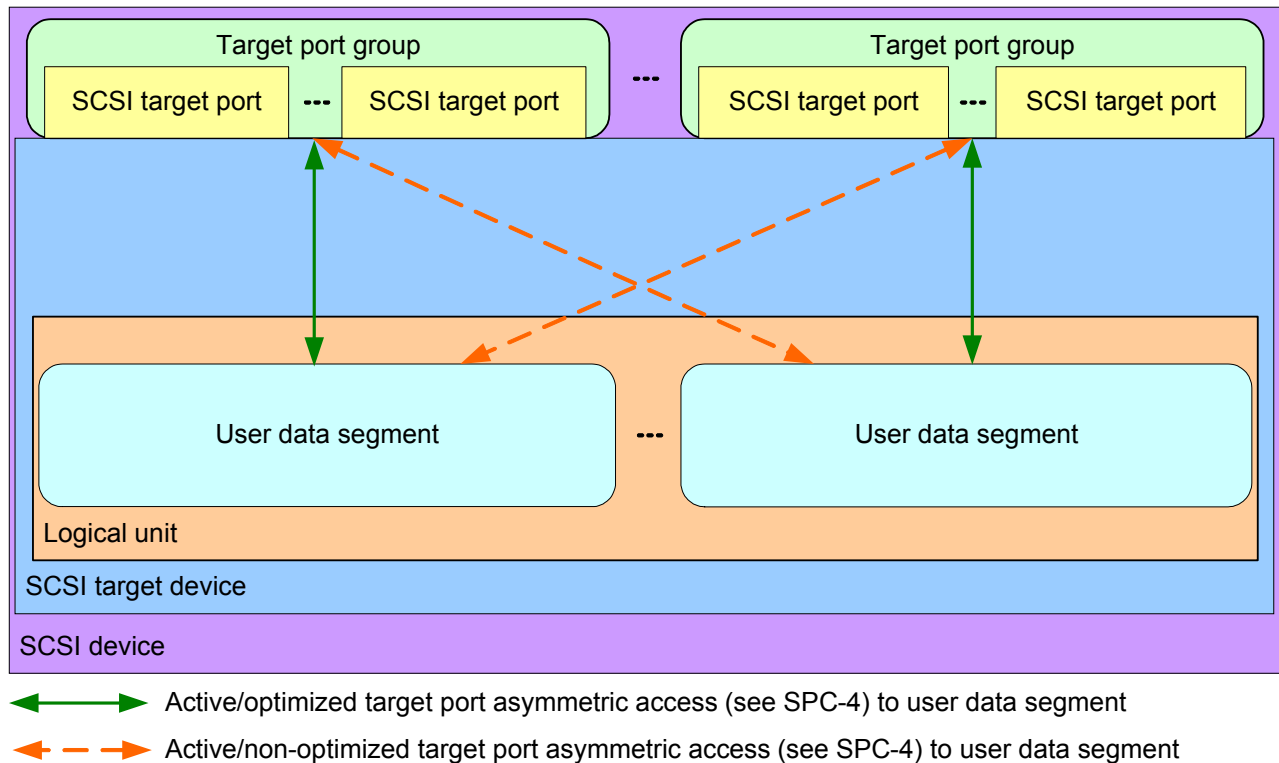


Figure 12 — Referrals

#### 4.28.2 Discovering referrals

An application client may determine referrals information on a logical unit by:

- 1) determining if the `R_SUP` bit is set to one (i.e., the logical unit supports referrals) in the Extended INQUIRY Data VPD page (see SPC-5);
- 2) requesting the user data segment information from the Referrals VPD page (see 6.6.7);
- 3) requesting a list of target port groups by issuing a REPORT TARGET PORT GROUPS command (see SPC-5); and
- 4) either:
  - A) requesting referrals information by issuing a REPORT REFERRALS command (see 5.23); or
  - B) monitoring for referral information in sense data returned by the device server (see 4.28.3).

The following calculation is used to determine the first LBA for each user data segment within the range of LBAs indicated by the user data segment referral descriptors (see table 18) returned in:

- a) the REPORT REFERRALS parameter data (see table 93); or
- b) the user data segment referrals sense data descriptor (see 4.18.4):

$$\text{first LBA of the current user data segment} = \text{first LBA} + (\text{segment size} \times \text{segment multiplier})$$

where:

first LBA	the initial value is the first user data segment LBA specified in the user data segment referral descriptor (see table 18). Subsequent values, if any, are the first LBA of the previous user data segment;
segment size	the content of the USER DATA SEGMENT SIZE field (see 6.6.7); and
segment multiplier	the content of the USER DATA SEGMENT MULTIPLIER field (see 6.6.7).

If the content of the USER DATA SEGMENT SIZE field is greater than zero, and the content of the USER DATA SEGMENT MULTIPLIER field is greater than zero, then the following calculation may be used to determine the last LBA for each user data segment within the range of LBAs indicated by the user data segment referral descriptors (see table 18) returned in:

- a) the REPORT REFERRALS parameter data (see table 93); or
- b) the user data segment referrals sense data descriptor (see 4.18.4):

$$\text{last LBA of the current user data segment} = \text{first LBA} + (\text{segment size} - 1)$$

where:

first LBA      the first LBA of the current user data segment;  
 segment size      the content of the USER DATA SEGMENT SIZE field (see 6.6.7).

If the content of the USER DATA SEGMENT SIZE field is zero, then there is only one user data segment, and the last LBA of that user data segment is equal to the last LBA specified in the last USER DATA SEGMENT LBA field (see table 18).

See annex G for examples for discovering referrals.

#### 4.28.3 Referrals in sense data

Returning referral information in sense data is enabled if the:

- a) R\_SUP bit in the Extended INQUIRY Data VPD page is set to one (i.e., the logical unit supports referrals) (see SPC-5); and
- b) D\_SENSE bit in the Control mode page is set to one (i.e., returning descriptor formatted sense data is enabled) (see SPC-5).

If reporting of referrals in sense data is enabled, a command completes without error, no other sense data is available within the logical unit, and the device server has an alternate I\_T\_L nexus that an application client should use to access at least one of the specified logical blocks, then the device server shall complete the command with GOOD status with the sense key set to COMPLETED, the additional sense code set to INSPECT REFERRALS SENSE DESCRIPTORS, and a user data segment referrals sense data descriptor (see 4.18.4).

The user data segment referral sense data descriptor (see 4.18.4) shall define the description of as many complete user data segments (i.e., one user data segment referral descriptor contains one complete user data segment) that fit in the maximum number of bytes allowed for sense data (i.e., 244 bytes or the maximum supported sense data length indicated in the Extended INQUIRY Data VPD page (see SPC-5)). If all the user data segments do not fit within the maximum number of bytes allowed for sense data, then:

- a) the device server shall set the NOT\_ALL\_R bit to one in the user data segment referral sense data descriptor (see 4.18.4); and
- b) the selection of which user data segments to include in the user data segment referral sense data descriptor is vendor specific.

Each user data segment referral sense data descriptor (see 4.18.4) contains information on alternate I\_T\_L nexuses to user data segments that the application client should use to access LBAs within the LBA range(s) indicated by the user data segments.

If reporting of referrals in sense data is enabled, the device server receives a command for which the device server is not able to access user data associated with the requested command, and the inaccessible user data is accessible through another target port group, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ABORTED COMMAND, the additional sense code set to INSPECT REFERRALS SENSE DESCRIPTORS, and a user data segment referral sense data descriptor. The user data segment referral sense data descriptor shall, at a minimum, indicate the user data segment that contains the LBA of the first inaccessible logical block. Any other type of error that occurs while processing the command shall take precedence and be reported as described in this standard. If any other type of error occurs while the device server is processing the command, then processing that error shall take precedence over processing the command, and the device server shall report the error as described in this standard.

If reporting of referrals in sense data is disabled (see 4.28.1), the device server receives a command for which the device server is not able to access user data associated with the requested command, and the inaccessible user data is accessible through another target port group, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to HARDWARE ERROR and the additional sense code set to INTERNAL TARGET FAILURE.

## 4.29 ORWRITE commands

### 4.29.1 ORWRITE commands overview

The ORWRITE commands (see 5.7 and 5.8) provide a mechanism for an application client to manipulate bitmap structures on direct access block devices.

An ORWRITE command shall be processed by the device server performing the following as an uninterrupted sequence of actions (see 4.27):

- 1) perform read operations from the LBAs specified by this command;
- 2) transfer the specified number of logical blocks from the Data-Out Buffer;
- 3) perform the specified Boolean arithmetic function on:
  - A) the user data contained in the logical blocks from read operations; and
  - B) the user data contained in the logical blocks transferred from the Data-Out Buffer;
- 4) store the results of the Boolean arithmetic function in a bitmap buffer;
- 5) generate new protection information, if any, from the stored results;
- 6) store the generated protection information, if any, into the bitmap buffer; and
- 7) perform write operations using the updated logical block data from the bitmap buffer.

If the check of the protection information from the read operations is successful (see table 55), and the check of the protection information transferred from the Data-Out Buffer is successful (see table 56), then the device server shall generate the new protection information (see 4.22) as follows:

- a) set the LOGICAL BLOCK GUARD field to the CRC (see 4.22.4) generated from the bitmap buffer by the device server;
- b) set the LOGICAL BLOCK REFERENCE TAG field to the LOGICAL BLOCK REFERENCE TAG field received from the Data-Out Buffer; and
- c) set the LOGICAL BLOCK APPLICATION TAG field to the LOGICAL BLOCK APPLICATION TAG field received from the Data-Out Buffer.

In order to support the manipulation of bitmap structures:

- a) the ORWRITE (16) command supports the set operation (see 4.29.4); and
- b) the ORWRITE (32) command supports:
  - A) the set operation; and
  - B) the change generation and clear operation (see 4.29.3).

### 4.29.2 ORWgeneration code

#### 4.29.2.1 ORWgeneration code overview

The ORWRITE commands use a generation code for synchronization. The device server shall establish and maintain the following generation codes:

- a) a current ORWgeneration code; and
- b) a previous ORWgeneration code.

Subsequent ORWRITE command processing by the device server is dependent on comparisons involving the ORWgeneration codes. Changes in these ORWgeneration codes define a synchronization point in the management of the bitmap.

#### 4.29.2.2 ORWgeneration code processing

The device server shall maintain at least one current ORWRITE processing policy. The device server may support more than one ORWRITE processing policy (see table 30 in 4.29.4).

When processing an ORWRITE (32) command (see 5.8), the device server compares the value in the EXPECTED ORWGENERATION field in the CDB to the current ORWgeneration code (see 4.29.2.1), and:

- a) if the two values are equal, then the device server continues processing the ORWRITE (32) command as described in table 30 for the set operation and as described in 4.29.3 for the change generation and clear operation; or
- b) if the two values are not equal, then:
  - A) for a set operation, the current ORWRITE processing policy (see table 30) determines how the device server continues processing the ORWRITE (32) command; and
  - B) for a change generation and clear operation, the device server terminates the ORWRITE (32) command (see 4.29.3).

If the device server supports both the ORWRITE (16) command (see 5.7) and the ORWRITE (32) command, then the device server shall process all ORWRITE (16) commands as if they contained an EXPECTED ORWGENERATION field set to zero.

#### 4.29.3 Change generation and clear operation

The change generation portion of the change generation and clear operation is used to establish a point of synchronization. The clear portion of the change generation and clear operation is used to set zero or more bits in the bitmap structure to zero.

The device server performs a change generation and clear operation if:

- a) the BMOP field (see 5.8) is set to 001b; and
- b) the value in the EXPECTED ORWGENERATION field is equal to the current ORWgeneration code in the device server.

If the device server performs a change generation and clear operation, then the device server shall perform the following as an uninterrupted sequence:

- 1) perform read operations from the LBAs specified by this command;
- 2) transfer the specified logical blocks from the Data-Out Buffer;
- 3) perform an AND operation (see 3.1.4) on the user data contained in the logical blocks from the read operations and the user data contained in the logical blocks transferred from the Data-Out Buffer;
- 4) store the results of the AND operation in a bitmap buffer;
- 5) generate new protection information, if any, from the stored results;
- 6) store the generated protection information, if any, into the bitmap buffer;
- 7) perform write operations using the updated logical block data from the bitmap buffer;
- 8) set the current ORWRITE processing policy to the value in the PREVIOUS GENERATION PROCESSING field (see 5.8);
- 9) set the previous ORWGeneration code (see 4.29.2) to the current ORWgeneration code in the device server; and
- 10) set the current ORWGeneration code (see 4.29.2) to the value in the NEW ORWGENERATION field(see 5.8).

If the value in the EXPECTED ORWGENERATION field is not equal to the current ORWgeneration code, then the device server shall terminate the ORWRITE (32) command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to ORWRITE GENERATION DOES NOT MATCH.

If a power on or hard reset condition occurs, then the device server shall set:

- a) the current ORWGeneration code to zero;
- b) the previous ORWGeneration code to zero; and
- c) the current ORWRITE processing policy to 7h.

The device server shall preserve the following across a logical unit reset:

- a) the current ORWGeneration code;
- b) the previous ORWGeneration code; and
- c) the current ORWRITE processing policy.

#### 4.29.4 Set operation

The set operation is used to set zero or more bits in the bitmap structure to one.

The device server performs a set operation for an ORWRITE command (see 5.7 and 5.8) if the BMOP field in the CDB is set to 000b.

The device server shall perform a set operation by performing the actions specified in table 30, which shows the current ORWgeneration code, the previous ORWgeneration code, and the device server's current ORWRITE processing policy (see 4.29.3).

**Table 30 — Performing an ORWRITE set operation**

Current ORWRITE processing policy	The value in the EXPECTED ORWGENERATION field matches		
	Current ORWgeneration code	Previous ORWgeneration code	Any other value
0h	PA	PA	CCG
1h	Reserved		
2h	PA	DN	CCG
3h	PA	PA	PA
4h	Reserved		
5h	PA	DN	DN
6h	Reserved		
7h	PA	CCG	CCG
8h to Fh	Reserved		
<p>Key:</p> <p>PA = the device server shall perform the following as an uninterrupted sequence:</p> <ul style="list-style-type: none"><li>1) perform read operations from the LBAs specified by the command;</li><li>2) transfer the specified logical blocks from the Data-Out Buffer;</li><li>3) perform an OR operation on the logical blocks from the read operations and the user data contained in the logical blocks transferred from the Data-Out Buffer;</li><li>4) store the results of the OR operation in a bitmap buffer;</li><li>5) generate new protection information, if any, from the stored results;</li><li>6) store the generated protection information, if any, into the bitmap buffer; and</li><li>7) perform write operations using the updated logical block data (i.e., those containing the user data resulting from the OR operation and the generated protection information, if any) from the bitmap buffer.</li></ul> <p>DN = the device server shall discard the contents of the Data-Out Buffer and shall complete the command with GOOD status.</p> <p>CCG = the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to ORWRITE GENERATION DOES NOT MATCH</p>			

## 4.30 Block device ROD token operations

### 4.30.1 Block device ROD token operations overview

Application clients request that block device ROD token operations (see SPC-5) be performed using the commands summarized in this subclause or the commands specified in SPC-5.

Copy managers (see SPC-5) that implement the POPULATE TOKEN command (see 5.9) or the WRITE USING TOKEN command (see 5.52) shall implement the following:

- a) the POPULATE TOKEN command;
- b) the WRITE USING TOKEN command;
- c) the RECEIVE COPY STATUS (LID4) command (see SPC-5);
- d) the RECEIVE ROD TOKEN INFORMATION command (see SPC-5 and 5.22); and
- e) the Third-party Copy VPD page (see 6.6.8) containing at least one Block Device ROD Token Limits descriptor (see 6.6.8.3).

The POPULATE TOKEN command may cause the copy manager to create zero or one point in time ROD tokens. If the POPULATE TOKEN command causes one point in time ROD token to be created, then this point in time ROD token may be retrieved by an application client using the RECEIVE ROD TOKEN INFORMATION command.

The WRITE USING TOKEN command causes the copy manager to transfer the data represented by the specified ROD token (i.e., the data represented by the ROD token retrieved using the RECEIVE ROD TOKEN INFORMATION command or the data represented by the block device zero ROD token).

The copy manager manages the point in time ROD token.

After the copy manager begins processing a POPULATE TOKEN command or a WRITE USING TOKEN command, the copy manager shall preserve information for return in response to a RECEIVE ROD TOKEN INFORMATION command as defined in SPC-5.

Block device range descriptor lists (see 5.9.3) contain non-overlapping block device range descriptors and are used by the application client to specify:

- a) the logical blocks to include in the data represented by the ROD token;
- b) the sequence of the logical blocks in the data represented by the ROD token (e.g., the first logical block represented by the LBA described in the first block device range descriptor is placed at the beginning of the data represented by the ROD token, and the first logical block represented by the LBA described in the second block device range descriptor is placed in the data represented by the ROD token immediately following the last logical block represented by the LBA described in the first block device range descriptor);
- c) the logical blocks to be written from the data represented by the ROD token; and
- d) the sequence of the logical blocks written from the data represented by the ROD token (e.g., the first logical block represented by the LBA described in the first block device range descriptor is written from the beginning of the data represented by the ROD token, and the first logical block represented by the LBA described in the second block device range descriptor is written from data represented by the ROD token immediately following the data written to the last logical block represented by the LBA described in the first block device range descriptor).

If the copy manager uses out of order transfers to create the representation of data for the ROD token, then the TRANSFER COUNT field in the parameter data returned in the response to a RECEIVE ROD TOKEN INFORMATION command with a list identifier that specifies a POPULATE TOKEN command (see 5.22.2) shall be based only on the contiguous transfers that complete without error starting at the first LBA specified by the first block device range descriptor (i.e., any transfers completed without error beyond the first incomplete or unsuccessful transfer shall not contribute to the computation of the value in the TRANSFER COUNT field).

If the copy manager uses out of order transfers to write from the data represented by the ROD token, then the TRANSFER COUNT field in the parameter data returned in response to a RECEIVE ROD TOKEN INFORMATION command with a list identifier that specifies a WRITE USING TOKEN command (see 5.22.3)



shall be based only on the contiguous transfers that complete without error starting at the first LBA specified by the first block device range descriptor (i.e., any transfers completed without error beyond the first incomplete or unsuccessful transfer shall not contribute to the computation of the value in the TRANSFER COUNT field).

#### 4.30.2 POPULATE TOKEN command and WRITE USING TOKEN command completion

As part of completing a block device token operation originated by a POPULATE TOKEN command (see 5.9) or a WRITE USING TOKEN command (see 5.52), the copy manager shall compute the residual by subtracting the sum of the contents of the NUMBER OF LOGICAL BLOCK fields in all of the complete block device range descriptors of the parameter list (see 5.9.3) from the TRANSFER COUNT field in the parameter data returned in response to a RECEIVE ROD TOKEN INFORMATION command (see 5.22.2 and 5.22.3).

If the POPULATE TOKEN command was received with the IMMED bit set to zero, and the residual is negative, then the copy manager shall:

- a) terminate the command with CHECK CONDITION status, with the additional sense code set to COPY TARGET DEVICE DATA UNDERRUN, the sense key set to:
  - A) COPY ABORTED, if the transfer count is not zero; or
  - B) ILLEGAL REQUEST, if the transfer count is zero,
 and report the transfer count in the INFORMATION field (see SPC-5); or
- b) complete the command with GOOD status and return parameter data for the RECEIVE ROD TOKEN INFORMATION command received on the same I\_T nexus with a matching LIST IDENTIFIER field with:
  - A) the COPY OPERATION STATUS field set to 03h or 60h (see SPC-5);
  - B) the EXTENDED COPY COMPLETION STATUS field set to CHECK CONDITION (see SAM-5); and
  - C) the SENSE DATA field with the additional sense code set to COPY TARGET DEVICE DATA UNDERRUN, the sense key set to:
    - a) COPY ABORTED, if the transfer count is not zero; or
    - b) ILLEGAL REQUEST, if the transfer count is zero;
 and report the transfer count in the INFORMATION field (see SPC-5).

If the WRITE USING TOKEN command was received with the IMMED bit set to zero, and the residual is negative, then the copy manager shall terminate the command with CHECK CONDITION status, the additional sense code set to COPY TARGET DEVICE DATA UNDERRUN, the sense key set to:

- a) COPY ABORTED, if the transfer count is not zero; or
- b) ILLEGAL REQUEST, if the transfer count is zero,

and report the transfer count in the INFORMATION field (see SPC-5).

If the POPULATE TOKEN command or WRITE USING TOKEN command was received with the IMMED bit set to one, and the residual is negative, then the copy manager shall return parameter data for the RECEIVE ROD TOKEN INFORMATION command received on the same I\_T nexus with a matching LIST IDENTIFIER field with:

- a) the COPY OPERATION STATUS field set to 03h or 60h (see SPC-5);
- b) the EXTENDED COPY COMPLETION STATUS field set to CHECK CONDITION status (see SAM-5); and
- c) the SENSE DATA field with the additional sense code set to COPY TARGET DEVICE DATA UNDERRUN, the sense key set to:
  - A) COPY ABORTED, if the transfer count is not zero; or
  - B) ILLEGAL REQUEST, if the transfer count is zero,
 and report the transfer count in the INFORMATION field (see SPC-5).

#### 4.30.3 Block device specific ROD tokens

Block device specific ROD token types (see SPC-5) are shown in table 31.

**Table 31 — ROD token type values**

ROD token type	Description	Reference
FF00_0000h to FFFF_0000h	Reserved	
FFFF_0001h	Block device zero ROD token	4.30.4
FFFF_0002h to FFFF_FFEFh	Reserved	

#### 4.30.4 Block device zero ROD token

The block device zero ROD token represents user data in which all bits are set to zero and protection information, if any, is set to FFFF\_FFFF\_FFFF\_FFFFh. If user data with all bits set to zero and protection information, if any, set to FFFF\_FFFF\_FFFF\_FFFFh, is represented by a ROD token, then, in response to a RECEIVE ROD TOKEN INFORMATION command in which the LIST IDENTIFIER field specifies a POPULATE TOKEN command, the copy manager may or may not return a ROD token that is the block device zero ROD token. The block device zero ROD token format is shown in table 32.

**Table 32 — Block device zero ROD token format**

Bit	7	6	5	4	3	2	1	0
Byte								
0	(MSB)							
...	ROD TOKEN TYPE (FFFF_0001h)							
3								
4	Reserved							
5								
6	(MSB)							
7	ROD TOKEN LENGTH (01F8h)							
8	Reserved							
...								
511								

The logical block length associated with the block device zero ROD token is that of the direct access block device to which the data is being written (e.g., if a block device zero ROD token is used to write to a logical unit that has 642-byte logical blocks, then the logical block length of the block device zero ROD token is 642 bytes).

The ROD TOKEN TYPE field is defined in SPC-5 and shall be set to the value shown in table 32 for the block device zero ROD token.

The ROD TOKEN LENGTH field is defined in SPC-5 and shall be set to the value shown in table 32 for the block device zero ROD token.

#### 4.30.5 ROD token device type specific data

The device type specific data for ROD tokens (see SPC-5) created by a copy manager for a direct access block devices (see 6.7):

- a) provides information about the logical unit at the time that the ROD token was created; and

- b) is a subset of the parameter data returned by the READ CAPACITY (16) command (see 5.18) for the logical unit that contains the copy manager that created the ROD token.

If the READ CAPACITY (16) parameter data changes so that the copy manager that created the ROD token is no longer able to access the data represented by the ROD token, then that copy manager shall invalidate the ROD token.

## 4.31 Atomic writes

### 4.31.1 Atomic writes overview

An atomic write command performs one or more atomic write operations. The following write commands are atomic write commands:

- a) WRITE ATOMIC (16) (see 5.43); and
- b) WRITE ATOMIC (32) (see 5.44).

To perform an atomic write operation, the device server:

- a) ensures that any subsequent read operation returns either all of the write data or none of the write data for the atomic write operation as described in 4.31.2;
- b) performs operations as described in 4.31.3; and
- c) processes ACA conditions as described in 4.31.4.

All commands that are not specifically identified as atomic commands are non-atomic commands.

The MAXIMUM ATOMIC TRANSFER LENGTH field (see 6.6.4) indicates the maximum transfer length for atomic write commands that specify an atomic boundary set to zero. If an atomic write command specifies an atomic boundary set to zero and a transfer length greater than the maximum atomic transfer length, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB, and the field pointer pointing to the transfer length field in the CDB.

The MAXIMUM ATOMIC TRANSFER LENGTH WITH ATOMIC BOUNDARY field (see 6.6.4) indicates the maximum transfer length for atomic write commands that specify an atomic boundary set to a non-zero value. If an atomic write command specifies a non-zero atomic boundary and the transfer length is greater than the maximum atomic transfer length with atomic boundary, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST, the additional sense code set to INVALID FIELD IN CDB, and the field pointer pointing to the transfer length field in the CDB.

The MAXIMUM ATOMIC BOUNDARY SIZE field (see 6.6.4) indicates the maximum atomic boundary for atomic write commands. If an atomic write command specifies a non-zero atomic boundary that is greater than the maximum atomic boundary size, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST, the additional sense code set to INVALID FIELD IN CDB, and the field pointer pointing to the ATOMIC BOUNDARY field in the CDB.

If the starting LBA of an atomic write command does not meet the requirements of the ATOMIC ALIGNMENT field (see 6.6.4), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

If the ATOMIC TRANSFER LENGTH GRANULARITY field (see 6.6.4) contains a non-zero value and:

- a) the TRANSFER LENGTH field in an atomic write command is not a multiple of the atomic transfer length granularity; or
- b) the ATOMIC BOUNDARY field in an atomic write command contains a non-zero value that is not a multiple of the atomic transfer length granularity,

then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

If an atomic write command specifies a non-zero atomic boundary that is less than or equal to the maximum atomic boundary size and a transfer length that is less than or equal to the maximum atomic transfer length

with atomic boundary, then the device server shall perform one or more atomic write operations each of which is a size that is a multiple of the size specified in the ATOMIC BOUNDARY field.

If multiple atomic write operations are performed as a result of a non-zero atomic boundary, then:

- a) the atomic write operations may be performed in any order; and
- b) if the device server is unable to complete all atomic write operations, then the device server shall terminate the command with the sense key set to ABORTED COMMAND and the additional sense code set to INCOMPLETE MULTIPLE ATOMIC WRITE OPERATIONS.

All operations that are not specifically identified as atomic operations are non-atomic operations.

If the device server supports atomic write commands, then the device server shall process commands as described in SAM-5 with the additional restrictions described in 4.31.

#### **4.31.2 Atomic write operations that do not complete**

If the device server is not able to successfully complete an atomic write operation (e.g., the command is terminated or aborted), then the device server shall ensure that none of the LBAs specified by the atomic write operation have been altered by any logical block data from the atomic write operation (i.e., the specified LBAs return logical block data as if the atomic write operation had not occurred).

If a power loss causes loss of logical block data from an atomic write operation in a volatile write cache that has not yet been stored on the medium, then the device server shall ensure that none of the LBAs specified by the atomic write operation have been altered by any logical block data from the atomic write operation (i.e., the specified LBAs return logical block data as if the atomic write operation had not occurred and writes from the cache to the medium preserve the specified atomicity).

EXAMPLE - If the device server processes an atomic write command that specifies a non-zero atomic boundary using five atomic write operations in such a way that:

- 1) two atomic write operations are performed successfully;
- 2) an atomic write operation encounters an error; and
- 3) the remaining two atomic write operations are not performed,

then subsequent read commands of the LBAs written by the two successful atomic write operations return new logical block data and subsequent read commands of the LBAs specified by the other three atomic write operations return old logical block data.

#### **4.31.3 Performing operations with respect to atomic operations**

##### **4.31.3.1 Performing operations before and after an atomic write operation**

Before an atomic write operation completes, the device server shall not use any logical block data from the atomic write operation for any other operations (e.g., read operations return old logical block data until the atomic write operation completes).

After an atomic write operation completes, the device server shall use all of the logical block data from the atomic write operation as the basis for subsequent operations (e.g., read operations return logical block data from the atomic write operation).

#### 4.31.3.2 Performing operations during an atomic write operation

Table 33 defines the device server requirements for starting to perform an atomic write operation that has an overlapping LBA with an operation that the device server is already performing.

**Table 33 — Performing atomic operations with overlapping LBAs during current operations**

Operation A that is currently being processed by the device server is a	Device server is requested to perform an atomic write operation B that accesses an LBA that is being accessed by operation A
read or verify	The device server shall perform one of the following: a) should complete operation A before starting operation B; or b) may terminate operation A with CHECK CONDITION status with a sense key of ABORTED COMMAND and an additional sense code set to OVERLAPPING ATOMIC COMMAND IN PROGRESS before starting operation B.
write or unmap	
format or sanitize	The device server shall abort operation B as specified in 3.4.11.2 and 5.3.1
Uninterruptible sequence of operations	The device server shall wait for operation A to complete before it starts performing operation B.
atomic write	

If the device server is performing an atomic write operation and is requested to perform another operation that has an LBA that overlaps with the atomic operation being performed, then the device server shall complete the current atomic write operation before performing the requested operation.

EXAMPLE - While the device server is performing an atomic write operation that accesses LBAs 0 to 15, it may perform a second atomic write operation accessing LBAs 16 to 31 and shall not perform a third atomic write operation accessing LBAs 7 to 15. After the first atomic write operation completes, the device server may perform the third atomic write operation.

EXAMPLE - While the device server is performing an atomic write operation that accesses LBAs 0 to 15, if a second write operation accessing LBAs 7 to 31 is available for performing, then the device server may perform the write operation's accesses to LBAs 16 to 31 and shall not perform the write operation's accesses to LBAs 7 to 15. After the atomic write operation completes, the device server may perform the write operation's accesses to LBAs 7 to 15.

Performing an atomic write operation does not prevent the device server from performing, at any time, other operations that only contain non-overlapping LBAs.

#### 4.31.4 Processing ACA conditions during atomic write commands

If another command causes an ACA condition while an atomic write command is being processed, then the device server:

- a) shall ensure that none of the LBAs specified by the atomic write operation have been altered by any logical block data from the atomic write operation (i.e., the specified LBAs return logical block data, as if the atomic write operation had not occurred); and
- b) after the ACA condition is cleared:
  - A) may terminate the atomic write command with CHECK CONDITION status with a sense key of ABORTED COMMAND and an additional sense code set to ATOMIC COMMAND ABORTED DUE TO ACA; or
  - B) may process the atomic write operation.

## 4.32 IO Advice Hints

### 4.32.1 IO Advice Hints Overview

IO advice hints allow application clients to provide device servers with information about anticipated LBA usage patterns. A device server uses IO advice hints to provide access to logical blocks that is appropriate to the specified usage patterns and optimize other aspects of the direct access block device (e.g., power consumption and performance). The degree to which the device server optimizes the direct access block device based on anticipated LBA usage patterns may not be measurable by the application client.

EXAMPLE - A device server that implements the IO advice hints features described in this subclause may provide improved performance for cases where the access pattern matches what was specified by the applicable IO advice hint (e.g., LBAs that are read in a sequential manner may exhibit better performance when the IO advice hint indicates sequential access).

IO advice hints are associated with individual commands that contain a GROUP NUMBER field (see 4.32.2). The device server is not required to retain this association in a way that produces repeatable results (e.g., recorded metadata associated with each LBA specified by the command). Device servers may provide the features described in this subclause by other means (e.g., a device server may direct the transferred data and the associated LBAs to an area of the media that has properties that match the specified IO advice hints).

### 4.32.2 Specifying IO Advice Hints

An application client may specify an IO advice hint by applying the IO advice hints extension of the grouping function using the GROUP NUMBER field in the CDB of certain commands. The GROUP NUMBER field in each command determines the IO advice hint used during the processing of that command. At the time a command that specifies an IO advice hint is received, the device server should process that command in a manner that produces the effects described in 4.32.1. This standard does not define a method for the device server to save or return IO advice hints after a command that specifies an IO advice hint has been processed.

IO advice hints are provided in most data transfer commands. If a device server supports IO advice hints, then the device server shall implement the:

- a) IO Advice Hints Grouping mode page (see 6.5.7); and
- b) grouping function (see 4.23.1) and extensions for IO advice hints (see 4.23.2).

## 4.33 Background Operation Control

Background Operation Control is managed by information in the Background Control Mode page (see 6.5.4). A device server may have requirements to perform advanced background operations if the percentage of device resources available for allocation reaches the minimum percentage indicated in the MINIMUM PERCENTAGE field (see 6.6.6). The device server may provide an indication that these resources have reached a value (e.g., minimum percentage plus 10%) that is specified by an armed decreasing threshold percentage (see 4.7.3.7.3). If the device server does not process a BACKGROUND CONTROL command that requests advanced background operations and the percentage of device resources available for allocation reach the minimum percentage indicated in the MINIMUM PERCENTAGE field (see 6.6.6), then the device server may perform advanced background operations without a request from the application client.

EXAMPLE - Advanced background operation may include NAND block erase operations, media read operations, and media write operations (e.g., garbage collection), which may impact response time for normal read requests or write requests from the application client.

If the application client is able to predict idle time when there are few read requests and few write requests to the device server, then the application client may notify the device server about this idle time so that the device server may perform advanced background operations. As a result, advanced background operations are minimally overlapped with normal read commands and normal write commands from the application

client. The BACKGROUND CONTROL command (see 5.2) provides the mechanism for the application client to communicate this information to a logical unit.

The logical block provisioning thresholds (see 4.7.3.7) provides a mechanism for the device server to establish a unit attention condition to notify application clients when a related logical block provisioning threshold is crossed.

A device server that supports Background Operation Control as described in this subclause:

- a) shall be a resource provisioned device as described in 4.7.3.2 or a thin provisioned device as described in 4.7.3.3;
- b) shall set the BOCS bit to one in the Block Device Characteristics VPD page (see 6.6.2);
- c) shall support the Logical Block Provisioning VPD page (see 6.6.6);
- d) should support logical block provisioning thresholds with the threshold type field set to 001b and threshold resource value set to 01h (see 4.7.3.7.3);
- e) shall support the Background Operation log page (see 6.4.3);
- f) shall support the Background Control mode page (see 6.5.4); and
- g) shall support the BACKGROUND CONTROL command (see 5.2).

The device server performs notification to the application client of decreasing provisioning resource percentage using logical block provisioning thresholds (see 4.7.3.7.3). Performing this notification informs the application client that the application client should determine a time to perform advanced background operations and then request the device server to perform these advanced background operations using the BACKGROUND CONTROL command.

If the device server reaches the minimum percentage of resources available as indicated in the MINIMUM PERCENTAGE field (see 6.6.6), then the device server may begin performing device server initiated advanced background operations. As a result, the application client should set the threshold percentage sufficiently high to allow the application client to receive a notification and specify to the device server to perform host initiated advanced background operations before the device reaches the minimum percentage.

If the device server processes a BACKGROUND CONTROL command with the BO\_CTL field set to 01b (i.e., start host initiated background control operations), then the device server shall:

- 1) initialize the bo\_time timer to the value specified in the BO\_TIME field (see 5.2) and start the bo\_time timer; and
- 2) perform host initiated advanced background operations until:
  - A) the bo\_time timer expires;
  - B) the device server determines that all necessary advanced background operations are completed;
  - or
  - C) the device server processes a BACKGROUND CONTROL command with the BO\_CTL field set to 10b (i.e., stop host initiated background control operations).

If the device server processes a BACKGROUND CONTROL command with the BO\_CTL field set to 10b (i.e., stop host initiated background control operations) and the device server is not performing host initiated advanced background operations, then the device server shall not considered this an error.

If the device server processes a BACKGROUND CONTROL command with the BO\_CTL field set to 10b (i.e., stop host initiated background control operations) and the device server is performing host initiated advanced background operations, then the device server shall stop performing host initiated advanced background operations.

If the device server processes a BACKGROUND CONTROL command with the BO\_CTL field set to 00b, then the device server shall:

- a) not change any host initiated advanced background operations;
- b) ignore the value in the BO\_TIME field; and
- c) not considered this an error.

### 4.34 Stream Control

Stream control is limited by the MAXIMUM NUMBER OF STREAMS field, the OPTIMAL STREAM WRITE SIZE field, and the STREAM GRANULARITY SIZE field (see 6.6.5). Logical units that support stream control shall be resource provisioned logical units as described in 4.7.3.2 or thin provisioned logical units as described in 4.7.3.3.

Logical units that implement the Stream Control model are comprised of physical blocks that are arranged in a manner such that background operations to prepare resources for future allocations to LBAs are performed on those physical blocks. STREAM WRITE commands with a transfer length that matches the value specified in the OPTIMAL STREAM WRITE SIZE field (see 6.6.5), and with an initial LBA that is either zero or a multiple of the value specified in the OPTIMAL STREAM WRITE SIZE field provide optimal performance from the device.

See annex E for an example method in which application clients may use alignment information to determine optimal performance for stream writes.

A device server that supports Stream Control as described in this subclause shall:

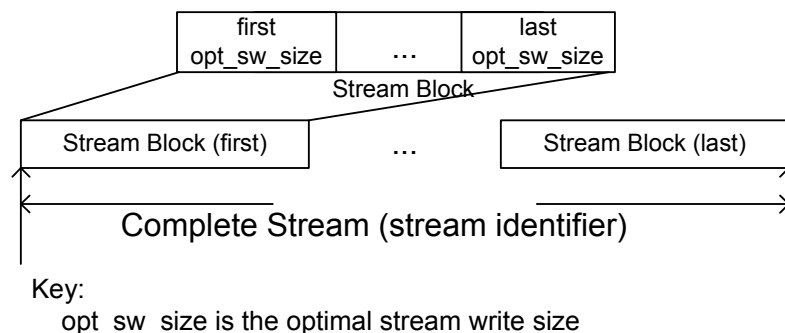
- a) support the STREAM CONTROL command (see 5.27);
- b) Support the WRITE STREAM (16) command (see 5.50) or the WRITE STREAM (32) command (see 5.51);
- c) support the GET STREAM STATUS command (see 5.6); and
- d) support the Block Limits Extension VPD page (see 6.6.5).

The number of optimal stream write size blocks that are prepared as a unit for future allocation to LBAs is indicated in the STREAM GRANULARITY SIZE field (see 6.6.5). The application client should write an integer multiple of stream blocks to the logical unit before closing a stream.

Stream Control provides control of streams, each of which are comprised of user data associated with a single data structure (e.g., a single file or a complete database) or data that has the same lifetime. An application client opens a stream using the STREAM CONTROL command and then uses the stream identifier provided by the device server for all subsequent writes associated with that stream. Each stream is associated with a stream identifier that is associated with the stream from when the stream is opened until the stream is closed. Valid stream identifiers are 0001h to FFFFh. A stream allows the device server to place the data associated with a stream identifier into locations that:

- a) are part of one or more stream blocks; and
- b) do not contain any data not associated with that stream identifier.

For optimal performance, writes to the stream should be aligned to and a multiple of the length specified by the OPTIMAL STREAM WRITE SIZE field indicated in the Block Limits Extension VPD page (see 6.6.5). Figure 13 shows the relationship of data blocks of optimal stream write size within a stream block, and multiple stream blocks within a complete stream.



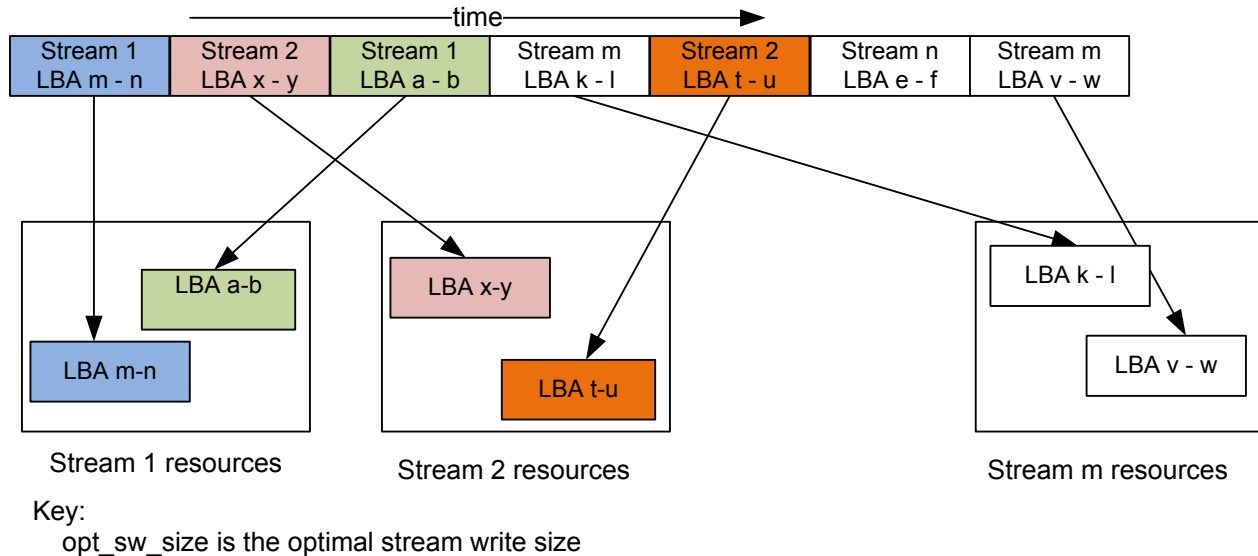
**Figure 13 — Stream Block Relationships**

Figure 14 shows an example assignment of data blocks received for four different streams to associated device server resources where:

- a) Stream 1 LBA m - n is one optimal stream write size in length;



- b) Stream 2 LBA x - y is one optimal stream write size in length;
- c) Stream 1 LBA a - b is two optimal stream write sizes in length;
- d) Stream 2 LBA t - u is one optimal stream write size in length;
- e) Stream m LBA k - l is one optimal stream write size in length;
- f) Stream m LBA e - f is one optimal stream write size in length; and
- g) Stream n mapping is not shown.



**Figure 14 — Multiple streams example**

The application client opens and closes streams using the STREAM CONTROL command (see 5.27) and requests write operations to a stream using the WRITE STREAM (16) command (see 5.50) or the WRITE STREAM (32) command (see 5.51). To use a stream for data the application client:

- 1) sends a STREAM CONTROL command with the STR\_CTL field set to 01b (i.e., open);
- 2) waits for the stream identifier value returned in the ASSIGNED\_STR\_ID field (see 5.27);
- 3) sends one or more WRITE STREAM commands with the STR\_ID field set to the associated stream identifier;
- 4) waits for responses to all WRITE STREAM commands for the stream associated with the stream identifier; and
- 5) sends a STREAM CONTROL command with the STR\_CTL field set to 10b (i.e., close) and the STR\_ID field set to the associated stream identifier.

If the device server processes a WRITE STREAM command with the STR\_ID field set to a stream identifier of a stream that is not open, then the device server shall terminate the WRITE STREAM command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to STREAM NOT OPEN.

If the device server processes a STREAM CONTROL command with the STR\_CTL field set to 10b (i.e., close) and the STR\_ID field set to the stream identifier of a stream that is not open, then the device server shall terminate the STREAM CONTROL command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to STREAM NOT OPEN.

If the device server processes a STREAM CONTROL command with the STR\_CTL field set to 01b (i.e., open) and the maximum number of streams as defined by the MAXIMUM NUMBER OF STREAMS field (see 6.6.5) are already open, then the device server shall terminate the STREAM CONTROL command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to MAXIMUM NUMBER OF STREAMS OPEN.

The application client may discover the state of open streams using the GET STREAM STATUS command (see 5.6).

Resources (e.g., stream identifier) allocated to maintaining a specific stream are released if:

- a) the device server processes a STREAM CONTROL command (see 5.27) with the stream identifier of the specified stream and the STR\_CTL field set to 10b (i.e., close);
- b) a hard reset occurs;
- c) a logical unit reset occurs;
- d) a power-on occurs;
- e) a FORMAT UNIT command is processed; or
- f) a sanitize operation is processed.

If physical blocks that are part of a stream block do not contain logical block data when the associated stream is closed, then logical block data should not be written to those physical blocks until all LBAs in that stream block are unmapped.

## 4.35 [Format operations](#)

### 4.35.1 [Format operations overview](#)

[A format operation results in the device server:](#)

- a) [configuring the logical block length and number of logical blocks of the logical unit as specified by the block descriptor \(see 6.5.2\); and](#)
- b) [performing the following as specified by a FORMAT UNIT command \(see 5.4\):](#)
  - A) [configure protection information;](#)
  - B) [perform defect management;](#)
  - C) [initialize LBAs; and](#)
  - D) [vendor specific medium certification.](#)

[The degree that the medium is altered by a format operation is vendor specific. A format operation is requested by a FORMAT UNIT command.](#)

### 4.35.2 [Performing a format operation](#)

[Before performing a format operation, the device server shall stop all:](#)

- a) [enabled power condition timers \(see SPC-5\);](#)
- b) [timers for enabled background scan operations \(see 4.24\); and](#)
- c) [timers or counters enabled for device-specific background functions.](#)

[After the format operation is complete, the device server shall reinitialize and restart all enabled timers and counters for power conditions and background functions.](#)

[While performing a format operation, the device server shall:](#)

- a) [process commands already in a task set when a FORMAT UNIT command is received in a vendor specific manner;](#)
- b) [process an INQUIRY command by returning parameter data based on the condition of the logical unit before beginning the FORMAT UNIT command \(i.e., INQUIRY data shall not change until after successful completion of a format operation\);](#)
- c) [process a REQUEST SENSE command by returning parameter data containing sense data with the sense key set to NOT READY, the additional sense code set to LOGICAL UNIT NOT READY, FORMAT IN PROGRESS, and the PROGRESS INDICATION field in the sense data \(see SPC-5\) set to indicate the progress of the format operation;](#)
- d) [process REPORT LUNS commands;](#)
- e) [terminate all commands, except INQUIRY commands, REPORT LUNS commands, and REQUEST SENSE commands, with CHECK CONDITION status with the sense key set to NOT READY and the additional sense code set to LOGICAL UNIT NOT READY, FORMAT IN PROGRESS; and](#)

- f) remove all Background Scan Results log parameters (see 6.4.2.3) from the Background Scan Results log page, if supported, and remove all Pending Defect log parameters (see 6.4.8.3) from the Pending Defects log page, if supported.

The application client may specify:

- a) that the device server clear the existing GLIST;
- b) a list of address descriptors that the device server adds to the GLIST;
- c) that the device server enable a certification operation that adds address descriptors for physical blocks with medium defects discovered during the certification operation to the GLIST; and
- d) the behavior of the device server if it is not able to:
  - A) access the PLIST or GLIST; or
  - B) determine whether the PLIST or GLIST exists.

#### **4.35.3 Completing a format operation**

##### **4.35.3.1 Completing a format operation overview**

If a format operation completes without error, then:

- a) stream resources, if any, shall be released;
  - b) if the logical unit is fully provisioned (i.e., the LBPME bit (see 5.18.2) is set to zero), then all LBAs in the logical unit are mapped (see 4.7.2); or
  - c) if the logical unit supports logical block provisioning management (i.e., the LBPME bit is set to one), then:
    - A) if the LBPRZ field (see 6.6.6) is set to 000b, then each LBA in the logical unit shall be either:
      - a) mapped, if an initialization pattern was specified that does not match the vendor-specific data returned by a read command for an unmapped LBA (see 4.7.4.4); or
      - b) unmapped (see 4.7.3.2 and 4.7.3.3), if no initialization pattern was specified or an initialization pattern was specified that matches the vendor-specific data returned by a read command for an unmapped LBA (see 4.7.4.4);
    - B) if the LBPRZ field (see 6.6.6) is set to xx1b, then each LBA in the logical unit:
      - a) shall be mapped, if the format operation did not initialize the user data to all zeroes for the logical block referenced by that LBA;
      - b) shall be unmapped (see 4.7.3.2 and 4.7.3.3), if the format operation initialized the user data to all zeroes for the logical blocks referenced by all valid LBAs in the logical unit; or
      - c) may be unmapped (see 4.7.3.2 and 4.7.3.3), if the format operation initialized the user data to all zeroes for the logical block referenced by that LBA, and the format operation did not initialize the user data to all zeroes for the logical blocks referenced by all valid LBAs in the logical unit;
- and
- C) if the LBPRZ field (see 6.6.6) is set to 010b, then each LBA in the logical unit:
    - a) shall be mapped, if an initialization pattern was specified that does not match the provisioning initialization pattern; or
    - b) shall be unmapped (see 4.7.3.2 and 4.7.3.3), if no initialization pattern was specified or an initialization pattern was specified that matches the provisioning initialization pattern.

If a format operation is aborted or completes with an error, then the logical unit may become format corrupt. Format corrupt shall be cleared by a format operation that completes without error. If the logical unit is format corrupt, then the device server shall terminate any medium access command with CHECK CONDITION status, with the sense key set to MEDIUM ERROR and the additional sense code set to MEDIUM FORMAT CORRUPTED.

#### 4.35.3.2 Completing read commands after a successful format operation

##### 4.35.3.2.1 Completing read commands overview

Following a successful format operation and before a write operation to an LBA, a read command or verify command that specifies that LBA shall be processed by the device server as described in 4.35.3.2.2, 4.35.3.2.3, and 4.35.3.2.4.

##### 4.35.3.2.2 With FFMT field set to 00b

If the FFMT field (see table 40) was set to 00b in the most recent successful FORMAT UNIT command, then subsequent read commands or verify commands (see 4.35.3.2.1) that complete without error are processed using:

- a) the user data set as specified by:
  - A) the initialization pattern, if any;
  - B) the provisioning initialization pattern, if applicable; or
  - C) the manufacturer's default initialization pattern;
- and
- b) the protection information, if any, set to FFFF\_FFFF\_FFFF\_FFFFh.

##### 4.35.3.2.3 With FFMT field set to 01b

If the FFMT field (see table 40) was set to 01b in the most recent successful FORMAT UNIT command, then subsequent read commands or verify commands (see 4.35.3.2.1):

- a) with unrecovered medium errors are processed as described in 4.18.1;
- b) with pseudo unrecovered errors are processed as described in 4.18.2;
- c) should be processed using unspecified logical block data and complete without error, if protection information is disabled; or
- d) if protection information is enabled, then:
  - A) may be processed using unspecified logical block data and complete without error; or
  - B) may terminate with CHECK CONDITION status with sense data that indicates that the protection information check fails as defined in table 67.

##### 4.35.3.2.4 With FFMT field set to 10b

If the FFMT field (see table 40) was set to 10b in the most recent successful FORMAT UNIT command, then the device server may:

- a) return unspecified logical block data and complete subsequent read commands (see 4.35.3.2.1) without error;
- b) process unspecified logical block data and complete subsequent verify commands without error; or
- c) terminate subsequent read commands or verify commands with CHECK CONDITION status with the sense key set to HARDWARE ERROR, MEDIUM ERROR, or ABORTED COMMAND.

## 5 Commands for direct access block devices

### 5.1 Commands for direct access block devices overview

The commands for direct access block devices are listed in table 34.

**Table 34 — Commands for direct access block devices (part 1 of 5)**

Command	Operation code <sup>a</sup>	Type	LBACT	Reference
<del>ACCESS CONTROL IN</del>	<del>86h</del>	<del>Ø</del>	<del>n/a</del>	<del>SPC-4</del>
<del>ACCESS CONTROL OUT</del>	<del>87h</del>	<del>Ø</del>	<del>n/a</del>	<del>SPC-4</del>
ATA PASS-THROUGH (12)	A1h	O	n/a	SAT-3
ATA PASS-THROUGH (16)	85h	O	n/a	SAT-3
BACKGROUND CONTROL	9Eh/15h	O	n/a	5.2
CHANGE ALIASES	A4h/0Bh	O	n/a	SPC-5
CLOSE ZONE	94h/01h	X	n/a	ZBC
COMPARE AND WRITE	89h	O	R, W	5.3
<u>COPY OPERATION ABORT</u>	<u>83h/1Ch</u>	<u>O</u>	<u>n/a</u>	<u>SPC-5</u>
EXTENDED COPY	83h	O	n/a	SPC-5
FINISH ZONE	94h/02h	X	n/a	ZBC
FORMAT UNIT	04h	M	Z	5.4
GET LBA STATUS	9Eh/12h	O	n/a	5.5
GET STREAM STATUS	9Eh/16h	O	n/a	5.6
INQUIRY	12h	M	n/a	SPC-5
LOG SELECT	4Ch	O	n/a	SPC-5
LOG SENSE	4Dh	O	n/a	SPC-5
MAINTENANCE IN	A3h/00h to 04h A3h/06h to 09h	X	n/a	SPC-5 SCC-2
MAINTENANCE OUT	A4h/00h to 05h A4h/07h to 09h	X	n/a	SPC-5 SCC-2
MODE SELECT (6)	15h	O	n/a	SPC-5
MODE SELECT (10)	55h	O	n/a	SPC-5
MODE SENSE (6)	1Ah	O	n/a	SPC-5
MODE SENSE (10)	5Ah	O	n/a	SPC-5
Key: O = optional M = mandatory X = implementation requirements are defined in the reference R = read command U = unmap command V = verify command W = write command Z = other command <del>PI</del> = <del>protection information</del> LBACT= logical block access command type (see 4.2.2)				
<sup>a</sup> If a command is defined by a combination of operation code and service action, then the operation code value is shown preceding a slash and the service action value is shown after the slash.				

Table 34 — Commands for direct access block devices (part 2 of 5)

Command	Operation code <sup>a</sup>	Type	LBACT	Reference
OPEN ZONE	94h/03h	X	n/a	ZBC
ORWRITE (16)	8Bh	O	R, W	5.7
ORWRITE (32)	7Fh/000Eh	O	R, W	5.8
PERSISTENT RESERVE IN	5Eh	O	n/a	SPC-5
PERSISTENT RESERVE OUT	5Fh	O	n/a	SPC-5
POPULATE TOKEN	83h/10h	O	n/a	5.9
PRE-FETCH (10)	34h	O	R	5.10
PRE-FETCH (16)	90h	O	R	5.11
PREVENT ALLOW MEDIUM REMOVAL	1Eh	O	n/a	5.12
READ (10)	28h	M	R	5.13
READ (12)	A8h	O	R	5.14
READ (16)	88h	M	R	5.15
READ (32)	7Fh/0009h	O	R	5.16
READ ATTRIBUTE	8Ch	O	n/a	SPC-5
READ BUFFER (10)	3Ch	O	n/a	SPC-5
<a href="#">READ BUFFER (16)</a>	<a href="#">9Eh/10h</a>	<a href="#">O</a>	<a href="#">n/a</a>	<a href="#">SPC-5</a>
READ CAPACITY (10)	25h	M	n/a	5.17
READ CAPACITY (16)	9Eh/10h	M	n/a	5.18
READ DEFECT DATA (10)	37h	O	n/a	5.19
READ DEFECT DATA (12)	B7h	O	n/a	5.20
REASSIGN BLOCKS	07h	O	Z	5.21
<a href="#">READ MEDIA SERIAL NUMBER</a>	<a href="#">ABh/01h</a>	<a href="#">O</a>	<a href="#">n/a</a>	<a href="#">SPC-5</a>
<a href="#">RECEIVE COPY DATA</a>	<a href="#">84h/06h</a>	<a href="#">O</a>	<a href="#">n/a</a>	<a href="#">SPC-5</a>
RECEIVE COPY <a href="#">STATUS</a> <del>RESULTS</del>	<a href="#">84h/05h</a>	O	n/a	SPC-5
RECEIVE DIAGNOSTIC RESULTS	1Ch	X	n/a	SPC-5 6.3 <del>4.30</del>
RECEIVE ROD TOKEN INFORMATION	84h/07h	X	n/a	SPC-5 <a href="#">4.30</a> 5.22
REDUNDANCY GROUP IN	BAh	X	n/a	SCC-2
REDUNDANCY GROUP OUT	BBh	X	n/a	SCC-2
REMOVE I_T NEXUS	A4h/0Ch	O	n/a	SPC-5
Key: O = optional M = mandatory X = implementation requirements are defined in the reference R = read command U = unmap command V = verify command W = write command Z = other command <del>P</del> = <del>protection information</del> LBACT= logical block access command type (see 4.2.2)				
<sup>a</sup> If a command is defined by a combination of operation code and service action, then the operation code value is shown preceding a slash and the service action value is shown after the slash.				

Table 34 — Commands for direct access block devices (part 3 of 5)

Command	Operation code <sup>a</sup>	Type	LBACT	Reference
REPORT ALIASES	A3h/0Bh	O	n/a	SPC-5
<a href="#">REPORT ALL ROD TOKENS</a>	<a href="#">84h/08h</a>	<a href="#">O</a>	<a href="#">n/a</a>	<a href="#">SPC-5</a>
REPORT IDENTIFYING INFORMATION	A3h/05h	O	n/a	SPC-5
REPORT LUNS	A0h	M	n/a	SPC-5
REPORT PRIORITY	A3h/0Eh	O	n/a	SPC-5
REPORT PROVISIONING INITIALIZATION PATTERN	A3h/1Dh	O	n/a	5.24
REPORT REFERRALS	9Eh/13h	O	n/a	5.23
REPORT SUPPORTED OPERATION CODES	A3h/0Ch	O	n/a	SPC-5
REPORT SUPPORTED TASK MANAGEMENT FUNCTIONS	A3h/0Dh	O	n/a	SPC-5
REPORT TARGET PORT GROUPS	A3h/0Ah	O	n/a	SPC-5
<a href="#">REPORT TIMESTAMP</a>	<a href="#">A3h/0Fh</a>	<a href="#">O</a>	<a href="#">n/a</a>	<a href="#">SPC-5</a>
REPORT ZONES	95h/00h	X	n/a	ZBC
REQUEST SENSE	03h	M	n/a	SPC-5
RESET WRITE POINTER	94h/04h	X	n/a	ZBC
SANITIZE	48h	O	Z	5.25
SECURITY PROTOCOL IN	A2h	O	n/a	SPC-5
SECURITY PROTOCOL OUT	B5h	O	n/a	SPC-5
SEND DIAGNOSTIC	1Dh	O	n/a	SPC-5
SET IDENTIFYING INFORMATION	A4h/06h	O	n/a	SPC-5
SET PRIORITY	A4h/0Eh	O	n/a	SPC-5
SET TARGET PORT GROUPS	A4h/0Ah	O	n/a	SPC-5
<a href="#">SET TIMESTAMP</a>	<a href="#">A4h/0Fh</a>	<a href="#">O</a>	<a href="#">n/a</a>	<a href="#">SPC-5</a>
SPARE IN	BCh	X	n/a	SCC-2
SPARE OUT	BDh	X	n/a	SCC-2
START STOP UNIT	1Bh	O	n/a	5.26
STREAM CONTROL	9Eh/14h	O	n/a	5.27
SYNCHRONIZE CACHE (10)	35h	O	W	5.28
SYNCHRONIZE CACHE (16)	91h	O	W	5.29
TEST UNIT READY	00h	M	n/a	SPC-5
Key: O = optional M = mandatory X = implementation requirements are defined in the reference R = read command U = unmap command V = verify command W = write command Z = other command <del>P</del> = <del>protection information</del> LBACT= logical block access command type (see 4.2.2)				
<sup>a</sup> If a command is defined by a combination of operation code and service action, then the operation code value is shown preceding a slash and the service action value is shown after the slash.				

Table 34 — Commands for direct access block devices (part 4 of 5)

Command	Operation code <sup>a</sup>	Type	LBACT	Reference
UNMAP	42h	X	U	5.30 4.7
VERIFY (10)	2Fh	O	V, W	5.31
VERIFY (12)	AFh	O	V, W	5.32
VERIFY (16)	8Fh	O	V, W	5.33
VERIFY (32)	7Fh/000Ah	O	V, W	5.34
VOLUME SET IN	BEh	X	n/a	SCC-2
VOLUME SET OUT	BFh	X	n/a	SCC-2
WRITE (10)	2Ah	O	W	5.35
WRITE (12)	AAh	O	W	5.36
WRITE (16)	8Ah	O	W	5.37
WRITE (32)	7Fh/000Bh	O	W	5.38
WRITE AND VERIFY (10)	2Eh	O	V, W	5.39
WRITE AND VERIFY (12)	AEh	O	V, W	5.40
WRITE AND VERIFY (16)	8Eh	O	V, W	5.41
WRITE AND VERIFY (32)	7Fh/000Ch	O	V, W	5.42
WRITE ATOMIC (16)	9Ch	O	W	5.43
WRITE ATOMIC (32)	7Fh/000Fh	O	W	5.44
WRITE ATTRIBUTE	8Dh	O	n/a	SPC-5
WRITE BUFFER	3Bh	O	n/a	SPC-5
WRITE LONG (10)	3Fh	O	Z	5.45
WRITE LONG (16)	9Fh/11h	O	Z	5.46
WRITE SAME (10)	41h	X	U, W	5.47 4.7
WRITE SAME (16)	93h	X	U, W	5.48 4.7
WRITE SAME (32)	7Fh/000Dh	X	U, W	5.49 4.7
WRITE STREAM (16)	9Ah	O	W	5.50
WRITE STREAM (32)	7Fh/0010h	O	W	5.51
WRITE USING TOKEN	83h/11h	X	Z	5.52 4.30
XDWRITEREAD (10)	53h	O	R, W	5.53
<b>Key:</b> O = optional M = mandatory X = implementation requirements are defined in the reference R = read command U = unmap command V = verify command W = write command Z = other command <del>PI</del> = <del>protection information</del> LBACT= logical block access command type (see 4.2.2)				
<sup>a</sup> If a command is defined by a combination of operation code and service action, then the operation code value is shown preceding a slash and the service action value is shown after the slash.				



**Table 34 — Commands for direct access block devices** (part 5 of 5)

Command	Operation code <sup>a</sup>	Type	LBACT	Reference																				
XDWRITEREAD (32)	7Fh/0007h	O	R, W	5.54																				
XPWRITE (10)	51h	O	R, W	5.55																				
XPWRITE (32)	7Fh/0006h	O	R, W	5.56																				
<p>Note 1 - The following operation codes are obsolete: 01h, 08h, 0Ah, 0Bh, 16h, 17h, 18h, 2Bh, 30h, 31h, 32h, 33h, 36h, 39h, 3Ah, 3Eh, 40h, 50h, 52h, 56h, 57h, 80h, 81h, 82h, 92h, 9Eh/11h, A7h, B3h, and B4h.</p> <p>Note 2 - The following operation codes are vendor specific: 02h, 05h, 06h, 09h, 0Ch, 0Dh, 0Eh, 0Fh, 10h, 11h, 13h, 14h, 19h, 20h, 21h, 22h, 23h, 24h, 26h, 27h, 29h, 2Ch, 2Dh, and C0h to FFh.</p> <p>Note 3 - A complete summary of operation codes is available at <a href="http://www.t10.org/lists/2op.htm">http://www.t10.org/lists/2op.htm</a>. The summary includes information about obsolete commands.</p>																								
Key:																								
<table><tr><td>O</td><td>= optional</td><td>V</td><td>= verify command</td></tr><tr><td>M</td><td>= mandatory</td><td>W</td><td>= write command</td></tr><tr><td>X</td><td>= implementation requirements are defined in the reference</td><td>Z</td><td>= other command</td></tr><tr><td>R</td><td>= read command</td><td><del>PI</del></td><td><del>= protection information</del></td></tr><tr><td>U</td><td>= unmap command</td><td>LBACT</td><td>= logical block access command type (see 4.2.2)</td></tr></table>					O	= optional	V	= verify command	M	= mandatory	W	= write command	X	= implementation requirements are defined in the reference	Z	= other command	R	= read command	<del>PI</del>	<del>= protection information</del>	U	= unmap command	LBACT	= logical block access command type (see 4.2.2)
O	= optional	V	= verify command																					
M	= mandatory	W	= write command																					
X	= implementation requirements are defined in the reference	Z	= other command																					
R	= read command	<del>PI</del>	<del>= protection information</del>																					
U	= unmap command	LBACT	= logical block access command type (see 4.2.2)																					
<sup>a</sup> If a command is defined by a combination of operation code and service action, then the operation code value is shown preceding a slash and the service action value is shown after the slash.																								

## 5.2 BACKGROUND CONTROL command

### 5.2.1 BACKGROUND CONTROL command overview

The BACKGROUND CONTROL command (see table 35) is used to request that the device server start or stop host initiated advanced background operations (see 4.33), if any.

This command uses the SERVICE ACTION IN (16) CDB format (see A.2).

**Table 35 — BACKGROUND CONTROL command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (9Eh)							
1		Reserved			SERVICE ACTION (15h)				
2		BO_CTL		Reserved					
3		BO_TIME							
4		Reserved							
...									
14									
15		CONTROL							

The OPERATION CODE field and SERVICE ACTION field are defined in SPC-5 and shall be set to the value shown in table 35 for the BACKGROUND CONTROL command.

The background operation control (BO\_CTL) field specifies that the device server shall control host initiated advanced background operations as described in 4.33. The BO\_CTL field is described in table 36.

**Table 36 — BO\_CTL field**

Code	Description
00b	Do not change host initiated advanced background operations.
01b	Start host initiated advanced background operations.
10b	Stop host initiated advanced background operations.
11b	Reserved

The background operation time (BO\_TIME) field specifies the maximum time that the device server shall have to perform host initiated advanced background operations in units of 100 ms (see 4.33). The BO\_TIME field is ignored if the BO\_CTL field is not set to 01b. A BO\_TIME field set to 00h specifies that there is no limit to the time that the device server may perform host initiated advanced background operations.

The CONTROL byte is defined in SAM-5.

### 5.3 COMPARE AND WRITE command

The COMPARE AND WRITE command (see table 37) requests that the device server perform the following as an uninterrupted sequence of actions (see 4.27):

- 1) perform read operations from the specified LBAs;
- 2) perform a compare operation on only the user data (i.e., not on the protection information) from:
  - A) the read operations; and
  - B) the compare instance transferred from the Data-Out Buffer;
- 3) if the compare operation indicates a match, then perform the following operations:
  - 1) check the protection information, if any, in the write instance transferred from the Data-Out Buffer based on the contents of the WRPROTECT field as defined in table 119; and
  - 2) perform write operations to the LBAs specified by this command using the write instance;
 and
- 4) if the compare operation does not indicate a match, then terminate the command with CHECK CONDITION status with the sense key set to MISCOMPARE and the additional sense code set to MISCOMPARE DURING VERIFY OPERATION. In the sense data (see 4.18 and SPC-5) the offset from the start of the Data-Out Buffer to the first byte of data that was not equal shall be reported in the INFORMATION field.

The Data-Out Buffer contains two instances of logical block data:

- 1) the compare instance, in which:
  - A) the user data is used for the compare operation; and
  - B) the protection information, if any, is not used;
 and
- 2) the write instance, in which:
  - A) the user data is used for the write operations; and
  - B) the protection information, if any, is used for the write operations.

**Table 37 — COMPARE AND WRITE command**

Bit Byte	7	6	5	4	3	2	1	0
0	OPERATION CODE (89h)							
1	WRPROTECT			DPO	FUA	Reserved		
2	(MSB)							
...	LOGICAL BLOCK ADDRESS							
9	(LSB)							
10	Reserved							
...								
12								
13	NUMBER OF LOGICAL BLOCKS							
14	Reserved		GROUP NUMBER					
15	CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 37 for the COMPARE AND WRITE command.

See the WRITE (10) command for the definition of the WRPROTECT field.

See the READ (10) command (see 5.13) for the definition of the DPO bit. See the READ (10) command (see 5.13) for the definition of the FUA bit specifying behavior for the read operations. See the WRITE (10) command (see 5.35) for the definition of the FUA bit specifying behavior for the write operations.

The LOGICAL BLOCK ADDRESS field specifies the LBA of the first logical block (see 4.5) to be accessed by the device server (e.g., the first LBA accessed by both a read operation and a write operation). If the specified LBA exceeds the capacity of the medium, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to LOGICAL BLOCK ADDRESS OUT OF RANGE.

The NUMBER OF LOGICAL BLOCKS field specifies:

- the number of contiguous logical blocks on which read operations shall be performed, starting with the LBA specified by the LOGICAL BLOCK ADDRESS field;
- the number of contiguous logical blocks that shall be transferred from the Data-Out Buffer for the compare operation; and
- if the compare operation indicates a match, then the number of contiguous logical blocks that shall be transferred from the Data-Out Buffer and on which write operations shall be performed, starting with the LBA specified by the LOGICAL BLOCK ADDRESS field.

A NUMBER OF LOGICAL BLOCKS field set to zero specifies that no read operations shall be performed, no logical block data shall be transferred from the Data-Out Buffer, no compare operations shall be performed, and no write operations shall be performed. This condition shall not be considered an error.

If the specified LBA and the specified number of logical blocks exceed the capacity of the medium (see 4.5), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to LOGICAL BLOCK ADDRESS OUT OF RANGE.

If the number of logical blocks exceeds the value in the MAXIMUM COMPARE AND WRITE LENGTH field ~~in the Block Limits VPD page~~ (see 6.6.4), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

See the PRE-FETCH (10) command (see 5.10) and 4.23 for the definition of the GROUP NUMBER field.

The CONTROL byte is defined in SAM-5.

## 5.4 FORMAT UNIT command

### 5.4.1 FORMAT UNIT command overview

The FORMAT UNIT command (see table 38) requests that the device server [perform a format operation](#). ~~format the medium into application client accessible logical blocks as specified in the number of logical blocks and logical block length values received in the last mode parameter block descriptor (see 6.5.2) in a MODE SELECT command (see SPC 5). In addition, the device server may certify the medium and create control structures for the management of the medium and defects. The degree that the medium is altered by this command is vendor specific.~~

~~If a device server receives a FORMAT UNIT command before receiving a MODE SELECT command with a mode parameter block descriptor, then the device server shall use the number of logical blocks and logical block length at which the logical unit is currently formatted (i.e., no change is made to the number of logical blocks and the logical block length of the logical unit during the format operation).~~

The device server shall handle any deferred microcode as specified in 4.26.

~~Before performing the operation specified by this command, the device server shall stop all:~~

- ~~a) enabled power condition timers (see SPC 5);~~
- ~~b) timers for enabled background scan operations (see 4.24); and~~
- ~~c) timers or counters enabled for device specific background functions.~~

~~After the operation is complete, the device server shall reinitialize and restart all enabled timers and counters for power conditions and background functions.~~

~~While performing a format operation, the device server shall:~~

- ~~a) process commands already in a task set when a FORMAT UNIT command is received in a vendor-specific manner;~~
- ~~b) process an INQUIRY command by returning parameter data based on the condition of the logical unit before beginning the FORMAT UNIT command (i.e., INQUIRY data shall not change until after successful completion of a format operation);~~
- ~~c) process a REQUEST SENSE command by returning parameter data containing sense data with the sense key set to NOT READY, the additional sense code set to LOGICAL UNIT NOT READY, FORMAT IN PROGRESS, and the PROGRESS INDICATION field in the sense data (see SPC 5) set to indicate the progress of the format operation;~~
- ~~d) process REPORT LUNS commands;~~
- ~~e) terminate all commands, except INQUIRY commands, REPORT LUNS commands, and REQUEST SENSE commands, with CHECK CONDITION status with the sense key set to NOT READY and the additional sense code set to LOGICAL UNIT NOT READY, FORMAT IN PROGRESS; and~~
- ~~f) remove all Background Scan Results log parameters (see 6.4.2.3) from the Background Scan Results log page, if supported, and remove all Pending Defect log parameters (see 6.4.8.3) from the Pending Defects log page, if supported.~~

**NOTE 5** - ~~The progress indication is available to the application client during a format operation as a means to determine the logical unit's format progress.~~

~~The simplest form of the FORMAT UNIT command (i.e., a FORMAT UNIT command with no parameter data) accomplishes medium formatting with little application client control over medium defect management. The device server implementation determines the degree of medium defect management that is to be performed. Additional forms of this command increase the application client's control over medium defect management (see 4.13). The application client may specify:~~

- ~~a) that the device server clear the existing GLIST;~~
- ~~b) a list of address descriptors that the device server adds to the GLIST;~~

- c) ~~that the device server enable a certification operation that adds address descriptors for physical blocks with medium defects discovered during the certification operation to the GLIST; and~~
- d) ~~the behavior of the device server if it is not able to access the PLIST or GLIST or determine whether one or both of them exists.~~

Following a successful format operation:

- a) ~~stream resources, if any, shall be released;~~
- b) ~~if the logical unit is fully provisioned (i.e., the LBPME bit (see 5.18.2) is set to zero), then all LBAs in the logical unit are mapped (see 4.7.2); or~~
- c) ~~if the logical unit supports logical block provisioning management (i.e., the LBPME bit is set to one), then:~~
  - A) ~~if the LBPRZ field (see 6.6.6) is set to 000b, then each LBA in the logical unit shall be either:~~
    - a) ~~mapped, if an initialization pattern was specified that does not match the vendor specific data returned by a read command for an unmapped LBA (see 4.7.4.4); or~~
    - b) ~~unmapped (see 4.7.3.2 and 4.7.3.3), if no initialization pattern was specified or an initialization pattern was specified that matches the vendor specific data returned by a read command for an unmapped LBA (see 4.7.4.4);~~
  - B) ~~if the LBPRZ field (see 6.6.6) is set to xx1b, then each LBA in the logical unit:~~
    - a) ~~shall be mapped, if the format operation did not initialize the user data to all zeroes for the logical block referenced by that LBA;~~
    - b) ~~shall be unmapped (see 4.7.3.2 and 4.7.3.3), if the format operation initialized the user data to all zeroes for the logical blocks referenced by all valid LBAs in the logical unit; or~~
    - c) ~~may be unmapped (see 4.7.3.2 and 4.7.3.3), if the format operation initialized the user data to all zeroes for the logical block referenced by that LBA, and the format operation did not initialize the user data to all zeroes for the logical blocks referenced by all valid LBAs in the logical unit;~~
- ~~and~~
- C) ~~if the LBPRZ field (see 6.6.6) is set to 010b, then each LBA in the logical unit:~~
  - a) ~~shall be mapped, if an initialization pattern was specified that does not match the provisioning initialization pattern; or~~
  - b) ~~shall be unmapped (see 4.7.3.2 and 4.7.3.3), if no initialization pattern was specified or an initialization pattern was specified that matches the provisioning initialization pattern.~~

Table 38 defines the FORMAT UNIT command.

Table 38 — FORMAT UNIT command

Bit Byte	7	6	5	4	3	2	1	0
0	OPERATION CODE (04h)							
1	FMTPINFO		LONGLIST	FMTDATA	CMPLST	DEFECT LIST FORMAT		
2	Vendor specific							
3	Reserved							
4	<a href="#">Reserved</a>						<a href="#">FFMT</a>	
5	CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 38 for the FORMAT UNIT command.

The combination (see table 44) of the format protection information (FMTPINFO) field and the PROTECTION FIELD USAGE field (see 5.4.2.2) specifies whether or not the device server enables or disables the use of protection information.

A **LONGLIST** bit set to zero specifies that the parameter list, if any, contains a short parameter list header as defined in table 42. A **LONGLIST** bit set to one specifies that the parameter list, if any, contains a long parameter list header as defined in table 43. If the **FMTDATA** bit is set to zero, then the **LONGLIST** bit shall be ignored.

A format data (**FMTDATA**) bit set to one specifies that the **FORMAT UNIT** parameter list (see 5.4.2) shall be transferred from the Data-Out Buffer. A **FMTDATA** bit set to zero specifies that no parameter list be transferred from the Data-Out Buffer. If the **FMTDATA** bit is set to zero and the **FMTPINFO** field is not set to zero, then the device server shall terminate the command with **CHECK CONDITION** status with the sense key set to **ILLEGAL REQUEST** and the additional sense code set to **INVALID FIELD IN CDB**.

Following a successful format operation, the **PROT\_EN** bit and the **P\_TYPE** field (see 5.18.2) indicate the type of protection currently in effect on the logical unit.

If protection information is written during a format operation (i.e., the **FMTPINFO** field is set to a value greater than zero), then protection information shall be written to a default value of **FFFF\_FFFF\_FFFF\_FFFFh**.

A complete list (**CMPLST**) bit set to zero specifies that the device server shall add the defect list included in the **FORMAT UNIT** parameter list to the existing **GLIST** (see 4.13). A **CMPLST** bit set to one specifies that the device server shall replace the existing **GLIST** with the defect list, if any, included in the **FORMAT UNIT** parameter list.

If the **FMTDATA** bit is set to zero, then the **CMPLST** bit shall be ignored.

If the **FMTDATA** bit is set to one, then the **DEFECT LIST FORMAT** field specifies the format of the address descriptors in the defect list in the **FORMAT UNIT** parameter list.

Table 39 defines support requirements for address descriptors based on the combinations of the FMTDATA bit, the Cmplst bit, the DEFECT LIST FORMAT field, and the DEFECT LIST LENGTH field.

**Table 39 — FORMAT UNIT command address descriptor support requirements**

Field in the FORMAT UNIT CDB			DEFECT LIST LENGTH <sup>a</sup>	Support	Comments
FMTDATA	CMPLST	DEFECT LIST FORMAT			
0	any	000b	Not available	M	Vendor specific defect information
1	0	Either: a) 000b (i.e., short block address descriptor)(see 6.2.2); b) 001b (i.e., extended bytes from index address descriptor)(see 6.2.3); c) 010b (i.e., extended physical sector address descriptor)(see 6.2.4); d) 011b (i.e., long block address descriptor)(see 6.2.5); e) 100b (i.e., bytes from index address descriptor)(see 6.2.6); or f) 101b (i.e., physical sector address descriptor)(see 6.2.7)	Zero	O	See <sup>b</sup> and <sup>d</sup>
	1			O	See <sup>b</sup> and <sup>e</sup>
	0		Non-zero	O	See <sup>c</sup> and <sup>d</sup>
	1	O		See <sup>c</sup> and <sup>e</sup>	
	0		110b (i.e., vendor specific)		Zero
	Non-zero	O		See <sup>c</sup> and <sup>d</sup>	
	1	Zero		O	See <sup>b</sup> and <sup>e</sup>
		Non-zero		O	See <sup>c</sup> and <sup>e</sup>
All other combinations				Reserved	
<sup>a</sup> This field is in the parameter list header. <sup>b</sup> No defect list is included in the parameter list. <sup>c</sup> A defect list is included in the parameter list. <sup>d</sup> The device server retains the existing GLIST. <sup>e</sup> The device server discards the existing GLIST.					

The fast format (FFMT) feild is described in table 40

**Table 40 — FFMT field description**

<u>Code</u>	<u>Description</u>	<u>Support</u>
<u>00b</u>	<u>The device server initializes the medium (see 4.10) as specified in the CDB and parameter list before completing the format operation. After successful completion of the format operation, read commands and verify commands are processed as described in 4.35.3.2.1 and 4.35.3.2.2.</u>	<u>Mandatory</u>
<u>01b</u>	<u>The device server initializes the medium (see 4.10) without overwriting the medium (i.e., resources for managing medium access are initialized and the medium is not written) before completing the format operation. After successful completion of the format operation, read commands and verify commands are processed as described in 4.35.3.2.1 and 4.35.3.2.3.</u>  <u>If the device server determines that the options specified in this FORMAT UNIT command are incompatible with the read command and verify command requirements described in 4.35.3.2.3, then the device server shall not perform the format operation and shall terminate the FORMAT UNIT command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FAST FORMAT COMBINATION.</u>	<u>Optional</u>
<u>10b</u>	<u>The device server initializes the medium (see 4.10) without overwriting the medium (i.e., resources for managing medium access are initialized and the medium is not written) before completing the format operation. After successful completion of the format operation, read commands and verify commands are processed as described in 4.35.3.2.1 and 4.35.3.2.4.</u>	<u>Optional</u>
<u>10b</u>	<u>Reserved</u>	

The CONTROL byte is defined in SAM-5.

#### 5.4.2 FORMAT UNIT parameter list

##### 5.4.2.1 FORMAT UNIT parameter list overview

Table 41 defines the FORMAT UNIT parameter list.

**Table 41 — FORMAT UNIT parameter list**

Bit	7	6	5	4	3	2	1	0
Byte								
0	Parameter list header (see table 42 or table 43 in 5.4.2.2)							
	Initialization pattern descriptor (if any) (see table 45 in 5.4.2.3)							
	Defect list (if any)							

The parameter list header is defined in 5.4.2.2.



The initialization pattern descriptor, if any, is defined in 5.4.2.3.

The defect list, if any, contains address descriptors (see 6.2) each specifying a location on the medium to which the device server shall not assign LBAs. The device server shall maintain the current logical block to physical block alignment (see 4.6) for logical blocks not specified in the defect list.

Short block format address descriptors and long block format address descriptors should be in ascending order. Bytes from index format address descriptors and physical sector format address descriptors shall be in ascending order. If the address descriptors are not in the required order, then the device server may terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

#### 5.4.2.2 Parameter list header

The parameter list headers (see table 42 and table 43) provide several optional format control parameters. If the application client requests a function that is not implemented by the device server, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

If the **LOONGLIST** bit is set to zero in the **FORMAT UNIT CDB**, then the short parameter list header (see table 42) is used.

**Table 42 — Short parameter list header**

Byte	Bit	7	6	5	4	3	2	1	0
0	Reserved						PROTECTION FIELD USAGE		
	FOV	Format options bits				Obsolete	IMMED	Vendor specific	
1		DPRY	DCRT	STPF	IP				
2	(MSB)								
3		DEFECT LIST LENGTH (LSB)							

If the **LOONGLIST** bit is set to one in the **FORMAT UNIT CDB**, then the long parameter list header (see table 43) is used.

**Table 43 — Long parameter list header**

Byte	Bit	7	6	5	4	3	2	1	0
0		Reserved					PROTECTION FIELD USAGE		
		FOV	Format options bits				Obsolete	IMMED	Vendor specific
1			DPRY	DCRT	STPF	IP			
2		Reserved							
3		P_I_INFORMATION				PROTECTION INTERVAL EXPONENT			
4		(MSB)							
...		DEFECT LIST LENGTH							
7		(LSB)							

The combination (see table 44) of the PROTECTION FIELD USAGE field and the FMTPINFO field (see 5.4.1) specifies the requested protection type (see 4.22.2).

**Table 44 — FMTPINFO field and PROTECTION FIELD USAGE field (part 1 of 2)**

Fields indicated by the device server		Fields specified by the application client		Description
SPT <sup>a</sup>	PROTECT <sup>b</sup>	FMTPINFO	PROTECTION FIELD USAGE	
xxxb	0	00b	000b	The logical unit shall be formatted to type 0 protection <sup>c</sup> (see 4.22.2.2) resulting in the PROT_EN bit <sup>d</sup> being set to zero and the P_TYPE field <sup>d</sup> being set to 000b.
			>000b	Illegal <sup>e</sup>
		01b	xxxb	Illegal <sup>f</sup>
		1xb	xxxb	Illegal <sup>f</sup>
xxxb	1	00b	000b	The logical unit shall be formatted to type 0 protection <sup>c</sup> (see 4.22.2.2) resulting in the PROT_EN bit <sup>d</sup> being set to zero and the P_TYPE field <sup>d</sup> being set to 000b.
			>000b	Illegal <sup>e</sup>
xxxb	1	01b	xxxb	Illegal <sup>f</sup>
000b 001b 011b 111b	1	10b	000b	The logical unit shall be formatted to type 1 protection <sup>g</sup> (see 4.22.2.3) resulting in the PROT_EN bit <sup>d</sup> being set to one and the P_TYPE field <sup>d</sup> being set to 000b.
			>000b	Illegal <sup>e</sup>
010b 100b 101b	1	10b	xxxb	Illegal <sup>f</sup>
000b	1	11b	xxxb	Illegal <sup>f</sup>
<p><sup>a</sup> See the Extended INQUIRY Data VPD page (see <del>SPC-4</del>SPC-5) for the definition of the SPT field.</p> <p><sup>b</sup> See the standard INQUIRY data (see SPC-5) for the definition of the PROTECT bit.</p> <p><sup>c</sup> The device server shall format the medium to the logical block length specified in the mode parameter block descriptor of the mode parameter header (see SPC-5).</p> <p><sup>d</sup> See the READ CAPACITY (16) parameter data (see 5.18.2) for the definition of the PROT_EN bit and P_TYPE field.</p> <p><sup>e</sup> The device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.</p> <p><sup>f</sup> The device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.</p> <p><sup>g</sup> The device server shall format the medium to the logical block length specified in the mode parameter block descriptor of the mode parameter header plus eight bytes for each protection information interval (e.g., if the logical block length is 2 048 and the PROTECTION INTERVAL EXPONENT field set to two, then there are four 512 byte protection information intervals each followed by eight bytes of protection information resulting in a formatted logical block length of 2 080 bytes). Following a successful format operation, the PROT_EN <del>bit in the READ CAPACITY (16) parameter data</del> (see 5.18.2) indicates whether protection information (see 4.22) is enabled.</p>				

Table 44 — FMTPINFO field and PROTECTION FIELD USAGE field (part 2 of 2)

Fields indicated by the device server		Fields specified by the application client		Description
SPT <sup>a</sup>	PROTECT <sup>b</sup>	FMTPINFO	PROTECTION FIELD USAGE	
001b 010b 101b 111b	1	11b	000b	The logical unit shall be formatted to type 2 protection <sup>g</sup> (see 4.22.2.4) resulting in the PROT_EN bit <sup>d</sup> being set to one and the P_TYPE field <sup>d</sup> being set to 001b.
001b 010b	1	11b	>000b	Illegal <sup>e</sup>
011b 100b	1	11b	000b	Illegal <sup>e</sup>
011b 100b 101b 111b	1	11b	001b	The logical unit shall be formatted to type 3 protection. <sup>g</sup> (see 4.22.2.5) resulting in the PROT_EN bit <sup>d</sup> being set to one and the P_TYPE field <sup>d</sup> being set to 010b.
			>001b	Illegal <sup>e</sup>
110b	1	10b 11b	xxxb	Reserved
<sup>a</sup> See the Extended INQUIRY Data VPD page (see <b>SPC-4</b> SPC-5) for the definition of the SPT field. <sup>b</sup> See the standard INQUIRY data (see SPC-5) for the definition of the PROTECT bit. <sup>c</sup> The device server shall format the medium to the logical block length specified in the mode parameter block descriptor of the mode parameter header (see SPC-5). <sup>d</sup> See the READ CAPACITY (16) parameter data (see 5.18.2) for the definition of the PROT_EN bit and P_TYPE field. <sup>e</sup> The device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST. <sup>f</sup> The device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB. <sup>g</sup> The device server shall format the medium to the logical block length specified in the mode parameter block descriptor of the mode parameter header plus eight bytes for each protection information interval (e.g., if the logical block length is 2 048 and the PROTECTION INTERVAL EXPONENT field set to two, then there are four 512 byte protection information intervals each followed by eight bytes of protection information resulting in a formatted logical block length of 2 080 bytes). Following a successful format operation, the PROT_EN bit in the <b>READ CAPACITY (16) parameter data</b> (see 5.18.2) indicates whether protection information (see 4.22) is enabled.				

A format options valid (FOV) bit set to zero specifies that the device server shall use its default settings for the functionality represented by the DPRY bit, the DCRT bit, the STPF bit, and IP bit (i.e., the format options bits). If the FOV bit is set to zero, then the application client should set each of the format options bits to zero. If the FOV bit is set to zero, and any of the format options bits are not set to zero, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

A FOV bit set to one specifies that the device server shall process the format options bits as follows:

- a) a disable primary (DPRY) bit:
  - A) set to zero specifies that the device server shall not assign LBAs to parts of the medium identified as defective in the PLIST; or
  - B) set to one specifies that the device server shall not use the PLIST to identify defective areas of the medium, and the PLIST shall not be deleted;

- b) a disable certification (DCRT) bit:
    - A) set to zero specifies that the device server shall perform a vendor specific medium certification operation and add address descriptors for defects that it detects during the certification operation to the GLIST; or
    - B) set to one specifies that the device server shall not perform any vendor specific medium certification process or format verification operation;
  - c) the stop format (STPF) bit controls the behavior of the device server if the device server has been requested to use the PLIST (i.e., the DPRY bit is set to zero) or the GLIST (i.e., the Cmplst bit is set to zero) and one or more of the following occurs:
    - A) **list locate error**: the device server is not able to locate a specified defect list or determine whether a specified defect list exists; or
    - B) **list access error**: the device server encounters an error while accessing a specified defect list;
  - d) a STPF bit set to zero specifies that:
    - A) if a list locate error or a list access error occurs, then the device server shall continue to process the FORMAT UNIT command; and
    - B) after the format operation is complete, if:
      - a) a list locate error and a list access error both occurred, then the device server shall terminate the FORMAT UNIT command with CHECK CONDITION status at the completion of the command with the sense key set to RECOVERED ERROR and the additional sense code set to DEFECT LIST NOT FOUND or DEFECT LIST ERROR;
      - b) a list locate error occurred and a list access error did not occur, then the device server shall terminate the FORMAT UNIT command with CHECK CONDITION status with the sense key set to RECOVERED ERROR with the additional sense code set to DEFECT LIST NOT FOUND; and
      - c) a list access error occurred and a list locate error did not occur, then the device server shall terminate the FORMAT UNIT command with CHECK CONDITION status with the sense key set to RECOVERED ERROR with the additional sense code set to DEFECT LIST ERROR;
  - e) a STPF bit set to one specifies that:
    - A) if a list locate error occurs, then the device server shall terminate the FORMAT UNIT command with CHECK CONDITION status with the sense key set to MEDIUM ERROR and the additional sense code set to either DEFECT LIST NOT FOUND; or
    - B) if a list access error occurs, then the device server shall terminate the FORMAT UNIT command with CHECK CONDITION status with the sense key set to MEDIUM ERROR with the additional sense code set to DEFECT LIST ERROR;
- and
- f) an initialization pattern (IP) bit:
    - A) set to zero specifies that an initialization pattern descriptor (see 5.4.2.3) is not included and that the device server shall use its default initialization pattern; or
    - B) set to one specifies that:
      - a) an initialization pattern descriptor is included in the FORMAT UNIT parameter list following the parameter list header; and
      - b) if the device server does not support initialization pattern descriptors, then the device server shall terminate the FORMAT UNIT command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

An immediate (IMMED) bit set to zero specifies that the device server shall return status after the format operation has completed. An IMMED bit set to one specifies that the device server shall return status after the entire parameter list has been transferred.

The P\_I\_INFORMATION field, if any (i.e., if the long parameter list header is used), should be set to 0h. If the P\_I\_INFORMATION field is not set to zero, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

For a type 0 or a type 1 protection information request, if the PROTECTION INTERVAL EXPONENT field, if any, is not set to 0h, then the device server shall terminate the command with CHECK CONDITION status with the

sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

For a type 2 protection or a type 3 protection format request, the protection interval exponent determines the length of user data to be transferred before protection information is transferred (i.e., the protection information interval).

The protection information interval is calculated as follows:

$$\text{protection information interval} = \text{logical block length} \div 2^{(\text{protection interval exponent})}$$

where:

logical block length is the number of bytes of user data in a logical block (see 4.5)

protection interval exponent is zero if the short parameter list header (see table 42) is used or the contents of the PROTECTION INTERVAL EXPONENT field if the long parameter list header (see table 43) is used

If the protection information interval calculates to a value that is not an even number (e.g.,  $520 \div 2^3 = 65$ ) or not a whole number (e.g.,  $520 \div 2^4 = 32.5$  and  $520 \div 2^{10} = 0.508$ ), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

The DEFECT LIST LENGTH field specifies the total length in bytes of the defect list (i.e., the address descriptors) that follow and does not include the length of the initialization pattern descriptor, if any. The formats for the address descriptor(s) are shown in 6.2.

#### 5.4.2.3 Initialization pattern descriptor

The initialization pattern descriptor specifies that the device server initialize logical blocks to a specified pattern. The initialization pattern descriptor (see table 45) is transferred to the device server as part of the FORMAT UNIT parameter list.

**Table 45 — Initialization pattern descriptor**

Bit	7	6	5	4	3	2	1	0
Byte								
0	Obsolete		SI	Reserved				
1	INITIALIZATION PATTERN TYPE							
2	(MSB)							
3	INITIALIZATION PATTERN LENGTH (n - 3)							
4	(LSB)							
...	INITIALIZATION PATTERN							
n								

A security initialize (SI) bit set to one specifies that the device server shall attempt to write the initialization pattern to all areas of the medium including those that may have been reassigned (i.e., are in a defect list). An SI bit set to one specifies that the device server shall ignore:

- the FMTPINFO field;
- the FMTDATA bit;
- the CMPLIST bit;
- the DEFECT LIST FORMAT field;
- all the bits and fields in the parameter list header, except the IMMED bit; and
- any defect list data.

The device server shall write the initialization pattern using a security erasure write technique. The security erasure write technique requirement and procedure is outside the scope of this standard. The device server is not required to write the initialization pattern over the header and other parts of the medium not previously accessible to the application client. If the device server is unable to write over any part of the medium that is currently accessible to the application client or may be made accessible to the application client in the future (e.g., by clearing the defect list), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to MEDIUM ERROR and the additional sense code set to the appropriate value for the condition. The device server shall attempt to rewrite all remaining parts of the medium even if some parts are not able to be rewritten.

NOTE 6 - The intent of the security erasure write is to render any previous user data unrecoverable by any analog or digital technique.

NOTE 7 - Migration from the SI bit to the SANITIZE command (see 5.25) is recommended for all implementations.

An SI bit set to zero specifies that the device server shall initialize the application client accessible part of the medium. The device server is not required to initialize other areas of the medium. The device server shall format the medium as defined in the FORMAT UNIT command.

The INITIALIZATION PATTERN TYPE field (see table 46) specifies the type of pattern the device server shall use to initialize each logical block within the application client accessible part of the medium. All bytes within a logical block shall be written with the initialization pattern.

**Table 46 — INITIALIZATION PATTERN TYPE field**

Code	Description
00h	Use a default initialization pattern <sup>a</sup>
01h	Repeat the pattern specified in the INITIALIZATION PATTERN field as required to fill the logical block <sup>b</sup>
02h to 7Fh	Reserved
80h to FFh	Vendor specific
<sup>a</sup> If the INITIALIZATION PATTERN LENGTH field is not set to zero, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST. <sup>b</sup> If the INITIALIZATION PATTERN LENGTH field is set to zero, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.	

The INITIALIZATION PATTERN LENGTH field specifies the number of bytes contained in the INITIALIZATION PATTERN field. If the initialization pattern length exceeds the current logical block length, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

The INITIALIZATION PATTERN field specifies the initialization pattern.

## 5.5 GET LBA STATUS command

### 5.5.1 GET LBA STATUS command overview

The GET LBA STATUS command (see table 47) requests that the device server transfer parameter data describing the logical block provisioning status (see 4.7) for the specified LBA and zero or more subsequent LBAs to the Data-In Buffer.

The device server may or may not process this command as an uninterrupted sequence of actions (e.g., if concurrent operations are occurring that affect the logical block provisioning status, then the returned parameter data may be inconsistent or out of date).

This command uses the SERVICE ACTION IN (16) CDB format (see A.2).

**Table 47 — GET LBA STATUS command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (9Eh)							
1		Reserved			SERVICE ACTION (12h)				
2	(MSB)	STARTING LOGICAL BLOCK ADDRESS							
...									
9									
10	(MSB)	ALLOCATION LENGTH							
...									
13									
14		Reserved							
15		CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 47 for the GET LBA STATUS command.

The SERVICE ACTION field is defined in SPC-5 and shall be set to the value shown in table 47 for the GET LBA STATUS command.

The STARTING LOGICAL BLOCK ADDRESS field specifies the LBA of the first logical block addressed by this command. If the specified starting LBA exceeds the capacity of the medium (see 4.5), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to LOGICAL BLOCK ADDRESS OUT OF RANGE.

The ALLOCATION LENGTH field is defined in SPC-5. In response to a GET LBA STATUS command, the device server may send less data to the Data-In Buffer than is specified by the allocation length. If, in response to a single GET LBA STATUS command, the device server does not send sufficient data to the Data-In Buffer to satisfy the requirement of the application client, then, to retrieve additional information, the application client may send additional GET LBA STATUS commands with different starting LBA values.

The CONTROL byte is defined in SAM-5.

## 5.5.2 GET LBA STATUS parameter data

### 5.5.2.1 GET LBA STATUS parameter data overview

The GET LBA STATUS parameter data (see table 48) contains an eight-byte header followed by one or more LBA status descriptors.

**Table 48 — GET LBA STATUS parameter data**

Byte	Bit	7	6	5	4	3	2	1	0
0	(MSB)	PARAMETER DATA LENGTH (n - 3)							
...									
3									
4		Reserved							
...									
7									
LBA status descriptors									
8		LBA status descriptor [first] (see 5.5.2.2)							
...									
23									
⋮									
n - 15		LBA status descriptor [last] (see 5.5.2.2) (if any)							
...									
n									

The PARAMETER DATA LENGTH field indicates the number of bytes of parameter data that follow. The value in the PARAMETER DATA LENGTH field shall be:

- a) at least 20 (i.e., the available parameter data shall contain at least one LBA status descriptor); and
- b) four added to a multiple of 16 (i.e., the available parameter data shall end on a boundary between LBA Status descriptors).

Due to processing considerations outside the scope of this standard, two GET LBA STATUS commands with identical values in all CDB fields may result in two different values in the PARAMETER DATA LENGTH field.

The relationship between the PARAMETER DATA LENGTH field and the ALLOCATION LENGTH field in the CDB is defined in SPC-5.



### 5.5.2.2 LBA status descriptor

The LBA status descriptor (see table 49) contains LBA status information for one or more LBAs.

**Table 49 — LBA status descriptor format**

Bit	7	6	5	4	3	2	1	0
Byte								
0	(MSB) LBA STATUS LOGICAL BLOCK ADDRESS (LSB)							
...								
7								
8	(MSB) NUMBER OF LOGICAL BLOCKS (LSB)							
...								
11								
12	Reserved			PROVISIONING STATUS				
13	Reserved							
...								
15								

The LBA STATUS LOGICAL BLOCK ADDRESS field contains the first LBA of the LBA extent for which this descriptor reports LBA status.

The NUMBER OF LOGICAL BLOCKS field contains the number of logical blocks in that LBA extent. The device server should return the largest possible value in the NUMBER OF LOGICAL BLOCKS field.

The PROVISIONING STATUS field is defined in table 50.

**Table 50 — PROVISIONING STATUS field**

Code	Description
0h	Each LBA in the LBA extent is mapped (see 4.7.4.5) or has an unknown state.
1h	Each LBA in the LBA extent is deallocated (see 4.7.4.6).
2h	Each LBA in the LBA extent is anchored (see 4.7.4.7).
All others	Reserved

If the logical unit is fully provisioned (see 4.7.2), then the PROVISIONING STATUS field for all LBAs shall be set to 0h (i.e., mapped or unknown).

### 5.5.2.3 LBA status descriptor relationships

The LBA STATUS LOGICAL BLOCK ADDRESS field in the first LBA status descriptor returned in the GET LBA STATUS parameter data shall contain the value specified in the STARTING LOGICAL BLOCK ADDRESS field of the CDB. For subsequent LBA status descriptors, the contents of the LBA STATUS LOGICAL BLOCK ADDRESS field shall contain the sum of the values in:

- the LBA STATUS LOGICAL BLOCK ADDRESS field in the previous LBA status descriptor; and
- the NUMBER OF LOGICAL BLOCKS field in the previous LBA status descriptor.

Adjacent LBA status descriptors may have the same values for the PROVISIONING STATUS field.

## 5.6 GET STREAM STATUS command

### 5.6.1 GET STREAM STATUS command overview

The GET STREAM STATUS command (see table 51) requests that the device server transfer parameter data describing the status of streams (see 4.34) for the logical unit to the Data-In Buffer.

The device server may or may not process this command as an uninterrupted sequence of actions (e.g., if concurrent operations are occurring that affect the status of streams, then the returned parameter data may be inconsistent or out of date).

This command uses the SERVICE ACTION IN (16) CDB format (see A.2).

**Table 51 — GET STREAM STATUS command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (9Eh)							
1		Reserved			SERVICE ACTION (16h)				
2		Reserved							
3									
4		STARTING STREAM IDENTIFIER							
5									
6		Reserved							
...									
9									
10	(MSB)	ALLOCATION LENGTH							
...									
13									
14		Reserved							
15		CONTROL							

The OPERATION CODE field and the SERVICE ACTION field are defined in SPC-5 and shall be set to the values shown in table 51 for the GET STREAM STATUS command.

The STARTING STREAM IDENTIFIER field specifies the stream identifier of the first stream addressed by this command (see 5.6.2.3). If the specified starting stream identifier exceeds the value indicated by the MAXIMUM NUMBER OF STREAMS field of the Block Limits VPD page (see 6.6.4), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to ILLEGAL FIELD IN CDB.

The ALLOCATION LENGTH field is defined in SPC-5. If, in response to a single GET STREAM STATUS command, the device server does not send sufficient data to the Data-In Buffer to satisfy the requirement of the application client, then the application client may send additional GET STREAM STATUS commands with different starting stream identifier values to retrieve additional information.

The CONTROL byte is defined in SAM-5.

## 5.6.2 GET STREAM STATUS parameter data

### 5.6.2.1 GET STREAM STATUS parameter data overview

The GET STREAM STATUS parameter data (see table 52) contains an eight-byte header followed by zero or more stream status descriptors.

**Table 52 — GET STREAM STATUS parameter data**

Byte	Bit	7	6	5	4	3	2	1	0
0		PARAMETER DATA LENGTH (n-7)							
...									
3									
4		Reserved							
5									
6	(MSB)	NUMBER OF OPEN STREAMS							
7									
stream status descriptors									
8		stream status descriptor [first] (see 5.6.2.2)							
...									
15									
⋮									
n-8		stream status descriptor [last] (see 5.6.2.2)							
...									
n									

The PARAMETER DATA LENGTH field shall contain the length in bytes of the stream list. The stream list length is the number of open streams in the logical unit multiplied by eight. The contents of the STREAM LIST LENGTH field are not altered based on the allocation length.

Due to processing considerations outside the scope of this standard, two GET STREAM STATUS commands with identical values in all CDB fields may result in two different values in the PARAMETER DATA LENGTH field.

The relationship between the PARAMETER DATA LENGTH field and the ALLOCATION LENGTH field in the CDB is defined in SPC-5.

The NUMBER OF OPEN STREAMS field indicates the number of streams that are currently open in the logical unit.

### 5.6.2.2 Stream status descriptor

The stream status descriptor (see table 53) contains stream status information for one open stream.

**Table 53 — Stream status descriptor format**

Byte	Bit	7	6	5	4	3	2	1	0
0		Reserved							
1									
2	(MSB)	STREAM IDENTIFIER							
3									
4		Reserved							
...									
7									

The STREAM IDENTIFIER field contains the stream identifier of an open stream.

### 5.6.2.3 Stream status descriptor relationships

The STREAM IDENTIFIER field in the first stream status descriptor returned in the GET STREAM STATUS parameter data shall contain:

- a) the value specified in the STARTING STREAM IDENTIFIER field of the CDB if that stream is open; or
- b) the value of the next greater stream identifier of an open stream.

If the value specified in the STARTING STREAM IDENTIFIER field of the CDB is greater than the highest stream identifier of an open stream, then the device server shall not return any stream status descriptors.

For subsequent stream status descriptors, the contents of the STREAM IDENTIFIER field shall contain the value of the next greater stream identifier of an open stream.

## 5.7 ORWRITE (16) command

The ORWRITE (16) command (see table 54) requests that the device server perform an ORWRITE command ~~set~~ operation (see 4.29.4).

**Table 54 — ORWRITE (16) command**

Bit	7	6	5	4	3	2	1	0
Byte								
0	OPERATION CODE (8Bh)							
1	ORPROTECT			DPO	FUA	Reserved		
2	(MSB)							
...	LOGICAL BLOCK ADDRESS							
9	(LSB)							
10	(MSB)							
...	TRANSFER LENGTH							
13	(LSB)							
14	Reserved		GROUP NUMBER					
15	CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 54 for the ORWRITE (16) command.

See the READ (10) command (see 5.13) for the definition of the FUA bit specifying behavior for read operations. See the WRITE (10) command (see 5.35) for the definition of the FUA bit specifying behavior for write operations. See the READ (10) command (see 5.13) for the definition of the DPO bit. See the PRE-FETCH (10) command (see 5.10) for the definition of the LOGICAL BLOCK ADDRESS field. See the PRE-FETCH (10) command (see 5.10) and 4.23 for the definition of the GROUP NUMBER field.

The TRANSFER LENGTH field specifies the number of contiguous logical blocks of data that are read, transferred from the Data-Out Buffer, and ORed into a bitmap buffer, starting with the logical block referenced by the LBA specified by the LOGICAL BLOCK ADDRESS field. If the specified LBA and the specified transfer length exceed the capacity of the medium (see 4.5), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to LOGICAL BLOCK ADDRESS OUT OF RANGE. The TRANSFER LENGTH field is constrained by the MAXIMUM TRANSFER LENGTH field (see 6.6.4).

The CONTROL byte is defined in SAM-5.

The device server shall:

- check protection information from the read operations based on the ORPROTECT field as described in table 55; and
- check protection information transferred from the Data-Out Buffer based on the ORPROTECT field as described in table 56.

The order of the user data and protection information checks and comparisons is vendor specific.

The device server shall check the protection information from the read operations based on the ORPROTECT field as described in table 55. All footnotes for table 55 are at the end of the table.

**Table 55 — ORPROTECT field - checking protection information from the read operations (part 1 of 3)**

Code	Logical unit formatted with protection information	Field in protection information <sup>g</sup>	Extended INQUIRY Data VPD page bit value <sup>f</sup>	If check fails <sup>d e</sup> , additional sense code
000b	Yes <sup>ij</sup>	LOGICAL BLOCK GUARD	GRD_CHK = 1	LOGICAL BLOCK GUARD CHECK FAILED
			GRD_CHK = 0	No check performed
		LOGICAL BLOCK APPLICATION TAG	APP_CHK = 1 <sup>c</sup>	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
			APP_CHK = 0	No check performed
		LOGICAL BLOCK REFERENCE TAG	REF_CHK = 1 <sup>h</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
			REF_CHK = 0	No check performed
	No	No protection information on the medium to check.		
001b 101b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	GRD_CHK = 1	LOGICAL BLOCK GUARD CHECK FAILED
			GRD_CHK = 0	No check performed
		LOGICAL BLOCK APPLICATION TAG	APP_CHK = 1 <sup>c</sup>	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
			APP_CHK = 0	No check performed
		LOGICAL BLOCK REFERENCE TAG	REF_CHK = 1 <sup>h</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
			REF_CHK = 0	No check performed
	No	Error condition <sup>a</sup>		
010b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	No check performed	
		LOGICAL BLOCK APPLICATION TAG	APP_CHK = 1 <sup>c</sup>	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
			APP_CHK = 0	No check performed
		LOGICAL BLOCK REFERENCE TAG	REF_CHK = 1 <sup>h</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
			REF_CHK = 0	No check performed
	No	Error condition <sup>a</sup>		

**Table 55 — ORPROTECT field - checking protection information from the read operations (part 2 of 3)**

Code	Logical unit formatted with protection information	Field in protection information <sup>g</sup>	Extended INQUIRY Data VPD page bit value <sup>f</sup>	If check fails <sup>d e</sup> , additional sense code
011b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	No check performed	
		LOGICAL BLOCK APPLICATION TAG	No check performed	
		LOGICAL BLOCK REFERENCE TAG	No check performed	
	No	Error condition <sup>a</sup>		
100b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	GRD_CHK = 1	LOGICAL BLOCK GUARD CHECK FAILED
			GRD_CHK = 0	No check performed
		LOGICAL BLOCK APPLICATION TAG	No check performed	
		LOGICAL BLOCK REFERENCE TAG	No check performed	
	No	Error condition <sup>a</sup>		

**Table 55 — ORPROTECT field - checking protection information from the read operations (part 3 of 3)**

Code	Logical unit formatted with protection information	Field in protection information <sup>g</sup>	Extended INQUIRY Data VPD page bit value <sup>f</sup>	If check fails <sup>d e</sup> , additional sense code
110b to 111b	Reserved			
<p><sup>a</sup> If the logical unit supports protection information (see 4.22) and has not been formatted with protection information, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.</p> <p><sup>b</sup> If the logical unit does not support protection information, then the device server should terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.</p> <p><sup>c</sup> If the device server has knowledge of the contents of the logical block application tag field, then the device server shall check the logical block application tag. If the ato bit in the Control mode page (see SPC-5) is set to one, then this knowledge is acquired from:</p> <p>a) the Application Tag mode page (see 6.5.3), if the atmpe bit in the Control mode page (see SPC-5) is set to one; or.</p> <p>b) a method not defined by this standard, if the ATMPE bit is set to zero.</p> <p><sup>d</sup> If the device server terminates the command with CHECK CONDITION status, then the device server shall set the sense key to ABORTED COMMAND.</p> <p><sup>e</sup> If multiple errors occur while the device server is processing the command, then the selection by the device server of which error to report is not defined by this standard.</p> <p><sup>f</sup> See the Extended INQUIRY Data VPD page (see SPC-5) for the definitions of the GRD_CHK bit, the APP_CHK bit, and the REF_CHK bit.</p> <p><sup>g</sup> If the device server detects:</p> <p>a) a LOGICAL BLOCK APPLICATION TAG field set to FFFFh, and type 1 protection (see 4.22.2.3) is enabled; or</p> <p>b) a LOGICAL BLOCK APPLICATION TAG field set to FFFFh, the LOGICAL BLOCK REFERENCE TAG field set to FFFF_FFFFh, and type 3 protection (see 4.22.2.5) is enabled, then the device server shall not check any protection information in the associated protection information interval.</p> <p><sup>h</sup> If type 1 protection is enabled, then the device server shall check the logical block reference tag by comparing it to the lower four bytes of the LBA associated with the logical block. If type 3 protection is enabled, then the device server shall check each logical block reference tag only if the device server has knowledge of the contents of the LOGICAL BLOCK REFERENCE TAG field. The method for acquiring this knowledge is not defined by this standard.</p> <p><sup>i</sup> If the RWWP bit in the Control mode page (see SPC-5) is set to one, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.</p> <p><sup>j</sup> If the DPICZ bit in the Control mode page (see SPC-5) is set to one, and the RWWP bit in the Control mode page is set to zero, then protection information shall not be checked.</p>				



The device server shall check the protection information transferred from the Data-Out Buffer based on the ORPROTECT field as described in table 56. All footnotes for table 56 are at the end of the table.

**Table 56 — ORPROTECT field - checking protection information from the Data-Out Buffer (part 1 of 2)**

Code	Logical unit formatted with protection information	Field in protection information	Device server check	If check fails <sup>d h</sup> , additional sense code
000b	Yes <sup>e f g</sup>	No protection information in the Data-Out buffer to check		
	No	No protection information in the Data-Out buffer to check		
001b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	Shall	LOGICAL BLOCK GUARD CHECK FAILED
		LOGICAL BLOCK APPLICATION TAG	Dependent on RWWP <sup>c</sup>	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
		LOGICAL BLOCK REFERENCE TAG	Shall (except for type 3) <sup>i</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
	No	Error condition <sup>a</sup>		
010b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	Shall not	No check performed
		LOGICAL BLOCK APPLICATION TAG	Dependent on RWWP <sup>c</sup>	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
		LOGICAL BLOCK REFERENCE TAG	May <sup>i</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
	No	Error condition <sup>a</sup>		
011b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	Shall not	No check performed
		LOGICAL BLOCK APPLICATION TAG	Shall not	No check performed
		LOGICAL BLOCK REFERENCE TAG	Shall not	No check performed
	No	Error condition <sup>a</sup>		
100b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	Shall	LOGICAL BLOCK GUARD CHECK FAILED
		LOGICAL BLOCK APPLICATION TAG	Shall not	No check performed
		LOGICAL BLOCK REFERENCE TAG	Shall not	No check performed
	No	Error condition <sup>a</sup>		

**Table 56 — ORPROTECT field - checking protection information from the Data-Out Buffer (part 2 of 2)**

Code	Logical unit formatted with protection information	Field in protection information	Device server check	If check fails <sup>d h</sup> , additional sense code
101b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	Shall	LOGICAL BLOCK GUARD CHECK FAILED
		LOGICAL BLOCK APPLICATION TAG	Dependent on RWWP <sup>c</sup>	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
		LOGICAL BLOCK REFERENCE TAG	May <sup>i</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
	No	Error condition <sup>a</sup>		
110b to 111b	Reserved			

<sup>a</sup> If a logical unit supports protection information (see 4.22) and has not been formatted with protection information, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

<sup>b</sup> If the logical unit does not support protection information, then the device server should terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

<sup>c</sup> If the ATO bit is set to one in the Control mode page (see SPC-5), and the device server has knowledge of the contents of the LOGICAL BLOCK APPLICATION TAG field, then the device server:

a) may check each logical block application tag if the RWWP bit is set to zero in the Control mode page (see SPC-5); and

b) shall check each logical block application tag if the RWWP bit is set to one in the Control mode page. If the ATMPE bit in the Control mode page (see SPC-5) is set to one, then this knowledge is acquired from the Application Tag mode page. If the ATMPE bit is set to zero, then the method for acquiring this knowledge is not defined by this standard.

<sup>d</sup> If the device server terminates the command with CHECK CONDITION status, then the device server shall set the sense key to ABORTED COMMAND.

<sup>e</sup> The device server shall write a generated CRC (see 4.22.4.2) into each LOGICAL BLOCK GUARD field.

<sup>f</sup> If the RWWP bit in the Control mode page (see SPC-5) is set to one, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB. If the RWWP bit is set to zero, and:

a) type 1 protection is enabled, then the device server shall write the least significant four bytes of each LBA into the LOGICAL BLOCK REFERENCE TAG field of each of the written logical blocks; or

b) type 3 protection is enabled, then the device server shall write a value of FFFF\_FFFFh into the LOGICAL BLOCK REFERENCE TAG field of each of the written logical blocks.

<sup>g</sup> If the ATO bit is set to one in the Control mode page (see SPC-5), then the device server shall write FFFFh into each LOGICAL BLOCK APPLICATION TAG field. If the ATO bit is set to zero, then the device server may write any value into each LOGICAL BLOCK APPLICATION TAG field.

<sup>h</sup> If multiple errors occur while the device server is processing the command, then the selection by the device server of which error to report is not defined by this standard.

<sup>i</sup> If type 1 protection is enabled, then the device server shall check the logical block reference tag by comparing it to the lower four bytes of the LBA associated with the logical block. If type 3 protection is enabled, then the device server may check each logical block reference tag if the ATO bit is set to one in the Control mode page (see SPC-5), and the device server has knowledge of the contents of the LOGICAL BLOCK REFERENCE TAG. The method for acquiring this knowledge is not defined by this standard.

## 5.8 ORWRITE (32) command

The ORWRITE (32) command (see table 57) requests that the device server perform one of the following ORWRITE command (see 4.29) operations:

- a) a change generation and clear operation (see 4.29.3); or
- b) a set operation (see 4.29.4).

**Table 57 — ORWRITE (32) command**

Bit	7	6	5	4	3	2	1	0
Byte								
0	OPERATION CODE (7Fh)							
1	CONTROL							
2	Reserved					BMOP		
3	Reserved				PREVIOUS GENERATION PROCESSING			
4	Reserved							
5	Reserved							
6	Reserved		GROUP NUMBER					
7	ADDITIONAL CDB LENGTH (18h)							
8	(MSB)							
9	SERVICE ACTION (000Eh)							
10	ORPROTECT			DPO	FUA	Reserved		
11	Reserved							
12	(MSB)							
...	LOGICAL BLOCK ADDRESS							
19	(LSB)							
20	(MSB)							
...	EXPECTED ORWGENERATION							
23	(LSB)							
24	(MSB)							
...	NEW ORWGENERATION							
27	(LSB)							
28	(MSB)							
...	TRANSFER LENGTH							
31	(LSB)							

The OPERATION CODE field, the ADDITIONAL CDB LENGTH field, and the SERVICE ACTION field are defined in SPC-5 and shall be set to the values shown in table 57 for the ORWRITE (32) command.

The CONTROL byte is defined in SAM-5.

The bitmap operation (BMOP) field specifies the operation as described in table 58.

**Table 58 — BMOP field**

Code	Description
000b	The device server shall perform a set operation (see 4.29.4), and the contents of the PREVIOUS GENERATION PROCESSING field and NEW ORWGENERATION field shall be ignored.
001b	The device server shall perform a change generation and clear operation (see 4.29.3).
All others	Reserved

The PREVIOUS GENERATION PROCESSING field specifies the policy for performing future set operations that is to be established in the device server by a successful change generation and clear operation (see 4.29.2.2).

See the ORWRITE (16) command (see 5.7) for the definitions of the FUA bit, the DPO bit, the ORPROTECT field, the LOGICAL BLOCK ADDRESS field, the TRANSFER LENGTH field, and the GROUP NUMBER field.

The EXPECTED ORWGENERATION field contains a code that is compared with generation codes established and maintained by the device server.

The NEW ORWGENERATION field specifies the current ORWgeneration code that is to be established in the device server by a successful change generation and clear operation (see 4.29.3).

The device server shall:

- a) check protection information from the read operations based on the ORPROTECT field as described in table 55; and
- b) check protection information transferred from the Data-Out Buffer based on the ORPROTECT field as described in table 56.

The order of the user data and protection information checks and comparisons is vendor specific.

## 5.9 POPULATE TOKEN command

### 5.9.1 POPULATE TOKEN command overview

The POPULATE TOKEN command (see table 59) requests that the copy manager (see SPC-5) create a point in time ROD token that represents the specified logical blocks (see 4.30).

Each logical block represented by the point in time ROD token includes logical block data.

**Table 59 — POPULATE TOKEN command**

Bit	7	6	5	4	3	2	1	0
Byte								
0	OPERATION CODE (83h)							
1	Reserved			SERVICE ACTION (10h)				
2	Reserved							
...								
5								
6	LIST IDENTIFIER							
...								
9								
10	PARAMETER LIST LENGTH							
...								
13								
14	Reserved		GROUP NUMBER					
15	CONTROL							

The OPERATION CODE field and the SERVICE ACTION field are defined in SPC-5 and shall be set to the values shown in table 59 for the POPULATE TOKEN command.

The LIST IDENTIFIER field is defined in SPC-5. The list identifier shall be processed as if the LIST ID USAGE field in the parameter data for an EXTENDED COPY(LID4) command (see SPC-5) is set to 00b.

The PARAMETER LIST LENGTH field specifies the length in bytes of the parameter list that is available to be transferred from the Data-Out Buffer. If the parameter list length is greater than zero and less than 00000010h (i.e., 16), then the copy manager shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to PARAMETER LIST LENGTH ERROR. A PARAMETER LIST LENGTH field set to zero specifies that no data shall be transferred. This shall not be considered an error.

See the PRE-FETCH (10) command (see 5.10) and 4.23 for the definition of the GROUP NUMBER field.

The CONTROL byte is defined in SAM-5.

### 5.9.2 POPULATE TOKEN parameter list

The parameter list for the POPULATE TOKEN command is shown in table 60.

**Table 60 — POPULATE TOKEN parameter list**

Byte	Bit	7	6	5	4	3	2	1	0
0	(MSB)	POPULATE TOKEN DATA LENGTH (n - 1)							
1									
2		Reserved						RTV	IMMED
3		Reserved							
4	(MSB)	INACTIVITY TIMEOUT							
...									
7									
8	(MSB)	ROD TYPE							
...									
11									
12		Reserved							
13									
14	(MSB)	BLOCK DEVICE RANGE DESCRIPTOR LENGTH (n - 15)							
15									
Block device range descriptor list									
16		Block device range descriptor [first] (see 5.9.3)							
...									
31									
⋮									
n - 15		Block device range descriptor [last] (see 5.9.3) (if any)							
...									
n									

The POPULATE TOKEN DATA LENGTH field specifies the length in bytes of the data that is available to be transferred from the Data-Out Buffer. The populate token data length does not include the number of bytes in the POPULATE TOKEN DATA LENGTH field. If the POPULATE TOKEN DATA LENGTH field is less than 001Eh (i.e., 30), then the copy manager shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

A ROD type valid (RTV) bit set to zero specifies that the copy manager may create a ROD token with any point in time copy ROD type and shall ignore the contents of the ROD TYPE field. An RTV bit set to one specifies that the copy manager shall use the contents of the ROD TYPE field to create the point in time copy ROD.

The immediate (IMMED) bit specifies when the copy manager shall return status for the POPULATE TOKEN command. If the IMMED bit is set to zero, then the copy manager shall process the POPULATE TOKEN command until all specified operations are complete or an error is detected. If the IMMED bit is set to one, then the copy manager:

- 1) shall validate the CDB (i.e., detect and report all errors in the CDB);
- 2) shall transfer all the parameter data to the copy manager;

- 3) may validate the parameter data;
- 4) shall complete the POPULATE TOKEN command with GOOD status; and
- 5) shall complete performing of all specified operations as a background operation (see SPC-5).

If the INACTIVITY TIMEOUT field is not set to zero, then the INACTIVITY TIMEOUT field specifies the number of seconds to use for the ROD token inactivity timeout (see SPC-5). If the INACTIVITY TIMEOUT field is set to a value larger than the value in the MAXIMUM INACTIVITY TIMEOUT field (see 6.6.8.3), then the copy manager shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

If the INACTIVITY TIMEOUT field is set to zero, then the DEFAULT INACTIVITY TIMEOUT field (see 6.6.8.3) specifies the number of seconds to use for the ROD token inactivity timeout (see SPC-5).

If the RTV bit is set to one, then the ROD TYPE field specifies the ROD type (see SPC-5) for creating the point in time copy ROD token. The copy manager shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST, if:

- a) the copy manager does not support the specified ROD type for use with the POPULATE TOKEN command; or
- b) the ROD TYPE field specifies a ROD type (see SPC-5) that is not a point in time copy ROD.

The BLOCK DEVICE RANGE DESCRIPTOR LENGTH field specifies the length in bytes of the block device range descriptor list. The block device range descriptor list length should be a multiple of 16. If the block device range descriptor list length is not a multiple of 16, then the last block device range descriptor is incomplete and shall be ignored. If the block device range descriptor list length is less than 0010h (i.e., 16), then the copy manager shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

If any block device range descriptors in the block device range descriptor list are truncated due to the parameter list length in the CDB, then those block device range descriptors shall be ignored.

If the number of complete block device range descriptors is larger than the maximum range descriptors value in the Block Device ROD Token Limits descriptor (see 6.6.8.3), then the copy manager shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to TOO MANY SEGMENT DESCRIPTORS.

If the same LBA is included in more than one block device range descriptor, then the copy manager shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

If the number of bytes of user data represented by the sum of the contents of the NUMBER OF LOGICAL BLOCKS fields in all of the complete block device range descriptors is larger than:

- a) the MAXIMUM BYTES IN BLOCK ROD field in the block ROD device type specific features descriptor in the ROD token features third-party copy descriptor in the Third-party Copy VPD page (see SPC-5) and that field is set to a nonzero value; or
- b) the MAXIMUM TOKEN TRANSFER SIZE field (see 6.6.8.3) and that field is set to a nonzero value,

then the copy manager shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

### 5.9.3 Block device range descriptor

The block device range descriptor is defined in table 61.

**Table 61 — Block device range descriptor**

Bit	7	6	5	4	3	2	1	0
Byte								
0	(MSB)							
...	LOGICAL BLOCK ADDRESS							
7	(LSB)							
8	(MSB)							
...	NUMBER OF LOGICAL BLOCKS							
11	(LSB)							
12	Reserved							
...								
15								

The LOGICAL BLOCK ADDRESS field specifies the first LBA on which the copy manager shall operate for this block device range descriptor.

The NUMBER OF LOGICAL BLOCKS field specifies the number of logical blocks on which the copy manager shall operate for this block device range descriptor beginning with the LBA specified by the LOGICAL BLOCK ADDRESS field.

Processing of block device range descriptors with a number of logical blocks that is not a multiple of the OPTIMAL BLOCK ROD LENGTH GRANULARITY field in the block ROD device type specific features descriptor in the ROD token features third-party copy descriptor in the Third-party Copy VPD page (see SPC-5) may incur delays in processing. If the OPTIMAL BLOCK ROD LENGTH GRANULARITY field in the block ROD device type specific features descriptor in the ROD token features third-party copy descriptor in the Third-party Copy VPD page is not reported, then the optimal transfer length granularity in the Block Limits VPD page (see 6.6.4) may indicate the granularity.

For a POPULATE TOKEN command, processing of block device range descriptors where the number of bytes of user data contained in the number of logical blocks exceeds the OPTIMAL BYTES TO TOKEN PER SEGMENT field in the block ROD device type specific features descriptor in the ROD token features third-party copy descriptor in the Third-party Copy VPD page may incur delays in processing.

For a WRITE USING TOKEN command (see 5.52), processing of block device range descriptors where the number of bytes of user data contained in the number of logical blocks exceeds the OPTIMAL BYTES FROM TOKEN PER SEGMENT field in the block ROD device type specific features descriptor in the ROD token features third-party copy descriptor in the Third-party Copy VPD page may incur delays in processing.

If the number of bytes of user data contained in the number of logical blocks is greater than:

- 1) the value in the MAXIMUM BYTES IN BLOCK ROD field in the block ROD device type specific features descriptor in the ROD token features third-party copy descriptor in the Third-party Copy VPD page, and the MAXIMUM BYTES IN BLOCK ROD field is set to a nonzero value; or
- 2) the value in the MAXIMUM TRANSFER LENGTH field (see 6.6.4), the MAXIMUM TRANSFER LENGTH field is set to a nonzero value, and the MAXIMUM BYTES IN BLOCK ROD field in the block ROD device type specific features descriptor is not reported,

then the copy manager shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

If the NUMBER OF LOGICAL BLOCKS field is set to zero, then the copy manager shall perform no operation for this block device range descriptor. This condition shall not be considered an error.



If the specified LBA and the specified number of logical blocks exceed the capacity of the medium (see 4.5), then the copy manager shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to LOGICAL BLOCK ADDRESS OUT OF RANGE.

## 5.10 PRE-FETCH (10) command

The PRE-FETCH (10) command (see table 62) requests that the device server:

- a) for any mapped LBAs specified by the command that are not already contained in cache, perform read medium operations and write cache operations (see 4.15); and
- b) for any unmapped LBAs specified by the command, update the volatile cache and/or non-volatile cache to prevent retrieval of stale data.

No data shall be transferred to the Data-In Buffer.

NOTE 8 - Migration from the PRE-FETCH (10) command to the PRE-FETCH (16) command is recommended for all implementations.

**Table 62 — PRE-FETCH (10) command**

Byte	Bit	7	6	5	4	3	2	1	0
0	OPERATION CODE (34h)								
1	Reserved							IMMED	Obsolete
2	(MSB)								
...	LOGICAL BLOCK ADDRESS								
5	(LSB)								
6	Reserved			GROUP NUMBER					
7	(MSB)								
8	PREFETCH LENGTH							(LSB)	
9	CONTROL								

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 62 for the PRE-FETCH (10) command.

An immediate (IMMED) bit set to zero specifies that status shall be returned after the operation is complete. An IMMED bit set to one specifies that the device server shall:

- a) validate the CDB;
- b) if the cache has:
  - A) sufficient capacity to accept all of the specified logical blocks, then complete the command with CONDITION MET status; or
  - B) insufficient capacity to accept all of the specified logical blocks, then complete the command with GOOD status;
 and
- c) if one or more of the specified logical blocks are not successfully transferred to the cache for reasons other than lack of cache capacity, then report a deferred error (see SPC-5).

The LOGICAL BLOCK ADDRESS field specifies the LBA of the first logical block (see 4.5) accessed by this command.

A GROUP NUMBER field set to a non-zero value specifies the group into which attributes associated with the command should be collected (see 4.23). A GROUP NUMBER field set to zero specifies that any attributes associated with the command shall not be collected into any group.

The PREFETCH LENGTH field specifies the number of contiguous logical blocks that shall be pre-fetched (i.e., transferred to the cache from the medium), starting with the LBA specified by the LOGICAL BLOCK ADDRESS field. A PREFETCH LENGTH field set to zero specifies that all logical blocks starting with the LBA specified in the LOGICAL BLOCK ADDRESS field to the last logical block on the medium shall be pre-fetched. Any other value specifies the number of logical blocks that shall be pre-fetched. If the specified LBA and the specified prefetch length exceed the capacity of the medium (see 4.5), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to LOGICAL BLOCK ADDRESS OUT OF RANGE.

The CONTROL byte is defined in SAM-5.

If the IMMED bit is set to zero, and the specified logical blocks were transferred to the cache without error, then the device server shall complete the command with CONDITION MET status.

If the IMMED bit is set to zero and the cache does not have sufficient capacity to accept all of the specified logical blocks, then the device server shall transfer to the cache as many of the specified logical blocks that fit. If these logical blocks are transferred without error, then the device server shall complete the command with GOOD status.

## 5.11 PRE-FETCH (16) command

The PRE-FETCH (16) command (see table 63) requests that the device server perform the actions defined for the PRE-FETCH (10) command (see 5.10).

No data shall be transferred to the Data-In Buffer.

**Table 63 — PRE-FETCH (16) command**

Bit	7	6	5	4	3	2	1	0
Byte								
0	OPERATION CODE (90h)							
1	Reserved						IMMED	Reserved
2	(MSB)							
...	LOGICAL BLOCK ADDRESS							
9	(LSB)							
10	(MSB)							
...	PREFETCH LENGTH							
13	(LSB)							
14	Reserved		GROUP NUMBER					
15	CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 63 for the PRE-FETCH (16) command.

See the PRE-FETCH (10) command (see 5.10) for the definitions of the other fields in this command.

## 5.12 PREVENT ALLOW MEDIUM REMOVAL command

The PREVENT ALLOW MEDIUM REMOVAL command (see table 64) requests that the logical unit enable or disable the removal of the medium. If medium removal is prevented on any I\_T nexus that has access to the logical unit, then the logical unit shall not allow medium removal.

**Table 64 — PREVENT ALLOW MEDIUM REMOVAL command**

Byte	Bit	7	6	5	4	3	2	1	0
0	OPERATION CODE (1Eh)								
1	Reserved								
2	Reserved								
3	Reserved								
4	Reserved							PREVENT	
5	CONTROL								

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 64 for the PREVENT ALLOW MEDIUM REMOVAL command.

Table 65 defines the PREVENT field.

**Table 65 — PREVENT field**

Value	Description
00b	Medium removal is allowed.
01b	Medium removal shall be prevented.
10b to 11b	Obsolete

The CONTROL byte is defined in SAM-5.

The prevention of medium removal shall begin when any application client issues a PREVENT ALLOW MEDIUM REMOVAL command with the PREVENT field set to 01b (i.e., medium removal prevented). The prevention of medium removal for the logical unit shall no longer be prevented after:

- a) one of the following occurs for each I\_T nexus through which medium removal had been prevented:
  - A) receipt of a PREVENT ALLOW MEDIUM REMOVAL command with the PREVENT field set to 00b; or
  - B) an I\_T nexus loss;
- b) a power on;
- c) a hard reset; or
- d) a logical unit reset.

If possible, the device server shall perform a synchronize cache operation before ending the prevention of medium removal.

If a persistent reservation or registration is being preempted by a PERSISTENT RESERVE OUT command with PREEMPT AND ABORT service action (see SPC-5) or PERSISTENT RESERVE OUT command with CLEAR service action (see SPC-5), then the equivalent of a PREVENT ALLOW MEDIUM REMOVAL command with the PREVENT field set to 00b shall be processed for each the I\_T nexuses associated with the persistent reservation or registrations being preempted allowing an application client to override the

prevention of medium removal function for a SCSI initiator port (e.g., an initiator port is not operating correctly).

While a prevention of medium removal condition is in effect, the logical unit shall inhibit mechanisms that allow removal of the medium by an operator.

### 5.13 READ (10) command

The READ (10) command (see table 66) requests that the device server:

- a) perform read operations from the specified LBAs: and
- b) transfer the requested logical block data to the Data-In Buffer.

The logical block data transferred to the Data-In Buffer shall include protection information based on the value in the RDPROTECT field (see table 67) and the medium format.

NOTE 9 - Migration from the READ (10) command to the READ (16) command is recommended for all implementations.

**Table 66 — READ (10) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (28h)							
1		RDPROTECT			DPO	FUA	RARC	Obsolete	Obsolete
2	(MSB)	LOGICAL BLOCK ADDRESS							
...									
5									
6		Reserved		GROUP NUMBER					
7	(MSB)	TRANSFER LENGTH							
8									
9		CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 66 for the READ (10) command.

The device server shall check the protection information from the read operations before returning status for the command based on the RDPROTECT field as described in table 67. All footnotes for table 67 are at the end of the table.

**Table 67 — RDPROTECT field (part 1 of 3)**

Code	Logical unit formatted with protection information	Shall device server transmit protection information?	Field in protection information <sup>h</sup>	Extended INQUIRY Data VPD page bit value <sup>g</sup>	If check fails <sup>d f</sup> , additional sense code
000b	Yes <sup>j</sup>	No	LOGICAL BLOCK GUARD	GRD_CHK = 1	LOGICAL BLOCK GUARD CHECK FAILED
				GRD_CHK = 0	No check performed
			LOGICAL BLOCK APPLICATION TAG	APP_CHK = 1 <sup>c</sup>	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
				APP_CHK = 0	No check performed
			LOGICAL BLOCK REFERENCE TAG	REF_CHK = 1 <sup>i</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
				REF_CHK = 0	No check performed
	No	No protection information available to check			
001b 101b <sup>b</sup>	Yes	Yes <sup>e</sup>	LOGICAL BLOCK GUARD	GRD_CHK = 1	LOGICAL BLOCK GUARD CHECK FAILED
				GRD_CHK = 0	No check performed
			LOGICAL BLOCK APPLICATION TAG	APP_CHK = 1 <sup>c</sup>	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
				APP_CHK = 0	No check performed
			LOGICAL BLOCK REFERENCE TAG	REF_CHK = 1 <sup>i</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
				REF_CHK = 0	No check performed
	No <sup>a</sup>	No protection information available to transmit to the Data-In Buffer or for checking			

Table 67 — RDPROTECT field (part 2 of 3)

Code	Logical unit formatted with protection information	Shall device server transmit protection information?	Field in protection information <sup>h</sup>	Extended INQUIRY Data VPD page bit value <sup>g</sup>	If check fails <sup>d f</sup> , additional sense code
010b <sup>b</sup>	Yes	Yes <sup>e</sup>	LOGICAL BLOCK GUARD	No check performed	
			LOGICAL BLOCK APPLICATION TAG	APP_CHK = 1 <sup>c</sup>	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
				APP_CHK = 0	No check performed
			LOGICAL BLOCK REFERENCE TAG	REF_CHK = 1 <sup>i</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
				REF_CHK = 0	No check performed
	No <sup>a</sup>	No protection information available to transmit to the Data-In Buffer or for checking			
011b <sup>b</sup>	Yes	Yes <sup>e</sup>	LOGICAL BLOCK GUARD	No check performed	
			LOGICAL BLOCK APPLICATION TAG	No check performed	
			LOGICAL BLOCK REFERENCE TAG	No check performed	
	No <sup>a</sup>	No protection information available to transmit to the Data-In Buffer or for checking			
100b <sup>b</sup>	Yes	Yes <sup>e</sup>	LOGICAL BLOCK GUARD	GRD_CHK = 1	LOGICAL BLOCK GUARD CHECK FAILED
				GRD_CHK = 0	No check performed
			LOGICAL BLOCK APPLICATION TAG	No check performed	
			LOGICAL BLOCK REFERENCE TAG	No check performed	
	No <sup>a</sup>	No protection information available to transmit to the Data-In Buffer or for checking			
110b to 111b	Reserved				

Table 67 — RDPROTECT field (part 3 of 3)

Code	Logical unit formatted with protection information	Shall device server transmit protection information?	Field in protection information <sup>h</sup>	Extended INQUIRY Data VPD page bit value <sup>g</sup>	If check fails <sup>d,f</sup> , additional sense code
<p><sup>a</sup> If the logical unit supports protection information (see 4.22) and has not been formatted with protection information, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.</p> <p><sup>b</sup> If the logical unit does not support protection information, then the device server should terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.</p> <p><sup>c</sup> If the device server has knowledge of the contents of the LOGICAL BLOCK APPLICATION TAG field, then the device server shall check each logical block application tag. If the ATO bit in the Control mode page (see SPC-5) is set to one, then this knowledge is acquired from:</p> <ul style="list-style-type: none"> <li>a) the EXPECTED LOGICAL BLOCK APPLICATION TAG field and the LOGICAL BLOCK APPLICATION TAG MASK field in the CDB, if a READ (32) command (see 5.16) is received by the device server;</li> <li>b) the Application Tag mode page (see 6.5.3), if a command other than READ (32) is received by the device server, and the ATMPE bit in the Control mode page (see SPC-5) is set to one; or</li> <li>c) a method not defined by this standard, if a command other than READ (32) is received by the device server, and the ATMPE bit is set to zero.</li> </ul> <p><sup>d</sup> If the device server terminates the command with CHECK CONDITION status, then the device server shall set the sense key to ABORTED COMMAND.</p> <p><sup>e</sup> The device server shall transmit protection information to the Data-In Buffer.</p> <p><sup>f</sup> If multiple errors occur while the device server is processing the command, then the selection by the device server of which error to report is not defined by this standard.</p> <p><sup>g</sup> See the Extended INQUIRY Data VPD page (see SPC-5) for the definitions of the GRD_CHK bit, the APP_CHK bit, and the REF_CHK bit.</p> <p><sup>h</sup> If the device server detects:</p> <ul style="list-style-type: none"> <li>a) a LOGICAL BLOCK APPLICATION TAG field set to FFFFh, and type 1 protection (see 4.22.2.3) or type 2 protection (see 4.22.2.4) is enabled; or</li> <li>b) a LOGICAL BLOCK APPLICATION TAG field set to FFFFh, the LOGICAL BLOCK REFERENCE TAG field set to FFFF_FFFFh, and type 3 protection (see 4.22.2.5) is enabled,</li> </ul> <p>then the device server shall not check any protection information in the associated protection information interval.</p> <p><sup>i</sup> If type 1 protection is enabled, then the device server shall check the logical block reference tag by comparing it to the lower four bytes of the LBA associated with the logical block. If type 2 protection or type 3 protection is enabled, and the device server has knowledge of the contents of the LOGICAL BLOCK REFERENCE TAG field, then the device server shall check each logical block reference tag. If type 2 protection is enabled, then this knowledge may be acquired through the EXPECTED INITIAL LOGICAL BLOCK REFERENCE TAG field in a READ (32) command (see 5.16). If type 3 protection is enabled, then the method for acquiring this knowledge is not defined by this standard.</p> <p><sup>j</sup> If the DPICZ bit in the Control mode page (see SPC-5) is set to one, then protection information shall not be checked.</p>					

A disable page out (DPO) bit set to zero specifies that the retention priority shall be determined by the RETENTION PRIORITY fields in the Caching mode page (see 6.5.6). A DPO bit set to one specifies that the device server shall assign the logical blocks accessed by this command the lowest retention priority for being fetched into or retained by the cache (see 4.15). A DPO bit set to one overrides any retention priority specified in the Caching mode page. All other aspects of the algorithm implementing the cache replacement strategy are not defined by this standard.

NOTE 10 - The DPO bit is used to control replacement of logical blocks in the cache when the application client has information on the future usage of the logical blocks. If the DPO bit is set to one, then the application

client is specifying that the logical blocks accessed by the command are not likely to be accessed again in the near future and are not to be put in the cache nor retained by the cache. If the DPO bit is set to zero, then the application client is specifying that the logical blocks accessed by this command are likely to be accessed again in the near future.

A force unit access (FUA) bit set to one specifies that the device server shall read the logical blocks from:

- a) the specified data pattern for that LBA (e.g., the data pattern for unmapped data (see 4.7.4.4));
- b) the non-volatile cache, if any; or
- c) the medium.

If the FUA bit is set to one and a volatile cache contains a more recent version of a logical block than the non-volatile cache, if any, or the medium, then, before reading the logical block, the device server shall write the logical block to:

- a) the non-volatile cache, if any; or
- b) the medium.

An FUA bit set to zero specifies that the device server may read the logical blocks from:

- a) the volatile cache, if any;
- b) the specified data pattern for that LBA (e.g., the data pattern for unmapped data (see 4.7.4.4))
- c) the non-volatile cache, if any; or
- d) the medium.

If rebuild assist mode (see 4.20) is supported and not enabled, then the device server shall ignore the rebuild assist recovery control (RARC) bit. If rebuild assist mode is supported and enabled, then the RARC bit specifies that the device server shall perform read medium operations as defined in 4.20.3.2 and 4.20.3.3.

If the rebuild assist mode is not supported and the RARC bit is set to one, then the device server should terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

See the PRE-FETCH (10) command (see 5.10) for the definition of the LOGICAL BLOCK ADDRESS field.

See the PRE-FETCH (10) command and 4.23 for the definition of the GROUP NUMBER field.

The TRANSFER LENGTH field specifies the number of contiguous logical blocks of data that shall be read and transferred to the Data-In Buffer, starting with the logical block referenced by the LBA specified by the LOGICAL BLOCK ADDRESS field. A TRANSFER LENGTH field set to zero specifies that no logical blocks shall be read or transferred. This condition shall not be considered an error. Any other value specifies the number of logical blocks that shall be read and transferred. If the specified LBA and the specified transfer length exceed the capacity of the medium (see 4.5), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to LOGICAL BLOCK ADDRESS OUT OF RANGE. The TRANSFER LENGTH field is constrained by the MAXIMUM TRANSFER LENGTH field (see 6.6.4).

The CONTROL byte is defined in SAM-5.

## 5.14 READ (12) command

The READ (12) command (see table 68) requests that the device server perform the actions defined for the READ (10) command (see 5.13).

NOTE 11 - Migration from the READ (12) command to the READ (16) command is recommended for all implementations.



**Table 68 — READ (12) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (A8h)							
1		RDPROTECT			DPO	FUA	RARC	Obsolete	Obsolete
2	(MSB)	LOGICAL BLOCK ADDRESS							
...									
5		(LSB)							
6	(MSB)	TRANSFER LENGTH							
...									
9		(LSB)							
10	Restricted for MMC-6	Reserved	GROUP NUMBER						
11		CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 68 for the READ (12) command.

The CONTROL byte is defined in SAM-5.

See the READ (10) command (see 5.13) for the definitions of the other fields in this command.

## 5.15 READ (16) command

The READ (16) command (see table 69) requests that the device server perform the actions defined for the READ (10) command (see 5.13).

**Table 69 — READ (16) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (88h)							
1		RDPROTECT			DPO	FUA	RARC	Obsolete	DLD2
2	(MSB)	LOGICAL BLOCK ADDRESS							
...									
9									
10	(MSB)	TRANSFER LENGTH							
...									
13									
14		DLD1	DLD0	GROUP NUMBER					
15		CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 69 for the READ (16) command.

The command duration time (see SAM-5) is specified by the command duration limit descriptor (see SPC-5) specified by the DLD2 bit, the DLD1 bit, and the DLD0 bit, as shown in table 70. The CDLP field in the REPORT SUPPORTED OPERATION CODES parameter data (see SPC-5) indicates that the command duration limit descriptor is in the Command Duration Limit A mode page or the Command Duration Limit B mode page (see SPC-5).

**Table 70 — Duration limit value bits specifying command duration limit descriptor**

Duration limit descriptor value bits			Command duration limit descriptor specifying command duration time
DLD2	DLD1	DLD0	
0b	0b	0b	Command is not a duration limited command (see SAM-5)
0b	0b	1b	First command duration limit descriptor
0b	1b	0b	Second command duration limit descriptor
0b	1b	1b	Third command duration limit descriptor
1b	0b	0b	Fourth command duration limit descriptor
1b	0b	1b	Fifth command duration limit descriptor
1b	1b	0b	Sixth command duration limit descriptor
1b	1b	1b	Seventh command duration limit descriptor

The CONTROL byte is defined in SAM-5.

See the READ (10) command (see 5.13) for the definitions of the other fields in this command.

## 5.16 READ (32) command

The READ (32) command (see table 71) requests that the device server perform the actions defined for the READ (10) command (see 5.13).

The device server shall only process a READ (32) command if type 2 protection is enabled (see 4.22.2.4).

**Table 71 — READ (32) command**

Bit	7	6	5	4	3	2	1	0
Byte								
0	OPERATION CODE (7Fh)							
1	CONTROL							
2	Reserved							
...								
5								
6	Reserved	GROUP NUMBER						
7	ADDITIONAL CDB LENGTH (18h)							
8	(MSB)	SERVICE ACTION (0009h)						
9								
10	RDPROTECT			DPO	FUA	RARC	Obsolete	Reserved
11	Reserved							
12	(MSB)	LOGICAL BLOCK ADDRESS						
...								
19								
20	(MSB)	EXPECTED INITIAL LOGICAL BLOCK REFERENCE TAG						
...								
23								
24	(MSB)	EXPECTED LOGICAL BLOCK APPLICATION TAG						
25								
26	(MSB)	LOGICAL BLOCK APPLICATION TAG MASK						
27								
28	(MSB)	TRANSFER LENGTH						
...								
31								

The OPERATION CODE field, the ADDITIONAL CDB LENGTH field, and the SERVICE ACTION field are defined in SPC-5 and shall be set to the values shown in table 71 for the READ (32) command.

The CONTROL byte is defined in SAM-5.

See the READ (10) command (see 5.13) for the definitions of the GROUP NUMBER field, the RDPROTECT field, the DPO bit, the FUA bit, the RARC bit, the LOGICAL BLOCK ADDRESS field, and the TRANSFER LENGTH field.

If checking of the LOGICAL BLOCK REFERENCE TAG field is enabled (see table 67 in 5.13), then the EXPECTED INITIAL LOGICAL BLOCK REFERENCE TAG field contains the value of the LOGICAL BLOCK REFERENCE TAG field

expected in the protection information of the first logical block accessed by the command instead of a value based on the LBA (see 4.22.3).

If the ATO bit is set to one in the Control mode page (see SPC-5), and checking of the LOGICAL BLOCK APPLICATION TAG field is enabled (see table 67 in 5.13), then the LOGICAL BLOCK APPLICATION TAG MASK field contains a value that is a bit mask for enabling the checking of the LOGICAL BLOCK APPLICATION TAG field in every instance of protection information for each logical block accessed by the command. A LOGICAL BLOCK APPLICATION TAG MASK field bit set to one enables the checking of the corresponding bit of the EXPECTED LOGICAL BLOCK APPLICATION TAG field with the corresponding bit of the LOGICAL BLOCK APPLICATION TAG field in every instance of protection information. A LOGICAL BLOCK APPLICATION TAG MASK field bit set to zero disables the checking of the corresponding bit of the EXPECTED LOGICAL BLOCK APPLICATION TAG field with the corresponding bit of the LOGICAL BLOCK APPLICATION TAG field in every instance of protection information.

If the ATO bit is set:

- a) to zero; or
- b) to one in the Control mode page (see SPC-5), and checking of the LOGICAL BLOCK APPLICATION TAG field is disabled (see table 67),

then the LOGICAL BLOCK APPLICATION TAG MASK field and the EXPECTED LOGICAL BLOCK APPLICATION TAG field shall be ignored

## 5.17 READ CAPACITY (10) command

### 5.17.1 READ CAPACITY (10) overview

The READ CAPACITY (10) command (see table 72) requests that the device server transfer eight bytes of parameter data describing the capacity and medium format of the direct access block device to the Data-In Buffer. This command may be processed as if it has a HEAD OF QUEUE task attribute (see 4.16). If the logical unit supports protection information (see 4.22) or logical block provisioning management (see 4.7), then the application client should use the READ CAPACITY (16) command (see 5.18) instead of the READ CAPACITY (10) command.

NOTE 12 - Migration from the READ CAPACITY (10) command to the READ CAPACITY (16) command is recommended for all implementations.

**Table 72 — READ CAPACITY (10) command**

Bit	7	6	5	4	3	2	1	0
Byte								
0	OPERATION CODE (25h)							
1	Reserved							Obsolete
2	Obsolete							
...								
5	Reserved							
6								
7	Reserved							
8								
9	CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 72 for the READ CAPACITY (10) command.

The CONTROL byte is defined in SAM-5.

### 5.17.2 READ CAPACITY (10) parameter data

The READ CAPACITY (10) parameter data is defined in table 73. Any time the READ CAPACITY (10) parameter data changes, the device server should establish a unit attention condition as described in 4.10.

**Table 73 — READ CAPACITY (10) parameter data**

Byte	Bit	7	6	5	4	3	2	1	0
0	(MSB)	RETURNED LOGICAL BLOCK ADDRESS							
...									
3									
4	(MSB)	LOGICAL BLOCK LENGTH IN BYTES							
...									
7									

The device server shall set the RETURNED LOGICAL BLOCK ADDRESS field to the lower of:

- a) the LBA of the last logical block on the direct access block device; or
- b) FFFF\_FFFFh, if the LBA of the last logical block on the direct access block device is greater than the maximum value that is able to be specified in the RETURNED LOGICAL BLOCK ADDRESS field.

If the RETURNED LOGICAL BLOCK ADDRESS field is set to FFFF\_FFFFh, then the application client should issue a READ CAPACITY (16) command (see 5.18) to request that the device server transfer the READ CAPACITY (16) parameter data to the Data-In Buffer.

The LOGICAL BLOCK LENGTH IN BYTES field contains the number of bytes of user data in a logical block.

## 5.18 READ CAPACITY (16) command

### 5.18.1 READ CAPACITY (16) command overview

The READ CAPACITY (16) command (see table 74) requests that the device server transfer parameter data describing the capacity and medium format of the direct access block device to the Data-In Buffer. This command is mandatory if the logical unit supports protection information (see 4.22) or logical block provisioning management (see 4.7) and is optional otherwise. This command may be processed as if it has a HEAD OF QUEUE task attribute (see 4.16).

This command uses the SERVICE ACTION IN (16) CDB format (see A.2).

**Table 74 — READ CAPACITY (16) command**

Bit Byte	7	6	5	4	3	2	1	0
0	OPERATION CODE (9Eh)							
1	Reserved			SERVICE ACTION (10h)				
2	Obsolete							
...								
9								
10	(MSB)							
...	ALLOCATION LENGTH							
13								(LSB)
14	Reserved							Obsolete
15	CONTROL							

The OPERATION CODE field and SERVICE ACTION field are defined in SPC-5 and shall be set to the values shown in table 74 for the READ CAPACITY (16) command.

The ALLOCATION LENGTH field is defined in SPC-5.

The CONTROL byte is defined in SAM-5.

### 5.18.2 READ CAPACITY (16) parameter data

The READ CAPACITY (16) parameter data is defined in table 75. Any time the READ CAPACITY (16) parameter data changes, the device server should establish a unit attention condition as described in 4.10.

**Table 75 — READ CAPACITY (16) parameter data**

Byte	Bit	7	6	5	4	3	2	1	0				
0	(MSB)	RETURNED LOGICAL BLOCK ADDRESS											
...													
7										(LSB)			
8	(MSB)	LOGICAL BLOCK LENGTH IN BYTES											
...													
11										(LSB)			
12		Reserved		Restricted for ZBC		P_TYPE		PROT_EN					
13		P_I_EXPONENT				LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT							
14		LBPME	LBPRZ	(MSB)	LOWEST ALIGNED LOGICAL BLOCK ADDRESS								
15		(LSB)											
16		Reserved											
...													
31													

The RETURNED LOGICAL BLOCK ADDRESS field and LOGICAL BLOCK LENGTH IN BYTES field of the READ CAPACITY (16) parameter data are defined in the READ CAPACITY (10) parameter data (see 5.17). The maximum value that shall be returned in the RETURNED LOGICAL BLOCK ADDRESS field is FFFF\_FFFF\_FFFF\_FFEh.

The protection type (P\_TYPE) field and the protection enable (PROT\_EN) bit (see table 76) indicate the logical unit's current type of protection.

**Table 76 — P\_TYPE field and PROT\_EN bit**

P_TYPE	PROT_EN	Description
n/a	0	The logical unit is formatted to type 0 protection (see 4.22.2.2).
000b	1	The logical unit is formatted to type 1 protection (see 4.22.2.3).
001b		The logical unit is formatted to type 2 protection (see 4.22.2.4).
010b		The logical unit is formatted to type 3 protection (see 4.22.2.5).
011b to 111b		Reserved

The P\_I\_EXPONENT field may be used to determine the number of protection information intervals placed within each logical block (see 5.4.2).

The number of protection information intervals is calculated as follows:

$$\text{number of protection information intervals} = 2^{(p\_i \text{ exponent})}$$

where:

$p\_i \text{ exponent}$  is the contents of the  $P\_I \text{ EXPONENT}$  field.

The LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT field is defined in table 77.

**Table 77 — LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT field**

Code	Description
0	One or more physical blocks per logical block <sup>a</sup>
$n > 0$	$2^n$ logical blocks per physical block
<sup>a</sup> The number of physical blocks per logical block is not reported.	

A logical block provisioning management enabled (LBPME) bit set to one indicates that the logical unit implements logical block provisioning management (i.e., is resource provisioned or thin provisioned) (see 4.7.3). An LBPME bit set to zero indicates that the logical unit does not implement logical block provisioning management (e.g., is fully provisioned (see 4.7.2)).

The logical block provisioning read zeros (LBPRZ) bit shall be set to one if the LBPRZ field (see 6.6.6) is set to  $xx1b$ . The LBPRZ bit shall be set to zero if the LBPRZ field is not set to  $xx1b$ .

The LOWEST ALIGNED LOGICAL BLOCK ADDRESS field indicates the LBA of the first logical block that is located at the beginning of a physical block (see 4.6).

## 5.19 READ DEFECT DATA (10) command

### 5.19.1 READ DEFECT DATA (10) command overview

The READ DEFECT DATA (10) command (see table 78) requests that the device server transfer parameter data (see 5.19.2) containing a four-byte header, the PLIST, and/or the GLIST to the Data-In Buffer.

If the device server is unable to access a specified defect list due to a medium error, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to MEDIUM ERROR and the additional sense code set to DEFECT LIST NOT FOUND.

If the device server is unable to access a specified defect list due to an error other than a medium error or because a specified defect list does not exist, then the device server shall either:

- 1) terminate the command with CHECK CONDITION status with the sense key set to NO SENSE and the additional sense code set to DEFECT LIST NOT FOUND; or
- 2) return only the READ DEFECT DATA parameter data header, with the DEFECT LIST LENGTH field set to zero.

Device servers may or may not return a defect list until after a successful completion of a FORMAT UNIT command (see 5.3).

NOTE 13 - Migration from the READ DEFECT DATA (10) command to the READ DEFECT DATA (12) command is recommended for all implementations.



**Table 78 — READ DEFECT DATA (10) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (37h)							
1		Reserved							
2		Reserved			REQ_PLIST	REQ_GLIST	DEFECT LIST FORMAT		
3		Reserved							
...									
6									
7	(MSB)	ALLOCATION LENGTH							
8		(LSB)							
9		CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 78 for the READ DEFECT DATA (10) command.

Table 79 defines the request PLIST (REQ\_PLIST) bit and the request GLIST (REQ\_GLIST) bit.

**Table 79 — REQ\_PLIST bit and REQ\_GLIST bit**

REQ_PLIST	REQ_GLIST	Description
0	0	The device server shall return only the first four bytes of the READ DEFECT DATA parameter data (i.e., the parameter data header), with the DEFECT LIST LENGTH field set to zero.
	1	The device server shall return the READ DEFECT DATA parameter data header and include the GLIST, if any, in the defect list.
1	0	The device server shall return the READ DEFECT DATA parameter data header and include the PLIST, if any, in the defect list.
	1	The device server shall return the READ DEFECT DATA parameter data header and include the both the PLIST, if any, and the GLIST, if any, in the defect list. Whether the PLIST and GLIST are merged or not is vendor specific.

The DEFECT LIST FORMAT field specifies the address descriptor format type (see 6.2) that the device server should use for the defect list. A device server unable to return the requested address descriptor format shall return the address descriptors in their default format and indicate that format type in the DEFECT LIST FORMAT field in (see 5.19.2 and 5.20.2).

If the requested defect list format and the returned defect list format are not the same, then the device server shall transfer the defect data and then terminate the command with CHECK CONDITION status with the sense key set to RECOVERED ERROR and the additional sense code set to DEFECT LIST NOT FOUND.

The ALLOCATION LENGTH field is defined in SPC-5. If the length of the address descriptors that the device server has to report is greater than the maximum value that is able to be specified by the ALLOCATION LENGTH field, then the device server shall transfer no data and shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

The CONTROL byte is defined in SAM-5.

### 5.19.2 READ DEFECT DATA (10) parameter data

The READ DEFECT DATA (10) parameter data (see table 80) contains a four-byte header, followed by zero or more address descriptors.

**Table 80 — READ DEFECT DATA (10) parameter data**

Byte	Bit	7	6	5	4	3	2	1	0
Parameter data header									
0	Reserved								
1	Reserved			PLISTV		GLISTV		DEFECT LIST FORMAT	
2	(MSB) _____								
3	DEFECT LIST LENGTH (n - 3) _____ (LSB)								
Defect list (if any)									
4	Address descriptor(s) (if any) _____								
...									
n									

A PLIST valid (PLISTV) bit set to zero indicates that the defect list does not contain the PLIST. A PLISTV bit set to one indicates that the defect list contains the PLIST.

A GLIST valid (GLISTV) bit set to zero indicates that the defect list does not contain the GLIST. A GLISTV bit set to one indicates that the defect list contains the GLIST.

The DEFECT LIST FORMAT field indicates the format of the address descriptors returned in the defect list. This field is defined in 6.2.

If the device server returns short block format address descriptors (see 6.2.2) or long block format address descriptors (see 6.2.5), then the address descriptors contain vendor-specific values.

The DEFECT LIST LENGTH field indicates the length in bytes of the defect list. The DEFECT LIST LENGTH is equal to four or eight times the number of the address descriptors, depending on the format of the returned address descriptors (see 6.2).

The defect list contains address descriptors (see 6.2).

## 5.20 READ DEFECT DATA (12) command

### 5.20.1 READ DEFECT DATA (12) command overview

The READ DEFECT DATA (12) command (see table 81) requests that the device server transfer parameter data (see 5.20.2) containing a four-byte header, the PLIST, and/or the GLIST to the Data-In Buffer.

An application client determines the length of the defect list by sending a READ DEFECT DATA (12) command with an allocation length field set to eight and the address descriptor index field set to 0000\_0000h. The device server returns the defect list header that contains the length of the defect list.

**Table 81 — READ DEFECT DATA (12) command**

Bit	7	6	5	4	3	2	1	0
Byte								
0	OPERATION CODE (B7h)							
1	Reserved			REQ_PLIST	REQ_GLIST	DEFECT LIST FORMAT		
2	(MSB)							
...	ADDRESS DESCRIPTOR INDEX							
5	(LSB)							
6	(MSB)							
...	ALLOCATION LENGTH							
9	(LSB)							
10	Reserved							
11	CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 81 for the READ DEFECT DATA (12) command.

See the READ DEFECT DATA (10) command (see 5.19) for the definitions of the REQ\_PLIST bit, the REQ\_GLIST bit, and the DEFECT LIST FORMAT field.

The ADDRESS DESCRIPTOR INDEX field specifies the index of the first address descriptor (see 6.2) in the defect list that the device server shall return. If the ADDRESS DESCRIPTOR INDEX field is set to:

- a value less than the number of available address descriptors, then the device server shall transfer a defect list beginning with the address descriptor that is at the ADDRESS DESCRIPTOR INDEX field value multiplied by the size of the address descriptor; or
- a value greater than or equal to the number of available address descriptors, then the device server shall return a zero length defect list.

The ALLOCATION LENGTH field is defined in SPC-5, however if the length of all the address descriptors that are available is greater than FFFF\_FFFFh, then the device server shall transfer the length of address descriptors specified by the allocation length or the DEFECT LIST LENGTH field value plus eight, whichever is less, and complete the command with GOOD status.

The CONTROL byte is defined in SAM-5.

### 5.20.2 READ DEFECT DATA (12) parameter data

The READ DEFECT DATA (12) parameter data (see table 82) contains an eight-byte header, followed by zero or more address descriptors.

**Table 82 — READ DEFECT DATA (12) parameter data**

Byte	Bit	7	6	5	4	3	2	1	0
Parameter data header									
0	Reserved								
1	Reserved			PLISTV		GLISTV		DEFECT LIST FORMAT	
2	(MSB) _____								
3	GENERATION CODE _____ (LSB)								
4	(MSB) _____								
...	DEFECT LIST LENGTH (n - 7) _____								
7	(LSB)								
Defect list (if any)									
8	_____								
...	Address descriptor(s) (if any) _____								
n	_____								

The GENERATION CODE field is a two-byte counter that shall be incremented by one by the device server every time the defect list is changed. A GENERATION CODE field set to 0000h indicates the generation code is not supported. If the GENERATION CODE field is supported, then the GENERATION CODE field shall be initialized to at least 0001h at power on and the device server shall wrap this field to 0001h as the next increment after reaching its maximum value (i.e., FFFFh).

Application clients that use the GENERATION CODE field should read this field often enough to ensure that the contents of this field do not increment a multiple of 65 535 times between readings.

The DEFECT LIST LENGTH field indicates the length in bytes of address descriptors from the beginning address descriptor specified by the ADDRESS DESCRIPTOR INDEX field to the last address descriptor available to be returned. A value of FFFF\_FFFFh in the DEFECT LIST LENGTH field indicates that more than FFFF\_FFFEh bytes are available.

See the READ DEFECT DATA (10) command (see 5.19) for the definitions of the other fields in the READ DEFECT DATA (12) parameter data.

## 5.21 REASSIGN BLOCKS command

### 5.21.1 REASSIGN BLOCKS command overview

The REASSIGN BLOCKS command (see table 83) requests that the device server perform a reassign operation on one or more LBAs (e.g., LBAs referencing logical blocks on which unrecovered read errors occurred) to another area on the medium set aside for this purpose and to add the physical blocks containing those logical blocks to the GLIST. This command shall not alter the contents of the PLIST (see 4.13).

The parameter list provided in the Data-Out Buffer contains a reassign LBA list that contains the LBAs of the logical blocks to be reassigned. The device server shall reassign the parts of the medium used for each logical block referenced by an LBA in the reassign LBA list. More than one physical block may be reassigned by each LBA. If the device server recovers logical block data from the original logical block, then the device server shall perform a write medium operation to that LBA using the recovered logical block data, which writes to the logical block referenced by the reassigned LBA.

The device server shall invalidate any of the specified LBAs that are in cache.

If the device server does not recover logical block data in a fully provisioned logical unit (see 4.7.2), then the device server shall:

- a) write vendor specific data as the user data; and
- b) write a default value of FFFF\_FFFF\_FFFF\_FFFFh as the protection information, if enabled (see 4.22.2).

If the device server does not recover logical block data in a resource provisioned logical unit (see 4.7.3.2) or a thin provisioned logical unit (see 4.7.3.3), then the device server shall, for each specified LBA, either:

- a) unmap the specified LBA; or
- b) perform the following operations:
  - A) write vendor-specific data as the user data; and
  - B) write a default value of FFFF\_FFFF\_FFFF\_FFFFh as the protection information, if enabled.

The vendor-specific data written as user data may contain remnants of the original logical block (e.g., partially or fully recovered user data).

The data in all other logical blocks on the medium shall be preserved.

Specifying an LBA to be reassigned that previously has been reassigned causes the device server to reassign that LBA again.

If the device server terminates the REASSIGN BLOCKS command with CHECK CONDITION status, and the sense data COMMAND-SPECIFIC INFORMATION field contains a valid LBA, then the application client should remove all LBAs from the reassign LBA list prior to the one returned in the COMMAND-SPECIFIC INFORMATION field. If the sense key is set to MEDIUM ERROR and the INFORMATION field contains the valid LBA, then the application client should insert that LBA into the reassign LBA list and reissue the REASSIGN BLOCKS command with the new reassign LBA list. Otherwise, the application client should perform any corrective action indicated by the sense data and then reissue the REASSIGN BLOCKS command with the new reassign LBA list.

**Table 83 — REASSIGN BLOCKS command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (07h)							
1		Reserved						LONGLBA	ONGLIST
2		Reserved							
...									
4									
5		CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 83 for the REASSIGN BLOCKS command.

A long LBA (LONGLBA) bit set to zero specifies that the reassign LBA list in the REASSIGN BLOCKS parameter list (see 5.21.2) contains four-byte LBAs. A LONGLBA bit set to one specifies that the reassign LBA list in the REASSIGN BLOCKS parameter list contains eight-byte LBAs.

A long list (ONGLIST) bit set to zero specifies the REASSIGN BLOCKS short parameter list header (see table 85) is used. A ONGLIST bit set to one specifies the REASSIGN BLOCKS long parameter list header (see table 86) is used.

The CONTROL byte is defined in SAM-5.

### 5.21.2 REASSIGN BLOCKS parameter list

The REASSIGN BLOCKS parameter list (see table 84) contains a four-byte parameter list header followed by a reassign LBA list containing one or more LBAs.

**Table 84 — REASSIGN BLOCKS parameter list**

Byte	Bit	7	6	5	4	3	2	1	0
0		Parameter list header (see table 85 or table 86)							
...									
3									
Reassign LBA list (if any)									
4		Reassign LBA [first] (see table 87 or table 88)							
...									
7 or 11									
n-4 or n-7		Reassign LBA [last] (see table 87 or table 88)							
...									
n									

The REASSIGN BLOCKS short parameter list header is defined in table 85.

**Table 85 — REASSIGN BLOCKS short parameter list header**

Bit	7	6	5	4	3	2	1	0
Byte								
0	Reserved							
1								
2	(MSB)	REASSIGN LBA LENGTH						
3								(LSB)

The REASSIGN BLOCKS long parameter list header is defined in table 86.

**Table 86 — REASSIGN BLOCKS long parameter list header**

Bit	7	6	5	4	3	2	1	0
Byte								
0	(MSB)							
...	REASSIGN LBA LENGTH							
3								

The REASSIGN LBA LENGTH field specifies the total length in bytes of the reassign LBA list. The REASSIGN LBA LENGTH field does not include the parameter list header length and is equal to:

- a) four times the number of LBAs, if the `LONGLBA` bit (see 5.21) is set to zero; or
- b) eight times the number of LBAs, if the `LONGLBA` bit is set to one.

The REASSIGN LBA LIST field contains a list of LBAs to be reassigned. The LBAs shall be sorted in ascending order.

If the `LONGLBA` bit is set to zero, then table 87 defines the format of the reassigned LBA.

**Table 87 — Reassign LBA if the `LONGLBA` bit is set to zero**

Bit	7	6	5	4	3	2	1	0
Byte								
0	(MSB)							
...	REASSIGN LOGICAL BLOCK ADDRESS							
3								

The REASSIGN LOGICAL BLOCK ADDRESS field specifies an LBA to be reassigned.

If the LONGLBA bit is set to one, then table 88 defines the reassign LBA.

**Table 88 — Reassign LBA if the LONGLBA bit is set to one**

Bit	7	6	5	4	3	2	1	0
Byte								
0	(MSB)							
...	REASSIGN LOGICAL BLOCK ADDRESS							
7	(LSB)							

The REASSIGN LOGICAL BLOCK ADDRESS field specifies an LBA to be reassigned.

If a specified LBA exceeds the capacity of the medium (see 4.5), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code should be set to LOGICAL BLOCK ADDRESS OUT OF RANGE or may be set to INVALID FIELD IN PARAMETER LIST.

If the direct access block device has insufficient capacity to reassign all of the specified LBAs, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to HARDWARE ERROR and the additional sense code set to NO DEFECT SPARE LOCATION AVAILABLE.

If the direct access block device is unable to complete a REASSIGN BLOCKS command without error, then the device server shall terminate the command with CHECK CONDITION status with the appropriate sense data (see 4.18 and **SPC-4**SPC-5).

If one or more LBAs are not reassigned, then the device server shall report the first LBA not reassigned in the COMMAND-SPECIFIC INFORMATION field of the sense data (see SPC-5). If:

- a) information about the first LBA not reassigned is not available;
- b) all the LBAs have been reassigned; or
- c) the first LBA not reassigned does not fit in the COMMAND-SPECIFIC INFORMATION field, then the device server shall report the following value in the COMMAND-SPECIFIC INFORMATION field of the sense data (see SPC-5):
  - A) FFFF\_FFFFh if fixed format sense data is being used; or
  - B) FFFF\_FFFF\_FFFF\_FFFFh if descriptor format sense data is being used.

If the REASSIGN BLOCKS command failed due to an unexpected unrecovered read error that would cause the loss of data in a logical block not specified in the reassign LBA list, then the LBA of the logical block with the unrecovered read error is reported in the INFORMATION field of the sense data(see 4.18.1).



## 5.22 RECEIVE ROD TOKEN INFORMATION

### 5.22.1 RECEIVE ROD TOKEN INFORMATION overview

The RECEIVE ROD TOKEN INFORMATION command (see SPC-5) provides a method for an application client to receive information about the results of a previous or current block device ROD token operation. Table 89 shows the operations and a reference to the subclause where each topic is described.

**Table 89 — RECEIVE ROD TOKEN INFORMATION reference**

Command originating the operation	Command reference	RECEIVE ROD TOKEN INFORMATION returned parameter data reference
POPULATE TOKEN	5.9	5.22.2
WRITE USING TOKEN	5.52	5.22.3

### 5.22.2 RECEIVE ROD TOKEN INFORMATION parameter data for POPULATE TOKEN command

If a RECEIVE ROD TOKEN INFORMATION command (see ~~SPC-4~~SPC-5) specifies a list identifier that matches the list identifier specified in a previous POPULATE TOKEN command (see 5.9) received on the same I\_T nexus, then table 90 shows the parameter data returned by the copy manager.

**Table 90 — RECEIVE ROD TOKEN INFORMATION parameter data for POPULATE TOKEN**

Byte	Bit	7	6	5	4	3	2	1	0	
0	(MSB)	AVAILABLE DATA (n - 3)								
...										
3										(LSB)
4		Reserved			RESPONSE TO SERVICE ACTION (10h)					
5	Reserved	COPY OPERATION STATUS								
6	(MSB)	OPERATION COUNTER								
7										(LSB)
8	(MSB)									
...		ESTIMATED STATUS UPDATE DELAY								
11										(LSB)
12		EXTENDED COPY COMPLETION STATUS								
13		LENGTH OF THE SENSE DATA FIELD (m - 31)								
14		SENSE DATA LENGTH								
15		TRANSFER COUNT UNITS (F1h)								
16	(MSB)	TRANSFER COUNT								
...										
23										(LSB)
24	(MSB)	SEGMENTS PROCESSED (0000h)								
25										(LSB)
26										
...		Reserved								
31										
32										
...		SENSE DATA (if any)								
m										
m + 1	(MSB)									
...		ROD TOKEN DESCRIPTOR LENGTH (n - (m + 4))								
m + 4										(LSB)
m + 5		Restricted (see SPC-5)								
m + 6										
m + 7		ROD TOKEN (if any)								
...										
n										

The AVAILABLE DATA field, the COPY OPERATION STATUS field, the OPERATION COUNTER field, the ESTIMATED STATUS UPDATE DELAY field, the EXTENDED COPY COMPLETION STATUS field, the LENGTH OF THE SENSE DATA FIELD field, SENSE DATA LENGTH field, the SENSE DATA field, and the ROD TOKEN field are defined in SPC-5.

The RESPONSE TO SERVICE ACTION field is defined in SPC-5 and shall be set to the value shown in table 90 in response to a RECEIVE ROD TOKEN INFORMATION command in which the LIST IDENTIFIER field specifies a POPULATE TOKEN command.

The TRANSFER COUNT UNITS field is defined in SPC-5 and shall be set to the value shown in table 90 in response to a RECEIVE ROD TOKEN INFORMATION command in which the LIST IDENTIFIER field specifies a POPULATE TOKEN command.

The TRANSFER COUNT field indicates the number of contiguous logical blocks represented by the ROD token that were read without error starting at the LBA specified in the first block device range descriptor and including the LBAs described in all complete block device range descriptors of the POPULATE TOKEN command to which this response applies.

If the value in the TRANSFER COUNT field is not equal to the sum of the contents of the NUMBER OF LOGICAL BLOCKS fields in all of the complete block device range descriptors of the POPULATE TOKEN command to which this response applies, then the COPY OPERATION STATUS field shall be set to 3h. Other values in the COPY OPERATION STATUS field are defined in SPC-5.

The SEGMENTS PROCESSED field is defined in SPC-5 and shall be set to the value shown in table 90 in response to a RECEIVE ROD TOKEN INFORMATION command in which the LIST IDENTIFIER field specifies a POPULATE TOKEN command.

The ROD TOKEN DESCRIPTOR LENGTH field is defined in SPC-5 and shall be set to the size of the ROD TOKEN field plus two in response to a RECEIVE ROD TOKEN INFORMATION command in which the LIST IDENTIFIER field specifies a POPULATE TOKEN command.

### 5.22.3 RECEIVE ROD TOKEN INFORMATION parameter data for WRITE USING TOKEN command

If a RECEIVE ROD TOKEN INFORMATION command (see SPC-5) specifies a list identifier that matches the list identifier specified in a previous WRITE USING TOKEN command (see 5.52) received on the same I\_T nexus, then table 91 shows the parameter data returned by the copy manager.

**Table 91 — RECEIVE ROD TOKEN INFORMATION parameter data for WRITE USING TOKEN**

Byte	Bit	7	6	5	4	3	2	1	0
0	(MSB)	AVAILABLE DATA (n - 3)							
3									
4		Reserved			RESPONSE TO SERVICE ACTION (11h)				
5	Reserved	COPY OPERATION STATUS							
6	(MSB)	OPERATION COUNTER							
7									
8	(MSB)	ESTIMATED STATUS UPDATE DELAY							
...									
11		(LSB)							
12		EXTENDED COPY COMPLETION STATUS							
13		LENGTH OF THE SENSE DATA FIELD ((n - 4) - 31)							
14		SENSE DATA LENGTH							
15		TRANSFER COUNT UNITS (F1h)							
16	(MSB)	TRANSFER COUNT							
...									
23		(LSB)							
24	(MSB)	SEGMENTS PROCESSED (0000h)							
25									
26		Reserved							
...									
31									
32		SENSE DATA (if any)							
...									
n - 4									
n - 3		Restricted (see <span style="color: red;">SPC-4</span> SPC-5)							
n									

The AVAILABLE DATA field, the COPY OPERATION STATUS field, the OPERATION COUNTER field, the ESTIMATED STATUS UPDATE DELAY field, the EXTENDED COPY COMPLETION STATUS field, the LENGTH OF THE SENSE DATA FIELD field, the SENSE DATA LENGTH field, and the SENSE DATA field are defined in [SPC-4](#)SPC-5.

The RESPONSE TO SERVICE ACTION field is defined in SPC-5 and shall be set to the value shown in table 91 in response to a RECEIVE ROD TOKEN INFORMATION command in which the LIST IDENTIFIER field specifies a WRITE USING TOKEN command.

The TRANSFER COUNT UNITS field is defined in SPC-5 and shall be set to the value shown in table 91 in response to a RECEIVE ROD TOKEN INFORMATION command in which the LIST IDENTIFIER field specifies a WRITE USING TOKEN command.

The TRANSFER COUNT field indicates the number of contiguous logical blocks that were written without error starting with the LBA specified in the first block device range descriptor and including the LBAs specified in all block device range descriptors of the WRITE USING TOKEN command to which this response applies.

If the value in the TRANSFER COUNT field is not equal to the sum of the contents of the NUMBER OF LOGICAL BLOCKS fields in all of the complete block device range descriptors of the WRITE USING TOKEN command to which this response applies, then the COPY OPERATION STATUS field shall be set to 3h. Other values in the COPY OPERATION STATUS field are defined in SPC-5.

The SEGMENTS PROCESSED field is defined in SPC-5 and shall be set to the value shown in table 91 in response to a RECEIVE ROD TOKEN INFORMATION command in which the LIST IDENTIFIER field specifies a POPULATE TOKEN command.

## 5.23 REPORT REFERRALS command

### 5.23.1 REPORT REFERRALS command overview

The REPORT REFERRALS command (see table 92) requests that the device server transfer parameter data indicating the user data segment(s) on the logical unit and the SCSI target ports through which those user data segments may be accessed (see 4.28) to the Data-In Buffer. This command shall be supported by a logical unit that reports in the Extended INQUIRY Data VPD page (see SPC-5) that it supports referrals (i.e., the R\_SUP bit set to one).

This command uses the SERVICE ACTION IN (16) CDB format (see A.2).

**Table 92 — REPORT REFERRALS command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (9Eh)							
1		Reserved			SERVICE ACTION (13h)				
2		(MSB)							
...		LOGICAL BLOCK ADDRESS							
9									
10		(MSB)							
...		ALLOCATION LENGTH							
13									
14		Reserved							ONE_SEG
15		CONTROL							

The OPERATION CODE field and SERVICE ACTION field are defined in SPC-5 and shall be set to the values shown in table 92 for the REPORT REFERRALS command.

The LOGICAL BLOCK ADDRESS field specifies an LBA in the first user data segment that the device server shall report in the REPORT REFERRALS parameter data. If the specified LBA exceeds the capacity of the medium (see 4.5), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to LOGICAL BLOCK ADDRESS OUT OF RANGE.

The ALLOCATION LENGTH field is defined in SPC-5.

A one segment (ONE\_SEG) bit set to zero specifies that the device server shall return information on all user data segments starting with the user data segment that contains the LBA specified in the LOGICAL BLOCK ADDRESS field and ending with the user data segment that contains the last LBA of the logical unit. A ONE\_SEG bit set to one specifies the device server shall only return information on the user data segment that contains the LBA specified in the LOGICAL BLOCK ADDRESS field.

The CONTROL byte is defined in SAM-5.

### 5.23.2 REPORT REFERRALS parameter data

The REPORT REFERRALS parameter data (see table 93) contains information indicating the user data segment(s) on the logical unit and the SCSI target port groups through which those user data segments may be accessed (see 4.28).

**Table 93 — REPORT REFERRALS parameter data**

Bit Byte	7	6	5	4	3	2	1	0
0	Reserved							
1								
2	(MSB)	USER DATA SEGMENT REFERRAL DESCRIPTOR LENGTH (y - 3)						
3								(LSB)
User data segment referral descriptor list								
4	User data segment referral descriptor [first] (if any)							
...								
4 + n								
⋮								
y - m	User data segment referral descriptor [last] (if any)							
...								
y								

The USER DATA SEGMENT REFERRAL DESCRIPTOR LENGTH field indicates the number of bytes that follow in the REPORT REFERRALS parameter data.

The user data segment referral descriptor (see table 18) is defined in the user data segment referral sense data descriptor (see 4.18.4).

## 5.24 REPORT PROVISIONING INITIALIZATION PATTERN command

The REPORT PROVISIONING INITIALIZATION PATTERN command (see table 94) requests that the device server transfer the provisioning initialization pattern to the Data-In Buffer.

**Table 94 — REPORT PROVISIONING INITIALIZATION PATTERN command**

Bit	7	6	5	4	3	2	1	0
Byte								
0	OPERATION CODE (3A3h)							
1	Reserved				SERVICE ACTION (1Dh)			
2	Reserved							
...								
5								
6	(MSB)							
...	ALLOCATION LENGTH							
9								
10	Reserved							
11	CONTROL							

The OPERATION CODE field and SERVICE ACTION field are defined in SPC-5 and shall be set to the value shown in table 94 for the REPORT PROVISIONING INITIALIZATION PATTERN command.

The ALLOCATION LENGTH field is defined in SPC-5.

The CONTROL byte is defined in SAM-5.

## 5.25 SANITIZE command

### 5.25.1 SANITIZE command overview

The SANITIZE command (see table 95) requests that the device server perform a sanitize operation (see 4.11). This device server shall process this command as if it has a HEAD OF QUEUE task attribute (see 4.16).

**Table 95 — SANITIZE command**

Bit	7	6	5	4	3	2	1	0
Byte								
0	OPERATION CODE (48h)							
1	IMMED	ZNR	AUSE	SERVICE ACTION				
2	Reserved							
...								
6								
7	(MSB)	PARAMETER LIST LENGTH						(LSB)
8								
9	CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 95 for the SANITIZE command.

If the immediate (IMMED) bit is set to zero, then the device server shall return status after the sanitize operation is completed. If the IMMED bit set to one, then the device server shall return status as soon as the CDB and parameter data, if any, have been validated. The REQUEST SENSE command may be used to poll for progress of the sanitize operation regardless of the value of the IMMED bit.

For a zoned block device, a zoned no reset (ZNR) bit set to zero specifies that as part of completing a sanitize operation (see 4.11), the device server shall perform the equivalent of a RESET WRITE POINTER command (see ZBC) with the ALL bit set to one.

For a zoned block device, a ZNR bit set to one specifies that as part of completing a sanitize operation, the device server:

- a) unless otherwise specified (e.g., a write pointer zone is set to offline), shall not modify the write pointer (see ZBC) for any write pointer zone sanitized by a sanitize cryptographic erase operation; and
- b) shall perform the equivalent of a FINISH ZONE command (see ZBC) with the ALL bit set to one as part of a sanitize overwrite operation or a sanitize block erase operation.

For a logical unit that is not a zoned block device the ZNR bit shall be ignored.

If the allow unrestricted sanitize exit (AUSE) bit is set to one, and the specified sanitize operation fails, then the device server shall process a subsequent EXIT FAILURE MODE service action as if the previous sanitize operation had completed without error (see 4.11.4)

If:

- 1) the AUSE bit is set to zero in the SANITIZE command that requested a sanitize operation;
- 2) that sanitize operation completes with an error (see 4.11.4); and
- 3) a subsequent SANITIZE command is received with:
  - A) the EXIT FAILURE MODE service action; or
  - B) any service action with the AUSE bit set to one,

then the device server shall terminate that subsequent SANITIZE command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

The SERVICE ACTION field is defined in 5.25.2.

The PARAMETER LIST LENGTH field specifies the length in bytes of the parameter data that is available to be transferred from the Data-Out Buffer. A PARAMETER LIST LENGTH field set to zero specifies that no data shall be transferred.

The CONTROL byte is defined in SAM-5.

## 5.25.2 SANITIZE command service actions

### 5.25.2.1 SANITIZE command service actions overview

The SANITIZE command service actions are defined in table 96. At least one service action shall be supported if the SANITIZE command is supported. If the service action specified in the CDB is not supported, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.



If deferred microcode has been saved and not activated (see SPC-5), then the device server shall terminate this command with CHECK CONDITION status with the sense key set to NOT READY and the additional sense code set to LOGICAL UNIT NOT READY, MICROCODE ACTIVATION REQUIRED.

**Table 96 — SANITIZE service action codes**

Code	Name	Description	PARAMETER LIST LENGTH requirement <sup>a</sup>	Reference
01h	OVERWRITE	Perform a sanitize overwrite operation	Set to > 0004h and < (logical block length + 5)	5.25.2.2
02h	BLOCK ERASE	Perform a sanitize block erase operation	Set to 0000h	5.25.2.3
03h	CRYPTOGRAPHIC ERASE	Perform a sanitize cryptographic erase operation	Set to 0000h	5.25.2.4
1Fh	EXIT FAILURE MODE	Exit the sanitize failure mode	Set to 0000h	5.25.2.5
all others	Reserved			
<sup>a</sup> If the requirement is not met, then the SANITIZE command is terminated with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.				

#### 5.25.2.2 OVERWRITE service action

The OVERWRITE service action (see table 96) requests that the device server perform a sanitize overwrite operation (see 4.11).

While performing a sanitize overwrite operation, the device server shall remove all Background Scan Results log parameters (see 6.4.2.3) from the Background Scan Results log page, if supported, and remove all Pending Defect log parameters (see 6.4.8.3) from the Pending Defects log page, if supported.

The parameter list format for the OVERWRITE service action is defined in table 97.

**Table 97 — OVERWRITE service action parameter list**

Bit	7	6	5	4	3	2	1	0
Byte								
0	INVERT	TEST		OVERWRITE COUNT				
1	Reserved							
2	(MSB)	INITIALIZATION PATTERN LENGTH (n - 3)						LSB)
3								
4		INITIALIZATION PATTERN						
...								
n								

If the INVERT bit is set to zero, then on each overwrite pass:

- the user data shall be written as specified in the INITIALIZATION PATTERN field; and
- the protection information, if any, shall be set to FFFF\_FFFF\_FFFF\_FFFFh.

If the INVERT bit is set to one, then the user data and protection information bytes, if any, shall be inverted (i.e., each bit XORed with one) between consecutive overwrite passes.

The TEST field is defined in table 98.

**Table 98 — TEST field**

Code	Description
00b	Shall not cause any changes in the defined behavior of the SANITIZE command.
01b to 11b	Vendor specific <sup>a</sup>
<sup>a</sup> Setting the TEST field to one of these values may adversely affect security properties of the OVERWRITE service action.	

The OVERWRITE COUNT field specifies the number of overwrite passes to be performed. The value of 00h is reserved.

The INITIALIZATION PATTERN LENGTH field specifies the length in bytes of the INITIALIZATION PATTERN field. If the INITIALIZATION PATTERN LENGTH field is set to zero or a value greater than the logical block length, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

The INITIALIZATION PATTERN field specifies the data pattern to be used to write the user data. This data pattern is repeated as necessary to fill each logical block. For each logical block, the first byte of the user data shall begin with the first byte of the initialization pattern.

If the INVERT bit is set to one and:

- a) the OVERWRITE COUNT field is set to an even number, then the pattern used for the first write pass shall consist of:
  - A) the user data set to the inversion of the INITIALIZATION PATTERN data; and
  - B) the protection information, if any, set to 0000\_0000\_0000\_0000h;
 or
- b) the OVERWRITE COUNT field is set to an odd number, then the pattern used for the first write pass shall consist of:
  - A) the user data set to the INITIALIZATION PATTERN data; and
  - B) the protection information, if any, set to FFFF\_FFFF\_FFFF\_FFFFh.

After a sanitize overwrite operation completes without error:

- a) the device server completes read commands for which no other error occurs during processing with GOOD status and read medium operations return the data written by the sanitize overwrite operation; and
- b) protection information, if any, shall be set to FFFF\_FFFF\_FFFF\_FFFFh in all logical blocks on the medium.

### 5.25.2.3 BLOCK ERASE service action

The BLOCK ERASE service action (see table 96) requests that the device server perform a sanitize block erase operation (see 4.11).

After a sanitize block erase operation completes without error:

- a) the device server may terminate commands that request read operations specifying mapped LBAs (see 4.7.1) based on the setting of the WABEREQ field ~~in the Block Device Characteristics VPD page~~ (see 6.6.2); and
- b) if the logical unit is formatted with protection information, then:
  - A) the protection information for each mapped LBA may be indeterminate; and

- B) if the device server terminates a command that requests read operations specifying mapped LBAs as a result of a protection information error, then the device server shall terminate that command with CHECK CONDITION status with the sense key set to ABORTED COMMAND and the appropriate additional sense code for the condition (e.g., for READ commands, the additional sense code defined in table 67).

#### 5.25.2.4 CRYPTOGRAPHIC ERASE service action

The CRYPTOGRAPHIC ERASE service action (see table 96) requests that the device server perform a sanitize cryptographic erase operation (see 4.11).

After a sanitize cryptographic erase operation completes without error:

- a) the device server may terminate commands that request read operations specifying mapped LBAs (see 4.7.1) based on the setting of the WACERREQ field ~~in the Block Device Characteristics VPD page~~ (see 6.6.2); and
- b) if the logical unit is formatted with protection information, then:
  - A) the protection information for each mapped LBA may be indeterminate; and
  - B) if the device server terminates a command that requests read operations specifying mapped LBAs as a result of a protection information error, then the device server shall terminate that command with CHECK CONDITION status with the sense key set to ABORTED COMMAND and the appropriate additional sense code for the condition (e.g., for READ commands, the additional sense code defined in table 67).

#### 5.25.2.5 EXIT FAILURE MODE service action

The EXIT FAILURE MODE service action (see table 96) requests that the device server complete a sanitize operation which completed with an error as if the sanitize operation completed without an error (see 4.11). If the most recent sanitize operation, if any, has completed without error, then the EXIT FAILURE MODE service action completes without error.

After successful completion of a SANITIZE command with the EXIT FAILURE MODE service action:

- a) if any LBA is mapped (see 4.7.1), and the logical unit is formatted with protection information, then:
  - A) the protection information for each mapped LBA may be indeterminate; and

if the device server terminates a command that requests read operations specifying mapped LBAs as a result of a protection information error, then the device server shall terminate that command with CHECK CONDITION status with the sense key set to ABORTED COMMAND and the appropriate additional sense code for the condition (e.g., for READ commands, the additional sense code defined in table 67);
- b) the device server should complete reads to unmapped LBAs without error (see 4.7.4.6.1 and 4.7.4.7.1);
- c) If the logical unit is a zoned block device and the ZNR bit in the CDB is set to zero then:
  - A) the device server shall perform the equivalent of a RESET WRITE POINTER command (see ZBC) with the ALL bit set to one; and
  - B) for each write pointer zone, if the reset write pointer operation is not successful, then the zone condition (see ZBC) shall be set to offline;

and
- d) If the logical unit is a zoned block device and the ZNR bit in the CDB is set to one then, for any write pointer zone where the write pointer was not reset as part of the failed SANITIZE command, the device server shall not modify the write pointer (see ZBC).

## 5.26 START STOP UNIT command

The START STOP UNIT command (see table 99) requests that the device server change the power condition of the logical unit (see 4.21) or load or eject the medium. This includes specifying that the device server enable or disable the direct access block device for medium access operations by controlling power conditions and timers.

The device server shall handle any deferred microcode as specified in 4.26.

**Table 99 — START STOP UNIT command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (1Bh)							
1		Reserved							IMMED
2		Reserved							
3		Reserved				POWER CONDITION MODIFIER			
4		POWER CONDITION				Reserved	NO_FLUSH	LOEJ	START
5		CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 99 for the START STOP UNIT command.

If the immediate (IMMED) bit is set to zero, then the device server shall return status after the operation is completed. If the IMMED bit is set to one, then the device server shall return status as soon as the CDB has been validated.

The combinations of values in the POWER CONDITION field and the POWER CONDITION MODIFIER field are defined in table 100. If the POWER CONDITION field is supported and is set to a value other than 0h, then the device server shall ignore the START bit and the LOEJ bit.

**Table 100 — POWER CONDITION and POWER CONDITION MODIFIER field (part 1 of 2)**

POWER CONDITION value	POWER CONDITION name	POWER CONDITION MODIFIER value	Device server action
0h	START_ VALID	0h	Process the START bit and the LOEJ bit.
1h	ACTIVE	0h	Cause the logical unit to transition to the active power condition. <sup>a</sup>
2h	IDLE	0h	Cause the logical unit to transition to the idle_a power condition. <sup>a b</sup>
		1h	Cause the logical unit to transition to the idle_b power condition. <sup>a b c</sup>
		2h	Cause the logical unit to transition to the idle_c power condition. <sup>a b d</sup>
3h	STANDBY	0h	Cause the logical unit to transition to the standby_z power condition. <sup>a b</sup>
		1h	Cause the logical unit to transition to the standby_y power condition. <sup>a b</sup>
5h	Obsolete	0h to Fh	Obsolete
7h	LU_ CONTROL	0h	Initialize and start all of the idle condition timers that are enabled (see SPC-5), and initialize and start all of the standby condition timers that are enabled (see SPC-5).

<sup>a</sup> Process the following actions:

- 1) The device server shall comply with any SCSI transport protocol specific power condition transition restrictions (e.g., the NOTIFY (ENABLE SPINUP) requirement (see SPL-2));
- 2) the logical unit shall transition to the specified power condition; and
- 3) the device server shall disable all of the idle condition timers that are enabled (see SPC-5) and disable all of the standby condition timers that are enabled (see SPC-5) until another START STOP UNIT command is processed that returns control of the power condition to the logical unit or a logical unit reset occurs.

<sup>b</sup> If a timer for a background scan operation expires, or a device specific event occurs, then the logical unit shall not leave this power condition to perform the background function associated with the timer or event.

<sup>c</sup> The device server shall cause the direct access block device to increase its tolerance of external physical forces (e.g., causes a device that has movable read/write heads to move those heads to a safe position).

<sup>d</sup> The device server shall cause the direct access block device to increase its tolerance of external physical forces and reduce its power consumption to use less power than when the logical unit is in the idle\_b power condition (e.g., cause a device that has rotating medium to rotate the medium at a lower revolutions per minute).

<sup>e</sup> If the specified timer is supported and enabled, then the device server shall:

- 1) force the specified timer to expire, which may cause the logical unit to transition to the specified power condition;
- 2) initialize and start all of the idle condition timers that are enabled (see SPC-5); and
- 3) initialize and start all of the standby condition timers that are enabled (see SPC-5), otherwise the device server shall terminate the START STOP UNIT command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

**Table 100 — POWER CONDITION and POWER CONDITION MODIFIER field (part 2 of 2)**

POWER CONDITION value	POWER CONDITION name	POWER CONDITION MODIFIER value	Device server action
Ah	FORCE_ IDLE_0	0h	Force the idle_a condition timer to be initialized to zero. <sup>e</sup>
		1h	Force the idle_b condition timer to be initialized to zero. <sup>e</sup>
		2h	Force the idle_c condition timer to be initialized to zero. <sup>e</sup>
Bh	FORCE_ STANDBY_0	0h	Force the standby_z condition timer to be initialized to zero. <sup>e</sup>
		1h	Force the standby_y condition timer to be initialized to zero. <sup>e</sup>
All other combinations			Reserved
<div><sup>a</sup> Process the following actions:<div><div>1) The device server shall comply with any SCSI transport protocol specific power condition transition restrictions (e.g., the NOTIFY (ENABLE SPINUP) requirement (see SPL-2));</div><div>2) the logical unit shall transition to the specified power condition; and</div><div>3) the device server shall disable all of the idle condition timers that are enabled (see SPC-5) and disable all of the standby condition timers that are enabled (see SPC-5) until another START STOP UNIT command is processed that returns control of the power condition to the logical unit or a logical unit reset occurs.</div></div><div><sup>b</sup> If a timer for a background scan operation expires, or a device specific event occurs, then the logical unit shall not leave this power condition to perform the background function associated with the timer or event.</div><div><sup>c</sup> The device server shall cause the direct access block device to increase its tolerance of external physical forces (e.g., causes a device that has movable read/write heads to move those heads to a safe position).</div><div><sup>d</sup> The device server shall cause the direct access block device to increase its tolerance of external physical forces and reduce its power consumption to use less power than when the logical unit is in the idle_b power condition (e.g., cause a device that has rotating medium to rotate the medium at a lower revolutions per minute).</div><div><sup>e</sup> If the specified timer is supported and enabled, then the device server shall:<div><div>1) force the specified timer to expire, which may cause the logical unit to transition to the specified power condition;</div><div>2) initialize and start all of the idle condition timers that are enabled (see SPC-5); and</div><div>3) initialize and start all of the standby condition timers that are enabled (see SPC-5),</div></div>otherwise the device server shall terminate the START STOP UNIT command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.</div></div>			

If the START STOP UNIT command specifies a power condition that conflicts with an operation in progress then, after the START STOP UNIT command completes with GOOD status, the logical unit may not be in the power condition that was requested by the command.

It is not an error to specify that the logical unit transition to its current power condition.

If no START STOP UNIT command is being processed by the device server, then the device server shall process any received START STOP UNIT command.

If a START STOP UNIT command is being processed by the device server and a subsequent START STOP UNIT command for which the CDB is validated requests that the logical unit change to the same power condition that was specified by the START STOP UNIT command being processed, then the device server shall process the subsequent command.

If the NO\_FLUSH bit is set to zero, then the device server shall write all logical block data in cache that is newer than logical block data on the medium to the medium (e.g., as if in response to a SYNCHRONIZE CACHE command (see 5.28 and 5.29) with the LOGICAL BLOCK ADDRESS field set to zero and the NUMBER OF LOGICAL

BLOCKS field set to zero) prior to entering into any power condition that prevents accessing the medium (e.g., before the rotating medium spindle motor is stopped during transition to the stopped power condition). If the NO\_FLUSH bit is set to one, then the device server should not write any cached logical blocks to the medium prior to entering into any power condition that prevents accessing the medium.

If the load eject (LOEJ) bit is set to zero and the POWER CONDITION field is set to zero, then the logical unit shall take no action regarding loading or ejecting the medium. If the LOEJ bit is set to one and the POWER CONDITION field is set to zero, then the logical unit shall unload the medium if the START bit is set to zero. If the LOEJ bit is set to one, the POWER CONDITION field is set to zero, and the START bit is set to one, then the logical unit shall load the medium.

If the START bit is set to zero and the POWER CONDITION field is set to zero, then the device server shall:

- a) cause the logical unit to transition to the stopped power condition;
- b) stop any idle condition timer that is enabled (see SPC-5); and
- c) stop any standby condition timer that is enabled (see SPC-5).

If the START bit set to one and the POWER CONDITION field is set to zero, then the device server shall:

- 1) comply with requirements defined in SCSI transport protocol standards (e.g., the NOTIFY (ENABLE SPINUP) requirement (see SPL-2));
- 2) cause the logical unit to transition to the active power condition;
- 3) initialize and start any idle condition timer that is enabled; and
- 4) initialize and start any standby condition timer that is enabled.

The CONTROL byte is defined in SAM-5.

## 5.27 STREAM CONTROL command

### 5.27.1 STREAM CONTROL command overview

The STREAM CONTROL command (see table 101) requests the device server to open a stream and return the stream identifier in the return parameter data or close the stream specified in the STR\_ID field in the CDB.

This command uses the SERVICE ACTION IN (16) CDB format (see A.2).

**Table 101 — STREAM CONTROL command**

Byte	Bit	7	6	5	4	3	2	1	0	
0	OPERATION CODE (9Eh)									
1	Reserved	STR_CTL			SERVICE ACTION (14h)					
2	Reserved									
3										
4	(MSB)	STR_ID							(LSB)	
5	Reserved									
6										
...										
9	ALLOCATION LENGTH									
10										(MSB)
...										
13										(LSB)
14	Reserved									
15	CONTROL									

The OPERATION CODE field and SERVICE ACTION field are defined in SPC-5 and shall be set to the values shown in table 101 for the STREAM CONTROL command.

The ALLOCATION LENGTH field is defined in SPC-5.

The stream control (STR\_CTL) field specifies the operation to be performed as described in table 102.

**Table 102 — STR\_CTL field**

Code	Description
01b	Open a stream and return the stream identifier in the ASSIGNED_STR_ID field in the returned parameter data.
10b	Close the stream associated with the STR_ID field.
all others	Reserved

If the STR\_CTL field is set to 10b, then the stream identifier (STR\_ID) field specifies the stream identifier associated with the requested operation. If the STR\_CTL field is not set to 10b, then the device server shall ignore the STR\_ID field.



### 5.27.2 STREAM CONTROL parameter data

The STREAM CONTROL parameter data is defined in table 103.

**Table 103 — STREAM CONTROL parameter data**

Byte	Bit	7	6	5	4	3	2	1	0
0		PARAMETER LENGTH (07h)							
1		Reserved							
...									
3									
4	(MSB)	ASSIGNED_STR_ID							
5		(LSB)							
6		Reserved							
7									

The PARAMETER LENGTH field indicates the length of the parameter data and shall be set as shown in table 103 for the STREAM CONTROL parameter data.

If the STR\_CTL field was set to 01b (i.e., close) in the STREAM CONTROL command, then the device server shall set the ASSIGNED\_STR\_ID field to a non zero value that is not currently assigned to an open stream by the device server and open that stream. If the STR\_CTL field was not set to 01b in the STREAM CONTROL command, then the ASSIGNED\_STR\_ID field is reserved.

### 5.28 SYNCHRONIZE CACHE (10) command

The SYNCHRONIZE CACHE (10) command (see table 104) requests that, for each logical block whose logical block data is in the volatile cache and has not already been written to the non-volatile cache, if any, or the medium, the device server either:

- perform a write medium operation to the LBA using the logical block data in volatile cache; or
- write the logical block to the non-volatile cache, if any.

NOTE 14 - Migration from the SYNCHRONIZE CACHE (10) command to the SYNCHRONIZE CACHE (16) command is recommended for all implementations.

**Table 104 — SYNCHRONIZE CACHE (10) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (35h)							
1		Reserved					Obsolete	IMMED	Obsolete
2	(MSB)	LOGICAL BLOCK ADDRESS							
...									
5									
6		Reserved		GROUP NUMBER					
7	(MSB)	NUMBER OF LOGICAL BLOCKS							
8									
9		CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 104 for the SYNCHRONIZE CACHE (10) command.

See the PRE-FETCH (10) command (see 5.10) for the definition of the LOGICAL BLOCK ADDRESS field.

See the PRE-FETCH (10) command (see 5.10) and 4.23 for the definition of the GROUP NUMBER field.

An immediate (IMMED) bit set to zero specifies that the device server shall not return status until the synchronize cache operation has been completed. An IMMED bit set to one specifies that the device server shall return status as soon as the CDB has been validated. If the IMMED bit is set to one, and the device server does not support the IMMED bit, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

If the IMMED bit is set to one and the synchronize cache operation has not completed, then the SYNC\_PROG field (see 6.5.6) defines the device server behavior while the synchronize cache command is being processed.

A NUMBER OF LOGICAL BLOCKS field set to a non-zero value specifies the number of logical blocks that shall be synchronized, starting with the logical block referenced by the LBA specified by the LOGICAL BLOCK ADDRESS field. A NUMBER OF LOGICAL BLOCKS field set to zero specifies that all logical blocks starting with the one referenced by the LBA specified in the LOGICAL BLOCK ADDRESS field to the last logical block on the medium shall be synchronized. If the specified LBA and the specified number of logical blocks exceed the capacity of the medium (see 4.5), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to LOGICAL BLOCK ADDRESS OUT OF RANGE.

A logical block within the range that is not in cache is not considered an error.

The CONTROL byte is defined in SAM-5.

## 5.29 SYNCHRONIZE CACHE (16) command

The SYNCHRONIZE CACHE (16) command (see table 105) requests that the device server perform the actions defined for the SYNCHRONIZE CACHE (10) command (see 5.28).

**Table 105 — SYNCHRONIZE CACHE (16) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (91h)							
1		Reserved					Obsolete	IMMED	Reserved
2		(MSB) _____ LOGICAL BLOCK ADDRESS _____ (LSB)							
...									
9									
10		(MSB) _____ NUMBER OF LOGICAL BLOCKS _____ (LSB)							
...									
13									
14		Reserved		GROUP NUMBER					
15		CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 105 for the SYNCHRONIZE CACHE (16) command.

See the SYNCHRONIZE CACHE (10) command (see 5.28) for the definitions of the other fields in this command.

## 5.30 UNMAP command

### 5.30.1 UNMAP command overview

The UNMAP command (see table 106) requests that the device server cause one or more LBAs to be unmapped (see 4.7.3).

**Table 106 — UNMAP command**

Bit	7	6	5	4	3	2	1	0
Byte								
0	OPERATION CODE (42h)							
1	Reserved							ANCHOR
2	Reserved							
...								
5								
6	Reserved	GROUP NUMBER						
7	(MSB)	PARAMETER LIST LENGTH						
8								(LSB)
9	CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 106 for the UNMAP command.

For a thin provisioned logical unit (see 4.7.3.3) with the ANC\_SUP bit set to one in the Logical Block

Provisioning VPD page (see 6.6.6):

- if the ANCHOR bit is set to zero, then any LBA on which an unmap operation (see 4.7.3.4) is performed shall either become deallocated (see 4.7.4.6) or anchored (see 4.7.4.7) and should become deallocated; and
- if the ANCHOR bit is set to one, then any LBA on which an unmap operation is performed shall become anchored.

For a thin provisioned logical unit (see 4.7.3.3) with the ANC\_SUP bit set to zero in the Logical Block

Provisioning VPD page:

- if the ANCHOR bit is set to zero, then any LBA on which an unmap operation is performed shall become deallocated; and
- if the ANCHOR bit is set to one, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

For a resource provisioned logical unit (see 4.7.3.2), the ANCHOR bit shall be ignored and any LBA on which an unmap operation is performed shall become anchored (i.e., the command is processed as if the ANCHOR bit is set to one).

For a thin provisioned logical unit or a resource provisioned logical unit any LBA on which an unmap operation is not performed does not change logical block provisioning state.

See the PRE-FETCH (10) command (see 5.10) and 4.23 for the definition of the GROUP NUMBER field.

The PARAMETER LIST LENGTH field specifies the length in bytes of the UNMAP parameter list that is available to be transferred from the Data-Out Buffer. If the parameter list length is greater than zero and less than eight, then the device server shall terminate the command with CHECK CONDITION status with the sense key set

to ILLEGAL REQUEST and the additional sense code set to PARAMETER LIST LENGTH ERROR. A PARAMETER LIST LENGTH set to zero specifies that no data shall be transferred.

The CONTROL byte is defined in SAM-5.

### 5.30.2 UNMAP parameter list

The UNMAP parameter list (see table 107) contains an UNMAP parameter list header and block descriptors. Each UNMAP block descriptor specifies a range of LBAs for which each LBA is processed as specified by the ANCHOR bit in the CDB (see 5.30.1). The LBAs specified in the block descriptors may contain overlapping LBA extents, and may be in any order.

c)

**Table 107 — UNMAP parameter list**

Byte	Bit	7	6	5	4	3	2	1	0	
0	(MSB)	UNMAP DATA LENGTH (n - 1)								(LSB)
1										
2	(MSB)	UNMAP BLOCK DESCRIPTOR DATA LENGTH (n - 7)								(LSB)
3										
4		Reserved								
...										
7										
UNMAP block descriptor list (if any)										
8		UNMAP block descriptor [first] (see table 108) (if any)								
...										
23										
⋮										
n - 15		UNMAP block descriptor [last] (see table 108) (if any)								
...										
n										

The UNMAP DATA LENGTH field specifies the length in bytes of the following data that is available to be transferred from the Data-Out Buffer. The unmap data length does not include the number of bytes in the UNMAP DATA LENGTH field.

The UNMAP BLOCK DESCRIPTOR DATA LENGTH field specifies the length in bytes of the UNMAP block descriptors that are available to be transferred from the Data-Out Buffer. The unmap block descriptor data length should be a multiple of 16. If the unmap block descriptor data length is not a multiple of 16, then the last unmap block descriptor is incomplete and shall be ignored. If the UNMAP BLOCK DESCRIPTOR DATA LENGTH is set to zero, then no unmap block descriptors are included in the UNMAP parameter list. This condition shall not be considered an error.

If any UNMAP block descriptors in the UNMAP block descriptor list are truncated due to the parameter list length in the CDB, then that UNMAP block descriptor shall be ignored.

Table 108 defines the UNMAP block descriptor.

**Table 108 — UNMAP block descriptor**

Byte	Bit	7	6	5	4	3	2	1	0
0	(MSB)	UNMAP LOGICAL BLOCK ADDRESS							
...									
7									
8	(MSB)	NUMBER OF LOGICAL BLOCKS							
...									
11									
12		Reserved							
...									
15									

The UNMAP LOGICAL BLOCK ADDRESS field specifies the first LBA that is requested to be unmapped (see 4.7.3.4.2) for this UNMAP block descriptor.

The NUMBER OF LOGICAL BLOCKS field specifies the number of LBAs that are requested to be unmapped beginning with the LBA specified by the UNMAP LOGICAL BLOCK ADDRESS field.

If the NUMBER OF LOGICAL BLOCKS is set to zero, then no LBAs shall be unmapped for this UNMAP block descriptor. This condition shall not be considered an error.

If the specified LBA and the specified number of logical blocks exceeds the capacity of the medium (see 4.5), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to LOGICAL BLOCK ADDRESS OUT OF RANGE.

If the total number of logical blocks specified in the UNMAP block descriptor data exceeds the value indicated in the MAXIMUM UNMAP LBA COUNT field (see 6.6.4) or if the number of UNMAP block descriptors exceeds the value of the MAXIMUM UNMAP BLOCK DESCRIPTOR COUNT field (see 6.6.4), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

### 5.31 VERIFY (10) command

The VERIFY (10) command (see table 110) requests that the device server:

- 1) perform verify operations from the specified LBAs; and
- 2) if specified, perform a compare operation on:
  - A) the logical block data transferred from the Data-Out Buffer; and
  - B) the logical block data from the verify operations.

The device server may process the LBAs in any order but shall perform this sequence in the specified order for a given LBA.

The application client uses the BYTCHK field in the CDB to specify the contents of the Data-Out Buffer as shown in table 109.

**Table 109 — Data-Out Buffer contents for the VERIFY (10) command**

BYTCHK field	Data-Out Buffer contents
00b	Not used
01b	Logical block data for the number of logical blocks specified in the VERIFICATION LENGTH field
10b <sup>a</sup>	Not defined
11b	Logical block data for a single logical block
<sup>a</sup> A BYTCHK field set to 10b is reserved.	

NOTE 15 - Migration from the VERIFY (10) command to the VERIFY 16) command is recommended for all implementations.

**Table 110 — VERIFY (10) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (2Fh)							
1		VRPROTECT			DPO	Reserved	BYTCHK		Obsolete
2	(MSB)	LOGICAL BLOCK ADDRESS							
...									
5									
6	Restricted for MMC-6	Reserved	GROUP NUMBER						
7	(MSB)	VERIFICATION LENGTH							
8									
9		CONTROL							

The OPERATION CODE field is defined in ~~SPC-4~~SPC-5 and shall be set to the value shown in table 110 for the VERIFY (10) command.

See the READ (10) command (see 5.13) for the definition of the DPO bit. See the PRE-FETCH (10) command (see 5.10) for the definition of the LOGICAL BLOCK ADDRESS field. See the PRE-FETCH (10) command (see 5.10) and 4.23 for the definition of the GROUP NUMBER field.

If the byte check (BYTCHK) field is set to 00b, then:

- no Data-Out Buffer transfer shall occur;
- for any mapped LBA specified by the command, the device server shall check the protection information from the verify operation based on the VRPROTECT field as defined in table 111; and
- for any unmapped LBA specified by the command, the verify operation shall complete without error.

If:

- the BYTCHK field is set to 01b or 11b;

- b) the VBULS bit is set to zero in the Block Device Characteristics VPD page (see 6.6.2); and
- c) any LBA specified by the command is unmapped (i.e., deallocated or anchored),

then the device server shall terminate the command with CHECK CONDITION status with the sense key set to MISCOMPARE and the additional sense code set to MISCOMPARE VERIFY OF UNMAPPED LBA.

If:

- a) the BYTCHK field is set to 01b or 11b; and
- b) either:
  - A) the VBULS bit is set to one in the Block Device Characteristics VPD page; or
  - B) all LBAs specified by the command are mapped,

then:

- a) if the BYTCHK field is set to 01b, then the Data-Out Buffer transfer shall include the number of logical blocks specified by the VERIFICATION LENGTH field;
  - b) if the BYTCHK field is set to 11b, then:
    - A) the Data-Out Buffer transfer shall include one logical block; and
    - B) the device server shall:
      - 1) duplicate the single logical block, as described in the WRITE SAME command (see 5.47), the number of times required to satisfy the VERIFICATION LENGTH field; and
      - 2) place the duplicated data in the Data-Out Buffer;
  - c) the device server shall check the protection information transferred from the Data-Out Buffer based on the VRPROTECT field as defined in table 113;
  - d) for any mapped LBA specified by the command, the device server shall perform the verify operation and check the protection information from the verify operation based on the VRPROTECT field as defined in table 112;
- and
- e) the device server shall perform:
    - A) a compare operation of:
      - a) user data from the verify operations; and
      - b) user data from the Data-Out Buffer;
    - and
    - B) a compare operation based on the VRPROTECT field as defined in table 114 of:
      - a) protection information from the verify operations; and
      - b) protection information from the Data-Out Buffer.

The order of the user data and protection information checks and compare operations is vendor specific.

If a compare operation indicates a miscompare, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to MISCOMPARE and the additional sense code set to the appropriate value for the condition.

The VERIFICATION LENGTH field specifies the number of contiguous logical blocks that shall be verified, starting with the logical block referenced by the LBA specified by the LOGICAL BLOCK ADDRESS field. A VERIFICATION LENGTH field set to zero specifies that no logical blocks shall be transferred or verified. This condition shall not be considered an error. If the specified LBA and the specified verification length exceed the capacity of the medium (see 4.5), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to LOGICAL BLOCK ADDRESS OUT OF RANGE. The VERIFICATION LENGTH field is constrained by the MAXIMUM TRANSFER LENGTH field ~~in the Block Limits VPD page~~ (see 6.6.4).

If the BYTCHK field is set to 01b, then the VERIFICATION LENGTH field also specifies the number of logical blocks that the device server shall transfer from the Data-Out Buffer.

The CONTROL byte is defined in SAM-5.



If the BYTCHK field is set to 00b, then table 111 defines the checks that the device server shall perform on the protection information from the verify operations based on the VRPROTECT field. All footnotes for table 111 are at the end of the table.

**Table 111 — VRPROTECT field with the BYTCHK field set to 00b – checking protection information from the verify operations (part 1 of 3)**

Code	Logical unit formatted with protection information	Field in protection information <sup>g</sup>	Extended INQUIRY Data VPD page bit value <sup>f</sup>	If check fails <sup>d e</sup> , additional sense code
000b	Yes <sup>i</sup>	LOGICAL BLOCK GUARD	GRD_CHK = 1	LOGICAL BLOCK GUARD CHECK FAILED
			GRD_CHK = 0	No check performed
		LOGICAL BLOCK APPLICATION TAG	APP_CHK = 1 <sup>c</sup>	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
			APP_CHK = 0	No check performed
		LOGICAL BLOCK REFERENCE TAG	REF_CHK = 1 <sup>h</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
			REF_CHK = 0	No check performed
	No	No protection information on the medium to check.		
001b 101b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	GRD_CHK = 1	LOGICAL BLOCK GUARD CHECK FAILED
			GRD_CHK = 0	No check performed
		LOGICAL BLOCK APPLICATION TAG	APP_CHK = 1 <sup>c</sup>	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
			APP_CHK = 0	No check performed
		LOGICAL BLOCK REFERENCE TAG	REF_CHK = 1 <sup>h</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
			REF_CHK = 0	No check performed
	No	Error condition <sup>a</sup>		
010b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	No check performed	
		LOGICAL BLOCK APPLICATION TAG	APP_CHK = 1 <sup>c</sup>	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
			APP_CHK = 0	No check performed
		LOGICAL BLOCK REFERENCE TAG	REF_CHK = 1 <sup>h</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
			REF_CHK = 0	No check performed
	No	Error condition <sup>a</sup>		

**Table 111 — VRPROTECT field with the BYTCHK field set to 00b – checking protection information from the verify operations (part 2 of 3)**

Code	Logical unit formatted with protection information	Field in protection information <sup>g</sup>	Extended INQUIRY Data VPD page bit value <sup>f</sup>	If check fails <sup>d e</sup> , additional sense code
011b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	No check performed	
		LOGICAL BLOCK APPLICATION TAG	No check performed	
		LOGICAL BLOCK REFERENCE TAG	No check performed	
	No	Error condition <sup>a</sup>		
100b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	GRD_CHK = 1	LOGICAL BLOCK GUARD CHECK FAILED
			GRD_CHK = 0	No check performed
		LOGICAL BLOCK APPLICATION TAG	No check performed	
		LOGICAL BLOCK REFERENCE TAG	No check performed	
	No	Error condition <sup>a</sup>		
110b to 111b	Reserved			

**Table 111 — VRPROTECT field with the BYTCHK field set to 00b – checking protection information from the verify operations (part 3 of 3)**

Code	Logical unit formatted with protection information	Field in protection information <sup>g</sup>	Extended INQUIRY Data VPD page bit value <sup>f</sup>	If check fails <sup>d e</sup> , additional sense code
<p><sup>a</sup> If the logical unit supports protection information (see 4.22) and has not been formatted with protection information, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.</p> <p><sup>b</sup> If the logical unit does not support protection information, then the device server should terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.</p> <p><sup>c</sup> If the device server has knowledge of the contents of the LOGICAL BLOCK APPLICATION TAG field, then the device server shall check each logical block application tag. If the ATO bit in the Control mode page (see SPC-5) is set to one, then this knowledge is acquired from:</p> <ul style="list-style-type: none"> <li>a) the EXPECTED LOGICAL BLOCK APPLICATION TAG field and the LOGICAL BLOCK APPLICATION TAG MASK field in the CDB, if the command is a VERIFY (32) command (see 5.34);</li> <li>b) the Application Tag mode page (see 6.5.3), if the command is not a VERIFY (32) command and the ATMPE bit in the Control mode page (see SPC-5) is set to one; or</li> <li>c) a method not defined by this standard, if the command is not a VERIFY (32) command and the ATMPE bit is set to zero.</li> </ul> <p><sup>d</sup> If the device server terminates the command with CHECK CONDITION status, then the device server shall set the sense key to ABORTED COMMAND.</p> <p><sup>e</sup> If multiple errors occur while the device server is processing the command, then the selection by the device server of which error to report is not defined by this standard.</p> <p><sup>f</sup> See the Extended INQUIRY Data VPD page (see SPC-5) for the definitions of the GRD_CHK bit, the APP_CHK bit, and the REF_CHK bit.</p> <p><sup>g</sup> If the device server detects:</p> <ul style="list-style-type: none"> <li>a) a LOGICAL BLOCK APPLICATION TAG field set to FFFFh and type 1 protection (see 4.22.2.3) or type 2 protection (see 4.22.2.4) is enabled; or</li> <li>b) a LOGICAL BLOCK APPLICATION TAG field set to FFFFh, LOGICAL BLOCK REFERENCE TAG field set to FFFF_FFFFh, and type 3 protection (see 4.22.2.5) is enabled,</li> </ul> <p>then the device server shall not check any protection information in the associated protection information interval.</p> <p><sup>h</sup> If type 1 protection is enabled, then the device server shall check each logical block reference tag by comparing it to the lower four bytes of the LBA associated with the logical block. If type 2 protection or type 3 protection is enabled, and the device server has knowledge of the contents of the LOGICAL BLOCK REFERENCE TAG field, then the device server shall check the logical block reference tag. If type 2 protection is enabled, then this knowledge may be acquired through the EXPECTED INITIAL LOGICAL BLOCK REFERENCE TAG field in a VERIFY (32) command (see 5.34). If type 3 protection is enabled, then the method for acquiring this knowledge is not defined by this standard.</p> <p><sup>i</sup> If the DPICZ bit in the Control mode page (see SPC-5) is set to one, then protection information shall not be checked.</p>				

If the BYTCHK field is set to 01b or 11b, then table 112 defines the checks that the device server shall perform on the protection information from the verify operations based on the VRPROTECT field. All footnotes for table 112 are at the end of the table.

**Table 112 — VRPROTECT field with the BYTCHK field set to 01b or 11b – checking protection information from the verify operations (part 1 of 2)**

Code	Logical unit formatted with protection information	Field in protection information <sup>g</sup>	Extended INQUIRY Data VPD page bit value <sup>f</sup>	If check fails <sup>d e</sup> , additional sense code
000b	Yes <sup>i</sup>	LOGICAL BLOCK GUARD	GRD_CHK = 1	LOGICAL BLOCK GUARD CHECK FAILED
			GRD_CHK = 0	No check performed
		LOGICAL BLOCK APPLICATION TAG	APP_CHK = 1 <sup>c g</sup>	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
			APP_CHK = 0	No check performed
		LOGICAL BLOCK REFERENCE TAG	REF_CHK = 1 <sup>h</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
			REF_CHK = 0	No check performed
	No	No protection information available to check		
001b 010b 011b 100b 101b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	No check performed	
		LOGICAL BLOCK APPLICATION TAG	No check performed	
		LOGICAL BLOCK REFERENCE TAG	No check performed	
	No	Error condition <sup>a</sup>		
110b to 111b	Reserved			

**Table 112 — VRPROTECT field with the BYTCHK field set to 01b or 11b – checking protection information from the verify operations (part 2 of 2)**

Code	Logical unit formatted with protection information	Field in protection information <sup>g</sup>	Extended INQUIRY Data VPD page bit value <sup>f</sup>	If check fails <sup>d e</sup> , additional sense code
<p><sup>a</sup> If the logical unit supports protection information (see 4.22) and has not been formatted with protection information, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.</p> <p><sup>b</sup> If the logical unit does not support protection information, then the device server should terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.</p> <p><sup>c</sup> If the device server has knowledge of the contents of the LOGICAL BLOCK APPLICATION TAG field, then the device server shall check each logical block application tag. If the ATO bit in the Control mode page (see SPC-5) is set to one, then this knowledge is acquired from:</p> <ul style="list-style-type: none"> <li>a) the EXPECTED LOGICAL BLOCK APPLICATION TAG field and the LOGICAL BLOCK APPLICATION TAG MASK field in the CDB, if the command is a VERIFY (32) command (see 5.34);</li> <li>b) the Application Tag mode page (see 6.5.3), if the command is not a VERIFY (32) command and the ATMPE bit in the Control mode page (see SPC-5) is set to one; or</li> <li>c) a method not defined by this standard, if the command is not a VERIFY (32) command and the ATMPE bit is set to zero.</li> </ul> <p><sup>d</sup> If the device server terminates the command with CHECK CONDITION status, then the device server shall set the sense key to ABORTED COMMAND.</p> <p><sup>e</sup> If multiple errors occur while the device server is processing the command, then the selection by the device server of which error to report is not defined by this standard.</p> <p><sup>f</sup> See the Extended INQUIRY Data VPD page (see SPC-5) for the definitions of the GRD_CHK bit, the APP_CHK bit, and the REF_CHK bit.</p> <p><sup>g</sup> If the device server detects:</p> <ul style="list-style-type: none"> <li>a) a LOGICAL BLOCK APPLICATION TAG field set to FFFFh and type 1 protection (see 4.22.2.3) or type 2 protection (see 4.22.2.4) is enabled; or</li> <li>b) a LOGICAL BLOCK APPLICATION TAG field set to FFFFh, LOGICAL BLOCK REFERENCE TAG field set to FFFF_FFFFh, and type 3 protection (see 4.22.2.5) is enabled,</li> </ul> <p>then the device server shall not check any protection information in the associated protection information interval.</p> <p><sup>h</sup> If type 1 protection is enabled, then the device server shall check each logical block reference tag by comparing it to the lower four bytes of the LBA associated with the logical block. If type 2 protection or type 3 protection is enabled, and the device server has knowledge of the contents of the LOGICAL BLOCK REFERENCE TAG field, then the device server shall check the logical block reference tag. If type 2 protection is enabled, then this knowledge may be acquired through the EXPECTED INITIAL LOGICAL BLOCK REFERENCE TAG field in a VERIFY (32) command (see 5.34). If type 3 protection is enabled, then the method for acquiring this knowledge is not defined by this standard.</p> <p><sup>i</sup> If the DPICZ bit in the Control mode page (see SPC-5) is set to one, then protection information shall not be checked.</p>				

If the BYTCHK field is set to 01b or 11b, then table 113 defines the checks that the device server shall perform on the protection information transferred from the Data-Out Buffer based on the VRPROTECT field. All footnotes for table 113 are at the end of the table.

**Table 113 — VRPROTECT field with the BYTCHK field set to 01b or 11b – checking protection information from the Data-Out Buffer (part 1 of 2)**

Code	Logical unit formatted with protection information	Field in protection information <sup>g</sup>	Device server check	If check fails <sup>d e</sup> , additional sense code
000b	Yes	No protection information in the Data-Out Buffer to check		
	No	No protection information in the Data-Out Buffer to check		
001b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	Shall	LOGICAL BLOCK GUARD CHECK FAILED
		LOGICAL BLOCK APPLICATION TAG	May <sup>c</sup>	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
		LOGICAL BLOCK REFERENCE TAG	Shall (except for type 3) <sup>f</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
	No	Error condition <sup>a</sup>		
010b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	Shall not	No check performed
		LOGICAL BLOCK APPLICATION TAG	May <sup>c</sup>	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
		LOGICAL BLOCK REFERENCE TAG	May <sup>f</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
	No	Error condition <sup>a</sup>		
011b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	Shall not	No check performed
		LOGICAL BLOCK APPLICATION TAG	Shall not	No check performed
		LOGICAL BLOCK REFERENCE TAG	Shall not	No check performed
	No	Error condition <sup>a</sup>		
100b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	Shall	LOGICAL BLOCK GUARD CHECK FAILED
		LOGICAL BLOCK APPLICATION TAG	Shall not	No check performed
		LOGICAL BLOCK REFERENCE TAG	Shall not	No check performed
	No	Error condition <sup>a</sup>		

**Table 113 — VRPROTECT field with the BYTCHK field set to 01b or 11b – checking protection information from the Data-Out Buffer (part 2 of 2)**

Code	Logical unit formatted with protection information	Field in protection information <sup>g</sup>	Device server check	If check fails <sup>d e</sup> , additional sense code
101b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	Shall	LOGICAL BLOCK GUARD CHECK FAILED
		LOGICAL BLOCK APPLICATION TAG	May <sup>c</sup>	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
		LOGICAL BLOCK REFERENCE TAG	May <sup>f</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
	No	Error condition <sup>a</sup>		
110b to 111b	Reserved			

<sup>a</sup> If the logical unit supports protection information (see 4.22) and has not been formatted with protection information, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

<sup>b</sup> If the logical unit does not support protection information, then the device server should terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

<sup>c</sup> If the device server has knowledge of the contents of the LOGICAL BLOCK APPLICATION TAG field and the ATO bit is set to one in the Control mode page (see SPC-5), then the device server may check each logical block application tag. If the ATO bit is set to one, then this knowledge is acquired from:

a) the EXPECTED LOGICAL BLOCK APPLICATION TAG field and the LOGICAL BLOCK APPLICATION TAG MASK field in the CDB, if the command is a VERIFY (32) command;

b) the Application Tag mode page (see 6.5.3), if the command is not a VERIFY (32) command and the ATMPE bit in the Control mode page (see SPC-5) is set to one; or

c) a method not defined by this standard, if the command is not a VERIFY (32) command and the ATMPE bit is set to zero.

<sup>d</sup> If the device server terminates the command with CHECK CONDITION status, then the device server shall set the sense key to ABORTED COMMAND.

<sup>e</sup> If multiple errors occur while the device server is processing the command, then the selection by the device server of which error to report is not defined by this standard.

<sup>f</sup> If type 1 protection is enabled, then the device server shall check the logical block reference tag by comparing it to the lower four bytes of the LBA associated with the logical block. If type 2 protection is enabled and the device server has knowledge of the contents of the LOGICAL BLOCK REFERENCE TAG field, then the device server shall check each logical block reference tag. If type 2 protection is enabled, then this knowledge may be acquired through the EXPECTED INITIAL LOGICAL BLOCK REFERENCE TAG field in a VERIFY (32) command (see 5.34). If type 3 protection is enabled, the ATO bit is set to one in the Control mode page (see SPC-5), and the device server has knowledge of the contents of the LOGICAL BLOCK REFERENCE TAG field, then the device server may check each logical block reference tag. If type 3 protection is enabled, then the method for acquiring this knowledge is not defined by this standard.

<sup>g</sup> If the NO\_PI\_CHK bit is set to one in the Extended INQUIRY Data VPD page (see SPC-5) and the device server detects:

a) a LOGICAL BLOCK APPLICATION TAG field set to FFFFh and type 1 protection (see 4.22.2.3) or type 2 protection (see 4.22.2.4) is enabled; or

b) a LOGICAL BLOCK APPLICATION TAG field set to FFFFh, LOGICAL BLOCK REFERENCE TAG field set to FFFF\_FFFFh, and type 3 protection (see 4.22.2.5) is enabled, then the device server shall not check any protection information in the associated protection information interval.

If the BYTCHK field is set to 01b or 11b, then table 114 defines the processing by the device server of the protection information during the compare operation based on the VRPROTECT field. All footnotes for table 114 are at the end of the table.

**Table 114 — VRPROTECT field with the BYTCHK field set to 01b or 11b – compare operation requirements**  
(part 1 of 3)

Code	Logical unit formatted with protection information	Field in protection information	Compare operation	If compare fails <sup>c d</sup> , additional sense code
000b	Yes	No protection information in the Data-Out Buffer to compare. Only user data is compared within each logical block.		
	No	No protection information from the verify operations or in the Data-Out Buffer to compare. Only user data is compared within each logical block.		
001b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	Shall	LOGICAL BLOCK GUARD CHECK FAILED
		LOGICAL BLOCK APPLICATION TAG (ATO = 1) <sup>e</sup>	Shall	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
		LOGICAL BLOCK APPLICATION TAG (ATO = 0) <sup>f</sup>	Shall not	No compare performed
		LOGICAL BLOCK REFERENCE TAG (not type 3)	Shall <sup>g</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
		LOGICAL BLOCK REFERENCE TAG (type 3 and ATO = 1)	Shall <sup>h</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
		LOGICAL BLOCK REFERENCE TAG (type 3 and ATO = 0)	Shall not	No compare performed
	No	Error condition <sup>a</sup>		



**Table 114 — VRPROTECT field with the BYTCHK field set to 01b or 11b – compare operation requirements**  
(part 2 of 3)

Code	Logical unit formatted with protection information	Field in protection information	Compare operation	If compare fails <sup>c d</sup> , additional sense code
010b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	Shall not	No compare performed
		LOGICAL BLOCK APPLICATION TAG (ATO = 1) <sup>e</sup>	Shall	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
		LOGICAL BLOCK APPLICATION TAG (ATO = 0) <sup>f</sup>	Shall not	No compare performed
		LOGICAL BLOCK REFERENCE TAG (not type 3)	Shall <sup>g</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
		LOGICAL BLOCK REFERENCE TAG (type 3 and ATO = 1)	Shall <sup>h</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
		LOGICAL BLOCK REFERENCE TAG (type 3 and ATO = 0)	Shall not	No compare performed
	No	Error condition <sup>a</sup>		
011b 100b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	Shall	LOGICAL BLOCK GUARD CHECK FAILED
		LOGICAL BLOCK APPLICATION TAG (ATO = 1) <sup>e</sup>	Shall	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
		LOGICAL BLOCK APPLICATION TAG (ATO = 0) <sup>f</sup>	Shall not	No compare performed
		LOGICAL BLOCK REFERENCE TAG (not type 3)	Shall <sup>g</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
		LOGICAL BLOCK REFERENCE TAG (type 3 and ATO = 1)	Shall <sup>h</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
		LOGICAL BLOCK REFERENCE TAG (type 3 and ATO = 0)	Shall not	No compare performed
	No	Error condition <sup>a</sup>		

**Table 114 — VRPROTECT field with the BYTCHK field set to 01b or 11b – compare operation requirements**  
(part 3 of 3)

Code	Logical unit formatted with protection information	Field in protection information	Compare operation	If compare fails <sup>c d</sup> , additional sense code
101b <sup>b</sup>	Yes	LOGICAL BLOCK GUARD	Shall	LOGICAL BLOCK GUARD CHECK FAILED
		LOGICAL BLOCK APPLICATION TAG (ATO = 1) <sup>e</sup>	Shall	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
		LOGICAL BLOCK APPLICATION TAG (ATO = 0) <sup>f</sup>	Shall not	No compare performed
		LOGICAL BLOCK REFERENCE TAG	Shall not	No compare performed
	No	Error condition <sup>a</sup>		
110b to 111b	Reserved			

<sup>a</sup> If the logical unit supports protection information (see 4.22) and has not been formatted with protection information, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

<sup>b</sup> If the logical unit does not support protection information, then the device server should terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

<sup>c</sup> If the device server terminates the command with CHECK CONDITION status, then the device server shall set the sense key to MISCOMPARE.

<sup>d</sup> If multiple errors occur while the device server is processing the command, then the selection by the device server of which error to report is not defined by this standard.

<sup>e</sup> If the ATO bit is set to one in the Control mode page (see SPC-5), then the device server shall not modify the logical block application tag.

<sup>f</sup> If the ATO bit is set to zero in the Control mode page (see SPC-5), then the device server may modify any logical block application tag.

<sup>g</sup> If the BYTCHK field is set to 11b, then the device server shall compare the value from each LOGICAL BLOCK REFERENCE TAG field received in the single logical block data from the Data-Out Buffer with the corresponding LOGICAL BLOCK REFERENCE TAG field in the first logical block from the verify operations, and the device server shall compare the value of the previous LOGICAL BLOCK REFERENCE TAG field plus one with each of the subsequent LOGICAL BLOCK REFERENCE TAG fields (see 4.22.3).

<sup>h</sup> If the BYTCHK field is set to 11b, then the device server shall compare the value from each LOGICAL BLOCK REFERENCE TAG field received in the single logical block data from the Data-Out Buffer with the corresponding LOGICAL BLOCK REFERENCE TAG field in each logical block from the verify operations (see 4.22.3).

### 5.32 VERIFY (12) command

The VERIFY (12) command (see table 115) requests that the device server perform the actions defined for the VERIFY (10) command (see 5.31).

NOTE 16 - Migration from the VERIFY (12) command to the VERIFY (16) command is recommended for all implementations.

**Table 115 — VERIFY (12) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (AFh)							
1		VRPROTECT			DPO	Reserved	BYTCHK		Obsolete
2	(MSB)	LOGICAL BLOCK ADDRESS							
...									
5									
6	(MSB)	VERIFICATION LENGTH							
...									
9									
10		Reserved		GROUP NUMBER					
11		CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 115 for the VERIFY (12) command.

See the VERIFY (10) command (see 5.31) for the definitions of the other fields in this command.

### 5.33 VERIFY (16) command

The VERIFY (16) command (see table 116) requests that the device server perform the actions defined for the VERIFY (10) command (see 5.31).

**Table 116 — VERIFY (16) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (8Fh)							
1		VRPROTECT			DPO	Reserved	BYTCHK		Reserved
2	(MSB)	LOGICAL BLOCK ADDRESS							
...									
9									
10	(MSB)	VERIFICATION LENGTH							
...									
13									
14		Reserved		GROUP NUMBER					
15		CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 116 for the VERIFY (16) command.

See the VERIFY (10) command (see 5.31) for the definitions of the other fields in this command.

### 5.34 VERIFY (32) command

The VERIFY (32) command (see table 117) requests that the device server perform the actions defined for the VERIFY (10) command (see 5.31).

The device server shall process a VERIFY (32) command only if type 2 protection is enabled (see 4.22.2.4).

**Table 117 — VERIFY (32) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (7Fh)							
1		CONTROL							
2		Reserved							
...									
5									
6		Reserved	GROUP NUMBER						
7		ADDITIONAL CDB LENGTH (18h)							
8	(MSB)	SERVICE ACTION (000Ah)							
9									
10		VRPROTECT			DPO	Reserved	BYTCHK		Reserved
11		Reserved							
12	(MSB)	LOGICAL BLOCK ADDRESS							
...									
19		(LSB)							
20	(MSB)	EXPECTED INITIAL LOGICAL BLOCK REFERENCE TAG							
...									
23		(LSB)							
24	(MSB)	EXPECTED LOGICAL BLOCK APPLICATION TAG							
25									
26	(MSB)	LOGICAL BLOCK APPLICATION TAG MASK							
27									
28	(MSB)	VERIFICATION LENGTH							
...									
31		(LSB)							

The OPERATION CODE field, the ADDITIONAL CDB LENGTH field, and the SERVICE ACTION field are defined in SPC-5 and shall be set to the values shown in table 117 for the VERIFY (32) command.

See the VERIFY (10) command (see 5.31) for the definitions of the CONTROL byte, the GROUP NUMBER field, the VRPROTECT field, the DPO bit, the BYTCHK field, the LOGICAL BLOCK ADDRESS field, and the VERIFICATION LENGTH field.

If checking of the LOGICAL BLOCK REFERENCE TAG field is enabled (see table 111, table 112, table 113, and table 114 in 5.31), then the EXPECTED INITIAL LOGICAL BLOCK REFERENCE TAG field contains the value of the LOGICAL BLOCK REFERENCE TAG field expected in the protection information of the first logical block accessed by the command instead of a value based on the LBA (see 4.22.3).

If the ATO bit is set to one in the Control mode page (see SPC-5) and checking of the LOGICAL BLOCK APPLICATION TAG field is enabled (see table 111, table 112, table 113, and table 114 in 5.31), then the LOGICAL BLOCK APPLICATION TAG MASK field contains a value that is a bit mask for enabling the checking of the LOGICAL BLOCK APPLICATION TAG field in every instance of the protection information for each logical block accessed by the command. A LOGICAL BLOCK APPLICATION TAG MASK bit set to one enables the checking of the corresponding bit of the EXPECTED LOGICAL BLOCK APPLICATION TAG field with the corresponding bit of the LOGICAL BLOCK APPLICATION TAG field in every instance of the protection information. A LOGICAL BLOCK APPLICATION TAG MASK field bit set to zero disables the checking of the corresponding bit of the EXPECTED LOGICAL BLOCK APPLICATION TAG field with the corresponding bit of the LOGICAL BLOCK APPLICATION TAG field in every instance of protection information.

The LOGICAL BLOCK APPLICATION TAG MASK field and the EXPECTED LOGICAL BLOCK APPLICATION TAG field shall be ignored if:

- a) the ATO bit is set to zero; or
- b) the ATO bit is set to one in the Control mode page (see SPC-5) and checking of the LOGICAL BLOCK APPLICATION TAG field is disabled (see table 111, table 112, table 113, and table 114 in 5.31).

### 5.35 WRITE (10) command

The WRITE (10) command (see table 118) requests that the device server:

- a) transfer the specified logical block data from the Data-Out Buffer; and
- b) perform write operations to the specified LBAs using the transferred logical blocks.

NOTE 17 - Migration from the WRITE (10) command to the WRITE (16) command is recommended for all implementations.

**Table 118 — WRITE (10) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (2Ah)							
1		WRPROTECT			DPO	FUA	Reserved	Obsolete	Obsolete
2	(MSB)	LOGICAL BLOCK ADDRESS							
...									
5		(LSB)							
6		Reserved		GROUP NUMBER					
7	(MSB)	TRANSFER LENGTH							
8									
9		CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 118 for the WRITE (10) command.

The device server shall check the protection information, if any, transferred from the Data-Out Buffer based on the WRPROTECT field as described in table 119. All footnotes for table 119 are at the end of the table.

**Table 119 — WRPROTECT field (part 1 of 3)**

Code	Logical unit formatted with protection information	Field in protection information <sup>k</sup>	Device server check	If check fails <sup>d i</sup> , additional sense code
000b	Yes <sup>f g h</sup>	No protection information received from application client to check		
	No	No protection information received from application client to check		
001b <sup>b</sup>	Yes <sup>e</sup>	LOGICAL BLOCK GUARD	Shall	LOGICAL BLOCK GUARD CHECK FAILED
		LOGICAL BLOCK APPLICATION TAG	Dependent on RWWP <sup>c</sup>	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
		LOGICAL BLOCK REFERENCE TAG	Shall (except for type 3) <sup>j</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
	No <sup>a</sup>	No protection information available to check		
010b <sup>b</sup>	Yes <sup>e</sup>	LOGICAL BLOCK GUARD	Shall not	No check performed
		LOGICAL BLOCK APPLICATION TAG	Dependent on RWWP <sup>c</sup>	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
		LOGICAL BLOCK REFERENCE TAG	May <sup>j</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
	No <sup>a</sup>	No protection information available to check		
011b <sup>b</sup>	Yes <sup>e</sup>	LOGICAL BLOCK GUARD	Shall not	No check performed
		LOGICAL BLOCK APPLICATION TAG	Shall not	No check performed
		LOGICAL BLOCK REFERENCE TAG	Shall not	No check performed
	No <sup>a</sup>	No protection information available to check		
100b <sup>b</sup>	Yes <sup>e</sup>	LOGICAL BLOCK GUARD	Shall	LOGICAL BLOCK GUARD CHECK FAILED
		LOGICAL BLOCK APPLICATION TAG	Shall not	No check performed
		LOGICAL BLOCK REFERENCE TAG	Shall not	No check performed
	No <sup>a</sup>	No protection information available to check		

Table 119 — WRPROTECT field (part 2 of 3)

Code	Logical unit formatted with protection information	Field in protection information <sup>k</sup>	Device server check	If check fails <sup>d i</sup> , additional sense code
101b <sup>b</sup>	Yes <sup>e</sup>	LOGICAL BLOCK GUARD	Shall	LOGICAL BLOCK GUARD CHECK FAILED
		LOGICAL BLOCK APPLICATION TAG	Dependent on RWWP <sup>c</sup>	LOGICAL BLOCK APPLICATION TAG CHECK FAILED
		LOGICAL BLOCK REFERENCE TAG	May <sup>j</sup>	LOGICAL BLOCK REFERENCE TAG CHECK FAILED
	No <sup>a</sup>	No protection information available to check		
110b to 111b	Reserved			



Table 119 — WRPROTECT field (part 3 of 3)

Code	Logical unit formatted with protection information	Field in protection information <sup>k</sup>	Device server check	If check fails <sup>d i</sup> , additional sense code
<p><sup>a</sup> If a logical unit supports protection information (see 4.22) and has not been formatted with protection information, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.</p> <p><sup>b</sup> If the logical unit does not support protection information, then the device server should terminate the requested command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.</p> <p><sup>c</sup> If the device server has knowledge of the contents of the LOGICAL BLOCK APPLICATION TAG field and the ATO bit is set to one in the Control mode page (see SPC-5), then the device server:</p> <p>a) may check each logical block application tag if the RWWP bit is set to zero in the Control mode page (see SPC-5); and</p> <p>b) shall check each logical block application tag if the RWWP bit is set to one in the Control mode page. If the ATO bit in the Control mode page (see SPC-5) is set to one, then this knowledge is acquired from:</p> <p>a) the EXPECTED LOGICAL BLOCK APPLICATION TAG field and the LOGICAL BLOCK APPLICATION TAG MASK field in the CDB, if a WRITE (32) command (see 5.38), a WRITE ATOMIC (32) command (see 5.44), a WRITE SAME (32) command (see 5.49), or a WRITE STREAM (32) command (see 5.51) is received by the device server;</p> <p>b) the Application Tag mode page (see 6.5.3), if a command other than WRITE (32), WRITE ATOMIC (32), WRITE SAME (32), or WRITE STREAM (32) is received by the device server and the ATMPE bit in the Control mode page (see SPC-5) is set to one; or</p> <p>c) a method not defined by this standard, if a command other than WRITE (32), WRITE ATOMIC (32), WRITE SAME (32), or WRITE STREAM (32) is received by the device server, and the ATMPE bit is set to zero.</p> <p><sup>d</sup> If the device server terminates the command with CHECK CONDITION status, then the device server shall set the sense key to ABORTED COMMAND.</p> <p><sup>e</sup> The device server shall preserve the contents of protection information (e.g., write it to the medium or store it in non-volatile memory).</p> <p><sup>f</sup> The device server shall write a generated CRC (see 4.22.4.2) into each LOGICAL BLOCK GUARD field.</p> <p><sup>g</sup> If the RWWP bit in the Control mode page (see SPC-5) is set to one, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB. If the RWWP bit is set to zero and:</p> <p>a) type 1 protection is enabled, then the device server shall write the least significant four bytes of each LBA into the LOGICAL BLOCK REFERENCE TAG field of each of the written logical blocks; or</p> <p>b) type 2 protection or type 3 protection is enabled, then the device server shall write a value of FFFF_FFFFh into the LOGICAL BLOCK REFERENCE TAG field of each of the written logical blocks.</p> <p><sup>h</sup> If the ATO bit is set to one in the Control mode page (see SPC-5), then the device server shall write FFFFh into each LOGICAL BLOCK APPLICATION TAG field. If the ATO bit is set to zero, then the device server may write any value into each LOGICAL BLOCK APPLICATION TAG field.</p> <p><sup>i</sup> If multiple errors occur while the device server is processing the command, then the selection by the device server of which error to report is not defined by this standard.</p> <p><sup>j</sup> If type 1 protection is enabled, then the device server shall check the logical block reference tag by comparing it to the lower four bytes of the LBA associated with the logical block. If type 2 protection is enabled and the device server has knowledge of the contents of the LOGICAL BLOCK REFERENCE TAG field, then the device server shall check each logical block reference tag. If type 2 protection is enabled, then this knowledge may be acquired through the EXPECTED INITIAL LOGICAL BLOCK REFERENCE TAG field in a WRITE (32) command, a WRITE ATOMIC (32) command, a WRITE SAME (32) command, or a WRITE STREAM (32) command. If type 3 protection is enabled, the ATO bit is set to one in the Control mode page (see SPC-5), and the device server has knowledge of the contents of the LOGICAL BLOCK REFERENCE TAG field, then the device server may check each logical block reference tag. If type 3 protection is enabled, then the method for acquiring this knowledge is not defined by this standard.</p> <p><sup>k</sup> If the NO_PL_CHK bit is set to one in the Extended INQUIRY Data VPD page (see SPC-5) and the device server detects:</p> <p>a) a LOGICAL BLOCK APPLICATION TAG field set to FFFFh and type 1 protection (see 4.22.2.3) or type 2 protection (see 4.22.2.4) is enabled; or</p> <p>b) a LOGICAL BLOCK APPLICATION TAG field set to FFFFh, LOGICAL BLOCK REFERENCE TAG field set to FFFF_FFFFh, and type 3 protection (see 4.22.2.5) is enabled,</p> <p>then the device server shall not check any protection information in the associated protection information interval.</p>				

See the READ (10) command (see 5.13) for the definition of the DPO bit.

A force unit access (FUA) bit set to one specifies that the device server shall write the logical blocks to:

- a) the non-volatile cache, if any; or
- b) the medium.

An FUA bit set to zero specifies that the device server shall write the logical blocks to:

- a) volatile cache, if any;
- b) non-volatile cache, if any; or
- c) the medium.

See the PRE-FETCH (10) command (see 5.10) for the definition of the LOGICAL BLOCK ADDRESS field.

See the PRE-FETCH (10) command and 4.23 for the definition of the GROUP NUMBER field.

The TRANSFER LENGTH field specifies the number of contiguous logical blocks of data that shall be transferred from the Data-Out Buffer and written, starting with the logical block referenced by the LBA specified by the LOGICAL BLOCK ADDRESS field. A TRANSFER LENGTH field set to zero specifies that no logical blocks shall be transferred or written. This condition shall not be considered an error. Any other value specifies the number of logical blocks that shall be transferred and written. If the specified LBA and the specified transfer length exceed the capacity of the medium (see 4.5), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to LOGICAL BLOCK ADDRESS OUT OF RANGE. The TRANSFER LENGTH field is constrained by the MAXIMUM TRANSFER LENGTH field ~~in the Block Limits VPD page~~ (see 6.6.4).

The CONTROL byte is defined in SAM-5.

### 5.36 WRITE (12) command

The WRITE (12) command (see table 120) requests that the device server perform the actions defined for the WRITE (10) command (see 5.35).

NOTE 18 - Migration from the WRITE (12) command to the WRITE 16) command is recommended for all implementations.

**Table 120 — WRITE (12) command**

Bit	7	6	5	4	3	2	1	0
Byte								
0	OPERATION CODE (AAh)							
1	WRPROTECT			DPO	FUA	Reserved	Obsolete	Obsolete
2	(MSB)							
...	LOGICAL BLOCK ADDRESS							
5	(LSB)							
6	(MSB)							
...	TRANSFER LENGTH							
9	(LSB)							
10	Restricted for MMC-6	Reserved	GROUP NUMBER					
11	CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 120 for the WRITE (12) command.

See the WRITE (10) command (see 5.35) for the definitions of the other fields in this command.

### 5.37 WRITE (16) command

The WRITE (16) command (see table 121) requests that the device server perform the actions defined for the WRITE (10) command (see 5.35).

**Table 121 — WRITE (16) command**

Bit	7	6	5	4	3	2	1	0
Byte								
0	OPERATION CODE (8Ah)							
1	WRPROTECT			DPO	FUA	Reserved	Obsolete	DLD2
2	(MSB)							
...	LOGICAL BLOCK ADDRESS							
9	(LSB)							
10	(MSB)							
...	TRANSFER LENGTH							
13	(LSB)							
14	DLD1	DLD0	GROUP NUMBER					
15	CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 121 for the WRITE (16) command.

See the READ (16) command (see 5.15) for the definitions of the DLD2 bit, the DLD1 bit, and the DLD0 bit.

See the WRITE (10) command (see 5.35) for the definitions of the other fields in this command.

### 5.38 WRITE (32) command

The WRITE (32) command (see table 122) requests that the device server perform the actions defined for the WRITE (10) command (see 5.35).

The device server shall process a WRITE (32) command only if type 2 protection is enabled (see 4.22.2.4).

**Table 122 — WRITE (32) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (7Fh)							
1		CONTROL							
2		Reserved							
...									
5									
6		Reserved	GROUP NUMBER						
7		ADDITIONAL CDB LENGTH (18h)							
8	(MSB)	SERVICE ACTION (000Bh)							
9									
10		WRPROTECT			DPO	FUA	Reserved	Obsolete	Reserved
11		Reserved							
12	(MSB)	LOGICAL BLOCK ADDRESS							
...									
19		(LSB)							
20	(MSB)	EXPECTED INITIAL LOGICAL BLOCK REFERENCE TAG							
...									
23		(LSB)							
24	(MSB)	EXPECTED LOGICAL BLOCK APPLICATION TAG							
25									
26	(MSB)	LOGICAL BLOCK APPLICATION TAG MASK							
27									
28	(MSB)	TRANSFER LENGTH							
...									
31									

The OPERATION CODE field, the ADDITIONAL CDB LENGTH field, and the SERVICE ACTION field are defined in SPC-5 and shall be set to the values shown in table 122 for the WRITE (32) command.

See the WRITE (10) command (see 5.35) for the definitions of the CONTROL byte, the GROUP NUMBER field, the WRPROTECT field, the DPO bit, the FUA bit, the LOGICAL BLOCK ADDRESS field, and the TRANSFER LENGTH field.

If checking of the LOGICAL BLOCK REFERENCE TAG field is enabled (see table 119), then the EXPECTED INITIAL LOGICAL BLOCK REFERENCE TAG field contains the value of the LOGICAL BLOCK REFERENCE TAG field expected in the protection information of the first logical block accessed by the command instead of a value based on the LBA (see 4.22.3).

If the ATO bit is set to one in the Control mode page (see SPC-5) and checking of the LOGICAL BLOCK APPLICATION TAG field is enabled (see table 119), then the LOGICAL BLOCK APPLICATION TAG MASK field contains a value that is a bit mask for enabling the checking of the LOGICAL BLOCK APPLICATION TAG field in every instance of protection information for each logical block accessed by the command. A LOGICAL BLOCK APPLICATION TAG MASK bit set to one enables the checking of the corresponding bit of the EXPECTED LOGICAL BLOCK APPLICATION TAG field with the corresponding bit of the LOGICAL BLOCK APPLICATION TAG field in every instance of protection information. A LOGICAL BLOCK APPLICATION TAG MASK field bit set to zero disables the checking of the corresponding bit of the EXPECTED LOGICAL BLOCK APPLICATION TAG field with the corresponding bit of the LOGICAL BLOCK APPLICATION TAG field in every instance of protection information.

If the ATO bit is set to:

- a) zero; or
- b) one in the Control mode page (see SPC-5) and checking of the LOGICAL BLOCK APPLICATION TAG field is disabled (see table 119),

then the LOGICAL BLOCK APPLICATION TAG MASK field and the EXPECTED LOGICAL BLOCK APPLICATION TAG field shall be ignored

### 5.39 WRITE AND VERIFY (10) command

The WRITE AND VERIFY (10) command (see table 123) requests that the device server:

- 1) transfer the specified the logical block data for the command from the Data-Out Buffer;
- 2) perform write medium operations to the specified LBAs;
- 3) perform verify operations from the specified LBAs; and
- 4) if specified, perform a compare operation on:
  - A) the logical block data transferred from the Data-Out Buffer; and
  - B) the logical block data from the verify operations.

The device server may process the LBAs in any order but shall perform this sequence in the specified order for a given LBA.

NOTE 19 - Migration from the WRITE AND VERIFY (10) command to the WRITE AND VERIFY (16) command is recommended for all implementations.

**Table 123 — WRITE AND VERIFY (10) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (2Eh)							
1		WRPROTECT			DPO	Reserved	BYTCHK		Obsolete
2	(MSB)	LOGICAL BLOCK ADDRESS							
...									
5									
6		Reserved		GROUP NUMBER					
7	(MSB)	TRANSFER LENGTH							
8									
9		CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 123 for the WRITE AND VERIFY (10) command.

See the PRE-FETCH (10) command (see 5.10) for the definition of the LOGICAL BLOCK ADDRESS field. See the PRE-FETCH (10) command and 4.23 for the definition of the GROUP NUMBER field. See the WRITE (10) command (see 5.35) for the definitions of the CONTROL byte, the TRANSFER LENGTH field and the WRPROTECT field. See the READ (10) command (see 5.13) for the definition of the DPO bit.

See the VERIFY (10) command (see 5.31) for definition of the byte check (BYTCHK) field when set to 00b, 01b, and 10b. For a WRITE AND VERIFY (10) command, a BYTCHK field set to 11b is reserved.

## 5.40 WRITE AND VERIFY (12) command

The WRITE AND VERIFY (12) command (see table 124) requests that the device server perform the actions defined for the WRITE AND VERIFY (10) command (see 5.39).

NOTE 20 - Migration from the WRITE AND VERIFY (12) command to the WRITE AND VERIFY (16) command is recommended for all implementations.

**Table 124 — WRITE AND VERIFY (12) command**

Bit	7	6	5	4	3	2	1	0
Byte								
0	OPERATION CODE (AEh)							
1	WRPROTECT			DPO	Reserved	BYTCHK		Obsolete
2	(MSB)							
...	LOGICAL BLOCK ADDRESS							
5	(LSB)							
6	(MSB)							
...	TRANSFER LENGTH							
9	(LSB)							
10	Reserved		GROUP NUMBER					
11	CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 124 for the WRITE AND VERIFY (12) command.

See the WRITE AND VERIFY (10) command (see 5.39) for the definitions of the other fields in this command.

## 5.41 WRITE AND VERIFY (16) command

The WRITE AND VERIFY (16) command (see table 125) requests that the device server perform the actions defined for the WRITE AND VERIFY (10) command (see 5.39).

**Table 125 — WRITE AND VERIFY (16) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (8Eh)							
1		WRPROTECT			DPO	Reserved	BYTCHK		Reserved
2		(MSB) LOGICAL BLOCK ADDRESS (LSB)							
...									
9									
10		(MSB) TRANSFER LENGTH (LSB)							
...									
13									
14		Reserved		GROUP NUMBER					
15		CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 125 for the WRITE AND VERIFY (16) command.

See the WRITE AND VERIFY (10) command (see 5.39) for the definitions of the other fields in this command.

## 5.42 WRITE AND VERIFY (32) command

The WRITE AND VERIFY (32) command (see table 126) requests that the device server perform the actions defined for the WRITE AND VERIFY (10) command (see 5.39).

The device server shall process a WRITE AND VERIFY (32) command only if type 2 protection is enabled (see 4.22.2.4).

**Table 126 — WRITE AND VERIFY (32) command**

Bit	7	6	5	4	3	2	1	0
Byte								
0	OPERATION CODE (7Fh)							
1	CONTROL							
2	Reserved							
...								
5								
6	Reserved	GROUP NUMBER						
7	ADDITIONAL CDB LENGTH (18h)							
8	(MSB)	SERVICE ACTION (000Ch)						
9								
10	WRPROTECT			DPO	Reserved	BYTCHK		Reserved
11	Reserved							
12	(MSB)	LOGICAL BLOCK ADDRESS						
...								
19								
20	(MSB)	EXPECTED INITIAL LOGICAL BLOCK REFERENCE TAG						
...								
23								
24	(MSB)	EXPECTED LOGICAL BLOCK APPLICATION TAG						
25								
26	(MSB)	LOGICAL BLOCK APPLICATION TAG MASK						
27								
28	(MSB)	TRANSFER LENGTH						
...								
31								

The OPERATION CODE field, the ADDITIONAL CDB LENGTH field, and the SERVICE ACTION field are defined in SPC-5 and shall be set to the values shown in table 126 for the WRITE AND VERIFY (32) command.

See the WRITE AND VERIFY (10) command (see 5.39) for the definitions of the CONTROL byte, the GROUP NUMBER field, the WRPROTECT field, the DPO bit, the BYTCHK field, the LOGICAL BLOCK ADDRESS field, and the TRANSFER LENGTH field.

If checking of the LOGICAL BLOCK REFERENCE TAG field is enabled (see table 119), then the EXPECTED INITIAL LOGICAL BLOCK REFERENCE TAG field contains the value of the LOGICAL BLOCK REFERENCE TAG field expected in



the protection information of the first logical block accessed by the command instead of a value based on the LBA (see 4.22.3).

If the ATO bit is set to one in the Control mode page (see SPC-5) and checking of the LOGICAL BLOCK APPLICATION TAG field is enabled (see table 119 in 5.35), then the LOGICAL BLOCK APPLICATION TAG MASK field contains a value that is a bit mask for enabling the checking of the LOGICAL BLOCK APPLICATION TAG field in every instance of protection information for each logical block accessed by the command. A LOGICAL BLOCK APPLICATION TAG MASK bit set to one enables the checking of the corresponding bit of the EXPECTED LOGICAL BLOCK APPLICATION TAG field with the corresponding bit of the LOGICAL BLOCK APPLICATION TAG field in every instance of protection information. A LOGICAL BLOCK APPLICATION TAG MASK field bit set to zero disables the checking of the corresponding bit of the EXPECTED LOGICAL BLOCK APPLICATION TAG field with the corresponding bit of the LOGICAL BLOCK APPLICATION TAG field in every instance of protection information.

If the ATO bit is set to one in the Control mode page (see SPC-5) and checking of the LOGICAL BLOCK APPLICATION TAG field is disabled (see table 119), or if the ATO bit is set to zero, then the LOGICAL BLOCK application tag mask field and the expected logical block application tag field shall be ignored.

### 5.43 WRITE ATOMIC (16) command

The WRITE ATOMIC (16) command (see table 127) requests that the device server:

- a) transfer logical block data from the Data-Out Buffer; and
- b) perform one or more atomic write operations (see 4.31) of the LBAs specified by this command.

**Table 127 — WRITE ATOMIC (16) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (9Ch)							
1		WRPROTECT			DPO	FUA	Reserved		
2	(MSB)	LOGICAL BLOCK ADDRESS							
...									
9		(LSB)							
10		ATOMIC BOUNDARY							
11									
12	(MSB)	TRANSFER LENGTH							
13									
14		Reserved		GROUP NUMBER					
15		CONTROL							

The operation code field is defined in SPC-5 and shall be set to the value shown in table 127 for the WRITE ATOMIC (16) command.

The ATOMIC BOUNDARY field specifies whether multiple atomic write operations may be performed. If the ATOMIC BOUNDARY field is set to zero, then a single atomic write operation of the length specified in the TRANSFER LENGTH field shall be performed. If the ATOMIC BOUNDARY field is set to a non-zero value then multiple atomic write operations may be performed as described in 4.31.

See the WRITE (10) command (see 5.35) for the definitions of the other fields in this command.

## 5.44 WRITE ATOMIC (32) command

The WRITE ATOMIC (32) command (see table 128) requests that the device server perform the actions defined for the WRITE ATOMIC (16) command (see 5.43).

The device server shall process a WRITE ATOMIC (32) command only if type 2 protection is enabled (see 4.22.2.4).

**Table 128 — WRITE ATOMIC (32) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (7Fh)							
1		CONTROL							
2		Reserved							
3									
4	(MSB)	ATOMIC BOUNDARY							
5		(LSB)							
6		Reserved		GROUP NUMBER					
7		ADDITIONAL CDB LENGTH (18h)							
8	(MSB)	SERVICE ACTION (000Fh)							
9		(LSB)							
10		WRPROTECT			DPO	FUA	Reserved	Obsolete	Reserved
11		Reserved							
12	(MSB)	LOGICAL BLOCK ADDRESS							
...									
19		(LSB)							
20	(MSB)	EXPECTED INITIAL LOGICAL BLOCK REFERENCE TAG							
...									
23		(LSB)							
24	(MSB)	EXPECTED LOGICAL BLOCK APPLICATION TAG							
25									
26	(MSB)	LOGICAL BLOCK APPLICATION TAG MASK							
27									
28	(MSB)	TRANSFER LENGTH							
...									
31		(LSB)							

The operation code field, the additional cdb length field, and the service action field are defined in SPC-5 and shall be set to the values shown in table 128 for the WRITE ATOMIC (32) command.

See the WRITE ATOMIC (16) command (see 5.43) for the definition of the ATOMIC BOUNDARY field.

See the WRITE (32) command (see 5.38) for the definitions of the other fields in this command.

## 5.45 WRITE LONG (10) command

The WRITE LONG (10) command (see table 129) requests that the device server mark a logical block as containing a pseudo unrecovered error. If a cache contains the specified logical block, then the device server shall invalidate that logical block in the cache.

NOTE 21 - Migration from the WRITE LONG (10) command to the WRITE LONG (16) command is recommended for all implementations.

**Table 129 — WRITE LONG (10) command**

Byte	Bit	7	6	5	4	3	2	1	0
0	OPERATION CODE (3Fh)								
1	Obsolete	WR_UNCOR	Obsolete	Reserved				Obsolete	
2	(MSB)								
...	LOGICAL BLOCK ADDRESS								
5									
6	Reserved								
7	(MSB)								
8	Obsolete								
9	(LSB)								
	CONTROL								

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 129 for the WRITE LONG (10) command.

The write uncorrectable error (WR\_UNCOR) bit is defined in table 130.

**Table 130 — WR\_UNCOR bit**

WR_UNCOR	Description
0	Obsolete
1	Mark the specified logical block as containing a pseudo unrecovered error with correction disabled (see 4.18.2).  No data is transferred.

If the WRITE LONG command is supported, then the WU\_SUP bit in the Extended INQUIRY Data VPD page (see SPC-5) shall be set to one to indicate that the logical unit supports setting the WR\_UNCOR bit to one.

The LOGICAL BLOCK ADDRESS field specifies an LBA. If the specified LBA exceeds the capacity of the medium (see 4.5), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to LOGICAL BLOCK ADDRESS OUT OF RANGE.

The CONTROL byte is defined in SAM-5.

## 5.46 WRITE LONG (16) command

The WRITE LONG (16) command (see table 131) requests that the device server mark a logical block as containing an error. If a cache contains the specified logical block, then the device server shall invalidate that logical block in the cache.

This command uses the SERVICE ACTION OUT (16) CDB format (see A.2).

**Table 131 — WRITE LONG (16) command**

Byte	Bit	7	6	5	4	3	2	1	0
0	OPERATION CODE (9Fh)								
1	Obsolete	WR_UNCOR	Obsolete	SERVICE ACTION (11h)					
2	(MSB)								
...	LOGICAL BLOCK ADDRESS								
9									
10	Reserved								
11									
12	(MSB)								
13	Obsolete								
14	Reserved								
15									

The OPERATION CODE field and SERVICE ACTION field are defined in SPC-5 and shall be set to the values shown in table 131 for the WRITE LONG (16) command.

See the WRITE LONG (10) command (see 5.45) for the definitions of the fields in this command.

## 5.47 WRITE SAME (10) command

The WRITE SAME (10) command (see table 132) requests that the device server transfer a single logical block from the Data-Out Buffer and for each LBA in the specified range of LBAs:

- perform a write operation using the contents of that logical block; or
- perform an unmap operation.

The device server writes (i.e., subsequent read operations behave as if the device server wrote the single block of user data received from the Data-Out Buffer to each logical block without modification (see 4.7.3.4.4)).

If the medium is formatted with protection information and the WRPROTECT field is set to 000b, then the device server shall write the LOGICAL BLOCK GUARD field, APPLICATION TAG field, and LOGICAL BLOCK REFERENCE TAG field (see 4.22) for each logical block as described in table 119 (i.e., code equal to 000b row of table 119).

If:

- the medium is formatted with protection information;
- the WRPROTECT field is not set to 000b or a reserved value (see table 119); and
- the protection information from the Data-Out Buffer is set to FFFF\_FFFF\_FFFF\_FFFFh,

then the device server shall write FFFF\_FFFF\_FFFF\_FFFFh to the protection information for each logical block.

If:

- a) the medium is formatted with type 1 or type 2 protection information;
- b) the WRPROTECT field is not set to 000b or a reserved value (see table 119); and
- c) the protection information from the Data-Out Buffer is not set to FFFF\_FFFF\_FFFF\_FFFFh,

then:

- a) the device server shall write the value from the LOGICAL BLOCK REFERENCE TAG field (see 4.22) received in the logical block from the Data-Out Buffer into the corresponding LOGICAL BLOCK REFERENCE TAG field of the first logical block written. The device server shall write the value of the previous LOGICAL BLOCK REFERENCE TAG field plus one into each of the subsequent LOGICAL BLOCK REFERENCE TAG fields;
- b) if the ATO bit is set to one in the Control mode page (see SPC-5) and the and the ATMPE bit is set to zero in the Control mode page, then the device server shall write the logical block application tag received in the logical block from the Data-Out Buffer into the corresponding LOGICAL BLOCK APPLICATION TAG field (see 4.22) of each logical block;
- c) if the ATO bit is set to one in the Control mode page and the and the ATMPE bit is set to zero in the Control mode page, then the device server shall write the value defined in the Application Tag mode page (see 6.5.3) into the corresponding LOGICAL BLOCK APPLICATION TAG field of each logical block;
- d) if the ATO bit is set to zero in the Control mode page, then the device server may write any value into the LOGICAL BLOCK APPLICATION TAG field of each logical block; and
- e) the device server shall write the value from the LOGICAL BLOCK GUARD field (see 4.22) received in the logical block from the Data-Out Buffer into the corresponding LOGICAL BLOCK GUARD field of each logical block.

If:

- a) the medium is formatted with type 3 protection information;
- b) the WRPROTECT field is not set to 000b or a reserved value (see table 119); and
- c) the protection information from the Data-Out Buffer is not set to FFFF\_FFFF\_FFFF\_FFFFh,

then:

- a) if the ATO bit is set to one in the Control mode page (see SPC-5), then the device server shall write the value from the LOGICAL BLOCK REFERENCE TAG field and the LOGICAL BLOCK APPLICATION TAG field received in the logical block from the Data-Out Buffer into the corresponding LOGICAL BLOCK REFERENCE TAG field of each logical block;
- b) if the ATO bit is set to zero in the Control mode page, then the device server may write any value into the LOGICAL BLOCK REFERENCE TAG field of each logical block; and
- c) the device server shall write the value from the LOGICAL BLOCK GUARD field and the LOGICAL BLOCK APPLICATION TAG field received in the logical block from the Data-Out Buffer into the corresponding LOGICAL BLOCK GUARD field of each logical block.

NOTE 22 - Migration from the WRITE SAME (10) command to the WRITE SAME (16) command is recommended for all implementations.

**Table 132 — WRITE SAME (10) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (41h)							
1		WRPROTECT			ANCHOR	UNMAP	Obsolete		
2	(MSB)	LOGICAL BLOCK ADDRESS							
...									
5									
6		Reserved		GROUP NUMBER					
7	(MSB)	NUMBER OF LOGICAL BLOCKS							
8									
9		CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 132 for the WRITE SAME (10) command.

See the WRITE (10) command (see 5.35) for the definition of the WRPROTECT field.

If the logical unit supports logical block provisioning management (see 4.7.3), then the ANCHOR bit in the CDB, the UNMAP bit in the CDB, and the ANC\_SUP bit ~~in the Logical Block Provisioning VPD page~~ (see 6.6.6) determine how the device server processes the command as described in table 133.

**Table 133 — UNMAP bit, ANCHOR bit, and ANC\_SUP bit relationships**

UNMAP bit <sup>a</sup>	ANCHOR bit	ANC_SUP bit <sup>b</sup>	Action
0	0	0 or 1	Write <sup>c</sup>
	1	0 or 1	Error <sup>d</sup>
1	0	0 or 1	Deallocate request (see 4.7.3.4.4)
	1	0	Error <sup>d</sup>
		1	Anchor request (see 4.7.3.4.4)

<sup>a</sup> The device server in a logical unit that supports logical block provisioning management (see 4.7.3) may implement the UNMAP bit.

<sup>b</sup> See the Logical Block Provisioning VPD page (see 6.6.6).

<sup>c</sup> The device server shall perform the specified write operation to each LBA specified by the command.

<sup>d</sup> The device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

The device server shall ignore the UNMAP bit and the ANCHOR bit, or the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB if:

- a) the logical unit is fully provisioned (i.e., the LBPME bit is set to zero in the READ CAPACITY (16) parameter data (see 5.18.2)); and
- b) the UNMAP bit is set to one or the ANCHOR bit is set to one.

See the PRE-FETCH (10) command (see 5.10) for the definition of the LOGICAL BLOCK ADDRESS field.

See the PRE-FETCH (10) command (see 5.10) and 4.23 for the definition of the GROUP NUMBER field.

If the WSNZ bit is set to zero in the Block Limits VPD page (see 6.6.4), then a NUMBER OF LOGICAL BLOCKS field set to a non-zero value specifies the number of contiguous logical blocks that are requested be unmapped or written, starting with the logical block referenced by the LBA specified by the LOGICAL BLOCK ADDRESS field.

If the WSNZ bit is set to zero, then a NUMBER OF LOGICAL BLOCKS field set to zero specifies that the number of contiguous logical blocks that are requested to be unmapped or written includes all of the logical blocks starting with the LBA specified in the LOGICAL BLOCK ADDRESS field to the last logical block on the medium. If the WSNZ bit is set to one and the NUMBER OF LOGICAL BLOCKS field is set to zero, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

If the number of logical blocks specified to be unmapped or written exceeds the value indicated in the MAXIMUM WRITE SAME LENGTH field ~~in the Block Limits VPD page~~ (see 6.6.4), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

See the WRITE (10) command (see 5.35) for the definition of the CONTROL byte.

## 5.48 WRITE SAME (16) command

The WRITE SAME (16) command (see table 134) requests that the device server perform the actions defined for the WRITE SAME (10) command (see 5.47).

**Table 134 — WRITE SAME (16) command**

Bit Byte	7	6	5	4	3	2	1	0
0	OPERATION CODE (93h)							
1	WRPROTECT			ANCHOR	UNMAP	Obsolete		NDOB
2	(MSB)							
...	LOGICAL BLOCK ADDRESS							
9	(LSB)							
10	(MSB)							
...	NUMBER OF LOGICAL BLOCKS							
13	(LSB)							
14	Reserved		GROUP NUMBER					
15	CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 134 for the WRITE SAME (16) command.

A NDOB bit set to zero specifies that the device server shall process the command using logical block data from the Data-Out Buffer. A NDOB bit set to one specifies that:

- a) the device server shall not transfer data from the Data-Out Buffer;
  - b) if [the Logical Block Provisioning VPD page \(see 6.6.6\) is not supported or](#) the LBPRZ field ~~in the Logical Block Provisioning VPD page~~ (see 6.6.6) is set to 000b or xx1b, then the device server shall process the command as if the Data-Out Buffer contained user data set to all zeroes and protection information, if any, containing:
    - A) the logical block guard field set to FFFFh;
    - B) the logical block reference tag field set to FFFF\_FFFFh; and
    - C) the logical block application tag field set to FFFFh;
- and
- c) if the LBPRZ field ~~in the Logical Block Provisioning VPD page~~ is set to 010b, then the device server shall process the command as if the Data-Out Buffer contained user data set to the provisioning initialization pattern and protection information, if any, containing:
    - A) the logical block guard field set to FFFFh;
    - B) the logical block reference tag field set to FFFF\_FFFFh; and
    - C) the logical block application tag field set to FFFFh.

See the WRITE SAME (10) command (see 5.47) for the definitions of the other fields in this command.

## 5.49 WRITE SAME (32) command

The WRITE SAME (32) command (see table 135) requests that the device server perform the actions defined for the WRITE SAME (10) command (see 5.47).

The device server shall process a WRITE SAME (32) command only if type 2 protection is enabled (see 4.22.2.4).



**Table 135 — WRITE SAME (32) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (7Fh)							
1		CONTROL							
2		Reserved							
...									
5									
6		Reserved	GROUP NUMBER						
7		ADDITIONAL CDB LENGTH (18h)							
8	(MSB)	SERVICE ACTION (000Dh)							
9									
10		WRPROTECT			ANCHOR	UNMAP	Obsolete		NDOB
11		Reserved							
12	(MSB)	LOGICAL BLOCK ADDRESS							
...									
19		(LSB)							
20	(MSB)	EXPECTED INITIAL LOGICAL BLOCK REFERENCE TAG							
...									
23		(LSB)							
24	(MSB)	EXPECTED LOGICAL BLOCK APPLICATION TAG							
25									
26	(MSB)	LOGICAL BLOCK APPLICATION TAG MASK							
27									
28	(MSB)	NUMBER OF LOGICAL BLOCKS							
...									
31		(LSB)							

The OPERATION CODE field, the ADDITIONAL CDB LENGTH field, and the SERVICE ACTION field are defined in SPC-5 and shall be set to the values shown in table 135 for the WRITE SAME (32) command.

See the WRITE SAME (10) command (see 5.47) for the definitions of the CONTROL byte, the GROUP NUMBER field, the WRPROTECT field, the LOGICAL BLOCK ADDRESS field, the NUMBER OF LOGICAL BLOCKS field, the UNMAP bit, and the ANCHOR bit.

See the WRITE SAME (16) command (see 5.48) for the definition of the NDOB bit.

If checking of the LOGICAL BLOCK REFERENCE TAG field is enabled (see table 119 in 5.35), then the EXPECTED INITIAL LOGICAL BLOCK REFERENCE TAG field contains the value of the LOGICAL BLOCK REFERENCE TAG field expected in the protection information of the first logical block accessed by the command instead of a value based on the LBA (see 4.22.3).

If the ATO bit is set to one in the Control mode page (see SPC-5) and checking of the LOGICAL BLOCK APPLICATION TAG field is enabled (see table 119 in 5.35), then the LOGICAL BLOCK APPLICATION TAG MASK field contains a value that is a bit mask for enabling the checking of the LOGICAL BLOCK APPLICATION TAG field in every instance of protection information for each logical block accessed by the command. A LOGICAL BLOCK

APPLICATION TAG MASK bit set to one enables the checking of the corresponding bit of the EXPECTED LOGICAL BLOCK APPLICATION TAG field with the corresponding bit of the LOGICAL BLOCK APPLICATION TAG field in every instance of protection information. A LOGICAL BLOCK APPLICATION TAG MASK field bit set to zero disables the checking of the corresponding bit of the EXPECTED LOGICAL BLOCK APPLICATION TAG field with the corresponding bit of the LOGICAL BLOCK APPLICATION TAG field in every instance of protection information.

If the ATO bit is set to one in the Control mode page (see SPC-5) and checking of the LOGICAL BLOCK APPLICATION TAG field is disabled (see table 119 in 5.35), or if the ATO bit is set to zero, then the LOGICAL BLOCK APPLICATION TAG MASK field and the EXPECTED LOGICAL BLOCK APPLICATION TAG field shall be ignored.

## 5.50 WRITE STREAM (16) command

The WRITE STREAM (16) command (see table 136) requests that the device server perform the actions defined for the WRITE (10) command (see 5.35).

**Table 136 — WRITE STREAM (16) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (9Ah)							
1		WRPROTECT			DPO	FUA	Reserved		
2	(MSB)	LOGICAL BLOCK ADDRESS							
...									
9									
10	(MSB)	STR_ID							
11									
12	(MSB)	TRANSFER LENGTH							
13									
14		Reserved		GROUP NUMBER					
15		CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 136 for the WRITE STREAM (16) command.

The stream identifier (STR\_ID) field specifies the stream identifier associated with this command as described in 4.34.

See the WRITE (10) command (see 5.33) for the definitions of the other fields in this command.

## 5.51 WRITE STREAM (32) command

The WRITE STREAM (32) command (see table 137) requests that the device server perform the actions defined for the WRITE (32) command (see 5.38).

The device server shall process a WRITE STREAM (32) command only if type 2 protection is enabled (see 4.22.2.4).

**Table 137 — WRITE STREAM(32) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (7Fh)							
1		CONTROL							
2		Reserved							
3									
4		STR_ID							
5									
6		Reserved		GROUP NUMBER					
7		ADDITIONAL CDB LENGTH (18h)							
8	(MSB)	SERVICE ACTION (0010h)							
9									
10		WRPROTECT			DPO	FUA	Reserved		
11		Reserved							
12	(MSB)	LOGICAL BLOCK ADDRESS							
...									
19		(LSB)							
20	(MSB)	EXPECTED INITIAL LOGICAL BLOCK REFERENCE TAG							
...									
23		(LSB)							
24	(MSB)	EXPECTED LOGICAL BLOCK APPLICATION TAG							
25									
26	(MSB)	LOGICAL BLOCK APPLICATION TAG MASK							
27									
28	(MSB)	TRANSFER LENGTH							
...									
31		(LSB)							

The OPERATION CODE field, the ADDITIONAL CDB LENGTH field, and the SERVICE ACTION field are defined in SPC-5 and shall be set to the values shown in table 137 for the WRITE STREAM (32) command.

See the WRITE STREAM (16) command (see 5.50) for the definition of the STR\_ID field.

See the WRITE (32) command (see 5.38) for the definitions of the other fields in this command.

## 5.52 WRITE USING TOKEN command

### 5.52.1 WRITE USING TOKEN command overview

The WRITE USING TOKEN command (see table 138) requests that the copy manager (see SPC-5) write logical block data represented by the specified ROD token to the specified LBAs.

**Table 138 — WRITE USING TOKEN command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (83h)							
1		Reserved			SERVICE ACTION (11h)				
2		Reserved							
...									
5									
6	(MSB)	LIST IDENTIFIER							
...									
9									
10	(MSB)	PARAMETER LIST LENGTH							
...									
13									
14		Reserved		GROUP NUMBER					
15		CONTROL							

The OPERATION CODE field and the SERVICE ACTION field are defined in ~~SPC-4~~SPC-5 and shall be set to the values shown in table 138 for the WRITE USING TOKEN command.

The LIST IDENTIFIER field is defined in SPC-5. The list identifier shall be processed as if the LIST ID USAGE field is set to 00b in the parameter data for an EXTENDED COPY(LID4) command (see SPC-5).

The PARAMETER LIST LENGTH field specifies the length in bytes of the parameter data that is available to be transferred from the Data-Out Buffer. A PARAMETER LIST LENGTH set to zero specifies that no data shall be transferred. This shall not be considered an error.

See the PRE-FETCH (10)FETCH (10) command (see 5.10) and 4.23 for the definition of the GROUP NUMBER field.

The CONTROL byte is defined in SAM-5.

### 5.52.2 WRITE USING TOKEN parameter list

The parameter list for the WRITE USING TOKEN command is defined in table 139.

**Table 139 — WRITE USING TOKEN parameter list**

Byte	Bit	7	6	5	4	3	2	1	0
0	(MSB)	WRITE USING TOKEN DATA LENGTH (n - 1)							
1									
2		Reserved						DEL_TKN	IMMED
3		Reserved							
...									
7									
8	(MSB)	OFFSET INTO ROD							
...									
15									
16		ROD TOKEN							
...									
527									
528		Reserved							
...									
533									
534		BLOCK DEVICE RANGE DESCRIPTOR LENGTH (n - 535)							
535									
Block device range descriptor list (if any)									
536		Block device range descriptor [first] (see 5.9.3)							
...									
551									
⋮									
n - 15		Block device range descriptor [last] (see 5.9.3)							
...									
n									

The WRITE USING TOKEN DATA LENGTH field specifies the length in bytes of the data that is available to be transferred from the Data-Out Buffer. The write using token data length does not include the number of bytes in the WRITE USING TOKEN DATA LENGTH field.

If the WRITE USING TOKEN DATA LENGTH field is less than 0226h (i.e., 550), then the copy manager shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

If the contents of the WRITE USING TOKEN DATA LENGTH field is not equal to the contents of the BLOCK DEVICE RANGE DESCRIPTOR LENGTH field plus 534 then the device server should terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

The immediate (IMMED) bit specifies when the copy manager shall return status for the WRITE USING TOKEN command. If the IMMED bit is set to zero, then the copy manager shall process the WRITE USING TOKEN command until all specified operations are complete or an error is detected. If the IMMED bit is set to one, then the copy manager:

- 1) shall validate the CDB (i.e., detect and report all errors in the CDB);
- 2) shall transfer all the parameter data to the copy manager;
- 3) may validate the parameter data;
- 4) shall complete the WRITE USING TOKEN command with GOOD status; and
- 5) shall complete processing of all specified operations as a background operation (see **SPC-4** SPC-5).

If the operations specified by a WRITE USING TOKEN command are processed as a background operation (i.e., the IMMED bit is set to one) (see SPC-5), then the copy manager shall not generate deferred errors (see SAM-5) to report the errors encountered, if any, during this processing. The copy manager shall make error information available to an application client using a RECEIVE ROD TOKEN INFORMATION command (see 5.22).

A delete token (DEL\_TKN) bit set to one specifies that the ROD token specified in the ROD TOKEN field should be deleted when processing of the WRITE USING TOKEN command is complete. A DEL\_TKN bit set to zero specifies that the ROD token lifetime for the ROD token specified in the ROD TOKEN field shall be as described in SPC-5.

The OFFSET INTO ROD field specifies the offset into the data represented by the ROD token from the first byte represented by the ROD token to the first byte to be transferred. The offset is specified in number of blocks based on the logical block length of the logical unit to which the WRITE USING TOKEN command is to write data. The copy manager that processes the WRITE USING TOKEN command shall compute the byte offset into the ROD by multiplying the contents of the OFFSET INTO ROD field by the logical block length of the logical unit to which the WRITE USING TOKEN command is to write data.

EXAMPLE - To calculate an offset, a ROD token is created from LBAs 15 to 20 followed by LBAs 40 to 100 by a copy manager associated with a logical unit with a logical block length of 512 bytes per logical block. That ROD token is specified in a WRITE USING TOKEN command that transfers one logical block to a logical unit with a logical block length of 4 096 bytes per logical block. The subsequent RECEIVE ROD TOKEN INFORMATION command indicates the successful transfer of 4 096 bytes by setting the TRANSFER COUNT field to one. To create a WRITE USING TOKEN command that transfers bytes from the ROD token starting at the point where the previous WRITE USING TOKEN command stopped, the OFFSET INTO ROD field is set to one (i.e., the contents of the TRANSFER COUNT field) plus the value in the OFFSET INTO ROD field in the previous WRITE USING TOKEN command (i.e., zero). The copy manager multiplies one (i.e., the value in the OFFSET INTO ROD field) by 4 096 (i.e., the logical block length for the logical unit to which the data is being written) and the result is 4 096. As a result, the ROD token logical block that is the start of the transfer is LBA 8 (i.e., LBA 42 from the logical unit whose logical block length is 512 bytes per logical block that was used to create the ROD token).

If the computed byte offset into the ROD is greater than or equal to the number of bytes represented by the ROD token, then the copy manager shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

If the ROD TOKEN LENGTH field (see SPC-5) in the ROD TOKEN field is not set to 01F8h, then the copy manager shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense key set to INVALID TOKEN OPERATION, INVALID TOKEN LENGTH.

The ROD TOKEN field specifies the ROD token that represents the data from which logical block data is written. The ROD token is defined as follows:

- a) a ROD token returned by a RECEIVE ROD TOKEN INFORMATION command; or
- b) a block device zero ROD token (see 4.30.4).

If the ROD token does not match any known to the copy manager, then the copy manager shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID TOKEN OPERATION, TOKEN UNKNOWN.

The BLOCK DEVICE RANGE DESCRIPTOR LENGTH field specifies the length in bytes of the block device range descriptor list. The block device range descriptor list length should be a multiple of 16. If the block device range descriptor list length is not a multiple of 16, then the last block device range (see 5.9.3) descriptor is incomplete and shall be ignored. If the block device range descriptor list length is less than 16, then the copy manager shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

If the number of complete block device range descriptors is larger than the MAXIMUM RANGE DESCRIPTORS field ~~in the Block Device ROD Token Limits descriptor~~ (see 6.6.8.3), then the copy manager shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to TOO MANY SEGMENT DESCRIPTORS.

If the same LBA is included in more than one block device range descriptor, then the copy manager shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

If the number of bytes of user data represented by the sum of the contents of the NUMBER OF LOGICAL BLOCKS fields in all of the complete block device range descriptors is larger than:

- a) the MAXIMUM BYTES IN BLOCK ROD field in the block ROD device type specific features descriptor in the ROD token features third-party copy descriptor in the Third-party Copy VPD page (see SPC-5) and that field is set to a nonzero value; or
- b) the MAXIMUM TOKEN TRANSFER SIZE field ~~in the Block Device ROD Token Limits descriptor~~ (see 6.6.8.3) and that field is set to a nonzero value,

then the copy manager shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

If the number of bytes of user data represented by the sum of the contents of the NUMBER OF LOGICAL BLOCKS fields in all of the complete block device range descriptors is larger than the number of bytes in the data represented by the ROD token minus the computed byte offset into the ROD (i.e., the total requested length of the transfer exceeds the length of the data available in the data represented by the ROD token), then the copy manager shall:

- a) transfer as many whole logical blocks as possible; and
- b) if any portion of a logical block that is written by the copy manager corresponds to offsets into the ROD at or beyond the length of the data represented by the ROD token, then write that portion of the logical block with user data with all bits set to zero.

The copy manager may perform this check during the processing of each block device range descriptor.

## 5.53 XDWRITEREAD (10) command

The XDWRITEREAD (10) command (see table 140) requests that the device server perform the following as an uninterrupted sequence of actions (see 4.27):

- 1) perform read operations from the specified LBAs;
- 2) transfer the specified number of logical blocks from the Data-Out Buffer;
- 3) perform an XOR operation on:
  - A) the user data contained in the logical blocks from the read operations; and
  - B) the user data contained in the logical blocks transferred from the Data-Out Buffer;
- 4) store the results of the XOR operation (i.e., the XOR data) in a buffer;
- 5) if the DISABLE WRITE bit is set to zero, then perform write operations to the specified LBAs using the logical block data transferred from the Data-Out Buffer; and
- 6) transfer the results of the XOR operation to the Data-In Buffer.

The device server may process the LBAs in any order but shall perform this sequence in the specified order for a given LBA.

Logical block data for this command may include protection information, based on the WRPROTECT field, the XORPINFO bit, and the medium format. This command is only available on transport protocols supporting bidirectional commands.

**Table 140 — XDWRITEREAD (10) command**

Byte	Bit	7	6	5	4	3	2	1	0
0	OPERATION CODE (53h)								
1	WRPROTECT			DPO	FUA	DISABLE WRITE	Obsolete	XORPINFO	
2	(MSB)								
...	LOGICAL BLOCK ADDRESS								
5	(LSB)								
6	Reserved			GROUP NUMBER					
7	(MSB)								
8	TRANSFER LENGTH								
9	(LSB)								
	CONTROL								

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 140 for the XDWRITEREAD (10) command.

See the WRITE (10) command (see 5.35) for the definitions of the WRPROTECT field and the FUA bit.

See the READ (10) command (see 5.13) for the definition of the DPO bit.

A DISABLE WRITE bit set to zero specifies that the device server shall perform a write operation using data transferred from the Data-Out Buffer after the XOR operation is complete. A DISABLE WRITE bit set to one specifies that the device server shall not perform a write operation.

If the XOR protection information (XORPINFO) bit is set to zero, then the device server shall not check or transmit protection information.

If the XORPINFO bit is set to one, the device server supports protection information, and the medium has been formatted with protection information, then the device server shall transmit protection information but shall not check any of the protection information fields.

If the XORPINFO bit is set to one and the device server:

- a) supports protection information and the medium has not been formatted with protection information or
- b) does not support protection information,

then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

See the PRE-FETCH (10) command (see 5.10) for the definition of the LOGICAL BLOCK ADDRESS field.

See the PRE-FETCH (10) command (see 5.10) and 4.23 for the definition of the GROUP NUMBER field.

The TRANSFER LENGTH field specifies the number of contiguous logical blocks of data that the device server shall read, transfer from the Data-Out Buffer, and XOR into a buffer, starting with the logical block referenced by the LBA specified by the LOGICAL BLOCK ADDRESS field. If the specified LBA and the specified transfer length exceed the capacity of the medium (see 4.5), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to LOGICAL BLOCK ADDRESS OUT OF RANGE. The TRANSFER LENGTH field is constrained by the MAXIMUM TRANSFER LENGTH field ~~in the Block Limits VPD page~~ (see 6.6.4).



The CONTROL byte is defined in SAM-5.

## 5.54 XDWRITEREAD (32) command

The XDWRITEREAD (32) command (see table 141) requests that the device server perform the actions defined for the XDWRITEREAD (10) command (see 5.53).

**Table 141 — XDWRITEREAD (32) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (7Fh)							
1		CONTROL							
2		Reserved							
...									
5									
6		Reserved	GROUP NUMBER						
7		ADDITIONAL CDB LENGTH (18h)							
8	(MSB)	SERVICE ACTION (0007h)							
9									
10		WRPROTECT			DPO	FUA	DISABLE WRITE	Obsolete	XORPINFO
11		Reserved							
12	(MSB)	LOGICAL BLOCK ADDRESS							
...									
19									
20		Reserved							
...									
27									
28	(MSB)	TRANSFER LENGTH							
...									
31									

The OPERATION CODE field, the ADDITIONAL CDB LENGTH field, and the SERVICE ACTION field are defined in SPC-5 and shall be set to the values shown in table 141 for the XDWRITEREAD (32) command.

See the XDWRITEREAD (10) command (see 5.53) for the definitions of the other fields in this command.

## 5.55 XPWRITE (10) command

The XPWRITE (10) command (see table 142) requests that the device server perform the following as an uninterrupted sequence of actions (see 4.27):

- 1) perform read operations from the specified LBAs;
- 2) transfer the specified number of logical blocks from the Data-Out Buffer;
- 3) perform an XOR operation on:
  - A) the user data contained in the logical block data from the read operations; and
  - B) the user data contained in the logical block data transferred from the Data-Out Buffer;
 and
- 4) perform write operations to the specified LBAs using the results of the XOR operation (i.e., the XOR data).

The device server may process the LBAs in any order but shall perform this sequence in the specified order for a given LBA.

Logical block data for this command may include protection information, based on the XORPINFO bit and the medium format.

**Table 142 — XPWRITE (10) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (51h)							
1		Reserved			DPO	FUA	Reserved	Obsolete	XORPINFO
2	(MSB)	LOGICAL BLOCK ADDRESS							
...									
5		(LSB)							
6		Reserved		GROUP NUMBER					
7	(MSB)	TRANSFER LENGTH							
8									
9		CONTROL							

The OPERATION CODE field is defined in SPC-5 and shall be set to the value shown in table 142 for the XPWRITE (10) command.

See the READ (10) command (see 5.13) for the definition of the DPO bit. See the WRITE (10) command (see 5.35) for the definition of the FUA bit. See the PRE-FETCH (10) command (see 5.10) for the definition of the LOGICAL BLOCK ADDRESS field. See the PRE-FETCH (10) command (see 5.10) and 4.23 for the definition of the GROUP NUMBER field.

If the XOR protection information (XORPINFO) bit is set to zero, the device server supports protection information, and the medium has been formatted with protection information, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

If the XORPINFO bit is set to one, the device server supports protection information, and the medium has been formatted with protection information, then the device server shall XOR the logical block data transferred from the Data-Out Buffer with the logical block data read, and then write the resulting XOR data. The device server shall not check any of the protection information fields.

If the XORPINFO bit is set to one and the device server:

- a) supports protection information and the medium has not been formatted with protection information;  
or
- b) does not support protection information,

then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN CDB.

The TRANSFER LENGTH field specifies the number of contiguous logical blocks that shall be read, XORed with logical blocks transferred from the Data-Out Buffer, and written, starting with the logical block referenced by the LBA specified by the LOGICAL BLOCK ADDRESS field. If the specified LBA and the specified transfer length exceed the capacity of the medium (see 4.5), then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to LOGICAL BLOCK ADDRESS OUT OF RANGE. The TRANSFER LENGTH field is constrained by the MAXIMUM TRANSFER LENGTH field ~~in the Block Limits VPD page~~ (see 6.6.4).

The CONTROL byte is defined in SAM-5.

## 5.56 XPWRITE (32) command

The XPWRITE (32) command (see table 143) requests that the device server perform the actions defined for the XPWRITE (10) command (see 5.55).

**Table 143 — XPWRITE (32) command**

Byte	Bit	7	6	5	4	3	2	1	0
0		OPERATION CODE (7Fh)							
1		CONTROL							
2		Reserved							
...									
5									
6		Reserved	GROUP NUMBER						
7		ADDITIONAL CDB LENGTH (18h)							
8	(MSB)	SERVICE ACTION (0006h)							
9									
10		Reserved			DPO	FUA	Reserved	Obsolete	XORPINFO
11		Reserved							
12	(MSB)	LOGICAL BLOCK ADDRESS							
...									
19		(LSB)							
20		Reserved							
...									
27									
28	(MSB)	TRANSFER LENGTH							
...									
31		(LSB)							

The OPERATION CODE field, the ADDITIONAL CDB LENGTH field, and the SERVICE ACTION field are defined in SPC-5 and shall be set to the values shown in table 143 for the XPWRITE (32) command.

See the XPWRITE (10) command (see 5.55) for the definitions of the other fields in this command.

## 6 Parameters for direct access block devices

### 6.1 Parameters for direct access block devices introduction

Table 144 shows the parameters for direct access block devices defined in clause 6 and a reference to the subclause where each parameter type is defined.

**Table 144 — Parameters for direct access block devices**

Parameter type	Reference
Address descriptors	6.2
Diagnostic parameters	6.3
Log parameters	6.4
Mode parameters	6.5
Vital product data (VPD) parameters	6.6
Copy manager parameters	6.7
Logical block markup descriptors	6.8

### 6.2 Address descriptors

#### 6.2.1 Address descriptor overview

This subclause describes the address descriptors (see table 145) used for:

- a) the FORMAT UNIT command (see 5.4);
- b) the READ DEFECT DATA commands (see 5.19 and 5.20);
- c) the Translate Address Input diagnostic page (see 6.3.4) for the SEND DIAGNOSTIC command (see SPC-5); and
- d) the Translate Address Output diagnostic page (see 6.3.5) for the RECEIVE DIAGNOSTIC RESULTS command (see SPC-5).

The format type of an address descriptor is:

- a) specified in the DEFECT LIST FORMAT field (see 5.4.1);
- b) indicated in the DEFECT LIST FORMAT field (see 5.19.2 and 5.20.2);
- c) specified in the SUPPLIED FORMAT field and the TRANSLATE FORMAT field for the Translate Address Output diagnostic page; or
- d) indicated in the SUPPLIED FORMAT field and the TRANSLATE FORMAT field for the Translate Address Input diagnostic page.

Table 145 defines the types of address descriptors.

**Table 145 — Address descriptors**

Type	Description	Reference
000b	Short block format address descriptor	6.2.2
001b	Extended bytes from index format address descriptor <sup>a</sup>	6.2.3
010b	Extended physical sector format address descriptor <sup>a</sup>	6.2.4
011b	Long block format address descriptor	6.2.5
100b	Bytes from index format address descriptor <sup>a</sup>	6.2.6
101b	Physical sector format address descriptor <sup>a</sup>	6.2.7
110b	Vendor specific	
111b	Reserved	
<sup>a</sup> This address descriptor format type is defined for direct access block devices using rotating media (see 4.3.2).		

### 6.2.2 Short block format address descriptor

A format type of 000b specifies the short block format address descriptor (see table 146).

**Table 146 — Short block format address descriptor (000b)**

Bit	7	6	5	4	3	2	1	0
Byte								
0	(MSB)							
...	SHORT BLOCK ADDRESS							
3								

For the FORMAT UNIT parameter list, the SHORT BLOCK ADDRESS field specifies a four-byte LBA. If the physical block containing the logical block referenced by the specified LBA contains additional logical blocks, then the device server may consider the LBAs of those additional logical blocks to also have been specified.

For the READ DEFECT DATA parameter data, the SHORT BLOCK ADDRESS field indicates a vendor specific four-byte value.

For the Translate Address diagnostic pages, the SHORT BLOCK ADDRESS field contains:

- a) a four-byte LBA, if the value is less than or equal to the capacity of the medium; or
- b) a vendor specific four-byte value, if the value is greater than the capacity of the medium.

### 6.2.3 Extended bytes from index address descriptor

A format type of 001b specifies the extended bytes from index format address descriptor (see table 147). For the FORMAT UNIT parameter list and the READ DEFECT DATA parameter data, this address descriptor contains the location of a defect that:

- a) is the length of one track (see 4.3.2);
- b) is less than the length of a physical block; or
- c) starts from one address descriptor and extends to the next address descriptor.

For the Translate Address diagnostic pages, this address descriptor contains the location of an LBA. For the Translate Address Output diagnostic page (see 6.2.5), if the SUPPLIED FORMAT field is set to 001b and the MADS bit in the ADDRESS TO TRANSLATE field is set to one, then the device server shall terminate the SEND DIAGNOSTIC command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

**Table 147 — Extended bytes from index format address descriptor (001b)**

Bit	7	6	5	4	3	2	1	0
Byte								
0	(MSB)							
...	CYLINDER NUMBER							
2	(LSB)							
3	HEAD NUMBER							
4	MADS	Reserved			(MSB)			
...	BYTES FROM INDEX							
7	(LSB)							

The CYLINDER NUMBER field contains the cylinder number (see 4.3.2).

The HEAD NUMBER field contains the head number (see 4.3.2).

A multi-address descriptor start (MADS) bit set to one specifies that this address descriptor defines the beginning of a defect that spans multiple addresses. The defect may be a number of sequential physical blocks on the same cylinder and head (i.e., a track) or may span a number of sequential tracks on the same head. A MADS bit set to zero specifies that:

- this address descriptor defines the end of a defect if the previous address descriptor has the MADS bit set to one; or
- this address descriptor defines a single track that contains one or more defects (i.e., the BYTES FROM INDEX field contains FFF\_FFFFh) or a single defect (i.e., the BYTES FROM INDEX field does not contain FFF\_FFFFh).

See 4.13.2 for valid combinations of two address descriptors that describe a defect.

The BYTES FROM INDEX field:

- if not set to FFF\_FFFFh, contains the number of bytes from the index (e.g., from the start of the track) to the location being described; or
- if set to FFF\_FFFFh, specifies or indicates that the entire track is being described.

More than one logical block may be described by this address descriptor.

Table 148 defines the order of the fields used for sorting extended bytes from index format address descriptors if the command using the address descriptors specifies sorting.

**Table 148 — Sorting order for extended bytes from index format address descriptors**

Bit:	(MSB) 59	...	36	35	...	28	27	...	(LSB) 0
	CYLINDER NUMBER field			HEAD NUMBER field			BYTES FROM INDEX field		

### 6.2.4 Extended physical sector format address descriptor

A format type of 010b specifies the extended physical sector format address descriptor (see table 149). For the FORMAT UNIT parameter list and the READ DEFECT DATA parameter data, this address descriptor contains the location of a defect that:

- a) is the length of one track (see 4.3.2);
- b) is less than the length of a physical block; or
- c) starts from one address descriptor and extends to the next address descriptor.

For the Translate Address diagnostic pages, this address descriptor specifies the location of an LBA. For the Translate Address Output diagnostic page (see 6.2.5), if the SUPPLIED FORMAT field is set to 010b and the MADS bit in the ADDRESS TO TRANSLATE field is set to one, then the device server shall terminate the SEND DIAGNOSTIC command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

**Table 149 — Extended physical sector format address descriptor (010b)**

Byte	Bit	7	6	5	4	3	2	1	0
0	(MSB)								
...	CYLINDER NUMBER								
2									
3	HEAD NUMBER								
4	MADS	Reserved				(MSB)			
...	SECTOR NUMBER								
7									

The CYLINDER NUMBER field contains the cylinder number (see 4.3.2).

The HEAD NUMBER field contains the head number (see 4.3.2).

A multi-address descriptor start (MADS) bit set to one specifies that this address descriptor defines the beginning of a defect that spans multiple addresses. The defect may span a number of sequential physical blocks on the same cylinder and head (i.e., a track) or may span a number of sequential tracks on the same head. A MADS bit set to zero specifies that:

- a) this address descriptor defines the end of a defect if the previous address descriptor has the MADS bit set to one; or
- b) this address descriptor defines a single track that contains one or more defects (i.e., the SECTOR NUMBER field contains FFF\_FFFFh) or a single defect (i.e., the SECTOR NUMBER field does not contain FFF\_FFFFh).

See 4.13.2 for valid combinations of two address descriptors that describe a defect.

The SECTOR NUMBER field:

- a) if not set to FFF\_FFFFh, contains the sector number (see 4.3.2) of the location being described; or
- b) if set to FFF\_FFFFh, specifies or indicates that the entire track is being described.

More than one logical block may be described by this address descriptor.



Table 150 defines the order of the fields used for sorting extended physical sector format address descriptors if the command using the address descriptors specifies sorting.

**Table 150 — Sorting order for extended physical sector format address descriptors**

Bit:	(MSB) 59	...	36	35	...	28	27	...	(LSB) 0
	CYLINDER NUMBER field			HEAD NUMBER field			SECTOR NUMBER field		

### 6.2.5 Long block format address descriptor

A format type of 011b specifies the long block format address descriptor (see table 151).

**Table 151 — Long block format address descriptor (011b)**

Bit	7	6	5	4	3	2	1	0
Byte								
0	(MSB)							
...	LONG BLOCK ADDRESS							
7	(LSB)							

For the FORMAT UNIT parameter list, the LONG BLOCK ADDRESS field specifies an eight-byte LBA. If the physical block containing the logical block referenced by the specified LBA contains additional logical blocks, then the device server may consider the LBAs of those additional logical blocks to also have been specified.

For the READ DEFECT DATA parameter data, the LONG BLOCK ADDRESS field indicates a vendor specific eight-byte value.

For the Translate Address diagnostic pages, the LONG BLOCK ADDRESS field contains:

- a) an eight-byte LBA, if the value is less than or equal to the capacity of the medium; or
- b) a vendor specific eight-byte, if the value is greater than the capacity of the medium.

### 6.2.6 Bytes from index format address descriptor

A format type of 100b specifies the bytes from index format address descriptor (see table 152). This address descriptor contains the location of a track or an offset from the start of a track.

**Table 152 — Bytes from index format address descriptor (100b)**

Bit	7	6	5	4	3	2	1	0
Byte								
0	(MSB)							
...	CYLINDER NUMBER							
2	(LSB)							
3	HEAD NUMBER							
4	(MSB)							
...	BYTES FROM INDEX							
7	(LSB)							

The CYLINDER NUMBER field contains the cylinder number (see 4.3.2).

The HEAD NUMBER field contains the head number (see 4.3.2).

The BYTES FROM INDEX field contains the number of bytes from the index (e.g., from the start of the track) to the location being described. A BYTES FROM INDEX field set to FFFF\_FFFFh specifies or indicates that the entire track is being described.

More than one logical block may be described by this address descriptor.

Table 153 defines the order of the fields used for sorting bytes from index format address descriptors if the command using the address descriptors specifies sorting.

**Table 153 — Sorting order for bytes from index format address descriptors**

Bit:	(MSB) 63	...	40	39	...	32	31	...	(LSB) 0
	CYLINDER NUMBER field			HEAD NUMBER field			BYTES FROM INDEX field		

### 6.2.7 Physical sector format address descriptor

A format type of 101b specifies the physical sector format address descriptor (see table 154). This address descriptor contains the location of a track or a sector (see 4.3.2).

**Table 154 — Physical sector format address descriptor (101b)**

Bit Byte	7	6	5	4	3	2	1	0
0	(MSB) _____ CYLINDER NUMBER _____ (LSB)							
...								
2								
3	HEAD NUMBER							
4	(MSB) _____ SECTOR NUMBER _____ (LSB)							
...								
7								

The CYLINDER NUMBER field contains the cylinder number (see 4.3.2).

The HEAD NUMBER field contains the head number (see 4.3.2).

The SECTOR NUMBER field contains the sector number (see 4.3.2). A SECTOR NUMBER field set to FFFF\_FFFFh specifies or indicates that the entire track is being described.

More than one logical block may be described by this address descriptor.

Table 155 defines the order of the fields used for sorting physical sector format address descriptors if the command using the address descriptors specifies sorting.

**Table 155 — Sorting order for physical sector format address descriptors**

Bit:	(MSB) 63	...	40	39	...	32	31	...	(LSB) 0
	CYLINDER NUMBER field			HEAD NUMBER field			SECTOR NUMBER field		

## 6.3 Diagnostic parameters

### 6.3.1 Diagnostic parameters overview

See table 156 for references to the pages and descriptors for diagnostic parameters used by direct access block devices.

The diagnostic pages and their corresponding page codes for direct access block devices are defined in table 156.

**Table 156 — Diagnostic page codes for direct access block devices**

Diagnostic page name	Page code	Reference
Diagnostic pages assigned by <del>SPC-4</del> SPC-5	30h to 3Fh	SPC-5
Rebuild Assist Input diagnostic page	42h	6.3.2
Rebuild Assist Output diagnostic page		6.3.3
SCSI enclosure services diagnostic pages	01h to 2Fh	SES-3
Supported Diagnostic Page diagnostic page	00h	SPC-5
Translate Address Input diagnostic page	40h	6.3.4
Translate Address Output diagnostic page		6.3.5
Obsolete	41h	
Vendor specific diagnostic pages	80h to FFh	
Reserved for this standard	43h to 7Fh	

### 6.3.2 Rebuild Assist Input diagnostic page

An application client sends a RECEIVE DIAGNOSTIC RESULTS command to retrieve a Rebuild Assist Input diagnostic page (see table 157), which provides information about whether the rebuild assist mode (see 4.20) is enabled or not and a device server's rebuild assist mode capabilities.

**Table 157 — Rebuild Assist Input diagnostic page**

Byte	Bit	7	6	5	4	3	2	1	0
0		PAGE CODE (42h)							
1		Reserved							
2	(MSB)	PAGE LENGTH (4 + (2 × n))							
3									
4		Reserved							ENABLED
5		Reserved							
6									
7		PHYSICAL ELEMENT LENGTH (n)							
8		DISABLED PHYSICAL ELEMENT MASK (if any)							
...									
7 + n									
8 + n		DISABLED PHYSICAL ELEMENT (if any)							
...									
7 + (2 × n)									

The PAGE CODE field and the PAGE LENGTH field are defined in SPC-5 and shall be set to the values defined in table 157.

An ENABLED bit set to one indicates that the rebuild assist mode is enabled. An ENABLED bit set to zero indicates that the rebuild assist mode is disabled.

The PHYSICAL ELEMENT LENGTH field indicates the length in bytes of the DISABLED PHYSICAL ELEMENT MASK field and the length in bytes of the DISABLED PHYSICAL ELEMENT field.

The bits in the DISABLED PHYSICAL ELEMENT MASK field indicate the bits in the DISABLED PHYSICAL ELEMENT field that are supported. Each bit set to one in the DISABLED PHYSICAL ELEMENT MASK field indicates that the corresponding bit in the DISABLED PHYSICAL ELEMENT field is supported and may be set to one in a Rebuild Assist Output diagnostic page sent with a SEND DIAGNOSTIC command.

The bits in the DISABLED PHYSICAL ELEMENT field indicate the physical elements that are disabled in this logical unit. Each bit set to one indicates that a physical element is disabled, and the device server shall report predicted read errors and predicted write errors for the associated group of LBAs.

### 6.3.3 Rebuild Assist Output diagnostic page

An application client sends a SEND DIAGNOSTIC command to send a Rebuild Assist Output diagnostic page (see table 158) that:

- a) enables or disables rebuild assist mode (see 4.20.2); and/or
- b) puts the logical unit in a simulated failure mode by disabling physical elements in conjunction with rebuild assist mode (see 4.20.5).

**Table 158 — Rebuild Assist Output diagnostic page**

Byte	Bit	7	6	5	4	3	2	1	0
0		PAGE CODE (42h)							
1		Reserved							
2	(MSB)	PAGE LENGTH (4 + (2 × n))							
3									
4		Reserved							ENABLE
5		Reserved							
6									
7		PHYSICAL ELEMENT LENGTH (n)							
8		DISABLED PHYSICAL ELEMENT MASK (if any)							
...									
7 + n									
8 + n		DISABLE PHYSICAL ELEMENT (if any)							
...									
7 + (2 × n)									

The PAGE CODE field and the PAGE LENGTH field are defined in SPC-5 and shall be set to the values defined in table 158.

An ENABLE bit set to one specifies that, after all fields in this diagnostic page have been validated:

- a) a self-test of the physical elements in the logical unit may be performed; and
- b) rebuild assist mode is enabled.

An ENABLE bit set to zero specifies that:

- a) rebuild assist mode shall be disabled;
- b) the other fields in this page shall be ignored; and
- c) all physical elements shall be enabled.

The PHYSICAL ELEMENT LENGTH field shall be set to the same value that is returned in the PHYSICAL ELEMENT LENGTH field ~~in the Rebuild Assist Input diagnostic page~~ (see 6.3.2).

If the PHYSICAL ELEMENT LENGTH field is not set to the same value that is returned in the PHYSICAL ELEMENT LENGTH field, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

The device server shall ignore the DISABLED PHYSICAL ELEMENT MASK field.

Each bit in the DISABLE PHYSICAL ELEMENT field specifies a physical element that shall be disabled. A bit set to one in the DISABLE PHYSICAL ELEMENT field specifies that the device server shall respond to read commands

and write commands specifying LBAs associated with that physical element as if the associated LBAs have predicted errors. A bit set to zero in the DISABLE PHYSICAL ELEMENT field specifies that the device server shall respond to read commands and write commands specifying LBAs associated with that physical element as if the associated LBAs do not have predicted errors. If the ENABLE bit is set to one, and the DISABLE PHYSICAL ELEMENT field specifies:

- a) any bits set to one that are not supported by the logical unit;
- b) all bits that are supported by the logical unit are set to one; or
- c) setting to zero any bits that are set to one,

then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

#### 6.3.4 Translate Address Input diagnostic page

Table 159 defines the Translate Address Input diagnostic page sent by a device server in response to a RECEIVE DIAGNOSTIC RESULTS command after a Translate Address Output diagnostic page (see 6.3.4) has been sent by an application client with the SEND DIAGNOSTIC command. If a Translate Address Output diagnostic page has not yet been processed by the device server, the results of a RECEIVE DIAGNOSTIC RESULTS command requesting this diagnostic page are vendor specific.

**Table 159 — Translate Address Input diagnostic page**

Byte	Bit	7	6	5	4	3	2	1	0
0		PAGE CODE (40h)							
1		Reserved							
2	(MSB)	PAGE LENGTH (n - 3)							
3									
4		Reserved				SUPPLIED FORMAT			
5		RAREA	ALTSEC	ALTTRK	Reserved		TRANSLATED FORMAT		
Translated address list									
6	(MSB)	TRANSLATED ADDRESS 1 (if any)							
...									
13									
		⋮							
n - 7	(MSB)	TRANSLATED ADDRESS x (if any)							
...									
n									

The PAGE CODE field is defined in SPC-5 and shall be set to the value shown in table 159 for the Translate Address Input diagnostic page.

The PAGE LENGTH field is defined in SPC-5.

The SUPPLIED FORMAT field contains the value from the SUPPLIED FORMAT field in the previous Translate Address Output diagnostic page (see 6.3.5).

A reserved area (RAREA) bit set to zero indicates that no part of the translated address falls within a reserved area of the medium (e.g., speed tolerance gap, alternate sector, or vendor reserved area). A RAREA bit set to one indicates that all or part of the translated address falls within a reserved area of the medium. If the entire translated address falls within a reserved area, then the device server may not return a translated address.

An alternate sector (ALTSEC) bit set to zero indicates that no part of the translated address is located in an alternate sector of the medium or that the device server is unable to determine this information. An ALTSEC bit set to one indicates that the translated address is located in an alternate sector of the medium. If the device server is unable to determine if all or part of the translated address is located in an alternate sector, then the device server shall set this bit to zero.

An alternate track (ALTTRK) bit set to zero indicates that no part of the translated address is located on an alternate track of the medium. An ALTTRK bit set to one indicates that part or all of the translated address is located on an alternate track of the medium or the device server is unable to determine if all or part of the translated address is located on an alternate track.

The TRANSLATED FORMAT field contains the value from the TRANSLATE FORMAT field in the previous Translate Address Output diagnostic page (see 6.3.4).

The TRANSLATED ADDRESS field(s) contains the address descriptor (see 6.2) that the device server translated from the address descriptor supplied by the application client in the previous Translate Address Output diagnostic page (see 6.3.5). Each field shall be in the format (see 6.2) specified in the TRANSLATED FORMAT field. If the short block format address descriptor (see 6.2.2) is specified, then the first four bytes of the TRANSLATED ADDRESS field shall contain the short block format address descriptor and the last four bytes shall contain 0000\_0000h.

If the returned data is in short block format (see 6.2.2), long block format (see 6.2.5), or physical sector format (see 6.2.7) and the ADDRESS TO TRANSLATE field in the previous Translate Address Output diagnostic page covers more than one address after it has been translated (e.g., because of multiple physical sectors within a single logical block or multiple logical blocks within a single physical sector), then the device server shall return all possible addresses that are contained in the area specified by the address to be translated. If the returned data is in bytes from index format (see 6.2.6), the device server shall return a pair of translated values for each of the possible addresses that are contained in the area specified by the ADDRESS TO TRANSLATE field in the previous Translate Address Output diagnostic page. Of the pair of translated values returned, the first indicates the starting location and the second the ending location of the area.

### 6.3.5 Translate Address Output diagnostic page

The Translate Address diagnostic pages provides a method for an application client to have a device server translate an address descriptor (see 6.2) from one format to another. The address descriptor to be translated is sent to the device server in the Translate Address Output diagnostic page with a SEND DIAGNOSTIC command and the results are returned by the device server in the Translate Address Input diagnostic page sent in response to a RECEIVE DIAGNOSTIC RESULTS command.

Table 160 defines the format of the Translate Address Output diagnostic page sent with the SEND DIAGNOSTIC command. The translated address returned in the Translate Address Input diagnostic page is defined in 6.3.4.

**Table 160 — Translate Address Output diagnostic page**

Byte	Bit	7	6	5	4	3	2	1	0
0		PAGE CODE (40h)							
1		Reserved							
2	(MSB)	PAGE LENGTH (000Ah)							
3									
4		Reserved					SUPPLIED FORMAT		
5		Reserved					TRANSLATE FORMAT		
6	(MSB)	ADDRESS TO TRANSLATE							
...									
13									

The PAGE CODE field and PAGE LENGTH field are defined in SPC-5 and shall be set to the values shown in table 160 for the Translate Address Output diagnostic page.

The SUPPLIED FORMAT field specifies the format (see 6.2) of the ADDRESS TO TRANSLATE field. If the device server does not support the requested format, then the device server shall terminate the SEND DIAGNOSTIC command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

The TRANSLATE FORMAT field specifies the format (see 6.2) the device server shall use for the result of the address translation. If the device server does not support the specified format, then the device server shall terminate the SEND DIAGNOSTIC command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

The ADDRESS TO TRANSLATE field contains a single address descriptor that the application client is requesting the device server to translate. The format of this field depends on the value in the SUPPLIED FORMAT field. The formats are described in 6.2. If the short block format address descriptor is specified, then the first four bytes of the ADDRESS TO TRANSLATE field shall contain the short block format address descriptor and the last four bytes shall contain 0000\_0000h.



## **6.4 Log parameters**

### **6.4.1 Log parameters overview**

#### **6.4.1.1 Summary of log pages**

See table 161 for references to the log pages and their corresponding page codes and subpage codes for direct access block devices. See SPC-5 for a detailed description of logging operations.

Table 161 — Log page codes and subpage codes for direct access block devices

Log page name	Page code <sup>a</sup>	Subpage code <sup>a</sup>	Reference
Application Client <del>-log page</del>	0Fh	00h	SPC-5
ATA PASS-THROUGH Results	16h	00h	SAT-3
Background Scan Results <del>-log page</del>	15h	00h	6.4.2
Background Operation <del>-log page</del>	15h	02h	6.4.3
Buffer Over-Run/Under-Run <del>-log page</del>	01h	00h	SPC-5
<a href="#">Cache Memory Statistics</a>	<a href="#">19h</a>	<a href="#">20h</a>	<a href="#">SPC-5</a>
<a href="#">Environmental Limits</a>	<a href="#">0Dh</a>	<a href="#">02h</a>	<a href="#">SPC-5</a>
<a href="#">Environmental Reporting</a>	<a href="#">0Dh</a>	<a href="#">01h</a>	<a href="#">SPC-5</a>
Format Status <del>-log page</del>	08h	00h	6.4.4
<a href="#">General Statistics and Performance</a>	<a href="#">19h</a>	<a href="#">00h</a>	<a href="#">SPC-5</a>
<a href="#">Group Statistics and Performance (1 to 31)</a>	<a href="#">19h</a>	<a href="#">01h to 1Fh</a>	<a href="#">SPC-5</a>
Informational Exceptions <del>-log page</del>	2Fh	00h	SPC-5
Last n Deferred Errors Or Asynchronous Events <del>-log page</del>	0Bh	00h	SPC-5
Last n Error Events <del>-log page</del>	07h	00h	SPC-5
Logical Block Provisioning <del>-log page</del>	0Ch	00h	6.4.5
<a href="#">LPS Misalignment</a>	<a href="#">15h</a>	<a href="#">03h</a>	<a href="#">6.4.6</a>
Non-Medium Error <del>-log page</del>	06h	00h	SPC-5
Non-volatile Cache <del>-log page</del>	17h	00h	6.4.7
Pending Defects <del>-log page</del>	15h	01h	6.4.8
<a href="#">Power Condition Transitions</a>	<a href="#">1Ah</a>	<a href="#">00h</a>	<a href="#">SPC-5</a>
Protocol-Specific Port <del>-log pages</del>	18h	00h to FEh	SPC-5
Read Error Counters <del>-log page</del>	03h	00h	SPC-5
Self-Test Results <del>-log page</del>	10h	00h	SPC-5
Solid State Media <del>-log page</del>	11h	00h	6.4.9
Start-Stop Cycle Counter <del>-log page</del>	0Eh	00h	SPC-5
Supported Log Pages <del>-log page</del>	00h	FFh	SPC-5
Supported Log Pages and Subpages <del>-log page</del>	00h	00h	SPC-5
Supported Subpages	01h to 3Fh	FFh	SPC-5
Temperature <del>-log page</del>	0Dh	00h	SPC-5
Utilization <del>-log page</del>	0Eh	01h	6.4.10
Verify Error Counters <del>-log page</del>	05h	00h	SPC-5
Write Error Counters <del>-log page</del>	02h	00h	SPC-5
Vendor specific	30h to 3Eh	00h to FEh	n/a
<sup>a</sup> All page code and subpage code combinations not shown in this table are reserved.			

### 6.4.1.2 Setting and resetting log parameters

In a LOG SELECT command (see SPC-5), an application client may specify that:

- a) all the parameters in a log page or pages are to be reset (i.e., the PCR bit set to one and the PARAMETER LIST LENGTH field is set to zero); or
- b) individual parameters in log page are to be changed to specified new values (i.e., the PCR bit is set to zero and the PARAMETER LIST LENGTH field is not set to zero).

The device server processing of LOG SELECT commands (see SPC-5) that request changes to individual log parameters or reset all log parameters depend on the log parameter that is being changed or reset, and is specified in the table that defines the log parameter using the keywords defined in table 162 (also see SPC-5).

**Table 162 — Keywords for resetting or changing log parameters**

Keyword	Device server processing when:	
	PCR bit is set to one <sup>a</sup>	PCR bit is set to zero <sup>b</sup>
Always	Reset the log parameter.	Change the log parameter.
Reset Only	Reset the log parameter.	If any changes are requested in the PARAMETER VALUE field of the log parameter, then:
Never	Do not reset the log parameter; see the LOG SELECT command in SPC-5 for description of possible error conditions.	<ul style="list-style-type: none"> <li>a) terminate the command with CHECK CONDITION status, with the sense key set to ILLEGAL REQUEST, and the additional sense code set to INVALID FIELD IN PARAMETER LIST; and</li> <li>b) do not make any requested changes in any field in any log parameter in any log page</li> </ul>
<sup>a</sup> If the PCR bit is set to one, and the PARAMETER LIST LENGTH field is not set to zero, then the device server shall terminate the LOG SELECT command (see SPC-5). <sup>b</sup> If the PCR bit is set to zero, and the PARAMETER LIST LENGTH field is set to zero. then no log parameters are changed (see SPC-5).		

### 6.4.2 Background Scan log page

#### 6.4.2.1 Background Scan log page overview

Using the format shown in table 164, the Background Scan log page reports information about:

- a) background pre-scan operations (see 4.24.2) and background medium scan operations (see 4.24.3); and
- b) any logical blocks where an error was detected during a background scan operation.

The parameter codes for the Background Scan log page are defined in table 163.

**Table 163 — Background Scan log page parameter codes**

Parameter code	Description	Resettable or Changeable <sup>a</sup>	Reference	Support- <del>required</del>
0000h	Background Scan Status	Never	6.4.2.2	Mandatory
0001h to 0800h	Background Scan Results	Reset Only	6.4.2.3	Optional <sup>b</sup>
8000h to AFFFh	Vendor specific		n/a	Optional
All others	Reserved			
<sup>a</sup> The keywords in this column – Always, Reset Only, and Never – are defined in 6.4.1.2. <sup>b</sup> If the Background Scan log page is supported, then at least one Background Scan Results log parameter shall be supported.				

The Background Scan log page has the format defined in table 164.

**Table 164 — Background Scan log page**

Byte	Bit	7	6	5	4	3	2	1	0
0		DS (1b)	SPF (0b)	PAGE CODE (15h)					
1		SUBPAGE CODE (00h)							
2		(MSB)	PAGE LENGTH (n - 3)						
3									
Background scan parameters									
4		Background scan parameter [first] (if any)							
...									
		⋮							
		Background scan parameter [last] (if any)							
...									
n									

The disable save (DS) bit, the subpage format (SPF) bit, the PAGE CODE field, the SUBPAGE CODE field, and the PAGE LENGTH field are described in SPC-5.

The DS bit, the SPF bit, the PAGE CODE field, and the SUBPAGE CODE field shall be set to the values shown in table 164 for the Background Scan log page.

If the device server processes a LOG SELECT command with the PCR bit set to one (see SPC-5), then the device server shall:

- a) not change the values in the Background Scan Status log parameter; and
- b) delete all Background Scan Results log parameters.

### 6.4.2.2 Background Scan Status log parameter

The Background Scan Status log parameter for the Background Scan log page has the format defined in table 165.

**Table 165 — Background Scan Status log parameter format**

Byte	Bit	7	6	5	4	3	2	1	0
0	(MSB)	PARAMETER CODE (0000h)							
1									
2	Parameter control byte – binary format list log parameter (see SPC-5)								
	DU	Obsolete	TSD	Obsolete			FORMAT AND LINKING		
3		PARAMETER LENGTH (0Ch)							
4	(MSB)	ACCUMULATED POWER ON MINUTES							
...									
7									
8		Reserved							
9		BACKGROUND SCAN STATUS							
10	(MSB)	NUMBER OF BACKGROUND SCANS PERFORMED							
11									
12	(MSB)	BACKGROUND SCAN PROGRESS							
13									
14	(MSB)	NUMBER OF BACKGROUND MEDIUM SCANS PERFORMED							
15									

The PARAMETER CODE field is described in SPC-5 and shall be set to the value shown in table 165 for the Background Scan Status log parameter.

The DU bit, the TSD bit, and the FORMAT AND LINKING field for the Background Scan Status log parameter shall be set for a binary format list log parameter as described in SPC-5.

The PARAMETER LENGTH field is described in SPC-5 and shall be set to the value shown in table 165 for the Background Scan Status log parameter.

The ACCUMULATED POWER ON MINUTES field indicates the number of minutes the device server has been powered on since manufacturing.

Table 166 defines the BACKGROUND SCAN STATUS field.

**Table 166 — BACKGROUND SCAN STATUS field**

Code	Description
00h	No background scan operation is active.
01h	A background medium scan operation is active.
02h	A background pre-scan operation is active.
03h	A background scan operation was halted due to a fatal error.
04h	A background scan operation was halted due to a vendor specific pattern of errors.
05h	A background scan operation was halted due to the medium being formatted without the PLIST.
06h	A background scan operation was halted due to a vendor specific cause.
07h	A background scan operation was halted due to the temperature being out of the allowed range.
08h	Background medium scan operations are enabled (i.e., the EN_BMS bit is set to one in the Background Control mode page (see 6.5.4)), and no background medium scan operation is active (i.e., the device server is waiting for Background Medium Scan Interval timer expiration before starting the next background medium scan operation).
09h	A background scan operation was halted due to the S_L_FULL bit being set to one in the Background Control mode page (see 6.5.4) and the background scan results list being full.
0Ah	A background pre-scan operation was halted due to the Background Pre-scan Time Limit timer expiring.
0Bh to FFh	Reserved

The NUMBER OF BACKGROUND SCANS PERFORMED field indicates the number of background scan operations (i.e., the total number of background pre-scan operations plus the number of background medium scan operations) that have been performed since the SCSI target device was shipped by the manufacturer.

The BACKGROUND SCAN PROGRESS field indicates the percent complete of a background scan operation in progress. The returned value is a numerator that has 65 536 (i.e., 1\_0000h) as its denominator. If there is no background scan operation in progress (i.e., no background scan operation has been initiated since power on or the most recent background scan operation has completed), then the device server shall set the BACKGROUND SCAN PROGRESS field to 0000h.

The NUMBER OF BACKGROUND MEDIUM SCANS PERFORMED field indicates the number of background medium scan operations that have been performed since the SCSI target device was shipped by the manufacturer. If the NUMBER OF BACKGROUND MEDIUM SCANS PERFORMED field contains 0000h, then the number of background medium scan operations is unknown.

The total number of background pre-scan operations that have been performed is the value in the NUMBER OF BACKGROUND SCANS PERFORMED field minus the value in the NUMBER OF BACKGROUND MEDIUM SCANS PERFORMED field.

### 6.4.2.3 Background Scan Results log parameter

The Background Scan Results log parameter for the Background Scan log page has the format defined in table 167. If the Background Scan log page is reset, then all Background Scan Results log parameters are discarded. If no errors have occurred during a background scan or the Background Scan log page has been reset, then no Background Scan Results log parameters shall be present.

**Table 167 — Background Scan Results log parameter format**

Byte	Bit	7	6	5	4	3	2	1	0	
0	(MSB)	PARAMETER CODE (0001h to 0800h)								
1										(LSB)
2	Parameter control byte – binary format list log parameter (see SPC-5)									
	DU	Obsolete	TSD	Obsolete			FORMAT AND LINKING			
3	PARAMETER LENGTH (14h)									
4	(MSB)	ACCUMULATED POWER ON MINUTES								
...										
7										(LSB)
8	REASSIGN STATUS					SENSE KEY				
9	ADDITIONAL SENSE CODE									
10	ADDITIONAL SENSE CODE QUALIFIER									
11	Vendor specific									
...										
15										
16	(MSB)	LOGICAL BLOCK ADDRESS								
...										
23										(LSB)

The PARAMETER CODE field is described in SPC-5 and shall be set to a value from 0001h through 0800h in sequence as errors are discovered during a background scan operation. When all of the supported parameter code values have been used, and a new error is discovered during a background scan operation, the oldest Background Scan Results log parameter in the list (i.e., the Background Scan Results log parameter with the smallest value in the ACCUMULATED POWER ON MINUTES field) shall be discarded, and the PARAMETER CODE field (see 6.4.2.3) for the new defect shall be set to the parameter code value of the discarded Background Scan Results log parameter.

The DU bit, the TSD bit, and the FORMAT AND LINKING field for a Background Scan Results log parameter shall be set for a binary format list log parameter as described in SPC-5.

The PARAMETER LENGTH field is defined in SPC-5 and shall be set to the value shown in table 167 for the Background Scan Results log parameter.

The ACCUMULATED POWER ON MINUTES field indicates the number of minutes that the device server has been powered on since manufacturing at the time the background scan error reported in the Background Scan Results log parameter occurred.

Table 168 defines the REASSIGN STATUS field.

**Table 168 — REASSIGN STATUS field (part 1 of 2)**

Code	Reason			
	LOWIR bit <sup>a</sup>		Original error <sup>b</sup>	Additional conditions
	0	1		
1h	Yes	Yes	Recovered or unrecovered	The LBA has not yet been reassigned. <sup>c</sup>
2h	Yes	No	Recovered	The device server performed automatic read reassignment for the LBA (i.e., performed a reassign operation for the LBA and a write operation with recovered logical block data). <sup>d</sup>
4h	Yes	Yes	Recovered	The device server's attempt to perform automatic read reassignment failed. The logical block may or may not now have an uncorrectable error. <sup>c</sup>
5h	Yes	No	Recovered	The error was corrected by the device server rewriting the logical block without performing a reassign operation.
6h	Yes	Yes	Recovered or unrecovered	Either: a) an application client caused automatic write reassignment for the LBA with a command performing a write operation; or b) the LBPRZ field is set to xx1b in the Logical Block Provisioning VPD page (see 6.6.6), and an application client caused an unmap operation for the LBA. <sup>c</sup>
7h	Yes	Yes	Recovered or unrecovered	Either: a) an application client caused a reassign operation for the LBA with a REASSIGN BLOCKS command; or b) the LBPRZ field is set to 000b or is set to 010b in the Logical Block Provisioning VPD page (see 6.6.6), and an application client caused an unmap operation for the LBA. <sup>c</sup>
Key:				
Yes = specifies that a Background Scan Results log parameter shall be generated for the error.				
No = specifies that a Background Scan Results log parameter shall not be generated for the error				
<sup>a</sup> The LOWIR bit (see 6.5.4). <sup>b</sup> Type of error detected while reading the logical block referenced by the LBA specified by the LOGICAL BLOCK ADDRESS field(see 6.4.2.3) during a background scan operation. <sup>c</sup> The REASSIGN STATUS field in a given log parameter changes from 1h or 4h to 6h, 7h, or 8h when a reassign operation, write operation, or unmap operation on the LBA succeeds or when a reassign operation on the LBA fails. After the LBA is reassigned, any subsequent medium error occurring for the LBA is reported in a new log parameter with the same value in the LOGICAL BLOCK ADDRESS field as the value in the LOGICAL BLOCK ADDRESS field in the log parameter for the previous medium error for the LBA. <sup>d</sup> The ARRE bit (see 6.5.10) controls automatic read reassignment based on errors detected during all read medium operations, including those that are part of background scan operations.				



**Table 168 — REASSIGN STATUS field (part 2 of 2)**

Code	Reason			
	LOWIR bit <sup>a</sup>		Original error <sup>b</sup>	Additional conditions
	0	1		
8h	Yes	Yes	Recovered or unrecovered	An application client's request for a reassign operation for the LBA with a REASSIGN BLOCKS command failed. The logical block referenced by the LBA may or may not still have an uncorrectable error.
All others	Reserved			
Key:				
Yes = specifies that a Background Scan Results log parameter shall be generated for the error.				
No = specifies that a Background Scan Results log parameter shall not be generated for the error				
<sup>a</sup> The LOWIR bit (see 6.5.4).				
<sup>b</sup> Type of error detected while reading the logical block referenced by the LBA specified by the LOGICAL BLOCK ADDRESS field(see 6.4.2.3) during a background scan operation.				
<sup>c</sup> The REASSIGN STATUS field in a given log parameter changes from 1h or 4h to 6h, 7h, or 8h when a reassign operation, write operation, or unmap operation on the LBA succeeds or when a reassign operation on the LBA fails. After the LBA is reassigned, any subsequent medium error occurring for the LBA is reported in a new log parameter with the same value in the LOGICAL BLOCK ADDRESS field as the value in the LOGICAL BLOCK ADDRESS field in the log parameter for the previous medium error for the LBA.				
<sup>d</sup> The ARRE bit (see 6.5.10) controls automatic read reassignment based on errors detected during all read medium operations, including those that are part of background scan operations.				

If sense data is available, then the device server shall set the SENSE KEY field, the ADDITIONAL SENSE CODE field, and the ADDITIONAL SENSE CODE QUALIFIER field to a hierarchy of additional information relating to error conditions that occurred during the background scan operation. The content of these fields is represented in the same format used by the sense data (see SPC-5).

The LOGICAL BLOCK ADDRESS field indicates the LBA associated with the medium error.

### 6.4.3 Background Operation log page

#### 6.4.3.1 Background Operation log page overview

Using the format shown in table 169, the Background Operation log page reports parameters that are specific to background operations.

**Table 169 — Background Operation log page**

Byte	Bit	7	6	5	4	3	2	1	0
0		DS	SPF (1b)	PAGE CODE (15h)					
1		SUBPAGE CODE (02h)							
2		(MSB)	PAGE LENGTH (n - 3)						
3									
Background Operation parameters									
4		Background Operation parameter [first] (if any)							
...									
		⋮							
		Background Operation parameter [last] (if any)							
...									
n									

The disable save (DS) bit, the subpage format (SPF) bit, the PAGE CODE field, the SUBPAGE CODE field, and the PAGE LENGTH field are described in SPC-5.

The SPF bit, the PAGE CODE field, and the SUBPAGE CODE field shall be set to the values shown in table 169 for the Background Operation log page.

The parameter codes for the Background Operation log page are listed in table 170.

**Table 170 — Background Operation log page parameter codes**

Parameter code	Description	Resettable or Changeable <sup>a</sup>	Reference	Support-Required
0000h	Background Operation	Never	6.4.3.2	Mandatory
All others	Reserved			
<sup>a</sup> The keywords in this column – Always, Reset Only, and Never – are defined in 6.4.1.2.				

### 6.4.3.2 Background Operation log parameter

The Background Operation log parameter of the Background Operation log page has the format defined in table 171.

**Table 171 — Background Operation log parameter format**

Byte	Bit	7	6	5	4	3	2	1	0
0	(MSB)	PARAMETER CODE (0000h)							
1		(LSB)							
2	Parameter control byte – binary format list log parameter (see SPC-5)								
	DU	Obsolete	TSD	Obsolete			FORMAT AND LINKING		
3	PARAMETER LENGTH (4h)								
4	BO_STATUS								
5	Reserved								
...									
n									

The PARAMETER CODE field is described in SPC-5 and shall be set to the value shown in table 171 for the Background Operation log parameter.

The DU bit, the TSD bit, and the FORMAT AND LINKING field for the Background Operation log parameter shall be set for a binary format list log parameter as described in SPC-5.

The PARAMETER LENGTH field is described in SPC-5 and shall be set to the value shown in table 171 for the Background Operation log parameter.

The background operation status (BO\_STATUS) field indicates the type of background operation, if any, that is being performed by the device server as defined in table 172.

**TABLE 172 — BO\_STATUS DEFINITIONS**

Code	Description
00h	No indication
01h	No advanced background operation being performed
02h	Host initiated advanced background operation being performed
03h	Device initiated advanced background operation being performed
All others	Reserved

## 6.4.4 Format Status log page

### 6.4.4.1 Format Status log page overview

Using the format shown table 174, the Format Status log page reports information about the most recent successful format operation and the state of the direct access block device since that operation was performed. The parameter codes for the Format Status log page are listed in table 173.

**Table 173 — Format Status log page parameter codes**

Parameter code	Description	Resettable or Changeable <sup>a</sup>	Reference	Support-Required
0000h	Format Data Out	Never	6.4.4.2	Mandatory
0001h	Grown Defects During Certification	Never	6.4.4.3	Mandatory
0002h	Total Blocks Reassigned During Format	Never	6.4.4.4	Mandatory
0003h	Total New Blocks Reassigned	Never	6.4.4.5	Mandatory
0004h	Power On Minutes Since Format	Never	6.4.4.6	Mandatory
0005h to 7FFFh	Reserved			
8000h to FFFFh	Vendor specific			Optional
<sup>a</sup> The keywords in this column – Always, Reset Only, and Never – are defined in 6.4.1.2.				

The Format Status log page has the format defined in table 174.

**Table 174 — Format Status log page**

Byte	Bit	7	6	5	4	3	2	1	0
0		DS (1b)	SPF (0b)	PAGE CODE (08h)					
1		SUBPAGE CODE (00h)							
2		(MSB)	PAGE LENGTH (n - 3)						
3									(LSB)
Format Status log parameters									
4		Format Status log parameter [first]							
...									
		Format Status log parameter [last]							
...									
n									

The disable save (DS) bit, the subpage format (SPF) bit, the PAGE CODE field, the SUBPAGE CODE field, and the PAGE LENGTH field are described in SPC-5.

The DS bit, the SPF bit, the PAGE CODE field, and the SUBPAGE CODE field shall be set to the values shown in table 174 for the Format Status log page.

If a format operation has never been performed by the logical unit, then the log parameter for each Format Status log parameter listed in table 173 is not defined by this standard. If a device server begins a format operation, then the device server shall set each byte of the log parameter data (i.e., bytes four to n of the log parameter), if any, to FFh for each Format Status log parameter (e.g., if the PARAMETER LENGTH field is set to 02h, then the log parameter data is set to FFFFh).

If the most recent format operation failed or the information for a Format Status log parameter is not available, then the device server shall return FFh in each byte of the log parameter data (i.e., bytes four to n of the log parameter), if any, for the Format Status log parameter (e.g., if the PARAMETER LENGTH field is set to 04h, then the log parameter data shall be set to FFFF\_FFFFh). The device server shall set each Format Status log parameter to be a multiple of four bytes.

#### 6.4.4.2 Format Data Out log parameter

The Format Data Out log parameter of the Format Status log page has the format defined in table 175.

**Table 175 — Format Data Out log parameter format**

Bit	7	6	5	4	3	2	1	0
Byte								
0	(MSB) _____							
1	PARAMETER CODE (0000h) _____ (LSB)							
2	Parameter control byte – binary format list log parameter (see SPC-5)							
	DU	Obsolete	TSD	Obsolete		FORMAT AND LINKING		
3	PARAMETER LENGTH (n - 3)							
4	(MSB) _____							
...	FORMAT DATA OUT _____							
n	(LSB)							

The PARAMETER CODE field is described in SPC-5 and shall be set to the value shown in table 175 for the Format Data Out log parameter.

The DU bit, and the FORMAT AND LINKING field for the Format Data Out log parameter shall be set for a binary format list log parameter as described in SPC-5.

The target save disable (TSD) bit (see SPC-5) shall be set to zero for the Format Data Out log parameter, indicating that the logical unit saves the Format Data Out log parameter at vendor specific intervals without any request from an application client.

The PARAMETER LENGTH field is described in SPC-5.

After a successful format operation, the FORMAT DATA OUT field contains the FORMAT UNIT parameter list (see 5.4.2).

#### 6.4.4.3 Grown Defects During Certification log parameter

The Grown Defects During Certification log parameter for the Format Status log page has the format defined in table 176.

**Table 176 — Grown Defects During Certification log parameter format**

Byte	Bit	7	6	5	4	3	2	1	0
0	(MSB)	PARAMETER CODE (0001h)							
1									
2	Parameter control byte – binary format list log parameter (see SPC-5)								
	DU	Obsolete	TSD	Obsolete			FORMAT AND LINKING		
3	PARAMETER LENGTH (08h)								
4	(MSB)	GROWN DEFECTS DURING CERTIFICATION							
...									
11	(LSB)								

The PARAMETER CODE field is described in SPC-5 and shall be set to the value shown in table 176 for the Grown Defects During Certification log parameter.

The DU bit, and the FORMAT AND LINKING field for the Grown Defects During Certification log parameter shall be set for a binary format list log parameters as described in SPC-5.

The target save disable (TSD) bit (see SPC-5) shall be set to zero for the Grown Defects During Certification log parameter, indicating that the logical unit saves the Grown Defects During Certification log parameter at vendor specific intervals without any request from an application client.

The PARAMETER LENGTH field is described in SPC-5 and shall be set to the value shown in table 176 for the Grown Defects During Certification log parameter.

After a successful format operation during which certification was performed, the GROWN DEFECTS DURING CERTIFICATION field shall indicate the number of defects detected as a result of performing the certification. The value in the GROWN DEFECTS DURING CERTIFICATION field count reflects only those defects detected and replaced during the successful format operation that were not already part of the PLIST or GLIST.

After a successful format operation during which certification was not performed, the GROWN DEFECTS DURING CERTIFICATION field shall be set to zero.

#### 6.4.4.4 Total Blocks Reassigned During Format log parameter

The Total Blocks Reassigned During Format log parameter for the Format Status log page has the format defined in table 177.

**Table 177 — Total Blocks Reassigned During Format log parameter format**

Byte	Bit	7	6	5	4	3	2	1	0
0	(MSB)	PARAMETER CODE (0002h)							
1									
2	Parameter control byte – binary format list log parameter (see SPC-5)								
	DU	Obsolete	TSD	Obsolete			FORMAT AND LINKING		
3	PARAMETER LENGTH (08h)								
4	(MSB)	TOTAL BLOCKS REASSIGNED DURING FORMAT							
...									
11									

The PARAMETER CODE field is described in SPC-5 and shall be set to the value shown in table 177 for the Total Blocks Reassigned During Format log parameter.

The DU bit, and the FORMAT AND LINKING field for the Total Blocks Reassigned During Format log parameter shall be set for a binary format list log parameters described in SPC-5.

The target save disable (TSD) bit (see SPC-5) shall be set to zero for the Total Blocks Reassigned During Format log parameter, indicating that the logical unit saves the Total Blocks Reassigned During Format log parameter at vendor specific intervals without any request from an application client.

The PARAMETER LENGTH field is described in SPC-5 and shall be set to the value shown in table 177 for the Total Blocks Reassigned During Format log parameter.

The TOTAL BLOCKS REASSIGNED DURING FORMAT field contains the count of the total number of logical blocks that were reassigned during the most recent successful format operation.

#### 6.4.4.5 Total New Blocks Reassigned log parameter

The Total New Blocks Reassigned log parameter for the Format Status log page has the format defined in table 178.

**Table 178 — Total New Blocks Reassigned log parameter format**

Byte	Bit	7	6	5	4	3	2	1	0
0	(MSB)	PARAMETER CODE (0003h)							
1		(LSB)							
2	Parameter control byte – binary format list log parameter (see SPC-5)								
	DU	Obsolete	TSD	Obsolete			FORMAT AND LINKING		
3	PARAMETER LENGTH (08h)								
4	(MSB)	TOTAL NEW BLOCKS REASSIGNED							
...									
11									

The PARAMETER CODE field is described in SPC-5 and shall be set to the value shown in table 178 for the Total New Blocks Reassigned log parameter.

The DU bit, and the FORMAT AND LINKING field for the Total New Blocks Reassigned log parameter shall be set for a binary format list log parameters described in SPC-5.

The target save disable (TSD) bit (see SPC-5) shall be set to zero for the Total New Blocks Reassigned log parameter, indicating that the logical unit saves the Total New Blocks Reassigned log parameter at vendor specific intervals without any request from an application client.

The PARAMETER LENGTH field is described in SPC-5 and shall be set to the value shown in table 178 for the Total New Blocks Reassigned log parameter.

The TOTAL NEW BLOCKS REASSIGNED field contains a count of the total number of logical blocks that have been reassigned since the completion of the most recent successful format operation.



#### 6.4.4.6 Power On Minutes Since Format log parameter

The Power On Minutes Since Format log parameter for the Format Status log page has the format defined in table 179.

**Table 179 — Power On Minutes Since Format log parameter format**

Bit	7	6	5	4	3	2	1	0
Byte								
0	(MSB) _____							
1	PARAMETER CODE (0004h) _____							(LSB)
2	Parameter control byte – binary format list log parameter (see SPC-5)							
	DU	Obsolete	TSD	Obsolete		FORMAT AND LINKING		
3	PARAMETER LENGTH (04h) _____							
4	(MSB) _____							
...	POWER ON MINUTES SINCE FORMAT _____							(LSB)
7								

The PARAMETER CODE field is described in SPC-5 and shall be set to the value shown in table 179 for the Power On Minutes Since Format log parameter.

The DU bit, and the FORMAT AND LINKING field for the Power On Minutes Since Format log parameter shall be set for a binary format list log parameter as described in SPC-5.

The target save disable (TSD) bit (see SPC-5) shall be set to zero for the Power On Minutes Since Format log parameter, indicating that the logical unit saves the Power On Minutes Since Format log parameter at vendor specific intervals without any request from an application client.

The PARAMETER LENGTH field is described in SPC-5 and shall be set to the value shown in table 179 for the Power On Minutes Since Format log parameter.

The POWER ON MINUTES SINCE FORMAT field contains the unsigned number of usage minutes (i.e., minutes with power applied regardless of power state) that have elapsed since the most recent successful format operation.

## 6.4.5 Logical Block Provisioning log page

### 6.4.5.1 Logical Block Provisioning log page overview

Using the format defined in table 181, the Logical Block Provisioning log page reports the logical block provisioning status of the logical unit. The parameter codes for the Logical Block Provisioning log page are listed in table 180.

**Table 180 — Logical Block Provisioning log parameters**

Parameter code <sup>a</sup>	Description	Resettable or Changeable <sup>b</sup>	Reference	Support-Required
Resources that are associated with thresholds (0000h to 00FFh)				
0000h	Reserved			
0001h	Available LBA Mapping Resource Count	Never	6.4.5.2	Optional <sup>c</sup>
0002h	Used LBA Mapping Resource Count	Never	6.4.5.3	
0003h	Available Provisioning Resource Percentage	Never	6.4.5.4	
0003h to 00FFh	Reserved			
Resources that are not associated with thresholds (0000h to 00FFh)				
0100h	De-duplicated LBA Resource Count	Never	6.4.5.5	Optional
0101h	Compressed LBA Resource Count	Never	6.4.5.6	
0102h	Total Efficiency LBA Resource Count	Never	6.4.5.7	
0103h to FFEFh	Reserved			
FFF0h to FFFFh	Vendor specific			
<sup>a</sup> Parameter codes 0000h to 00FFh are coordinated with the THRESHOLD RESOURCE field (see 6.5.9). <sup>b</sup> The keywords in this column – Always, Reset Only, and Never – are defined in 6.4.1.2. <sup>c</sup> If this log page is supported, then at least one parameter shall be supported. A logical block provisioning log parameter in the range 0001h to 00FFh should be provided to report resource usage for each threshold resource for which a threshold descriptor in the Logical Block Provisioning mode page (see 6.5.9) is available.				

The Logical Block Provisioning log page has the format defined in table 181.

**Table 181 — Logical Block Provisioning log page**

Byte	Bit	7	6	5	4	3	2	1	0
0		DS (1b)	SPF (0b)	PAGE CODE (0Ch)					
1		SUBPAGE CODE (00h)							
2		(MSB)	PAGE LENGTH (n - 3)						
3									
Logical block provisioning parameter list									
4		Logical block provisioning log parameter [first]							
...									
⋮									
		Logical block provisioning log parameter [last]							
...									
n									

The disable save (DS) bit, the subpage format (SPF) bit, the PAGE CODE field, the SUBPAGE CODE field, and the PAGE LENGTH field are described in SPC-5.

The DS bit, the SPF bit, the PAGE CODE field, and the SUBPAGE CODE field shall be set to the values shown in table 181 for the Logical Block Provisioning log page.

### 6.4.5.2 Available LBA Mapping Resource Count log parameter

#### 6.4.5.2.1 Available LBA Mapping Resource Count log parameter overview

The Available LBA Mapping Resource Count log parameter of the Logical Block Provisioning log page has the format defined in table 182.

**Table 182 — Available LBA Mapping Resource Count log parameter format**

Bit	7	6	5	4	3	2	1	0
Byte								
0	(MSB) _____							
1	PARAMETER CODE (0001h) _____ (LSB)							
2	Parameter control byte – binary format list log parameter (see SPC-5)							
	DU	Obsolete	TSD	Obsolete			FORMAT AND LINKING	
3	PARAMETER LENGTH (08h) _____							
4	(MSB) _____							
...	RESOURCE COUNT _____							
7	_____ (LSB)							
8	Reserved						SCOPE	
9	_____							
...	Reserved _____							
11	_____							

The PARAMETER CODE field is described in SPC-5 and shall be set to the value shown in table 182 for the Available LBA Mapping Resource Count log parameter.

The DU bit, the TSD bit, and the FORMAT AND LINKING field for the Available LBA Mapping Resource Count shall be set for a binary format list log parameter as described in SPC-5.

The PARAMETER LENGTH field is described in SPC-5 and shall be set to the value shown in table 182 for the Available LBA Mapping Resource Count.

The RESOURCE COUNT field indicates an estimate of the number of available LBA mapping resources and is defined in 6.4.5.2.2.

The SCOPE field indicates the scope to which the RESOURCE COUNT field applies and is defined in table 183.

**Table 183 — SCOPE field**

Code	Description
00b	The scope of the resource count is not reported.
01b	The RESOURCE COUNT field indicates a resource that is dedicated to the logical unit. Usage of resources on other logical units does not impact the resource count.
10b	The resource count field indicates resources that may or may not be dedicated to any logical unit including the addressed logical unit. Usage of resources on other logical units may impact the resource count.
11b	Reserved

#### 6.4.5.2.2 RESOURCE COUNT field

The RESOURCE COUNT field indicates an estimate of the number of LBA resources expressed as a number of threshold sets for the threshold resource indicated by the parameter code value. The nominal number of LBA resources is calculated as follows:

$$\text{LBA resources} = \text{resource count} \times \text{threshold set size}$$

where:

resource count is the value in the RESOURCE COUNT field; and  
 threshold set size is the number of LBAs in each threshold set (i.e.,  $2^{\text{(threshold exponent)}}$  LBAs, where the threshold exponent is indicated in the Logical Block Provisioning VPD page (see 6.6.6)).

#### 6.4.5.3 Used LBA Mapping Resource Count log parameter

The Used LBA Mapping Resource Count log parameter of the Logical Block Provisioning log page has the format defined in table 184.

**Table 184 — Used LBA Mapping Resource Count log parameter format**

Byte	Bit	7	6	5	4	3	2	1	0	
0	(MSB)	PARAMETER CODE (0002h)								
1										(LSB)
2	Parameter control byte – binary format list log parameter (see SPC-5)									
	DU	Obsolete	TSD	Obsolete			FORMAT AND LINKING			
3	PARAMETER LENGTH (08h)									
4	(MSB)	RESOURCE COUNT								
...										
7	(LSB)									
8	Reserved							SCOPE		
9	Reserved									
...										
11										

The PARAMETER CODE field is described in ~~SPC-4~~SPC-5 and shall be set to the value shown in table 184 for the Used LBA Mapping Resource Count log parameter.

The DU bit, the TSD bit, and the FORMAT AND LINKING field for the Used LBA Mapping Resource Count log parameter shall be set for a binary format list log parameter as described in SPC-5.

The PARAMETER LENGTH field is described in SPC-5 and shall be set to the value shown in table 184 for the Used LBA Mapping Resource Count log parameter.

The RESOURCE COUNT field indicates an estimate of the number of used LBA mapping resources and is defined in 6.4.5.2.2.

The SCOPE field indicates the scope to which the RESOURCE COUNT field applies and is defined in table 183.

#### 6.4.5.4 Available Provisioning Resource Percentage log parameter

##### 6.4.5.4.1 Available Provisioning Resource Percentage log parameter overview

The Available Provisioning Resource Percentage log parameter of the Logical Block Provisioning log page has the format defined in table 185

**Table 185 — Available Provisioning Resource Percentage log parameter format**

Byte	Bit	7	6	5	4	3	2	1	0	
0	(MSB)	PARAMETER CODE (0003h)								
1		(LSB)								
2	Parameter control byte – binary format list log parameter (see SPC-5)									
	DU	Reserved	TSD	Reserved			FORMAT AND LINKING			
3	PARAMETER LENGTH (08h)									
4	(MSB)	RESOURCE COUNT								
5		(LSB)								
6	Reserved									
7										
8	Reserved						SCOPE			
9	Reserved									
...										
11										

The PARAMETER CODE field is described in SPC-5 and shall be set to the value shown in table 185 for the Available Provisioning Resource Percentage log parameter.

The DU bit, the TSD bit, and the FORMAT AND LINKING field for the Available Provisioning Resource Percentage log parameter shall be set for a binary format list log parameter as described in SPC-5.

The PARAMETER LENGTH field is described in SPC-5 and shall be set to the value shown in table 185 for the Available Provisioning Resource Percentage.

The RESOURCE COUNT field indicates an estimate of the percentage of available provisioning resources used for this logical block provisioning threshold and is defined in 6.4.5.4.2.

The SCOPE field indicates the scope to which the RESOURCE COUNT field applies and is defined in table 183.

##### 6.4.5.4.2 RESOURCE COUNT field

The RESOURCE COUNT field (see table 186) contains an estimate of the percentage of resources available for allocation to LBAs as a percentage of the manufacturer's total resources available for allocation. The units for the reported values are percent and range from 0% to 100%.

**Table 186 — RESOURCE COUNT field**

Code	Description
0 to 100	0% to 100% of the provisioning resources of the logical unit are available
All others	Reserved

#### 6.4.5.5 De-duplicated LBA Resource Count log parameter

The De-duplicated LBA Resource Count log parameter of the Logical Block Provisioning log page (see table 187) contains information about de-duplicated LBA resources.

**Table 187 — De-duplicated LBA Resource Count log parameter format**

Byte	Bit	7	6	5	4	3	2	1	0	
0	(MSB)	PARAMETER CODE (0100h)								
1										(LSB)
2	Parameter control byte – binary format list log parameter (see SPC-5)									
	DU	Obsolete	TSD	Obsolete			FORMAT AND LINKING			
3	PARAMETER LENGTH (08h)									
4	(MSB)	RESOURCE COUNT								
...										
7		(LSB)								
8	Reserved							SCOPE		
9	Reserved									
...										
11										

The PARAMETER CODE field is described in SPC-5 and shall be set to the value shown in table 187 for the De-duplicated LBA Resource Count log parameter.

The DU bit, the TSD bit, and the FORMAT AND LINKING field for the De-duplicated LBA Resource Count log parameter shall be set for a binary format list log parameter as described in SPC-5.

The PARAMETER LENGTH field is described in SPC-5 and shall be set to the value shown in table 187 for the De-duplicated LBA Resource Count log parameter.

The RESOURCE COUNT field indicates an estimate of the number of LBA resources made available as a result of de-duplication and is defined in 6.4.5.2.2.

The SCOPE field indicates the scope to which the RESOURCE COUNT field applies and is defined in table 183.

#### 6.4.5.6 Compressed LBA Resource Count log parameter

The Compressed LBA Resource Count log parameter of the Logical Block Provisioning log page (see table 188) contains information about compressed LBA resources.

**Table 188 — Compressed LBA Resource Count log parameter format**

Bit	7	6	5	4	3	2	1	0
Byte								
0	(MSB) _____							
1	PARAMETER CODE (0101h) _____ (LSB)							
2	Parameter control byte – binary format list log parameter (see SPC-5)							
	DU	Obsolete	TSD	Obsolete			FORMAT AND LINKING	
3	PARAMETER LENGTH (08h) _____							
4	(MSB) _____							
...	RESOURCE COUNT _____							
7	(LSB) _____							
8	Reserved						SCOPE	
9	_____							
...	Reserved _____							
11	_____							

The PARAMETER CODE field is described in SPC-5 and shall be set to the value shown in table 188 for the Compressed LBA Resource Count log parameter.

The DU bit, the TSD bit, and the FORMAT AND LINKING field for the Compressed LBA Resource Count log parameter shall be set for a binary format list log parameter as described in SPC-5.

The PARAMETER LENGTH field is described in SPC-5 and shall be set to the value shown in table 188 for the Compressed LBA Resource Count log parameter.

The RESOURCE COUNT field indicates an estimate of the number of LBA resources made available as a result of compression and is defined in 6.4.5.2.2.

The SCOPE field indicates the scope to which the RESOURCE COUNT field applies and is defined in table 183.



#### 6.4.5.7 Total Efficiency LBA Resource Count log parameter

The Total Efficiency LBA Resource Count log parameter of the Logical Block Provisioning log page (see table 189) contains information about the combined effects of all LBA resource efficiencies (e.g., the result of the combination of de-duplicated LBA resources and compressed LBA resources).

**Table 189 — Total Efficiency LBA Resource Count log parameter format**

Bit	7	6	5	4	3	2	1	0
Byte								
0	(MSB)							
1	PARAMETER CODE (0102h) (LSB)							
2	Parameter control byte – binary format list log parameter (see SPC-5)							
	DU	Obsolete	TSD	Obsolete		FORMAT AND LINKING		
3	PARAMETER LENGTH (08h)							
4	(MSB)							
...	RESOURCE COUNT							
7	(LSB)							
8	Reserved						SCOPE	
9								
...	Reserved							
11								

The PARAMETER CODE field is described in SPC-5 and shall be set to the value shown in table 189 for the Total Efficiency LBA Resource Count log parameter.

The DU bit, the TSD bit, and the FORMAT AND LINKING field for the Total Efficiency LBA Resource Count log parameter shall be set for a binary format list log parameter as described in SPC-5.

The PARAMETER LENGTH field is described in SPC-5 and shall be set to the value shown in table 189 for the Total Efficiency LBA Resource Count log parameter.

The RESOURCE COUNT field indicates an estimate of the number of LBA resources made available by the combined effects of all LBA resource efficiency methods (e.g., de-duplication and compression) and is defined in 6.4.5.2.2. The algorithm used to calculate this value is not defined by this standard.

The SCOPE field indicates the scope to which the RESOURCE COUNT field applies and is defined in table 183.

#### 6.4.6 [LPS Misalignment log page](#)

##### 6.4.6.1 [Overview](#)

Using the format defined in table 191, the LPS Misalignment log page reports misaligned write command information (see 4.6.2). The parameter codes for the LPS Misalignment log page are listed in table 190.

**Table 190 — [LPS Misalignment log page parameter codes](#)**

<a href="#">Parameter code</a>	<a href="#">Description</a>	<a href="#">Resettable or Changeable</a> <sup>a</sup>	<a href="#">Reference</a>	<a href="#">Support</a>
<a href="#">0000h</a>	<a href="#">LPS Misalignment Count</a>	<a href="#">Reset Only</a>	<a href="#">6.4.6.2</a>	<a href="#">Mandatory</a>
<a href="#">0001h to F000h</a>	<a href="#">LPS Misalignment</a>	<a href="#">Reset Only</a>	<a href="#">6.4.6.3</a>	<a href="#">Optional</a> <sup>b</sup>
<a href="#">All others</a>	<a href="#">Reserved</a>			
<sup>a</sup> <a href="#">The keywords in this column – Always, Reset Only, and Never – are defined in 6.4.1.2.</a> <sup>b</sup> <a href="#">If the LPS Misalignment log page is supported, then at least one LPS Misalignment log parameter shall be supported.</a>				

The LPS Misalignment log page has the format defined in table 191

**Table 191 — [LPS Misalignment log page](#)**

Byte	Bit	7	6	5	4	3	2	1	0
0	DS	SPF (1b)	PAGE CODE (15h)						
1	SUBPAGE CODE (03h)								
2	(MSB)								
3	PAGE LENGTH (n - 3)								
		(LSB)							
LPS Misalignment log parameters									
4	LPS misalignment log parameter [first] (see table 193)								
...									
⋮									
	LPS misalignment log parameter [last] (see table 193)								
...									
n									

The disable save (DS) bit, the subpage format (SPF) bit, the PAGE CODE field, and the SUBPAGE CODE field are described in SPC-4 and shall be set to the values defined in table 191 for the LPS Misalignment log page.

The PAGE LENGTH field is defined in SPC-5.

The contents of each LPS misalignment log parameter depends on the value in the PARAMETER CODE field (see table 190).

#### 6.4.6.2 LPS Misalignment Count log parameter

The LPS Misalignment Count log parameter has the format defined in table 192 and indicates the number of LPS Misalignment log parameters that are available.

**Table 192 — LPS Misalignment Count log parameter format**

Byte	Bit	7	6	5	4	3	2	1	0	
0	(MSB)	PARAMETER CODE (0000h)								(LSB)
1										
2	Parameter control byte – binary format list log parameter (see SPC-5)									
	DU	Obsolete	TSD	Obsolete			FORMAT AND LINKING			
3		PARAMETER LENGTH (04h)								
4	(MSB)	MAX_LPSM								(LSB)
5										
6	(MSB)	LPS MISALIGNMENT COUNT								(LSB)
7										

The PARAMETER CODE field is described in SPC-4 and shall be set as defined in table 192 for the LPS Misalignment Count log parameter.

The DU bit, the TSD bit, the ETC bit, the TMC field, and the FORMAT AND LINKING field for the LPS Misalignment Count log parameter shall be set for a binary format list log parameter as described in SPC-4.

The PARAMETER LENGTH field is described in SPC-4 and shall be set to the value defined in table 192 for the LPS Misalignment Count log parameter.

The MAX\_LPSM field indicates the maximum number of LPS Misalignment log parameters (see 6.4.6.3) supported by the device server. The device server may support any number of LPS Misalignment log parameters from 0001h to F000h inclusive.

The LPS MISALIGNMENT COUNT field indicates the number of LPS Misalignment log parameters that are available.

The LPS MISALIGNMENT COUNT field changes whenever the number of LPS Misalignment log parameters changes.

#### 6.4.6.3 LPS Misalignment log parameter

An LPS Misalignment log parameter has the format defined in table 193. If no misaligned write commands have been processed since the most recent reset of the LPS Misalignment log page then no LPS

Misalignment log parameters shall be present. LPS Misalignment log parameters are added as described in 4.6.2.

**Table 193 — LPS Misalignment log parameter format**

Byte	Bit	7	6	5	4	3	2	1	0
0	(MSB)	PARAMETER CODE (0001h to F000h)							
1		(LSB)							
2	Parameter control byte – binary format list log parameter (see SPC-5)								
	DU	Obsolete	TSD	Obsolete			FORMAT AND LINKING		
3	PARAMETER LENGTH (08h)								
4	(MSB)								
...		LBA OF MISALIGNED BLOCK							
11		(LSB)							

The PARAMETER CODE field is described in SPC-4. A PARAMETER CODE field set to 0001h indicates the oldest misaligned write command reported by the log, and successive values indicate successive misaligned writes.

The DU bit, the TSD bit, the ETC bit, the TMC field, and the FORMAT AND LINKING field for a LPS Misalignment log parameter shall be set for a binary format list log parameter as described in SPC-4.

The PARAMETER LENGTH field is defined in SPC-4 and shall be set to the value defined in table 193 for a LPS Misalignment log parameter.

The LBA OF MISALIGNED BLOCK field indicates the starting LBA associated with the misaligned write command.

#### 6.4.7 Non-volatile Cache log page

##### 6.4.7.1 Non-volatile Cache log page overview

Using the format shown in table 195, the Nonvolatile Cache log page reports the status of battery backup for a nonvolatile cache. The parameter codes for the Nonvolatile Cache log page are listed in table 194.

**Table 194 — Nonvolatile Cache log parameters**

Parameter code	Description	Resettable or Changeable <sup>a</sup>	Reference	Support-Required
0000h	Remaining Nonvolatile Time	Never	6.4.7.2	Mandatory
0001h	Maximum Nonvolatile Time	Never	6.4.7.3	Mandatory
All others	Reserved			
<sup>a</sup> The keywords in this column – Always, Reset Only, and Never – are defined in 6.4.1.2.				

The Nonvolatile Cache log page has the format defined in table 195.

**Table 195 — Nonvolatile Cache log page**

Byte	Bit	7	6	5	4	3	2	1	0
0		DS	SPF (0b)	PAGE CODE (17h)					
1		SUBPAGE CODE (00h)							
2		(MSB)	PAGE LENGTH (n - 3)						
3									
Nonvolatile cache log parameters									
4		Non-volatile cache log parameter [first] (see table 194)							
...									
⋮									
		Nonvolatile cache log parameter [last] (see table 194)							
...									
n									

The disable save (DS) bit, the subpage format (SPF) bit, the PAGE CODE field, the SUBPAGE CODE field, and the PAGE LENGTH field are described in SPC-5.

The SPF bit, the PAGE CODE field, and the SUBPAGE CODE field shall be set to the values shown in table 195 for the Nonvolatile Cache log page.

#### 6.4.7.2 Remaining Nonvolatile Time log parameter

The Remaining Nonvolatile Time log parameter of the Nonvolatile Cache log page has the format defined in table 196.

**Table 196 — Remaining Nonvolatile Time log parameter format**

Byte	Bit	7	6	5	4	3	2	1	0
0		(MSB)	PARAMETER CODE (0000h)						
1									(LSB)
2	Parameter control byte – binary format list log parameter (see SPC-5)								
	DU	Obsolete	TSD	Obsolete			FORMAT AND LINKING		
3	PARAMETER LENGTH (04h)								
4	Obsolete								
5		(MSB)							
...		REMAINING NONVOLATILE TIME							
7									(LSB)

The PARAMETER CODE field is described in SPC-5 and shall be set to the value shown in table 196 for the Remaining Nonvolatile Time log parameter.

The DU bit, the TSD bit, and the FORMAT AND LINKING field for the Remaining Nonvolatile Time log parameter shall be set for a binary format list log parameter as described in SPC-5.

The PARAMETER LENGTH field is described in SPC-5 and shall be set to the value shown in table 196 for the Remaining Nonvolatile Time log parameter.

The REMAINING NONVOLATILE TIME field is defined in table 197.

**Table 197 — REMAINING NONVOLATILE TIME field**

Code	Description
00_0000h	Nonvolatile cache is volatile, either permanently or temporarily (e.g., if batteries require recharging).
00_0001h	Nonvolatile cache is expected to remain nonvolatile for an unknown amount of time (e.g., if battery status is unknown)
00_0002h to FF_FFFEh	Nonvolatile cache is expected to remain nonvolatile for the number of minutes indicated (e.g., for the life of the battery supplying power to random access memory).
FF_FFFFh	Nonvolatile cache is indefinitely nonvolatile.

#### 6.4.7.3 Maximum Nonvolatile Time log parameter

The Maximum Nonvolatile Time log parameter of the Nonvolatile Cache log page has the format defined in table 198.

**Table 198 — Maximum Nonvolatile Time log parameter format**

Bit	7	6	5	4	3	2	1	0
Byte								
0	(MSB) _____							
1	PARAMETER CODE (0001h) _____ (LSB)							
2	Parameter control byte – binary format list log parameter (see SPC-5)							
	DU	Obsolete	TSD	Obsolete		FORMAT AND LINKING		
3	PARAMETER LENGTH (04h) _____							
4	Obsolete							
5	(MSB) _____							
...	MAXIMUM NONVOLATILE TIME _____							
7	(LSB)							

The PARAMETER CODE field is described in SPC-5 and shall be set to the value shown in table 198 for the Maximum Nonvolatile Time log parameter.

The DU bit, the TSD bit, and the FORMAT AND LINKING field for the Maximum Nonvolatile Time log parameter shall be set for a binary format list log parameter as described in SPC-5.

The PARAMETER LENGTH field is described in SPC-5 and shall be set to the value shown in table 198 for the Maximum Nonvolatile Time log parameter.

The MAXIMUM NONVOLATILE TIME field is defined in table 199.

**Table 199 — MAXIMUM NONVOLATILE TIME field**

Code	Description
00_0000h	Nonvolatile cache is volatile
00_0001h	Reserved
00_0002h to FF_FFFEh	Nonvolatile cache is capable of being nonvolatile for the estimated number of minutes indicated. If the time is based on batteries, then the time shall be based on the last full charge capacity rather than the design capacity of the batteries.
FF_FFFFh	Nonvolatile cache is indefinitely nonvolatile.

#### 6.4.8 Pending Defects log page

##### 6.4.8.1 Overview

Using the format defined in table 201, the Pending Defects log page reports an unsorted list of logical blocks for which the device server has detected an unrecovered medium error. The parameter codes for the Pending Defects log page are listed in table 200.

**Table 200 — Pending Defects log page parameter codes**

Parameter code	Description	Resettable or Changeable <sup>a</sup>	Reference	Support- <del>required</del>
0000h	Pending Defect Count	Never	6.4.8.2	Mandatory
0001h to F000h	Pending Defect	Never	6.4.8.3	Mandatory <sup>b</sup>
All others	Reserved			
<sup>a</sup> The keywords in this column – Always, Reset Only, and Never – are defined in 6.4.1.2. <sup>b</sup> If the Pending Defects log page is supported, then at least one Pending Defect log parameter shall be supported.				

The Pending Defects log page has the format defined in table 201

**Table 201 — Pending Defects log page**

Byte	Bit	7	6	5	4	3	2	1	0
0		DS	SPF (1b)	PAGE CODE (15h)					
1		SUBPAGE CODE (01h)							
2		(MSB)	PAGE LENGTH (n-3)						
3									
		Pending defect parameters							
4		Pending defect parameter [first] (see table 202)							
...									
		⋮							
		Pending defect parameter [last] (see table 202)							
...									
n									

The ~~disable save (DS) bit, the~~ subpage format (SPF) bit, the PAGE CODE field, the SUBPAGE CODE field, and the PAGE LENGTH field are described in SPC-5 and shall be set to the values defined in table 201 for the Pending Defects log page.

The disable save (DS) bit is described in SPC-5

The contents of each pending defect parameter depends on the value in its PARAMETER CODE field (see table 200).

#### 6.4.8.2 Pending Defect Count log parameter

The Pending Defect Count log parameter has the format defined in table 202 and indicates the number of Pending Defect log parameters that are available.

**Table 202 — Pending Defect Count log parameter format**

Byte	Bit	7	6	5	4	3	2	1	0
0		(MSB)	PARAMETER CODE (0000h)						
1									(LSB)
2		Parameter control byte – binary format list log parameter (see SPC-5)							
		DU	Obsolete	TSD	Obsolete		FORMAT AND LINKING		
3		PARAMETER <del>LIST</del> LENGTH (04h)							
4		(MSB)							
...		PENDING DEFECT COUNT							
7									(LSB)



The PARAMETER CODE field is described in SPC-5 and shall be set as defined in table 202 for the Pending Defect Count log parameter.

The DU bit, the TSD bit, and the FORMAT AND LINKING field for the Pending Defect Count log parameter shall be set for a binary format list log parameter as described in SPC-5.

The PARAMETER ~~LIST~~ LENGTH field is described in SPC-5 and shall be set to the value defined in table 202 for the Pending Defect Count log parameter.

The PENDING DEFECT COUNT field indicates the number of Pending Defect log parameters that are available. The PENDING DEFECT COUNT field changes whenever the number of Pending Defect log parameters (see 6.4.8.3) changes.

#### 6.4.8.3 Pending Defect log parameter

A Pending Defect log parameter has the format defined in table 203. If no unrecovered errors have occurred then no Pending Defect log parameters shall be present. A Pending Defect log parameter shall be added for each LBA for which the device server has detected an unrecovered medium error that is not:

- a) a pseudo unrecovered read error (see 4.18.2);
- b) a predicted unrecovered read error (see 4.20.3.3); or
- c) a predicted unrecovered write error (see 4.20.3.5).

If all of the supported parameter code values have been used and a new defect is discovered, then the device server shall not add more Pending Defect log parameters and the PENDING ERROR COUNT field shall not be changed.

Pending Defect log parameters may be duplicates of Background Scan Results parameters in the Background Scan log page (see 6.4.2).

A Pending Defect log parameter shall be removed if the indicated LBA:

- a) is reassigned without error;
- b) is written without error; or
- c) is read without error.

A Pending Defect log parameter may be removed if the indicated LBA is unmapped without error.

A sanitize overwrite operation (see 5.25.2.2) and a format operation (see 5.4.1) shall cause all Pending Defect log parameters to be removed..

**Table 203 — Pending Defect log parameter format**

Byte	Bit	7	6	5	4	3	2	1	0
0	(MSB)	PARAMETER CODE (0001h) to (F000h)							
1		(LSB)							
2	Parameter control byte – binary format list log parameter (see SPC-5)								
	DU	Obsolete	TSD	Obsolete			FORMAT AND LINKING		
3	PARAMETER <del>LIST</del> LENGTH (14h)								
4	(MSB)	ACCUMULATED POWER ON HOURS							
...									
7									
8	(MSB)	LOGICAL BLOCK ADDRESS							
...									
15									
		(LSB)							

The PARAMETER CODE field is described in SPC-5 and shall be set as defined in table 203 for a Pending Defect log parameter.

The DU bit, the TSD bit, and the FORMAT AND LINKING field for a Pending Defect log parameter shall be set for a binary format list log parameter as described in SPC-5.

The PARAMETER ~~LIST~~ LENGTH field is defined in SPC-5 and shall be set to the value defined in table 203 for a Pending Defect log parameter.

The ACCUMULATED POWER ON HOURS field indicates the number of hours that the device server has been powered on since manufacturing at the time the Pending Defect log parameter was created. A value of FFFF\_FFFF\_FFFF\_FFFFh indicates that the accumulated power on hours value is unknown.

The LOGICAL BLOCK ADDRESS field indicates the LBA associated with the unrecovered medium error.

## 6.4.9 Solid State Media log page

### 6.4.9.1 Solid State Media log page overview

Using the format shown in table 204, the Solid State media log page reports parameters that are specific to SCSI target devices that contain solid state media.

**Table 204 — Solid State Media log page**

Byte	Bit	7	6	5	4	3	2	1	0
0		DS	SPF (0b)	PAGE CODE (11h)					
1		SUBPAGE CODE (00h)							
2		(MSB)							
3		PAGE LENGTH (n - 3)							(LSB)
Solid state media log parameters									
4									
...		Solid state media parameter [first]							
⋮									
...		Solid state media parameter [last]							
n									

The disable save (DS) bit, the subpage format (SPF) bit, the PAGE CODE field, the SUBPAGE CODE field, and the PAGE LENGTH field are described in SPC-5.

The SPF bit, the PAGE CODE field, and the SUBPAGE CODE field shall be set to the values shown in table 204 for the Solid State Media log page.

The parameter codes for the Solid State Media log page are listed in table 205.

**Table 205 — Solid State Media log parameters**

Parameter code	Description	Resettable or Changeable <sup>a</sup>	Reference	Support <b>Required</b>
0001h	Percentage Used Endurance Indicator	Never	6.4.9.2	Mandatory
All others	Reserved			
<sup>a</sup> The keywords in this column – Always, Reset Only, and Never – are defined in 6.4.1.2.				

#### 6.4.9.2 Percentage Used Endurance Indicator log parameter

The Percentage Used Endurance Indicator log parameter of the Solid State Media log page has the format defined in table 206.

**Table 206 — Percentage Used Endurance Indicator log parameter format**

Byte	Bit	7	6	5	4	3	2	1	0
0		(MSB) _____							
1		PARAMETER CODE (0001h) _____ (LSB)							
2		Parameter control byte – binary format list log parameter (see SPC-5)							
		DU	Obsolete	TSD	Obsolete		FORMAT AND LINKING		
3		PARAMETER LENGTH (04h)							
4		Reserved _____							
...									
6									
7		PERCENTAGE USED ENDURANCE INDICATOR							

The PARAMETER CODE field is described in SPC-5 and shall be set to the value shown in table 206 for the Percentage Used Endurance Indicator log parameter.

The DU bit, the TSD bit, and the FORMAT AND LINKING field for the Percentage Used Endurance Indicator log parameter shall be set for a binary format list log parameter as described in SPC-5.

The PARAMETER LENGTH field is described in SPC-5 and shall be set to the value shown in table 206 for the Percentage Used Endurance Indicator log parameter.

The PERCENTAGE USED ENDURANCE INDICATOR field indicates an estimate of the percentage of a SCSI target device that contains solid state media life that has been used. The value in the field shall be set to zero at the time of manufacture. A value of 100 indicates that the estimated endurance of the SCSI target device that contains solid state media has been consumed, but may not indicate a SCSI target device that contains solid state media failure (e.g., minimum power-off data retention capability reached for SCSI target devices that contain solid state media using flash technology). The value is allowed to exceed 100. Values greater than 254 shall be reported as 255. The device server shall update the value at least once per power-on hour.

## 6.4.10 Utilization log page

### 6.4.10.1 Utilization log page overview

Using the format shown in table 208, the Utilization log page reports estimates of the rate at which device wear factors (e.g., damage to the recording medium) are being used. The parameter codes for the Utilization log page are defined in table 207.

**Table 207 — Utilization log page parameter codes**

Parameter code	Description	Resettable or Changeable <sup>a</sup>	Reference	Support <b>Required</b>
0000h	Workload Utilization	Never	6.4.10.2	Mandatory
0001h	Utilization Usage Rate Based on Date and Time	Never	6.4.10.3	Optional
All others	Reserved			
<sup>a</sup> The keywords in this column – Always, Reset Only, and Never – are defined in 6.4.1.2.				

The Utilization log page has the format defined in table 208.

**Table 208 — Utilization log page**

Byte	Bit	7	6	5	4	3	2	1	0
0		DS (1b)	SPF (1b)	PAGE CODE (0Eh)					
1		SUBPAGE CODE (01h)							
2		(MSB)							
3		PAGE LENGTH (n - 3)							(LSB)
Utilization log parameters									
4		Utilization log parameter [first]							
...									
⋮									
		Utilization log parameter [last]							
...									
n									

The disable save (DS) bit, the subpage format (SPF) bit, the PAGE CODE field, the SUBPAGE CODE field, and the PAGE LENGTH field are described in SPC-5 and shall be set to the values shown in table 208 for the Utilization log page.

#### 6.4.10.2 Workload Utilization log parameter

The Workload Utilization log parameter for the Utilization log page has the format defined in table 209.

**Table 209 — Workload Utilization log parameter format**

Byte	Bit	7	6	5	4	3	2	1	0
0	(MSB)	PARAMETER CODE (0000h)							
1									
2	Parameter control byte – bounded data counter log parameter (see SPC-5)								
	DU	Obsolete	TSD	Obsolete			FORMAT AND LINKING		
3		PARAMETER LENGTH (02h)							
4	(MSB)	WORKLOAD UTILIZATION							
5									

The PARAMETER CODE field is described in ~~SPC-4~~SPC-5 and shall be set to the value shown in table 209 for the Workload Utilization log parameter.

The DU bit, the TSD bit, and the FORMAT AND LINKING field for the Workload Utilization log parameter shall be set for a bounded data counter log parameter as described in SPC-5.

The PARAMETER LENGTH field is described in SPC-5 and shall be set to the value shown in table 209 for the Workload Utilization log parameter.

The WORKLOAD UTILIZATION field (see table 210) contains an estimate of the utilization associated with the logical unit as a percentage of the manufacturer's designs for various wear factors (e.g., wear of the medium, head load events), if any. The the units for the reported values are percent times 100 and range from 0.00% to 655.35%.

**Table 210 — WORKLOAD UTILIZATION field**

Code	Description
0 to 9 999	Less than (i.e., 0.00% to 99.99% of) the designed workload has been utilized.
10 000	Exactly the designed workload for the device has been utilized.
10 001 to 65 534	Greater than (i.e., 100.01% to 655.34% of) the designed workload has been utilized.
65 535	Greater than 655.34% of the designed workload has been utilized.

#### 6.4.10.3 Utilization Usage Rate Based on Date and Time

The Utilization Rate Based on Date and Time log parameter for the Utilization log page has the format defined in table 211. If the current date and time has not been initialized by a SET TIMESTAMP command (see SPC-5), then the Utilization Rate Based on Date and Time log parameter shall not be returned in the Utilization log page.

**Table 211 — Utilization Rate Based on Date and Time log parameter format**

Byte	Bit	7	6	5	4	3	2	1	0
0	(MSB)	PARAMETER CODE (0001h)							
1		(LSB)							
2	Parameter control byte – bounded data counter log parameter (see SPC-5)								
	DU	Obsolete	TSD	Obsolete			FORMAT AND LINKING		
3	PARAMETER LENGTH (02h)								
4	DATE AND TIME BASED UTILIZATION RATE								
5	Reserved								

The PARAMETER CODE field is described in SPC-5 and shall be set to the value shown in table 211 for the Utilization Rate Based on Date and Time log parameter.

The DU bit, the TSD bit, and the FORMAT AND LINKING field for the Utilization Rate Based on Date and Time log parameter shall be set for a bounded data counter log parameter as described in SPC-5.

The PARAMETER LENGTH field is described in SPC-5 and shall be set to the value shown in table 211 for the Utilization Rate Based on Date and Time log parameter.

The DATE AND TIME BASED UTILIZATION RATE field (see table 212) contains an estimate of the rate at which device wear factors (e.g., damage to the recording medium) associated with the logical unit have been used during the interval that begins at the date and time of manufacture and ends at the timestamp that is reported by a REPORT TIMESTAMP command (i.e., the current value of a device clock) (see SPC-5).

**Table 212 — DATE AND TIME BASED UTILIZATION RATE field**

Code	Description
0 to 99	The Workload Utilization usage rate has been less than (i.e., 0% to 99% of) the designed usage rate during the interval that begins at the date and time of manufacture and ends at the timestamp.
100	The Workload Utilization usage rate has been the exact designed usage rate during the interval that begins at the date and time of manufacture and ends at the timestamp.
101 to 254	The Workload Utilization usage rate has been greater than (i.e., 101% to 254% of) the designed usage rate during the interval that begins at the date and time of manufacture and ends at the timestamp.
255	The Workload Utilization usage rate has been greater than 254% of designed usage rate during the interval that begins at the date and time of manufacture and ends at the timestamp.

## 6.5 Mode parameters

### 6.5.1 Mode ~~pages~~~~parameters~~ overview

See table 213 for references to defines the mode pages and block descriptors (see 6.5.2) used by direct access block devices. The mode pages and their corresponding page codes and subpage codes for direct access block devices are shown in table 213. [See 6.5.2 for a description of block descriptors.](#)

**Table 213 — Mode page codes and subpage codes for direct access block devices (part 1 of 2)**

Mode page name	Page code	Subpage code	Reference
Application Tag <del>mode page</del>	0Ah	02h	6.5.3
ATA Power Condition	1Ah	F1h	SAT-3
Background Control <del>mode page</del>	1Ch	01h	6.5.4
Background Operation Control <del>mode page</del>	0Ah	06h	6.5.5
Caching <del>mode page</del>	08h	00h	6.5.6
<a href="#">Command Duration Limit A</a>	<a href="#">0Ah</a>	<a href="#">03h</a>	<a href="#">SPC-5</a>
<a href="#">Command Duration Limit B</a>	<a href="#">0Ah</a>	<a href="#">04h</a>	<a href="#">SPC-5</a>
Control <del>mode page</del>	0Ah	00h	SPC-5
Control Extension <del>mode page</del>	0Ah	01h	SPC-5
Disconnect-Reconnect <del>mode page</del>	02h	00h	SPC-5
Enclosure Services Management <del>mode page</del> <sup>a</sup>	14h	00h	SES-3
IO Advice Hints Grouping <del>mode page</del>	0Ah	05h	6.5.7
Informational Exceptions Control <del>mode page</del>	1Ch	00h	6.5.8
Logical Block Provisioning <del>mode page</del>	1Ch	02h	6.5.9
PATA Control	0Ah	F1h	SAT-3
Power Condition <del>mode page</del>	1Ah	00h	SPC-5
Power Consumption <del>mode page</del>	1Ah	01h	SPC-5
Protocol-Specific Logical Unit <del>mode page</del>	18h	00h	SPC-5
Protocol-Specific Port <del>mode page</del>	19h	00h	SPC-5
Read-Write Error Recovery <del>mode page</del>	01h	00h	6.5.10
Return all mode pages and subpages <sup>b</sup>	3Fh	FFh	SPC-5
Return all mode pages <del>only (i.e., not including subpages)</del> <sup>b</sup>	3Fh	00h	SPC-5
Return all subpages for the specified mode page code <sup>b</sup>	00h to 3Eh	FFh	SPC-5
Note: SPC-5 contains a listing of mode page and subpage codes in numeric order.			
<sup>a</sup> Valid only if the ENCSERV bit is set to one in the standard INQUIRY data (see SPC-5). <sup>b</sup> Valid only for the MODE SENSE command. <sup>c</sup> <a href="#">All Ssubpage codes</a> of the following mode page codes are obsolete: 03h, 04h, 05h, 09h, 0Bh, 0Ch, 0Dh, and 10h. <sup>d</sup> The following mode page code and subpage code combinations are vendor specific and do not require a page format: a) mode page code 00h with subpage code 00h; and b) mode page codes 20h to 3Eh with all subpage codes.			

**Table 213 — Mode page codes and subpage codes for direct access block devices (part 2 of 2)**

Mode page name	Page code	Subpage code	Reference
Verify Error Recovery- <del>mode page</del>	07h	00h	6.5.11
Obsolete <sup>c</sup>			
Vendor specific <sup>d</sup>			
Reserved	all other page and subpage code combinations for direct access block devices		
Note: SPC-5 contains a listing of mode page and subpage codes in numeric order.			
<div>a Valid only if the ENCSERV bit is set to one in the standard INQUIRY data (see SPC-5).</div> <div>b Valid only for the MODE SENSE command.</div> <div>c <u>All Subpage codes</u> of the following mode page codes are obsolete: 03h, 04h, 05h, 09h, 0Bh, 0Ch, 0Dh, and 10h.</div> <div>d The following mode page code and subpage code combinations are vendor specific and do not require a page format:<div>a) mode page code 00h with subpage code 00h; and</div><div>b) mode page codes 20h to 3Eh with all subpage codes.</div></div>			

The mode parameter list, including the mode parameter header, is described in SPC-5. Direct access block devices support zero or one mode parameter block descriptors (i.e., the block descriptor is shared by all the logical blocks on the medium) (see 6.5.2).

The MEDIUM TYPE field in the mode parameter header (see SPC-5) shall be set to 00h.

The DEVICE-SPECIFIC PARAMETER field in the mode parameter header (see SPC-5) for direct access block devices is defined in table 214.

**Table 214 — DEVICE-SPECIFIC PARAMETER field for direct access block devices**

Bit	7	6	5	4	3	2	1	0
	WP	Reserved		DPOFUA	Reserved			

If the medium is write-protected (i.e., due to mechanisms outside the scope of this standard), or the software write protect (SWP) bit in the Control mode page (see SPC-5) is set to one, then the device server shall set the WP bit to one when returning a DEVICE-SPECIFIC PARAMETER field in response to a MODE SENSE command. If the medium is not write-protected and the SWP bit is set to zero, then the device server shall set the WP bit to zero when returning a DEVICE-SPECIFIC PARAMETER field in response to a MODE SENSE command.

The write protect (WP) bit for mode data sent with a MODE SELECT command shall be ignored by the device server.

The DPOFUA bit is reserved for mode data sent with a MODE SELECT command.

If the device server does not support the DPO bit and the FUA bit (see 4.15) being set to one, then the device server shall set the DPO and FUA supported (DPOFUA) bit to zero when returning a DEVICE-SPECIFIC PARAMETER field in response to a MODE SENSE command. If the device server supports the DPO bit and the FUA bit being set to one, then the device server shall set the DPOFUA bit to one when returning a DEVICE-SPECIFIC PARAMETER field in response to a MODE SENSE command.



## 6.5.2 Mode parameter block descriptors

### 6.5.2.1 Mode parameter block descriptors overview

If a device server returns a mode parameter block descriptor, then the device server shall return a short LBA mode parameter block descriptor (see 6.5.2.2) in the mode parameter data in response to:

- a) a MODE SENSE (6) command; or
- b) a MODE SENSE (10) command with the LLBAA bit set to zero.

A device server may return a long LBA mode parameter block descriptor (see 6.5.2.3) in the mode parameter data in response to a MODE SENSE (10) command with the LLBAA bit set to one.

If an application client sends a mode parameter block descriptor in the mode parameter list, then the application client sends a short LBA mode parameter block descriptor (see 6.5.2.2) for a MODE SELECT (6) command.

If an application client sends a mode parameter block descriptor in the mode parameter list, then the application client may send a long LBA mode parameter block descriptor (see 6.5.2.3) for a MODE SELECT (10) command.

Support for the mode parameter block descriptors is optional.

If the device server supports changing the block descriptor parameters by a MODE SELECT command, the number of logical blocks is changed, and the Application Tag mode page is supported, then the device server shall set:

- a) the current value of the ATMPE bit to zero in the Control mode page (see SPC-5); and
- b) the saved value, if saving is implemented, of the ATMPE bit to zero in the Control mode page.

If the device server supports changing the block descriptor parameters by a MODE SELECT command and the number of logical blocks or the logical block length is changed, then the device server establishes a unit attention condition of:

- a) CAPACITY DATA HAS CHANGED as described in 4.10; and
- b) MODE PARAMETERS CHANGED as described in SPC-5.

### 6.5.2.2 Short LBA mode parameter block descriptor

Table 215 defines the short LBA mode parameter block descriptor for direct access block devices used:

- a) with the MODE SELECT (6) and MODE SENSE (6) commands; and
- b) with the MODE SELECT (10) and MODE SENSE (10) commands when the LONGLBA bit is set to zero in the mode parameter header (see SPC-5).

**Table 215 — Short LBA mode parameter block descriptor**

Bit	7	6	5	4	3	2	1	0
Byte								
0	(MSB)							
...	NUMBER OF LOGICAL BLOCKS							
3								
4	Reserved							
5	(MSB)							
...	LOGICAL BLOCK LENGTH							
7								

A device server shall respond to a MODE SENSE command (see SPC-5) by reporting the number of logical blocks specified in the NUMBER OF LOGICAL BLOCKS field sent in the last MODE SELECT command that contained a mode parameter block descriptor. If no MODE SELECT command with a mode parameter block descriptor has been received then the current number of logical blocks shall be returned. To determine the number of logical blocks at which the logical unit is currently formatted, the application client shall use a READ CAPACITY command (see 5.17 and 5.18) rather than the MODE SENSE command.

In response to a MODE SENSE command, the device server may return a value of zero indicating that it does not report the number of logical blocks in the short LBA mode parameter block descriptor.

In response to a MODE SENSE command, if the number of logical blocks on the medium exceeds the maximum value that is able to be specified in the NUMBER OF LOGICAL BLOCKS field, then the device server shall return a value of FFFF\_FFFFh.

If the logical unit does not support changing its capacity by changing the NUMBER OF LOGICAL BLOCKS field using the MODE SELECT command (see SPC-5), then the value in the NUMBER OF LOGICAL BLOCKS field is ignored. If the device supports changing its capacity by changing the NUMBER OF LOGICAL BLOCKS field, then the NUMBER OF LOGICAL BLOCKS field is interpreted as follows:

- a) if the NUMBER OF LOGICAL BLOCKS field is set to zero, then the logical unit shall retain its current capacity if the logical block length has not changed. If the NUMBER OF LOGICAL BLOCKS field is set to zero and the content of the LOGICAL BLOCK LENGTH field (i.e., new logical block length) is different than the current logical block length, then the logical unit shall be set to its maximum capacity when the new logical block length takes effect (i.e., after a successful FORMAT UNIT command);
- b) if the NUMBER OF LOGICAL BLOCKS field is greater than zero and less than or equal to its maximum capacity, then the logical unit shall be set to that number of logical blocks. If the content of the LOGICAL BLOCK LENGTH field is the same as the current logical block length, then the logical unit shall not become format corrupt (see 4.10). This capacity setting shall be retained through power cycles, hard resets, logical unit resets, and I\_T nexus losses. If the content of the LOGICAL BLOCK LENGTH field is the same as the current logical block length, then this capacity setting shall take effect on successful completion of the MODE SELECT command. If the content of the LOGICAL BLOCK LENGTH field (i.e., new logical block length) is different than the current logical block length, then this capacity setting shall take effect when the new logical block length takes effect (i.e., after a successful FORMAT UNIT command);
- c) if the NUMBER OF LOGICAL BLOCKS field is set to a value greater than the maximum capacity of the device and less than FFFF\_FFFFh, then the device server shall terminate the MODE SELECT command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST. The logical unit shall retain its previous logical block descriptor settings; or
- d) if the NUMBER OF LOGICAL BLOCKS field is set to FFFF\_FFFFh, then the logical unit shall be set to its maximum capacity. If the content of the LOGICAL BLOCK LENGTH field is the same as the current logical block length, then the logical unit shall not become format corrupt (see 4.10). This capacity setting shall be retained through power cycles, hard resets, logical unit resets, and I\_T nexus losses. If the content of the LOGICAL BLOCK LENGTH field is the same as the current logical block length, then this capacity setting shall take effect on successful completion of the MODE SELECT command. If the content of the LOGICAL BLOCK LENGTH field (i.e., new logical block length) is different than the current logical block length, then this capacity setting shall take effect when the new logical block length takes effect (i.e., after a successful FORMAT UNIT command).

If the device server supports changing its logical unit's capacity by changing the NUMBER OF LOGICAL BLOCKS field, is in a logical unit that supports logical block provisioning management, and the capacity is increased, then the additional LBAs shall be in the initial provisioning management condition as specified in 4.7.3.2 or 4.7.3.3.

The LOGICAL BLOCK LENGTH field specifies the length in bytes of each logical block. No change shall be made to any logical blocks on the medium until a format operation (see 5.4) is initiated by an application client.

A device server shall respond to a MODE SENSE command (see SPC-5) by reporting the length of the logical blocks as specified in the LOGICAL BLOCK LENGTH field sent in the last MODE SELECT command that contained a mode parameter block descriptor (e.g., if the logical block length is 512 bytes and a MODE

SELECT command occurs with the LOGICAL BLOCK LENGTH field set to 520 bytes, any MODE SENSE commands would return 520 in the LOGICAL BLOCK LENGTH field). If no MODE SELECT command with a block descriptor has been received then the current logical block length shall be returned. To determine the logical block length at which the logical unit is currently formatted, the application client shall use a READ CAPACITY command (see 5.17 and 5.18) rather than a MODE SENSE command.

### 6.5.2.3 Long LBA mode parameter block descriptor

Table 216 defines the long LBA mode parameter block descriptor for direct access block devices used with the MODE SELECT (10) command and MODE SENSE (10) command when the LONGLBA bit is set to one in the mode parameter header (see SPC-5).

**Table 216 — Long LBA mode parameter block descriptor**

Bit	7	6	5	4	3	2	1	0
Byte								
0	(MSB)							
...	NUMBER OF LOGICAL BLOCKS							
7	(LSB)							
8	Reserved							
...								
11								
12	(MSB)							
...	LOGICAL BLOCK LENGTH							
15	(LSB)							

A device server shall respond to a MODE SENSE command (see SPC-5) by reporting the number of logical blocks specified in the NUMBER OF LOGICAL BLOCKS field sent in the last MODE SELECT command that contained a mode parameter block descriptor. If no MODE SELECT command with a mode parameter block descriptor has been received then the current number of logical blocks shall be returned. To determine the number of logical blocks at which the logical unit is currently formatted, the application client shall use a READ CAPACITY command (see 5.17 and 5.18) rather than a MODE SENSE command.

In response to a MODE SENSE command, the device server may return a value of zero indicating that it does not report the number of logical blocks in the long LBA mode parameter block descriptor.

If the logical unit does not support changing its capacity by changing the NUMBER OF LOGICAL BLOCKS field using the MODE SELECT command (see SPC-5), then the value in the NUMBER OF LOGICAL BLOCKS field is ignored. If the device supports changing its capacity by changing the NUMBER OF LOGICAL BLOCKS field, then the NUMBER OF LOGICAL BLOCKS field is interpreted as follows:

- if the NUMBER OF LOGICAL BLOCKS field is set to zero, then the logical unit shall retain its current capacity if the logical block length has not changed. If the NUMBER OF LOGICAL BLOCKS field is set to zero and the content of the LOGICAL BLOCK LENGTH field (i.e., new logical block length) is different than the current logical block length, then the logical unit shall be set to its maximum capacity when the new logical block length takes effect (i.e., after a successful FORMAT UNIT command);
- if the NUMBER OF LOGICAL BLOCKS field is greater than zero and less than or equal to its maximum capacity, then the logical unit shall be set to that number of logical blocks. If the content of the LOGICAL BLOCK LENGTH field is the same as the current logical block length, then the logical unit shall not become format corrupt (see 4.10). This capacity setting shall be retained through power cycles, hard resets, logical unit resets, and I\_T nexus losses. If the content of the LOGICAL BLOCK LENGTH field is the same as the current logical block length, then this capacity setting shall take effect on successful completion of the MODE SELECT command. If the content of the LOGICAL BLOCK LENGTH field (i.e., new logical block length) is different than the current logical block length, then this capacity setting

shall take effect when the new logical block length takes effect (i.e., after a successful FORMAT UNIT command);

- c) if the NUMBER OF LOGICAL BLOCKS field is set to a value greater than the maximum capacity of the device and less than FFFF\_FFFF\_FFFF\_FFFFh, then the device server shall terminate the MODE SELECT command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST. The logical unit shall retain its previous block descriptor settings; or
- d) if the NUMBER OF LOGICAL BLOCKS field is set to FFFF\_FFFF\_FFFF\_FFFFh, then the logical unit shall be set to its maximum capacity. If the content of the LOGICAL BLOCK LENGTH field is the same as the current logical block length, then the logical unit shall not become format corrupt (see 4.10). This capacity setting shall be retained through power cycles, hard resets, logical unit resets, and I\_T nexus losses. If the content of the LOGICAL BLOCK LENGTH field is the same as the current logical block length, then this capacity setting shall take effect on successful completion of the MODE SELECT command. If the content of the LOGICAL BLOCK LENGTH field (i.e., new logical block length) is different than the current logical block length, then this capacity setting shall take effect when the new logical block length takes effect (i.e., after a successful FORMAT UNIT command).

If the device server supports changing its logical unit's capacity by changing the NUMBER OF LOGICAL BLOCKS field, supports logical block provisioning management, and the capacity is increased, then the additional LBAs shall be in the initial provisioning management condition as specified in 4.7.3.2 or 4.7.3.3.

The LOGICAL BLOCK LENGTH field specifies the length in bytes of each logical block. No change shall be made to any logical blocks on the medium until a format operation (see 5.3) is initiated by an application client.

A device server shall respond to a MODE SENSE command (see SPC-5) by reporting the length of the logical blocks as specified in the LOGICAL BLOCK LENGTH field sent in the last MODE SELECT command that contained a mode parameter block descriptor (e.g., if the logical block length is 512 bytes and a MODE SELECT command occurs with the LOGICAL BLOCK LENGTH field set to 520 bytes, any MODE SENSE commands would return 520 in the LOGICAL BLOCK LENGTH field). If no MODE SELECT command with a block descriptor has been received then the current logical block length shall be returned. To determine the logical block length at which the logical unit is currently formatted, the application client shall use a READ CAPACITY command (see 5.17 and 5.18) rather than a MODE SELECT command.

### 6.5.3 Application Tag mode page

#### 6.5.3.1 Introduction

The Application Tag mode page (see table 217) specifies the logical block application tag that a device server configured for protection information (see 4.22.2) shall use for each LBA range if:

- a) the ATO bit in the Control mode page (see SPC-5) is set to one;
- b) the ATMPE bit in the Control mode page (see SPC-5) is set to one; and
- c) the WRPROTECT field requirements (see table 119) specify use of the Application tag mode page.

The mode page policy (see SPC-5) for this page shall be shared.

If a method not defined by this standard changes the parameter data to be returned by the device server in the Application Tag mode page, then the device server shall establish a unit attention condition for the SCSI initiator port associated with every I\_T nexus with the additional sense code set to MODE PARAMETERS CHANGED.

**Table 217 — Application Tag mode page**

Byte	Bit	7	6	5	4	3	2	1	0
0		PS	SPF (1b)	PAGE CODE (0Ah)					
1		SUBPAGE CODE (02h)							
2	(MSB)	PAGE LENGTH (n - 3)							
3									
4		Reserved							
...									
15									
Application tag descriptors									
16		Application tag descriptor [first] (see 6.5.3.2)							
...									
39									
		⋮							
n - 24		Application tag descriptor [last] (see 6.5.3.2)							
...									
n									

The parameters saveable (PS) bit, the subpage format (SPF) bit, the PAGE CODE field, the SUBPAGE CODE field, and the PAGE LENGTH field are defined in SPC-5.

The SPF bit, the PAGE CODE field, and the SUBPAGE CODE field shall be set to the values shown in table 217 for the Application Tag mode page.

The application tag descriptor is defined in 6.5.3.2.

### 6.5.3.2 Application tag descriptor

The application tag descriptor format is defined in table 218.

**Table 218 — Application tag descriptor format**

Byte	Bit	7	6	5	4	3	2	1	0
0	LAST	Reserved							
1		Reserved							
...									
5									
6	(MSB)	LOGICAL BLOCK APPLICATION TAG							
7		(LSB)							
8	(MSB)	LOGICAL BLOCK ADDRESS							
...									
15		(LSB)							
16	(MSB)	LOGICAL BLOCK COUNT							
...									
23		(LSB)							

A LAST bit set to one specifies that this application tag descriptor is the last valid application tag descriptor in the Application Tag mode page. A LAST bit set to zero specifies that the application tag descriptor is not the last valid application tag descriptor in the Application Tag mode page.

The LOGICAL BLOCK APPLICATION TAG field specifies the value to be compared with the LOGICAL BLOCK APPLICATION TAG field associated with an LBA within the range specified by the LOGICAL BLOCK ADDRESS field and the LOGICAL BLOCK COUNT field within this descriptor.

The LOGICAL BLOCK ADDRESS field contains the starting LBA for this application tag descriptor. The LOGICAL BLOCK ADDRESS field in the first Application tag descriptor shall be set to 0000\_0000\_0000\_0000h. For subsequent application tag descriptors in which the LOGICAL BLOCK COUNT field is not set to zero, the contents of the LOGICAL BLOCK ADDRESS field shall contain the sum of the values in:

- the LOGICAL BLOCK ADDRESS field in the previous application tag descriptor; and
- the LOGICAL BLOCK COUNT field in the previous application tag descriptor.

The sum of the LOGICAL BLOCK ADDRESS field in the application tag descriptor with the LAST bit set to one and the LOGICAL BLOCK COUNT field in the application tag descriptor with the LAST bit set to one equals the RETURNED LOGICAL BLOCK ADDRESS field ~~in the READ CAPACITY (16) parameter data~~ (see 5.18.2).

If an invalid combination of the LAST bit, LOGICAL BLOCK APPLICATION TAG field, and LOGICAL BLOCK ADDRESS field are sent by the application client, then the device server shall terminate the MODE SELECT command (see SPC-5) with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.

The LOGICAL BLOCK COUNT field specifies the number of logical blocks to which this application tag descriptor applies.

A LOGICAL BLOCK COUNT field set to 0000\_0000\_0000\_0000h specifies that this application tag descriptor shall be ignored.

### 6.5.4 Background Control mode page

The Background Control mode page (see table 219) provides controls over background scan operations (see 4.24). The mode page policy (see SPC-5) for this subpage shall be shared.

**Table 219 — Background Control mode page**

Byte	Bit	7	6	5	4	3	2	1	0
0		PS	SPF (1b)	PAGE CODE (1Ch)					
1		SUBPAGE CODE (01h)							
2	(MSB)	PAGE LENGTH (000Ch)							
3									
4		Reserved					S_L_FULL	LOWIR	EN_BMS
5		Reserved							EN_PS
6	(MSB)	BACKGROUND MEDIUM SCAN INTERVAL TIME							
7									
8	(MSB)	BACKGROUND PRE-SCAN TIME LIMIT							
9									
10	(MSB)	MINIMUM IDLE TIME BEFORE BACKGROUND SCAN							
11									
12	(MSB)	MAXIMUM TIME TO SUSPEND BACKGROUND SCAN							
13									
14		Reserved							
15									

The parameters saveable (PS) bit, the subpage format (SPF) bit, the PAGE CODE field, the SUBPAGE CODE field, and the PAGE LENGTH field are defined in SPC-5.

The SPF bit, the PAGE CODE field, the SUBPAGE CODE field, and the PAGE LENGTH field shall be set to the values shown in table 219 for the Background Control mode page.

A suspend on log full (S\_L\_FULL) bit set to zero specifies that the device server shall continue running a background scan operation (see 4.24.1) even if the Background Scan Results log page (see 6.4.2) contains the maximum number of Background Scan log parameters (see 6.4.2.3) supported by the logical unit. A S\_L\_FULL bit set to one specifies that the device server shall suspend a background scan operation if the Background Scan Results log page contains the maximum number of Background scan log parameters supported by the logical unit.

A log only when intervention required (LOWIR) bit set to zero specifies that the device server shall log all suspected recoverable medium errors or unrecoverable medium errors that are identified during background scan operations in the Background Scan Results log page. A LOWIR bit set to one specifies that the device server shall only log medium errors identified during background scan operations in the Background Scan Results log page that require application client intervention.

An enable background medium scan (EN\_BMS) bit set to zero specifies that the device server shall disable background medium scan operations (see 4.24.3). An EN\_BMS bit set to one specifies that the device server shall enable background medium scan operations. If the EN\_PS bit is also set to one, and a background pre-scan operation is in progress, then the logical unit shall not start a background medium scan operation

until after the background pre-scan operation is halted or completed. If a background medium scan operation is in progress when the EN\_BMS bit is changed from one to zero, then the logical unit shall suspend the background medium scan operation before the device server completes the MODE SELECT command, and the background medium scan shall remain suspended until the EN\_BMS bit is set to one, at which time the logical unit shall resume the background medium scan operation beginning with the logical block being tested when the background medium scan operation was suspended.

An enable pre-scan (EN\_PS) bit set to zero specifies that the device server shall disable background pre-scan operations (see 4.24.2). If a background pre-scan operation is in progress when the EN\_PS bit is changed from a one to a zero, then the logical unit shall halt the background pre-scan operation before the device server completes the MODE SELECT command. An EN\_PS bit set to one specifies that the logical unit shall start a background pre-scan operation after the next power on. Once a logical unit has completed a background pre-scan operation, the logical unit shall not perform another background pre-scan operation unless the EN\_PS bit is set to zero, then set to one, and another power on occurs.

The BACKGROUND MEDIUM SCAN INTERVAL TIME field specifies the minimum time, in hours, between the start of one background scan operation and the start of the next background medium scan operation. If the current background scan operation takes longer than the value specified in the BACKGROUND MEDIUM SCAN INTERVAL TIME field, then the logical unit shall:

- a) continue the current background scan operation until that background scan operation is complete; and
- b) start the next background medium scan operation upon completion of the current background scan operation.

The BACKGROUND PRE-SCAN TIME LIMIT field specifies the maximum time, in hours, for a background pre-scan operation to complete. If the background pre-scan operation does not complete within the specified time then the device server shall halt the background pre-scan operation. A value of zero specifies an unlimited timeout value.

The MINIMUM IDLE TIME BEFORE BACKGROUND SCAN field specifies the time, in milliseconds, that the logical unit shall be idle before resuming a background scan operation (e.g., after the device server has completed all of the commands in all task sets).

The MAXIMUM TIME TO SUSPEND BACKGROUND SCAN field specifies the time, in milliseconds, that the device server should take to start processing a command received while a logical unit is performing a background scan operation.

### 6.5.5 Background Operation Control mode page

The Background Operation Control mode page (see table 220) provides controls of device server background operation.

**Table 220 — Background Operation Control mode page**

Byte	Bit	7	6	5	4	3	2	1	0
0		PS	SPF (1b)	PAGE CODE (0Ah)					
1		SUBPAGE CODE (06h)							
2		PAGE LENGTH (01FDh)							
3									
4		BO_MODE		Reserved					
5									
...		Reserved							
511									



The PS bit is described in SPC-5.

The SPF bit, PAGE CODE field, SUBPAGE CODE field, and PAGE LENGTH field are described in SPC-5 and shall be set as shown in table 220 for the Background Operation Control mode page.

The background operation mode (BO\_MODE) field specifies how host initiated advanced background operations shall operate during read operations or write operations as defined in table 221.

**Table 221 — BO\_MODE field**

Code	Description
00b	Host initiated advanced background operation shall be suspended during read operations and write operations and resume advanced background operation when read operations and write operations are complete.
01b	Host initiated advanced background operation shall continue during read operations and write operations.
All others	Reserved

### 6.5.6 Caching mode page

The Caching mode page (see table 222) defines the parameters that affect the use of the cache.

**Table 222 — Caching mode page**

Byte	Bit	7	6	5	4	3	2	1	0		
0		PS	SPF (0b)	PAGE CODE (08h)							
1		PAGE LENGTH (12h)									
2		IC	ABPF	CAP	DISC	SIZE	WCE	MF	RCD		
3		DEMAND READ RETENTION PRIORITY				WRITE RETENTION PRIORITY					
4		(MSB)	DISABLE PRE-FETCH TRANSFER LENGTH						(LSB)		
5											
6		(MSB)	MINIMUM PRE-FETCH						(LSB)		
7											
8		(MSB)	MAXIMUM PRE-FETCH						(LSB)		
9											
10		(MSB)	MAXIMUM PRE-FETCH CEILING						(LSB)		
11											
12		FSW	LBCSS	DRA	Vendor specific		SYNC_PROG		NV_DIS		
13		NUMBER OF CACHE SEGMENTS									
14		(MSB)	CACHE SEGMENT SIZE						(LSB)		
15											
16		Reserved									
17		Obsolete									
...											
19											

The parameters saveable (PS) bit, the subpage format (SPF) bit, the PAGE CODE field, and the PAGE LENGTH field are defined in SPC-5.

The SPF bit, the PAGE CODE field, and the PAGE LENGTH field shall be set to the values shown in table 222 for the Caching mode page.

An initiator control (IC) enable bit set to one specifies that the device server use one of the following fields to control the caching algorithm rather than the device server's own adaptive algorithm:

- the NUMBER OF CACHE SEGMENTS field, if the SIZE bit is set to zero; or
- the CACHE SEGMENT SIZE field, if the SIZE bit is set to one.

An abort pre-fetch (ABPF) bit set to one and a DRA bit is set to zero specify that the device server abort a pre-fetch upon receipt of a new command. An ABPF bit set to one takes precedence over the value specified in the MINIMUM PRE-FETCH field. An ABPF bit set to zero and a DRA bit set to zero specify that the termination of any active pre-fetch is dependent upon Caching mode page bytes 4 through 11 and is vendor specific.

A caching analysis permitted (CAP) bit set to one specifies that the device server perform caching analysis during subsequent operations. A CAP bit set to zero specifies that caching analysis be disabled (e.g., to reduce overhead time or to prevent nonpertinent operations from impacting tuning values).

A discontinuity (DISC) bit set to one specifies that the device server continue the pre-fetch across time discontinuities (e.g., across cylinders) up to the limits of the buffer, or segment, space available for the pre-fetch. A DISC bit set to zero specifies that pre-fetches be truncated or wrapped at time discontinuities.

A size enable (SIZE) bit set to one specifies that the CACHE SEGMENT SIZE field be used to control caching segmentation. A SIZE bit set to zero specifies that the NUMBER OF CACHE SEGMENTS field be used to control caching segmentation. Simultaneous use of both the number of segments and the segment size is vendor specific.

A writeback cache enable (WCE) bit set to one specifies that the device server shall perform write operations by using write cache operations to volatile write cache, write cache operations to non-volatile cache, or write medium operations (e.g., a write command is able to complete without error after logical block data has been written to volatile cache but has not necessarily been written to the medium) as described in 4.15. A WCE bit set to zero specifies that the device server shall complete a write command without error only after all logical block data has been written without error by performing write cache operations to non-volatile cache or by performing write medium operations to non-volatile medium. If an application client changes the WCE bit from one to zero via a MODE SELECT command, then the device server shall perform write medium operations to any LBAs in cache containing logical block data that is not the same as the logical block data referenced by the corresponding LBAs on the medium before completing the MODE SELECT command.

A multiplication factor (MF) bit set to zero specifies that the device server shall interpret the MINIMUM PRE-FETCH field and the MAXIMUM PRE-FETCH field in terms of the number of logical blocks for each of the respective types of pre-fetch. An MF bit set to one specifies that the device server shall interpret the MINIMUM PRE-FETCH field and the MAXIMUM PRE-FETCH field to be specified in terms of a scalar number that, when multiplied by the number of logical blocks to be transferred for the current command, yields the number of logical blocks for each of the respective types of pre-fetch.

A read cache disable (RCD) bit set to zero specifies that the device server shall perform read operations by using read cache operations or read medium operations as described in 4.15. An RCD bit set to one specifies that the device server shall perform read operations by only using read medium operations.

The DEMAND READ RETENTION PRIORITY field (see table 223) specifies the retention priority the device server should assign for data read into the cache that has also been transferred to the Data-In Buffer.

**Table 223 — DEMAND READ RETENTION PRIORITY field**

Code	Description
0h	The device server should not distinguish between retaining the indicated data and data placed into the cache by other means (e.g., pre-fetch).
1h	The device server should replace data put into the cache via a READ command sooner (i.e., read data has lower priority) than data placed into the cache by other means (e.g., pre-fetch).
2h to Eh	Reserved
Fh	The device server should replace data placed into the cache by other means (e.g., pre-fetch) sooner than data put into the cache via a READ command (i.e., read data has higher priority).

The WRITE RETENTION PRIORITY field (see table 224) specifies the retention priority the device server should assign for data written into the cache that has also been transferred from the cache to the medium.

**Table 224 — WRITE RETENTION PRIORITY field**

Code	Description
0h	The device server should not distinguish between retaining the indicated data and data placed into the cache by other means (e.g., pre-fetch).
1h	The device server should replace data put into the cache during a WRITE command or a WRITE AND VERIFY command sooner (i.e., has lower priority) than data placed into the cache by other means (e.g., pre-fetch).
2h to Eh	Reserved
Fh	The device server should replace data placed into the cache by other means (e.g., pre-fetch) sooner than data put into the cache during a WRITE command or a WRITE AND VERIFY command (i.e., has higher priority).

An anticipatory pre-fetch occurs when data is placed in the cache that has not been requested. This may happen in conjunction with the reading of data that has been requested. The DISABLE PRE-FETCH TRANSFER LENGTH field, the MINIMUM PRE-FETCH field, the MAXIMUM PRE-FETCH field, and the MAXIMUM PRE-FETCH CEILING field give an indication to the device server how it should manage the cache based on the most recent READ command. An anticipatory pre-fetch may occur based on other information. These fields are only recommendations to the device server and should not cause a CHECK CONDITION to occur if the device server is not able to satisfy the request.

The DISABLE PRE-FETCH TRANSFER LENGTH field specifies the selective disabling of anticipatory pre-fetch on long transfer lengths. The value in this field is compared to the transfer length requested by a READ command. If the transfer length is greater than the disable pre-fetch transfer length, then an anticipatory pre-fetch is not done for the command. Otherwise the device server should attempt an anticipatory pre-fetch. If the DISABLE PRE-FETCH TRANSFER LENGTH field is set to zero, then all anticipatory pre-fetching is disabled for any request for data, including those with a transfer length of zero.

The MINIMUM PRE-FETCH field specifies the number of logical blocks to pre-fetch regardless of the delays that may be incurred in processing subsequent commands. The field contains either:

- a) a number of logical blocks, if the MF bit is set to zero; or
- b) a scalar multiplier of the value in the TRANSFER LENGTH field, if the MF bit is set to one.

The pre-fetching operation begins at the logical block after the last logical block of a READ command. Pre-fetching shall always halt when it reaches the last logical block on the medium. Errors that occur during the pre-fetching operation shall not be reported to the application client unless the device server is unable to process subsequent commands correctly as a result of the error. In this case the error may be reported either as:

- a) an error for that subsequent command; or
- b) a deferred error,

at the discretion of the device server and according to the rules for reporting deferred errors (see SPC-5).

If the pre-fetch has read more than the amount of data specified by the MINIMUM PRE-FETCH field, then pre-fetching should be terminated whenever another command enters the enabled state (see SAM-5). This requirement is ignored if the MINIMUM PRE-FETCH field value is equal to the MAXIMUM PRE-FETCH field value.

The MAXIMUM PRE-FETCH field specifies the number of logical blocks to pre-fetch if the pre-fetch does not delay processing of subsequent commands. The field contains either:

- a) a number of logical blocks, if the MF bit is set to zero; or
- b) a scalar multiplier of the value in the TRANSFER LENGTH field, if the MF bit is set to one.

The MAXIMUM PRE-FETCH field contains the maximum amount of data to pre-fetch as a result of one READ command. The MAXIMUM PRE-FETCH field is used in conjunction with the DISABLE PRE-FETCH TRANSFER LENGTH field and MAXIMUM PRE-FETCH CEILING field to trade off pre-fetching new data with displacing old data already stored in the cache.

The MAXIMUM PRE-FETCH CEILING field specifies an upper limit on the number of logical blocks computed as the maximum pre-fetch. If this number of logical blocks is greater than the value in the MAXIMUM PRE-FETCH field, then the number of logical blocks to pre-fetch shall be truncated to the value stored in the MAXIMUM PRE-FETCH CEILING field.

NOTE 23 - If the MF bit is set to one, then the MAXIMUM PRE-FETCH CEILING field is useful in limiting the amount of data to be pre-fetched.

A force sequential write (FSW) bit set to one specifies that, for commands requesting write operations to more than one logical block, the device server shall write the logical blocks to the medium in ascending sequential order. An FSW bit set to zero specifies that the device server may reorder the sequence of writing logical blocks to the medium (e.g., in order to achieve faster command completion).

A logical block cache segment size (LBCSS) bit set to one specifies that the CACHE SEGMENT SIZE field units shall be interpreted as logical blocks. An LBCSS bit set to zero specifies that the CACHE SEGMENT SIZE field units shall be interpreted as bytes. The LBCSS shall not impact the units of other fields.

A disable read-ahead (DRA) bit set to one specifies that the device server shall not read into the pre-fetch buffer any logical blocks beyond the addressed logical block(s). A DRA bit set to zero specifies that the device server may continue to read logical blocks into the pre-fetch buffer beyond the addressed logical block(s).

The synchronize cache progress indication support (SYNC\_PROG) field (see table 225) specifies device server progress indication reporting while a SYNCHRONIZE CACHE command (see 5.28 and 5.29) is being processed.

**Table 225 — SYNC\_PROG field**

Code	Description
00b	The device server shall not terminate commands due to the synchronize cache operation and shall not provide pollable sense data.
01b	The device server: a) shall not terminate commands due to the synchronize cache operation; and b) shall provide pollable sense data with the sense key set to NO SENSE, the additional sense code set to SYNCHRONIZE CACHE OPERATION IN PROGRESS, and the PROGRESS INDICATION field set to indicate the progress of the synchronize cache operation.
10b	The device server: a) shall process INQUIRY commands, REPORT LUNS commands, REPORT TARGET PORT GROUPS commands, and REQUEST SENSE commands; b) may process commands that do not require resources used for the synchronize cache operation; c) shall terminate commands that require resources used for the synchronize cache operation with CHECK CONDITION status with the sense key set to NOT READY, the additional sense code set to LOGICAL UNIT NOT READY, SYNCHRONIZE CACHE OPERATION IN PROGRESS, and the PROGRESS INDICATION field set to indicate the progress of the synchronize cache operation; and d) shall provide pollable sense data with the sense key set to NOT READY, the additional sense code set to LOGICAL UNIT NOT READY, SYNCHRONIZE CACHE OPERATION IN PROGRESS, and the PROGRESS INDICATION field set to indicate the progress of the synchronize cache operation.
11b	Reserved

An NV\_DIS bit set to one specifies that the device server shall disable a non-volatile cache and indicates that a non-volatile cache is supported but disabled. An NV\_DIS bit set to zero specifies that the device server may use a non-volatile cache and indicates that a non-volatile cache may be present and enabled.

The NUMBER OF CACHE SEGMENTS field specifies the number of segments into which the device server shall divide the cache.

The CACHE SEGMENT SIZE field specifies the segment size in bytes if the LBCSS bit is set to zero or in logical blocks if the LBCSS bit is set to one. The CACHE SEGMENT SIZE field is valid only if the SIZE bit is set to one.

### 6.5.7 IO Advice Hints Grouping mode page

The IO Advice Hints Grouping mode page (see table 226) provides the application client with the means to obtain or modify the IO advice hints of the logical unit and the group number associated with those IO advice hints.

The mode page policy (see SPC-5) for this page shall be shared.

**Table 226 — IO Advice Hints Grouping mode page**

Byte	Bit	7	6	5	4	3	2	1	0
0		PS	SPF (1b)	PAGE CODE (0Ah)					
1		SUBPAGE CODE (05h)							
2	(MSB)	PAGE LENGTH (40Ch)							
3									
4		Reserved							
...									
15									
IO advice hints for group descriptor list									
16		IO advice hints group descriptor (group 0) (see table 227)							
...									
31									
		⋮							
1024		IO advice hints group descriptor (group 63) (see table 227)							
...									
1039									

The parameters saveable (PS) bit, the subpage format (SPF) bit, the PAGE CODE field, the SUBPAGE CODE field, and the PAGE LENGTH field are defined in SPC-5.

The SPF bit, the PAGE CODE field, the SUBPAGE CODE field, and the PAGE LENGTH field shall be set to the values shown in table 226 for the IO Advice Hints Grouping mode page.

The format of each IO advice hints group descriptor is shown in table 227. There is one IO advice hints group descriptor for each group number. The logical block markup descriptor (see 6.8) in each IO advice hints group descriptor affects the processing of commands as described in 4.23.2.

**Table 227 — IO advice hints group descriptor**

Bit	7	6	5	4	3	2	1	0
Byte								
0	IO ADVICE HINTS MODE		Reserved				CS_ENABLE	IC_ENABLE
1	Reserved							
...								
3								
4	(MSB)							
...	logical block markup descriptor (see 6.8)							
p-1	(LSB)							
p	Pad (if any)							
...								
15								

The IO ADVICE HINTS MODE field specifies the mode of the logical block markup descriptor and is described in table 228.

**Table 228 — IO ADVICE HINTS MODE field**

Code	IO advice hints usage
00b	The logical block markup descriptor is valid (see 4.23.2).
01b	The logical block markup descriptor is invalid (see 4.23.2).
all others	Reserved

The cache segment enable (CS\_ENABLE) bit specifies whether cache segments (see 4.15.2) are associated with the cache ID that is the group number associated with this IO advice hints descriptor. If the CS\_ENABLE bit is set to zero, then no independent cache segments are associated with the cache ID for the group number associated with this IO advice hints descriptor. If the CS\_ENABLE bit is set to one, then independent cache segments are associated with the cache ID for the group number associated with this IO advice hints descriptor.

The information collection enable (IC\_ENABLE) bit specifies whether the group's information collection function (see 4.23.1) is enabled. If the IC\_ENABLE bit is set to zero, then the information collection function is not enabled. If the IC\_ENABLE bit is set to one, then the information collection function is enabled.

If the Group Statistics and Performance (n) log pages are not supported, then in the IO advice hints group descriptors for group 0 to group 31, the IC\_ENABLE bit shall be set to zero and shall not be changeable.

In the IO advice hints group descriptors for group 32 to group 63, the IC\_ENABLE bit shall be set to zero, and shall not be changeable.

The logical block markup descriptor is described in 6.8.

### 6.5.8 Informational Exceptions Control mode page

The Informational Exceptions Control mode page (see table 229) defines the methods used by the device server to control the processing and reporting of informational exception conditions. Informational exception conditions are defined as any event that the device server reports or logs as failure predictions (i.e., with the ADDITIONAL SENSE CODE field set to 5Dh (e.g., FAILURE PREDICTION THRESHOLD EXCEEDED)) or warnings (i.e., with the ADDITIONAL SENSE CODE field set to 0Bh (e.g., WARNING)).

Informational exception conditions may occur while a logical unit is processing:

- a) a background self-test (see ~~SPC-4~~SPC-5);
- b) device specific background functions (see SPC-5);
- c) a command; or
- d) other device specific events.

An informational exception condition may occur at any time (e.g., the condition may be asynchronous to any commands issued by an application client).

The mode page policy for this mode page shall be shared or per I\_T nexus (see SPC-5).

NOTE 24 - Storage devices that support SMART (Self-Monitoring Analysis and Reporting Technology) for predictive failure software should use informational exception conditions.

**Table 229 — Informational Exceptions Control mode page**

Byte	Bit	7	6	5	4	3	2	1	0
0		PS	SPF (0b)	PAGE CODE (1Ch)					
1		PAGE LENGTH (0Ah)							
2		PERF	Reserved	EBF	EWASC	DEXCPT	TEST	EBACKERR	LOGERR
3		Reserved				MRIE			
4		(MSB)							
...		INTERVAL TIMER							
7		(LSB)							
8		(MSB)							
...		REPORT COUNT							
11		(LSB)							

The PS bit, the SPF bit, the PAGE CODE field, and the PAGE LENGTH field are described in SPC-5.

The SPF bit, the PAGE CODE field, and the PAGE LENGTH field shall be set to the values shown in table 229 for the Informational Exceptions Control mode page.

If the performance (PERF) bit is set to zero, then the device server may process informational exception conditions that cause delays in processing other operations (e.g., processing a command). If the PERF bit is set to one, then the device server shall not process informational exception conditions that cause delays in processing other operations. A PERF bit set to one may cause the device server to disable some or all of the processing of informational exception conditions, thereby limiting the reporting of informational exception conditions.

If device specific background functions (see SPC-5) are implemented by the logical unit, and the enable background function (EBF) bit is set to one, then the device server shall enable device specific background functions. If the EBF bit is set to zero, then the device server shall disable device specific background functions. Background functions with separate enable control bits (e.g., the background medium scan (see 4.24)) are not controlled by the EBF bit.



The enable warning (EWASC) bit specifies if the device server enables reporting of warnings (see table 230).

The disable exception control (DEXCPT) bit specifies if the device server disables reporting of failure predictions (see table 230).

The TEST bit specifies if the device server creates a test device failure prediction (see table 230).

If an informational exception condition occurs that is not the result of the logical unit processing a background self-test (see SPC-5) or device specific background function (see SPC-5), then the device server:

- a) shall use the definitions for the combination of the values in the EWASC bit, the DEXCPT bit, and the TEST bit shown in table 230 for processing informational exception conditions if the MRIE field is set from 2h to 6h;
- b) may use the definitions for the combination of the values in the EWASC bit, the DEXCPT bit, and the TEST bit shown table 230 for processing informational exception conditions if the MRIE field is set from Ch to Fh; and
- c) shall ignore the EWASC bit, the DEXCPT bit, and the TEST bit if the MRIE field is set to any value other than 2h to 6h or Ch to Fh.

**Table 230 — Definitions for the combinations of values in EWASC, DEXCPT, and TEST**

EWASC	DEXCPT	TEST	Description
0	0	0	The device server shall process informational exception conditions as follows: a) failure prediction processing shall be enabled <sup>a</sup> ; and b) warning processing shall be disabled.
1	0	0	The device server shall process informational exception conditions as follows: a) failure prediction processing shall be enabled <sup>a</sup> ; and b) warning processing shall be enabled <sup>a</sup> .
0	1	0	The device server shall process informational exception conditions as follows: a) failure prediction processing shall be disabled; and b) warning processing shall be disabled.
1	1	0	The device server shall process informational exception conditions as follows: a) failure prediction processing shall be disabled; and b) warning processing shall be enabled <sup>a</sup> .
0	0	1	The device server shall set the additional sense code to FAILURE PREDICTION THRESHOLD EXCEEDED (FALSE) <sup>a</sup> .
1	0	1	
0	1	1	The device server shall terminate the MODE SELECT command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.
1	1	1	
<sup>a</sup> If applicable based on the value in the MRIE field (e.g., 2h to 6h), then the values in the LOGERR bit, the INTERVAL TIMER field, and the REPORT COUNT field determine how the informational exception condition is processed.			

If an informational exception condition occurs while the logical unit is processing a background self-test (see SPC-5) or background function (see SPC-5), then the enable background error (EBACKERR) bit determines how the device server processes the informational exception as defined in the following:

- a) if the EBACKERR bit is set to zero, then the device server shall disable reporting of informational exception conditions that occur during the processing of background self-tests and background functions;
- b) if the EBACKERR bit is set to one, then, for informational exception conditions that occur during the processing of background self-tests and background functions, the device server shall:
  - A) enable reporting of the informational exception conditions;

- B) use the method for reporting the informational exception conditions as determined by contents of the MRIE field; and
- C) report the informational exception conditions as soon as the method specified in the MRIE field occurs (i.e., the INTERVAL TIMER field and REPORT COUNT field do not apply for background self-test errors and errors that occur during background functions);

and

- c) logging by the device server of informational exception conditions is determined by the value in the LOGERR bit.

A LOGERR bit set to zero specifies that the device server may log any informational exception conditions in the Informational Exceptions log page (see SPC-5). A LOGERR bit set to one specifies that the device server shall log informational exception conditions in the Informational Exceptions log page.

The method of reporting informational exceptions (MRIE) field (see table 231) specifies the method that shall be used by the device server to report:

- a) informational exception conditions if the specified code value is supported by the device server; and
- b) background self-test errors and device specific background function errors with the ADDITIONAL SENSE CODE field set to 0Bh or 5Dh if the EBACKERR bit is set to one and the specified code value is supported by the device server.

A device server that supports the Informational Exceptions Control mode page shall support at least one code value other than zero in the MRIE field.

The priority of reporting multiple informational exceptions is vendor specific.

**Table 231 — Method of reporting informational exceptions (MRIE) field (part 1 of 2)**

Code	Description
0h	<b>No reporting of informational exception condition:</b> The device server shall not report information exception conditions.
1h	Obsolete
2h	<b>Establish unit attention condition:</b> The device server shall report informational exception conditions by establishing a unit attention condition (see SAM-5) for the SCSI initiator port associated with every I_T nexus, with the additional sense code set to indicate the cause of the informational exception condition. <sup>a</sup>
3h	<b>Conditionally generate recovered error:</b> The device server shall report informational exception conditions, if the reporting of recovered errors is allowed <sup>b</sup> , by modifying the completion of the next command processed without encountering any errors, regardless of the I_T nexus on which the command was received. The modification shall be to terminate the command with CHECK CONDITION status with the sense key set to RECOVERED ERROR and the additional sense code set to indicate the cause of the informational exception condition.
<sup>a</sup> The device server terminates the command to report the unit attention condition for the informational exception condition (i.e., the device server does not process the command except to report the unit attention condition) (see SAM-5). <sup>b</sup> This is controlled by the PER bit <del>in the Read-Write-Error-Recovery mode page</del> (see 6.5.10) or the PER bit <del>in the Verify-Error-Recovery mode page</del> (see 6.5.11).	

Table 231 — Method of reporting informational exceptions (MRIE) field (part 2 of 2)

Code	Description
4h	<b>Unconditionally generate recovered error:</b> The device server shall report informational exception conditions, regardless of whether the reporting of recovered errors is allowed <sup>b</sup> , by modifying the completion of the next command processed without encountering any errors, regardless of the I_T nexus on which the command was received. The modification shall be to terminate the command with CHECK CONDITION status with the sense key set to RECOVERED ERROR and the additional sense code set to indicate the cause of the informational exception condition.
5h	<b>Generate no sense:</b> The device server shall report informational exception conditions by modifying the completion of the next command processed without encountering any errors, regardless of the I_T nexus on which the command was received. The modification shall be to terminate the command with CHECK CONDITION status with the sense key set to NO SENSE and the additional sense code set to indicate the cause of the informational exception condition.
6h	<b>Only report informational exception condition on request:</b> The device server shall provide pollable sense data (see SPC-5) with the sense key set to NO SENSE and the additional sense code set to indicate the cause of the informational exception condition. To find out about information exception conditions, the application client polls the device server by issuing a REQUEST SENSE command.
7h to Bh	Reserved
Ch to Fh	Vendor specific
<sup>a</sup> The device server terminates the command to report the unit attention condition for the informational exception condition (i.e., the device server does not process the command except to report the unit attention condition) (see SAM-5). <sup>b</sup> This is controlled by the PER bit <del>in the Read-Write Error Recovery mode page</del> (see 6.5.10) or the PER bit <del>in the Verify Error Recovery mode page</del> (see 6.5.11).	

The INTERVAL TIMER field specifies the period in 100 millisecond increments that the device server shall use for reporting that an informational exception condition has occurred (see table 232). After an informational exception condition has been reported, the interval timer shall be started. An INTERVAL TIMER field set to zero or FFFF\_FFFFh specifies that the period for reporting an informational exception condition is vendor specific.

The REPORT COUNT field specifies the maximum number of times the device server may report an informational exception condition to the application client. A REPORT COUNT field set to zero specifies that there is no limit on the number of times the device server may report an informational exception condition.

The device server shall use the values in the INTERVAL TIMER field and the REPORT COUNT field based on the value in the MRE field as shown in table 232.

**Table 232 — Use of the INTERVAL TIMER field and the REPORT COUNT field based on the MRE field**

MRE <sup>a</sup>	Description
2h to 6h	<p>If reporting of an informational exception condition is enabled (see table 231), then the device server shall:</p> <ol style="list-style-type: none"> <li>1) report an informational exception condition when the condition is first detected; and</li> <li>2) if the value in the REPORT COUNT field is not equal to one, then:               <ol style="list-style-type: none"> <li>1) if the INTERVAL TIMER field is not set to zero or FFFF_FFFFh, then wait the time specified in the INTERVAL TIMER field, and, if that informational exception condition still exists, report the informational exception again; and</li> <li>2) while the informational exception condition exists, continue to report the informational exception condition after waiting the time specified in the INTERVAL TIMER field until the condition has been reported the number of times specified by the REPORT COUNT field.</li> </ol> </li> </ol>
Ch to Fh	<p>The device server may use or may ignore the values in the INTERVAL TIMER field and the REPORT COUNT field to report the informational exception condition based on the device specific implementation.</p>
<p><sup>a</sup> For values in the MRE field (see table 231) not shown in this table, the INTERVAL TIMER field and the REPORT COUNT field shall be ignored.</p>	

Maintaining the interval timer and the report counter across power cycles, hard resets, logical unit resets, and I\_T nexus losses by the device server is vendor specific.

## 6.5.9 Logical Block Provisioning mode page

### 6.5.9.1 Introduction

The Logical Block Provisioning mode page (see table 233) specifies the parameters that a device server that supports logical block provisioning threshold values (see 4.7.3.7) shall use to report logical block provisioning threshold notifications (see 4.7.3.7.6). The mode page policy (see SPC-5) for this page shall be shared.

If a method not defined by this standard changes the parameter data to be returned by a device server in the Logical Block Provisioning mode page, then the device server shall establish a unit attention condition for the SCSI initiator port associated with every I\_T nexus with the additional sense code set to MODE PARAMETERS CHANGED.

**Table 233 — Logical Block Provisioning mode page**

Byte	Bit	7	6	5	4	3	2	1	0
0		PS	SPF (1b)	PAGE CODE (1Ch)					
1		SUBPAGE CODE (02h)							
2	(MSB)	PAGE LENGTH (n - 3)							
3									
4		Reserved							SITUA
5		Reserved							
...									
15									
Threshold descriptors									
16		Threshold descriptor [first] (see 6.5.9.2)							
...									
23									
		⋮							
n - 7		Threshold descriptor [last] (see 6.5.9.2)							
...									
n									

The parameters saveable (PS) bit, the subpage format (SPF) bit, the PAGE CODE field, the SUBPAGE CODE field, and the PAGE LENGTH field are defined in SPC-5.

The SPF bit, the PAGE CODE field, and the SUBPAGE CODE field shall be set to the values shown in table 233 for the Logical Block Provisioning mode page.

A single initiator threshold unit attention (SITUA) bit set to one specifies that the logical block provisioning threshold notification unit attention condition is established on a single I\_T nexus as described in 4.7.3.7.6. A SITUA bit set to zero specifies that the logical block provisioning threshold notification unit attention condition is established on multiple I\_T nexuses as described in 4.7.3.7.6.

The threshold descriptors are defined in 6.5.9.2.

### 6.5.9.2 Threshold descriptor format

The threshold descriptor format is defined in table 234.

**Table 234 — Threshold descriptor format**

Byte	Bit	7	6	5	4	3	2	1	0
0		ENABLED	Reserved	THRESHOLD TYPE			THRESHOLD ARMING		
1		THRESHOLD RESOURCE							
2		Reserved							
3									
4	(MSB)	THRESHOLD COUNT							
...									
7									

An ENABLED bit set to one specifies that the threshold is enabled. An ENABLED bit set to zero specifies that the threshold is disabled.

The THRESHOLD TYPE field (see table 235) specifies the type of this threshold.

**Table 235 — THRESHOLD TYPE field**

Code	Description
000b	If the THRESHOLD COUNT field specifies a soft threshold, the threshold is enabled, and that threshold is reached, then the device server shall establish a unit attention condition as described in 4.7.3.7.6.
001b	If the THRESHOLD COUNT field specifies a percentage threshold, the threshold is enabled, and that threshold is reached, then the device server shall establish a unit attention condition as described in 4.7.3.7.6
All others	Reserved

The THRESHOLD ARMING field (see table 236) specifies the arming method used for operation of this threshold.

**Table 236 — THRESHOLD ARMING field**

Code	Description	Reference
000b	The threshold operates as an armed decreasing threshold.	4.7.3.7.4
001b	The threshold operates as an armed increasing threshold.	4.7.3.7.5
All others	Reserved	

The THRESHOLD RESOURCE field specifies the resource of this threshold. The contents of this field are as defined for parameters codes 0000h to 00FFh (see table 180) in the Logical Block Provisioning log page (see 6.4.5).

The valid combinations of the THRESHOLD TYPE field, the THRESHOLD ARMING field, and the THRESHOLD RESOURCE field are shown in table 7.

The THRESHOLD COUNT field specifies the center of the threshold range for this threshold expressed as:

- a) a number of threshold sets (i.e., the number of LBA mapping resources expressed as a number of threshold sets), if the value in the THRESHOLD TYPE field is set to 000b; or
- b) a percentage value, if the value in the THRESHOLD TYPE field is set to 001b.

#### 6.5.10 Read-Write Error Recovery mode page

The Read-Write Error Recovery mode page (see table 237) specifies the error recovery parameters the device server shall use during:

- a) read medium operations; or
- b) write medium operations.

**Table 237 — Read-Write Error Recovery mode page**

Byte	Bit	7	6	5	4	3	2	1	0
0		PS	SPF (0b)	PAGE CODE (01h)					
1		PAGE LENGTH (0Ah)							
2		AWRE	ARRE	TB	RC	Error recovery bits			
						Obsolete	PER	DTE	Obsolete
3		READ RETRY COUNT							
4		Obsolete							
5		Obsolete							
6		Obsolete							
7		LBPERE	ReservedMWR		Reserved			Restricted for MMC-6	
8		WRITE RETRY COUNT							
9		Reserved							
10		(MSB)							
11		RECOVERY TIME LIMIT							
		(LSB)							

The parameters saveable (PS) bit, the subpage format (SPF) bit, the PAGE CODE field, and the PAGE LENGTH field are defined in SPC-5.

The SPF bit, the PAGE CODE field, and the PAGE LENGTH field shall be set to the values shown in table 237 for the Read-Write Error Recovery mode page.

An automatic write reassignment enabled (AWRE) bit set to zero specifies that the device server shall not perform automatic write reassignment.

An AWRE bit set to one specifies that the device server shall enable automatic write reassignment for LBAs referencing logical blocks for which a recovered error or unrecovered error occurs during a write medium operation. Automatic write reassignment shall be performed only if the device server has the valid data (e.g., original data in a buffer or recovered from the medium). The valid data shall be placed in the logical block referenced by the reassigned LBA. The device server shall report any failures that occur during the reassign operation. Error reporting as specified by the error recovery bits (i.e., the PER bit, and the DTE bit) shall be performed only after completion of the reassign operation.

An automatic read reassignment enabled (ARRE) bit set to zero specifies that the device server shall not perform automatic read reassignment.

An ARRE bit set to one specifies that the device server shall enable automatic read reassignment for LBAs referencing logical blocks for which a recovered error occurs during a read medium operation. All error recovery actions required by the error recovery bits shall be processed. Automatic read reassignment shall then be performed only if the device server recovers the data without error. The recovered data shall be placed in the logical block referenced by the reassigned LBA. The device server shall report any failures that occur during the reassign operation. Error reporting as specified by the error recovery bits shall be performed only after completion of the reassign operation.

A transfer block (TB) bit set to zero specifies that if an unrecovered read error occurs during a read medium operation, then the device server shall not transfer any data for that logical block to the Data-In Buffer. A TB bit set to one specifies that if an unrecovered read error occurs during a read medium operation, then the device server shall transfer pseudo read data (e.g., data already in a buffer or any other vendor specific data) for that logical block before returning CHECK CONDITION status. The data returned in this case is vendor specific. The value of the TB bit does not specify any action for recovered read errors.

A read continuous (RC) bit set to zero specifies that error recovery operations that cause delays during the data transfer are acceptable. Data shall not be fabricated.

An RC bit set to one specifies the device server shall transfer the entire requested length of data without adding delays during the data transfer to perform error recovery procedures. The device server may transfer pseudo read data in order to maintain a continuous flow of data. The device server shall assign priority to the RC bit over conflicting bits within this byte.

NOTE 25 - The RC bit set to one is useful for image processing, audio, or video applications.

A post error (PER) bit set to one specifies that if a recovered read error occurs while processing a read command or a write command, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to RECOVERED ERROR. A PER bit set to zero specifies that if a recovered read error occurs while processing a read command or a write command, then the device server shall perform error recovery procedures within the limits established by the error recovery parameters and not terminate the specified command with CHECK CONDITION status with the sense key set to RECOVERED ERROR (e.g., the device server may terminate the specified command with CHECK CONDITION status with the sense key set to MEDIUM ERROR if an uncorrectable error is detected based on the established limits during the error recovery process). If the DTE bit is set to one, then the PER bit shall be set to one.

A data terminate on error (DTE) bit set to one specifies that, upon detection of a recovered error, the device server shall terminate the data transfer to the Data-In Buffer for a read command or the data transfer to the Data-Out Buffer for a write command upon detection of a recovered error. A DTE bit set to zero specifies that, upon detection of a recovered error, the device server shall not terminate the data transfer to the Data-In Buffer for a read command or the data transfer to the Data-Out Buffer for a write command.



The combinations of the error recovery bits (i.e., the PER bit, and the DTE bit) are defined in table 238.

**Table 238 — Error recovery bit combinations (part 1 of 2)**

PER	DTE	Description <sup>b</sup>
0	0	<p>The device server shall perform the full number of retries as specified in the READ RETRY COUNT field for read medium operations, the WRITE RETRY COUNT field for write medium operations, and the VERIFY RETRY COUNT field (see 6.5.11) for verify medium operations and shall perform error correction in an attempt to recover the data.</p> <p>The device server shall not report recovered read errors or write errors. The device server shall terminate a command performing a read medium operation or a write medium operation with CHECK CONDITION status before the transfer count is exhausted only if an unrecovered error is detected.</p> <p>If an unrecovered read error occurs during a read medium operation, then the transfer block (TB) bit determines whether the data for the logical block with the unrecovered read error is transferred to the Data-In Buffer.</p>
		Invalid mode. The PER bit shall be set to one if the DTE bit is set to one. <sup>a</sup>
1	0	<p>The device server shall perform the full number of retries as specified in the READ RETRY COUNT field for read medium operations, the VERIFY RETRY COUNT field (see 6.5.11) for verify medium operations, the WRITE RETRY COUNT field for write medium operations, and shall perform error correction in an attempt to recover the data.</p> <p>The device server shall terminate a command performing a read medium operation or write medium operation with CHECK CONDITION status before the transfer count is exhausted only if an unrecovered error is detected.</p> <p>If an unrecovered read error occurs during a read medium operation, the transfer block (TB) bit determines whether the data for the logical block with the unrecovered read error is transferred to the Data-In Buffer.</p> <p>If a recovered error occurs while the device server is performing a read medium operation or write medium operation, then, after the operation is complete, the device server shall terminate the command with CHECK CONDITION status with the sense key set to RECOVERED ERROR. The INFORMATION field in the sense data contains the LBA of the last recovered error that occurred during the command (see 4.18.1).</p>
<p><sup>a</sup> If an invalid combination of the error recovery bits is sent by an application client, then the device server shall terminate the MODE SELECT command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.</p> <p><sup>b</sup> This table is used by both the Read Write Error Recovery mode page and the Verify Error Recovery mode page (see 6.5.11). When used for the Read Write Error Recovery mode page, rules about the VERIFY RETRY COUNT field are not applicable. When used for the Verify Error Recovery mode page, rules about the READ RETRY COUNT field, the WRITE RETRY COUNT field, and write medium operations are not applicable and rules about read medium operations are applicable to verify medium operations.</p>		

**Table 238 — Error recovery bit combinations (part 2 of 2)**

PER	DTE	Description <sup>b</sup>
1	1	<p>The device server shall perform the full number of retries as specified in the READ RETRY COUNT field for read medium operations, the VERIFY RETRY COUNT field (see 6.5.11) for verify medium operations, the WRITE RETRY COUNT field for write medium operations, and shall perform error correction in an attempt to recover the data.</p> <p>The device server shall terminate a command performing a read medium operation or write medium operation with CHECK CONDITION status before the transfer count is exhausted if any error, either recovered or unrecovered, is detected. The INFORMATION field in the sense data contains the LBA of the error (see 4.18.1).</p> <p>If an unrecovered read error occurs during a read medium operation, the transfer block (TB) bit determines whether the data for the logical block with the unrecovered read error is transferred to the Data-In Buffer.</p>
<p><sup>a</sup> If an invalid combination of the error recovery bits is sent by an application client, then the device server shall terminate the MODE SELECT command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to INVALID FIELD IN PARAMETER LIST.</p> <p><sup>b</sup> This table is used by both the Read Write Error Recovery mode page and the Verify Error Recovery mode page (see 6.5.11). When used for the Read Write Error Recovery mode page, rules about the VERIFY RETRY COUNT field are not applicable. When used for the Verify Error Recovery mode page, rules about the READ RETRY COUNT field, the WRITE RETRY COUNT field, and write medium operations are not applicable and rules about read medium operations are applicable to verify medium operations.</p>		

The READ RETRY COUNT field specifies the number of times that the device server shall attempt its recovery algorithm during read medium operations.

The WRITE RETRY COUNT field specifies the number of times that the device server shall attempt its recovery algorithm during write medium operations.

A logical block provisioning error reporting enabled (LBPERE) bit set to one specifies that logical block provisioning threshold notification is enabled. A LBPERE bit set to zero specifies that logical block provisioning threshold notification is disabled (see 4.7.3.7.6).

[A misaligned write reporting \(MWR\) field \(see table 239\) specifies the behavior of the device server with respect to physical block misalignment reporting \(see 4.6.2\).](#)

**Table 239 — [MWR field](#)**

<a href="#">Code</a>	<a href="#">Name</a>	<a href="#">Description</a>
<a href="#">00b</a>	<a href="#">DISABLED</a>	<a href="#">Complete and do not report misaligned write operations</a>
<a href="#">01b</a>	<a href="#">ENABLED</a>	<a href="#">Complete and report misaligned write operations</a>
<a href="#">10b</a>	<a href="#">TERMINATE</a>	<a href="#">Terminate amd report misaligned write operations</a>
<a href="#">11b</a>	<a href="#">Reserved</a>	

A RECOVERY TIME LIMIT field set to a non-zero value specifies in milliseconds the maximum time duration that the device server shall use for data error recovery procedures during read medium operations and during write medium operations. The device server may round this value as described in SPC-5. The limit in this field specifies the maximum error recovery time allowed for any individual logical block. A RECOVERY TIME LIMIT field set to zero specifies that the device server shall use its default value.

If both a retry count and a recovery time limit are specified, then the field that specifies the recovery action of least duration shall have priority.

To disable all types of correction and retries, the application client should set:

- a) the PER bit to one;
- b) the DTE bit to one;
- c) the READ RETRY COUNT field to 00h;
- d) the WRITE RETRY COUNT field to 00h; and
- e) the RECOVERY TIME LIMIT field to 0000h.

#### 6.5.11 Verify Error Recovery mode page

The Verify Error Recovery mode page (see table 240) specifies the error recovery parameters the device server shall use during verify medium operations (e.g., from VERIFY commands and the verify medium operations of the WRITE AND VERIFY commands). Verify medium operations do not trigger automatic read reassignment.

**Table 240 — Verify Error Recovery mode page**

Byte	Bit	7	6	5	4	3	2	1	0
0		PS	SPF (0b)	PAGE CODE (07h)					
1		PAGE LENGTH (0Ah)							
2		Reserved				Error recovery bits			
						Obsolete	PER	DTE	Obsolete
3		VERIFY RETRY COUNT							
4		Obsolete							
5		Reserved							
...									
9									
10		(MSB)	VERIFY RECOVERY TIME LIMIT						
11									(LSB)

The parameters saveable (PS) bit, the subpage format (SPF) bit, the PAGE CODE field, and the PAGE LENGTH field are defined in SPC-5.

The SPF bit, the PAGE CODE field, and the PAGE LENGTH field shall be set to the values shown in table 240 for the Verify Error Recovery mode page.

The PER bit, and the DTE bit (i.e., the error recovery bits) are defined in 6.5.10. The combinations of these bits are defined in table 238.

The VERIFY RETRY COUNT field specifies the number of times that the device server shall attempt its recovery algorithm during a verify medium operation.

The VERIFY RECOVERY TIME LIMIT field specifies in milliseconds the maximum time duration that the device server shall use error recovery procedures to recover data for an individual logical block during a verify medium operation. The device server may round this value as described in SPC-5.

If both a verify retry count and a verify recovery time limit are specified, then the one that requires the least time for data error recovery actions shall have priority.

To disable all types of correction and retries, the application client should set:

- a) the PER bit to one;
- b) the DTE bit to one;
- c) the VERIFY RETRY COUNT field to 00h; and
- d) the VERIFY RECOVERY TIME LIMIT field to 0000h.

## 6.6 Vital product data (VPD) parameters

### 6.6.1 VPD ~~parameters~~ pages overview

See table 241 for references to the VPD pages used ~~only~~ with direct access block devices. ~~See SPC-5 for VPD pages used with all device types. The VPD pages and their corresponding page codes specific to direct access block devices are defined in table 241.~~

Table 241 — VPD page codes for direct access block devices

VPD page name	Page code <sup>a</sup>	Reference	Support- requirements
<a href="#">ASCII Information</a>	<a href="#">01h to 7Fh</a>	<a href="#">SPC-5</a>	<a href="#">See SPC-5</a>
ATA Information- <del>VPD page</del>	89h	SAT-3	See SAT-3
Block Device Characteristics- <del>VPD page</del>	B1h	6.6.2	Optional
Block Device Characteristics Extension- <del>VPD page</del>	B5h	6.6.3	Optional
Block Limits- <del>VPD page</del>	B0h	6.6.4	Optional
Block Limits Extension- <del>VPD page</del>	B7h	6.6.5	Optional
<a href="#">CFA Profile Information</a>	<a href="#">8Ch</a>	<a href="#">SPC-5</a>	<a href="#">See SPC-5</a>
<a href="#">Device Constituents</a>	<a href="#">8Bh</a>	<a href="#">SPC-5</a>	<a href="#">See SPC-5</a>
<a href="#">Device Identification</a>	<a href="#">83h</a>	<a href="#">SPC-5</a>	<a href="#">See SPC-5</a>
<a href="#">Extended INQUIRY Data</a>	<a href="#">86h</a>	<a href="#">SPC-5</a>	<a href="#">See SPC-5</a>
Logical Block Provisioning- <del>VPD page</del>	B2h	6.6.6	Optional
<a href="#">Management Network Addresses</a>	<a href="#">85h</a>	<a href="#">SPC-5</a>	<a href="#">See SPC-5</a>
<a href="#">Mode Page Policy</a>	<a href="#">87h</a>	<a href="#">SPC-5</a>	<a href="#">See SPC-5</a>
<a href="#">Power Condition</a>		<a href="#">SPC-5</a>	<a href="#">See SPC-5</a>
Power Consumption- <del>VPD page</del>	8Dh	SPC-5	See SPC-5
<a href="#">Protocol Specific Logical Unit Information</a>	<a href="#">90h</a>	<a href="#">SPC-5</a>	<a href="#">See SPC-5</a>
<a href="#">Protocol Specific Port Information</a>	<a href="#">91h</a>	<a href="#">SPC-5</a>	<a href="#">See SPC-5</a>
Referrals- <del>VPD page</del>	B3h	6.6.7	Optional
<a href="#">SCSI Ports</a>	<a href="#">88h</a>	<a href="#">SPC-5</a>	<a href="#">See SPC-5</a>
<a href="#">Software Interface Identification</a>	<a href="#">84h</a>	<a href="#">SPC-5</a>	<a href="#">See SPC-5</a>
<a href="#">Supported VPD Pages</a>	<a href="#">00h</a>	<a href="#">SPC-5</a>	<a href="#">See SPC-5</a>
Third-Party Copy- <del>VPD page</del>	8Fh	SPC-5 and 6.6.8	<del>See</del> <sup>b</sup> <a href="#">Optional</a>
Supported Block Lengths And Protection Types- <del>VPD page</del>	B4h	6.6.9	Optional
<a href="#">Unit Serial Number</a>	<a href="#">80h</a>	<a href="#">SPC-5</a>	<a href="#">See SPC-5</a>
<a href="#">Zoned Block Device Characteristics</a>	<a href="#">B6h</a>	<a href="#">ZBC</a>	<a href="#">See ZBC</a>
Reserved for this standard	B8h to BFh		
<sup>a</sup> All page codes for direct access block devices not shown in this table are reserved. <sup>b</sup> <del>Mandatory if the POPULATE TOKEN command and the WRITE USING TOKEN command are supported.</del>			

### 6.6.2 Block Device Characteristics VPD page

The Block Device Characteristics VPD page (see table 242) contains parameters indicating characteristics of the logical unit.

**Table 242 — Block Device Characteristics VPD page**

Byte	Bit	7	6	5	4	3	2	1	0
0		PERIPHERAL QUALIFIER			PERIPHERAL DEVICE TYPE				
1		PAGE CODE (B1h)							
2	(MSB)	PAGE LENGTH (003Ch)							
3									
4	(MSB)	MEDIUM ROTATION RATE							
5									
6		PRODUCT TYPE							
7		WABEREQ		WACEREQ		NOMINAL FORM FACTOR			
8		Reserved		ZONED		Reserved	BOCS	FUAB	VBULS
9		Reserved							
...									
63									

The PERIPHERAL QUALIFIER field and PERIPHERAL QUALIFIER field are defined in SPC-5.

The PAGE CODE field and PAGE LENGTH field are defined in SPC-5 and shall be set to the values shown in table 242 for the Block Device Characteristics VPD page.

The MEDIUM ROTATION RATE field is defined in table 243.

**Table 243 — MEDIUM ROTATION RATE field**

Code	Description
0000h	Medium rotation rate is not reported
0001h	Non-rotating medium (e.g., solid state)
0002h to 0400h	Reserved
0401h to FFFEh	Nominal medium rotation rate in revolutions per minute (e.g., 7 200 rpm = 1C20h, 10 000 rpm = 2710h, and 15 000 rpm = 3A98h)
FFFFh	Reserved

The PRODUCT TYPE field (see table 244) defines the product type of the storage device.

**Table 244 — PRODUCT TYPE field**

Code	Description
00h	Not indicated
01h	CFast™ (see CFast)
02h	CompactFlash® (see CF)
03h	Memory Stick™ (see MS)
04h	MultiMediaCard (see e•MMC)
05h	Secure Digital Card (see SD Card)
06h	XQD™ (see XQD)
07h	Universal Flash Storage (see UFS)
08h to EFh	Reserved
F0h to FFh	Vendor specific

The write after block erase required (WABEREQ) field indicates the device server behavior (see table 245), if a write operation has not been performed to a mapped LBA since a sanitize block erase operation was performed and no other error occurs during the processing of a read command specifying that LBA.

**Table 245 — WABEREQ field**

Code	Description
00b	Not specified.
01b	The device server completes the read command specifying that LBA with GOOD status and any data transferred to the Data-In Buffer is indeterminate.
10b	The device server terminates the read command specifying that LBA with CHECK CONDITION status with sense key set to MEDIUM ERROR and the additional sense code set to an appropriate value other than WRITE AFTER SANITIZE REQUIRED (e.g., ID CRC OR ECC ERROR).
11b	The device server terminates the read command specifying that LBA with CHECK CONDITION status with sense key set to MEDIUM ERROR and the additional sense code set to WRITE AFTER SANITIZE REQUIRED.



The write after cryptographic erase required (WACREQ) field indicates the device server behavior (see table 246), if a write operation has not been performed to a mapped LBA since a sanitize cryptographic erase operation was performed and no other error occurs during the processing of a read command specifying that LBA, ~~then~~.

**Table 246 — WACREQ field**

Code	Description
00b	Not specified.
01b	The device server completes the read command specifying that LBA with GOOD status and any data transferred to the Data-In Buffer is indeterminate.
10b	The device server terminates the read command specifying that LBA with CHECK CONDITION status with sense key set to MEDIUM ERROR and the additional sense code set to an appropriate value other than WRITE AFTER SANITIZE REQUIRED (e.g., ID CRC OR ECC ERROR).
11b	The device server terminates the read command specifying that LBA with CHECK CONDITION status with sense key set to MEDIUM ERROR and the additional sense code set to WRITE AFTER SANITIZE REQUIRED.

The NOMINAL FORM FACTOR field indicates the nominal form factor of the device containing the logical unit and is defined in table 247.

**Table 247 — NOMINAL FORM FACTOR field**

Code	Description
0h	Nominal form factor is not reported
1h	5.25 inch
2h	3.5 inch
3h	2.5 inch
4h	1.8 inch
5h	Less than 1.8 inch
All others	Reserved

The ZONED field indicates the type of zoned block capabilities implemented by the device server as defined in table 248.

**Table 248 — ZONED field**

code	Description
00b	Not reported
01b	Device server implements the host aware zoned block device capabilities defined in ZBC (see ZBC)
10b	Device server implements device managed zoned block device capabilities
11b	Reserved

A background operation control supported (BOCS) bit set to one indicates that Background Operation Control is supported as described in 4.33. A BOCS bit set to zero indicates that Background Operation Control is not supported.

A force unit access behavior (FUAB) bit set to one indicates that the device server interprets the SYNCHRONIZE CACHE command and the FUA bit in read commands and write commands in compliance with this standard. An FUAB bit set to zero indicates that the device server interprets the SYNCHRONIZE CACHE command and the FUA bit in read commands and write commands in compliance with SBC-2.

A verify byte check unmapped LBA supported (VBULS) bit set to one indicates that the device server supports unmapped LBAs while processing VERIFY commands (see 5.31, 5.32, 5.33, and 5.34) and WRITE AND VERIFY commands (see 5.39, 5.40, 5.41, and 5.42) with the BYCHK field set to 01b. A VBULS bit set to zero indicates that the device server does not support unmapped LBAs while processing VERIFY commands and WRITE AND VERIFY commands with the BYCHK field set to 01b. The device server should set the VBULS bit to one.

### 6.6.3 Block Device Characteristics Extension VPD page

The Block Device Characteristics Extension VPD page (see table 249) contains parameters indicating characteristics of the logical unit.

**Table 249 — Block Device Characteristics Extension VPD page**

Byte	Bit	7	6	5	4	3	2	1	0
0		PERIPHERAL QUALIFIER			PERIPHERAL DEVICE TYPE (00000b)				
1		PAGE CODE (B5h)							
2	(MSB)	PAGE LENGTH (007Ch)							
3									
4		RESERVED							
5		UTILIZATION TYPE							
6		UTILIZATION UNITS							
7		UTILIZATION INTERVAL							
8	(MSB)	UTILIZATION B							
...									
11		(LSB)							
12	(MSB)	UTILIZATION A							
...									
15		(LSB)							
16		Reserved							
...									
127									

The PERIPHERAL QUALIFIER field is defined in SPC-5.

The PAGE CODE field, PERIPHERAL DEVICE TYPE field, and PAGE LENGTH field are defined in SPC-5 and shall be set to the values shown in table 249 for the Block Device Characteristics Extension VPD page.

The UTILIZATION TYPE field (see table 250) indicates the designed utilization characteristics for the direct access block device based on the contents of the UTILIZATION A field and the UTILIZATION B field evaluated

using the units indicated by the UTILIZATION UNITS field over the time interval indicated by the UTILIZATION INTERVAL field.

**Table 250 — UTILIZATION TYPE field**

Code	Description
01h	<b>Combined writes and reads:</b> the UTILIZATION A field contains designed number of host requested bytes transferred by write operations and host requested bytes transferred by read operations. The UTILIZATION B field is reserved.
02h	<b>Writes only:</b> the UTILIZATION A field contains designed number of host requested bytes transferred by write operations. The UTILIZATION B field is reserved.
03h	<b>Separate writes and reads:</b> the UTILIZATION A field contains designed number of host requested bytes transferred by write operations. The UTILIZATION B field contains designed number of host requested bytes transferred by read operations.
all others	Reserved

The UTILIZATION UNITS field (see table 251) indicates the units of measure for the values, if any, in the UTILIZATION A field and the UTILIZATION B field.

**Table 251 — UTILIZATION UNITS field**

Code	Description
02h	megabytes
03h	gigabytes
04h	terabytes
05h	petabytes
06h	exabytes
all others	Reserved

The UTILIZATION INTERVAL field (see table 252) indicates a nominal calendar time reference interval over which the values, if any, in the UTILIZATION A field and the UTILIZATION B field may be applied.

**Table 252 — UTILIZATION INTERVAL field**

Code	Description
0Ah	per day
0Eh	per year
all others	Reserved

The UTILIZATION A field and the UTILIZATION B field indicate the designed utilization characteristics for the direct access block device as:

- a) defined by the utilization type field;
- b) expressed in the units defined by the utilization units field; and
- c) over the time interval defined by the utilization interval field.

### 6.6.4 Block Limits VPD page

The Block Limits VPD page (see table 253) provides the application client with the means to obtain certain operating parameters of the logical unit.

**Table 253 — Block Limits VPD page (part 1 of 2)**

Byte	Bit	7	6	5	4	3	2	1	0
0		PERIPHERAL QUALIFIER			PERIPHERAL DEVICE TYPE				
1		PAGE CODE (B0h)							
2	(MSB)	PAGE LENGTH (003Ch)							
3									
4		Reserved							WSNZ
5		MAXIMUM COMPARE AND WRITE LENGTH							
6	(MSB)	OPTIMAL TRANSFER LENGTH GRANULARITY							
7									
8	(MSB)	MAXIMUM TRANSFER LENGTH							
...									
11									(LSB)
12	(MSB)	OPTIMAL TRANSFER LENGTH							
...									
15									(LSB)
16	(MSB)	MAXIMUM PREFETCH LENGTH							
...									
19									(LSB)
20	(MSB)	MAXIMUM UNMAP LBA COUNT							
...									
23									(LSB)
24	(MSB)	MAXIMUM UNMAP BLOCK DESCRIPTOR COUNT							
...									
27									(LSB)
28	(MSB)	OPTIMAL UNMAP GRANULARITY							
...									
31									(LSB)
32	UGAVALID	(MSB)	UNMAP GRANULARITY ALIGNMENT						
...									
35									(LSB)
36	(MSB)	MAXIMUM WRITE SAME LENGTH							
...									
43									(LSB)
44	(MSB)	MAXIMUM ATOMIC TRANSFER LENGTH							
...									
47									(LSB)

**Table 253 — Block Limits VPD page (part 2 of 2)**

Byte	Bit	7	6	5	4	3	2	1	0
48	(MSB)	ATOMIC ALIGNMENT							
...									
51									
52	(MSB)	ATOMIC TRANSFER LENGTH GRANULARITY							
...									
55									
56	(MSB)	MAXIMUM ATOMIC TRANSFER LENGTH WITH ATOMIC BOUNDARY							
...									
59									
60	(MSB)	MAXIMUM ATOMIC BOUNDARY SIZE							
...									
63									

The PERIPHERAL QUALIFIER field and PERIPHERAL QUALIFIER field are defined in SPC-5.

The PAGE CODE field and PAGE LENGTH field are defined in SPC-5 and shall be set to the values shown in table 253 for the Block Limits VPD page.

A write same non-zero (WSNZ) bit set to one indicates that the device server does not support a value of zero in the NUMBER OF LOGICAL BLOCKS field in the WRITE SAME command CDBs (see 5.47, 5.48, and 5.49). A WSNZ bit set to zero indicates that the device server may or may not support a value of zero in the NUMBER OF LOGICAL BLOCKS field of the WRITE SAME commands.

A MAXIMUM COMPARE AND WRITE LENGTH field set to a non-zero value indicates the maximum value that the device server accepts in the NUMBER OF LOGICAL BLOCKS field in the COMPARE AND WRITE command (see 5.3). A MAXIMUM COMPARE AND WRITE LENGTH field set to zero indicates that the device server does not support the COMPARE AND WRITE command. If the MAXIMUM TRANSFER LENGTH field is not set to zero, then the device server shall set the MAXIMUM COMPARE AND WRITE LENGTH field to a value less than or equal to the value in the MAXIMUM TRANSFER LENGTH field.

An OPTIMAL TRANSFER LENGTH GRANULARITY field set to a non-zero value indicates the optimal transfer length granularity size in logical blocks for a single command shown in the command column of table 254. If a device server receives one of these commands with a transfer size that is not equal to a multiple of this value, then the device server may incur delays in processing the command. An OPTIMAL TRANSFER LENGTH GRANULARITY field set to 0000h indicates that the device server does not report optimal transfer length granularity.

Table 254 — Transfer limits for commands

Command	Field that specifies the transfer size	Block Limits VPD page field(s) that indicate maximum limits	Additional sense code if the value in the specified field exceeds the maximum limit
COMPARE AND WRITE	CDB NUMBER OF LOGICAL BLOCKS field	MAXIMUM COMPARE AND WRITE LENGTH field	INVALID FIELD IN CDB
ORWRITE	CDB TRANSFER LENGTH field	MAXIMUM TRANSFER LENGTH field	
PRE-FETCH	CDB PREFETCH LENGTH field	MAXIMUM PREFETCH LENGTH field	
READ	CDB TRANSFER LENGTH field	MAXIMUM TRANSFER LENGTH field	
VERIFY	CDB VERIFICATION LENGTH field		
WRITE	CDB TRANSFER LENGTH field		
WRITE AND VERIFY			
XDWRITEREAD			
XPWRITE			
Any individual block device range descriptor in a POPULATE TOKEN command (see 5.9)	Block device range descriptor NUMBER OF LOGICAL BLOCKS field	MAXIMUM TRANSFER LENGTH field	INVALID FIELD IN PARAMETER LIST
Any individual block device range descriptor in a WRITE USING TOKEN command (see 5.52)			

A MAXIMUM TRANSFER LENGTH field set to a non-zero value indicates the maximum transfer length in logical blocks that the device server accepts for a single command shown in table 254. If a device server receives one of these commands with a transfer size greater than this value, then the device server shall terminate the command with CHECK CONDITION status with the sense key set to ILLEGAL REQUEST and the additional sense code set to the value shown in table 254. A MAXIMUM TRANSFER LENGTH field set to 0000\_0000h indicates that the device server does not report a limit on the transfer length.

An OPTIMAL TRANSFER LENGTH field set to a non-zero value indicates the optimal transfer size in logical blocks for a single command shown in table 254. If a device server receives one of these commands with a transfer size greater than this value, then the device server may incur delays in processing the command. An OPTIMAL TRANSFER LENGTH field set to 0000\_0000h indicates that the device server does not report an optimal transfer size.

The MAXIMUM PREFETCH LENGTH field indicates the maximum prefetch length in logical blocks that the device server accepts for a single PRE-FETCH command. If the MAXIMUM TRANSFER LENGTH field is not set to zero,

then the device server should set the MAXIMUM PREFETCH LENGTH field to a value less than or equal to the value in the MAXIMUM TRANSFER LENGTH field.

A MAXIMUM UNMAP LBA COUNT field set to a non-zero value indicates the maximum number of LBAs that may be unmapped by an UNMAP command (see 5.30). If the number of LBAs that may be unmapped by an UNMAP command is constrained only by the amount of data that may be contained in the UNMAP parameter list (see 5.30.2), then the device server shall set the MAXIMUM UNMAP LBA COUNT field to FFFF\_FFFFh. If the device server implements the UNMAP command, then the value in this field shall be greater than or equal to one. A MAXIMUM UNMAP LBA COUNT field set to 0000\_0000h indicates that the device server does not implement the UNMAP command.

A MAXIMUM UNMAP BLOCK DESCRIPTOR COUNT field set to a non-zero value indicates the maximum number of UNMAP block descriptors (see 5.30.2) that shall be contained in the parameter data transferred to the device server for an UNMAP command (see 5.30). If there is no limit on the number of UNMAP block descriptors contained in the parameter data, then the device server shall set the MAXIMUM UNMAP BLOCK DESCRIPTOR COUNT field to FFFF\_FFFFh. If the device server implements the UNMAP command, then the value in the MAXIMUM UNMAP BLOCK DESCRIPTOR COUNT field shall be greater than or equal to one. A MAXIMUM UNMAP BLOCK DESCRIPTOR COUNT field set to 0000\_0000h indicates that the device server does not implement the UNMAP command.

An OPTIMAL UNMAP GRANULARITY field set to a non-zero value indicates the optimal granularity in logical blocks for unmap requests (e.g., an UNMAP command or a WRITE SAME (16) command with the UNMAP bit set to one). An unmap request with a number of logical blocks that is not a multiple of this value may result in unmap operations on fewer LBAs than requested. An OPTIMAL UNMAP GRANULARITY field set to 0000\_0000h indicates that the device server does not report an optimal unmap granularity.

An unmap granularity alignment valid (UGAVALID) bit set to one indicates that the UNMAP GRANULARITY ALIGNMENT field is valid. A UGAVALID bit set to zero indicates that the UNMAP GRANULARITY ALIGNMENT field is not valid.

The UNMAP GRANULARITY ALIGNMENT field indicates the LBA of the first logical block to which the OPTIMAL UNMAP GRANULARITY field applies. The unmap granularity alignment is used to calculate an optimal unmap request starting LBA as follows:

$$\text{optimal unmap request starting LBA} = (n \times \text{optimal unmap granularity}) + \text{unmap granularity alignment}$$

where:

$n$  is zero or any positive integer value;

optimal unmap granularity is the value in the OPTIMAL UNMAP GRANULARITY field; and

unmap granularity alignment is the value in the UNMAP GRANULARITY ALIGNMENT field.

An unmap request with a starting LBA that is not optimal may result in unmap operations on fewer LBAs than requested.

A MAXIMUM WRITE SAME LENGTH field set to a non-zero value indicates the maximum number of contiguous logical blocks that the device server allows to be unmapped or written in a single WRITE SAME command. A MAXIMUM WRITE SAME LENGTH field set to zero indicates that the device server does not report a limit on the number of logical blocks that it allows to be unmapped or written in a single WRITE SAME command.

If the ATOMIC BOUNDARY field in the CDB (see 5.43 and 5.44) is set to zero, then a MAXIMUM ATOMIC TRANSFER LENGTH field set to a non-zero value indicates the maximum atomic transfer length in logical blocks that the device server supports for a single atomic write command (see 4.31). A MAXIMUM ATOMIC TRANSFER LENGTH field set to 0000\_0000h indicates that the device server does not indicate a maximum atomic transfer length. The maximum atomic transfer length indicated by the MAXIMUM ATOMIC TRANSFER LENGTH field shall be less than or equal to the maximum transfer length indicated by the MAXIMUM TRANSFER LENGTH field. The maximum atomic transfer length indicated by the MAXIMUM ATOMIC TRANSFER LENGTH field shall be a multiple of the value in the ATOMIC TRANSFER LENGTH GRANULARITY field. If the ATOMIC BOUNDARY field in the CDB is set to a non-zero value, then the MAXIMUM ATOMIC TRANSFER LENGTH field is ignored.

The ATOMIC ALIGNMENT field indicates the required alignment of the starting LBA in an atomic write command. If the ATOMIC ALIGNMENT field is set to 0000\_0000h, then there is no alignment requirement for atomic write commands.

If the ATOMIC ALIGNMENT field is non-zero, then the starting LBA of an atomic write request shall meet the following:

$$\text{atomic request starting LBA} = n \times \text{atomic alignment}$$

where:

$n$  is zero or any positive integer value; and  
 atomic alignment is the value in the ATOMIC ALIGNMENT field,

The ATOMIC TRANSFER LENGTH GRANULARITY field indicates the minimum transfer length for a WRITE ATOMIC command. Atomic write operations are required to have a transfer length that is a multiple of the atomic transfer length granularity. An ATOMIC TRANSFER LENGTH GRANULARITY field set to 0000\_0000h indicates that there is no atomic transfer length granularity requirement.

If the ATOMIC BOUNDARY field in the CDB (see 5.43 and 5.44) is set to a non-zero value, then a MAXIMUM ATOMIC TRANSFER LENGTH WITH ATOMIC BOUNDARY field set to a non-zero value indicates the maximum transfer length in logical blocks that the device server supports for a single atomic write command (see 4.31). A MAXIMUM ATOMIC TRANSFER LENGTH WITH ATOMIC BOUNDARY field set to 0000\_0000h indicates that the device server does not indicate a maximum atomic transfer length with atomic boundary. The maximum atomic transfer length with atomic boundary indicated by the MAXIMUM ATOMIC TRANSFER LENGTH WITH ATOMIC BOUNDARY field shall be less than or equal to the maximum transfer length indicated by the MAXIMUM TRANSFER LENGTH field. The maximum atomic transfer length with atomic boundary indicated by the MAXIMUM ATOMIC TRANSFER LENGTH WITH BOUNDARY field shall be a multiple of the value in the ATOMIC TRANSFER LENGTH GRANULARITY field. If the ATOMIC BOUNDARY field in the CDB is set to zero, then the MAXIMUM ATOMIC TRANSFER LENGTH WITH ATOMIC BOUNDARY field is ignored.

A MAXIMUM ATOMIC BOUNDARY SIZE field set to a non-zero value indicates that the device server supports atomic write commands performing more than one atomic operation. The maximum atomic boundary size indicates the maximum number of logical blocks on which the device server is able to perform atomic operations (see 4.31.1). A MAXIMUM ATOMIC BOUNDARY SIZE field set to 0000h indicates that the device server does not support atomic write commands performing more than one atomic operation. The MAXIMUM ATOMIC BOUNDARY SIZE field shall be a multiple of the value in the ATOMIC TRANSFER LENGTH GRANULARITY field.

### 6.6.5 Block Limits Extension VPD page

The Block Limits Extension VPD page (see table 255) provides the application client with the means to obtain certain operating parameters of the logical unit.

**Table 255 — Block Limits Extension VPD page (part 1 of 2)**

Bit	7	6	5	4	3	2	1	0
Byte								
0	PERIPHERAL QUALIFIER			PERIPHERAL DEVICE TYPE				
1	PAGE CODE (B7h)							
2	(MSB)							
3	PAGE LENGTH (n-4)						(LSB)	
4								
5	Reserved							



Table 255 — Block Limits Extension VPD page (part 2 of 2)

Byte	Bit	7	6	5	4	3	2	1	0
6	(MSB)	MAXIMUM NUMBER OF STREAMS							
7									
8	(MSB)	OPTIMAL STREAM WRITE SIZE							
9									
10	(MSB)	STREAM GRANULARITY SIZE							
...									
13									
14		Reserved							
...									
n									

The PERIPHERAL QUALIFIER field and PERIPHERAL DEVICE TYPE field are defined in SPC-5.

The PAGE CODE field and PAGE LENGTH field are defined in SPC-5 and shall be set to the values shown in table 255 for the Block Limits Extension VPD page.

The MAXIMUM NUMBER OF STREAMS field indicates the maximum number of streams that the device server supports and the maximum value for the stream identifier ~~plus one~~. A maximum number of streams field set to 0000h indicates that the device server does not support Stream Control.

The OPTIMAL STREAM WRITE SIZE field indicates the alignment and size of the optimal stream write as a number of logical blocks. The optimal stream write size is the same for all streams in the device server.

The STREAM GRANULARITY SIZE field indicates the stream granularity size in number of optimal stream write size blocks as described in 4.34.

### 6.6.6 Logical Block Provisioning VPD page

The Logical Block Provisioning VPD page (see table 256) provides the application client with logical block provisioning related operating parameters of the logical unit.

**Table 256 — Logical Block Provisioning VPD page**

Byte	Bit	7	6	5	4	3	2	1	0
0		PERIPHERAL QUALIFIER			PERIPHERAL DEVICE TYPE (00000b)				
1		PAGE CODE (B2h)							
2	(MSB)	PAGE LENGTH (0004h or (n - 3))							
3									
4		THRESHOLD EXPONENT							
5		LBPV	LBPWS	LBPWS10	LBPRZ			ANC_SUP	DP
6		MINIMUM PERCENTAGE					PROVISIONING TYPE		
7		THRESHOLD PERCENTAGE							
8		PROVISIONING GROUP DESCRIPTOR (if any)							
...									
n									

The PERIPHERAL QUALIFIER field is defined in SPC-5.

The PAGE CODE field and PERIPHERAL DEVICE TYPE field are defined in SPC-5 and shall be set to the value shown in table 256 for the Logical Block Provisioning VPD page.

The PAGE LENGTH field is defined in SPC-5. If the DP bit is set to zero, then the PAGE LENGTH field shall be set to 0004h. If the DP bit is set to one, then the PAGE LENGTH field shall be set to the value defined in table 256.

If the Logical Block Provisioning log page (see 6.4.5.1) is supported, then the logical unit shall support:

- a) logical block provisioning threshold sets; or
- b) logical block provisioning percentages.

The THRESHOLD EXPONENT field indicates the threshold set size as described in 4.7.3.7.2. A THRESHOLD EXPONENT field set to zero indicates that the logical unit does not support logical block provisioning threshold sets.

If logical block provisioning threshold sets are supported, then the threshold exponent shall be a non-zero value selected such that:

$$(\text{capacity} \div 2^{(\text{threshold exponent})}) < 2^{(32)}$$

where:

capacity is 1 + the LBA of the last logical block as returned in the READ CAPACITY (16) parameter data (see 5.18.2) (i.e., the number of logical blocks on the direct access block device);

threshold exponent is the contents of the THRESHOLD EXPONENT field; and

$2^{(32)}$  is the constant value 1\_0000\_0000h (i.e., 4 294 967 296).

A THRESHOLD PERCENTAGE field set to zero indicates that the logical unit does not support logical block provisioning percentages. If logical block provisioning percentages are supported, then the threshold

percentage shall be set to a non-zero value selected from the values in table 257. The units for the threshold percentage is tenths of a percent. This percentage represents the range over which logical block provisioning threshold percentages operates as described in 4.7.3.7.3.

**Table 257 — THRESHOLD PERCENTAGE field**

Code	Description
0	The logical unit does not support logical block provisioning percentages
1 to 255 (i.e., 01h to FFh)	0.1% to 25.5% of the total allocation resources

A MINIMUM PERCENTAGE field set to zero indicates that the logical unit does report a minimum percentage of resources required by the device. A MINIMUM PERCENTAGE field set to a non-zero indicates the minimum percentage of resources required by the device server as described in table 258. This value indicates the point where a device server may begin device initiated advanced background operations.

**Table 258 — MINIMUM PERCENTAGE field**

Code	Description
0	The logical unit does not report a minimum percentage of resources required
1 to 30 (i.e., 01h to 1Eh)	1% to 30% of the total allocation resources
All others	Reserved

A logical block provisioning UNMAP command (LBPU) bit set to one indicates that the device server supports the UNMAP command (see 5.30). An LBPU bit set to zero indicates that the device server does not support the UNMAP command.

A logical block provisioning WRITE SAME (16) command (LBPWS) bit set to one indicates that the device server supports the use of the WRITE SAME (16) command (see 5.48) to unmap LBAs. An LBPWS bit set to zero indicates that the device server does not support the use of the WRITE SAME (16) command to unmap LBAs.

A logical block provisioning WRITE SAME (10) command (LBPWS10) bit set to one indicates that the device server supports the use of the WRITE SAME (10) command (see 5.47) to unmap LBAs. An LBPWS10 bit set to zero indicates that the device server does not support the use of the WRITE SAME (10) command to unmap LBAs.

The logical block provisioning read zeros (LBPRZ) field is described in table 259. See table 10 for the definition of the logical block data returned by a read operation from an unmapped LBA for the different values of the LBPRZ field.

**Table 259 — LBPRZ field**

Code	Description
000b	The logical block data represented by unmapped LBAs (see 4.7.4.4) is vendor specific
xx1b	The logical block data represented by unmapped LBAs is set to zeros
010b	The logical block data represented by unmapped LBAs is set to the provisioning initialization pattern
all others	Reserved

An anchor supported (ANC\_SUP) bit set to one indicates that the device server supports anchored LBAs (see 4.7.1). An ANC\_SUP bit set to zero indicates that the device server does not support anchored LBAs.

A descriptor present (DP) bit set to one indicates a PROVISIONING GROUP DESCRIPTOR field is present. A DP bit set to zero indicates a PROVISIONING GROUP DESCRIPTOR field is not present.

The PROVISIONING TYPE field is defined in table 260.

**Table 260 — PROVISIONING TYPE field**

Code	Description
000b	The device server does not report a provisioning type or may be fully provisioned.
001b	The logical unit is resource provisioned (see 4.7.3.2).
010b	The logical unit is thin provisioned (see 4.7.3.3).
All others	Reserved

The PROVISIONING GROUP DESCRIPTOR field, if any, contains a designation descriptor (see SPC-5) for the LBA mapping resources used by this logical unit.

If a PROVISIONING GROUP DESCRIPTOR field is present:

- a) the ASSOCIATION field shall be set to 00b (i.e. logical unit); and
- b) the DESIGNATOR TYPE field shall be set to:
  - A) 1h (i.e., T10 vendor ID based); ~~or~~
  - B) 3h (i.e., NAA); or
  - C) Ah (i.e., UUID identifier).

### 6.6.7 Referrals VPD page

The Referrals VPD page (see table 261) contains parameters indicating characteristics of the user data segments contained within this logical unit.

**Table 261 — Referrals VPD page**

Byte	Bit	7	6	5	4	3	2	1	0
0		PERIPHERAL QUALIFIER			PERIPHERAL DEVICE TYPE (00000b)				
1		PAGE CODE (B3h)							
2	(MSB)	PAGE LENGTH (000Ch)							
3									
4		Reserved							
...									
7									
8	(MSB)	USER DATA SEGMENT SIZE							
...									
11		(LSB)							
12	(MSB)	USER DATA SEGMENT MULTIPLIER							
...									
15		(LSB)							

The PERIPHERAL QUALIFIER field is defined in SPC-5.

The PAGE CODE field, PERIPHERAL DEVICE TYPE field, and PAGE LENGTH field are defined in SPC-5 and shall be set to the values shown in table 261 for the Referrals VPD page.

A USER DATA SEGMENT SIZE field set to a non-zero value indicates the number of contiguous logical blocks in a user data segment (see 4.28.2). A USER DATA SEGMENT SIZE field set to zero indicates the user data segment size information (i.e., the first user data segment LBA to the last user data segment LBA) is as indicated in the user data segment referral descriptor (see table 18).

The USER DATA SEGMENT MULTIPLIER field is used by an application client to calculate the beginning LBA of each user data segment as described in 4.28.2.

## 6.6.8 Third-Party Copy VPD page

### 6.6.8.1 Third-Party Copy VPD page overview

The Third-Party Copy VPD page (see SPC-5) provides a means to retrieve third-party copy descriptors including a descriptor that describes operating parameters for the POPULATE TOKEN command (see 5.9) and the WRITE USING TOKEN command (see 5.52).

### 6.6.8.2 Block device third-party copy descriptor type codes

Block device third-party copy descriptor type codes (see table 262) indicate which third-party copy descriptor is being returned.

**Table 262 — Block device third-party copy descriptor type codes**

Descriptor code	Third-party copy descriptor name	Reference	Support requirements
0000h	Block Device ROD Token Limits	6.6.8.3	See <sup>a</sup>
All other codes	See <del>SPC-4</del> SPC-5	See SPC-5	See SPC-5
<sup>a</sup> Mandatory if the POPULATE TOKEN command and the WRITE USING TOKEN command are supported.			

### 6.6.8.3 Block Device ROD Token Limits descriptor

The Block Device ROD Token Limits descriptor (see table 263) is a third-party copy descriptor in the Third-Party Copy VPD page (see SPC-5) that provides the application client with a method to obtain operating parameters for direct access block device ROD token operations (see 4.30).

**Table 263 — Block Device ROD Token Limits descriptor**

Byte	Bit	7	6	5	4	3	2	1	0
0	(MSB)	THIRD-PARTY COPY DESCRIPTOR TYPE (0000h)							
1									
2	(MSB)	THIRD-PARTY COPY DESCRIPTOR LENGTH (0020h)							
3									
4		Vendor specific							
...									
9									
10	(MSB)	MAXIMUM RANGE DESCRIPTORS							
11									
12	(MSB)	MAXIMUM INACTIVITY TIMEOUT							
...									
15		(LSB)							
16	(MSB)	DEFAULT INACTIVITY TIMEOUT							
...									
19		(LSB)							
20	(MSB)	MAXIMUM TOKEN TRANSFER SIZE							
...									
27		(LSB)							
28	(MSB)	OPTIMAL TRANSFER COUNT							
...									
35		(LSB)							

The THIRD-PARTY COPY DESCRIPTOR TYPE field and the THIRD-PARTY COPY DESCRIPTOR LENGTH field are defined in ~~SPC-4~~SPC-5 and shall be set to the values shown in table 263 for the Block Device ROD Token Limits descriptor.

The MAXIMUM RANGE DESCRIPTORS field indicates the maximum number of block device range descriptors that may be specified in the parameter data of a POPULATE TOKEN command (see 5.9) and the parameter data of a WRITE USING TOKEN command (see 5.52). If the MAXIMUM RANGE DESCRIPTORS field is set to zero, then the copy manager does not report a maximum number of block device range descriptors.

The MAXIMUM INACTIVITY TIMEOUT field indicates the maximum value in the INACTIVITY TIMEOUT field of the parameter data of a POPULATE TOKEN command that is accepted by the copy manager. If the MAXIMUM INACTIVITY TIMEOUT field is set to zero, then the device server does not report a maximum inactivity timeout value. If the MAXIMUM INACTIVITY TIMEOUT field is set to FFFF\_FFFFh then there is no maximum value that may be specified in the INACTIVITY TIMEOUT field in the parameter data of the POPULATE TOKEN command.

The DEFAULT INACTIVITY TIMEOUT field indicates the inactivity timeout value that is used if the INACTIVITY TIMEOUT field in the parameter data of a POPULATE TOKEN command is set to zero. If the DEFAULT INACTIVITY TIMEOUT field is set to zero, then the copy manager does not report a default inactivity timeout value.

The MAXIMUM TOKEN TRANSFER SIZE field indicates the maximum size in logical blocks that may be specified by the sum of the NUMBER OF LOGICAL BLOCKS fields in all block device range descriptors for the following commands:

- a) POPULATE TOKEN; and
- b) WRITE USING TOKEN.

If the MAXIMUM TOKEN TRANSFER SIZE field is set to zero, then the copy manager does not report a maximum token transfer size.

If the MAXIMUM BYTES IN BLOCK ROD field in the block ROD device type specific features descriptor in the ROD token features third-party copy descriptor in the Third-party Copy VPD page (see SPC-5) is reported, then the MAXIMUM TOKEN TRANSFER SIZE field shall be set to the number of logical blocks that represents the value in the MAXIMUM BYTES IN BLOCK ROD field in the block ROD device type specific features descriptor in the ROD token features third-party copy descriptor in the Third-party Copy VPD page.

The OPTIMAL TRANSFER COUNT field indicates the optimal number of logical blocks that the copy manager is able to transfer. If the sum of the NUMBER OF LOGICAL BLOCKS fields in all block device range descriptors in the parameter data of a POPULATE TOKEN command or the parameter data of a WRITE USING TOKEN command exceeds this value then, a delay in processing the request may be incurred. If the field is set to zero, then the copy manager does not report an optimal transfer count.

If the OPTIMAL BYTES IN BLOCK ROD TRANSFER field in the block ROD device type specific features descriptor in the ROD token features third-party copy descriptor in the Third-party Copy VPD page is reported, then the OPTIMAL TRANSFER COUNT field shall be set to the number of logical blocks that represents the value in the OPTIMAL BYTES IN BLOCK ROD TRANSFER field in the block ROD device type specific features descriptor in the ROD token features third-party copy descriptor in the Third-party Copy VPD page.

### 6.6.9 Supported Block Lengths And Protection Types VPD page

The Supported Block Lengths And Protection Types VPD page (see table 264) contains parameters indicating the specific protection types (see 4.22) supported for each supported logical block length. If the PROTECT bit is set to zero in the Standard INQUIRY data (See SPC-5) or the SBLP bit is set to zero in the Control mode page (See SPC-5), then the device server shall not support this VPD page.

**Table 264 — Supported Block Lengths And Protection Types VPD page**

Bit Byte	7	6	5	4	3	2	1	0
0	PERIPHERAL QUALIFIER (000b)			PERIPHERAL DEVICE TYPE (00h)				
1	PAGE CODE (B4h)							
2	(MSB)							
3	PAGE LENGTH (n-3) (LSB)							
	Logical block length and protection types descriptor list							
4	Logical block length and protection types descriptor [first]							
...								
11								
	⋮							
n-7	Logical block length and protection types descriptor [last]							
...								
n								



The PERIPHERAL QUALIFIER field and PERIPHERAL DEVICE TYPE field are as defined in SPC-5 and shall be set as defined in table 264.

The PAGE CODE field is defined in SPC-5, and shall be set to the value shown in table 264 for the Supported Logical Block Lengths and Protection Types VPD page.

The PAGE LENGTH field is defined in SPC-5.

The logical block length and protection types descriptor list shall contain one logical block length and protection types descriptor for each logical block length that the device server supports.

Each logical block length and protection types descriptor describes the protection types supported for the logical block length specified (see table 265).

**Table 265 — Logical block length and protection types descriptor format**

Bit Byte	7	6	5	4	3	2	1	0
0	(MSB)							
...	LOGICAL BLOCK LENGTH							
3	(LSB)							
4	Reserved	P_I_I_SUP	Reserved	NO_PI_CHK	GRD_CHK	APP_CHK	REF_CHK	
5	Reserved				T3PS	T2PS	T1PS	T0PS
6	Reserved							
7	Reserved							

The LOGICAL BLOCK LENGTH field indicates the logical block length in bytes that is supported by the device server for which the device server supports the protection types identified in the P\_I\_I\_SUP bit, T3PS bit, T2PS bit, T1PS bit, T0PS bit, GRD\_CHK bit, APP\_CHK bit, and REF\_CHK bit.

A protection information interval supported (P\_I\_I\_SUP) bit set to one indicates that the logical unit supports protection information intervals for the specified logical block length. A P\_I\_I\_SUP bit set to zero indicates that the logical unit does not support protection information intervals for the specified logical block length.

A no protection information checking (NO\_PI\_CHK) bit set to one indicates that the device server disables checking of all protection information for the associated protection information interval when performing a write operation if:

- the LOGICAL BLOCK APPLICATION TAG field is set to FFFFh and type 1 protection (see SBC-4) is enabled;
- the LOGICAL BLOCK APPLICATION TAG field is set to FFFFh and type 2 protection (see SBC-4) is enabled; or
- the LOGICAL BLOCK APPLICATION TAG field is set to FFFFh, the LOGICAL BLOCK REFERENCE TAG field is set to FFFF FFFFh, and type 3 protection (see SBC-4) is enabled.

A NO\_PI\_CHK bit set to zero indicates that the device server checks protection information as specified by the WRPROTECT field (see SBC-4) when performing a write operation.

A guard check (GRD\_CHK) bit set to zero indicates that the device server does not check the LOGICAL BLOCK GUARD field in the protection information, if any for the specified logical block length. A GRD\_CHK bit set to one indicates that the device server checks the LOGICAL BLOCK GUARD field in the protection information, if any, for the specified logical block length.

An application tag check (APP\_CHK) bit set to zero indicates that the device server does not check the LOGICAL BLOCK APPLICATION TAG field in the protection information, if any, for the specified logical block length. An APP\_CHK bit set to one indicates that the device server checks the LOGICAL BLOCK APPLICATION TAG field in the protection information, if any, for the specified logical block length.

A reference tag check (REF\_CHK) bit set to zero indicates that the device server does not check the LOGICAL BLOCK REFERENCE TAG field in the protection information, if any, for the specified logical block length. A REF\_CHK bit set to one indicates that the device server checks the LOGICAL BLOCK REFERENCE TAG field in the protection information, if any, for the specified logical block length.

A type 3 protection supported (T3PS) bit set to one indicates that type 3 protection is supported for the indicated logical block length. A T3PS bit set to zero indicates that type 3 protection is not supported for the indicated logical block length.

A type 2 protection supported (T2PS) bit set to one indicates that type 2 protection is supported for the indicated logical block length. A T2PS bit set to zero indicates that type 2 protection is not supported for the indicated logical block length.

A type 1 protection supported (T1PS) bit set to one indicates that type 1 protection is supported for the indicated logical block length. A T1PS bit set to zero indicates that type 1 protection is not supported for the indicated logical block length.

A type 0 protection supported (T0PS) bit set to one indicates that type 0 protection is supported for the indicated logical block length. A T0PS bit set to zero indicates that type 0 protection is not supported for the indicated logical block length.

## 6.7 Copy manager parameters

The copy manager parameters are device type specific data for ROD tokens (see SPC-5) created by a copy manager for a direct access block device (see 4.30) and are defined in table 266.

**Table 266 — ROD token device type specific data**

Bit	7	6	5	4	3	2	1	0
Byte								
0	(MSB)							
...	LOGICAL BLOCK LENGTH IN BYTES							
3	(LSB)							
4	Reserved				P_TYPE		PROT_EN	
5	P_I_EXPONENT				LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT			
6	LBPME	LBPRZ	(MSB)					
7	LOWEST ALIGNED LOGICAL BLOCK ADDRESS						(LSB)	
8	Reserved							
...								
31								

The LOGICAL BLOCK LENGTH IN BYTES field, the P\_TYPE field, the PROT\_EN bit, the P\_I\_EXPONENT field, the LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT field, the LBPME bit, the LBPRZ bit, and the LOWEST ALIGNED LOGICAL BLOCK ADDRESS field are defined in 5.18.2.

## 6.8 Logical Block Markup descriptors

### 6.8.1 Summary of LBM descriptors

An LBM descriptor communicates information about anticipated usage of a logical block or range of logical blocks to the device server that manages access to that logical block. The requirements, if any, placed on a device server that receives an LBM descriptor are associated with the command that causes the device server to receive one or more LBM descriptors. LBM descriptor formats and the meanings of LBM descriptor fields are described in 6.8.

### 6.8.2 LBM descriptor formats and types

The format of an LBM descriptor is shown in table 267.

**Table 267 — Parameters for direct access block devices**

Bit Byte	7	6	5	4	3	2	1	0
0	LBM descriptor type specific data				LBM DESCRIPTOR TYPE			
1		LBM descriptor type specific data						
...								
n								

The contents of the LBM DESCRIPTOR TYPE field are defined in table 268.

**Table 268 — LBM DESCRIPTOR TYPE field**

Code	LBM descriptor type	Reference
0h	Access patterns	6.8.3
all others	Reserved	

### 6.8.3 Access patterns LBM descriptors

#### 6.8.3.1 Access patterns LBM descriptor format

The access patterns LBM descriptor format is shown in table 269.

**Table 269 — Access patterns LBM descriptor format**

Bit Byte	7	6	5	4	3	2	1	0
0	ACDLU	Reserved			LBM DESCRIPTOR TYPE (0h)			
1	OVERALL FREQUENCY		READ/WRITE FREQUENCY		WRITE SEQUENTIALITY		READ SEQUENTIALITY	
2	ReservedIO CLASS				SUBSEQUENT I/O		OSI PROXIMITY	
3	Reserved							

The LBM DESCRIPTOR TYPE field is defined in 6.8.2 and shall be set as shown in table 269 for the access patterns LBM descriptor format.

### 6.8.3.2 Accesses continue during low utilization (ACDLU) bit

If the ACDLU bit is set to one, then any logical blocks associated with this logical block markup have a high probability of being accessed during those time intervals, if any, in which most logical blocks are not being accessed by any read commands or write commands. If the ACDLU bit is set to zero, then any logical blocks associated with this logical block markup have no utilization related probability of being read or written.

### 6.8.3.3 WRITE SEQUENTIALITY field and READ SEQUENTIALITY field

Using the values shown in table 270:

- a) the WRITE SEQUENTIALITY field specifies whether sequential write operations or random write operations are more probable to a logical block in the range associated with this LBM descriptor; and
- b) the READ SEQUENTIALITY field specifies whether sequential read operations or random read operations are more probable to a logical block in the range associated with this LBM descriptor.

**Table 270 — WRITE SEQUENTIALITY field and READ SEQUENTIALITY field**

Code	Description
00b	Access sequentiality is unknown or equally probable.
01b	Random operations are more probable.
10b	Sequential operations are more probable.
11b	Reserved

### 6.8.3.4 READ/WRITE FREQUENCY field

The READ/WRITE FREQUENCY field (see table 271) specifies whether read operations or write operations are more probable to a logical block in the range associated with this LBM descriptor.

**Table 271 — READ/WRITE FREQUENCY field**

Code	Description
00b	Read operation frequency versus write operation frequency is unknown or equally probable.
01b	Read operations are more probable.
10b	Write operations are more probable.
11b	Reserved

### 6.8.3.5 OVERALL FREQUENCY field

The OVERALL FREQUENCY field (see table 272) specifies the probability that a logical block in the range associated with this LBM descriptor is likely to be accessed more frequently than other logical blocks stored on the same medium.

**Table 272 — OVERALL FREQUENCY field**

Code	Description
00b	Overall access frequency is unknown or equally probable.
01b	Overall accesses are less frequent than average.
10b	Overall accesses are more frequent than average.
11b	Reserved

**6.8.3.6 SUBSEQUENT I/O field**

The SUBSEQUENT I/O field (see table 273) specifies the probability that the application client may be delaying the sending of additional read commands or write commands to the device server until the completion of a read command or a write command for a logical block associated with this LBM descriptor (e.g., commands are being delayed by the application client's file system until the completion of write commands to logical blocks in a file system allocation bit map that are associated with this LBM descriptor).

**Table 273 — SUBSEQUENT I/O field**

Code	Description
00b	The probability is unknown whether the application client is delaying the sending of commands to the device server until completion of a read command or a write command for a logical block associated with this LBM descriptor.
01b	The probability is low that the application client is delaying the sending of commands to the device server until completion of a read command or a write command for a logical block associated with this LBM descriptor.
10b	The probability is high that the application client is delaying the sending of commands to the device server until completion of a read command or a write command for a logical block associated with this LBM descriptor.
11b	Reserved

**6.8.3.7 OSI PROXIMITY field**

The OSI PROXIMITY field (see table 274) specifies the probability that any logical block associated with this LBM descriptor is likely to be accessed during an operating system or file system initialization operation (e.g., a boot block).

**Table 274 — OSI PROXIMITY field**

Code	Description
00b	Access during initialization probability is unknown or equally probable.
01b	Accesses are not probable during an initialization process.
10b	Accesses are probable during an initialization process.
11b	Reserved

Methods to determine whether an operating system or file system initialization operation is in progress are outside the scope of this standard.

EXAMPLE - Possible sources of knowledge about when an operating system initialization operation is occurring depend on the system in which the device server is participating and include prior knowledge of which LBA accesses are associated with an operating system initialization or detection of a power on event.

**6.8.3.8** [IO CLASS field](#)

**6.8.3.9** [The IO CLASS field \(see table 275\) specifies the type that is associated with this LBM descriptor.](#)

**Table 275 — [IO CLASS field](#)**

<a href="#">Code</a>	<a href="#">Description</a>
<a href="#">0h</a>	<a href="#">None specified</a>
<a href="#">1h</a>	<a href="#">IOs that specify this LBM descriptor are related to user data that describes other user data (e.g. file system meta-data).<sup>a</sup></a>
<a href="#">4h</a>	<a href="#">IOs that specify this LBM descriptor are related to a small collection of user data where small is defined by the application client.<sup>a</sup></a>
<a href="#">5h</a>	<a href="#">IOs that specify this LBM descriptor are related to a large collection of user data where large is defined by the application client.<sup>a</sup></a>
<a href="#">all others</a>	<a href="#">Reserved</a>
<sup>a</sup> <a href="#">See Differentiated Storage Services</a>	

**6.8.3.10 Access patterns LBM descriptor usage considerations**

Device servers may ignore all or part of the information contained in an access patterns LBM descriptor.

EXAMPLE 1 - A device server that processes only the OVERALL FREQUENCY field.

EXAMPLE 2 - A device server that processes all LBM descriptor fields except the WRITE SEQUENTIALITY field and the READ SEQUENTIALITY field.

EXAMPLE 3 - In a product with a focus on sequential writes (e.g., a disk drive based on shingled magnetic recording), the device server ignores all values in all fields except a WRITE SEQUENTIALITY field ~~this~~[that](#) is set to 10b.

Although Annex H shows some possible combinations of the fields in an access patterns LBM descriptor, this standard places no requirements on how the device server interprets the interactions, if any, between the fields in an access patterns LBM descriptor.

## Annex A

(informative)

### Numeric order codes

#### A.1 Variable length CDBs

Commands that use operation code 7Fh in table 34 use the variable length command format defined in SPC-5 and are differentiated by service action codes as described in table A.1.

**Table A.1 — Variable length command service action code assignments**

Operation code/service action code	Description
7Fh/0000h	Reserved
7Fh/0001h	Reserved
7Fh/0002h	Reserved
7Fh/0003h	Obsolete
7Fh/0004h	Obsolete
7Fh/0005h	Reserved
7Fh/0006h	XPWRITE (32)
7Fh/0007h	XDWRITEREAD (32)
7Fh/0008h	Reserved
7Fh/0009h	READ (32)
7Fh/000Ah	VERIFY (32)
7Fh/000Bh	WRITE (32)
7Fh/000Ch	WRITE AND VERIFY (32)
7Fh/000Dh	WRITE SAME (32)
7Fh/000Eh	ORWRITE (32)
7Fh/000Fh	WRITE ATOMIC(32)
7FH/0010h	WRITE STREAM (32)
7Fh/0011h to 07FFh	Reserved
7Fh/0800h to FFFFh	See SPC-5

## A.2 Service action CDBs

Commands that use operation code 9Eh (i.e., SERVICE ACTION IN (16)) (see SPC-5) in table 34 are differentiated by service action codes as described in table A.2.

**Table A.2 — SERVICE ACTION IN (16) service actions**

Operation code/service action code	Description
9Eh/00h to 0Fh	Reserved for commands applicable to all device types (see SPC-5)
9Eh/10h	READ CAPACITY (16)
9Eh/11h	Obsolete
9Eh/12h	GET LBA STATUS
9Eh/13h	REPORT REFERRALS
9Eh/14h	STREAM CONTROL
9Eh/15h	BACKGROUND CONTROL
9Eh/16h	GET STREAM STATUS
9Eh/17h to 1Fh	Reserved

Commands that use operation code 9Fh (i.e., SERVICE ACTION OUT (16)) (see SPC-5) in table 34 are differentiated by service action codes as described in table A.3.

**Table A.3 — SERVICE ACTION OUT (16) service actions**

Operation code/service action code	Description
9Fh/00h to 0Fh	Reserved for commands applicable to all device types (see SPC-5)
9Fh/10h	Reserved
9Fh/11h	WRITE LONG (16)
9Fh/12h to 1Fh	Reserved



## **Annex B**

(informative)

### **XOR command examples**

#### **B.1 XOR command examples overview**

This annex provides XOR command examples in SCSI storage array device (see SCC-2) supervised configurations.

#### **B.2 Update write operation**

Figure B.1 shows an update write operation (see 4.19.2.2) operation supervised by a SCSI storage array device (see SCC-2). The example uses a supervising SCSI storage array device, a direct access block device holding XOR-protected data (i.e., the data disk), and a direct access block device holding check data (i.e., the parity disk) that use the following SCSI commands:

- a) an XDWRITE command (see SBC-2);
- b) an XDREAD command (see SBC-2); and
- c) an XPWRITE command.

An XDWRITEREAD command may be used in place of any sequence of an XDWRITE command followed by an XDREAD command.

The supervising SCSI storage array device begins by sending XOR-protected data to the data disk using an XDWRITE command.

The data disk reads old XOR-protected data, performs an XOR operation using the old XOR-protected data and the XOR-protected data from the supervising SCSI storage array device, stores the resulting intermediate XOR data in its buffer memory, and writes the XOR-protected data from the supervising SCSI storage array device. The supervising SCSI storage array device reads the resulting intermediate XOR data from the buffer memory by sending the data disk an XDREAD command.

The supervising SCSI storage array device makes the resulting intermediate XOR data (i.e., data read with the XDREAD command) available to the parity disk by sending an XPWRITE command. The parity disk performs an XOR operation using the intermediate XOR data and the old XOR data read from the parity disk. The resulting new XOR data is written.

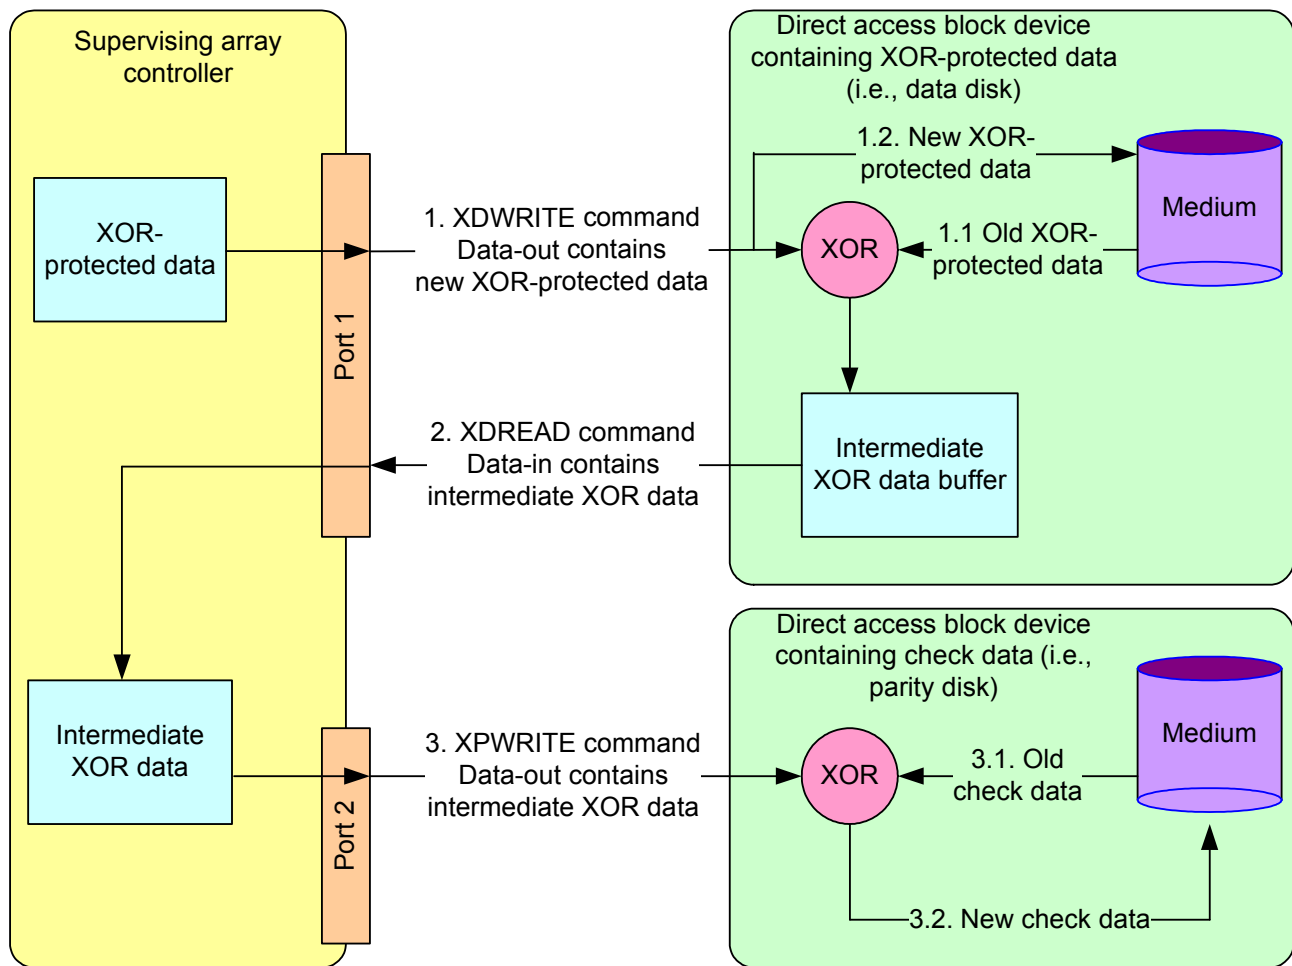


Figure B.1 — Update write operation (SCSI storage array device supervised)

### B.3 Regenerate operation

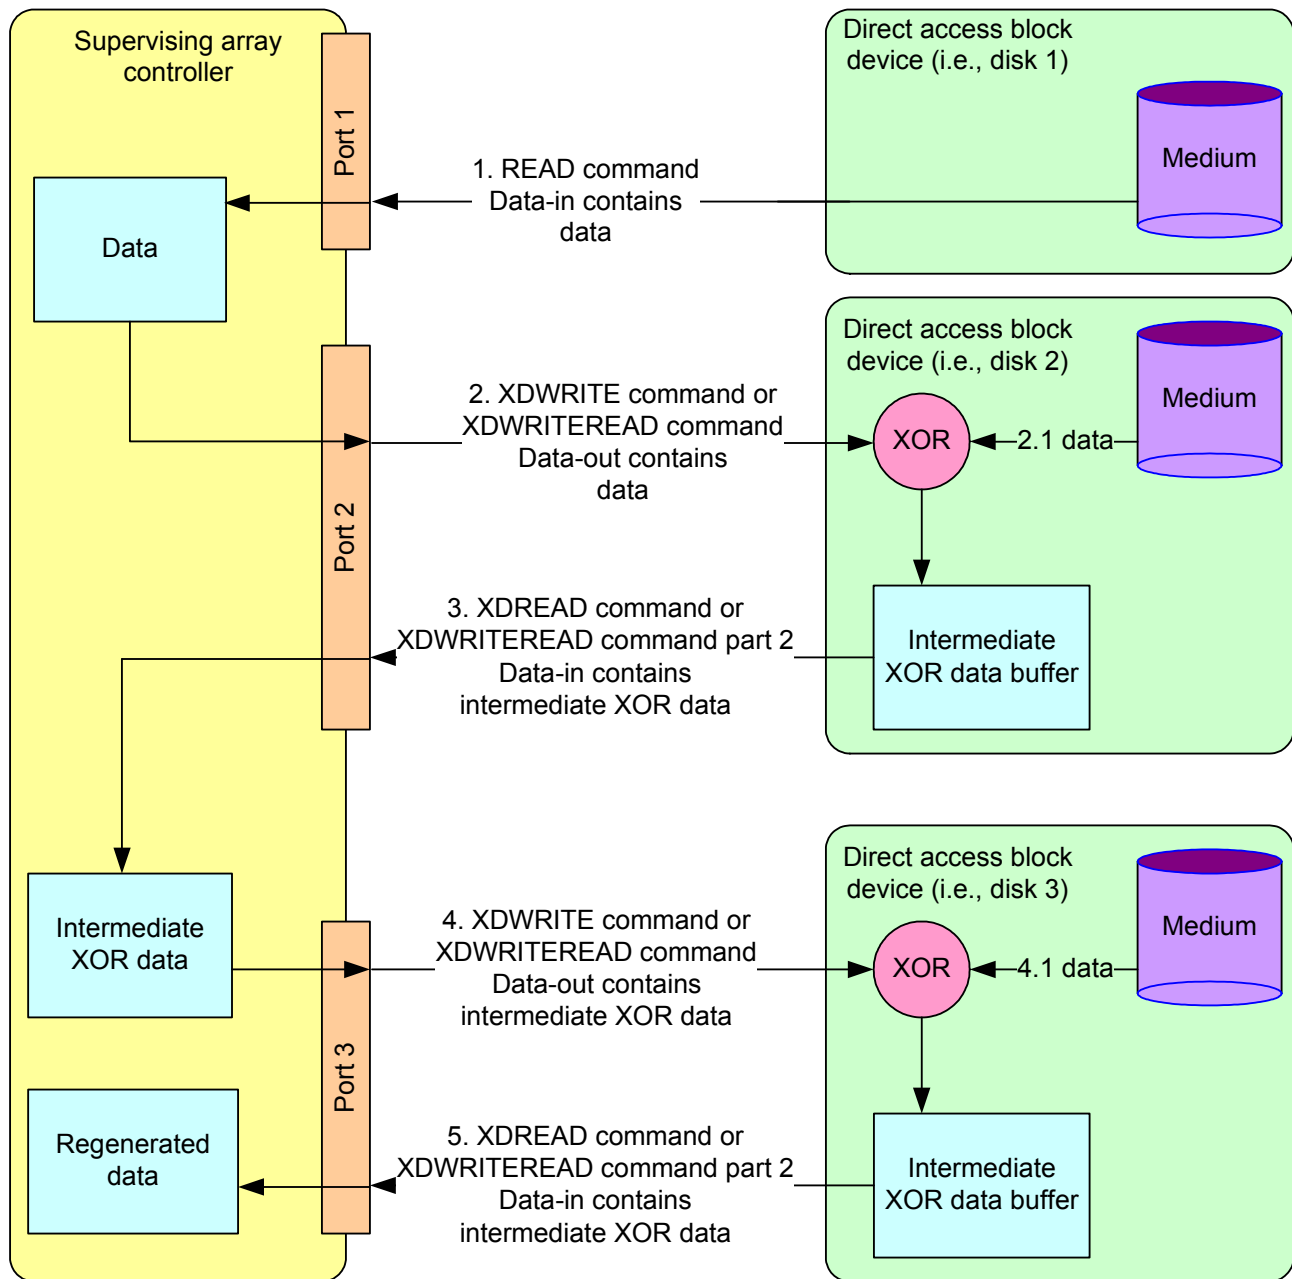
Figure B.2 shows a regenerate operation (see 4.19.2.3) supervised by a SCSI storage array device (see SCC-2). The example uses a supervising SCSI storage array device and three direct access block devices (i.e., disk 1, disk 2, and disk 3) that use the following SCSI commands:

- a READ command;
- an XDWRITE command (see SBC-2); and
- an XDREAD command (see SBC-2).

An XDWRITEREAD command may be used in place of any sequence of an XDWRITE command followed by an XDREAD command.

The supervising SCSI storage array device begins by issuing a READ command to disk 1. The data received from this command is sent by the supervising SCSI storage array device to disk 2 using an XDWRITE command with the DISABLE WRITE bit set to one. Disk 2 reads data, performs an XOR operation using that data and the data received from the supervising SCSI storage array device, and stores the resulting intermediate XOR data in its buffer memory. The supervising SCSI storage array device retrieves the intermediate XOR data from the buffer memory by issuing an XDREAD command to disk 2. The supervising SCSI storage array device issues an XDWRITE command and an XDREAD command in the same manner to disk 3.

The resulting data from disk 3 is the regenerated data.



**Figure B.2 — Regenerate operation (SCSI storage array device supervised)**

## B.4 Rebuild operation

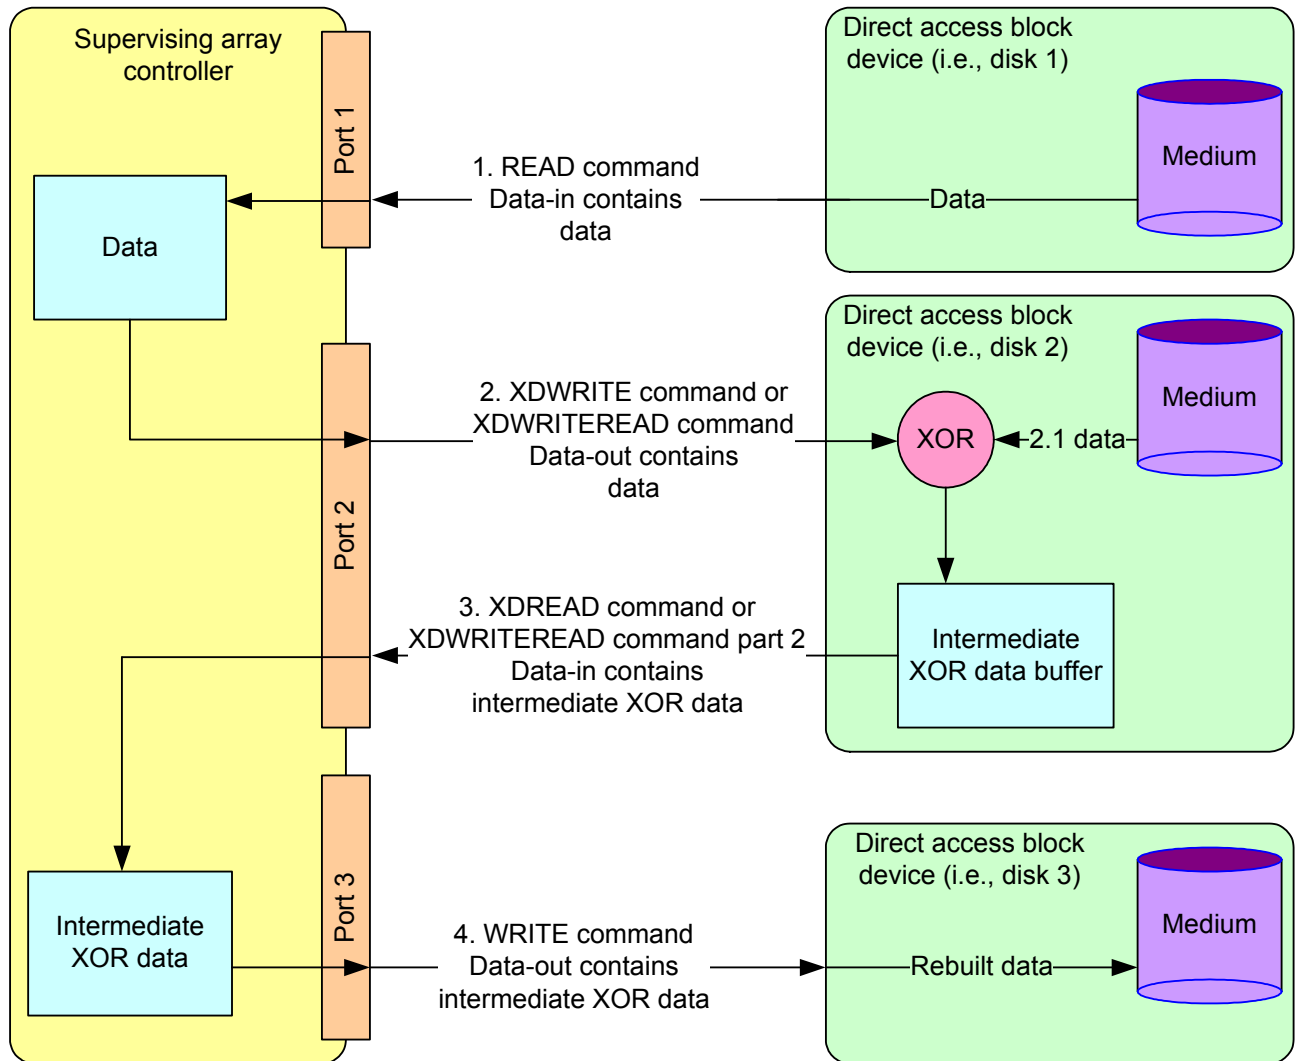
Figure B.3 shows a rebuild operation (see 4.19.2.4) supervised by a SCSI storage array device (see SCC-2). The example uses a supervising SCSI storage array device, two direct access block devices (i.e., disk 1 and disk 2) as the source devices, and one direct access block device (i.e., disk 3) as the rebuild device that use the following SCSI commands:

- a READ command;
- an XDWRITE command (see SBC-2);
- an XDREAD command (see SBC-2); and
- a WRITE command.

An XDWRITEREAD command may be used in place of any sequence of an XDWRITE command followed by an XDREAD command.

The supervising SCSI storage array device begins by issuing a READ command to disk 1. The data received from the READ command is sent by the supervising SCSI storage array device to disk 2 using an XDWRITE command with a DISABLE WRITE bit set to one. Disk 2 reads data, performs an XOR operation using that data and the data received from the supervising SCSI storage array device, and stores the resulting intermediate XOR data in its buffer memory. The supervising SCSI storage array device retrieves the intermediate XOR data by sending an XDREAD command to disk 2.

The resulting data from disk 2 is the rebuilt data and is sent to the direct access block device being rebuilt (i.e., disk 3) using a WRITE command.



**Figure B.3 — Rebuild operation (SCSI storage array device supervised)**

## Annex C

(informative)

### CRC example in C

The following is an example C program that generates the value for the LOGICAL BLOCK GUARD field in protection information (see 4.22).

```
// picrc.cpp : SCSI SBC-4 Protection Information CRC generator
#include "stdafx.h"
#include <stdio.h>
#include <malloc.h>

/* return crc value */
unsigned short calculate_crc(unsigned char *frame, unsigned long length) {
    unsigned short const poly = 0x8BB7L; /* Polynomial */
    unsigned const int poly_length = 16;
    unsigned short crc_gen;
    unsigned short x;
    unsigned int i, j, fb;
    unsigned const int invert = 0; /* 1=seed with 1s and invert the CRC */

    crc_gen = 0x0000;
    crc_gen ^= invert? 0xFFFF: 0x0000; /* seed generator */

    for (i = 0; i < length; i += 2) {
        /* assume little endian */
        x = (frame[i] << 8) | frame[i+1];

        /* serial shift register implementation */
        for (j = 0; j < poly_length; j++) {
            fb = ((x & 0x8000L) == 0x8000L) ^ ((crc_gen & 0x8000L) ==
0x8000L);
            x <<= 1;
            crc_gen <<= 1;
            if (fb)
                crc_gen ^= poly;
        }
    }

    return crc_gen ^ (invert? 0xFFFF: 0x0000); /* invert output */
} /* calculate_crc */

/* function prototype */
unsigned short calculate_crc(unsigned char *, unsigned long);

void main (void) {
    unsigned char *buffer;
    unsigned long buffer_size = 32;
    unsigned short crc;
    unsigned int i;

    /* 32 0x00 */
    buffer = (unsigned char *) malloc (buffer_size);
    for (i = 0; i < buffer_size; i++) {
        buffer[i] = 0x00;
    }
}
```

```

}
crc = calculate_crc(buffer, buffer_size);
printf ("Example CRC all-zeros is %04x\n", crc);
free (buffer);

/* 32 0xFF */
buffer = (unsigned char *) malloc (buffer_size);
for (i = 0; i < buffer_size; i++) {
    buffer[i] = 0xFF;
}
crc = calculate_crc(buffer, buffer_size);
printf ("Example CRC all-ones is %04x\n", crc);
free (buffer);

/* 0x00 incrementing to 0x1F */
buffer = (unsigned char *) malloc (buffer_size);
for (i = 0; i < buffer_size; i++) {
    buffer[i] = i;
}
crc = calculate_crc(buffer, buffer_size);
printf ("Example CRC incrementing is %04x\n", crc);
free (buffer);

/* 0xFF 0xFF then 30 zeros */
buffer = (unsigned char *) malloc (buffer_size);
buffer[0] = 0xff;
buffer[1] = 0xff;
for (i = 2; i < buffer_size; i++) {
    buffer[i] = 0x00;
}
crc = calculate_crc(buffer, buffer_size);
printf ("Example CRC FF FF then 30 zeros is %04x\n", crc);
free (buffer);

/* 0xFF decrementing to 0xE0 */
buffer = (unsigned char *) malloc (buffer_size);
for (i = 0; i < buffer_size; i++) {
    buffer[i] = 0xff - i;
}
crc = calculate_crc(buffer, buffer_size);
printf ("Example CRC FF decrementing to E0 is %04x\n", crc);
free (buffer);

} /* main */

```

## Annex D

(informative)

### Sense information for locked or encrypted logical units

A device server may complete some commands with CHECK CONDITION status under certain conditions while the logical unit is locked or encrypted. Table D.1 describes the conditions relative to the sense key and the additional sense code returned by the device server with the CHECK CONDITION status.

**Table D.1 — Sense information for locked or encrypted logical units**

Sense key	Additional sense code	Description
DATA PROTECT	ACCESS DENIED – NO ACCESS RIGHTS	The logical unit is locked. This condition may occur for read commands or write commands. This condition may occur for the entire logical unit or for a range of LBAs contained in the logical unit. To clear this condition, an application client performs a security protocol specific procedure to unlock access to the logical unit.
ABORTED COMMAND	LOGICAL BLOCK REFERENCE TAG CHECK FAILED	These conditions may occur for a read command. The additional sense codes may indicate that an encrypting logical unit has changed the encryption/decryption key, and the LBAs requested by the command have not yet been rewritten. Disabling protection information checking in a CDB may allow the command to complete successfully, but the data returned for the command may be invalid (i.e., not decrypted). To clear this condition, an application client writes the LBAs for which the condition occurred with new data.
ABORTED COMMAND	LOGICAL BLOCK APPLICATION TAG CHECK FAILED	
ABORTED COMMAND	LOGICAL BLOCK GUARD CHECK FAILED	

## **Annex E**

(informative)

### **Optimizing block access characteristics**

#### **E.1 Optimizing block access overview**

This annex describes example methods that application clients may use to achieve optimal performance for logical block access. These examples use the following information:

- a) the LOWEST ALIGNED LOGICAL BLOCK ADDRESS field (see 5.18.2);
- b) the LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT field (see 5.18.2);
- c) the OPTIMAL TRANSFER LENGTH GRANULARITY field (see 6.6.4);
- d) the OPTIMAL TRANSFER LENGTH field (see 6.6.4);
- e) the MAXIMUM TRANSFER LENGTH field (see 6.6.4);
- f) the OPTIMAL STREAM WRITE SIZE field (see 6.6.5); and
- g) the STREAM GRANULARITY SIZE field (see 6.6.5).

#### **E.2 Starting logical block offset**

The READ CAPACITY (16) command transfers parameter data which includes a value in the LOWEST ALIGNED LOGICAL BLOCK ADDRESS field. As shown in figure 4, the value in this field indicates the starting alignment of logical block addresses where optimal performance for logical block access begins.

#### **E.3 Optimal granularity sizes**

The READ CAPACITY (16) command transfers parameter data that includes a value in the LOGICAL BLOCKS PER PHYSICAL BLOCK EXPONENT field. As shown in figure 2 and in figure 4, the value in this field enables the application client to determine the number of logical blocks per physical block.

The Block Limits VPD page may include values in the OPTIMAL TRANSFER LENGTH GRANULARITY field, the OPTIMAL TRANSFER LENGTH field, and the MAXIMUM TRANSFER LENGTH field. These values may be used to determine optimum transfer sizes.

If the OPTIMAL TRANSFER LENGTH GRANULARITY field is valid (i.e., contains a value greater than zero), then the value in the OPTIMAL TRANSFER LENGTH GRANULARITY field is the optimal granularity size. If:

- a) the Block Limits VPD page is not supported; or
- b) the Block Limits VPD page is supported and the OPTIMAL TRANSFER LENGTH GRANULARITY field is set to zero,

then the value  $2^{(\text{logical blocks per physical block exponent})}$  is the optimal granularity size.

#### **E.4 Optimal stream granularity sizes**

The Block Limits Extension VPD page may include values in the OPTIMAL STREAM WRITE SIZE field and the STREAM GRANULARITY SIZE field. These values may be used to determine optimum transfer sizes and optimum transfer alignment to use with the WRITE STREAM commands (see 5.50 and 5.51).

If the OPTIMAL STREAM WRITE SIZE field is valid (i.e., contains a value greater than zero), then the value in the OPTIMAL STREAM WRITE SIZE field indicates the optimum transfer size and the optimum transfer alignment.



If:

- a) the Block Limits Extension VPD page is not supported; or
- b) the Block Limits Extension VPD page is supported and the OPTIMAL STREAM WRITE SIZE field is set to zero,

then the optimum transfer size and optimum transfer alignment is not reported for WRITE STREAM commands.

## E.5 Optimizing transfers

### E.5.1 Overview

Non-stream write commands may be optimized using the method described in E.5.2. WRITE STREAM commands may be optimized using the method described in E.5.3.

### E.5.2 Optimizing non-stream transfers

If:

- a) the Block Limits Extension VPD page is not supported;
- b) the OPTIMAL STREAM WRITE SIZE field (see 6.6.5)page is zero; or
- c) a write command other than a WRITE STREAM command is sent to the device server,

then to obtain optimal performance, the application client requests transfers with a starting LBA of the form calculated by the following formula:

$$\text{starting LBA} = \text{lowest aligned LBA} + (\text{optimal transfer length granularity} \times n)$$

where:

starting LBA      is the LBA of the first logical block accessed;  
 lowest aligned LBA      is the value in the LOWEST ALIGNED LOGICAL BLOCK ADDRESS field; and  
 n      is zero or a positive integer.

and using transfer lengths of the form:

$$\text{transfer length} = (\text{optimal granularity size} \times k)$$

where:

transfer length      is the number of contiguous logical blocks of data being accessed;  
 optimal granularity size      is the value described in E.3; and  
 k      is a positive integer.

To obtain optimal performance, the application client requests a transfer length, in logical blocks, that is no larger than the value in the MAXIMUM TRANSFER LENGTH field, and is:

- a) no larger than the optimal transfer length for logical units where the delay in processing transfers larger than the optimal transfer length is large; or
- b) not limited by the value in the OPTIMAL TRANSFER LENGTH field for logical units where the delay in processing transfers larger than the optimal transfer length is small (i.e., most direct access block devices exhibit this type of operation).

NOTE 26 - There is no method available to determine if the delay in processing for various transfer lengths is large or small.

It is more important that the application client meet the logical unit's starting and ending alignment boundary conditions than the maximum transfer length conditions. These considerations have larger impacts on write performance than read performance.

### E.5.3 Optimizing stream transfers

If the Block Limits Extension VPD page is supported, the OPTIMAL STREAM WRITE SIZE field (see 6.6.5) is non-zero, and WRITE STREAM commands are sent to the device server with a given stream identifier, then to obtain optimal performance, the application client requests transfers using transfer lengths of the form:

$$\text{transfer length} = (\text{optimal stream granularity} \times k)$$

where:

- transfer length is the number of contiguous logical blocks of data being accessed;
- optimal stream granularity is the optimal stream write size (see 6.6.5); and
- k is a positive integer,

and a stream length of the form:

$$\text{stream length} = (\text{optimal stream granularity} \times \text{stream granularity size} \times k)$$

where:

- stream length is the sum of the transfer length of all transfers for that stream (i.e., from when the stream is opened to when the stream is closed);
- optimal granularity size is the optimal stream write size;
- stream granularity size is the value in the STREAM GRANULARITY SIZE field; and
- k is a positive integer.

### E.6 Examples

In this first example, a logical unit reports the following information:

- a) the LOWEST ALIGNED LOGICAL BLOCK ADDRESS field set to 0003h in the READ CAPACITY (16) parameter data (see 5.18.2);
- b) the OPTIMAL TRANSFER LENGTH GRANULARITY field set to 0008h in the Block Limits VPD page (see 6.6.4);
- c) the MAXIMUM TRANSFER LENGTH field set to 0000\_0000h in the Block Limits VPD page; and
- d) the OPTIMAL TRANSFER LENGTH field set to 0000\_0080h (i.e., 128) in the Block Limits VPD page.

The starting LBA for optimal transfers on this logical unit should be of the form  $((8 \times n) + 3)$  where n is any integer greater than or equal to zero (e.g., starting LBAs of 3, 11, 19, 27, and 35). The transfer length for optimal transfers should be a multiple of eight logical blocks (e.g., transfer lengths of 8 blocks, 32 blocks, or 128 blocks).

A write command with the LOGICAL BLOCK ADDRESS field set to 19 and the TRANSFER LENGTH field set to 32 should exhibit improved performance over a write command with the LOGICAL BLOCK ADDRESS field set to 18 and the TRANSFER LENGTH field set to 32.

If the device has a delay in processing transfers larger than the optimal transfer length, some operations may exhibit improved performance if a single large request is broken into multiple smaller requests (e.g., rather than performing a single read of 248 logical blocks, the transfer may be optimized by setting the transfer length of one read command to 128 logical blocks and setting the transfer length of a second read command to 120 logical blocks).

In this second example, a logical unit reports the following information:

- a) the OPTIMAL STREAM WRITE SIZE field set to 0010h (i.e., 16) in the Block Limits Extension VPD page (see 6.6.5); and
- b) the STREAM GRANULARITY SIZE field set to 0040h (i.e., 64) in the Block Limits Extension VPD page.

The logical block address field in a WRITE STREAM command for optimal transfers should be of the form  $(16 \times j)$  where  $j$  is any integer greater than or equal to zero (e.g., logical block addresses of 16, 32, or 48). The transfer length in a WRITE STREAM command for optimal transfers should be of the form  $(16 \times k)$  where  $k$  is any integer greater than or equal to zero (e.g., transfer lengths of 16 blocks, 32 blocks, or 48 blocks). The stream length for optimal transfers should be of the form  $(16 \times 64 \times m)$  where  $m$  is any integer greater than or equal to zero (e.g., stream lengths of 1 024 blocks, 2 048 blocks, or 3 072 blocks).

The stream length is the total of all write transactions using the specified stream identifier (i.e., from the receipt of a STREAM CONTROL command with the STR\_CTL field set to 01b (i.e., open) to the receipt of a STREAM CONTROL command with the STR\_CTL field set to 10b (i.e., close) and the STR\_ID field set to the specified stream identifier).

A WRITE STREAM command with the LOGICAL BLOCK ADDRESS field set to 32 (i.e., an integer multiple of 16) and the TRANSFER LENGTH field set to 64 (i.e., an integer multiple of 16) should exhibit improved performance over a WRITE STREAM command with the LOGICAL BLOCK ADDRESS field set to 16 (i.e., an integer multiple of 16) and the TRANSFER LENGTH field set to 20 (i.e., not an integer multiple of 16).

A sequence of WRITE STREAM commands to the stream that specify a total of 2 048 logical blocks (i.e., an integer multiple of 1 024) should exhibit improved performance over a sequence of WRITE STREAM commands to the stream that specify a total of 1 536 logical blocks (i.e., not an integer multiple of 1 024).

## **Annex F**

(informative)

### **Logical block provisioning reporting examples**

#### **F.1 Logical block provisioning reporting examples overview**

Logical block provisioning reporting may be implemented using different methods. Implementations may include one or more of the following:

- a) use of dedicated LBA mapping resources (e.g., resources are associated with a specific logical unit);
- b) use of shared LBA mapping resources (e.g., resources are shared by multiple logical units);
- c) reporting based on dedicated LBA mapping resources (e.g., resources are reported specific to the logical unit);
- d) reporting based on shared LBA mapping resources (e.g., resources are reported for the resource pool as a whole);
- e) LBA mapping resource tracking based on logical blocks; and
- f) LBA mapping resource tracking based on threshold sets.

This annex describes examples of logical block provisioning reporting. Each example follows logical block provisioning resource usage and reporting over time as a specified set of operations occur.

#### **F.2 Interpreting log parameter counts**

Due to the variation of the threshold set size implementations, logical block usage and resource reporting may not have a direct relationship. The second example (see F.4) demonstrates an implementation where logical blocks are allocated on an individual LBA basis and reported using a larger threshold set basis. The reporting is a direct calculation from a logical block based count to a threshold set based count.

In implementations where a threshold set contains a set of contiguous logical blocks, the reporting may be substantially different. LUN 1 in the first example (see F.3) demonstrates such an implementation. At the initial conditions, two threshold sets are reported as being used. With a threshold set size of 1 024 blocks, these two threshold sets may contain as little as one logical block of application client data in each threshold set, or as many as 1 024 contiguous logical blocks in each threshold set. Which LBAs have been written by the application client has a substantial impact on how the usage of those resources is reported.

The relationship of the physical blocks to the logical blocks (see figure 4 and figure 5) may have an impact on the logical block provisioning log parameters. Which LBAs are written by the application may impact the number of physical blocks required to be allocated and therefore impact the reporting of the LBA mapping resource parameters.

The device server may not prioritize the maintenance of the values in the Logical Block Provisioning log page (see 6.4.5) above the completion of other operations (e.g., read operations or write operations). This may result in delays in updates to these values (e.g., after a request to unmap a large number of logical blocks). The logical block provisioning log parameters may also appear inaccurate for logical units where unmap operations cause LBA mapping resources to be released using a periodic background function.

Data de-duplication (see 4.8) and compression may impact the logical block provisioning log parameters. If the device server is able to perform data de-duplication or compression, then the number of LBA mapping resources:

- a) used for a successful write operation may not be the same as specified in the command that requested that operation; or
- b) made available after one or more successful unmap operations may not be the same as specified in the command that requested those operations.

EXAMPLE 1 - If a write command to LBAs that are all in the deallocated state (see 4.7.4.6) results in data that is all de-duplicated, then there may be no additional LBA mapping resources used when those LBAs transition to the mapped state (see 4.7.4.5). This may result in no change in the available LBA mapping resource count (see 6.4.5.2), or the used LBA mapping resource count (see 6.4.5.3).

EXAMPLE 2 -If a write command to LBAs that are all in the mapped state results in data that is no longer de-duplicated, then additional LBA mapping resources may be required even though all the LBAs addressed by the command remain in the mapped state. This may result in a reduction in the available LBA mapping resource count, or an increase in the used LBA mapping resource count. On a thin provisioned logical unit, such a write command may also be terminated as a result of a resource exhaustion condition (see 4.7.3.6.1).

EXAMPLE 3 -If a write command to LBAs that are all in the mapped state results in data that is more compressible than the previous data in those LBAs, then fewer LBA mapping resources may be required even though all the LBAs addressed by that command remain in the mapped state. This may result in an increase in the available LBA mapping resource count, or a decrease in the used LBA mapping resource count.

EXAMPLE 4 -If a write command to LBAs that are all in the mapped state results in data that is less compressible than the previous data in those LBAs, then additional LBA mapping resources may be required even though all the LBAs addressed by that command remain in the mapped state. This may result in a decrease in the available LBA mapping resource count, or an increase in the used LBA mapping resource count. On a thin provisioned logical unit, such a write command may also be terminated as a result of a resource exhaustion condition.

EXAMPLE 5 -If an unmap command addresses LBAs that are all in the mapped state and contain data that has been de-duplicated (i.e., additional logical blocks still contain that same de-duplicated data), then on a thin provisioned logical unit, there may be no change in the LBA mapping resources even though all of the LBAs addressed by that command transition to the deallocated state. This may result in no change to the available LBA mapping resource count or the used LBA mapping resource count.

As a result, application clients using logical block provisioning thresholds and examining logical block provisioning log parameters should not expect application client determined usage values or application client determined available space values to match log parameters or threshold events as reported by the logical unit.

### F.3 Dedicated resource, threshold set tracked example

#### F.3.1 Dedicated resource, threshold set tracked example overview

This example describes a method that reports dedicated logical block provisioning resources based on threshold sets. In this example, the values reported by the logical unit in the Logical Block Provisioning log page (see 6.4.5) reflect the usage for each logical unit and the available resources dedicated to each logical unit. Each threshold set is allocated to contain a set of contiguous logical blocks (e.g., LBAs 1024 to 2047 are contained in the same threshold set).

#### F.3.2 Dedicated resource, threshold set tracked example configuration

The configuration used for this example consists of two thin provisioned logical units, each with dedicated logical block provisioning resources. Table F.1 shows logical block provisioning related capacity values used in this example.

**Table F.1 — Dedicated resource, threshold set tracked example capacity information**

LUN	Capacity		THRESHOLD EXPONENT field <sup>c</sup>	Number of threshold sets <sup>d</sup>
	LBA <sup>a</sup>	Logical blocks <sup>b</sup>		
1	3FFF_FFFFh	1 Gi	0Ah (i.e., 512 KiB, 1 024 logical blocks)	0010_0000h (i.e., 1 Mi)
2	BFFF_FFFFh	3 Gi	0Ch (i.e., 2 MiB, 4 096 logical blocks)	000C_0000h (i.e., 768 Ki)
<sup>a</sup> RETURNED LOGICAL BLOCK ADDRESS field in READ CAPACITY parameter data (see 5.17.2 and 5.18.2). <sup>b</sup> The value returned in the RETURNED LOGICAL BLOCK ADDRESS field plus one. <sup>c</sup> In the Logical Block Provisioning VPD page (see 6.6.6). <sup>d</sup> Number of threshold sets = capacity ÷ 2 <sup>(threshold exponent)</sup> .				

Table F.2 shows LUN 1 with four enabled threshold descriptors and LUN 2 with two enabled threshold descriptors. The threshold descriptors in the Logical Block Provisioning mode page (see 6.5.9) for LUN 1 are configured to report a logical block provisioning threshold crossing (see 4.7.3.7) when:

- a) the percentage of available LBA mapping resources reaches 30% of reported capacity;
- b) the percentage of available LBA mapping resources reaches 20% of reported capacity;
- c) the percentage of available LBA mapping resources reaches 10% of reported capacity; or
- d) the percentage of used LBA mapping resources reaches 75% of reported capacity.

The threshold descriptors in the Logical Block Provisioning mode page for LUN 2 are configured to report a logical block provisioning threshold crossing when:

- a) the percentage of available LBA mapping resources reaches 50% of reported capacity; or
- b) the percentage of available LBA mapping resources reaches 10% of reported capacity.

**Table F.2 — Dedicated resource, threshold set tracked example capacity information**

LUN	Threshold resource <sup>a</sup>	Threshold count <sup>b</sup>	Description
1	0001h	0004_CCCCh	An available LBA mapping resource threshold set to 30% of the reported capacity (i.e., number of threshold sets from table F.1 $\times$ 0.30 = 0004_CCCCh threshold sets)
	0001h	0003_3333h	An available LBA mapping resource threshold set to 20% of the reported capacity
	0001h	0001_9999h	An available LBA mapping resource threshold set to 10% of the reported capacity
	0002h	000C_0000h	A used LBA mapping resource threshold set to 75% of the reported capacity (i.e., number of threshold sets from table F.1 (i.e., 0010_0000h) $\times$ 0.75 = 000C_0000h threshold sets)
2	0001h	0006_0000h	An available LBA mapping resource threshold set to 50% of the reported capacity (i.e., number of threshold sets from table F.1 (i.e., 000C_0000h) $\times$ 0.50 = 0006_0000h threshold sets)
	0001h	0001_3333h	An available LBA mapping resource threshold set to 10% of the reported capacity
<sup>a</sup> THRESHOLD RESOURCE field (see 6.5.9.2) and the PARAMETER CODE field (see 6.4.5.2). <sup>b</sup> THRESHOLD COUNT field (see 6.5.9.2).			

### F.3.3 Dedicated resource, threshold set tracked example sequence

The sequence of events for this example are:

- 1) initial conditions (see F.3.4);
- 2) operations that occur (see F.3.5); and
- 3) final values in the logical block provisioning log page (see F.3.6).

### F.3.4 Dedicated resource, threshold set tracked example initial conditions

Initially, LUN 1 has two threshold sets used and has 69 108 736 logical blocks available (i.e. 0001\_07A1h threshold sets). The application client has written at least one logical block into each of the two logical block ranges that correspond to those two threshold sets, therefore the application client may have written from two logical blocks to 2048 logical blocks. LUN 2 has 1 073 741 824 logical blocks available (i.e., 0004\_0000h threshold sets). LUN 2 does not report a used LBA mapping resource parameter. Table F.3 shows the values in the Logical Block Provisioning log page for the initial conditions in this example.

**Table F.3 — Dedicated resource, threshold set tracked example initial conditions**

LUN	Log page parameter <sup>a</sup>	Resource count <sup>b</sup>	Scope <sup>c</sup>	Description <sup>d</sup>
1	0001h	0001_07A1h	01b	The available LBA mapping resource parameter indicates that 69 108 736 logical blocks (i.e., 1_07A1h threshold sets × 1 024 logical blocks per threshold set) are available for LUN 1.
	0002h	0000_0002h	01b	The used LBA mapping resource parameter indicates that 2 048 logical blocks (i.e., 2h threshold sets × 1 024 logical blocks per threshold set) have been used (i.e., allocated) by LUN 1.
2	0001h	0004_0000h	01b	The available LBA mapping resource parameter indicates that 1 073 741 824 logical blocks (i.e., 4_0000h threshold sets × 4 096 logical blocks per threshold set) are available for LUN 2.
<sup>a</sup> THRESHOLD RESOURCE field (see 6.5.9.2) and the PARAMETER CODE field (see 6.4.5.2). <sup>b</sup> RESOURCE COUNT field (see 6.5.9.2). <sup>c</sup> SCOPE field (see 6.4.5.2). <sup>d</sup> LBA count = capacity × 2 <sup>(threshold exponent)</sup> .				

### F.3.5 Operations that occur

Write operations occur to LUN 1 that require one additional threshold set to be allocated when the application client writes 50 additional contiguous logical blocks. Used LBA mapping resources on LUN 1 are now 3 072 logical blocks (i.e., three threshold sets), and available LBA mapping resources are 69 107 712 logical blocks. Write operations also occur to LUN 2 that require no additional threshold sets when the application client writes an additional 100 logical blocks into a threshold set that was already allocated.



### F.3.6 Dedicated resource, threshold set tracked example final log page values

Table F.4 shows the values in the Logical Block Provisioning log page after the operations described in F.3.5 have occurred.

**Table F.4 — Dedicated resource, threshold set tracked example final log page values**

LUN	Log page parameter <sup>a</sup>	Resource count <sup>b</sup>	Scope <sup>c</sup>	Description <sup>d</sup>
1	0001h	0001_07A0h	01b	The available LBA mapping resource parameter indicates that 69 107 712 logical blocks (i.e., 1_07A0h threshold sets × 1 024 logical blocks per threshold set) are available for LUN 1.
	0002h	0000_0003h	01b	The used LBA mapping resource parameter indicates that 3 072 logical blocks (i.e., 3h threshold sets × 1 024 logical blocks per threshold set) have been used (i.e., allocated) by LUN 1.
2	0001h	0004_0000h	01b	The available LBA mapping resource parameter indicates that 1 073 741 824 (i.e., 4_0000h threshold sets × 4 096 logical blocks per threshold set) are available for LUN 2.
<sup>a</sup> THRESHOLD RESOURCE field (see 6.5.9.2) and the PARAMETER CODE field (see 6.4.5.2). <sup>b</sup> RESOURCE COUNT field (see 6.5.9.2). <sup>c</sup> SCOPE field (see 6.4.5.2). <sup>d</sup> LBA count = capacity × 2 <sup>(threshold exponent)</sup> .				

## F.4 Shared resource, logical block tracked example

### F.4.1 Shared resource, logical block tracked example overview

This example describes a method that tracks shared logical block provisioning resources based on logical blocks. The logical block provisioning resources are shared by multiple logical units. In this example, the values reported by each logical unit in its Logical Block Provisioning log page (see 6.4.5) reflect the combined usage of all logical units that share the logical block provisioning resources and the resources available for use by any of the logical units that share the logical block provisioning resources. Resources are allocated one logical block at a time but reported with a larger threshold set size.

#### F.4.2 Shared resource, logical block tracked example configuration

The configuration used for this example consists of two thin provisioned logical units, where the logical block provisioning resources are shared between both logical units. Table F.5 shows logical block provisioning related capacity values used in this example.

**Table F.5 — Shared resource, logical block tracked example capacity information**

LUN	Capacity		THRESHOLD EXPONENT field <sup>c</sup>	Number of threshold sets <sup>d</sup>
	LBA <sup>a</sup>	Logical blocks <sup>b</sup>		
1	3FFF_FFFFh	1 Gi	0Bh (i.e., 1 MiB, 2 048 logical blocks)	0008_0000h (i.e., 512 Ki)
2	BFFF_FFFFh	3 Gi	0Bh (i.e., 1 MiB, 2 048 logical blocks)	0018_0000h (i.e., 1 536 Ki)
<sup>a</sup> RETURNED LOGICAL BLOCK ADDRESS field in READ CAPACITY parameter data (see 5.17.2 and 5.18.2). <sup>b</sup> The value returned in the RETURNED LOGICAL BLOCK ADDRESS field plus one. <sup>c</sup> In the Logical Block Provisioning VPD page (see 6.6.6). <sup>d</sup> Number of threshold sets = capacity ÷ 2 <sup>(threshold exponent)</sup> .				

#### F.4.3 Shared resource, logical block tracked example time line

The sequence of events for this example are:

- 1) initial conditions (see F.4.4);
- 2) operations that occur (see F.4.5); and
- 3) final values in the logical block provisioning log page (see F.4.6).

#### F.4.4 Shared resource, logical block tracked example initial conditions

Initially, LUN 1 and LUN 2 have used a combined total of 57 000 logical blocks. LUN1 and LUN 2 have 1 073 741 900 logical blocks available for use by either LUN 1 or LUN 2. Table F.6 shows the values in the Logical Block Provisioning log page for the initial conditions in this example.

**Table F.6 — Shared resource, logical block tracked example initial conditions**

LUN	Log page parameter <sup>a</sup>	Resource count <sup>b</sup>	Scope <sup>c</sup>	Description
1	0001h	0008_0000h	10b	The available LBA mapping resource parameter indicates that from 1 073 741 824 logical blocks (i.e. 8_0000h threshold sets × 2 048 logical block per threshold set) to 1 073 743 871 logical blocks are available for LUN 1 or LUN 2. <sup>d</sup>
	0002h	0000_001Ch	10b	The used LBA mapping resource parameter indicates that from 55 297 logical blocks to 57 344 logical blocks (i.e., 1Ch threshold sets × 2 048 logical blocks per threshold set) have been used (i.e., allocated) by LUN 1 and LUN 2. <sup>e</sup>
2	0001h	0008_0000h	10b	The available LBA mapping resource parameter indicates that 1 073 741 824 logical blocks (i.e., 8_0000h threshold sets × 2 048 logical blocks per threshold set) are available for LUN 1 or LUN 2. <sup>d</sup>
<sup>a</sup> THRESHOLD RESOURCE field (see 6.5.9.2) and the PARAMETER CODE field (see 6.4.5.2). <sup>b</sup> RESOURCE COUNT field (see 6.5.9.2). <sup>c</sup> SCOPE field (see 6.4.5.2). <sup>d</sup> Minimum available LBA count = resource count × 2 <sup>(threshold exponent)</sup> . <sup>e</sup> Maximum used LBA count = resource count × 2 <sup>(threshold exponent)</sup> .				

#### F.4.5 Operations that occur

Write operations occur to LUN 1 that require 2 000 additional logical blocks to be used and write operations occur to LUN 2 that require 3 000 additional logical blocks to be used. Used LBA mapping resources on LUN 1 and LUN 2 are now 62 000 logical blocks, and the combined LBA mapping resources available to both LUN 1 and LUN 2 are 1 073 736 900 logical blocks (i.e., 1 073 741 900 minus 5 000).

#### F.4.6 Shared resource, logical block tracked example final log page values

Table F.7 shows the values in the Logical Block Provisioning log page after the operations described in F.4.5 have occurred.

**Table F.7 — Shared resource, logical block tracked example final log page values**

LUN	Log page parameter <sup>a</sup>	Resource count <sup>b</sup>	Scope <sup>c</sup>	Description
1	0001h	0007_FFFDh	10b	The available LBA mapping resource parameter indicates that from 1 073 735 680 logical blocks (i.e. 7_FFFDh threshold sets × 2 048 logical blocks per threshold set) to 1 073 737 727 logical blocks are available for LUN 1 or LUN 2. <sup>d</sup>
	0002h	0000_001Fh	10b	The used LBA mapping resource parameter indicates that from 61 441 logical blocks to 63 488 logical blocks (i.e., 1Fh threshold sets × 2 048 logical blocks per threshold set) have been used (i.e., allocated) by LUN 1 and LUN 2. <sup>e</sup>
2	0001h	0007_FFFDh	10b	The available LBA mapping resource parameter indicates that from 1 073 735 680 logical blocks (i.e. 7_FFFDh threshold sets × 2 048 logical blocks per threshold set) to 1 073 737 727 logical blocks are available for LUN 1 or LUN 2. <sup>d</sup>
<sup>a</sup> THRESHOLD RESOURCE field (see 6.5.9.2) and the PARAMETER CODE field (see 6.4.5.2). <sup>b</sup> RESOURCE COUNT field (see 6.5.9.2). <sup>c</sup> SCOPE field (see 6.4.5.2). <sup>d</sup> Minimum available LBA count = resource count × 2 <sup>(threshold exponent)</sup> . <sup>e</sup> Maximum used LBA count = resource count × 2 <sup>(threshold exponent)</sup> .				

## F.5 Shared available, dedicated used, logical block tracked example

### F.5.1 Shared available, dedicated used, logical block tracked example overview

This example describes a method that tracks available shared logical block provisioning resources based on logical blocks and dedicated used logical block provisioning resources based on logical blocks. The available logical block provisioning resources are shared by multiple logical units. In this example, the values reported by the logical unit in the available LBA mapping resource parameter of the Logical Block Provisioning log page (see 6.4.5) reflect the resources available for use by any of the logical units that share the logical block provisioning resources. The values reported by the logical unit in the used LBA mapping resource parameter of the Logical Block Provisioning log page reflect the usage for the individual logical unit.

### F.5.2 Shared available, dedicated used, logical block tracked example configuration

The configuration used for this example consists of two thin provisioned logical units, where the available logical block provisioning resources are shared between both logical units and used logical block provisioning resources are reported independently for each logical unit. Table F.8 shows logical block provisioning related capacity values used in this example.

**Table F.8 — Shared available, dedicated used example capacity information**

LUN	Capacity		THRESHOLD EXPONENT field <sup>c</sup>	Number of threshold sets <sup>d</sup>
	LBA <sup>a</sup>	Logical blocks <sup>b</sup>		
1	3FFF_FFFFh	1 Gi	0Bh (i.e., 1 MiB, 2 048 logical blocks)	0008_0000h (i.e., 512 Ki)
2	BFFF_FFFFh	3 Gi	0Bh (i.e., 1 MiB, 2 048 logical blocks)	0018_0000h (i.e., 1 536 Ki)
<sup>a</sup> RETURNED LOGICAL BLOCK ADDRESS field in READ CAPACITY parameter data (see 5.17.2 and 5.18.2). <sup>b</sup> The value returned in the RETURNED LOGICAL BLOCK ADDRESS field plus one. <sup>c</sup> In the Logical Block Provisioning VPD page (see 6.6.6). <sup>d</sup> Number of threshold sets = capacity ÷ 2 <sup>(threshold exponent)</sup> .				

### F.5.3 Shared available, dedicated used, logical block tracked example time line

The sequence of events for this example are:

- 1) initial conditions (see F.5.4);
- 2) operations that occur (see F.5.5); and
- 3) final values in the logical block provisioning log page (see F.5.6).

### F.5.4 Shared available, dedicated used, logical block tracked example initial conditions

Initially, LUN 1 has used 57 000 logical blocks and, LUN 2 has used 103 000 logical blocks. LUN 1 and LUN 2 have 1 073 741 900 logical blocks available for use by either LUN 1 or LUN 2. Table F.9 shows the values in the Logical Block Provisioning log page for the initial conditions in this example.

**Table F.9 — Shared resource, logical block tracked example initial conditions**

LUN	Log page parameter <sup>a</sup>	Resource count <sup>b</sup>	Scope <sup>c</sup>	Description
1	0001h	0008_0000h	10b	The available LBA mapping resource parameter indicates that from 1 073 741 824 logical blocks (i.e. 8_0000h threshold sets × 2 048 logical blocks per threshold set) to 1 073 743 871 logical blocks are available for LUN 1 or LUN 2. <sup>d</sup>
	0002h	0000_001Ch	01b	The used LBA mapping resource parameter indicates that from 55 297 logical blocks to 57 344 logical blocks (i.e., 1Ch threshold sets × 2 048 logical blocks per threshold set) have been used (i.e., allocated) by LUN 1. <sup>e</sup>
2	0001h	0008_0000h	10b	The available LBA mapping resource parameter indicates that from 1 073 741 824 logical blocks (i.e., 8_0000h threshold sets × 2 048 logical blocks per threshold set) to 1 073 743 871 logical blocks are available for LUN 1 or LUN 2. <sup>d</sup>
	0002h	0000_0033h	01b	The used LBA mapping resource parameter indicates that from 102 401 logical blocks to 104 448 (i.e., 33h threshold sets × 2 048 logical blocks per threshold set) have been used (i.e., allocated) by LUN 2. <sup>e</sup>
<sup>a</sup> THRESHOLD RESOURCE field (see 6.5.9.2) and the PARAMETER CODE field (see 6.4.5.2). <sup>b</sup> RESOURCE COUNT field (see 6.5.9.2). <sup>c</sup> SCOPE field (see 6.5.9.2). <sup>d</sup> Minimum available LBA count = resource count × 2 <sup>(threshold exponent)</sup> . <sup>e</sup> Maximum used LBA count = resource count × 2 <sup>(threshold exponent)</sup> .				

### F.5.5 Operations that occur

Write operations occur to LUN 1 that require 2 000 additional logical blocks to be used and write operations occur to LUN 2 that require 3 000 additional logical blocks to be used. Used LBA mapping resources on LUN 1 are now 59 000 logical blocks, used LBA mapping resources on LUN 2 are now 106 000 logical blocks, and the combined LBA mapping resources available to both LUN 1 and LUN 2 are 1 073 736 900 logical blocks.

### F.5.6 Shared available, dedicated used, example final log page values

Table F.10 shows the values in the Logical Block Provisioning log page after the operations described in F.5.5 have occurred.

**Table F.10 — Shared available, dedicated used example final log page values**

LUN	Log page parameter <sup>a</sup>	Resource count <sup>b</sup>	Scope <sup>c</sup>	Description
1	0001h	0007_FFFDh	10b	The available LBA mapping resource parameter indicates that from 1 073 735 680 logical blocks (i.e. 7_FFFDh threshold sets × 2 048 logical blocks per threshold set) to 1 073 737 727 logical blocks are available for LUN 1 or LUN 2. <sup>d</sup>
	0002h	0000_001Dh	01b	The used LBA mapping resource parameter indicates that from 57 345 logical blocks to 59 392 logical blocks (i.e., 1Dh threshold sets × 2 048 logical blocks per threshold set) have been used (i.e., allocated) by LUN 1. <sup>e</sup>
2	0001h	0007_FFFDh	10b	The available LBA mapping resource parameter indicates that from 1 073 735 680 logical blocks (i.e. 7_FFFDh threshold sets × 2 048 logical blocks per threshold set) to 1 073 737 727 logical blocks are available for LUN 1 or LUN 2. <sup>d</sup>
	0002h	0000_0034h	01b	The used LBA mapping resource parameter indicates that from 104 449 logical blocks to 106 496 logical blocks (i.e., 34h threshold sets × 2_048 logical blocks per threshold set) have been used (i.e., allocated) by LUN 2. <sup>e</sup>
<sup>a</sup> THRESHOLD RESOURCE field (see 6.5.9.2) and the PARAMETER CODE field (see 6.4.5.2). <sup>b</sup> RESOURCE COUNT field (see 6.5.9.2). <sup>c</sup> SCOPE field (see 6.4.5.2). <sup>d</sup> Minimum available LBA count = resource count × 2 <sup>(threshold exponent)</sup> . <sup>e</sup> Maximum used LBA count = resource count × 2 <sup>(threshold exponent)</sup> .				

## Annex G

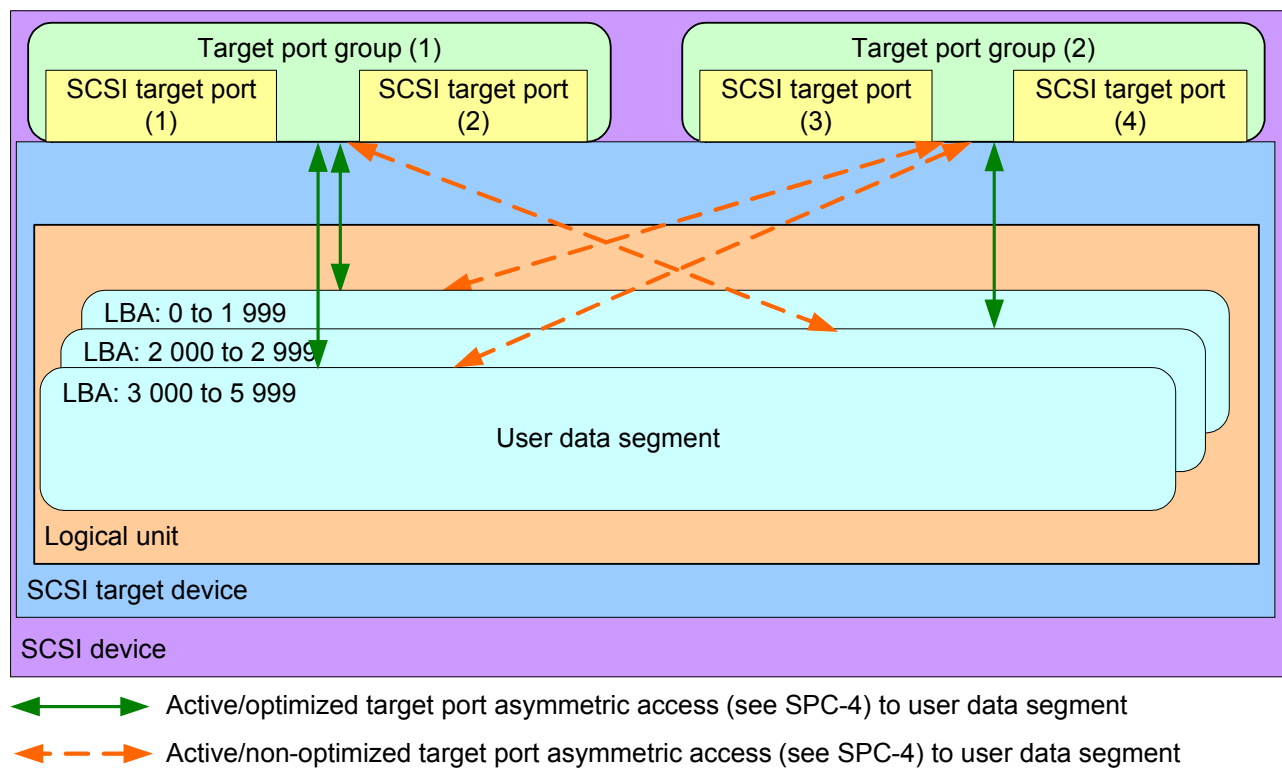
(informative)

### Discovering referrals examples

#### G.1 Referrals example with no user data segment multiplier

This subclause demonstrates a method an application client may use to determine the optimal target port group from which to access logical blocks using information sent from the device server when the user data segment multiplier is set to zero.

Figure G.1 shows an example of a SCSI device in which referrals have been implemented with a user data segment multiplier of zero.



**Figure G.1 — Referrals example with no user data segment multiplier**



In the example shown in figure G.1, the application client acquires the information from the logical unit as shown in table G.1.

**Table G.1 — Referrals application client information with no user data segment multiplier**

Referrals VPD page			
User data segment size		User data segment multiplier	
ignored		0	
REPORT TARGET PORT GROUPS command			
Asymmetric access state	Target port group	Relative target port identifier	
4h (i.e., logical block dependent)	1	1	
		2	
	2	3	
		4	
REPORT REFERRALS command or user data segment referral sense data descriptors			
First user data segment LBA	Last user data segment LBA	Asymmetric access state	Target port group
0	1 999	0 (i.e., active/optimized)	1
		1 (i.e., active/non-optimized)	2
2 000	2 999	0 (i.e., active/optimized)	2
		1 (i.e., active/non-optimized)	1
3 000	5 999	0 (i.e., active/optimized)	1
		1 (i.e., active/non-optimized)	2

The application may determine the user data segments that are optimally accessed through the two target port groups as shown in table G.2.

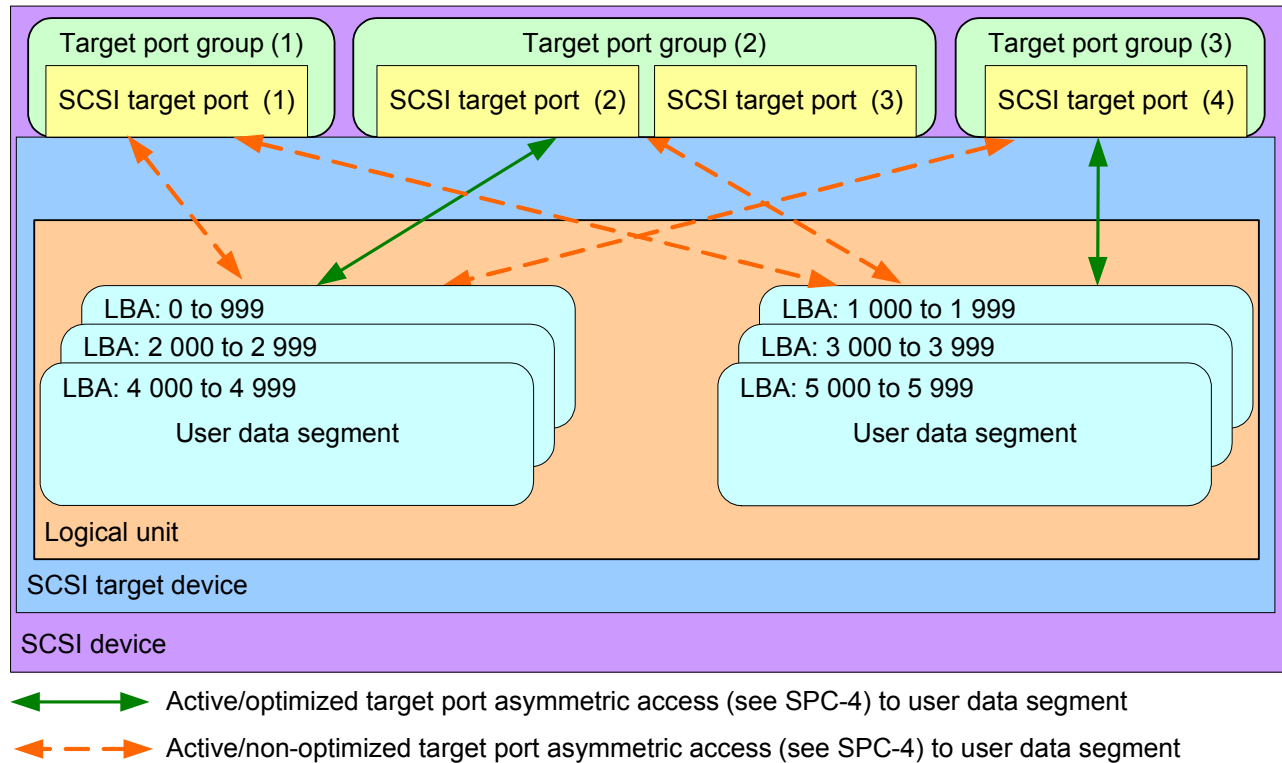
**Table G.2 — User data segment calculations with no user data segment multiplier**

First LBA of user data segment	Calculation (see 4.28.2)	Last LBA of user data segment	Calculation (see 4.28.2)
Target port group 1 user data segments in active/optimized asymmetric access state			
0	0 <sup>a</sup>	1 999	1 999 <sup>b</sup>
3 000	3 000 <sup>a</sup>	5 999	5 999 <sup>b</sup>
Target port group 2 user data segments in active/optimized asymmetric access state			
2 000	2 000 <sup>a</sup>	2 999	2 999 <sup>b</sup>
<sup>a</sup> The first user data segment LBA <sup>b</sup> The last user data segment LBA			

## G.2 Referrals example with non-zero user data segment multiplier

This subclause demonstrates a method that an application client may use to determine the optimal target port group from which to access logical blocks using information sent from the device server when the user data segment multiplier is set to a non-zero value.

Figure G.2 shows an example of a SCSI device in which referrals have been implemented with a user data segment multiplier of two and a user data segment size of 1 000.



**Figure G.2 — Referrals example with non-zero user data segment multiplier**

In the example shown in figure G.2, the application client acquires the information from the logical unit as shown in table G.3.

**Table G.3 — Referrals application client information with non-zero user data segment multiplier**

Referrals VPD page			
User data segment size		User data segment multiplier	
1 000		2	
REPORT TARGET PORT GROUPS command			
Asymmetric access state	Target port group	Relative target port identifier	
4h (i.e., logical block dependent)	1	1	
	2	2	
		3	
	3	4	
REPORT REFERRALS command or user data segment referral sense data descriptors			
First user data segment LBA	Last user data segment LBA	Asymmetric access state	Target port group
0	4 999	0 (i.e., active/optimized)	2
		1 (i.e., active/non-optimized)	1
		1 (i.e., active/non-optimized)	3
1 000	5 999	0 (i.e., active/optimized)	3
		1 (i.e., active/non-optimized)	1
		1 (i.e., active/non-optimized)	2

The application may determine the user data segments that are optimally accessed through the two target port groups as shown in table G.4.

**Table G.4 — User data segment calculations with non-zero user data segment multiplier**

First LBA of user data segment	Calculation (see 4.28.2)	Last LBA of user data segment	Calculation (see 4.28.2)
Target port group 2 user data segments in active/optimized asymmetric access state			
0	0 <sup>a</sup>	999	0 <sup>a</sup> + (1 000 – 1)
2 000	0 + (1 000 × 2)	2 999	2 000 + (1 000 – 1)
4 000	2 000 + (1 000 × 2)	4 999	4 000 + (1 000 – 1)
Target port group 3 user data segments in active/optimized asymmetric access state			
1 000	1 000 <sup>a</sup>	1 999	1 000 <sup>a</sup> + (1 000 – 1)
3 000	1 000 + (1 000 × 2)	3 999	3 000 + (1 000 – 1)
5 000	3 000 + (1 000 × 2)	5 999	5 000 + (1 000 – 1)
<sup>a</sup> The first user data segment LBA.			

## **Annex H**

(informative)

### **IO Advice Hints Usage**

#### **H.1 IO Advice Hints Overview**

This annex describes how application clients and device servers may use IO advice hints and their associated LBM descriptors to provide information to a storage device to enhance access to data.

#### **H.2 IO Advice Hints Grouping mode page**

The IO Advice Hints Grouping mode page (see 6.5.7) contains the definitions of the IO advice hints (see 4.23.2) that are associated with each group number.

A device server may implement pre-defined IO advice hints that are not changeable and indicate this by making one or more IO advice hints descriptors in the mode page not changeable.

The application client may use the MODE SENSE command to retrieve IO advice hints descriptors to determine the IO advice hints that are available, and determine which IO advice hints are associated with each group number.

If any IO advice hints descriptors in the mode page are changeable, the application client may use a MODE SELECT command to associate IO group numbers with advice hints that meet the application client needs.

A device server may implement any combination of changeable IO advice hints descriptors or predefined IO advice hints descriptors.

#### **H.3 Issuing I/O commands**

##### **H.3.1 Group numbers and I/O commands**

When the application client determines the IO advice hints descriptor appropriate for each particular type of command as described in H.2, the application client then determines the value to specify in the GROUP NUMBER field of each command as described in 4.23.2. The resulting group number is then placed into the GROUP NUMBER field of the CDB.

If cache segmentation (see 4.15.2) is enabled (see 6.5.7), then the use of different group numbers by components within the application client (e.g., file system, virtual memory swapping/paging, payroll application) allows the device server to associate an IO and the associated cached data with other IOs and other cached data from that same component.

##### **H.3.2 Possible constraints on IO advice hints**

There is no method to report how an IO advice hint has been processed by the device server. The implications of this model include the ability of the device server to ignore any, or even all, IO advice hints it receives. Since a device server that always ignores all IO advice hints begs the question of why implement the IO Advice Hints Grouping mode page (see 6.5.7) and then ignore everything it specifies, other constraints make more sense.

A more practical approach for a device server to constrain the complexity of its IO advice hints implementation is to ignore one or more of the values in the logical block markup descriptor (see 6.8) (e.g., to ignore the READ SEQUENTIALITY field, and the READ/WRITE FREQUENCY field).

If the complexity of specifying values or locating values in the IO Advice Hints Grouping mode page (see 6.5.7) is a concern, the device server may implement the page, but indicate that one or more IO advice hints groups descriptors (see 6.5.7) are not changeable (see SPC-5). A device server that does this should:

- a) define the unchangeable IO advice hints for groups descriptors in the lowest group numbers supported for usage by IO advice hints; and
- b) allow at least four IO advice hints group descriptors to be changeable.

## H.4 Logical Block Markup descriptor usage examples

### H.4.1 Example usage in tiered storage device implementations

Products (e.g., SCSI target devices and ATA devices) may use the contents of an LBM descriptor in ways that are not related to the contents of one LBM descriptor field in any way that is traceable to the way the field is defined. As an example, suppose a product contains the following storage tiers:

- a) volatile cache;
- b) non-volatile cache;
- c) solid state device storage; and
- d) magnetic media storage.

Table H.1 shows examples of how such a tiered product may interpret combinations of values in fields in an access patterns LBM descriptor (see 6.8.3.1).

**Table H.1 — Tiered product access patterns LBM descriptor examples**

Field	Code	Field	Code	Field	Code
Data that should not be written to the solid state device storage					
READ/WRITE FREQUENCY	10b	WRITE SEQUENTIALITY	see <sup>a</sup>	OSI PROXIMITY	see <sup>a</sup>
OVERALL FREQUENCY	10b	READ SEQUENTIALITY	see <sup>a</sup>	ACDLU	see <sup>a</sup>
SUBSEQUENT I/O	see <sup>a</sup>	<a href="#">IO CLASS</a>	<a href="#">see <sup>a</sup></a>		
Data that should be written to non-volatile media and not retained in cache					
READ/WRITE FREQUENCY	10b	WRITE SEQUENTIALITY	see <sup>a</sup>	OSI PROXIMITY	see <sup>a</sup>
OVERALL FREQUENCY	01b	READ SEQUENTIALITY	see <sup>a</sup>	ACDLU	see <sup>a</sup>
SUBSEQUENT I/O	see <sup>a</sup>	<a href="#">IO CLASS</a>	<a href="#">see <sup>a</sup></a>		
<sup>a</sup> The contents of this field do not contribute to this example case.					

### H.4.2 Example LBM descriptor values for software that sends read commands and write commands

Table H.2 shows example mappings to access patterns LBM descriptors (see 6.8.3.1) from:

- a) file types in common usage; and
- b) published file system and operating system specifications for usage characteristics.

Table H.2 — Sending device access patterns LBM descriptor examples (part 1 of 2)

Field	Code	Field	Code	Field	Code
<b>LBMd value:</b> 80h, A5h, <del>05</del> 2h, 00h <b>File type:</b> swap/page file		Applicable published specifications Microsoft®: FILE_TYPE_NOTIFICATION_GUID_PAGE_FILE Linux: None known			
READ/WRITE FREQUENCY	10b	WRITE SEQUENTIALITY	01b	OSI PROXIMITY	10b
OVERALL FREQUENCY	10b	READ SEQUENTIALITY	01b	ACDLU	1b
SUBSEQUENT I/O	00b	<a href="#">IO CLASS</a>	<a href="#">5h</a>		
<b>LBMd value:</b> 00h, 00h, <del>02</del> 56h, 00h <b>File type:</b> hibernate file		Applicable published specifications Microsoft®: FILE_TYPE_NOTIFICATION_GUID_HIBERNATION_FILE Linux: None known			
READ/WRITE FREQUENCY	00b	WRITE SEQUENTIALITY	00b	OSI PROXIMITY	10b
OVERALL FREQUENCY	00b	READ SEQUENTIALITY	00b	ACDLU	0b
SUBSEQUENT I/O	01b	<a href="#">IO CLASS</a>	<a href="#">5h</a>		
<b>LBMd value:</b> 00h, <del>AA</del> 55h, <del>20</del> 06h, 00h <b>File type:</b> system initialization (e.g., registry)		Applicable published specifications Microsoft®: None known Linux: None known			
READ/WRITE FREQUENCY	01b	WRITE SEQUENTIALITY	01b	OSI PROXIMITY	10b
OVERALL FREQUENCY	01b	READ SEQUENTIALITY	01b	ACDLU	0b
SUBSEQUENT I/O	01b	<a href="#">IO CLASS</a>	<a href="#">0h</a>		
<b>LBMd value:</b> 00h, 05h, <del>01</del> 4h, 00h <b>File type:</b> file system metadata (e.g., directory files)		Applicable published specifications Microsoft®: None known Linux: None known			
READ/WRITE FREQUENCY	00b	WRITE SEQUENTIALITY	01b	OSI PROXIMITY	00b
Key: LBMd = LBM descriptor					
Note 1 - Microsoft information obtained from <a href="http://msdn.microsoft.com/en-us/library/windows/desktop/hh404249%28v=vs.85%29.aspx">http://msdn.microsoft.com/en-us/library/windows/desktop/hh404249%28v=vs.85%29.aspx</a> . Note 2 - Linux information obtained from <a href="http://linux.die.net/man/2/fadvise">http://linux.die.net/man/2/fadvise</a> .					
a <a href="#">Where x represents the io_class value.</a> b <a href="#">The application client should select the io_class value (e.g., 4h or 5h) appropriate to the size of the object.</a>					

Table H.2 — Sending device access patterns LBM descriptor examples (part 2 of 2)

Field	Code	Field	Code	Field	Code
OVERALL FREQUENCY	00b	READ SEQUENTIALITY	01b	ACDLU	0b
SUBSEQUENT I/O	01b	<a href="#">IO CLASS</a>	<a href="#">1h</a>		
<b>LBMd value:</b> 00h, 05h, <a href="#">0x0h_a</a> , 00h <b>File type:</b> random access		Applicable published specifications Microsoft®: None known Linux: FADV_RANDOM			
READ/WRITE FREQUENCY	00b	WRITE SEQUENTIALITY	01b	OSI PROXIMITY	00b
OVERALL FREQUENCY	00b	READ SEQUENTIALITY	01b	ACDLU	0b
SUBSEQUENT I/O	00b	<a href="#">IO CLASS</a>	<a href="#">see_b</a>		
<b>LBMd value:</b> 00h, 0Ah, <a href="#">0x0h_a</a> , 00h <b>File type:</b> sequential access		Applicable published specifications Microsoft®: None known Linux: FADV_SEQUENTIAL			
READ/WRITE FREQUENCY	00b	WRITE SEQUENTIALITY	10b	OSI PROXIMITY	00b
OVERALL FREQUENCY	00b	READ SEQUENTIALITY	10b	ACDLU	0b
SUBSEQUENT I/O	00b	<a href="#">IO CLASS</a>	<a href="#">see_b</a>		
Key: LBMd = LBM descriptor					
Note 1 - Microsoft information obtained from <a href="http://msdn.microsoft.com/en-us/library/windows/desktop/hh404249%28v=vs.85%29.aspx">http://msdn.microsoft.com/en-us/library/windows/desktop/hh404249%28v=vs.85%29.aspx</a> . Note 2 - Linux information obtained from <a href="http://linux.die.net/man/2/fadvise">http://linux.die.net/man/2/fadvise</a> .					
<sup>a</sup> <a href="#">Where x represents the io_class value.</a> <sup>b</sup> <a href="#">The application client should select the io_class value (e.g., 4h or 5h) appropriate to the size of the object.</a>					

## Annex I

(informative)

### Bibliography

ISO/IEC 14776-321, *SCSI Block Commands (SBC)*

ISO/IEC 14776-322, *SCSI Block Commands-2 (SBC-2)*

ISO/IEC 14776-154, *Serial Attached SCSI-3 (SAS-3)*

~~ISO/IEC 14776-223~~ [ANSI INCITS 481-2012](#), *Fibre Channel Protocol for SCSI-3 (FCP-3)*

ANSI INCITS 481-2012, *Fibre Channel Protocol for SCSI - 4 (FCP-4)*

CFast (CFast)<sup>™</sup>, CompactFlash Association (see <http://www.compactflash.org>)

CompactFlash<sup>®</sup> (CF), CompactFlash Association (see <http://www.compactflash.org>)

Memory Stick<sup>™</sup> (MS). One Stop Site for Formats (see <https://www.oss-formats.org>)

MultiMediaCard (e•MMC), JEDEC<sup>®</sup> (see <http://www.jedec.org>)

NOTE 27 JEDEC<sup>®</sup> is a registered trademark of JEDEC Solid State Technology Association. This information is given for the convenience of users of this standard and does not constitute an endorsement by ISO, IEC, or ANSI.

Secure Digital Card (SD Card), SD Association (see <http://www.sdcard.org>)

XQD<sup>™</sup> (XQD), CompactFlash Association (see <http://www.compactflash.org>)

Universal Flash Storage (UFS), JEDEC<sup>®</sup> (see <http://www.jedec.org>)

[Differentiated Storage Services, ACM Special Interest Group on Operating Systems \(see http://sigops.org/sosp/sosp11/current/2011-Cascais/printable/05-mesnier.pdf\)](http://sigops.org/sosp/sosp11/current/2011-Cascais/printable/05-mesnier.pdf)